

Lecture 9

Ordered and Sorted Ranges

Algorithms and D.S. to Represent Sets



EECS 281: Data Structures & Algorithms

Ordered and Sorted Containers

Data Structures & Algorithms

Container Review

- Objects storing a variable number of other objects
- Allow for control/protection of data
- Can copy/edit/sort/move many objects at once
- Examples: vector, deque, stack, map, list, array

Unordered Container



Ordered Container



Nested Container?



Types of Containers

Type	Distinctive interfaces (not all methods listed)
Container	Supports add() and remove() operations
Searchable Container	Adds find() operation
Sequential Container	Allows iteration over elements in some order
Ordered Container	Sequential container which maintains current order. Can arbitrarily insert new elements anywhere. Example: Books on a shelf
Sorted Container	Sequential container with pre-defined order. Can NOT arbitrarily insert elements. Example: Students sorted by ID

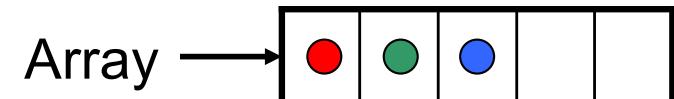
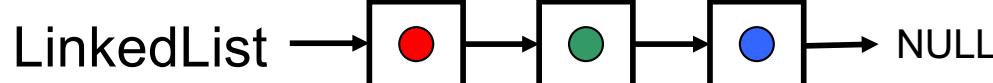
When would sorted containers be preferred over ordered?

Ordered Clarification

- Ordered container elements maintain their “relative” position unless they are removed
- Example: If A comes before Q, and then Z is inserted between them, their relative ordering has not changed
 - A still comes before Q
- If Q is then removed, the container is still ordered
 - Before and after delete, A comes before Z

Implementing Sorted and Ordered Containers

- Two implementation styles: connected, contiguous



- Preferred implementation dependent upon requirements of application
 - Know which operations will be called often
- Study multiple implementations
 - Know asymptotic complexity of each operation

When would a linked list be preferred over an array?

Asymptotic Complexities: Ordered Container

Operation	Array	Linked List
addElement(val)	$O(1)$	$O(1)$
remove(val)	$O(n)$	$O(n)$
remove(iterator)	$O(n)$	$O(n)$ or $O(1)$
find(val)	$O(n)$	$O(n)$
Iterator::operator*()	$O(1)$	$O(1)$
operator[](unsigned)	$O(1)$	$O(n)$
insertAfter(iterator, val)	$O(n)$	$O(1)$
insertBefore(iterator, val)	$O(n)$	$O(n)$ or $O(1)$

Asymptotic Complexities: Sorted Container

Operation	Array	Linked List
addElement(val)	$O(n)$	$O(n)$
remove(val)	$O(n)$	$O(n)$
remove(iterator)	$O(n)$	$O(n)$ or $O(1)$
find(val)	$O(\log n)$	$O(n)$
Iterator::operator*()	$O(1)$	$O(1)$
operator[](unsigned)	$O(1)$	$O(n)$
insertAfter(iterator, val)	N/A	N/A
insertBefore(iterator, val)	N/A	N/A

Ordered and Sorted Containers

Data Structures & Algorithms

Binary Search and Bounds

Data Structures & Algorithms

Binary Search Example: Find 21

21 > 13: Look to the right



2	3	5	7	13	21	25	31	42
---	---	---	---	----	----	----	----	----

21 < 25: Look to the left



21	25	31	42
----	----	----	----

21 is found



21

Binary search requires elements to be sorted

Binary Search

```
1 int bsearch(double a[], double val,
2             int left, int right) {
3     while (right > left) {           ← n is size of a[]
4         int mid = left + (right - left) / 2;   ← n = right - left
5         if (val == a[mid])               ← loop at most k times
6             return mid;                ← 1 step
7
8         if (val < a[mid])             ← 1 step
9             right = mid;            ← 1 step
10        else
11            left = mid + 1;          ← 1 step
12    } // while
13
14    return -1; // val not found
15 } // bsearch()
```

n is size of $a[]$
 $n = right - left$
loop at most k times
1 step
1 step
1 step
1 step
1 step
 n is split in half each loop
 $n = n / 2$
 $2^k = n$
Total: 5 k steps = $O(k)$
But what is k ? $\log(n)$

Asymptotic Complexity = $O(\log n)$

How do we compare elements that are objects?

Comparators (Function Objects)

Given elements a and b, tell if a "<" b

```
1  struct Point {  
2      int x, y;  
3      Point() : x(0), y(0) { }  
4      Point(int xx, int yy) : x(xx), y(yy) { };  
5  }; // Point{}  
  
6  struct CompareByX {  
7      bool operator()(const Point &a, const Point &b) const {  
8          return a.x > b.x;  
9      } // operator()()  
10 }; // CompareByX{}  
  
11 struct CompareByY {  
12     bool operator()(const Point &a, const Point &b) const {  
13         return a.y < b.y;  
14     } // operator()()  
15 }; // CompareByY{}
```

Speeding up Binary Search

- **Speed-up idea:** == rarely triggers,
 - check if < first
 - else if >
 - else must be ==
- More radical idea: move the == check out of the loop
 - Find a sharp *lower bound* for the sought element first
 - *First item >= what I'm looking for*
 - Check for the value == after the loop

Almost Always 3 Comparisons/Loop

```
1 int bsearch(double a[], double val,
2             int left, int right) {
3     while (right > left) { // ONE
4         int mid = left + (right - left) / 2;
5
6         if (val == a[mid]) // TWO
7             return mid;
8         else if (val < a[mid]) // THREE
9             right = mid;
10        else
11            left = mid + 1;
12    } // while
13
14    return -1; // val not found
15 } // bsearch()
```

2 Comparisons Half the Time

```
1 int bsearch(double a[], double val,
2             int left, int right) {
3     while (right > left) { // ONE
4         int mid = left + (right - left) / 2;
5
6         if (a[mid] < val) // TWO: check < not ==
7             left = mid + 1;
8         else if (val < a[mid]) // THREE
9             right = mid;
10        else // (val == a[mid])
11            return mid;
12    } // while
13
14    return -1; // val not found
15 } // bsearch()
```

Always 2 Comparisons/Loop

```
1 int lower_bound(double a[], double val,
2                  int left, int right) {
3     while (right > left) {                                // ONE
4         int mid = left + (right - left) / 2;
5
6         if (a[mid] < val)                                  // TWO
7             left = mid + 1;
8         else // (a[mid] >= val)
9             right = mid;
10    } // while
11
12    return left;
13 } // lower_bound()
```

Must check if `val` is present in `a[]` – do this before returning or (as in STL) require the caller to do this.

Binary Search in STL

`binary_search()` returns a bool, not the location

To find locations (iterators), use

- `lower_bound()` First item not less than target
- `upper_bound()` First item greater than target
- `equal_range()` Pair(lb, ub), all items equal to target

References

- http://en.wikipedia.org/wiki/Binary_search_algorithm
- <http://community.topcoder.com/tc?module=Static&d1=tutorials&d2=binarySearch>

Search Bounds

- `lower_bound(begin(v), end(v), 7)`

2	3	5	7	13	21	25	31	42
---	---	---	---	----	----	----	----	----

- `lower_bound(begin(v), end(v), 26)`

2	3	5	7	13	21	25	31	42
---	---	---	---	----	----	----	----	----

- `upper_bound(begin(v), end(v), 21)`

2	3	5	7	13	21	25	31	42
---	---	---	---	----	----	----	----	----

- `upper_bound(begin(v), end(v), 4)`

2	3	5	7	13	21	25	31	42
---	---	---	---	----	----	----	----	----

Binary Search and Bounds

Data Structures & Algorithms

Representing Sets

Data Structures & Algorithms

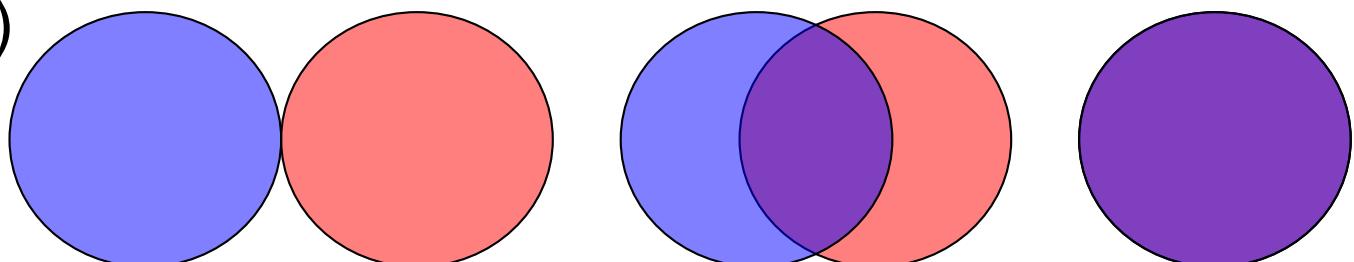
Searchable Containers as Sets

A set is well-defined if you can tell
if any given element is in the set

(Searchable containers well suited to finding elements for sets)

Set Operations (STL implements many of these)

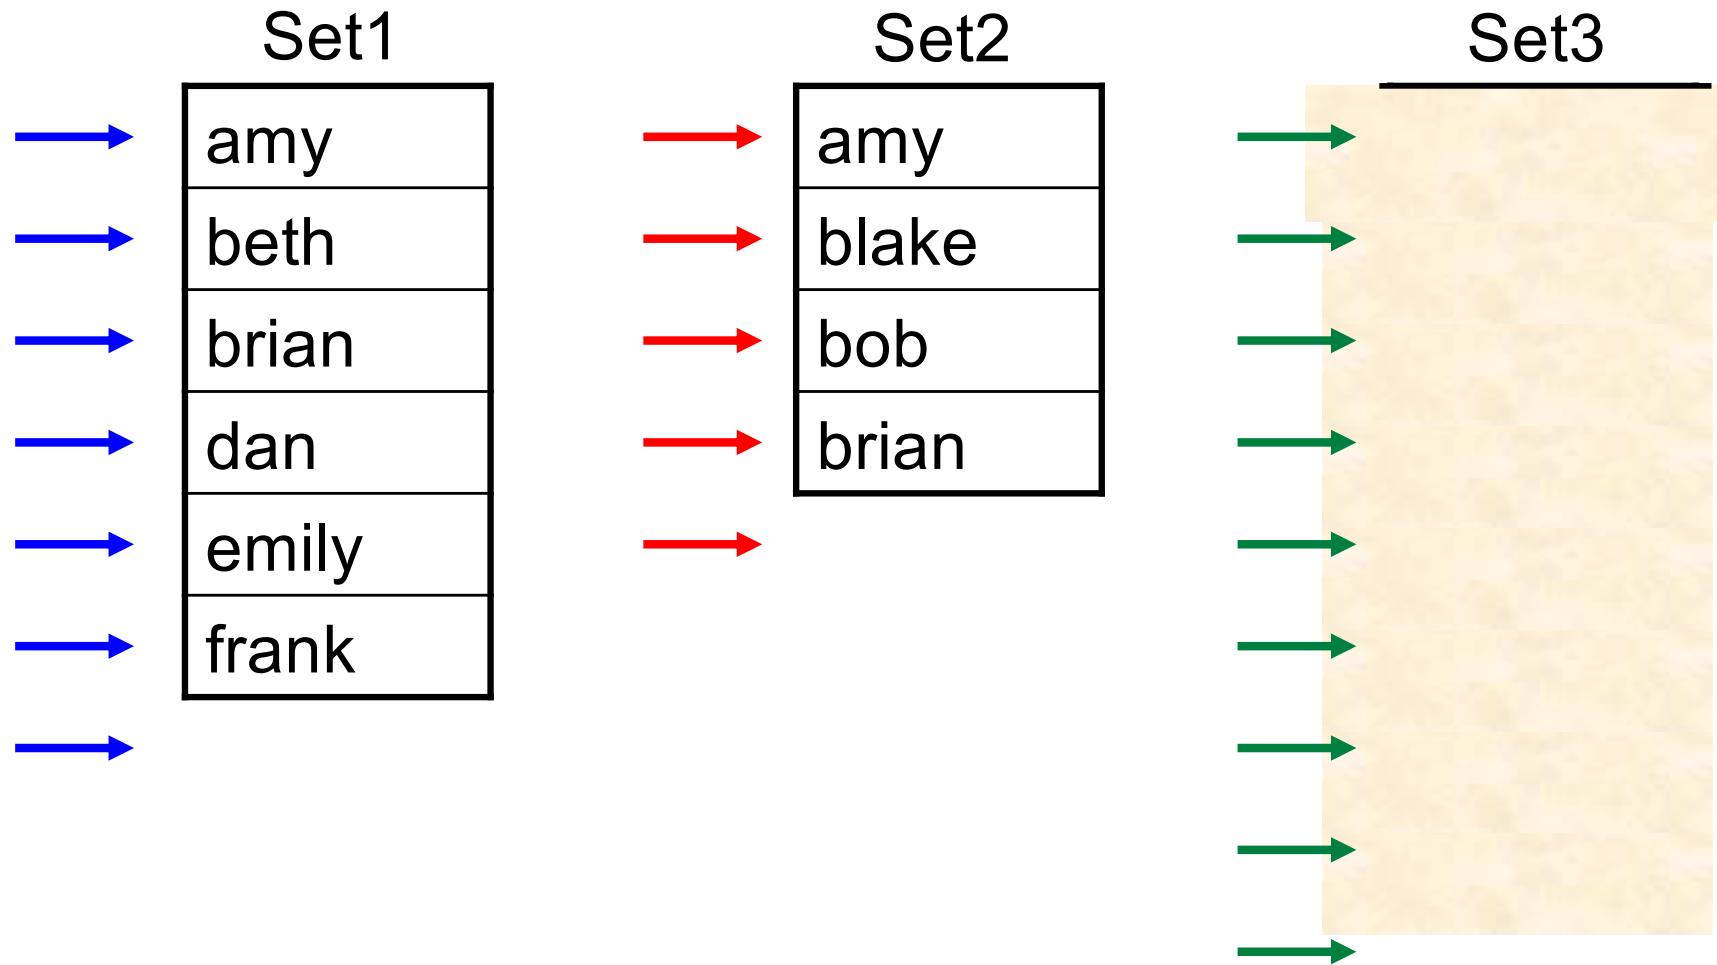
- Union (\cup): in one set or the other (OR)
- Intersection (\cap): in both sets (AND)
- Set Difference (\setminus): in one and not the other (AND-NOT)
- Symmetric Difference (\div): in only one (XOR)
- addElement (+)
- isElement (\in)
- isEmpty



set_union() Example

Set1 and Set2 are sorted ranges

Set3 is a union of Set1 and Set2



set_union() Example Code

```
1 template<class ForIterator1, class ForIterator2, class OutIterator,
2           class Compare>
3 OutIterator set_union(ForIterator1 first1, ForIterator1 last1,
4                       ForIterator2 first2, ForIterator2 last2,
5                       OutIterator result, Compare compare) {
6     while (first1 != last1 && first2 != last2) {
7         if (compare(*first1, *first2))
8             *result++ = *first1++; // set1 element less than set2 element
9         else if (compare(*first2, *first1))
10            *result++ = *first2++; // set2 element less than set1 element
11        else {
12            *result++ = *first1++; // set1 element == set2 element
13            ++first2;
14        } // else
15    } // while
16    while (first1 != last1)
17        *result++ = *first1++; // Remaining elements
18    while (first2 != last2)
19        *result++ = *first2++; // Remaining elements
20    return result; // returns end of sorted union of set1 and set2
21 } // set_union()
```

Implementing [sub]sets with ranges

Method	Asymptotic Complexity
initialize()	$O(1)$ or $O(n \log n)$
clear()	$O(1)$ or $O(n)$
isMember()	$O(\log n)$
copy()	$O(n)$
set_union()	$O(n)$
set_intersection()	$O(n)$

Universe: set of all elements that may be in a set

Job Interview Problems

- Given a sorted array with n elements and a number z
- Do the following in $O(n)$ time
 - Find pairs (x, y) such that $x - y = z$
 - Find pairs (x, y) such that $|x - y|$ is closest to z
 - Count all pairs (x, y) such that $x + y < z$
- What if the array was not sorted ?

Representing Sets

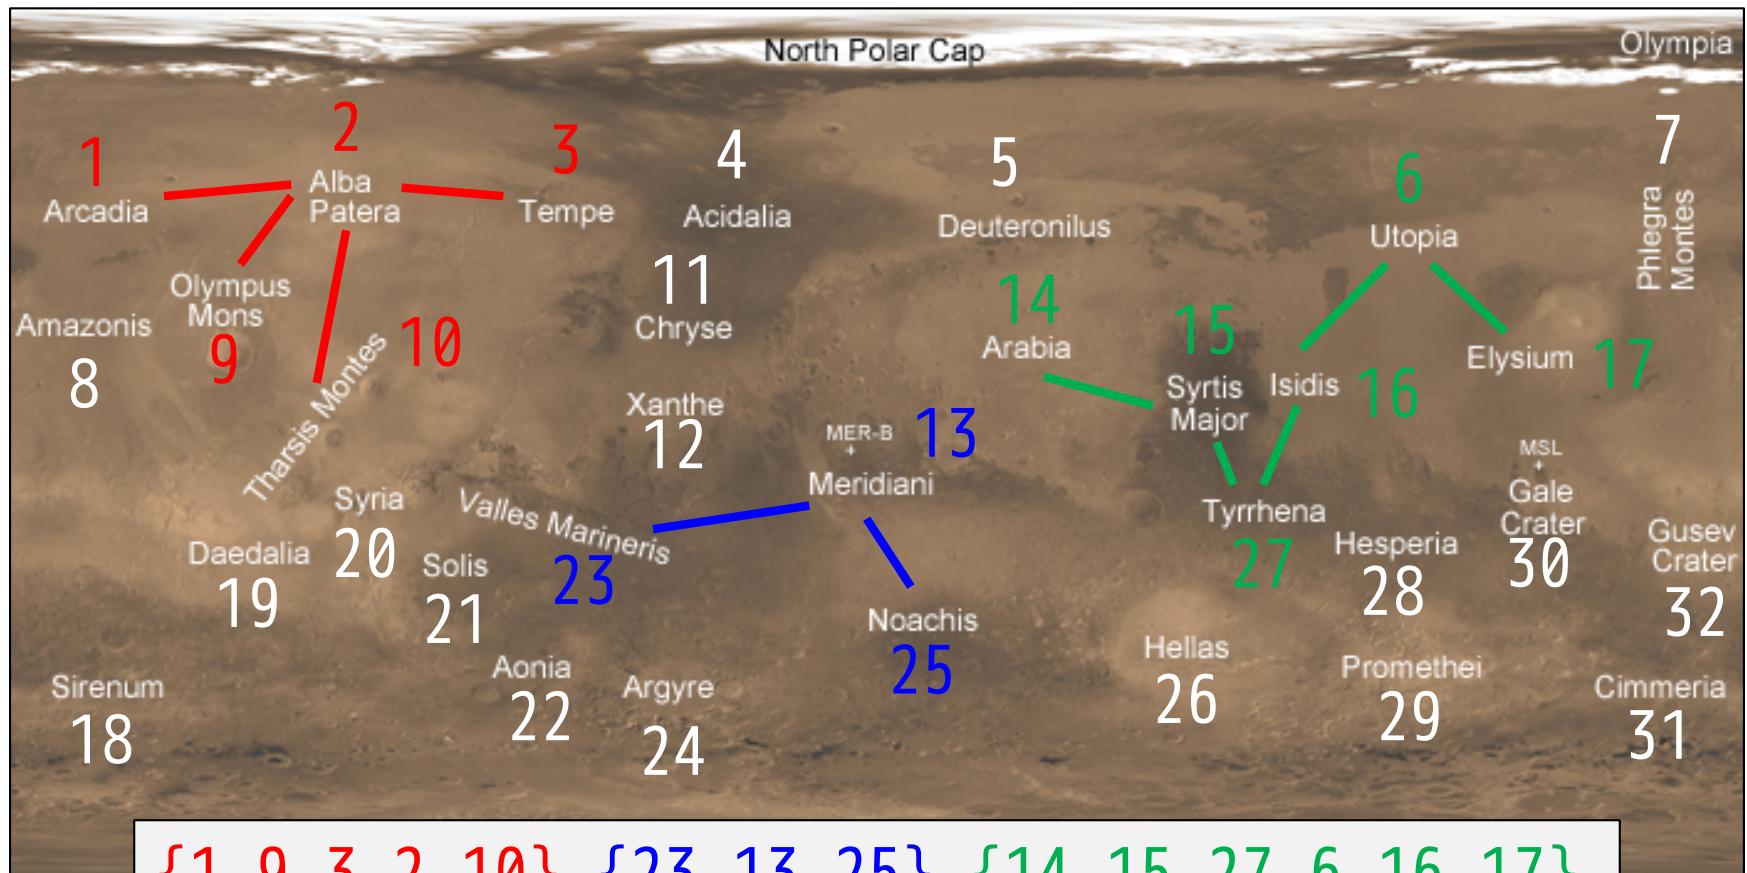
Data Structures & Algorithms

Union-Find

Data Structures & Algorithms

Representing Disjoint Sets

- Sets are **disjoint** if they do not share any elements.
(i.e. an element may only belong to one set)
- Many applications require representing and operating on disjoint sets.
For example, graph connectivity:

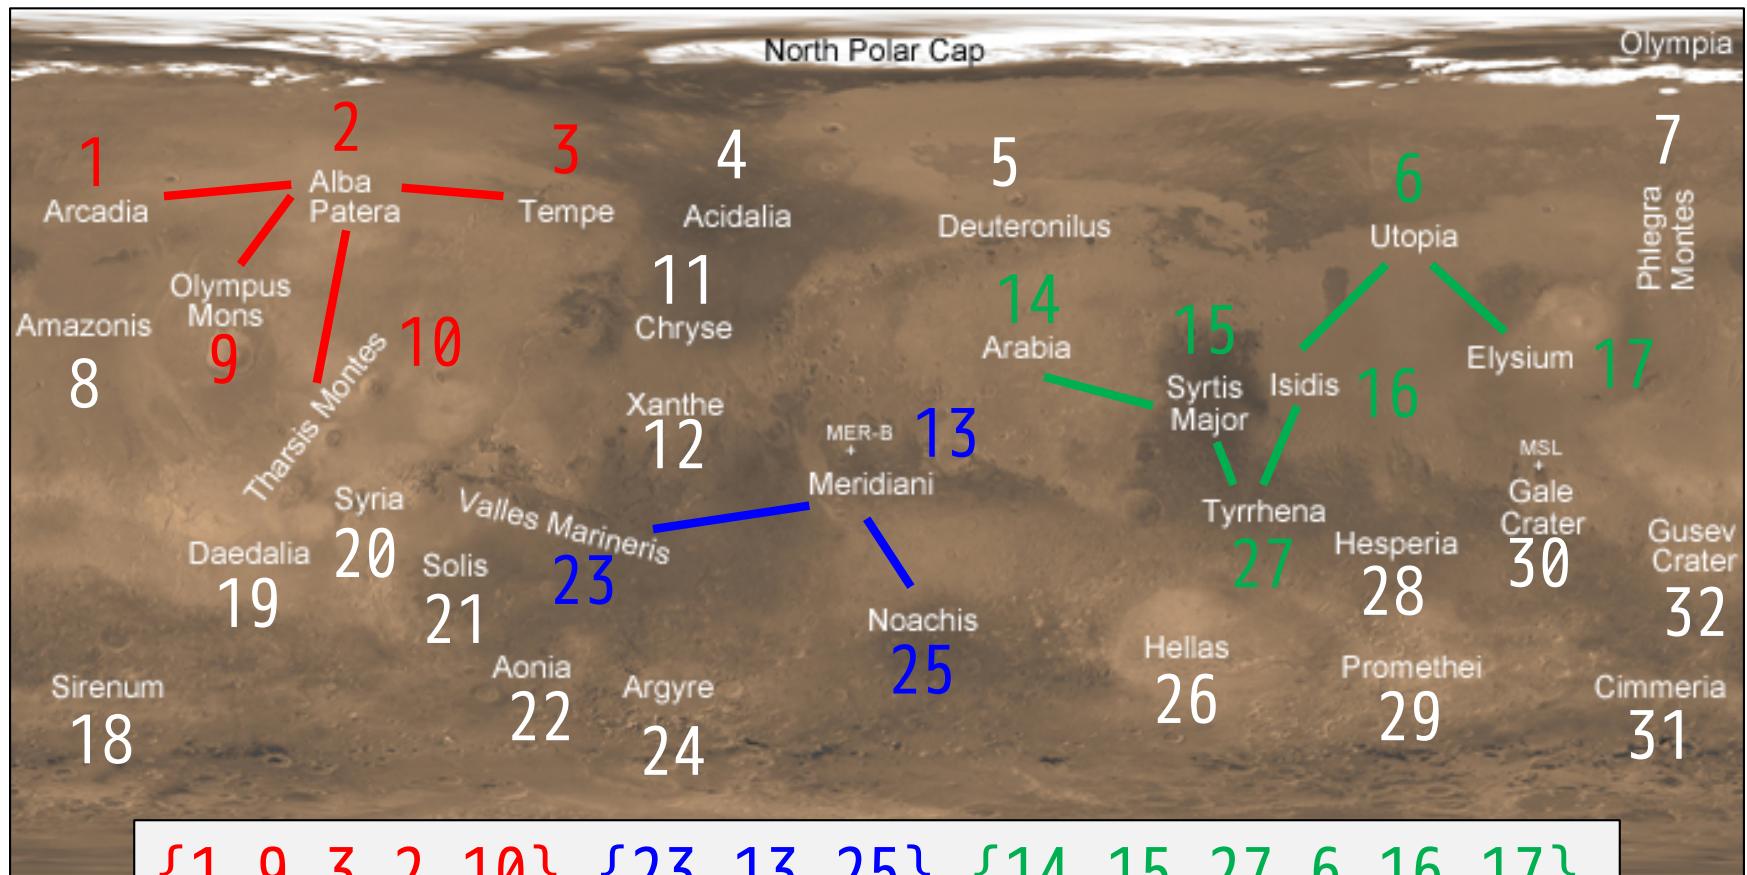


{1,9,3,2,10} {23,13,25} {14,15,27,6,16,17}

Also, single element sets: {8} {4} {5} {11} ...

Representing Disjoint Sets

- In this context, only two set operations make sense:
 - Find:** check if an element belongs to a particular set (or if two belong to the same set)
 - Union:** merge two sets together into a single set



Also, single element sets: {8} {4} {5} {11} ...

Representing Disjoint Sets

- In this context, only two set operations make sense:
 - **Find**: check if an element belongs to a particular set (or if two belong to the same set)
 - **Union**: merge two sets together into a single set
- Consequently, data structures for representing disjoint sets are often referred to as "**union-find**".
 - How can we implement this so that the two operations are as efficient as possible?
 - We also consider the amortized complexity over the entire lifetime of operations. (i.e. start with n single element sets, end up with one big set)

Union-Find: Example 1

Separate Containers

Items: 1 2 3 4 5 6 7 8 9 10

Group 1: 1 4 5 8 10

Group 2: 2 6 7

Are 1 and 8 connected? (i.e. in the same set?)

If we add a connection between 1 and 6,
how long does it take to merge the two sets?

Union(): $O(n)$

Find(): $O(n)$

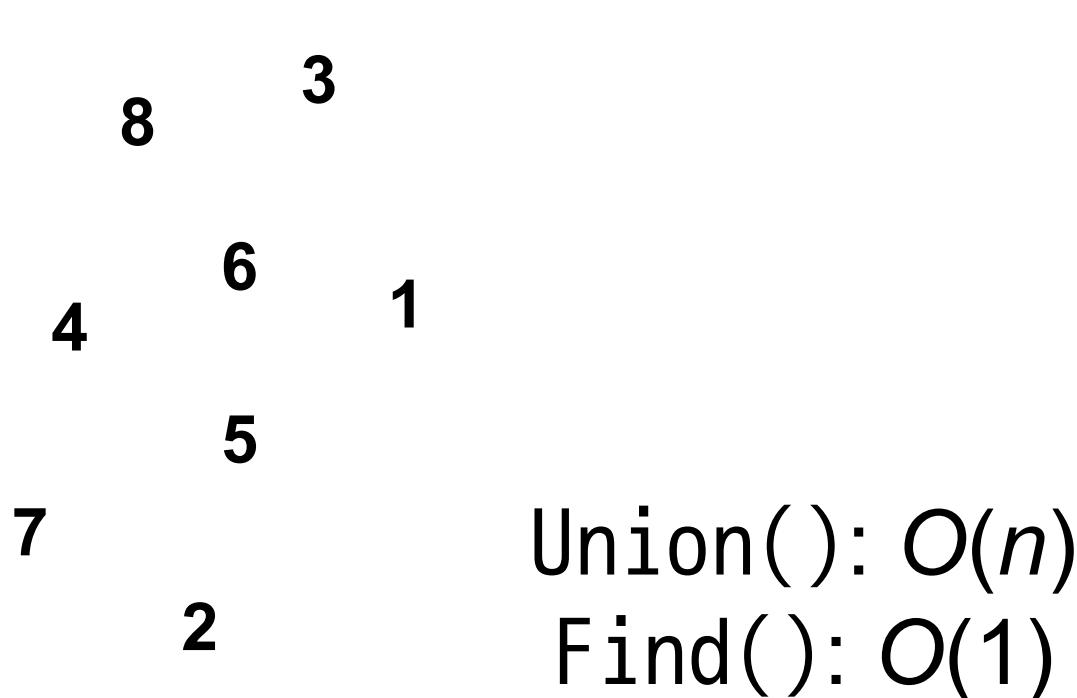
Union-Find Data Structure

- **Idea 1:** every *disjoint* set should have its unique representative (selected element)
 - Every set element k must know its representative j
- **Therefore:** to tell if k and m are in the same set, *compare their representatives*
 - Find() operation becomes quite fast!

Union-Find: Example 2

Representatives

Item	1	2	3	4	5	6	7	8
Representative								



- 1) Join 1 & 3
- 2) Join 3 & 8
- 3) Join 1 & 5
- 4) Join 7 & 4
- 5) Join 7 & 2
- 6) Join 2 & 5

Making Union-Find Faster

- **Idea 2:** When performing union of two sets, update the smaller set (less work)
- Measure complexity of all unions throughout the lifecycle (together)
 - We call Union() exactly $n - 1$ times
 - If we connect to a disjoint element every time, it will take n time total (best case)
 - Merging large sets, say $n/2$ and $n/2$ elements, will take $O(n)$ time for one Union() – too slow!

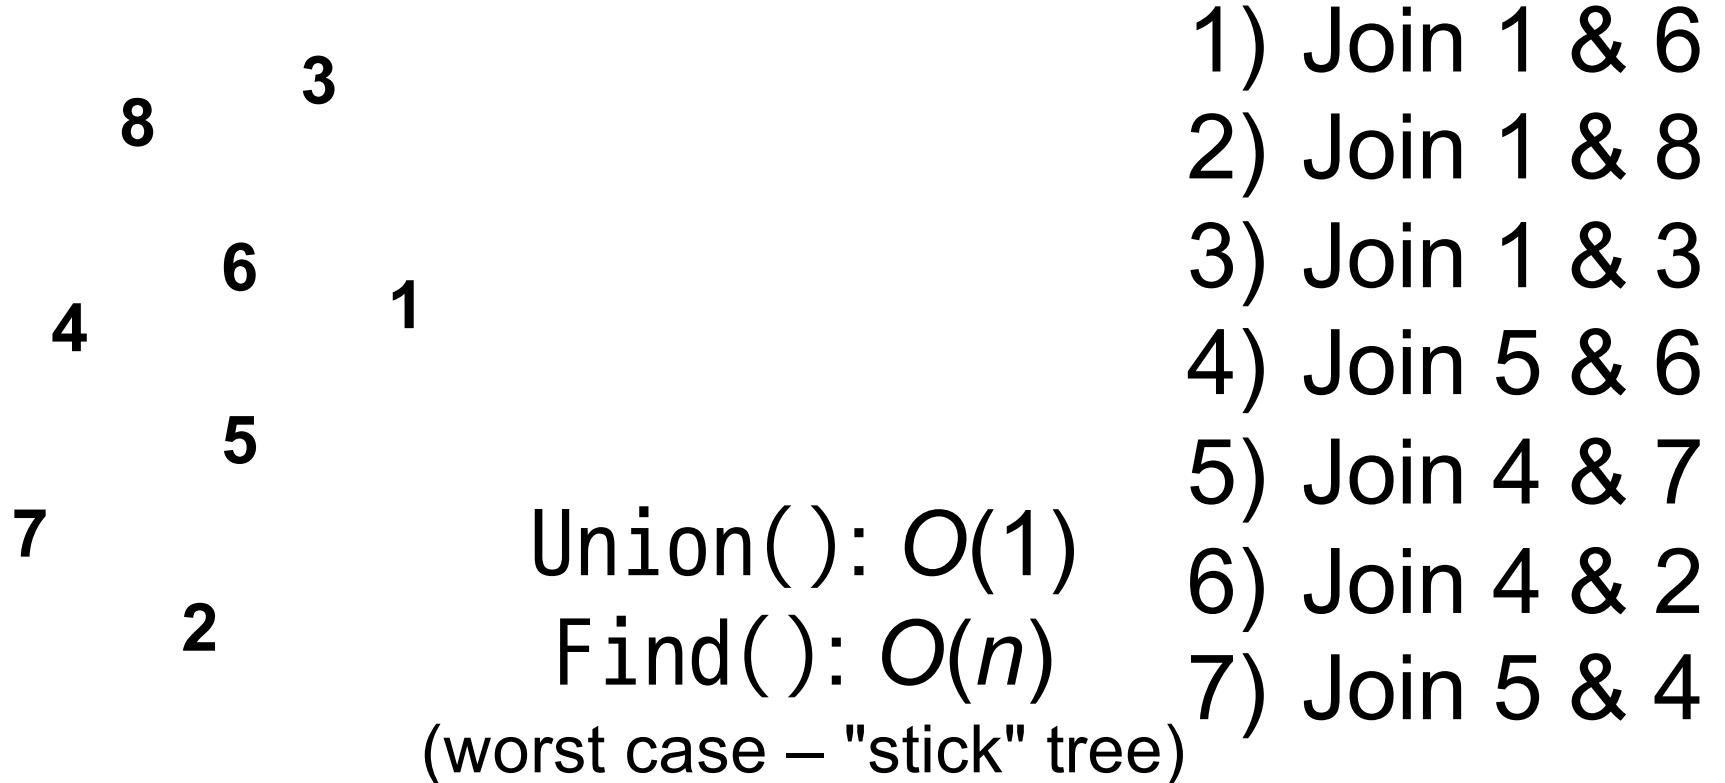
Smarter Union-Find

- **Idea 3:** No need to store the actual representative for each element, as long as you can find it quickly
 - Each element knows someone who knows the representative (may need more steps)
 - Union() becomes very fast: one of representatives will need to know the other
 - Find() becomes slower

Union-Find: Example 3

Hierarchical Representatives

Item	1	2	3	4	5	6	7	8
Representative								



Path Compression

- So far, `Find()` was read-only
 - For element j , finds the representative k
 - Traverses other elements on the way (for which k is also the representative)
- **Idea 4:** We can tell j that its representative is now k
 - Same for other elements on path from $j \rightarrow k$
 - Doubles workload of `Find()`, but same Big-O

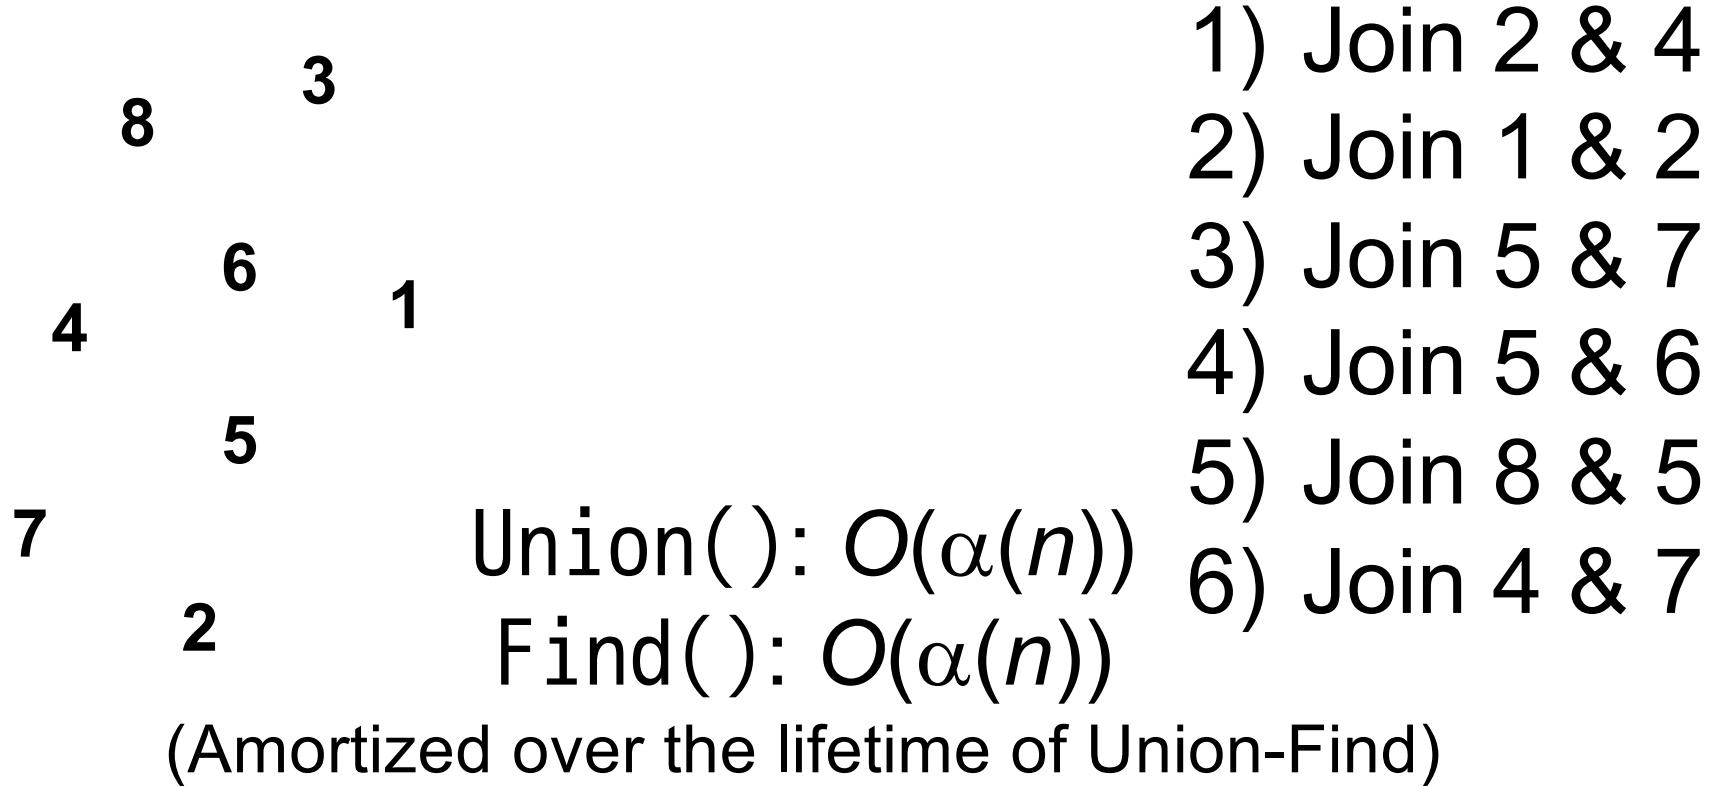
Complexity with Path Compression

- Must use amortized analysis over the life cycle of union-find (starting with n disjoint sets, and merging until there is one set containing all elements)
- Result is surprising
 - $O(n^*\alpha(n))$, where $\alpha()$ grows very slowly
 - $\alpha()$ is the inverse-Ackerman function
 - In practice, almost-linear-time performance
- Details taught in more advanced courses

Union-Find: Example 4

Path Compression

Item	1	2	3	4	5	6	7	8
Representative								



Union-Find

Data Structures & Algorithms

Study Questions



- What is the difference between a sorted and an ordered container?
- When should you implement a sorted container with an array instead of a linked list?
- When should you implement an ordered container with an array instead of a linked list?
- What is binary search? Study STL's interface to it.
- What are comparison operators and comparator objects?
- How are searchable containers and sets related?
- What is a universe set?
- Give an example of a universe set and a subset of it.
- Implement `set_intersection()`.
- How would you implement a Union-Find data structure?