

Lectures 17-18

Trees, Binary Search Trees



EECS 281: Data Structures & Algorithms

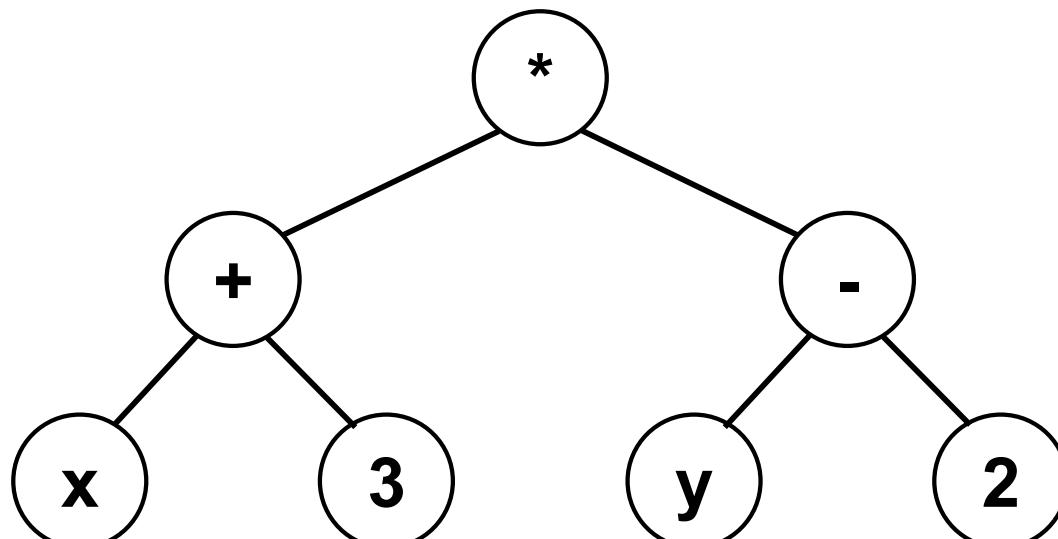
Tree Fundamentals

Data Structures & Algorithms

Context for Trees

Tree: a mathematical abstraction that

- Captures common properties of data
- Critical to the design and analysis of algorithms



$$(x + 3) * (y - 2)$$

Trees



A **graph** consists of **nodes** (sometimes called vertices) connected together by **edges**.

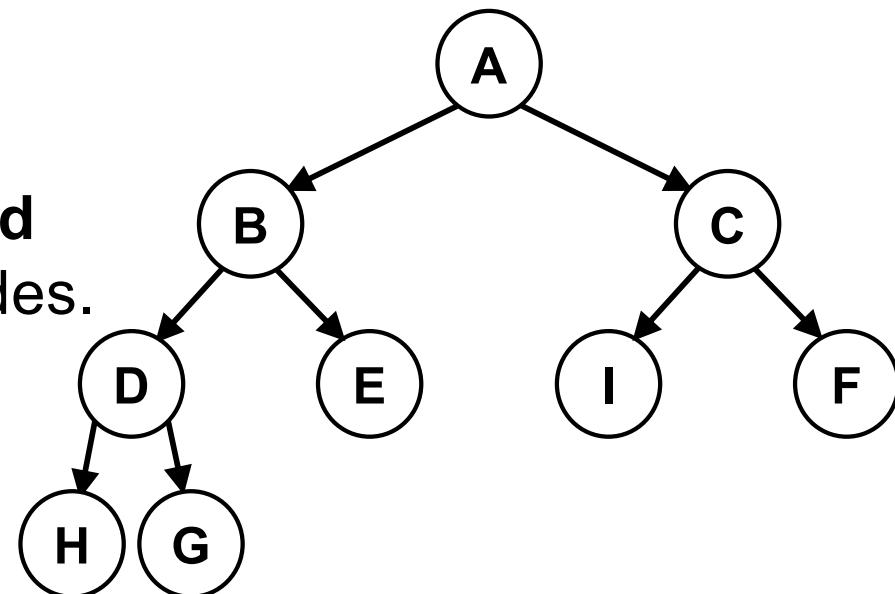
Each node can contain some data.

A **tree** is:

- (1) a connected graph (nodes + edges) w/o cycles.
- (2) a graph where any 2 nodes are connected by a unique shortest path.
(the two definitions are equivalent)

In a directed tree, we can identify **child** and **parent** relationships between nodes.

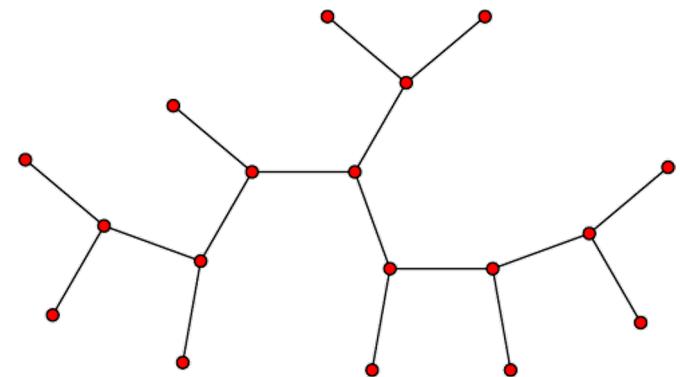
In a **binary tree**, a node has at most two children.



Types of Trees

1. (Simple) tree

- Acyclic connected graph
- Considered *undirected*



2. Rooted tree

- A simple tree with a selected node (root)
- All edges are *directed* away from root

- Any node could be selected as root

Tree Terminology

Root: the “topmost” node in the tree

Parent: Immediate predecessor of a node

Child: Immediate successor of a node

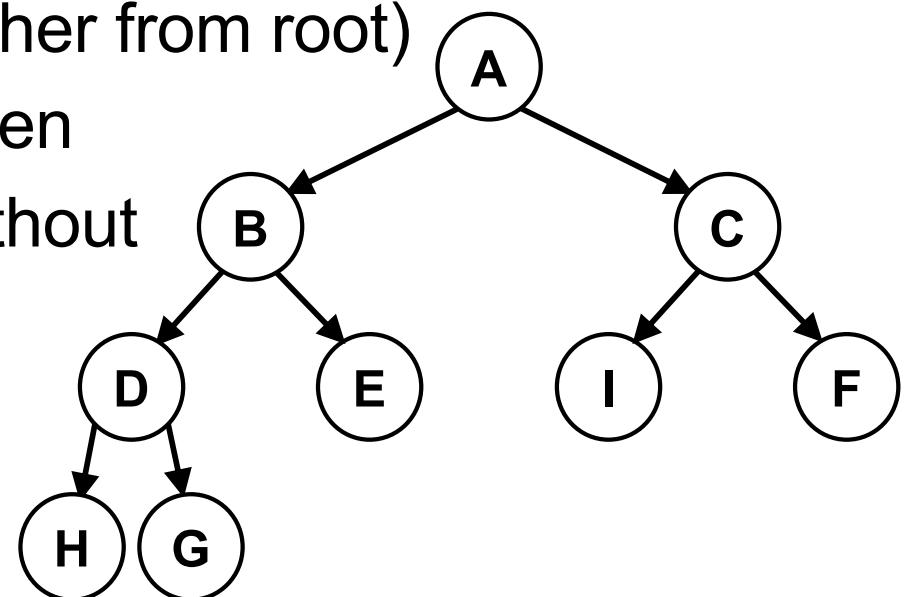
Siblings: Children of the same parent

Ancestor: parent of a parent (closer to root)

Descendent: child of a child (further from root)

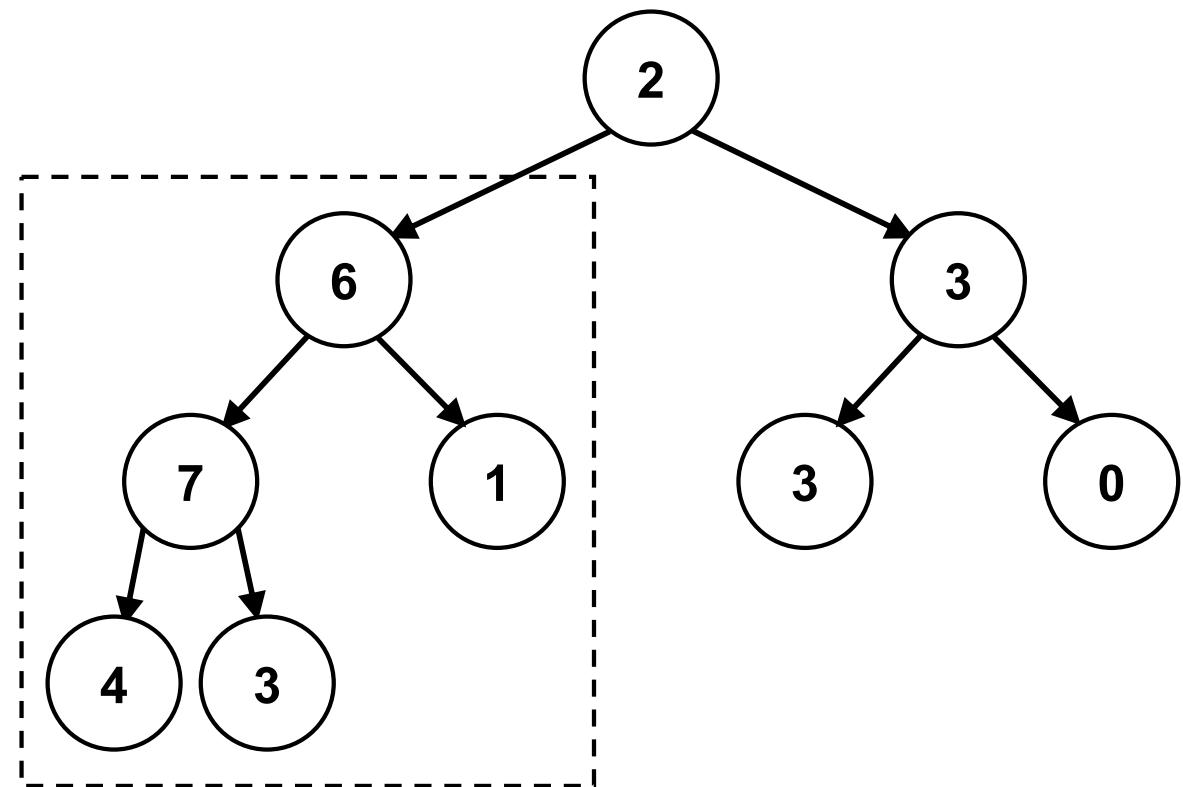
Internal node: a node with children

External (Leaf) node: a node without
children



Trees are Recursive Structures

Any subtree is just
as much a "tree"
as the original!



Tree Properties

Height:

$\text{height}(\text{empty}) = 0$

$\text{height}(\text{node}) = \max(\text{height}(\text{left_child}), \text{height}(\text{right_child})) + 1$

Size:

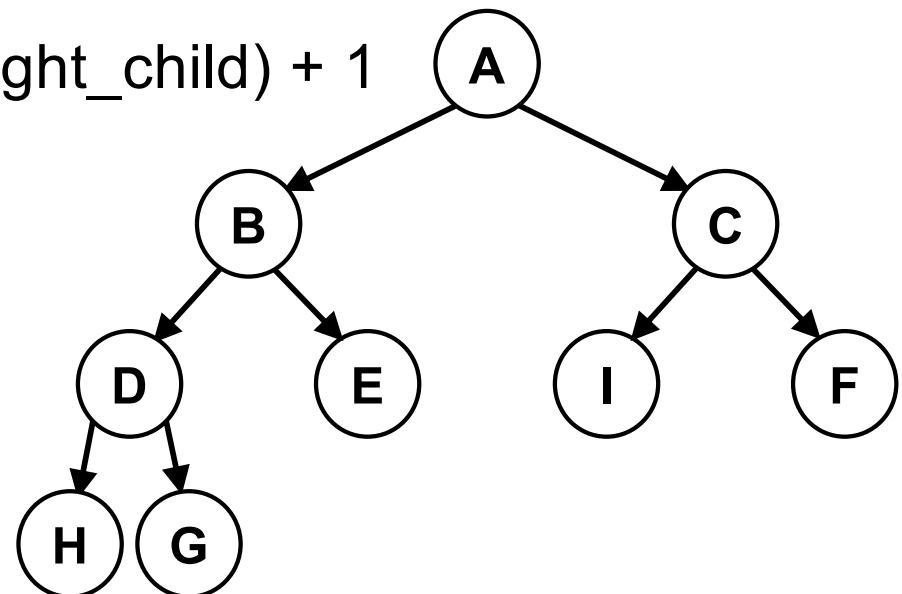
$\text{size}(\text{empty}) = 0$

$\text{size}(\text{node}) = \text{size}(\text{left_child}) + \text{size}(\text{right_child}) + 1$

Depth:

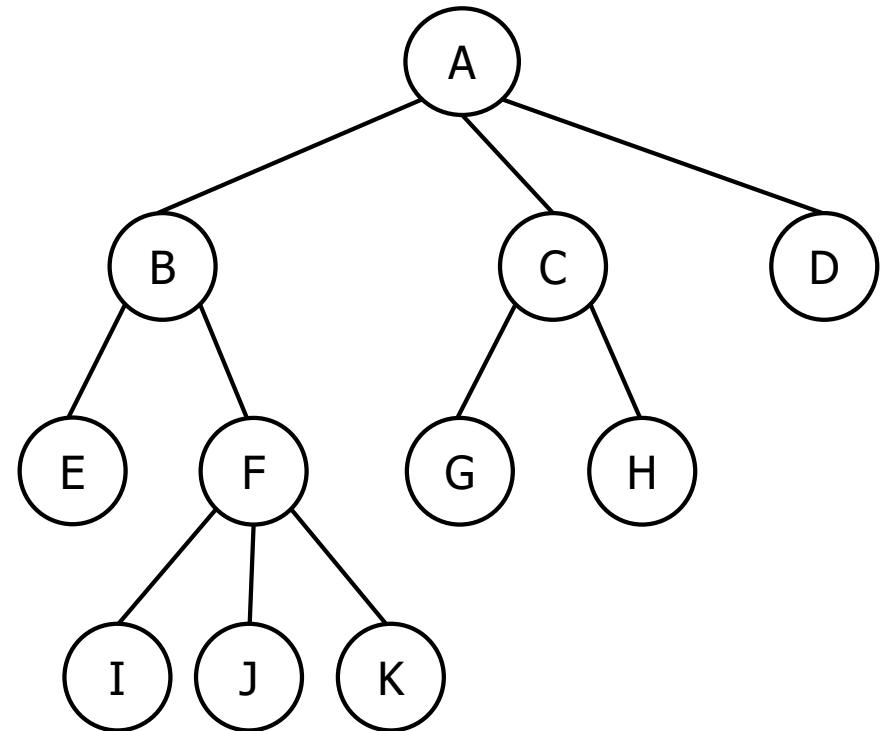
$\text{depth}(\text{empty}) = 0$

$\text{depth}(\text{node}) = \text{depth}(\text{parent}) + 1$



Tree Quiz

- Root:
- Internal nodes:
- Leaf nodes:
- Max depth:
- Label one subtree
- Is this a binary tree?



Tree Fundamentals

Data Structures & Algorithms

Binary Trees

Data Structures & Algorithms

Binary Trees

- **Ordered Tree:** linear ordering for the children of each node
- **Binary Tree:** ordered tree in which every node has at most two children
- Multiple binary tree implementation styles

Complete Binary Tree Property

Definition: ***complete binary tree***

- A binary tree with depth d where
 - Tree depths 1, 2, ..., $d - 1$ have the max number of nodes possible
 - All internal nodes are to the left of the external nodes at depth $d - 1$
 - That is, all leaves are at $d - 1$ or leftmost at depth d

Binary Tree Implementation

Binary tree array implementation

- Root at index 1
- Left child of node i at $2 * i$
- Right child of node i at $(2 * i) + 1$
- Some indices may be skipped
- Can be space prohibitive for sparse trees

Binary Tree Implementation

Complexity of array implementation

- Insert key (best-case) $O(1)$
- Insert key (worst-case) $O(n)$
- Remove key (worst-case) $O(n)$
- Parent $O(1)$
- Child $O(1)$
- Space (best-case) $O(n)$
- Space (worst-case) $O(2^n)$

Binary Tree Implementation

- Pointer-based binary tree implementation

```
1 template <class KEY>
2 struct Node {
3     KEY key;
4     Node *left = nullptr;
5     Node *right = nullptr;
6     Node(const KEY &k) : key{k} {}
7 }; // Node{}
```

- A node contains some information, and points to its left child node and right child node
- Efficient for moving *down* a tree from parent to child

Binary Tree Implementation

Complexity of pointer implementation

- Insert key (best-case) $O(1)$
- Insert key (worst-case) $O(n)$
- Remove key (worst-case) $O(n)$
- Parent $O(n)$
- Child $O(1)$
- Space (best-case) $O(n)$
- Space (worst-case) $O(n)$

Binary Tree++

Another way to do it (not common)

```
1 template <class KEY>
2 struct Node {
3     KEY key;
4     Node *parent, *left, *right;
5 }; // Node()
```

- If node is root, then ***parent** is **nullptr**
- If node is external, then ***left** and ***right** are **nullptr**

Translating General Trees into Binary Trees

T : General tree

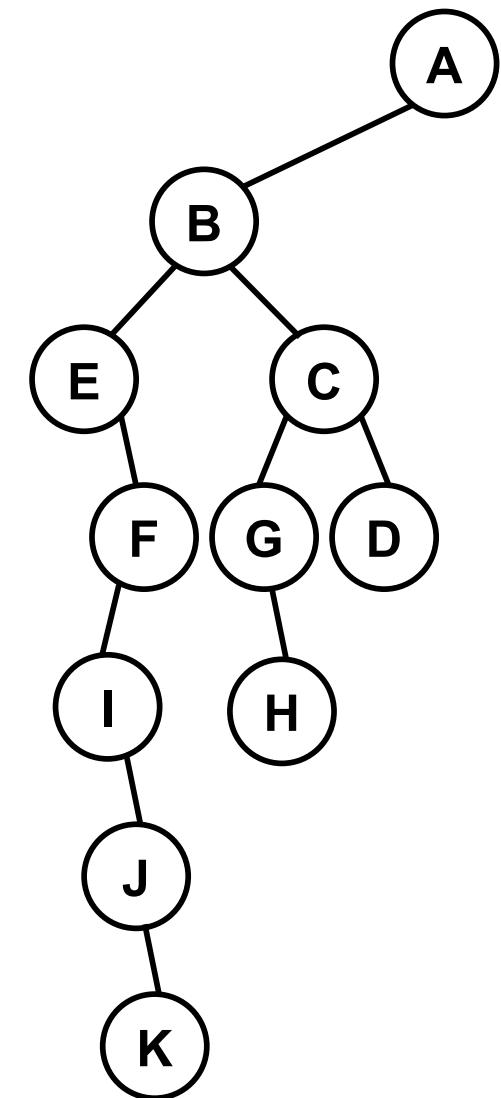
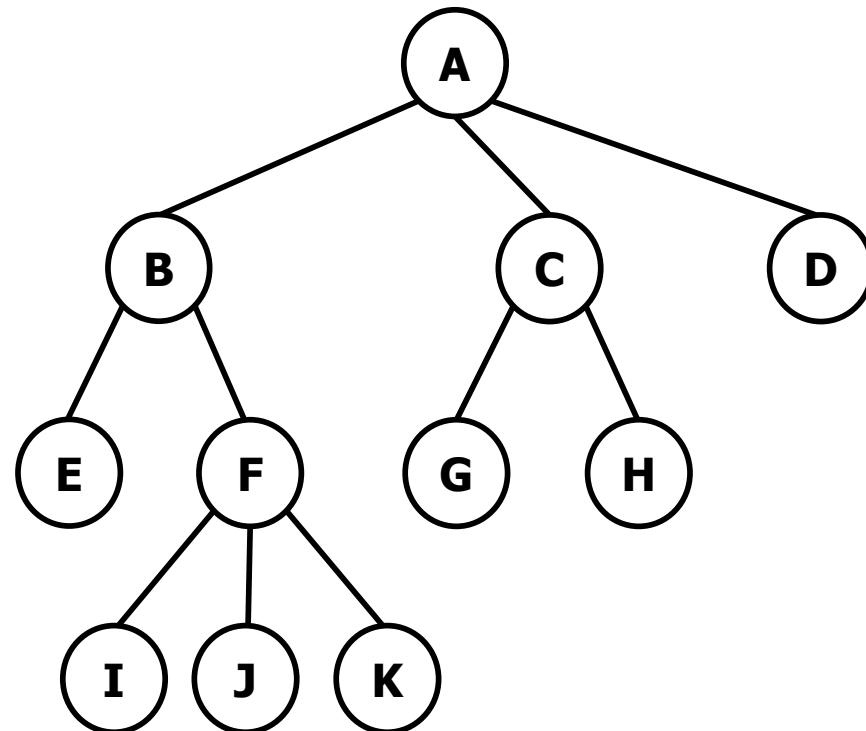
T' : Binary tree

Intuition:

- Take set of siblings $\{v_1, v_2, \dots, v_k\}$ in T , that are children of v
- v_1 becomes left child of v in T'
- $v_2 \dots v_k$ become chain of right children of v_1 , in T'
- Recurse from v_2

Left: new “generation”; Right: sibling

Tree Quiz



Binary Trees

Data Structures & Algorithms

Tree Traversals

Data Structures & Algorithms

Tree Traversal

Systematic method of processing every node in a tree

- Preorder
 1. Visit node
 2. Recursively visit left subtree
 3. Recursively visit right subtree
- Inorder
 1. Recursively visit left subtree
 2. Visit node
 3. Recursively visit right subtree

Tree Traversal

Systematic method of processing every node in a tree

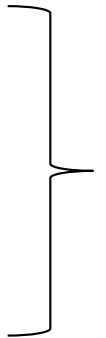
- Postorder
 - 1. Recursively visit left subtree
 - 2. Recursively visit right subtree
 - 3. Visit node
- Level order
 - Visit nodes in order of increasing depth in tree

Recursive Implementations

```
1 template <class KEY>
2 struct Node {
3     KEY key;
4     Node *left = nullptr;
5     Node *right = nullptr;
6 }; // Node{}
7 void preorder(Node *p) {
8     if (!p) return;
9     visit(p->key);
10    preorder(p->left);
11    preorder(p->right);
12 } // preorder()
13 void postorder(Node *p) {
14     if (!p) return;
15     postorder(p->left);
16     postorder(p->right);
17     visit(p->key);
18 } // postorder()
19 void inorder(Node *p) {
20     if (!p) return;
21     inorder(p->left);
22     visit(p->key);
23     inorder(p->right);
24 } // inorder()
```

Summary of Tree Traversals

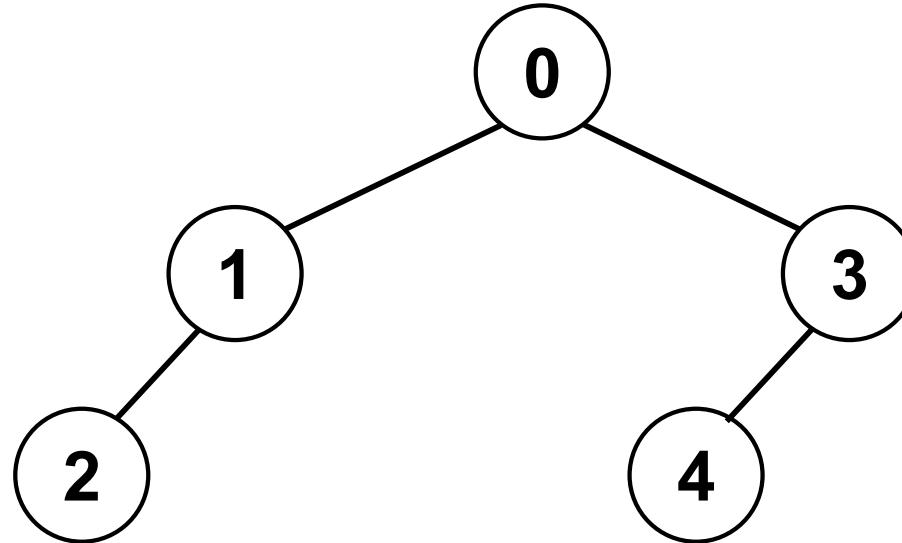
Tree traversal: Systematically process all nodes in a tree

- Preorder
 - Inorder
 - Postorder
 - Level order (breadth-first traversal)
- 
- All are depth-first traversals

Exercise

In what order are nodes visited?

- Preorder
- Inorder
- Postorder
- Level order



Applied Tree Traversals

Draw the binary tree from traversals

- Preorder: 7 3 6 9 8 13 27
- Inorder: 3 6 7 8 9 13 27

An Iterative Traversal

Write a function to do a level order traversal,
printing the key at each node

```
void levelorder(Node *p) {  
  
} // levelorder()
```

Tree Traversals

Data Structures & Algorithms

Search

- Retrieval of a particular piece of information from large volumes of previously stored data
- Purpose is typically to access information within the item (not just the key)
- Recall that arrays, linked lists are worst-case $O(n)$ for either searching or inserting
- Even a hash table has worst-case $O(n)$

Need a data structure with optimal efficiency for searching and inserting.

What if order is important?

Symbol Table: ADT

- **insert** a new item
- **search** for an item (items) with a given key
- **remove** an item with a specified key
- **sort** the symbol table
- **select** the item with the k^{th} largest key
- **join** two symbol tables

Also may want construct, test if empty, destroy, copy...

Binary Search Trees

Data Structures & Algorithms

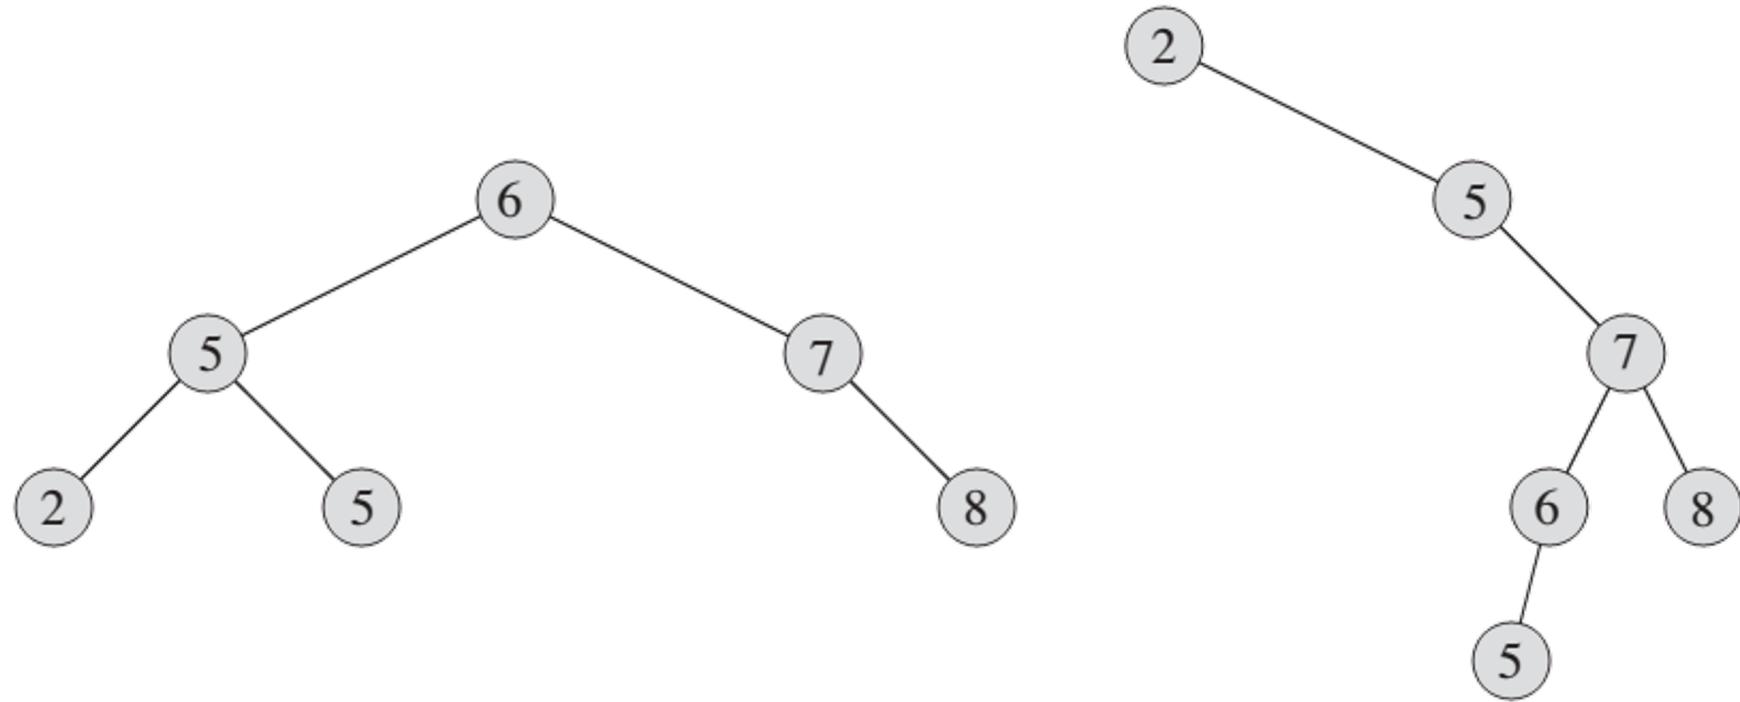
Binary Search Tree

- The keys in a binary search tree satisfy the Binary Search Tree Property
 - The key of any node is:
 - > keys of all nodes in its left subtree and
 - \leq keys of all nodes in its right subtree
- Essential property of BST is that `insert()` is as easy to implement as `search()`

Binary Search Tree Property

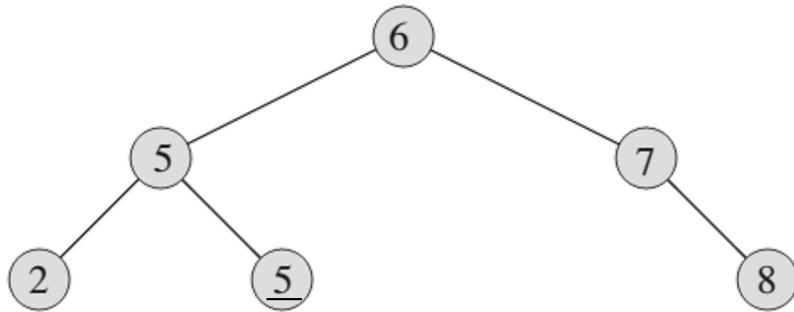
The key of any node is:

- > keys of all nodes in its left subtree and
- \leq keys of all nodes in its right subtree



Exercise

Write the output for inorder, preorder and post-order traversals of this BST



inorder: 2, 5, 5, 6, 7, 8

preorder: 6, 5, 2, 5, 7, 8

postorder: 2, 5, 5, 8, 7, 6

```
1 void inorder(Node *x) {  
2     if (!x) return;  
3     inorder(x->left);  
4     print(x->key);  
5     inorder(x->right);  
6 } // inorder()  
7  
8 void preorder(Node *x) {  
9     if (!x) return;  
10    print(x->key);  
11    preorder(x->left);  
12    preorder(x->right);  
13 } // preorder()  
14  
15 void postorder(Node *x) {  
16     if (!x) return;  
17     postorder(x->left);  
18     postorder(x->right);  
19     print(x->key);  
20 } // postorder()
```

Sort: Binary Search Tree

Can you think of an easy method of sorting using a binary search tree?

Search

- How can we find a key in a binary search tree?

```
// return a pointer to node with key k if  
// one exists; otherwise, return nullptr  
Node *tree_search(Node *x, KEY k);
```

- BST Property - the key of any node is:
 - $>$ keys of all nodes in its left subtree
 - \leq keys of all nodes in its right subtree
- What are the average- and worst-case complexities?

Search

```
1 // return a pointer to node with key k if
2 // one exists; otherwise, return nullptr
3 Node *tree_search(Node *x, KEY k) {
4     while (x != nullptr && k != x->key) {
5         if (k < x->key)
6             x = x->left;
7         else
8             x = x->right;
9     } // while
10    return x;
11 } // tree_search()
```

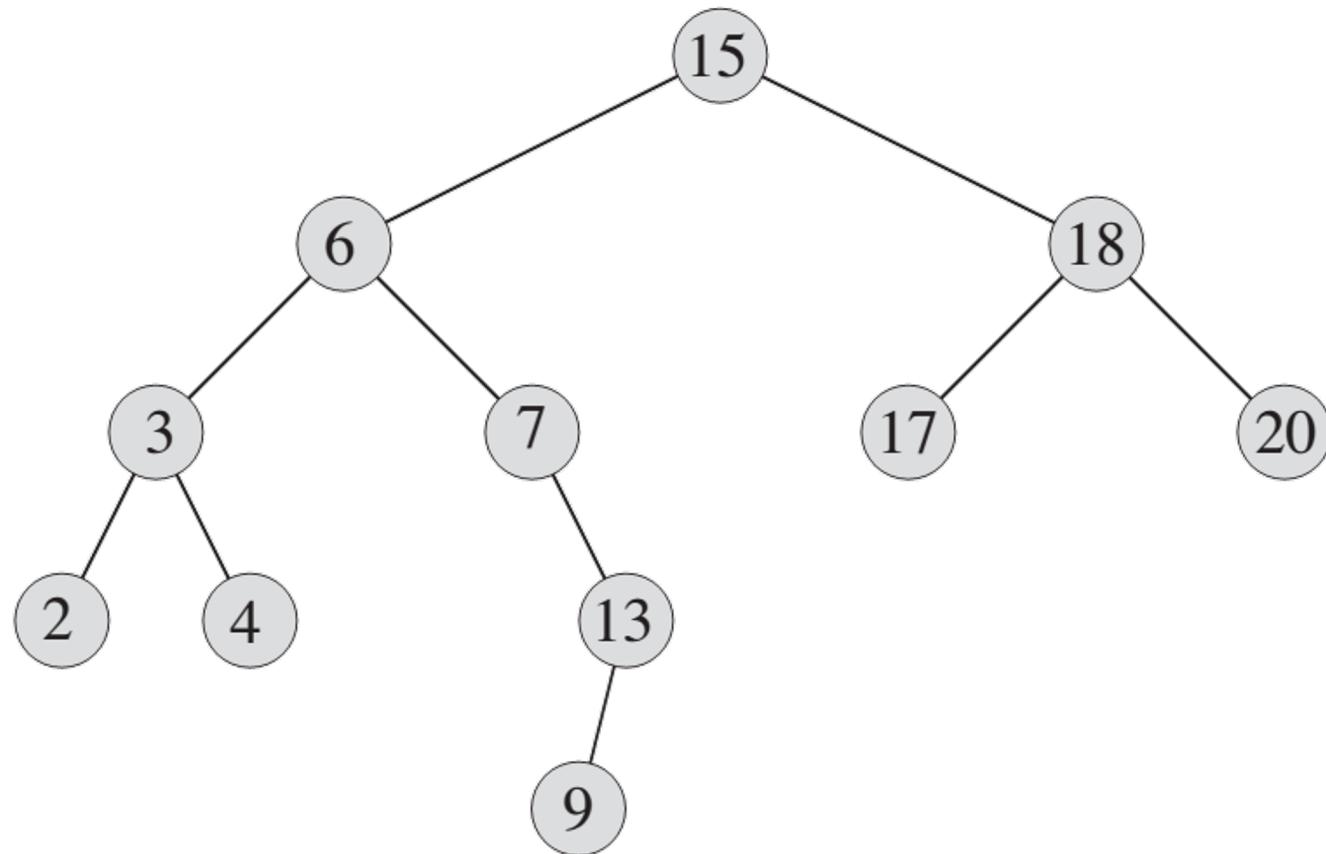
Search

```
1 // return a pointer to node with key k if
2 // one exists; otherwise, return nullptr
3 Node *tree_search(Node *x, KEY k) {
4     if (x == nullptr || x->key == k)
5         return x;
6     if (k < x->key)
7         return tree_search(x->left, k);
8     return tree_search(x->right, k);
9 } // tree_search()
```

- To search for a key k , we trace a downward path starting at the root
- The next node visited depends on the outcome of the comparison of k with the key of the current node

Search Example

```
tree_search(root, 9);
```



Search

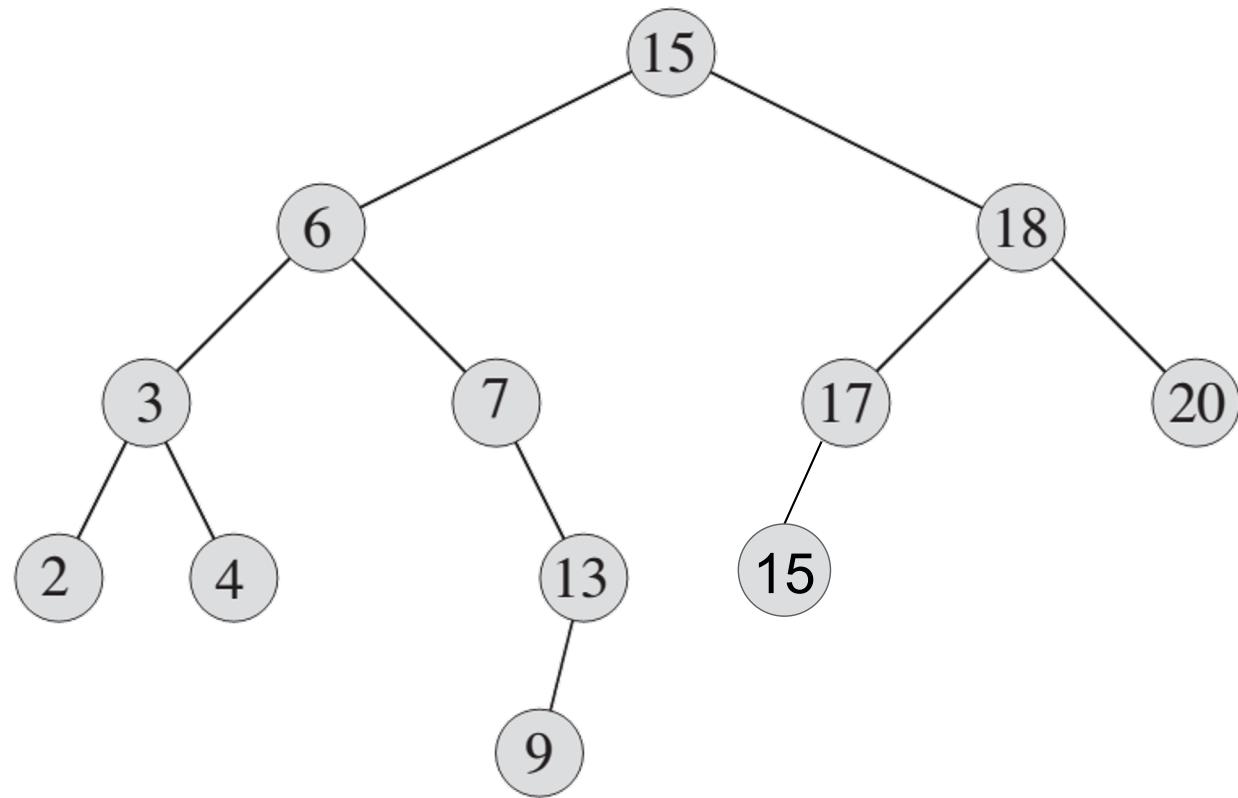
- Complexity is $O(h)$, where h is the (maximum) height of the tree
- Average-case complexity: $O(\log n)$
 - Balanced tree
- Worst-case complexity: $O(n)$
 - "Stick" tree

Insert

- How do we insert a new key into the tree?
- Similar to search
- Start at the root, and trace a path downwards, looking for a null pointer to append the node

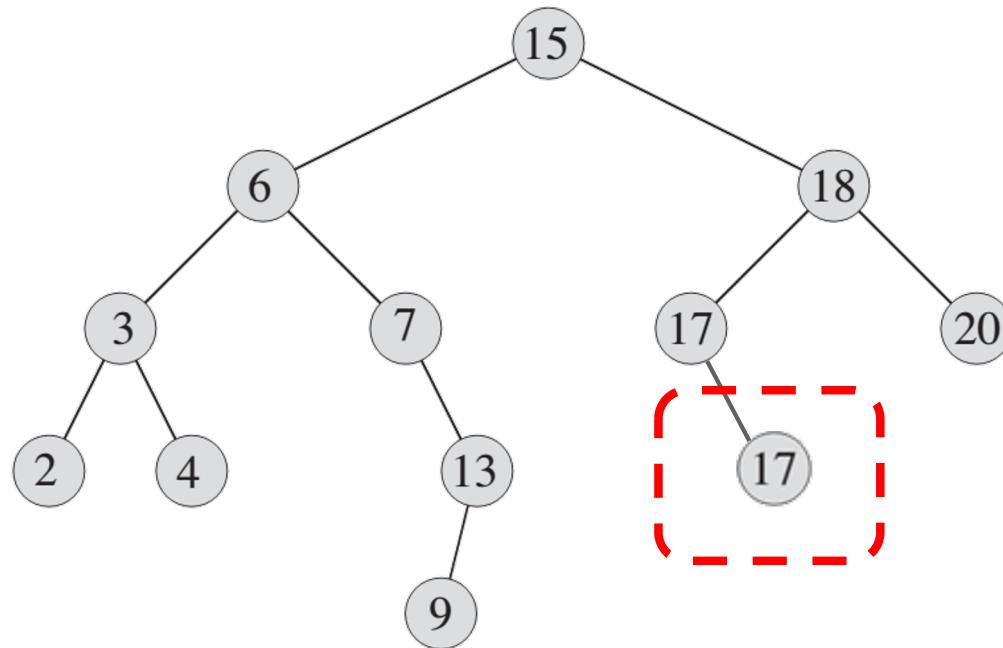
Insert Example

```
tree_insert(root, 15);
```



Insert with Duplicates

- For sets with no duplicates, use ($<$, $>$)
- For duplicates, we need a deterministic policy
 - Choose (\leq , $>$) or ($<$, \geq)... slides, STL use ($<$, \geq)



Insert

```
1 void tree_insert(Node *&x, KEY k) {  
2     if (x == nullptr)  
3         x = new Node(k);  
4     else if (k < x->key)  
5         tree_insert(x->left, k);  
6     else  
7         tree_insert(x->right, k);  
8 } // tree_insert()
```

- New node inserted at leaf
- Note the use of reference-to-pointer-to-Node
- Exercise: modify this code to set the parent pointer

Exercise

- Start with an empty tree
- Insert these keys, in this order:
12, 5, 18, 2, 9, 15, 19, 17, 13
- Draw the tree
- Write a new order to insert the same keys which generates a worst-case tree
- How many worst-case trees are possible for n unique values?

Complexity

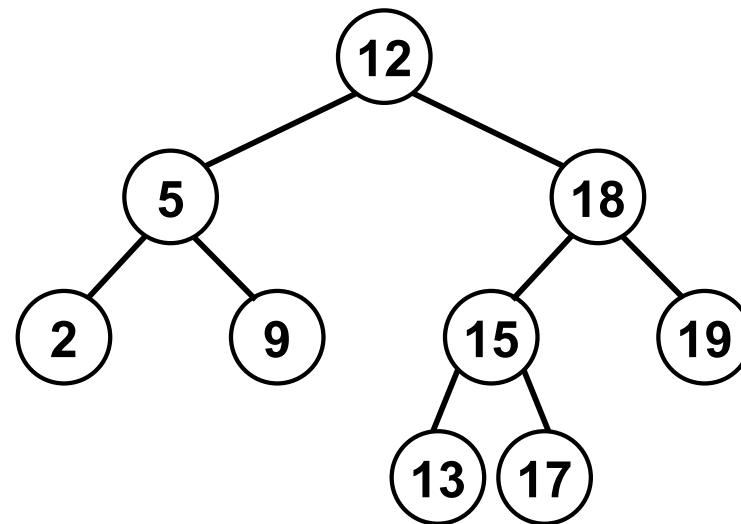
- The complexity of insert (and many other tree functions) depends on the height of the tree
- Average-case (balanced): $O(\log n)$
- Worst-case (unbalanced "stick"): $O(n)$
- Average-case:
 - Random data
 - Likely to be well-balanced

Exercise

- Write a function to find the Node with the smallest key
- What are the average- and worst-case complexities?

// Return a pointer to the Node with the min key

```
Node *tree_min(Node *x);
```

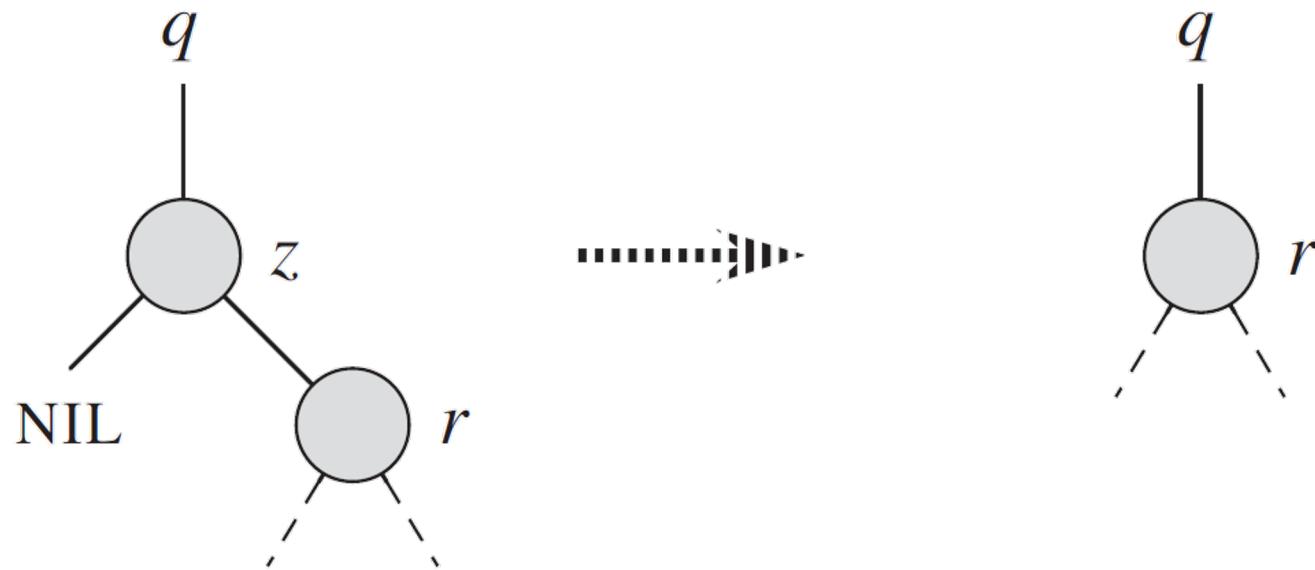


Remove

- What if we want to remove a node?
- To remove node z :
 1. z has no children (trivial)
 2. z has no left child
 3. z has no right child
 4. z has two children
- Complete algorithm is in CLRS 12.3

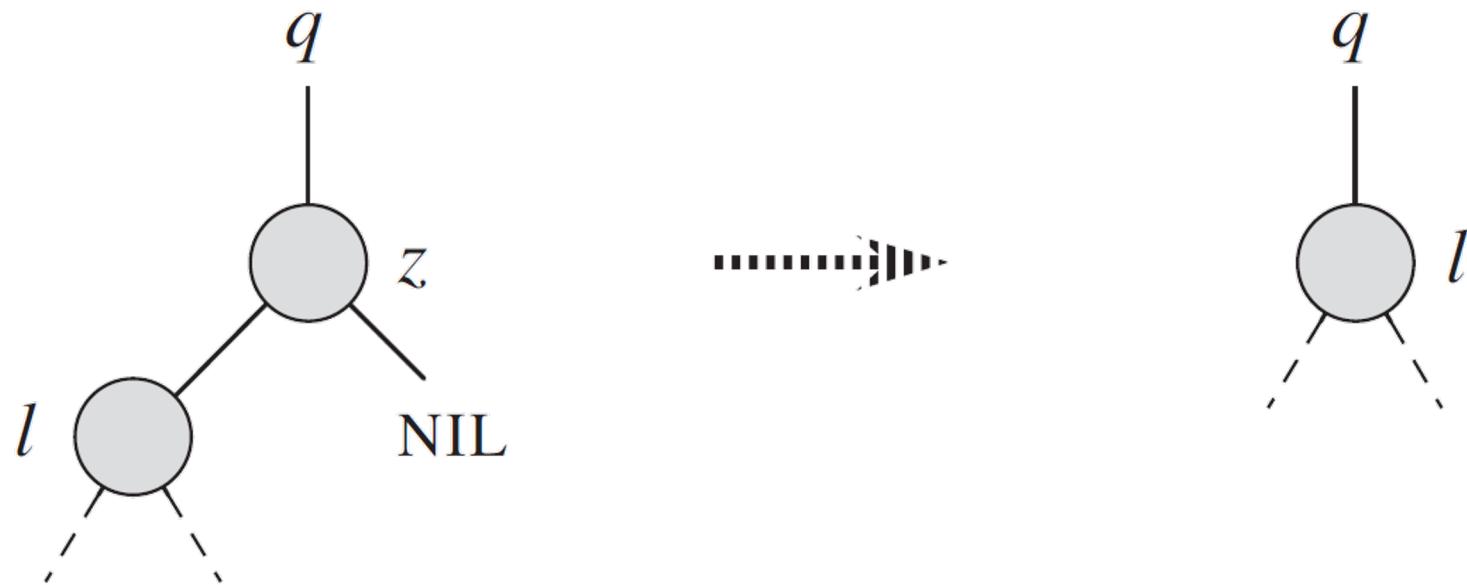
Remove: Easy Case #1

z has no left child: replace z by right child



Remove: Easy Case #2

z has no right child: replace z by left child

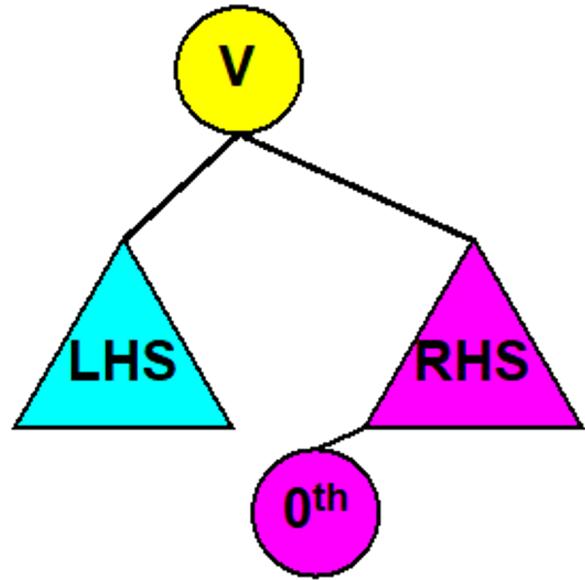


Remove: Hard Case

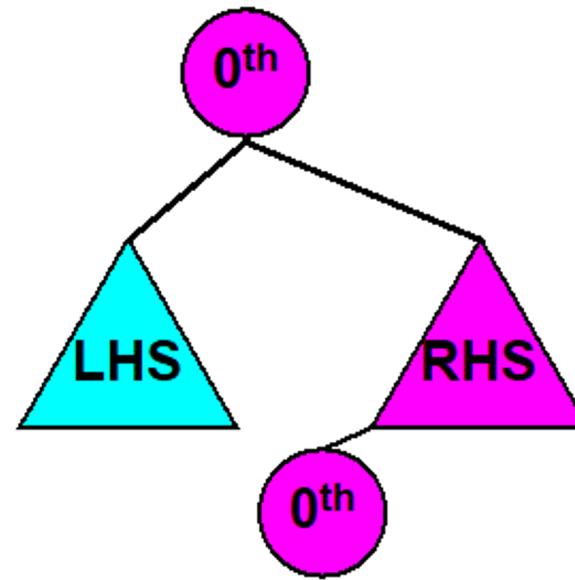
- z has left and right children
- Replace with a “combined” tree of both
- Key observation
 - All in LHS subtree < all in RHS subtree
 - Transplant smallest RHS node to root
 - Called the inorder successor
 - Must be some such node, since RHS is not empty
 - New root might have a right child, but no left child
 - Make new root’s left child the LHS subtree

Remove: Hard Case

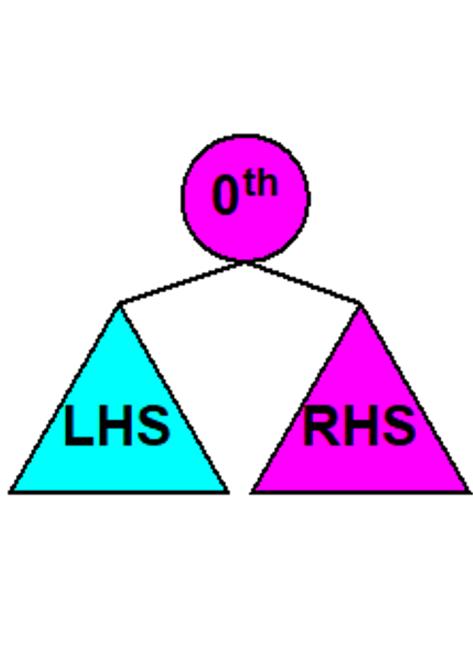
To Remove V:



Find smallest
node in RHS



Replace V
with smallest
node's value



Remove V
from RHS

Single-Function Remove 1/3

```
1 template <class T>
2 void BinaryTree<T>::remove(Node *&tree, const T &val) {
3     Node *nodeToDelete = tree;
4     Node *inorderSuccessor;
5
6     // Recursively find the node containing the value to remove
7     if (tree == nullptr)
8         return;
9     else if (val < tree->value)
10        remove(tree->left, val);
11     else if (tree->value < val)
12        remove(tree->right, val);
13     else {
```

Single-Function Remove 2/3

```
14 // Check for simple cases where at least one subtree is empty
15 if (tree->left == nullptr) {
16     tree = tree->right;
17     delete nodeToDelete;
18 } // if
19 else if (tree->right == nullptr) {
20     tree = tree->left;
21     delete nodeToDelete;
22 } // else if
```

Single-Function Remove 3/3

```
23     else {
24         // Node to delete has both left and right subtrees
25         inorderSuccessor = tree->right;
26
27         while (inorderSuccessor->left != nullptr)
28             inorderSuccessor = inorderSuccessor->left;
29
30         // Replace value with the inorder successor's value
31         nodeToDelete->value = inorderSuccessor->value;
32         // Remove the inorder successor from right subtree
33         remove(tree->right, inorderSuccessor->value);
34     } // else
35 } // else
36 } // BinaryTree::remove()
```

Summary: Binary Search Trees

- Each node points to two children (left, right), and possibly a parent
- All nodes are ordered: $\text{left} < \text{root} \leq \text{right}$
- Modification of nodes
 - External is easy
 - Internal is more complicated
- In general, operations on BSTs are:
 - $O(\log n)$ average-case
 - $O(n)$ worst-case

Binary Search Trees

Data Structures & Algorithms

AVL Trees

Data Structures & Algorithms

AVL Tree

- Self-balancing Binary Search Tree
- Named for Adelson-Velsky, and Landis
- Start with a BST
- Add the *Height Balance Property*
 - For every internal node v of T , the heights of the children of v differ by at most 1
 - Use *rotations* to correct imbalance
- Worst-case search/insert is now $O(\log n)$

Tree Height

- Measured upward from leaf nodes; all of which have height equal to 1
- Independent from depth
- Recursive formula for a recursive data structure
 - $\text{height}(\text{empty}) = 0;$
 - $\text{height}(\text{node}) = \max(\text{height}(\text{children})) + 1;$

Is this an AVL tree?

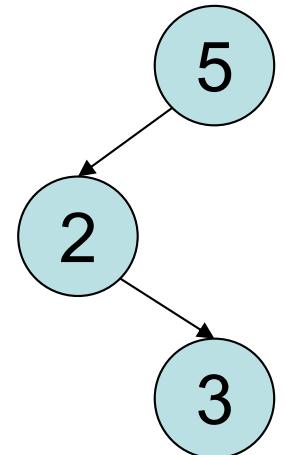
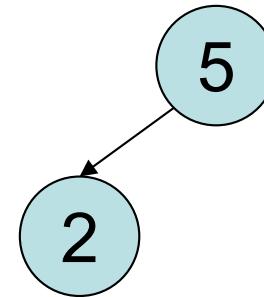
Height Balance Property: For every internal node v of T , the heights of the children of v differ by at most 1.

Tree 0 ✓

Tree 1 ✓

Tree 2 ✓

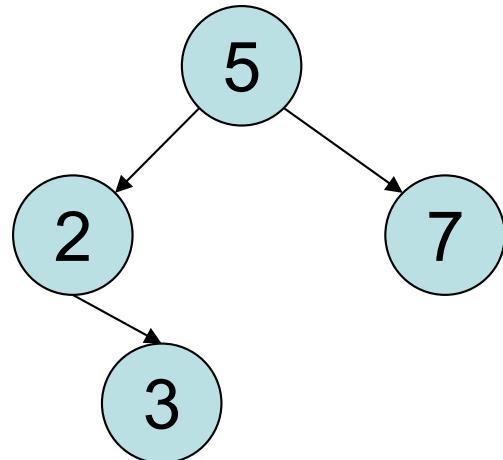
Tree 3 ✗



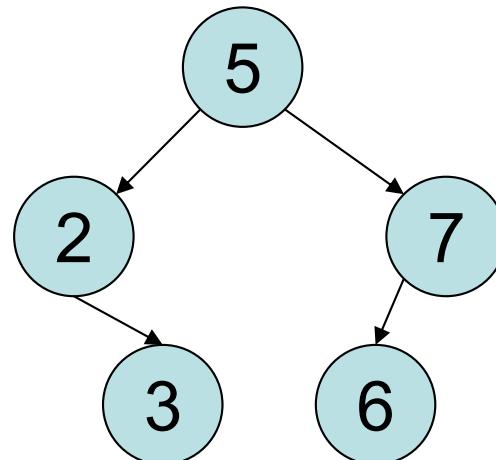
Is this an AVL tree?

Height Balance Property: For every internal node v of T , the heights of the children of v differ by at most 1.

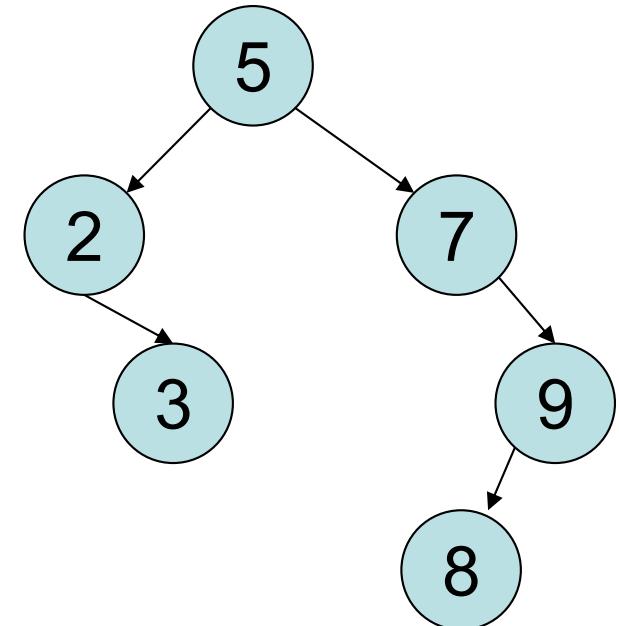
Tree 4 ✓



Tree 5 ✓

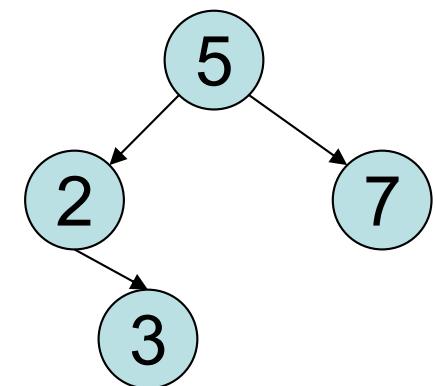


Tree 6 ✗



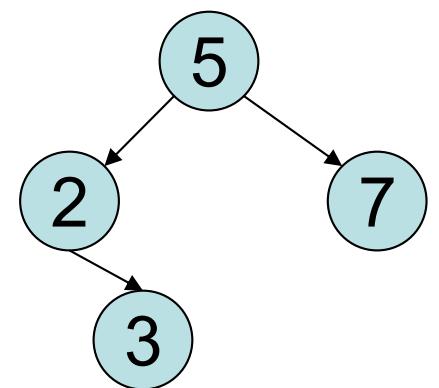
AVL Tree: Proof Setup

- h : height of tree
- $n(h)$: *minimum* number of nodes in AVL tree of height h
- $n(0) = 0$ (Empty tree)
- $n(1) = 1$ (Root-only tree)
- For $h > 1$, an AVL tree of height h contains:
 - Root Node
 - AVL Subtree of height $h - 1$
 - AVL Subtree of height $h - 2$
- Thus for $h > 1$, $n(h) = 1 + n(h - 1) + n(h - 2)$



Proof: Height Balance Property

- h : height of tree
 - $n(h)$: minimum number of nodes in a tree of height h
 - $n(h) = 1 + n(h - 1) + n(h - 2)$
-
- Knowing $n(h - 1) > n(h - 2)$ and $n(h) > 2n(h - 2)$,
then by induction, $n(h) > 2^i n(h - 2i)$
 - Closed form solution, $n(h) > 2^{h/2} - 1$
 - Taking logarithms: $h < 2 \log n(h) + 2$
 - Thus the height of the tree is **$O(\log n)$**



AVL Tree Algorithms

- Search is the same as a BST
- Sort (same as BST)
 - Insert nodes one at a time
 - What is worst-case complexity now?
 - Perform an inorder traversal
 - Still $O(n)$ for this step

AVL Tree Insert

- The basic idea:
 1. Insert like a BST
 2. Rearrange tree to balance height
- Each node records its height
- Can compute a node's *balance factor*
 - $\text{balance}(n) = \text{height}(n\rightarrow\text{left}) - \text{height}(n\rightarrow\text{right})$
- A node that is AVL-balanced:
 - $\text{balance}(n) = 0$: both subtrees equal
 - $\text{balance}(n) = +1$: left taller by one
 - $\text{balance}(n) = -1$: right taller by one
- $|\text{balance}(n)| > 1$: node is out of balance

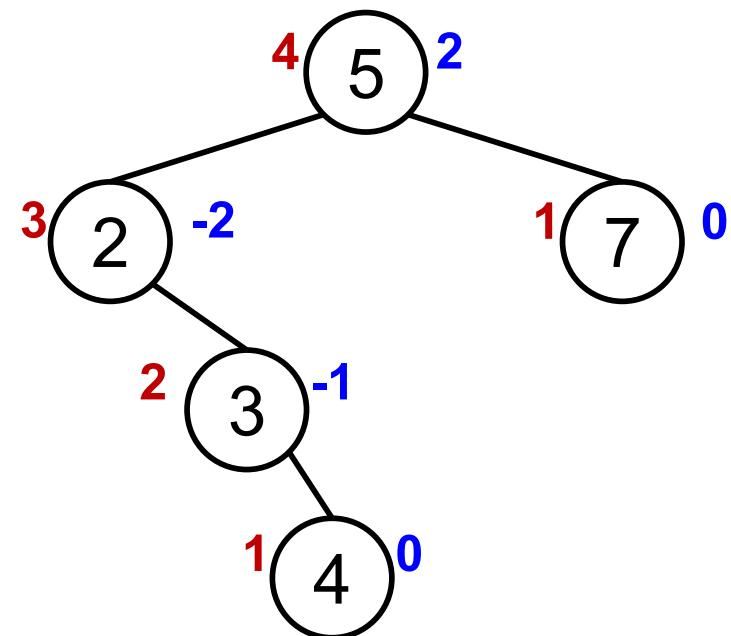
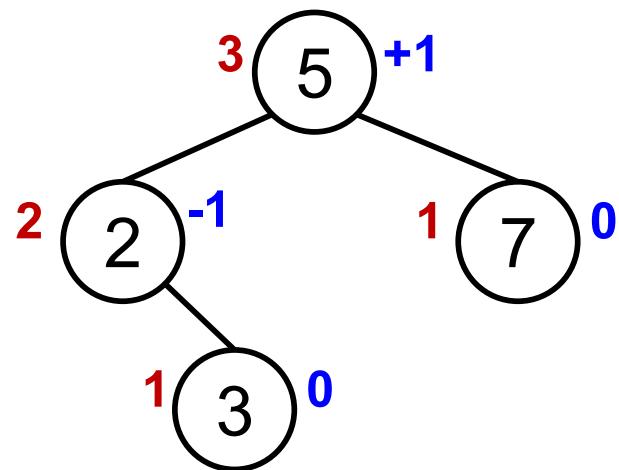
Balance Factor Example

- What is the **height** for each node?

$$\text{height}(\text{node}) = \max(\text{height}(\text{children})) + 1;$$

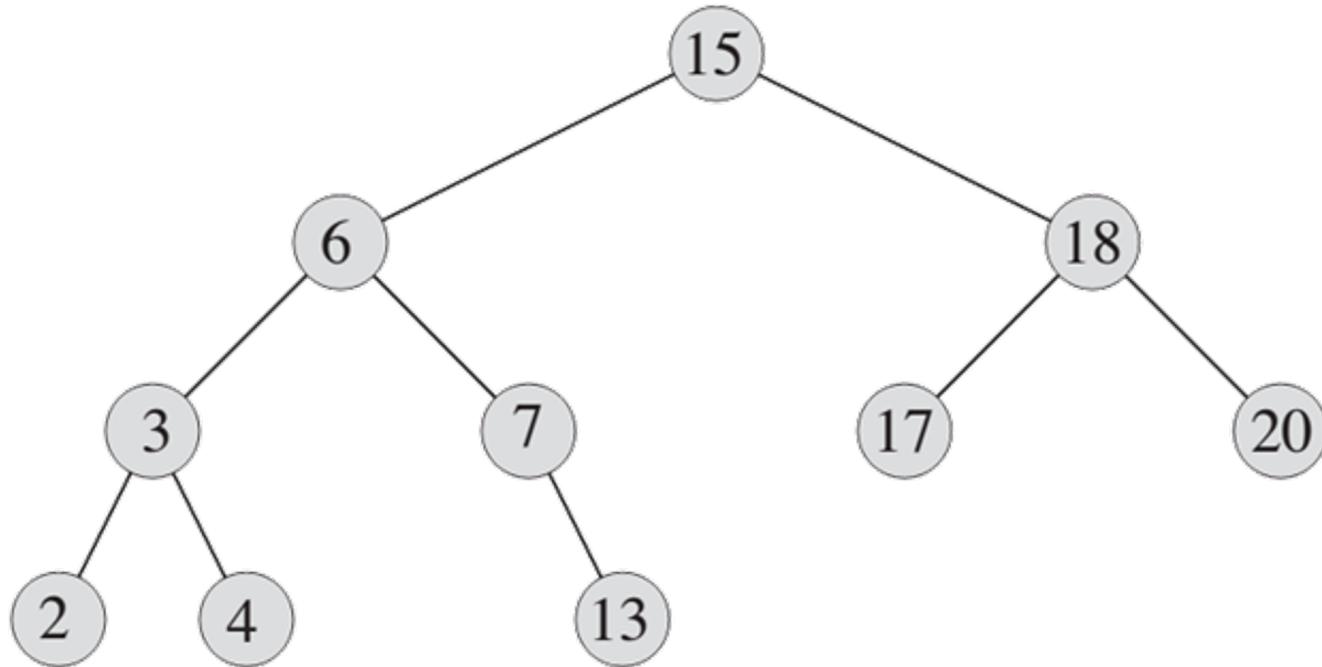
- What is the **balance factor**?

$$\text{balance}(n) = \text{height}(n\rightarrow\text{left}) - \text{height}(n\rightarrow\text{right})$$



Balance Factor Exercise

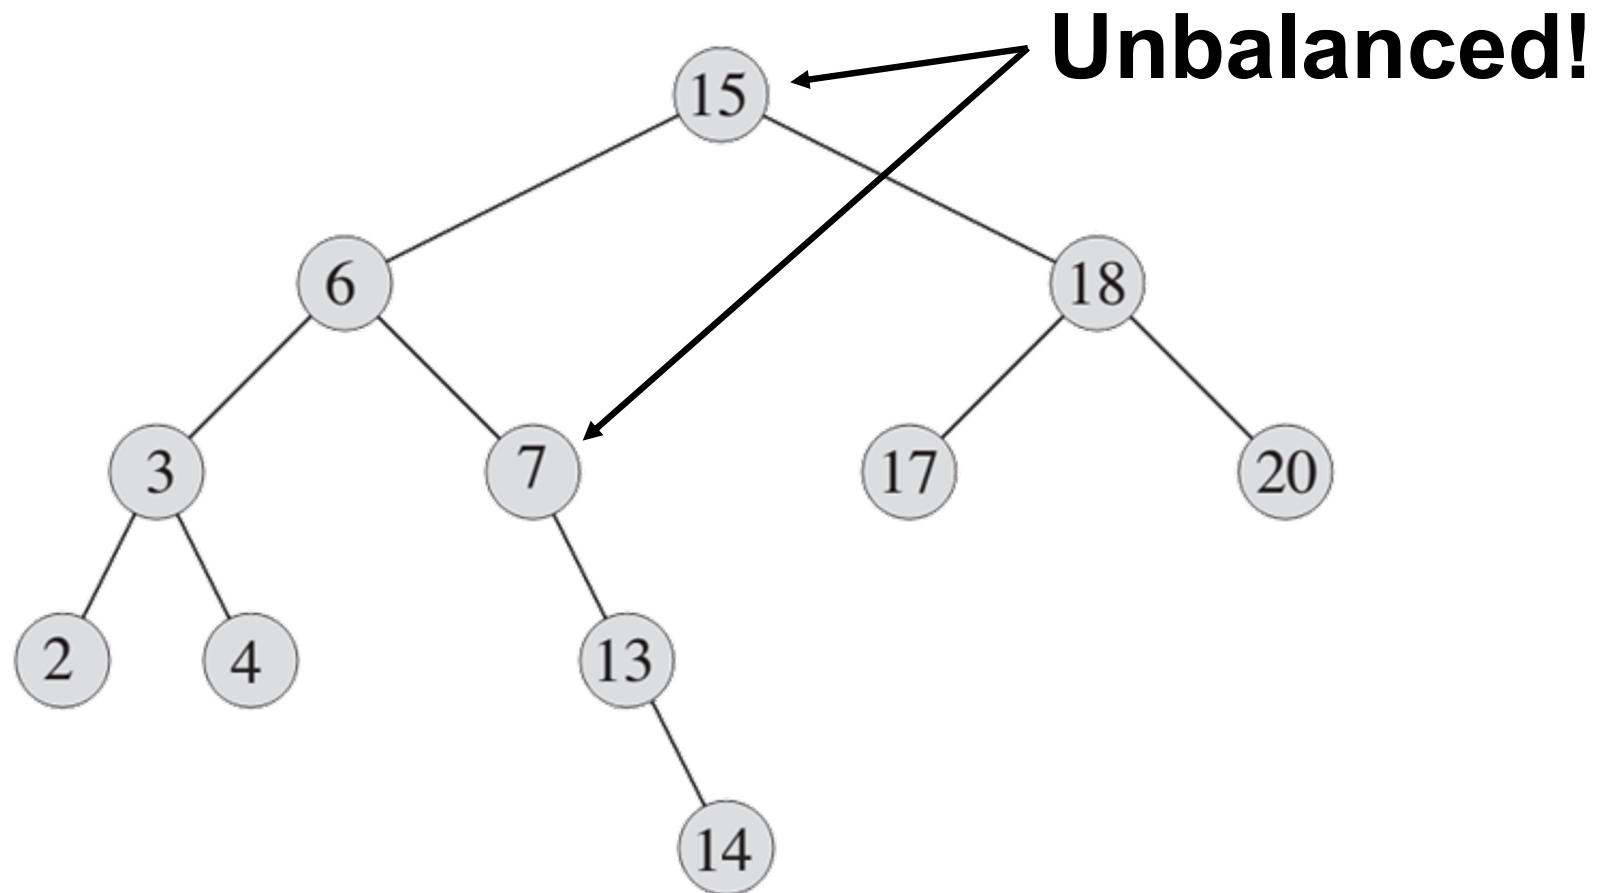
Label the balance factor on each node



$$\text{bal}(n) = \text{height}(n\rightarrow\text{left}) - \text{height}(n\rightarrow\text{right})$$

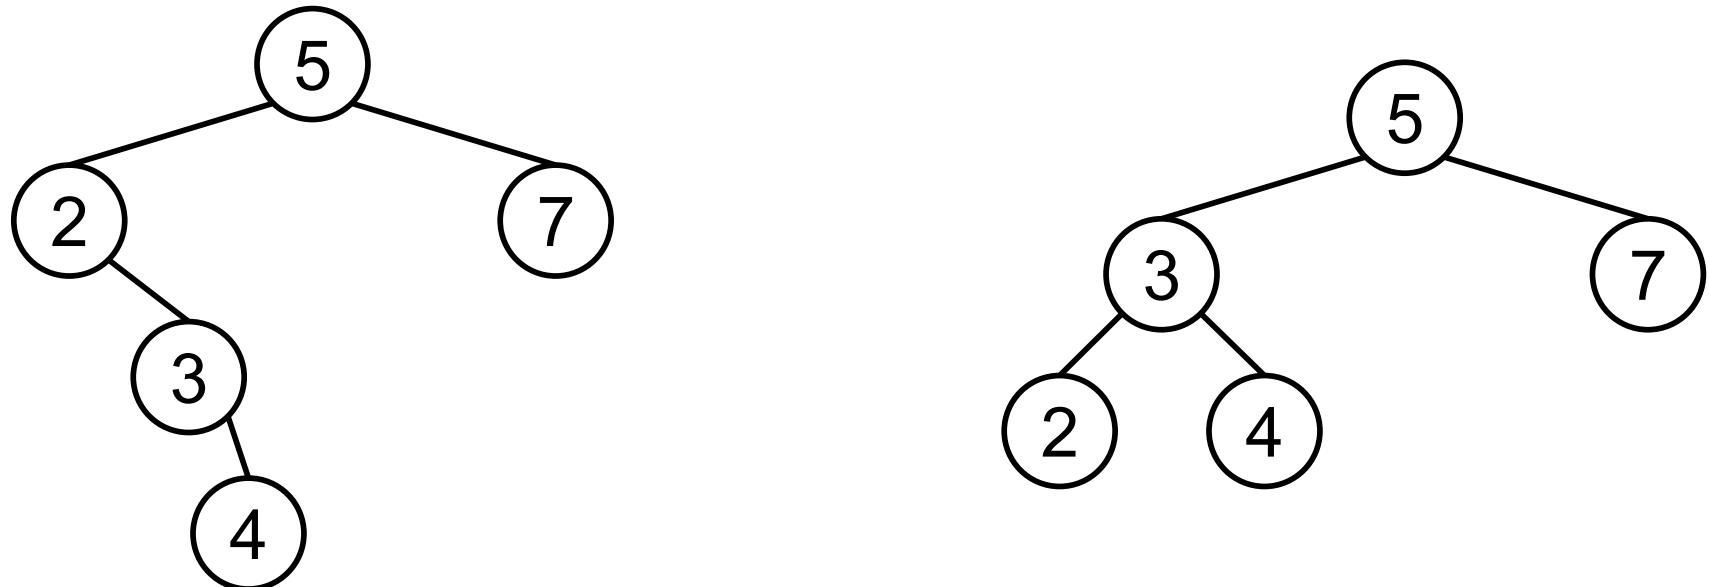
Insert (begins like BST)

```
tree_insert(root, 14);
```



Rotations

- We use *rotations* to rebalance the binary tree
 - Swap a parent and one of its children
 - BUT preserve the BST ordering property



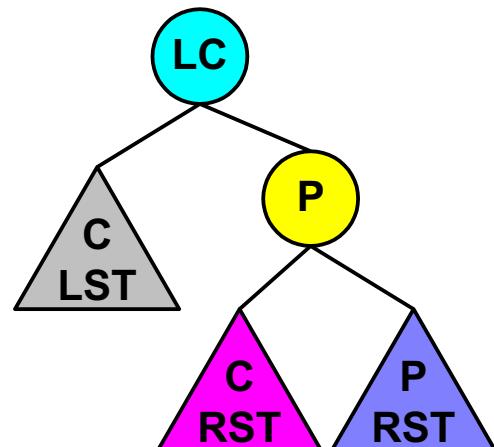
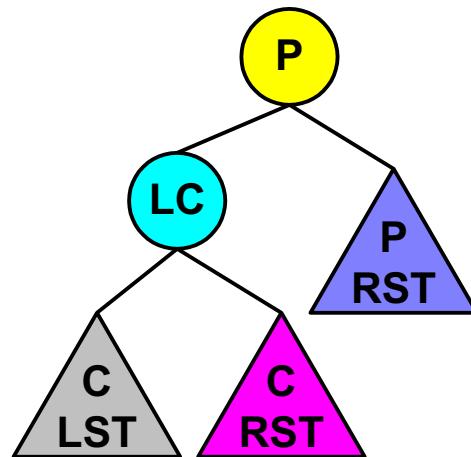
- *Rotation is a local change involving only three pointers and two nodes*

Rotations

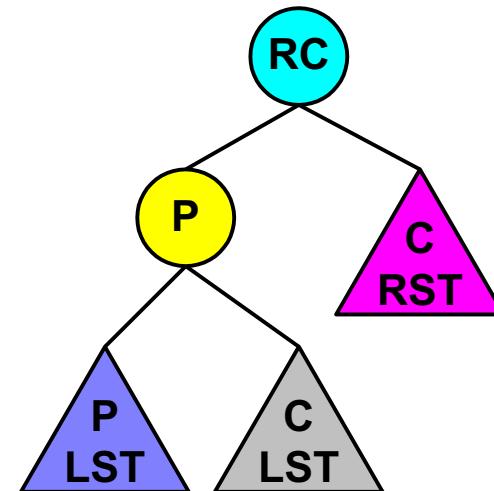
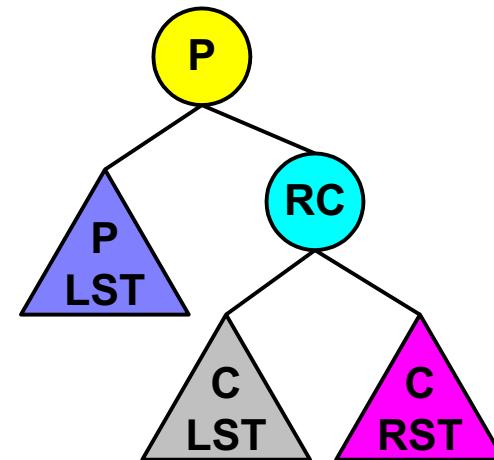
- We use *rotations* to rebalance the binary tree
 - Interchange the role of a parent and one of its children in a tree...
 - Preserve the BST ordering among the keys in the nodes
- The second part is tricky
 - Right rotation: copy the right pointer of the left child to be the left pointer of the old parent
 - Left rotation: copy the left pointer of the right child to be the right pointer of the old parent

Rotations

Rotate Right: RR(P)

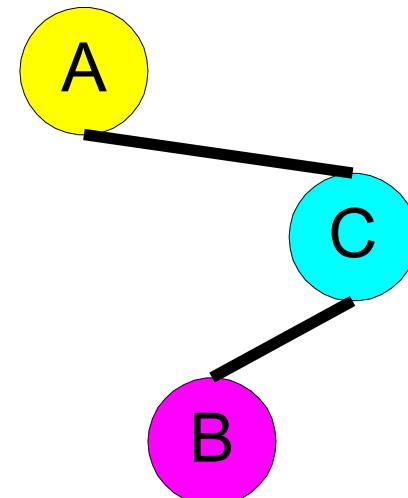
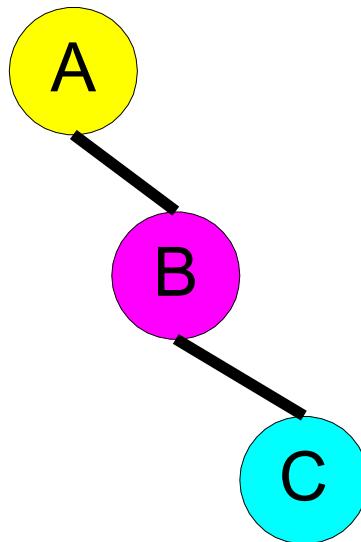


Rotate Left: RL(P)

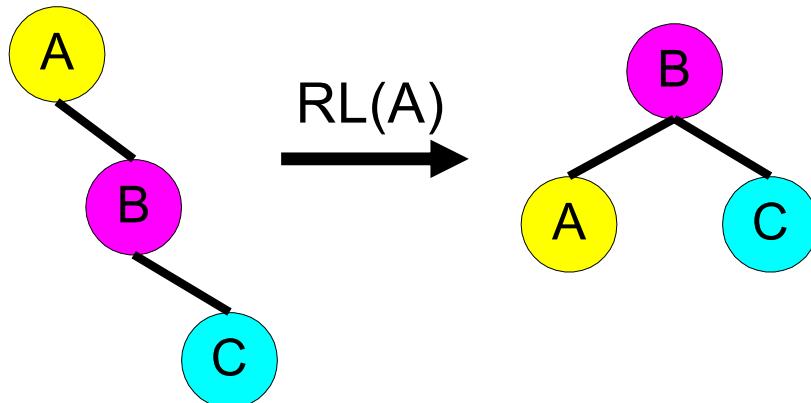


Rotation Exercise

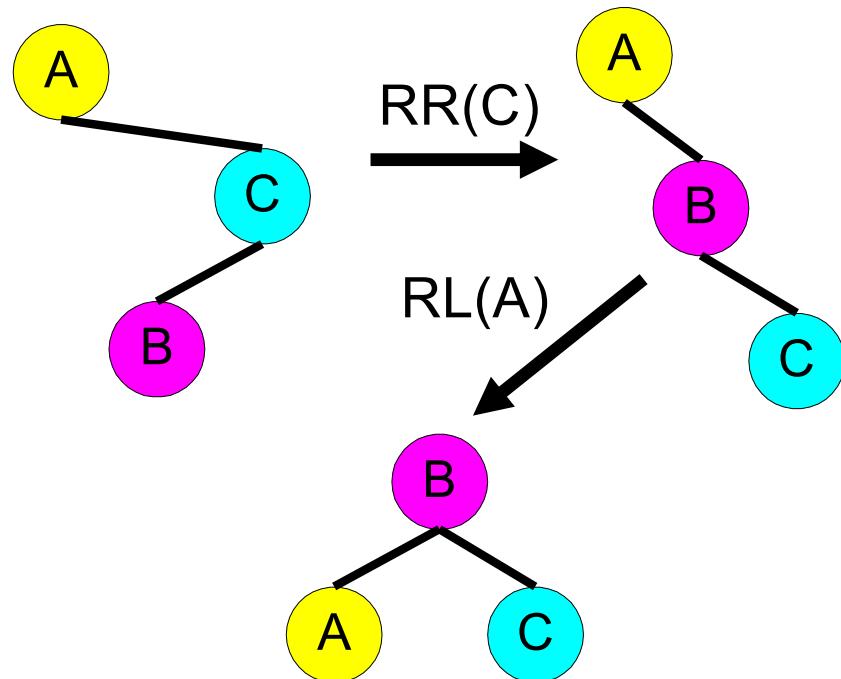
Use rotations to balance these trees



Solution



- First tree is easy:
 $\text{rotate-left}(A)$
- single rotation



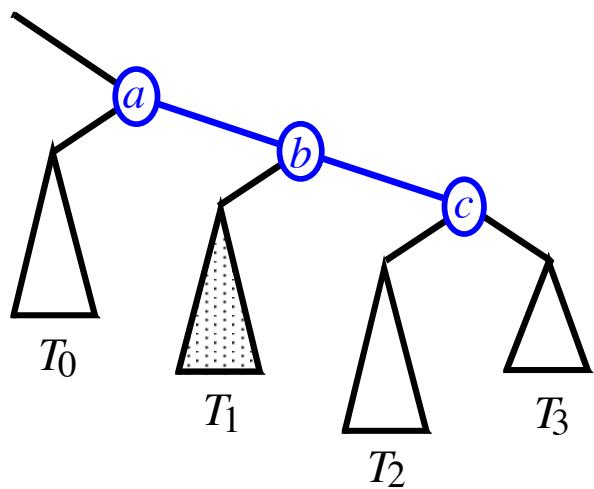
- Second is harder:
 $\text{rotate-right}(C)$
 $\text{rotate-left}(A)$
- double rotation

Insert

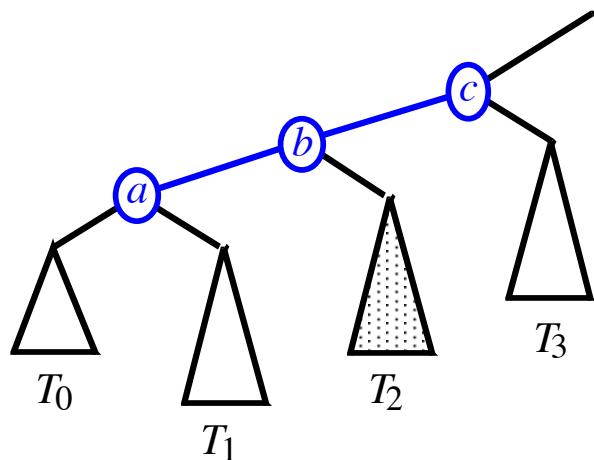
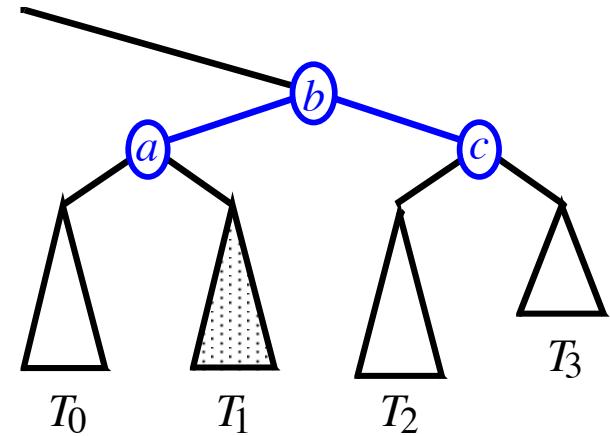
Four Cases

1. Single left rotation
 - RL(a)
2. Single right rotation
 - RR(a)
3. Double rotation right-left
 - a. RR(c)
 - b. RL(a)
4. Double rotation left-right
 - a. RL(a)
 - b. RR(c)

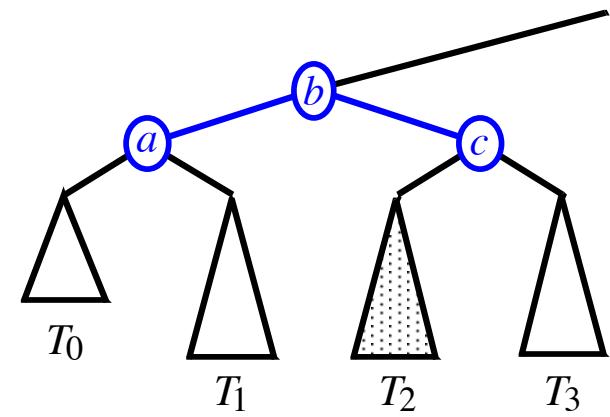
Single Rotations



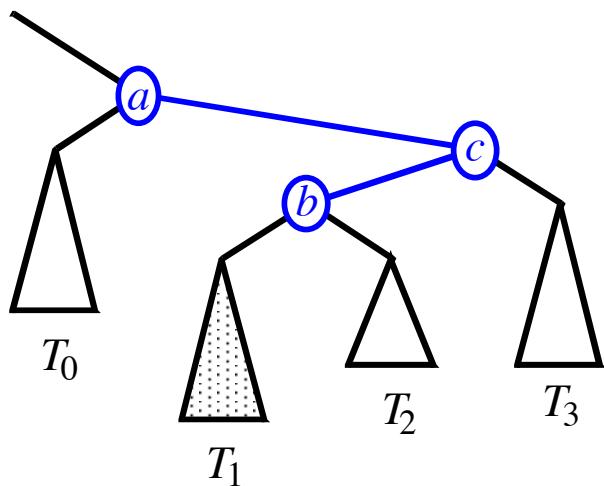
single rotation
→
RL(a)



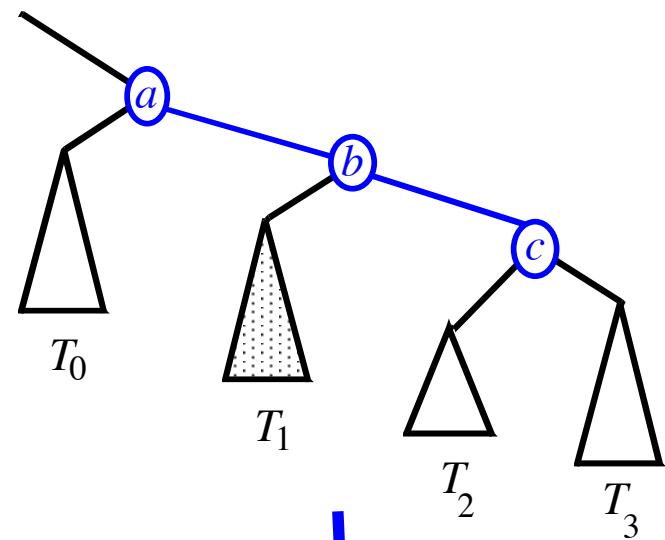
single rotation
→
RR(c)



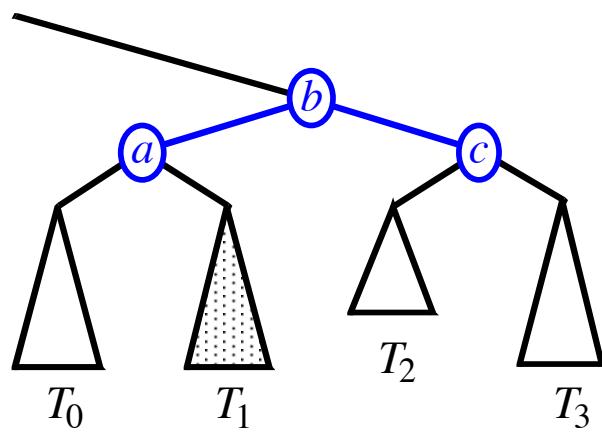
Double Rotations



$\textcolor{blue}{RR(c)}$

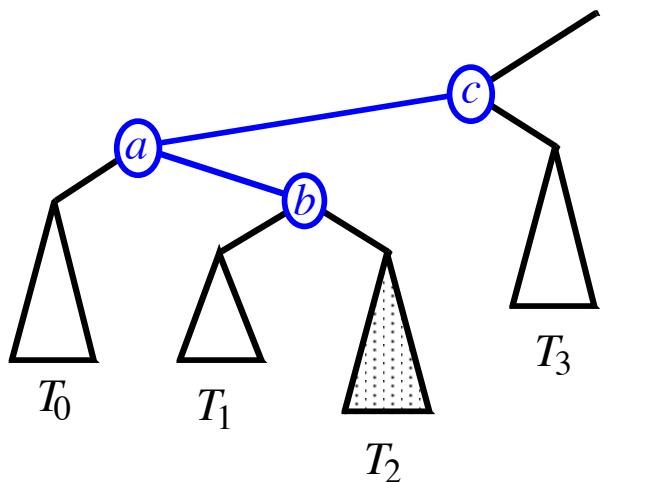


$\textcolor{blue}{RL(a)}$



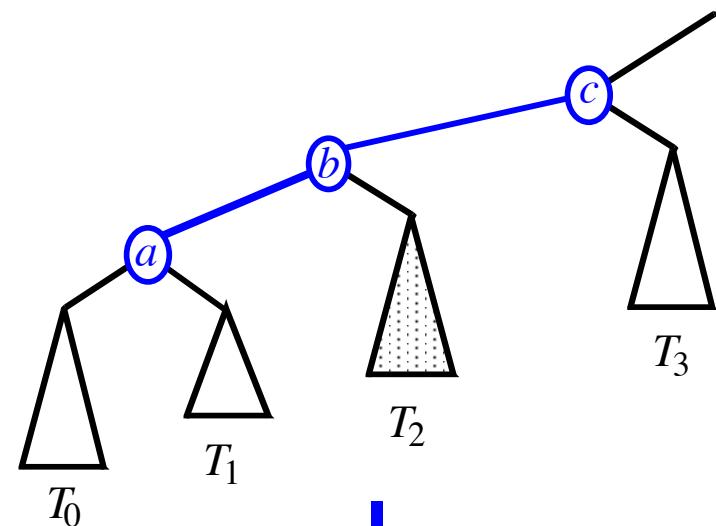
*double rotation
is needed*

Double Rotations



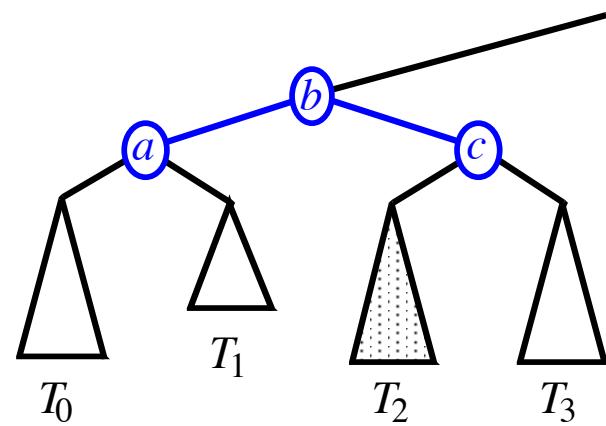
\rightarrow

$RL(a)$



\downarrow

$RR(c)$



*double rotation
is needed*

Left Left Case	Right Right Case	Left Right Case	Right Left Case
<p>Right Rotation</p>	<p>Left Rotation</p>	<p>Left Rotation</p>	<p>Right Rotation</p>

http://upload.wikimedia.org/wikipedia/commons/c/c4/Tree_Rebalancing.gif

Checking and Balancing

```
Algorithm checkAndBalance(Node *n)
    if balance(n) > +1
        if balance(n->left) < 0
            rotateL(n->left)
            rotateR(n)
        else if balance(n) < -1
            if balance(n->right) > 0
                rotateR(n->right)
            rotateL(n)
```

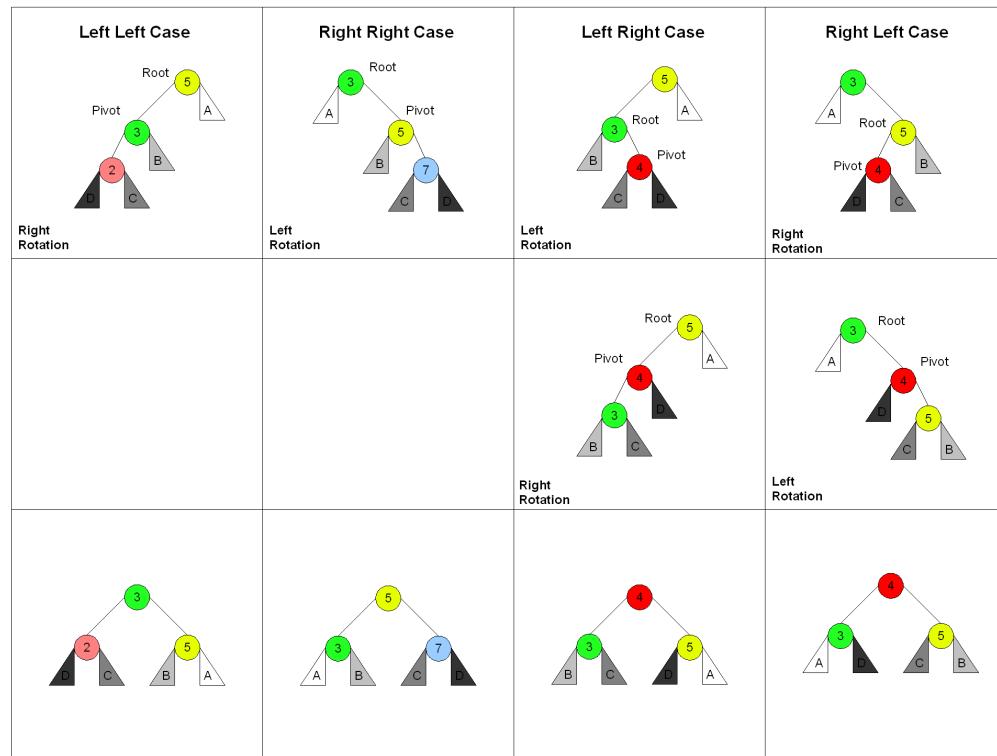
- Outermost if:
Q: Is node out of balance?
A: $> +1$: left too big
 < -1 : right too big
- Inner ifs:
Q: Do we need a double rotation?
A: Only if signs disagree

Checking and Balancing

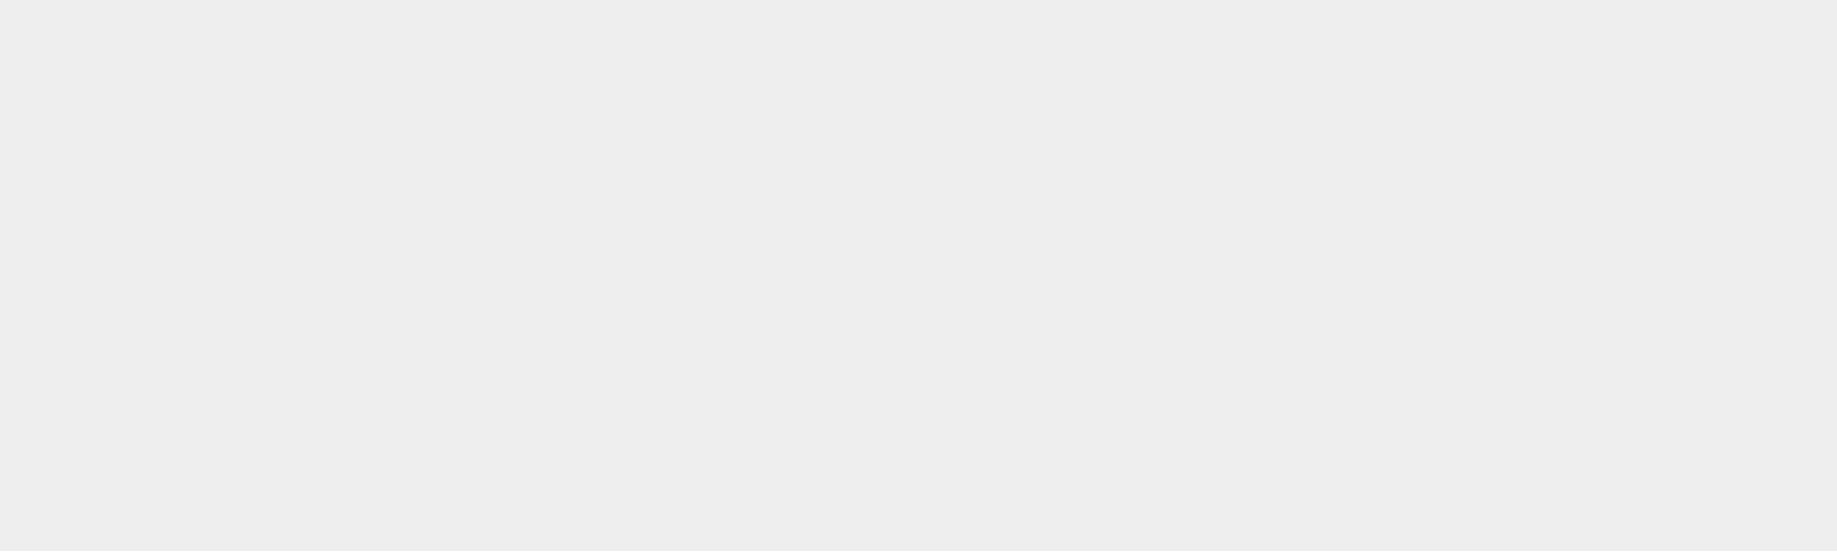
- As insert finishes, after every recursive call, update height of current node, then call checkAndBalance() on every node along the insertion path
- What is the time complexity of this?
- How many “fixes” are needed after insert?

Rotation Exercise

- Insert these keys into an AVL tree, rebalancing when necessary
- 3, 2, 1, 4, 5, 6, 7, 16, 15, 14



Rotation Exercise

- 
- Insert 3... 2... 1...
 - Unbalanced: RR(3)
 - Insert 4... 5...
 - Unbalanced: RL(3)
 - Insert 6...
 - Unbalanced: RL(2)
 - Insert 7...
 - Unbalanced: RL(5)
 - Insert 16... 15...
 - Unbalanced: RR(16), RL(7)
 - Insert 14...
 - Unbalanced: RR(15), RL(6)

Animation screen capture from visualgo.net

AVL Tree Remove

1. Remove like a BST

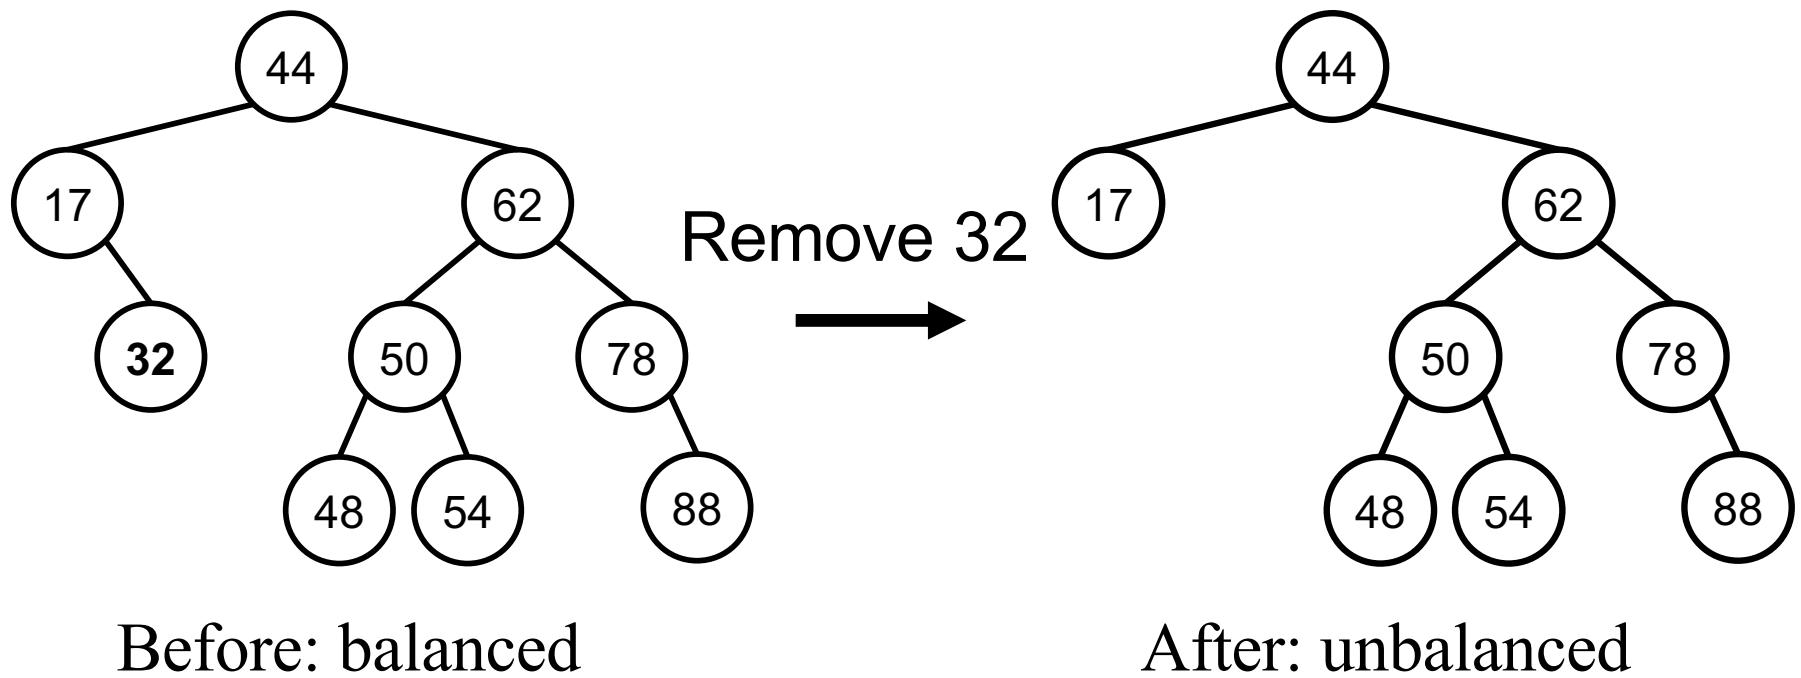
- Key observation: **All keys in LHS \leq all in keys RHS**
- Rearrange RHS so that its smallest node is its root
 - Must be some such node, since RHS is not empty
 - New RHS root has a right child, but no left child
- Make the RHS root's left child the LHS root

2. Rearrange tree to balance height

- Travel up the tree from the parent of removed node
- At every unbalanced node encountered, rotate as needed
- This restructuring may unbalance multiple ancestors, so continue check and rebalance up to the root

AVL Tree Remove

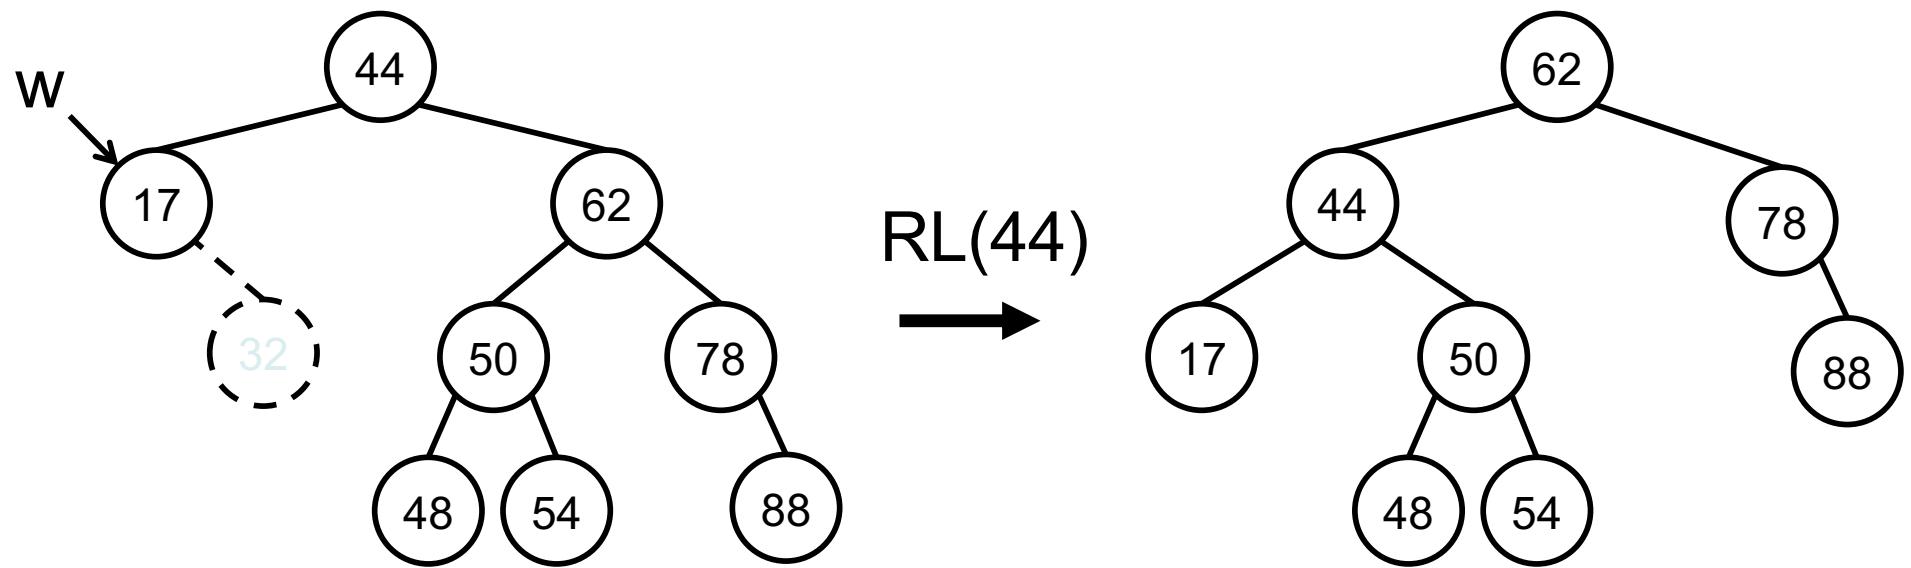
- Remove as in a binary search tree
 - Rebalance if an imbalance has been created



Rebalance Required!

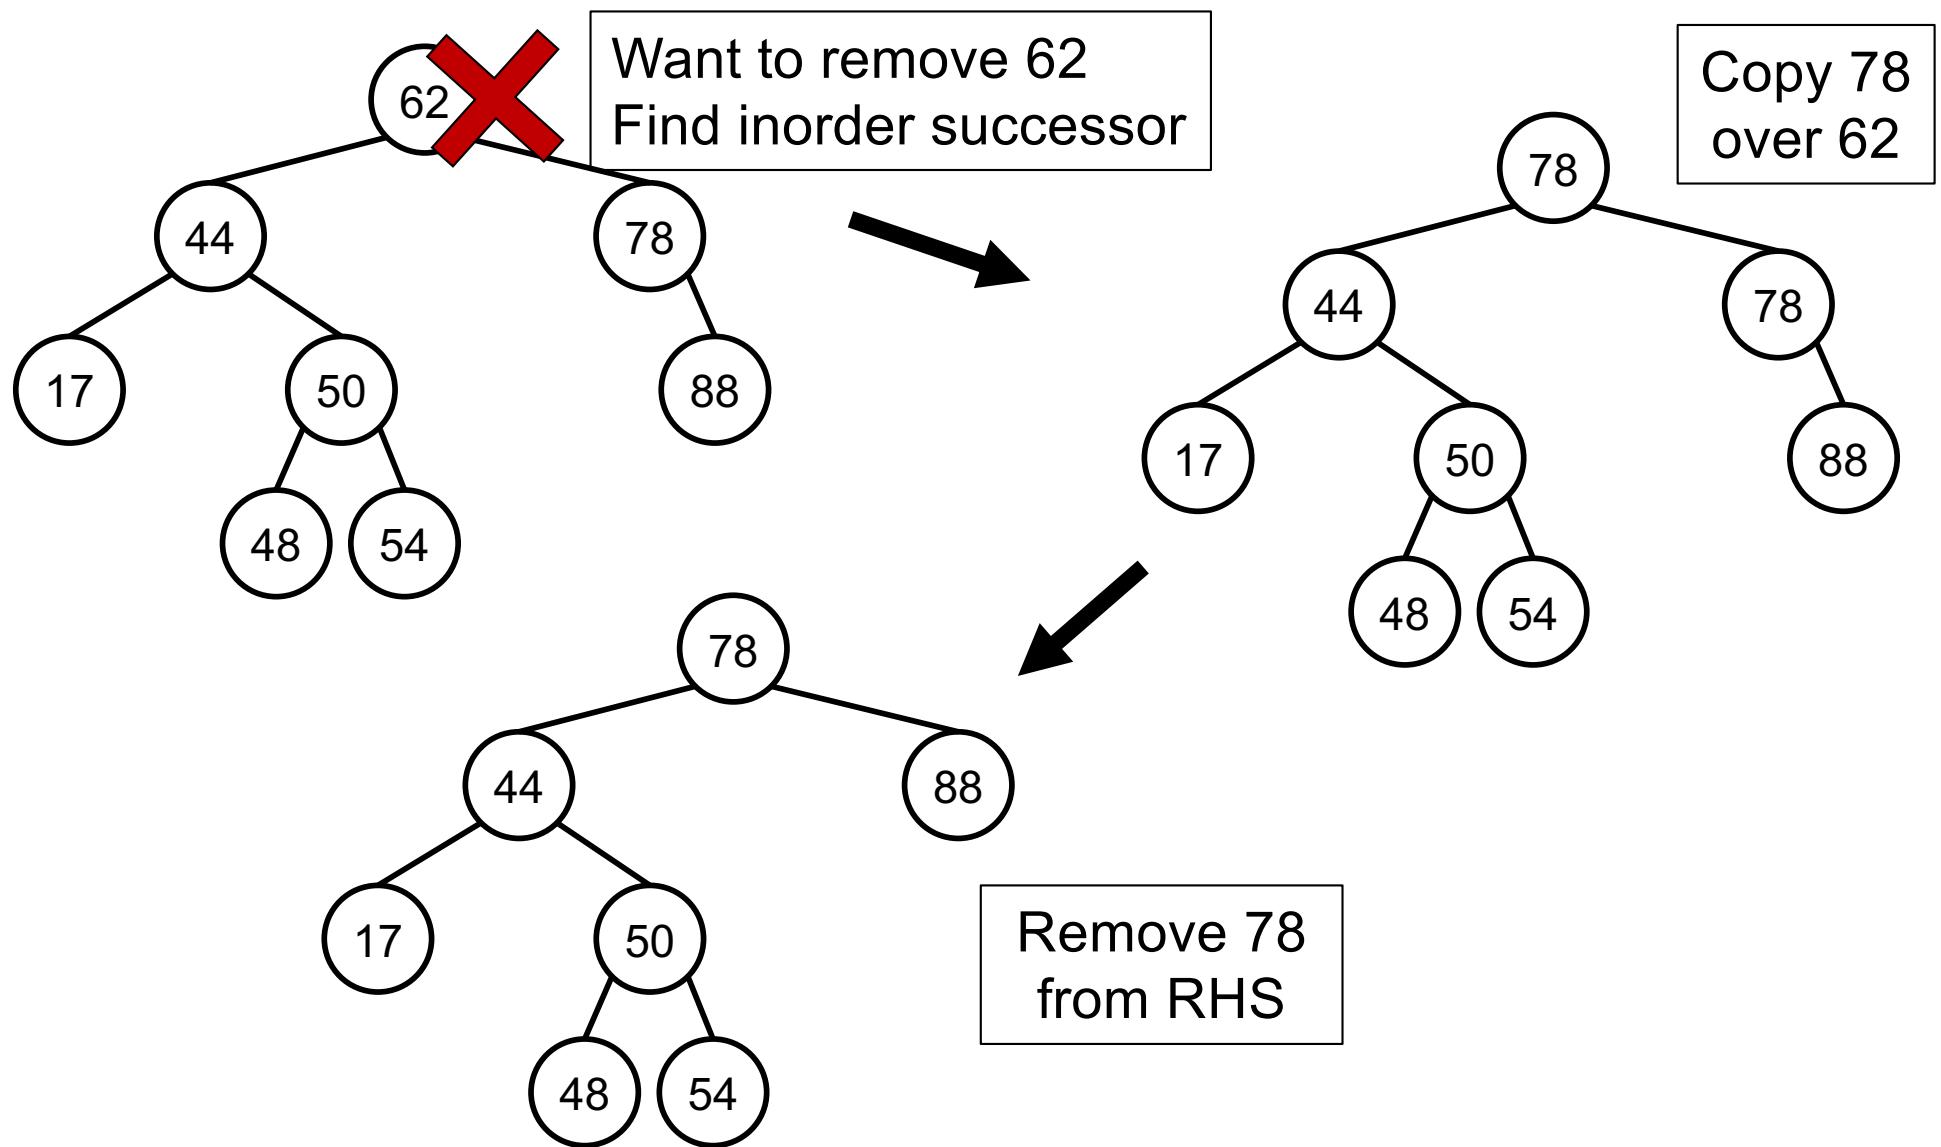
Rebalancing after a Remove

- Travel up the tree from w , the parent of the removed node
- At every unbalanced node encountered, rotate as needed
- This restructuring may unbalance multiple ancestors, so continue checking and rebalancing **up to the root**



How many “fixes” are needed after remove?

Rebalancing after a Remove



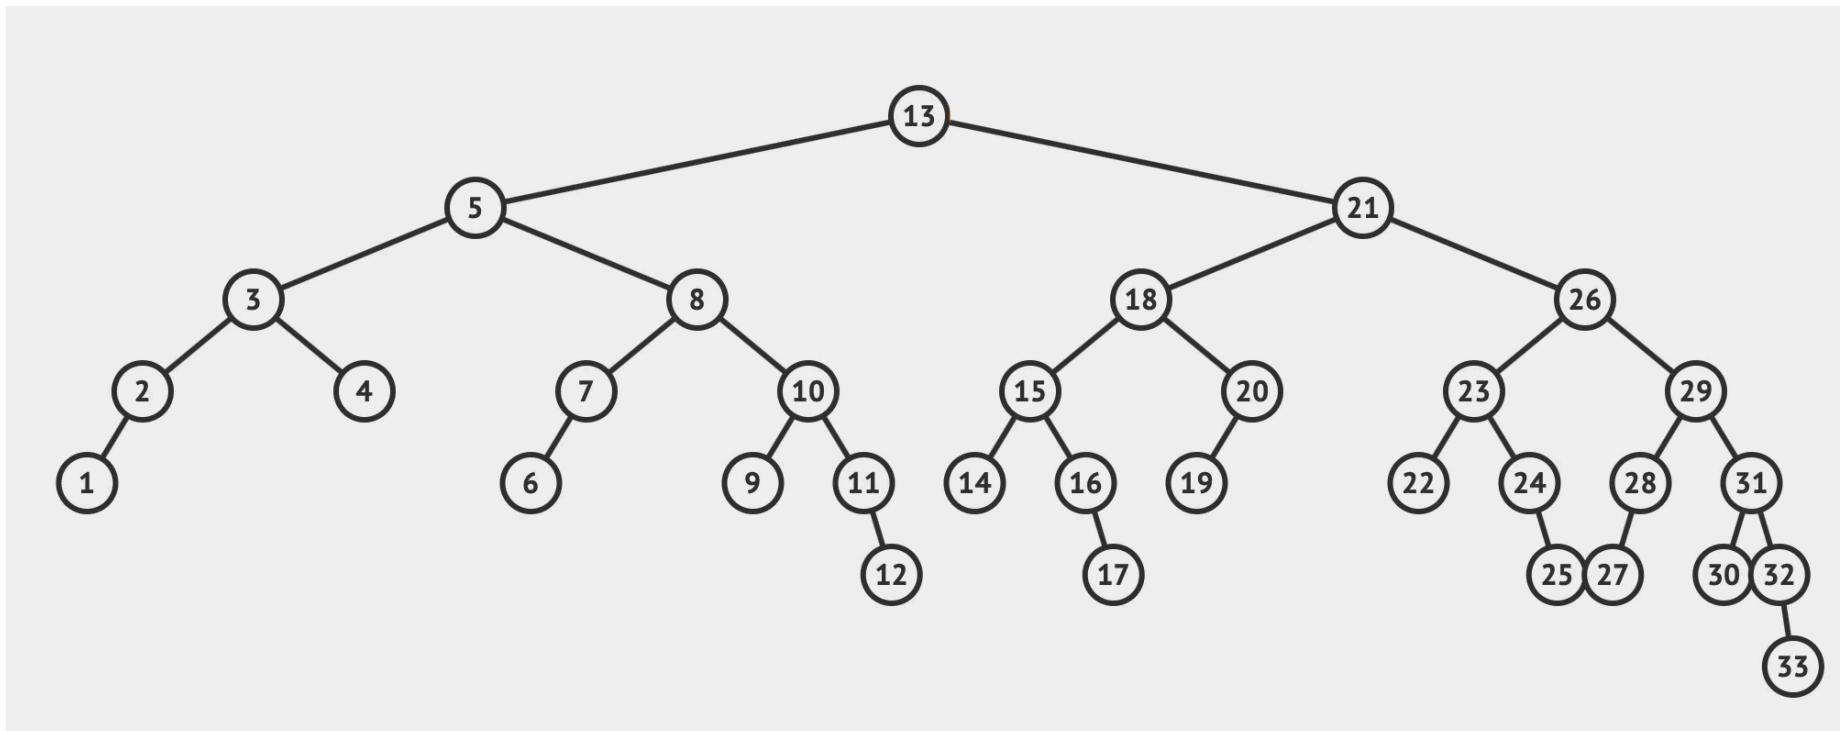
Rebalancing after a Remove The diagram illustrates the rebalancing of a binary search tree after a remove operation. It shows three stages of the tree's state: - Initial State:** A root node with value 78. Its left child is 44, and its right child is 88. Node 44 has left child 17 and right child 50. Node 50 has left child 48 and right child 54. - After Removal:** A root node with value 78. Its left child is 50, and its right child is 88. Node 50 has left child 44 and right child 54. Node 44 has left child 17 and right child 48. - Final State:** A root node with value 50. Its left child is 44, and its right child is 78. Node 44 has left child 17 and right child 48. Node 78 has left child 54 and right child 88. Annotations provide information about the rotations: - A box labeled "First: RL(44)" with an arrow points to the transition from the initial state to the intermediate state. - A box labeled "Second: RR(78)" with a downward arrow points to the transition from the intermediate state to the final state. - A bracket on the left groups two boxes: "78: +2 Bal" and "44: -1 Bal". - A bracket on the left groups the "78: +2 Bal" and "44: -1 Bal" boxes with the text "Need double rotation: RL(44), RR(78)". 99

Rebalancing: 1 or More?

- A “fix” (single or double rotation) shortens a subtree that is too tall
- When inserting, the new node is the source of any imbalance, and a single fix will counteract it and repair the entire tree
- When removing, a deleted node can only create an imbalance by making subtrees shorter
- If a shortened subtree was already shorter than its sibling, a fix is needed at a higher level, so multiple fixes may be required

Rebalancing: 1 or More?

- Values 1-33 inserted into an AVL in a particular order
- Resulting tree is balanced
- Multiple fixes required when 4 is removed



Animation screen capture from visualgo.net

Useful Website

- <https://visualgo.net/en/bst>
- Close the help
- Click on “AVL TREE” near the top
- You can insert/remove several nodes at once, or one at a time, slow down or speed up the demo

AVL Trees

Data Structures & Algorithms

Summary

- Binary Search Tree
 - Worst-case insert or search is $O(n)$
- AVL Tree
 - Worst-case insert or search is $O(\log n)$
 - Must guarantee **height balance property**
- Operations
 - Search: $O(\log n)$ (same algorithm as BST, but faster)
 - Insert: $O(\log n)$ (Starts like BST, then may rebalance)
 - Remove: $O(\log n)$ (Starts like BST, then may rebalance)
 - Sort: $O(n \log n)$ to build the tree, $O(n)$ to do inorder traversal
- Rotation:
 - $O(1)$ cost for single rotation, $O(1)$ cost for double rotation