

EECS 370 - Lecture 6

Function Calls



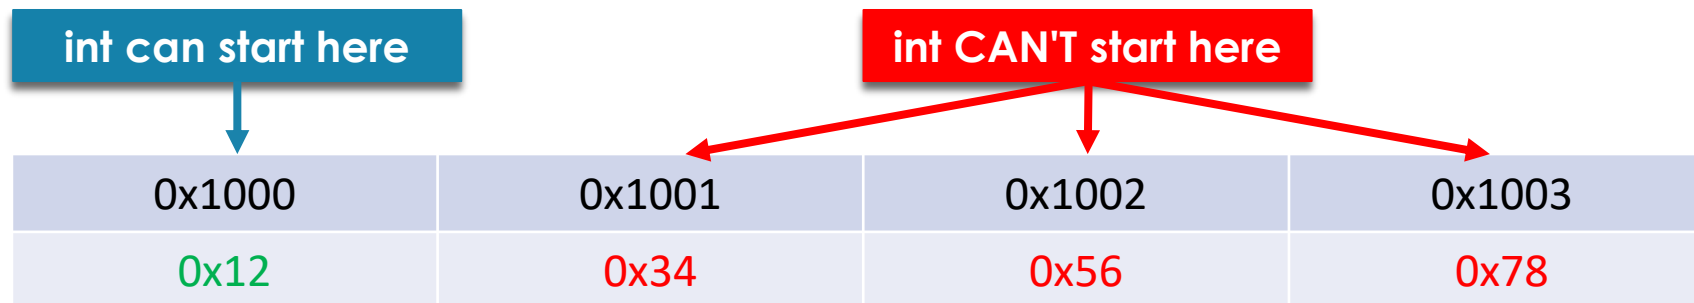
Announcements

- Project 1a due tonight
- Project 1s+m due next Thursday
- Let us know about exam conflicts in the **next week**
 - Form on website




Reminder: Memory Alignment

- Most modern ISAs require that data be aligned
 - An N-byte variable must start at an address A, such that $(A\%N) == 0$
- For example, starting address of a 32 bit **int** must be divisible by 4



- Starting address of a 16 bit **short** must be divisible by 2



Golden Rule of Alignment

```
char  c;  
short s;  
int   i;
```

- Every (primitive) object starts at an address divisible by its size
- "Padding" is placed in between objects if needed

0x1000	0x1001	0x1002	0x1003	0x1004	0x1005	0x1006	0x1007
[c]	[padding]	[s]		[i]			


- But what about non-primitive data types?
 - Arrays? Treat as independent objects
 - Structs? Trickier...

Problem with Structs

- If we align each element of a struct according to the Golden Rule, we can still run into issues
 - E.g.: An array of structs

```
char c;  
  
struct {  
    char c;  
    int i;  
} s[2];
```

Amount of padding
is different across
different instances



1000	1001	1002	1003	1004	1005	1006	1007	1008	1009	100A	100B	100C	100D	100E	100F
c	s[0].c	[pad]	[pad]	s[0].i				s[1].c	[pad]	[pad]	[pad]	s[1].i			

- Why is this bad?
- It makes "for" loops very difficult to write!
 - Offsets need to be different on each iteration

Structure Alignment

- Solution: in addition to laying out each field according to Golden Rule...
 - Identify largest (primitive) field
 - Starting address of overall struct is aligned based on the largest field
 - Padded in the back so total size is a multiple of the largest primitive

```
char c;  
  
struct {  
    char c;  
    int i;  
} s[2];
```

1000	1001	1002	1003	1004	1005	1006	1007	1008	1009	100A	100B	100C	100D	100E	100F
c	[pad]	[pad]	[pad]	s[0].c	[pad]	[pad]	[pad]	s[0].i				s[1].c	[pad]	[pad]	[pad]

Guaranteed to lay
out each instance
identically

Structure Example

```
struct {  
    char w;  
    int x[3];  
    char y;  
    short z;  
}
```

Poll: What boundary should this struct be aligned to?

- a) 1 byte
- b) 4 bytes
- c) 12 bytes
- d) 2 bytes
- e) 19 bytes

- Assume struct starts at location 1000,
 - char w → 1000
 - x[0] → 1004-1007, x[1] → 1008 – 1011, x[2] → 1012 – 1015
 - char y → 1016
 - short z → 1018 – 1019

Total size = 20 bytes!

Calculating Load/Store Addresses for Variables

Datatype	size (bytes)
char	1
short	2
int	4
double	8

```
short a[100];  
char b;  
int c;  
double d;  
short e;  
struct {  
    char f;  
    int g[1];  
    char h;  
} i;
```

- *Problem:* Assume data memory starts at address 100 and no reordering, calculate the total amount of memory needed

a = 200 bytes (100-299)

b = 1 byte (300-300)

c = 4 bytes (304-307)

d = 8 bytes (312-319)

e = 2 bytes (320-321)

struct: largest field is 4 bytes, start at 324

f = 1 byte (324-324)

g = 4 bytes (328-331)

h = 1 byte (332-332)

i = 12 bytes (324-335)

236 bytes total!! (compared to 221, originally)

Data Layout – Why?

- Does gcc (or another compiler) reorder variables in memory to avoid padding?
- Only outside structs
- C99 forbids reordering elements inside a struct
- The programmer (i.e., you) are expected to manage data layout of variables for your program and structs.
- Two optimal strategies:
 - Order fields in struct by datatype size, smallest first
 - Or by largest first

Instruction Set Architecture (ISA) Design Lectures

- Lecture 2: ISA - storage types, binary and addressing modes
- Lecture 3 : LC2K
- Lecture 4 : ARM
- Lecture 5 : Converting C to assembly – basic blocks
- **Lecture 6 : Converting C to assembly – functions**
- Lecture 7 : Translation software; libraries, memory layout



Agenda

- Using branches more generally
- **Function calls and the call stack**
- Assigning variables to memory locations
- Saving registers
- Caller/callee example

Implementing Functions

Poll: What's wrong
with this?

- Does this assembly code do what we need?

```
int mult_2(int x) {  
    int temp = x*2;  
    return temp;  
}  
  
int GLOBAL = 6;  
  
int main() {  
    int result = mult_2(GLOBAL+1);  
    printf(result);  
}
```

```
LDURSW X1, [XZR, GLOBAL]  
ADD     X2, X1, #1           // Inc GLOBAL  
STURW   X2, [XZR, X]         // Pass arg  
B       MULT_2               // Execute func  
  
RETURN:  
LDURSW X3, [XZR, TEMP]       // load result  
STURW   X3, [XZR, STRING]    // Pass arg  
B       PRINTF               // Execute func  
  
...  
MULT_2:  
LDURSW X1, [XZR, X]           // load arg  
ADD     X2, X1, X1            // mult by 2  
STURW   X2, [XZR, TEMP]      // return result  
B       RETURN               // return
```

Problem 1: Returning from Functions

- Branches so far have hard-coded destination

```
B.NE L1
ADD X1, X1, #1
B L2
L1: ADD X2, X2, #1
L2: ...
```

```
B.NE 3
ADD X1, X1, #1
B 2
ADD X2, X2, #1
```

- This is fine for if-statements, for-loops etc
- But functions can be called from multiple places
 - Meaning we'll return to different spots on each func call! Can't hardcode offset!

```
int func(int x) {
    printf(x * 10);
    return;
}
int helper() {
    func(7);
}
int main() {
    helper();
    func(13);
}
```

Should this return to
"helper" or "main"?

Solution: Indirect Jumps

- Indirect branches or "jumps" don't hardcode destination in instruction
- They index a register whose value holds destination

Unconditional branch	branch	B	2500	go to PC + 10000	Branch to target address; PC-relative
	branch to register	BR	X30	go to X30	For switch, procedure return
	branch with link	BL	2500	X30 = PC + 4; PC + 10000	For procedure call PC-relative

- Use "BL" to **call a function**
 - Destination is hardcoded
 - PC +4 (return address) stored in X30
- Use "BR" to **return from a function**
 - X30 is read for return address
 - Allows us to return to different places

Solution: Indirect Jumps

```
int mult_2(int x) {  
    int temp = x*2;  
    return temp;  
}
```

```
int GLOBAL = 6;
```

```
int main() {  
    int result = mult_2(GLOBAL+1);  
    printf(result);  
}
```

Also don't
need "return"
labels

Now MULT_2
can return to
whatever
function
called it

```
LDURSW X1, [XZR, GLOBAL]  
ADD     X2, X1, #1           // Inc GLOBAL  
STURW   X2, [XZR, X]         // Pass arg  
BL      MULT_2               // Execute func  
  
RETURN:  
LDURSW X3, [XZR, TEMP]       // load result  
STURW   X3, [XZR, STRING]    // Pass arg  
BL      PRINTF               // Execute func  
  
...  
MULT_2:  
LDURSW X1, [XZR, X]           // load arg  
ADD     X2, X1, X1            // mult by 2  
STURW   X2, [XZR, TEMP]       // return result  
BR      RETURN               // return
```

Problem 2: Passing Parameters

For any recursive functions, global variables will be overwritten

```
int mult_2(int x) {  
    int temp = x*2;  
    return temp;  
}  
  
int GLOBAL = 6;  
  
int main() {  
    int result = mult_2(GLOBAL+1);  
    printf(result);  
}
```

```
LDURSW X1, [XZR, GLOBAL]  
ADD     X2, X1, #1           // Inc GLOBAL  
STURW   X2, [XZR, X]         // Pass arg  
BL      MULT_2               // Execute func  
LDURSW  X3, [XZR, TEMP]      // load result  
STURW   X3, [XZR, STRING]    // Pass arg  
BL      PRINTF               // Execute func  
...  
MULT_2:  
LDURSW  X1, [XZR, X]         // load arg  
ADD     X2, X1, X1           // mult by 2  
STURW   X2, [XZR, TEMP]      // return result  
BR                               // return
```


Task 1: Passing parameters

- Where should you put all of the parameters?
 - Registers?
 - Fast access but few in number and wrong size for some objects
 - Memory?
 - Good general solution but slow
- ARMv8 solution—and the usual answer:
 - Both
 - Put the first few parameters in registers (if they fit) (X0 – X7)
 - Put the rest in memory on the call stack— **important concept**

Call stack

- ARM conventions (and most other processors) allocate a region of memory for the “call” stack
 - This memory is used to manage all the storage requirements to simulate function call semantics
 - Parameters (that were not passed through registers)
 - Local variables
 - Temporary storage (when you run out of registers and need somewhere to save a value)
 - Return address
 - Etc.
- Sections of memory on the call stack [**stack frames**] are allocated when you make a function call, and de-allocated when you return from a function


The stack grows as functions are called

FUNCTION CALLS

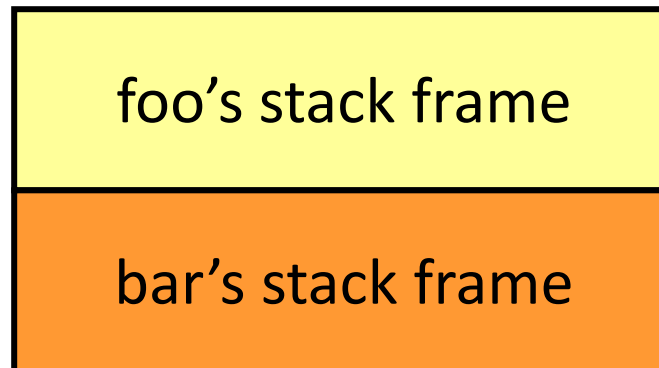
```
void foo()  
{  
    int x, y[2];  
    bar(x);  
}
```

```
void bar(int x)  
{  
    int a[3];  
    printf();  
}
```

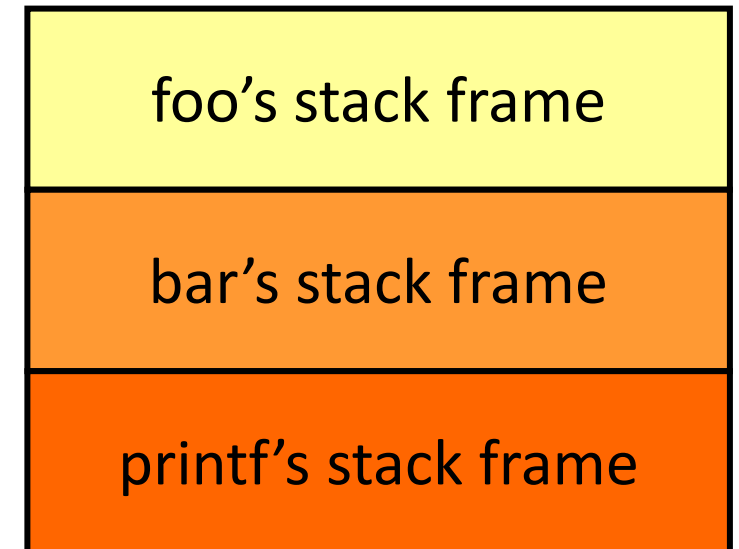
inside foo



foo calls bar



bar calls printf



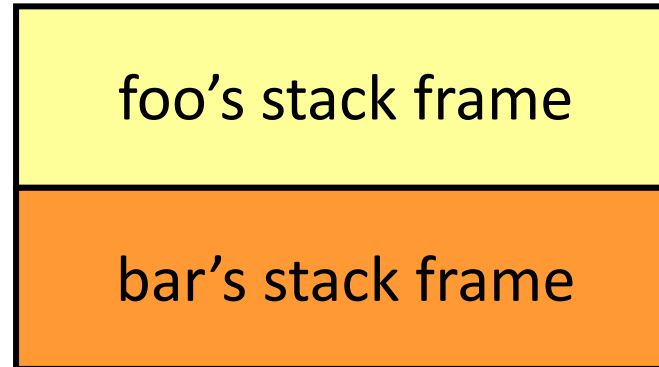
The stack shrinks as functions return

FUNCTION CALLS


```
void foo()  
{  
    int x, y[2];  
    bar(x);  
}
```

```
void bar(int x)  
{  
    int a[3];  
    printf();  
}
```

printf returns



bar returns



Stack frame contents

FUNCTION CALLS

```
void foo()  
{  
    int x, y[2];  
    bar(x);  
}  
  
void bar(int x)  
{  
    int a[3];  
    printf();  
}
```

foo's stack frame

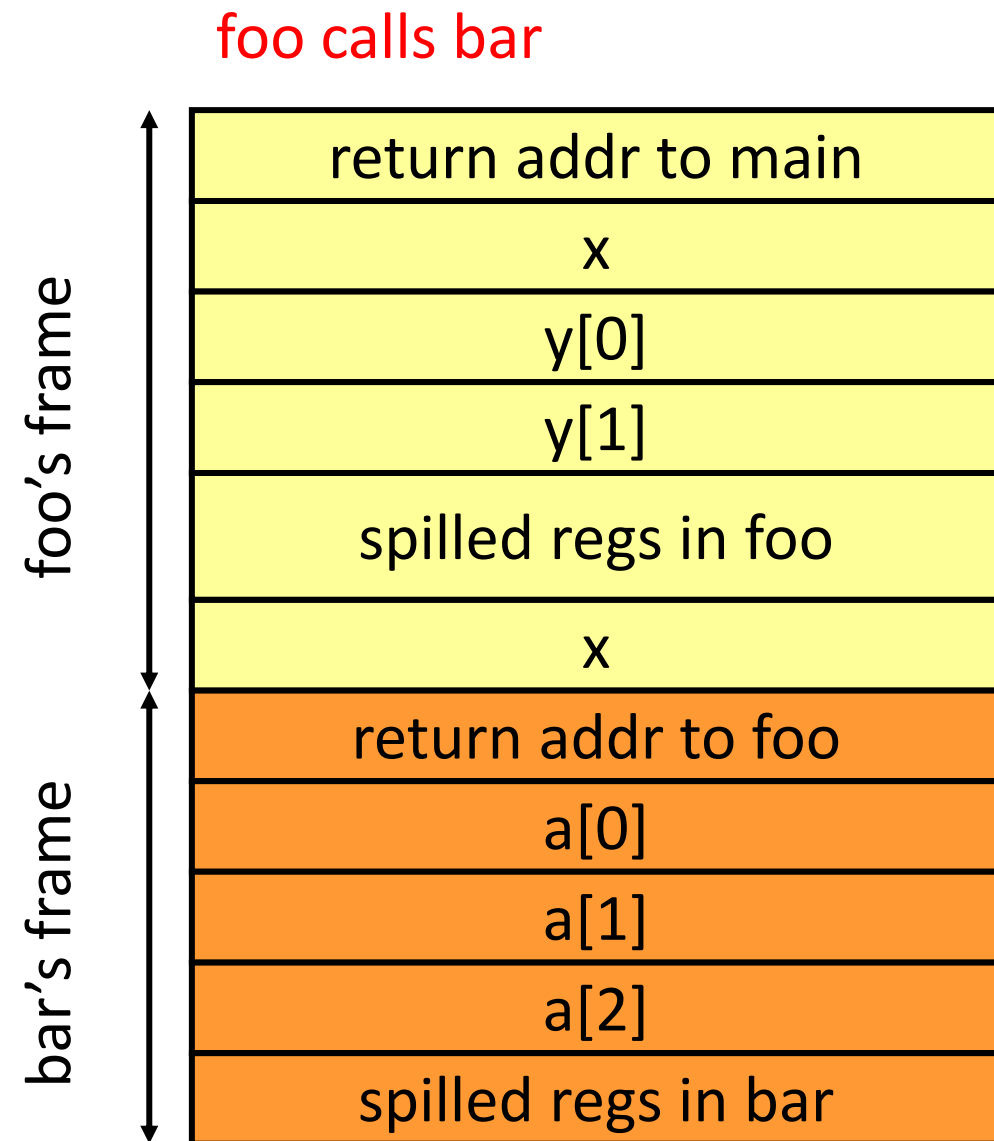
return addr to main
x
y[0]
y[1]
spilled registers in foo

Stack frame contents (2)

```
void foo()
{
    int x, y[2];
    bar(x);
}

void bar(int x)
{
    int a[3];
    printf();
}
```

Spill data—not enough room in x0-x7 for
params and also caller and callee saves



Agenda


- Using branches more generally
- Function calls and the call stack
- **Assigning variables to memory locations**
- Saving registers
- Caller/callee example

Review: Where do the variables go?

Assigning variables to memory spaces

FUNCTION CALLS

```
int w;  
void foo(int x)  
{  
    static int y[4];  
    char* p;  
    p = malloc(10);  
    //...  
    printf("%s\n", p);  
}
```



Assigning variables to memory spaces

FUNCTION CALLS

```
int w;  
void foo(int x)  
{  
    static int y[4];  
    char* p;  
    p = malloc(10);  
    //...  
    printf("%s\n", p);  
}
```

w goes in static, as it's a global

x goes on the stack, as it's a parameter

y goes in static, 1 copy of this!!

p goes on the stack


allocate 10 bytes on heap, ptr

set to the address

string literal "%s\n" goes in static,

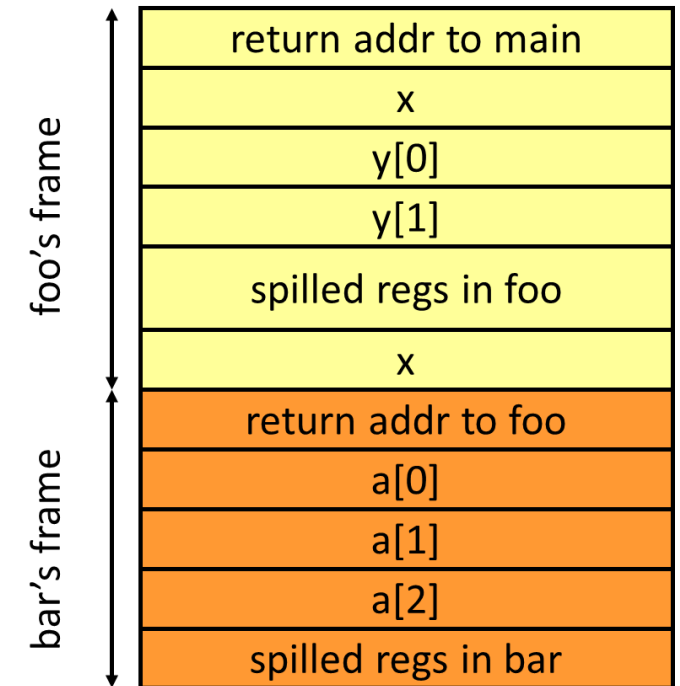
implicit pointer to string on stack, p goes on stack

The addresses of local variables will be different depending on where we are in the call stack



Accessing Local Variables

- Stack pointer (SP):
 - register that keeps track of current top of stack
- Compiler (or assembly writer) knows relative offsets of objects in stack
- Can access using lw/sw offsets

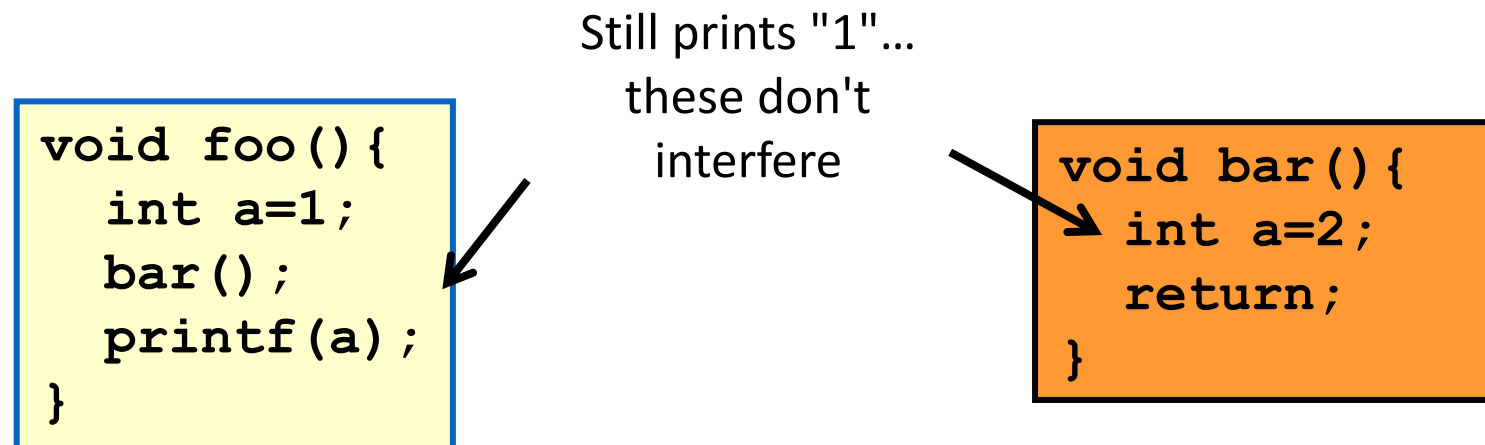


Agenda

- Using branches more generally
- Function calls and the call stack
- Assigning variables to memory locations
- **Saving registers**
- Caller/callee example

Problem 3: Reusing registers

- Higher level languages (like C/C++) provide many abstractions that don't exist at the assembly level
- E.g. in C, each function has its own local variables
 - Even if different function have local variables with the same name, they are independent and guaranteed not to interfere with each other!



What about registers?

- But in assembly, all functions share a small set (e.g. 32) of registers
 - Called functions will overwrite registers needed by calling functions

```
main: movz X0, #1  
      bl  foo  
      bl  printf
```

foo() overwrites
X0 if we don't
do something!!

```
foo: movz X0, #2  
     br  X30
```

- "Someone" needs to save/restore values when a function is called to ensure this doesn't happen

Two Possible Solutions

- Either the **called** function **saves** register values before it overwrites them and **restores** them before the function returns (**callee** saved)...

```
main: movz X0, #1
      bl  foo
      bl  printf
```

```
foo:  stur X0, [stack]
      movz X0, #2
      ldur X0, [stack]
      br  X30
```

- Or the **calling** function **saves** register values before the function call and **restores** them after the function call (**caller** saved)...

```
main: movz X0, #1
      stur X0, [stack]
      bl  foo
      ldur X0, [stack]
      bl  printf
```

```
foo:  movz X0, #2
      br  X30
```

Another example

Original C Code

```
void foo() {
    int a,b,c,d;

    a = 5; b = 6;
    c = a+1; d=c-1;

    bar();

    d = a+d;
    return();
}
```

No need to
save r2/r3.
Why?

Additions for Caller-save

```
void foo() {
    int a,b,c,d;

    a = 5; b = 6;
    c = a+1; d=c-1;
    save r1 to stack
    save r4 to stack
    bar();
    restore r4
    restore r1
    d = a+d;
    return();
}
```

Assume bar() will
overwrite registers
holding a,d

Additions for Callee-save

```
void foo() {
    int a,b,c,d;
    save r1
    save r2
    save r3
    save r4
    a = 5; b = 6;
    c = a+1; d=c-1;
    bar();
    d = a+d;
    restore r4
    restore r3
    restore r2
    restore r1
    return();
}
```

bar() will save a,b, but
now foo() must save
main's variables

“caller-save” vs. “callee-save”

- Caller-save

- What if bar() doesn't use r1/r4?
- No harm done, but wasted work

```
void foo(){
    int a,b,c,d;

    a = 5; b = 6;
    c = a+1; d=c-1;
    save r1 to stack
    save r4 to stack
    bar();
    restore r1
    restore r4
    d = a+d;
    return();
}
```

- Callee-save

- What if main() doesn't use r1-r4?
- No harm done, but wasted work

```
void foo(){
    int a,b,c,d;
    save r1
    save r2
    save r3
    save r4
    a = 5; b = 6;
    c = a+1; d=c-1;
    bar();
    d = a+d;
    restore r1
    restore r2
    restore r3
    restore r4
    return();
}
```


Saving/Restoring Optimizations

CALLER-CALLEE

- Where can we avoid loads/stores?
- **Caller-saved**
 - Only needs saving if value is “live” across function call
 - **Live** = contains a useful value: Assign value before function call, use that value after the function call
 - In a leaf function (a function that calls no other function), caller saves can be used without saving/restoring

a, d are live

b, c are NOT
live

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

Saving/Restoring Optimizations

CALLER-CALLEE

- Where can we avoid loads/stores?
- Callee-saved
 - Only needs saving at beginning of function and restoring at end of function
 - Only save/restore it if function overwrites the register

Only use r1-
r4

No need to
save other
registers

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

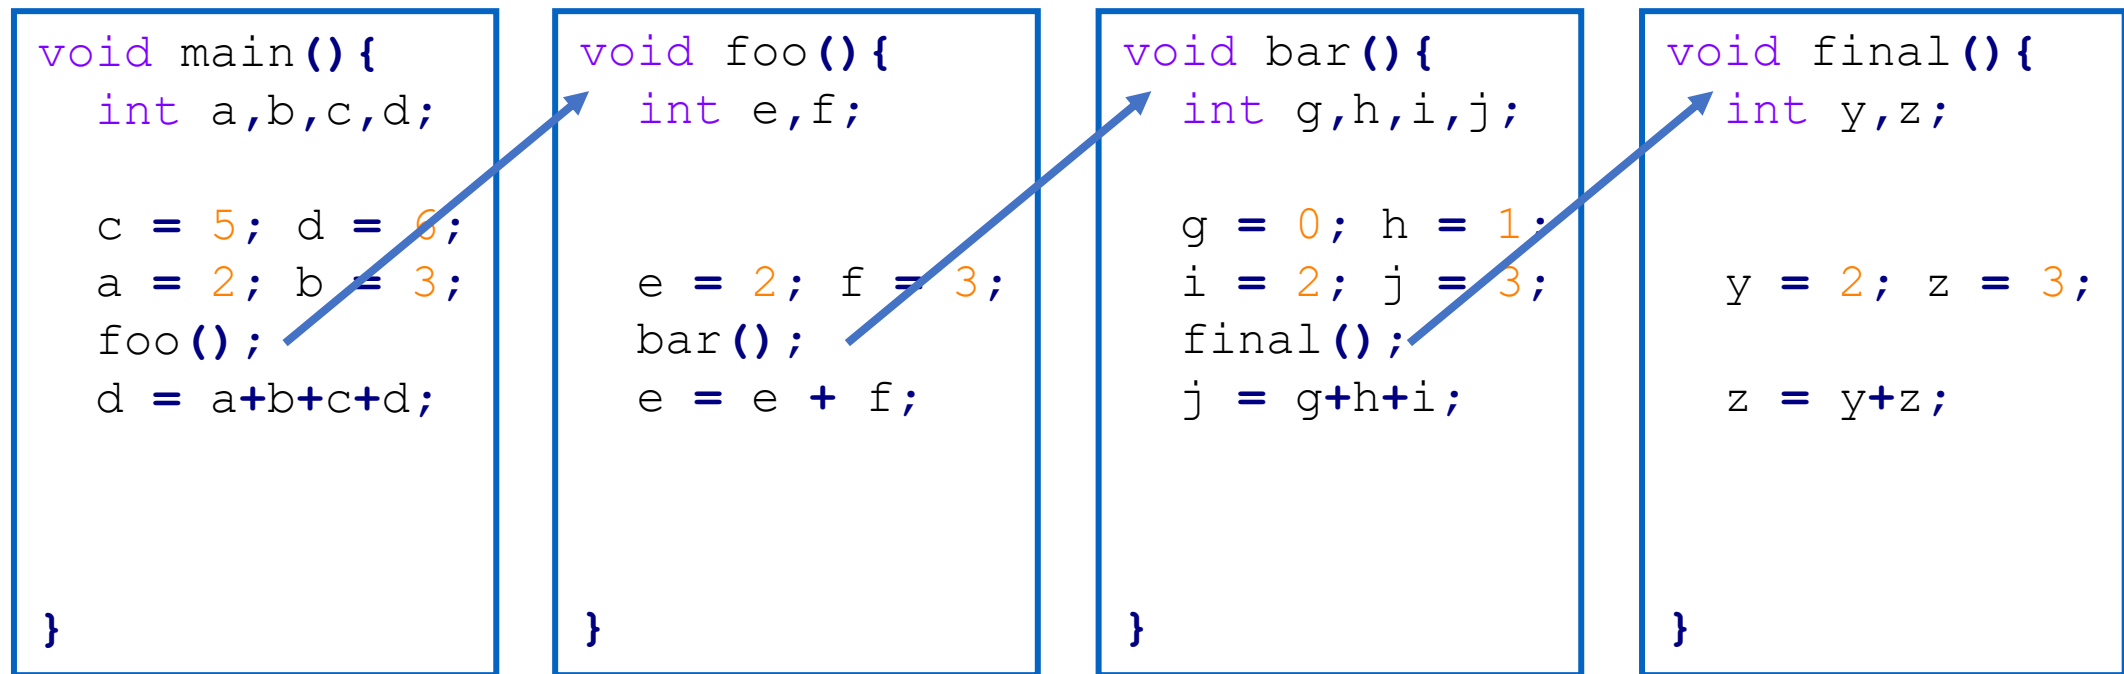
Agenda

- Using branches more generally
- Function calls and the call stack
- Assigning variables to memory locations
- Saving registers
- **Caller/callee example**

Caller versus Callee

- Which is better??
- Let's look at some examples...
- Simplifying assumptions:
 - A function can be invoked by many different call sites in different functions.
 - Assume no inter-procedural analysis (hard problem)
 - A function has no knowledge about which registers are used in either its caller or callee
 - Assume `main()` is not invoked by another function
- Implication
 - Any register allocation optimization is done using function local information

Caller-saved vs. callee saved – Multiple function case



Note: assume main does not have to save any callee registers

Caller-saved vs. callee saved – Multiple function case

- Questions:

1. How many registers need to be saved/restored if we use a **caller-save** convention?
2. How many registers need to be saved/restored if we use a **callee-save** convention?
3. How many registers need to be saved/restored if we use a mix of **caller-save** and **callee-save**?

Question 1: Caller-save

```
void main() {  
    int a,b,c,d;  
    c = 5; d = 6;  
    a = 2; b = 3;  
    [4 STUR]  
    foo();  
    [4 LDUR]  
    d = a+b+c+d;  
}
```

```
void foo() {  
    int e,f;  
  
    e = 2; f = 3;  
    [2 STUR]  
    bar();  
    [2 LDUR]  
    e = e + f;  
}
```

```
void bar() {  
    int g,h,i,j;  
    g = 0; h = 1;  
    i = 2; j = 3;  
    [3 STUR]  
    final();  
    [3 LDUR]  
    j = g+h+i;  
}
```

```
void final() {  
    int y,z;  
  
    y = 2; z = 3;  
  
    z = y+z;  
}
```

Total: 9 STUR / 9 LDUR

Question 2: Callee-save

Poll: How many ld/st pairs are needed?

```
void main() {  
    int a,b,c,d;  
  
    c = 5; d = 6;  
    a = 2; b = 3;  
    foo();  
    d = a+b+c+d;  
  
}
```

```
void foo() {  
    [2 STUR]  
    int e,f;  
  
    e = 2; f = 3;  
    bar();  
    e = e + f;  
  
    [2 LDUR]  
}
```

```
void bar() {  
    [4 STUR]  
    int g,h,i,j;  
    g = 0; h = 1;  
    i = 2; j = 3;  
    final();  
    j = g+h+i;  
  
    [4 LDUR]  
}
```

```
void final() {  
    [2 STUR]  
    int y,z;  
  
    y = 2; z = 3;  
  
    z = y+z;  
  
    [2 LDUR]  
}
```

Total: 8 STUR / 8 LDUR

Is one better?

- **Caller-save** works best when we don't have many live values across function call
- **Callee-save** works best when we don't use many registers overall
- We probably see functions of both kinds across an entire program
- Solution:
 - Use both!
 - E.g. if we have 6 registers, use some (say r0-r2) as **caller-save** and others (say r3-r5) as **callee-save**
 - Now each function can optimize for each situation to reduce saving/restoring
 - Not discussed further in this class

LEGv8 ABI- Application Binary Interface

- The ABI is an agreement about how to use the various registers
- Not enforced by hardware, just a convention by programmers / compilers
- If you won't your code to work with other functions / libraries, **follow these**
- Some register conventions in ARMv8
 - X30 is the **link register** – used to hold return address
 - X28 is **stack pointer** – holds address of top of stack
 - X19-X27 are **callee-saved** – function must save these before writing to them
 - X0-15 are **caller-saved** –function must save live values before call
 - X0-X7 used for **arguments** (memory used if more space is needed)
 - X0 used for **return value**

Next Time

- Finish Up Function Calls
- Talks about linking – the final puzzle piece of software