

# EECS 370 - Lecture 8

## Combinational Logic



# Announcements

- P2
  - P1 s+m due today
  - P2 posted by tomorrow
- HW 1
  - Due Monday
- Lab 4 meets Fr/M
  - Don't forget the pre-lab quiz tonight!

# What do object files look like?

```
extern int X;  
extern void foo();  
int Y;  
  
void main() {  
    Y = X + 1;  
    foo();  
}
```

"extern" means  
defined in another  
file

```
extern int Y;  
int X;  
  
void foo() {  
    Y *= 2;  
}
```

Compile →

```
.main:  
LDUR    X1, [XZR, X]  
ADDI    X9, X1, #1  
STUR    X9, [XZR, Y]  
BL      foo  
HALT
```

Compile →

```
.foo:  
LDUR    X1, [XZR, Y]  
LSL     X9, X1, #1  
STUR    X9, [XZR, Y]  
BR      X30
```

Uh-oh!  
Don't know  
address of X, Y,  
or foo!

# Linking

```
.main:  
LDUR    X1, [XZR, X]  
ADDI    X9, X1, #1  
STUR    X9, [XZR, Y]  
BL      foo  
HALT
```

```
.foo:  
LDUR    X1, [XZR, Y]  
LSL     X9, X1, #1  
STUR    X9, [XZR, Y]  
BR      X30
```

Assemble

???

Assemble

???

What needs to go  
in this intermediate  
"object file"?

LINK

LINK

NOTE: this will  
actually be in  
machine code, not  
assembly

```
LDUR    X1, [XZR, #40]  
ADDI    X9, X1, #1  
STUR    X9, [XZR, #36]  
BL      #2  
HALT  
LDUR    X1, [XZR, #36]  
LSL     X9, X1, #1  
STUR    X9, [XZR, #36]  
BR      X30  
// Addr #36 starts here
```

# Linking

.main:

```
LDUR    X1, [XZR, X]  
ADDI    X9, X1, #1  
STUR    X9, [XZR, Y]  
BL      foo  
HALT
```

Assemble

???

LINK

```
LDUR    X1, [XZR, #40]  
ADDI    X9, X1, #1  
STUR    X9, [XZR, #36]  
BL      #2  
HALT  
LDUR    X1, [XZR, #36]  
LSL     X9, X1, #1  
STUR    X9, [XZR, #36]  
BR      X30
```

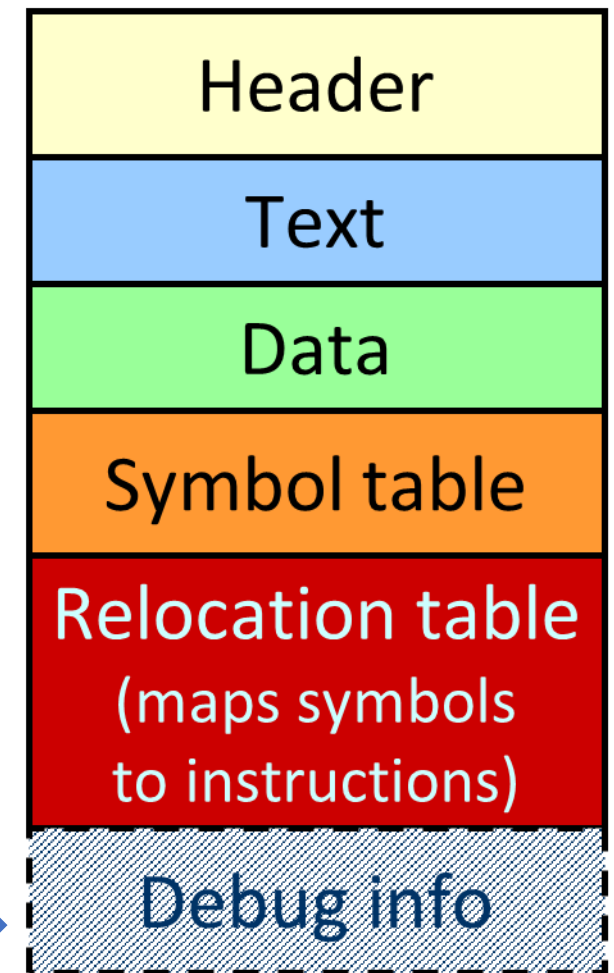
We need:

- the assembled machine code:
- list of instructions that need to be updated once addresses are resolved
- list of symbols to cross-ref

# What do object files look like?

- Since we can't make executable, we make an object file
- Basically, includes the machine code that will go in the executable
  - Plus extra information on what we need to modify once we stitch all the other object files together
- Looks like this ->

## Object code format



We won't discuss "Debug" much. Get's included when you compile with "-g" in gcc



# Assembly → Object file - example

```
extern int G;
extern void B();
int X = 3;
main() {
    Y = G + 1;
    B();
}
```

```
LDUR    X1, [XZR, G]
ADDI    X9, X1, #1
BL      B
```

<b>Header</b>	Name	foo	
	Text size	0x0C //probably bigger	
	Data size	0x04 //probably bigger	
<b>Text</b>	Address	Instruction	
	0	LDUR X1, [XZR, G]	
	4	ADDI X9, X1, #1	
	8	BL B	
<b>Data</b>	0	X	3
<b>Symbol table</b>	Label	Address	
	X	0	
	B	-	
	main	0	
	G	-	
<b>Reloc table</b>	Addr	Instruction type	Dependency
	0	LDUR	G
	8	BL	B

# Assembly → Object file - example

```
extern int X;
extern void B();
int X = 3;
main() {
    Y = G + 1;
    B();
}
```

**Header:**  
keeps track of  
size of each  
section

```
LDUR    X1, [XZR, G]
ADDI    X9, X1, #1
BL      B
```

<b>Header</b>	Name	foo	
	Text size	0x0C //probably bigger	
	Data size	0x04 //probably bigger	
<b>Text</b>	Address	Instruction	
	0	LDUR X1, [XZR, G]	
	4	ADDI X9, X1, #1	
	8	BL B	
<b>Data</b>	0	X	3
<b>Symbol table</b>	Label	Address	
	X	0	
	B	-	
	main	0	
	G	-	
<b>Reloc table</b>	Addr	Instruction type	Dependency
	0	LDUR	G
	8	BL	B



# Assembly → Object file - example

```
extern int G;
extern void B();
int X = 3;
main() {
    Y = G +
    B();
}
```

**Text:**  
machine code

```
LDUR    X1, [XZR, G]
ADDI    X9, X1, #1
BL      B
```

<b>Header</b>	Name	foo	
	Text size	0x0C //probably bigger	
	Data size	0x04 //probably bigger	
<b>Text</b>	Address	Instruction	
	0	LDUR X1, [XZR, G]	
	4	ADDI X9, X1, #1	
	8	BL B	
<b>Data</b>	0	X	3
<b>Symbol table</b>	Label	Address	
	X	0	
	B	-	
	main	0	
	G	-	
<b>Reloc table</b>	Addr	Instruction type	Dependency
	0	LDUR	G
	8	BL	B

Simplifying Assumption for EECS370

All globals and static locals (initialized or not) go in the data segment

Assembly → Object file - example

```
extern int G;
extern void B();
int X = 3;
main() {
    Y = G + 1;
    B();
}
```

**Data:**  
initialized globals  
and static locals

```
LDUR    X1, [XZR, G]
ADDI    X9, X1, #1
BL      B
```

Header	Name	foo	
	Text size	0x0C //probably bigger	
	Data size	0x04 //probably bigger	
Text	Address	Instruction	
	0	LDUR X1, [XZR, G]	
	4	ADDI X9, X1, #1	
	8	BL B	
Data	0	X	3
Symbol table	Label	Address	
	X	0	
	B	-	
	main	0	
Reloc table	G	-	
	Addr	Instruction type	Dependency
	0	LDUR	G
	8	BL	B



# Assembly → Object file - example

```
extern int G;
extern void B();
int X = 3;
main() {
    Y = G + 1;
    B();
}
```

LDUR  
ADDI  
BL

**Symbol table:**  
Lists all labels  
visible outside this file  
(i.e. function names  
and global variables)

<b>Header</b>	Name	foo	
	Text size	0x0C //probably bigger	
	Data size	0x04 //probably bigger	
<b>Text</b>	Address	Instruction	
	0	LDUR X1, [XZR, G]	
	4	ADDI X9, X1, #1	
	8	BL B	
<b>Data</b>	0	X	3
<b>Symbol table</b>	Label	Address	
	X	0	
	B	-	
	main	0	
	G	-	
<b>Reloc table</b>	Addr	Instruction type	Dependency
	0	LDUR	G
	8	BL	B

# Assembly → Object file - example

```
extern int G;
extern void B();
int X = 3;
main() {
    Y = G + 1;
    B();
}
```

LDUR X1, [XZR, G]

**Relocation Table:**  
list of instructions and data words that must be updated if things are moved in memory

<b>Header</b>	Name	foo	
	Text size	0x0C //probably bigger	
	Data size	0x04 //probably bigger	
<b>Text</b>	Address	Instruction	
	0	LDUR X1, [XZR, G]	
	4	ADDI X9, X1, #1	
	8	BL B	
<b>Data</b>	0	X	3
<b>Symbol table</b>	Label	Address	
	X	0	
	B	-	
	main	0	
	G	-	
<b>Reloc table</b>	Addr	Instruction type	Dependency
	0	LDUR	G
	8	BL	B

# Class Problem 1

**Poll: Which symbols will be put in the symbol table?** (i.e. which "things" should be visible to all files?)

```
file1.c
extern void bar(int);
extern char c[];
int a;
int foo (int x) {
    int b;
    a = c[3] + 1;
    bar(x);
    b = 27;
}
```

file 1 – symbol table

sym	loc
a	data
foo	text
c	-
bar	-

```
file2.c
extern int a;
char c[100];
void bar (int y) {
    char e[100];
    a = y;
    c[20] = e[7];
}
```

file 2 – symbol table

sym	loc
c	data
bar	text
a	-

**Poll: Which lines / instructions are in the relocation table?** (i.e. which "things" need to be updated after linking?)

## Class Problem 2

file1.c

```
1 extern void bar(int);
2 extern char c[];
3 int a;
4 int foo (int x) {
5     int b;
6     a = c[3] + 1;
7     bar(x);
8     b = 27;
9 }
```

file 1 - relocation table

line	type	dep
6	ldur	c
6	stur	a
7	bl	bar

file2.c

```
1 extern int a;
2 char c[100];
3 void bar (int y) {
4     char e[100];
5     a = y;
6     c[20] = e[7];
7 }
```

file 2 - relocation table

line	type	dep
5	stur	a
6	stur	c

Note: in a real relocation table, the "line" would really be the address in "text" section of the instruction we need to update.

# Linker

- Stitches independently created object files into a single executable file (i.e., a.out)
  - Step 1: Take text segment from each .o file and put them together.
  - Step 2: Take data segment from each .o file, put them together, and concatenate this onto end of text segments.
- What about libraries?
  - Libraries are just special object files.
  - You create new libraries by making lots of object files (for the components of the library) and combining them (see ar and ranlib on Unix machines).
- Step 3: Resolve cross-file references to labels
  - Make sure there are no undefined labels

# Linker - Continued

- Determine the memory locations the code and data of each file will occupy
  - Each function could be assembled on its own
  - Thus, the relative placement of code/data is not known up to this point
  - **Must relocate absolute references to reflect placement by the linker**
    - PC-Relative Addressing (beq, bne): never relocate
    - Absolute Address (mov 27, #X): always relocate
    - External Reference (usually bl): always relocate
    - Data Reference (often movz/movk): always relocate
- Executable file contains no relocation info or symbol table  
these just used by assembler/linker
  - (Not necessarily true today)



# Loader

- Executable file is sitting on the disk
- Puts the executable file code image into memory and asks the operating system to schedule it as a new process
  - Creates new address space for program large enough to hold text and data segments, along with a stack segment
  - Copies instructions and data from executable file into the new address space
  - Initializes registers (PC and SP most important)
- Take operating systems class (EECS 482) to learn more!

# Summary

- Compiler converts a single source code file into a single assembly language file
- Assembler handles directives (.fill), converts what it can to machine language, and creates a checklist for the linker (relocation table). This changes each .s file into a .o file
- Assembler does 2 passes to resolve addresses, handling internal forward references
- Linker combines several .o files and resolves absolute addresses
- Linker enables separate compilation: Thus unchanged files, including libraries need not be recompiled.
- Linker resolves remaining addresses.
- Loader loads executable into memory and begins execution

# Floating Point Arithmetic

See end of slides for  
bonus material (not  
covered in HW or  
exams)



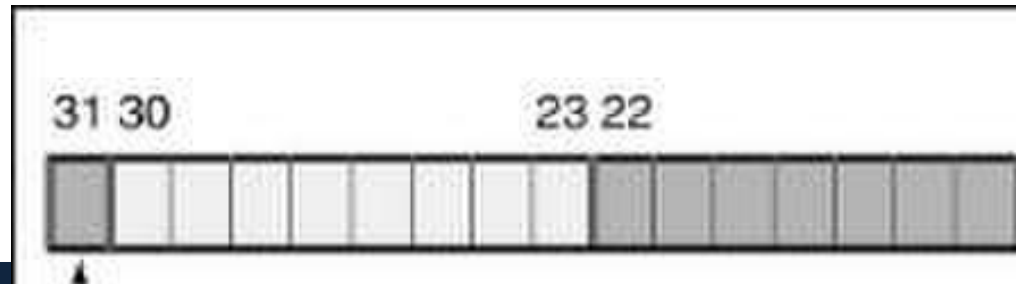
# Why floating point

- Have to represent real numbers somehow
- Rational numbers
  - Ok, but can be cumbersome to work with
- Fixed point
  - Do everything in thousandths (or millionths, etc.)
  - Not always easy to pick the right units
  - Different scaling factors for different stages of computation
- **Scientific notation: this is good!**
  - Exponential notation allows HUGE dynamic range
  - Constant (approximately) relative precision across the whole range

# IEEE Floating point format (single precision)

- Sign bit: (0 is positive, 1 is negative)
- Significand: (also called the *mantissa*; stores the 23 most significant bits after the decimal point)
- Exponent: used biased base 127 encoding
  - Add 127 to the value of the exponent to encode:
  - -127 → 00000000      1 → 10000000
  - -126 → 00000001      2 → 10000001
  - ...      ...
  - 0 → 01111111      128 → 11111111
- How do you represent zero ? Special convention:
  - Exponent: -127 (all zeroes ), Significand 0 (all zeroes), Sign + or -

Some other exception cases (e.g. NaN) we won't cover



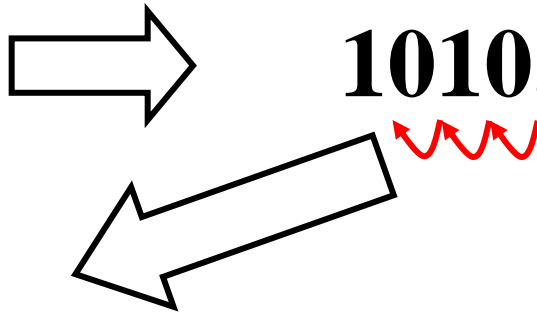
# Floating Point Representation

$$10.625_{10} \Rightarrow 1010.101_2$$

- Step 1: convert from decimal to binary
  - 1<sup>st</sup> bit after "binary" point represents 0.5 (i.e.  $2^{-1}$ )
  - 2<sup>nd</sup> bit represents 0.25 (i.e.  $2^{-2}$ )
  - etc.



# Floating Point Representation

$$10.625_{10} \quad \Rightarrow \quad 1010.101_2$$


$$1.010101 \times 2^3$$

- Step 2: normalize number by shifting binary point until you get  $1.XXX \times 2^Y$

# Floating Point Representation

$$10.625_{10} \longrightarrow 1010.101_2$$

$$1.010101 \times 2^3$$

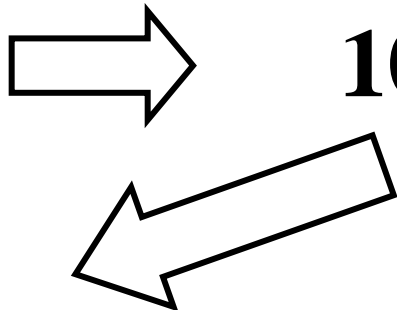
This must be a 1!  
So don't store it.



- Step 3: store relevant numbers in proper location (ignoring initial 1 of significand)



# Floating Point Representation

$$10.625_{10} \longrightarrow 1010.101_2$$


$$1.010101 \times 2^3$$

**This must be a 1!  
So don't store it.**



$$10.625_{10} = 0 \ 10000010 \ 010101000000000000000000$$

# Class Problem

Poll

- What is the value of the following IEEE 754 floating point encoded number?

1 = -  
10000101 = 133 - 127 -> exponent 6  
01011001 = mantissa

$-1.01011001 \times 2^6$

-1010110.01  
 $-(2^6 + 2^4 + 2^2 + 2^1 + 2^{-2})$   
 $-(64 + 16 + 4 + 2 + 1/4)$

-86.25

1	10000101	010110010000000000000000
---	----------	--------------------------

# What matters to a CS person?

- What happens if you add a big number to a small number?

- E.g.

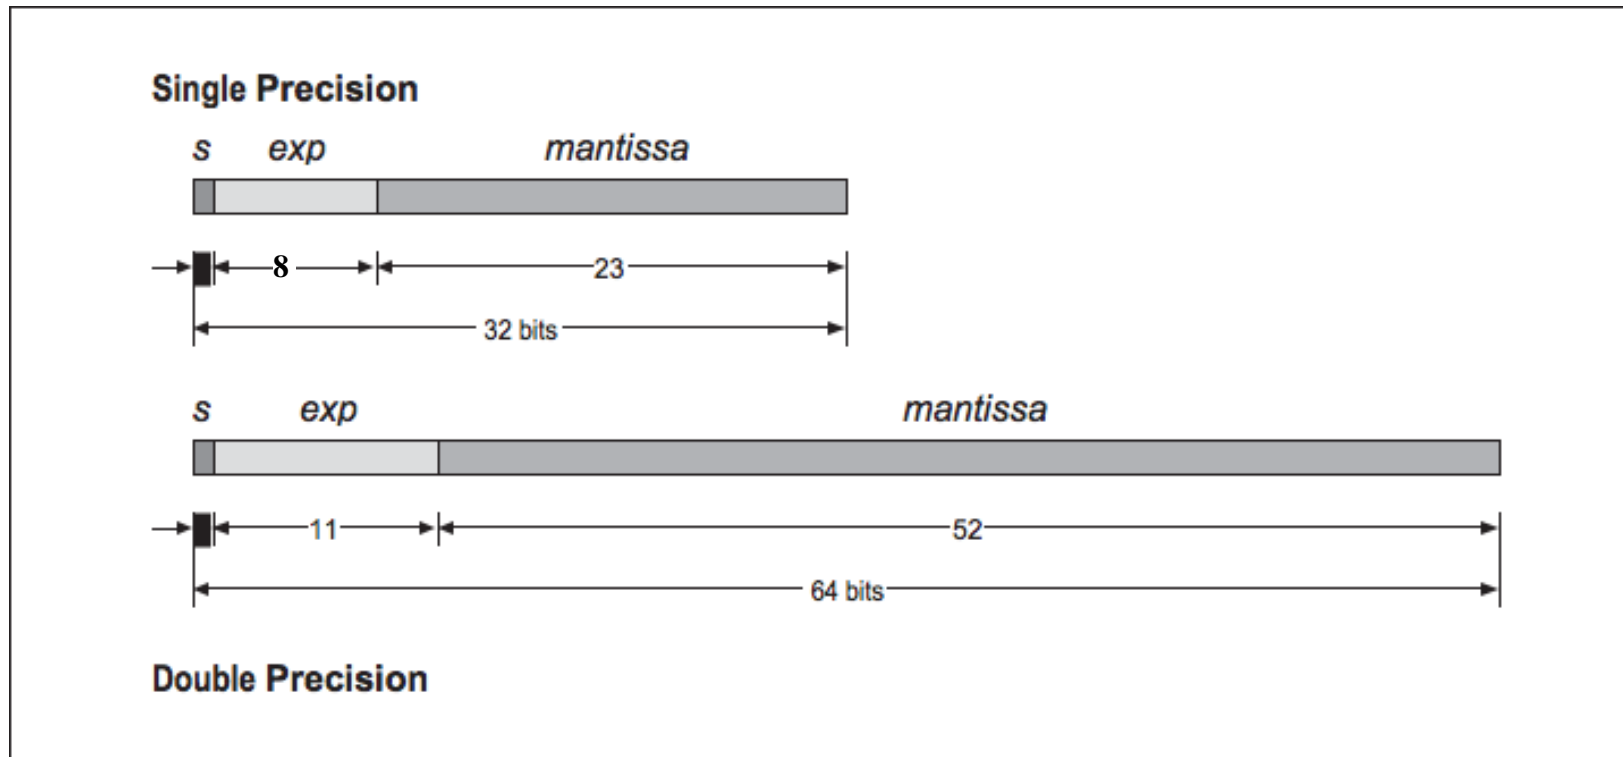
$$1000 + .00001$$

- The larger the exponent, the larger the "gap" between numbers that can be represented
- When the smaller number is added to the larger one, it can't be represented so precisely
  - Rounded down to zero: addition has no effect
  - Imagine having a loop do the above a million times: you'd end up with 1000 instead of 1010
  - Can be a big problem in scientific code
- You need to be aware of the issue.
  - This is why most people use "double" instead of "float"
  - The problem can still exist, it's just less likely.

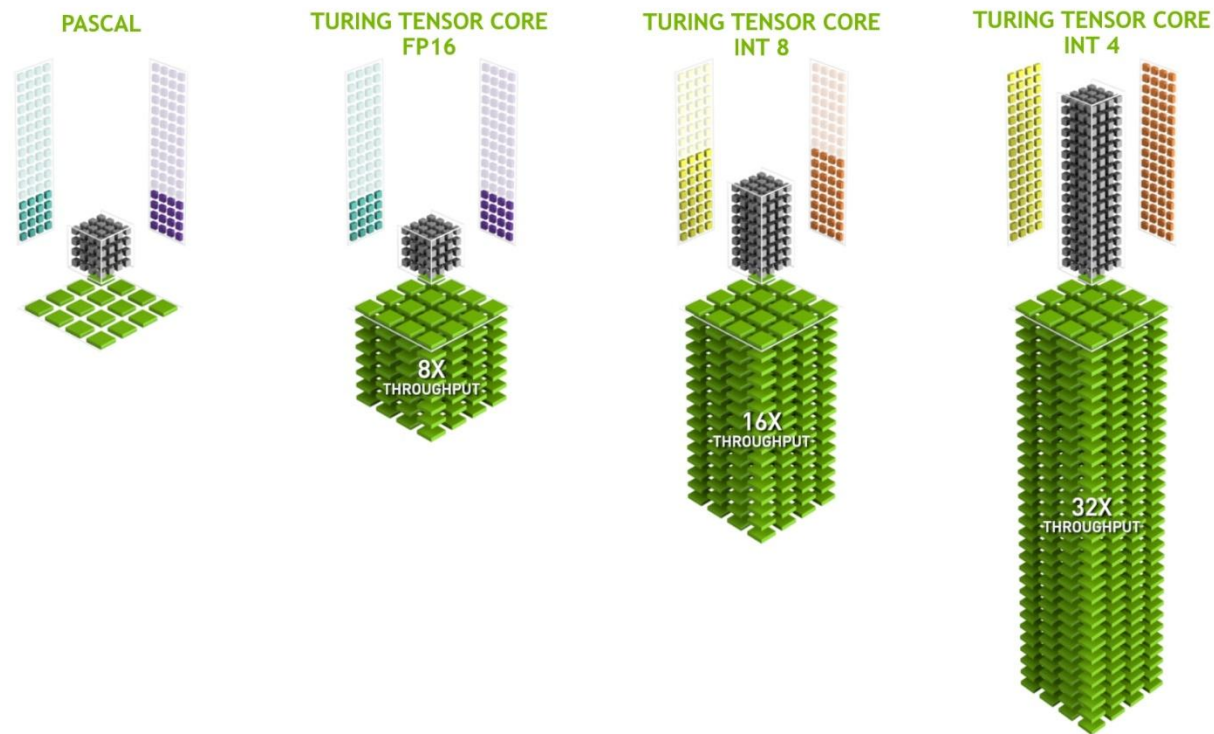
# More precision and range

- We've described IEEE-754 binary32 floating point format, i.e. “single precision” (“float” in C/C++)
  - 24 bits precision; equivalent to about 7 decimal digits
  - $3.4 * 10^{38}$  maximum value
  - Good enough for most but not all calculations
- IEEE-754 also defines larger binary64 format, “double precision” (“double” in C/C++)
  - 53 bits precision, equivalent to about 16 decimal digits
  - $1.8 * 10^{308}$  maximum value
  - Most accurate physical values currently known only to about 47 bits precision, about 14 decimal digits
- Interestingly, there's been a surge in popularity for **lower** precision floating point lately. Why?

# Single (“float”) precision



# Low Precision FP Popular in accelerating AI



# Next few lectures: Digital Logic

- Lectures 1-7:
  - LC2K and ARMv8/LEGV8 ISAs
  - Converting C to Assembly
  - Function Calls
  - Linking
- **Today:**
  - **Floating Point**
  - **Combinational Logic**
- Next lecture:
  - Sequential Logic



# Up Until Now...

- We've covered high-level C code to an executable
  - Compilation
  - Assembly
  - Linking
  - Loading
- Now, we'll talk about the hardware that runs this code
  - First step: the basics of digital logic



# Next 3 Lectures

## 1. Combinational Logic:

- Basics of electronics; logic gates, muxes, decoders

## 2. Sequential Logic

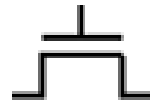
- Clocks, latches and flip-flops

## 3. State Machines

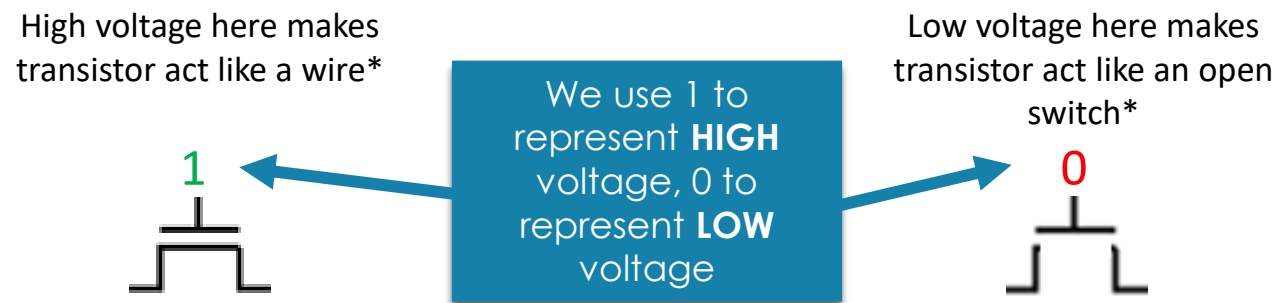
- Building a simple processor

# Transistors

- ❑ At the heart of digital logic is the transistor
- ❑ Electrical engineers draw it like this



- ❑ The physics is complicated, but at the end of the day, all it is a **really** small and **really** fast electric switch



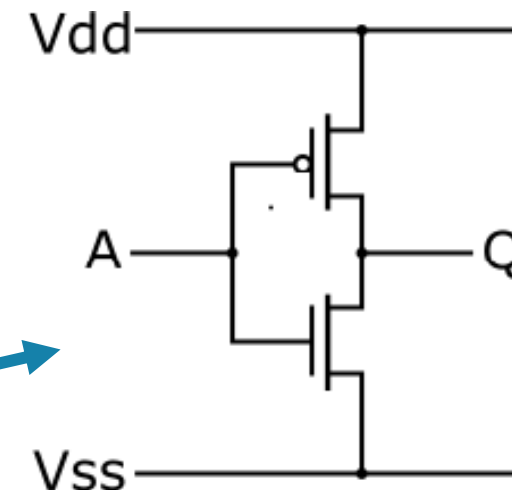
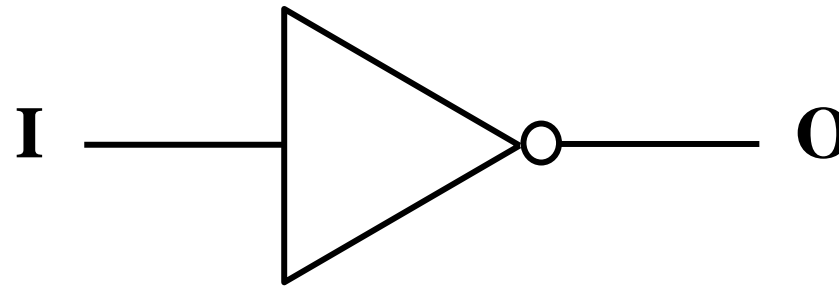
# Basic gate: Inverter

**CS abstraction**  
**- logic function**

Truth Table

I	O
0	1
1	0

**Schematic symbol (CS/EE)**



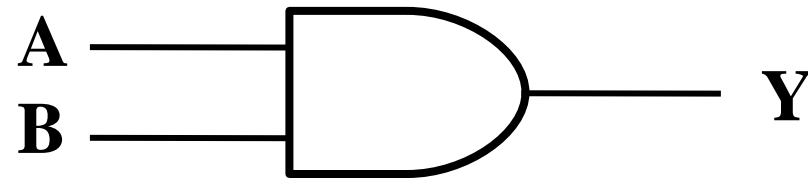
We won't focus on  
this part in 370. Take  
270 and 312 to  
learn more

# Basic gates: AND and OR

AND

Truth Table

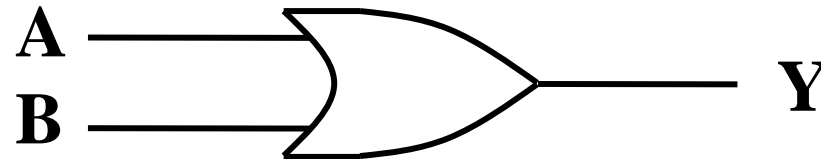
A	B	Y
0	0	0
0	1	0
1	0	0
1	1	1



OR

Truth Table

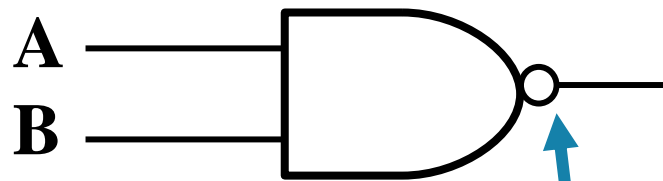
A	B	Y
0	0	0
0	1	1
1	0	1
1	1	1



# Basic gate: NAND

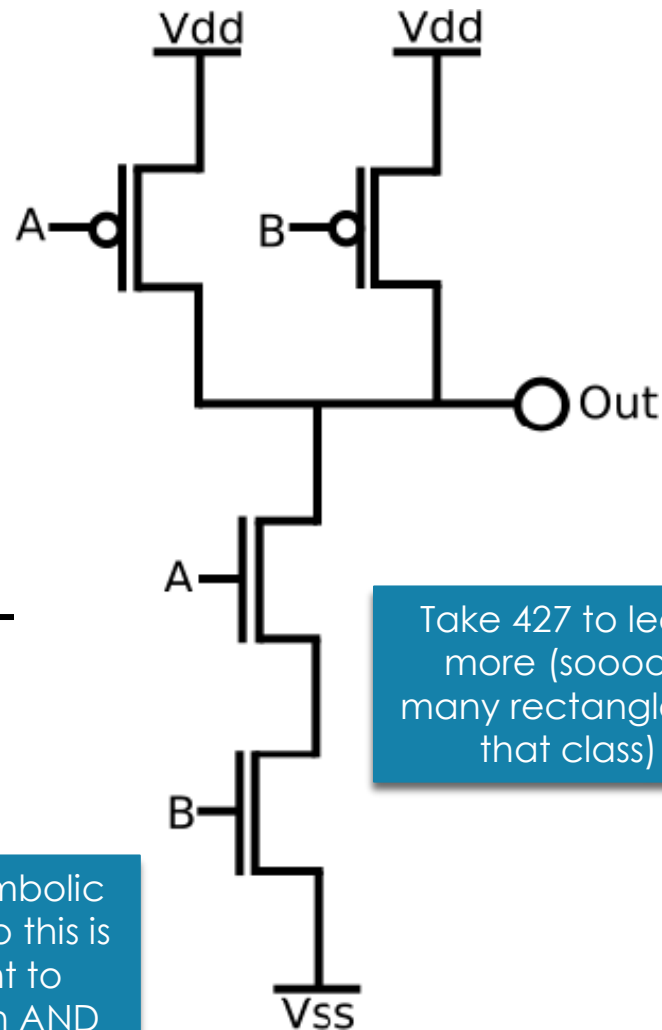
Truth Table

A	B	Y
0	0	1
0	1	1
1	0	1
1	1	0



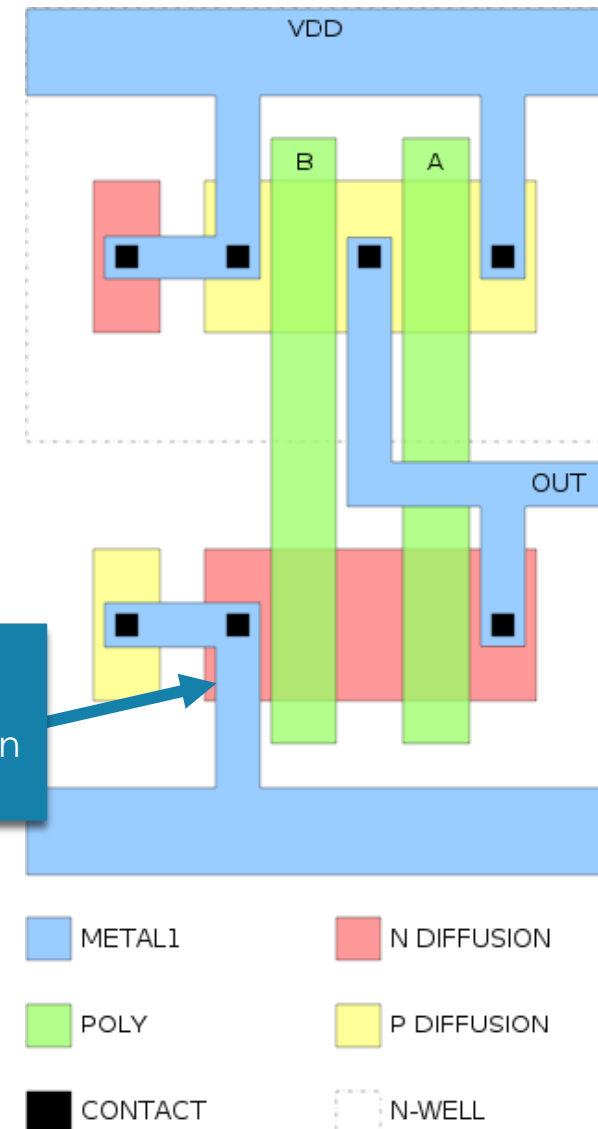
Bubble is symbolic for "invert", so this is equivalent to negating an AND gate

Transistor-level schematic



Take 427 to learn more (sooooo many rectangles in that class)

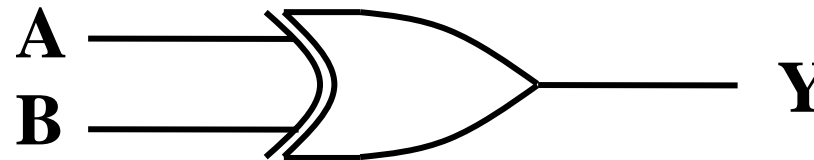
Layout schematic



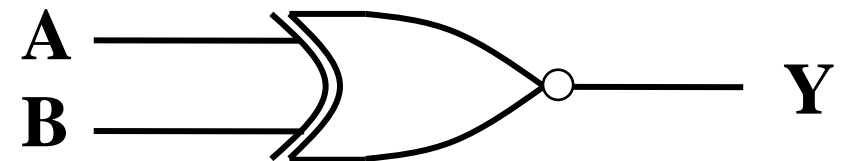
# Basic gate: XOR (eXclusive OR)

Truth Table

A	B	Y
0	0	
0	1	
1	0	
1	1	



XNOR Gate



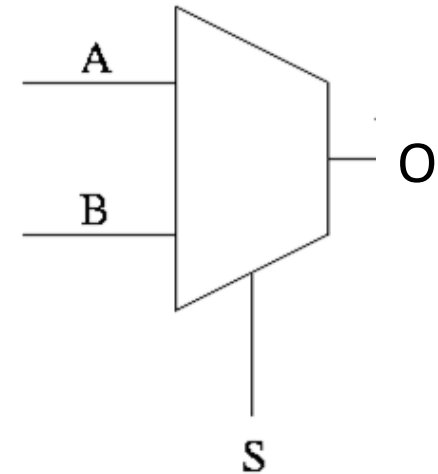
# Building Complexity: Selecting

- We want to design a circuit that can select between two inputs (multiplexer or **mux**)
- Let's do a one-bit version
  1. Draw a truth table

**Poll**: How do we fill in the truth table for this?

A	B	S	O
0	0	0	
0	0	1	
0	1	0	
0	1	1	
1	0	0	
1	0	1	
1	1	0	
1	1	1	

## Symbol



$$O = S ? B : A$$



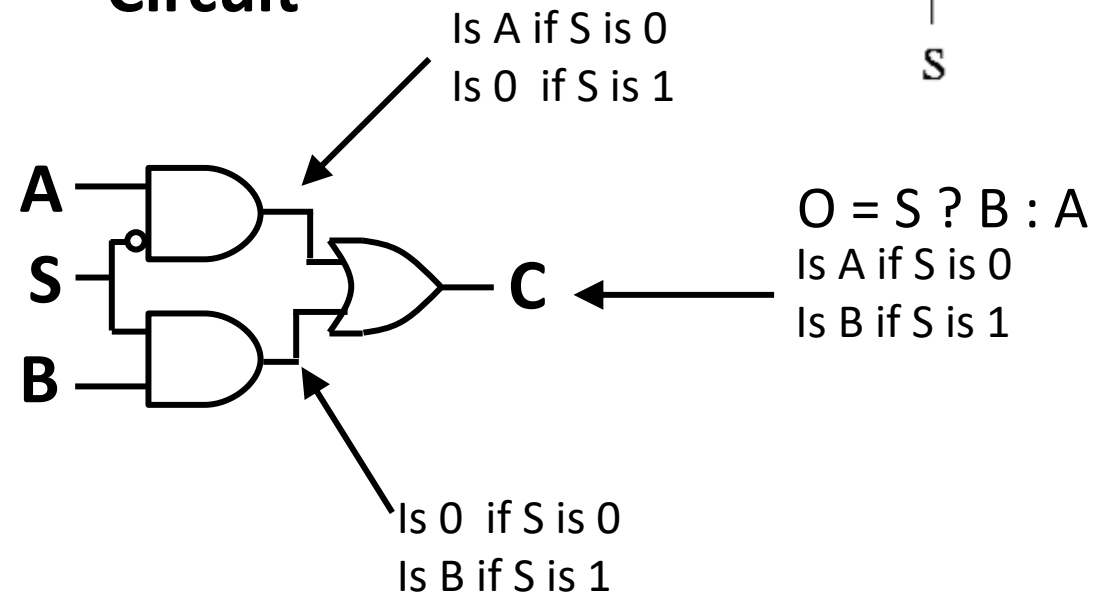
# Building Complexity: Selecting

- We want to design a circuit that can select between two inputs (multiplexor or **mux**)
- Let's do a one-bit version
  1. Draw a truth table

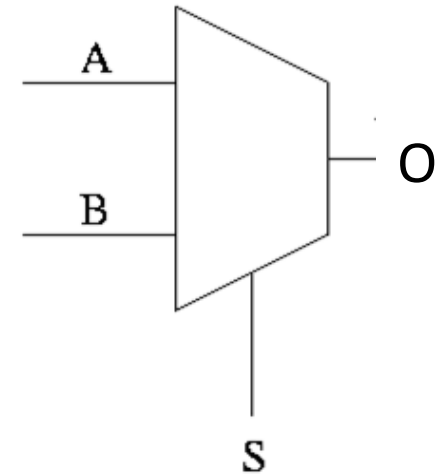
Muxes are universal! A  $2^N$  entry truth table can be implemented by passing each output value into an input of a  $2^N$ -to-1 mux

A	B	S	O
0	0	0	0
0	0	1	0
0	1	0	0
0	1	1	1
1	0	0	1
1	0	1	0
1	1	0	1
1	1	1	1

## Circuit



## Symbol





# Building Complexity: Addition

- We want to design a circuit that performs binary addition
- Let's start by adding two bits
  - Design a circuit that takes two bits (**A** and **B**) as input
    - Generates a sum and carry bit (S and C)
      1. Make a truth table
      2. Design a circuit

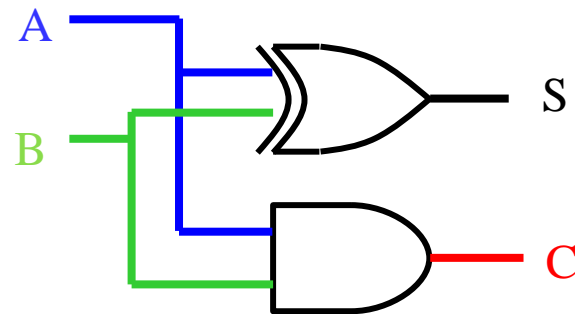
$$\begin{array}{r} 10011 \\ + 00110 \\ \hline \end{array}$$

A	B	C	S

# Building Complexity: Addition

- We want to design a circuit that performs binary addition
- Let's start by adding two bits
  - Design a circuit that takes two bits (**A** and **B**) as input
    - Generates a sum and carry bit (S and C)
  - 1. Make a truth table
  - 2. Design a circuit

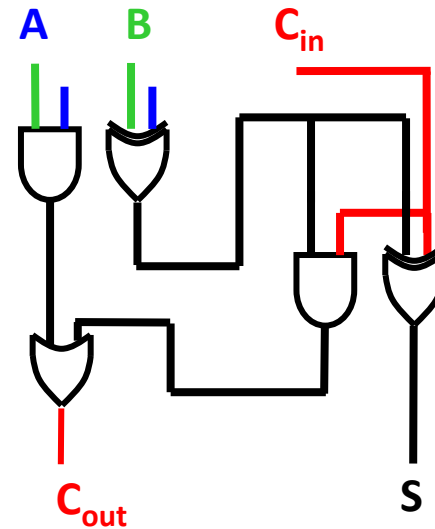
$$\begin{array}{r} 0\ 1\ 1\ 0 \\ 1\ 0\ 0\ 1\ 1 \\ + 0\ 0\ 1\ 1\ 0 \\ \hline 1\ 1\ 0\ 0\ 1 \end{array}$$



A	B	C	S
0	0	0	0
0	1	0	1
1	0	0	1
1	1	1	0

# Building Complexity: Addition

- Now we can add two bits, but how do we deal with carry bits?
- This is a **full adder**
  - We have to design a circuit that can add three bits
    - Inputs: A, B, C<sub>in</sub>
    - Outputs: S, C<sub>out</sub>
- 1. Design a truth table
- 2. Circuit
- This is a **full adder**

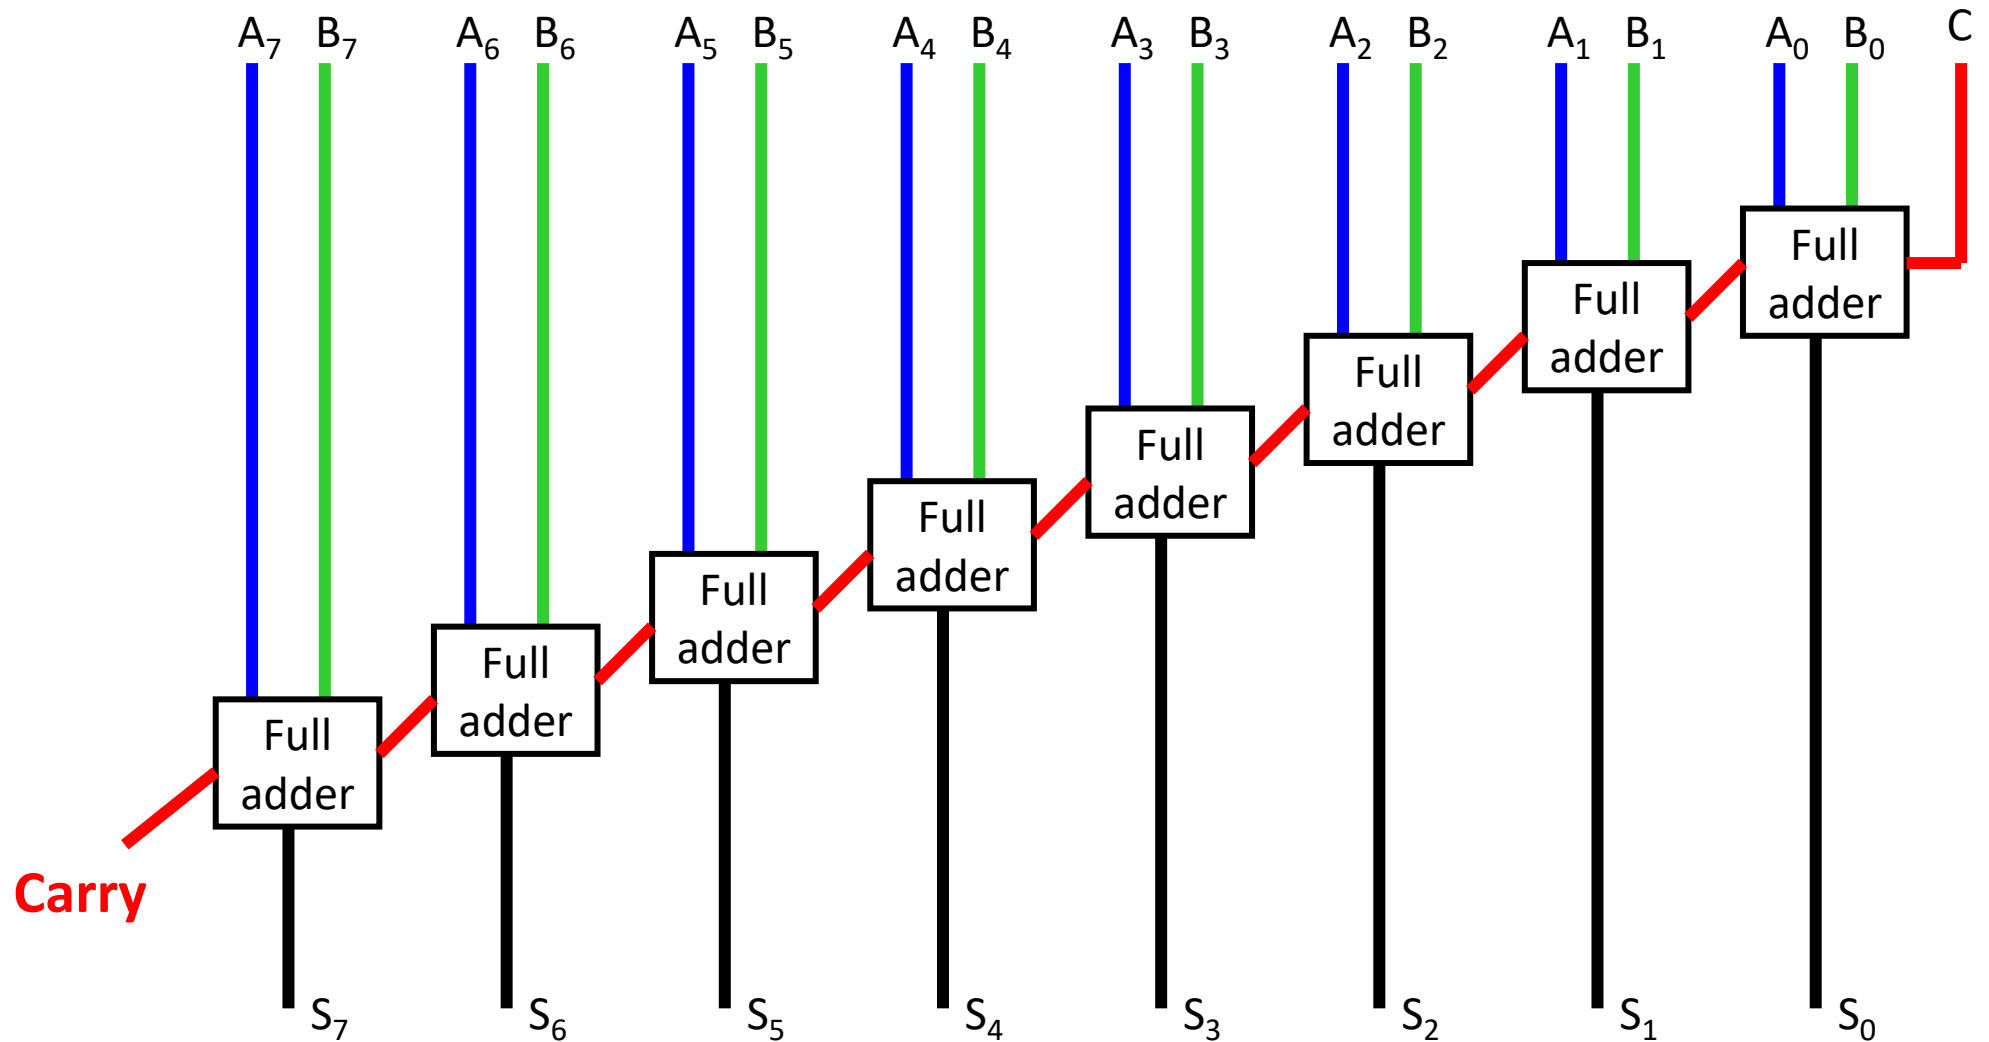


$$\begin{array}{r} 0110 \\ 10011 \\ + 00110 \\ \hline 11001 \end{array}$$

C <sub>in</sub>	A	B	C <sub>out</sub>	S
0	0	0	0	0
0	0	1	0	1
0	1	0	0	1
0	1	1	1	0
1	0	0	0	1
1	0	1	1	0
1	1	0	1	0
1	1	1	1	1

# 8-bit Ripple Carry Adder

If we invert B's bits and set C to 1, we also have a subtractor! Why?

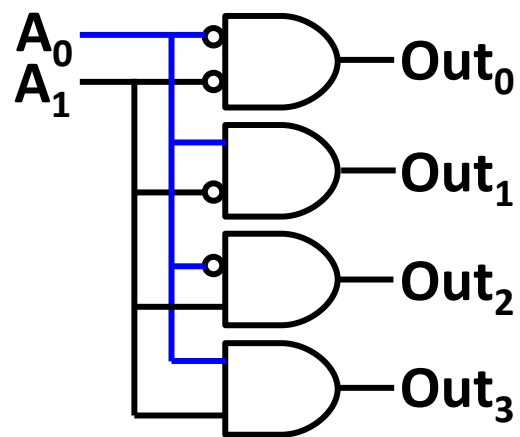


This will be very slow for 32 or 64 bit adds, but is sufficient for our needs

# Building Complexity: Decoding

- Another common device is a decoder
  - Input: N-bit binary number
  - Output:  $2^N$  bits, exactly one of which will be high
  - Allows us to **index** into things (like a register file)

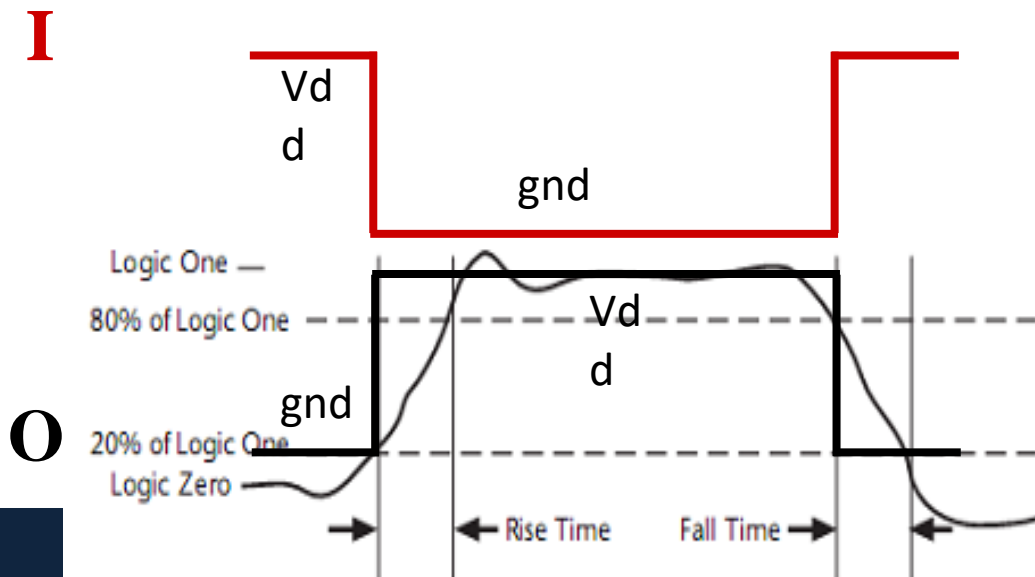
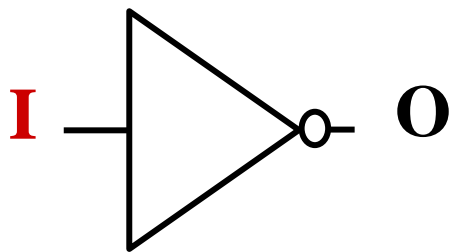
Decoder



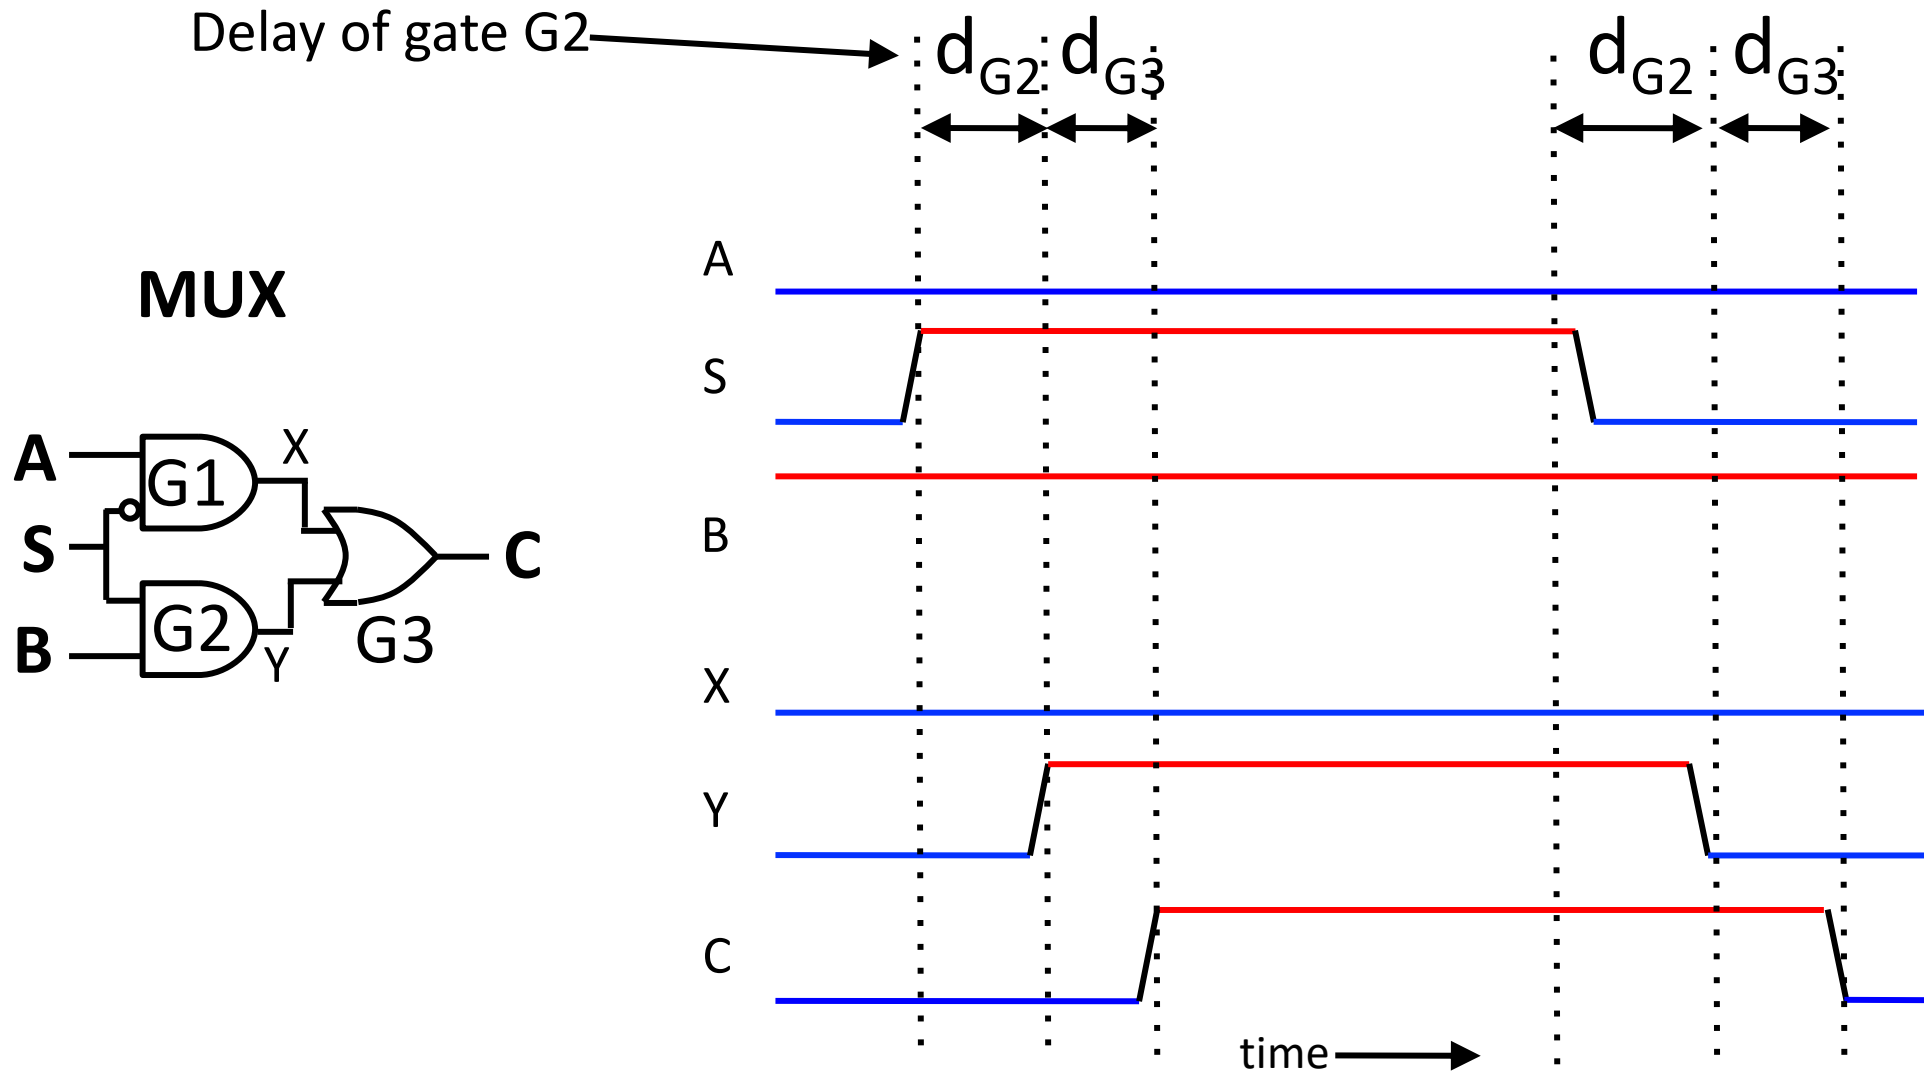
Poll: What will be the output for 101?

# Propagation delay in combinational gates

- Gate outputs do not change exactly when inputs do.
    - Transmission time over wires ( $\sim$ speed of light)
    - Saturation time to make transistor gate switch
- ⇒ Every combinatorial circuit has a propagation delay  
(time between input and output stabilization)



# Timing in Combinational Circuits

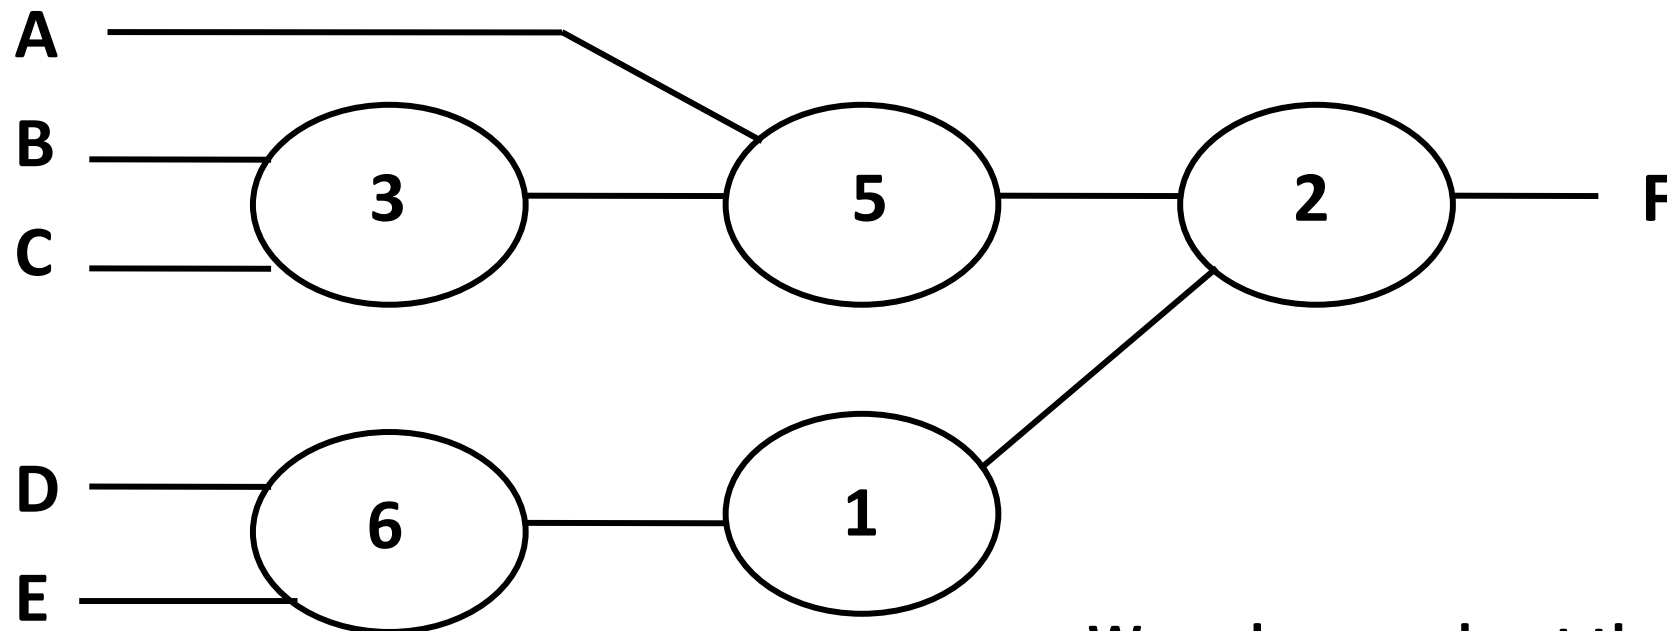


What is the input/output delay (or simply, delay) of the MUX?

# What is the delay of this Circuit?

Each oval represents one gate,  
the type does not matter

Poll : What is the delay?

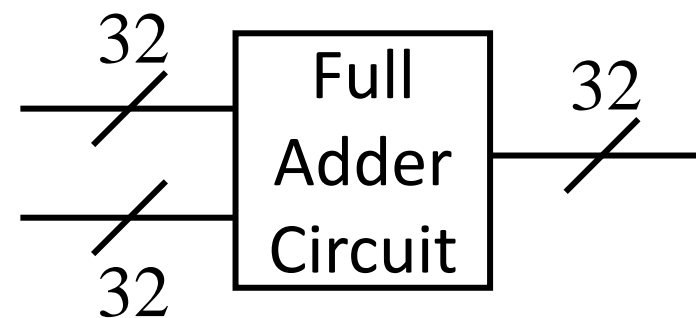
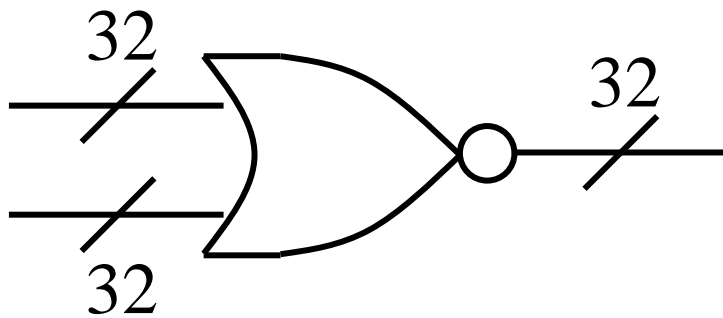


We only care about the longest  
path, or critical path



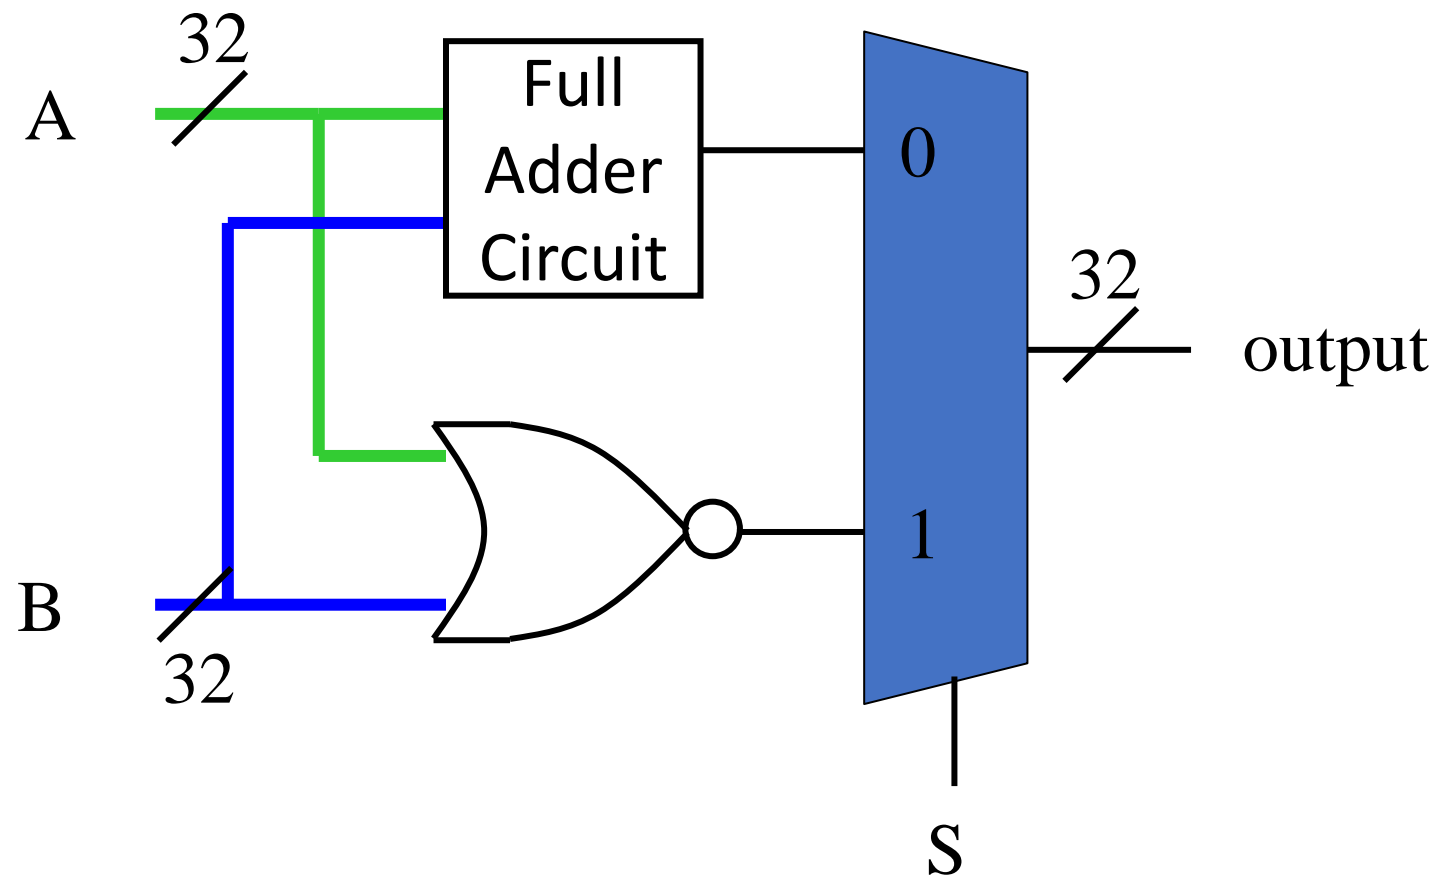
# Exercise

- Use the blocks we have learned about so far (full adder, NOR, mux) to build this circuit
  - Input A, 32 bits
  - Input B, 32 bits
  - Input S, 1 bit
  - Output, 32 bits
  - When S is low, the output is  $A+B$ , when S is high, the output is  $\text{NOR}(a,b)$
- Hint: you can express multi-bit gates like this:



# Exercise

- This is a basic ALU (Arithmetic Logic Unit)
- It is the heart of a computer processor!



# Next Time

- Logic circuits that "remember"
  - Aka "sequential logic"

# BONUS Floating Point Slides

Not  
testable

# Bonus slides – this material is not testable

- This material is here for those folks that may care.
  - **You *may* find it useful when considering the gap between representations**
  - But the material isn't directly testable.
- It is interesting if you are into that kind of thing.
- It can be useful if you are going to do scientific programming for a living.
- So it is provided as a reference, but isn't part of the class (we may cover a bit of it in lecture if we have time)

*Not  
testable*

# Floating point multiplication

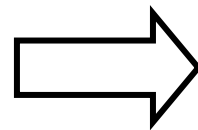
- Add exponents (don't forget to account for the bias of 127)
- Multiply significands (don't forget the implicit **1** bits)
- Renormalize if necessary
- Compute sign bit (simple exclusive-or)

# Floating point multiply

Not testable

$$10.625_{10} = 1010.101_2 \Rightarrow$$

$$10_{10} = 1010_2$$



0	10000010	010101000000000000000000
	+	×
0	10000010	010000000000000000000000
	-127	

$$\begin{array}{r}
 1010101 \\
 \times 101 \\
 \hline
 1010101 \\
 101010100 \\
 \hline
 110101001
 \end{array}$$

$$0 \ 10000101 \ 101010010000000000000000$$

$$1101010.01_2 = 106.25_{10}$$

*Not  
testable*

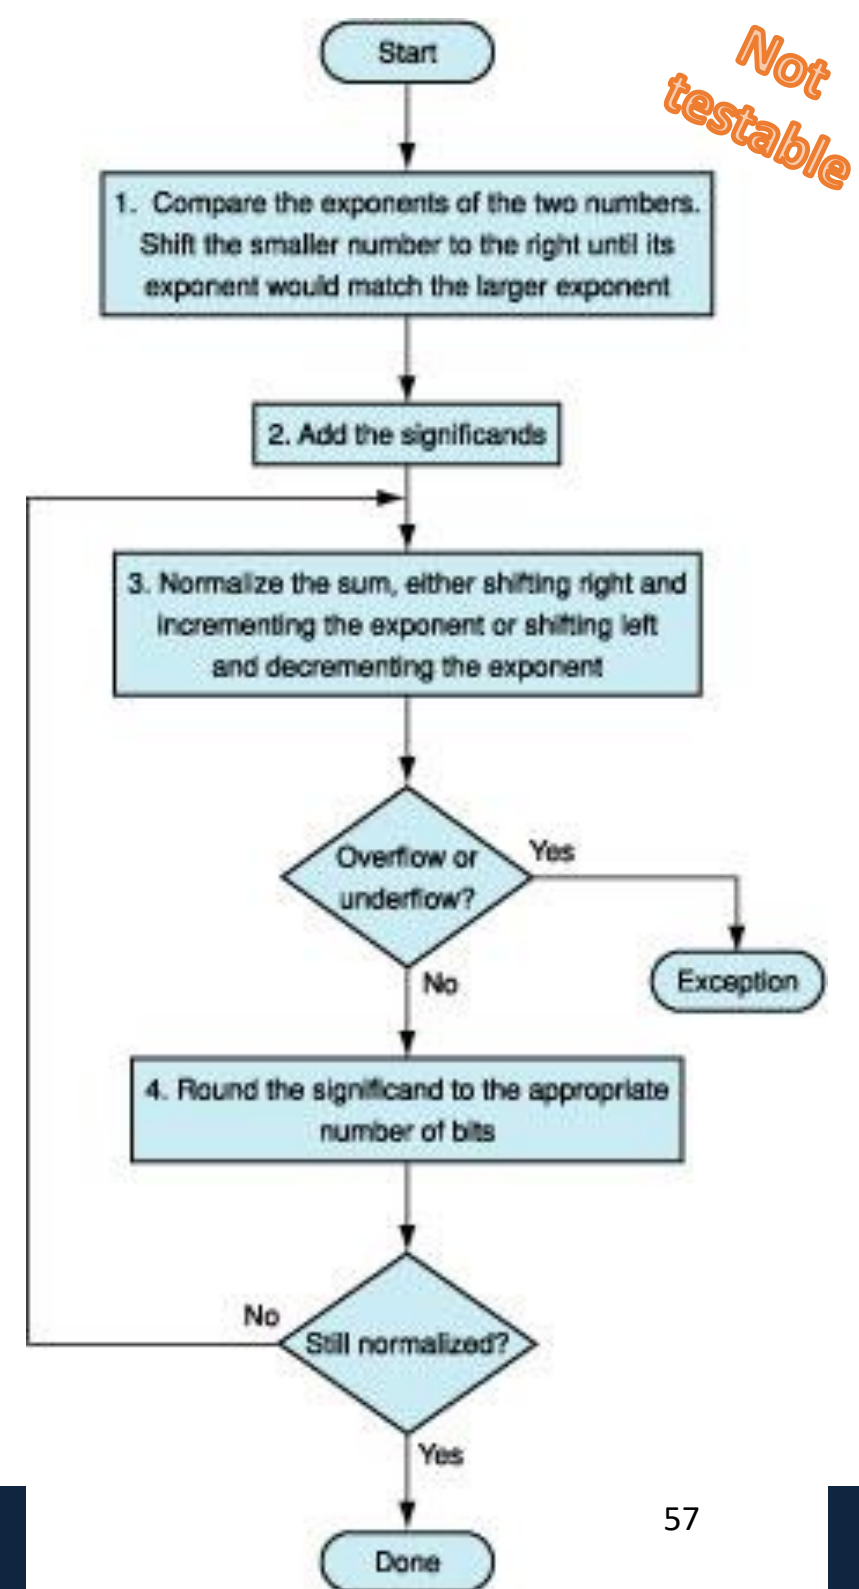
# Floating point addition

- More complicated than floating point multiplication!
- If exponents are unequal, must shift the significand of the smaller number to the right to align the corresponding place values
- Once numbers are aligned, simple addition (could be subtraction, if one of the numbers is negative)
- Renormalize (which could be messy if the numbers had opposite signs; for example, consider addition of +1.5000 and  $-1.4999$ )
- Added complication: rounding to the correct number of bits to store could denormalize the number, and require one more step



# Floating point Addition

1. Shift smaller exponent right to match larger.
2. Add significands
3. Normalize and update exponent
4. Check for “out of range”



*Not  
testable*

# Class Problem

**Show how to add the following 2 numbers using IEEE floating point addition:  $101.125 + 13.75$**

Not testable

# Class Problem

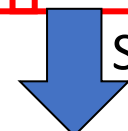
101.125



13.75



Shift by  $6-3 = 3$

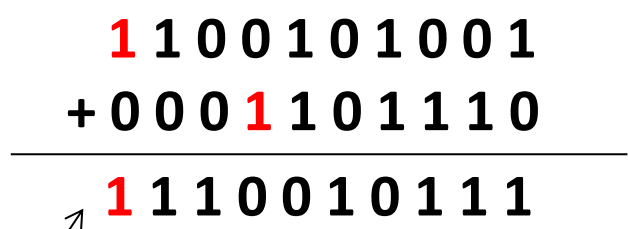


Shift mantissa by difference in exponent



Note: When shifting to the right, the first shift should put the implicit 1, then 0's

Sum Significands



Sum didn't overflow, so no re-normalization needed



= 114.875



# Class Problem

*Not  
testable*

**Show how to add the following 2 numbers using IEEE floating point addition:  $117.125 + 13.75$**

Not testable

# Class Problem

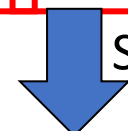
117.125



13.75



Shift by  $6-3 = 3$

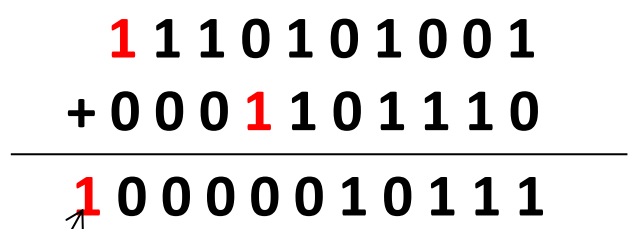


Shift mantissa by difference in exponent



Note: When shifting to the right, the first shift should put the implicit 1, then 0's

Sum Significands



Sum overflows, re-normalize by adding one to exponent and shifting mantissa by one



= 130.875

