

# EECS 370

## Multi-Level Page Tables

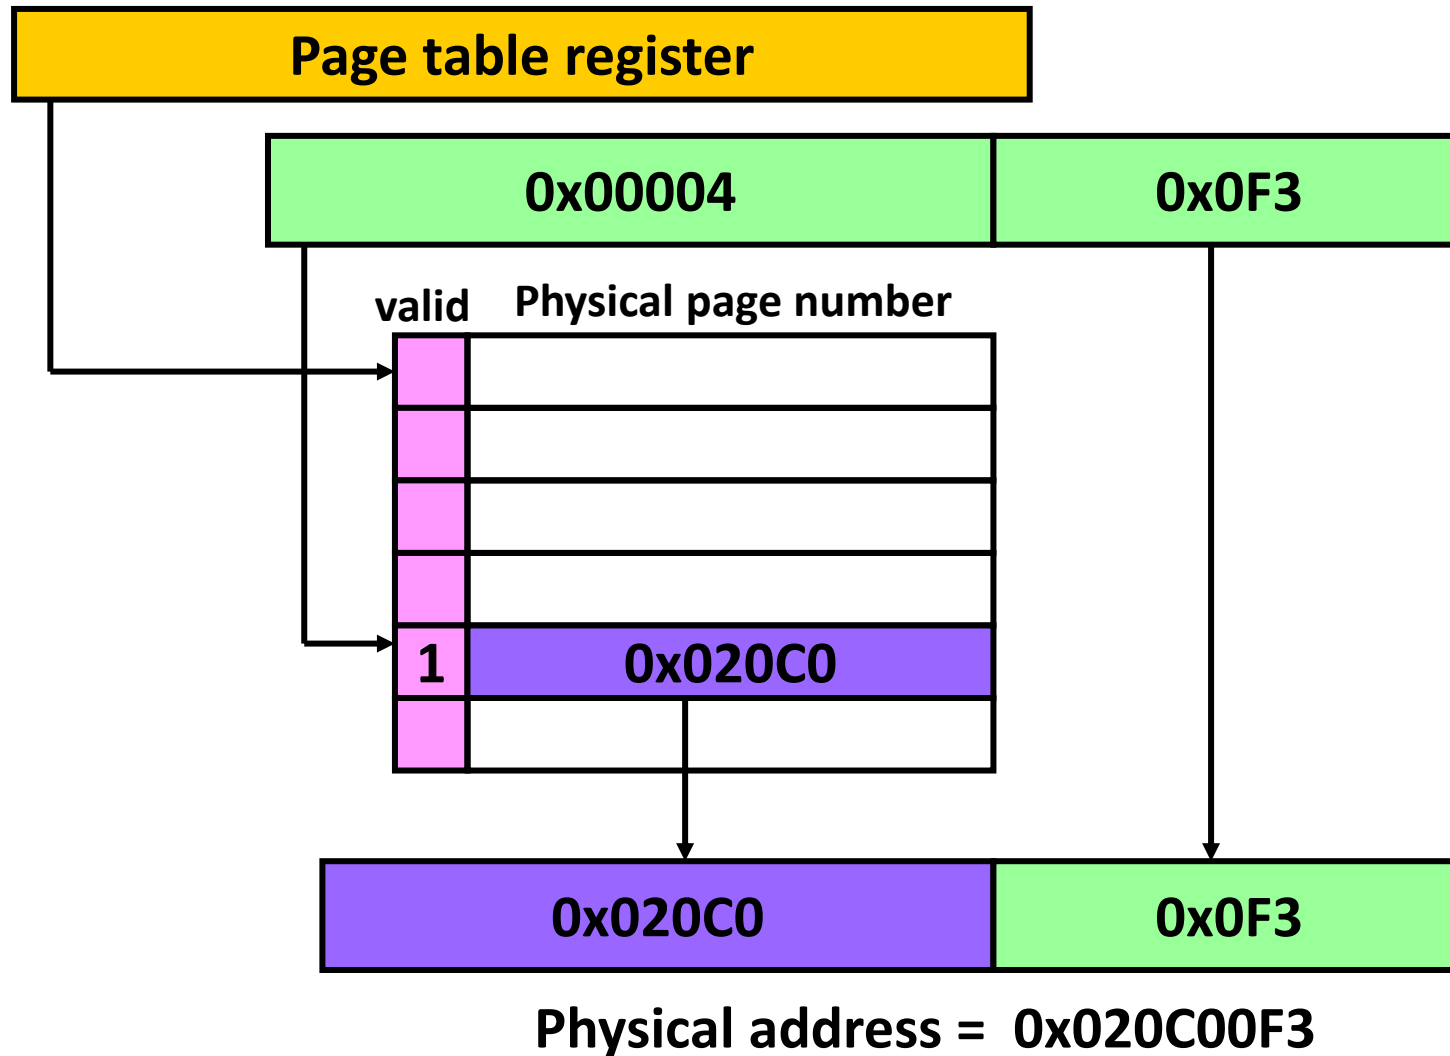


# Announcements

- Lab
  - Assignment due tonight
  - Pre-lab quiz due tomorrow
- Project 4
  - Out now, due after Thanksgiving
- Final exam
  - Previous exams posted

# Reminder: Page tables

Virtual address = 0x000040F3



# Class Problem

- Given the following:
  - 4KB page size, physical memory of 16KB, page table stored in physical page 0 and can never be evicted, 20 bit, byte-addressable virtual address space.
  - The page table initially has virtual page 0 in physical page 1, virtual page 1 in physical page 2 and no valid data in other physical pages.
- Fill in the table on the next slide for each reference
  - Note: like caches we'll use LRU when we need to replace a page.

4KB page size,  
physical memory of 16KB,  
page table stored in physical  
page 0 and can never be  
evicted, 20 bit, byte-  
addressable virtual address  
space.

The page table initially has  
virtual page 0 in physical  
page 1, virtual page 1 in  
physical page 2 and no valid  
data in other physical pages.

**Poll:** How many hex digits  
should the page number  
be?

# Class Problem (continued)

Virt addr	Virt page	Page fault?	Phys addr
0x00F0C			
0x01F0C			
0x20F0C			
0x00100			
0x00200			
0x30000			
0x01FFF			
0x00200			



# Class Problem (continued)

4KB page size,  
physical memory of 16KB,  
page table stored in physical  
page 0 and can never be  
evicted, 20 bit, byte-  
addressable virtual address  
space.

The page table initially has  
virtual page 0 in physical  
page 1, virtual page 1 in  
physical page 2 and no valid  
data in other physical pages.

Virt addr	Virt page	Page fault?	Phys addr
0x00F0C	0x0	N	0x1F0C
0x01F0C	0x1	N	0x2F0C
0x20F0C	0x20	Y (into 3)	0x3F0C
0x00100	0x0	N	0x1100
0x00200	0x0	N	0x1200
0x30000	0x30	Y (into 2)	0x2000
0x01FFF	0x1	Y (into 3)	0x3FFF
0x00200	0x0	N	0x1200

# Performance of System

- Every memory access requires an additional access to memory
  - Page table must be read
- Can really slow down our system (potentially multiple accesses to slow DRAM)
- Solution: create a cache of previous translations
  - Called a "Translation-Lookaside Buffer" (TLB)
  - Will be discussed in lab 12 and lecture next week

# Size of the page table

- How big is a page table entry?
  - For 32-bit virtual address:
    - If the machine can support  $1\text{GB} = 2^{30}$  bytes of physical memory and we use pages of size  $4\text{KB} = 2^{12}$ ,
    - then the physical page number is  $30-12 = 18$  bits.  
Plus another **valid bit** + other useful stuff (**read only, dirty, etc.**)
    - Let say about 3 bytes.
- How many entries in the page table?
  - 1 entry per virtual page
  - ARM virtual address is 32 bits – 12 bit page offset = 20
  - Total number of virtual pages =  $2^{20}$
- Total size of page table = Number of virtual pages
  - \* Size of each page table entry
  - =  $2^{20} \times 3$  bytes  $\sim 3$  MB



# Agenda

- **Motivation for Multi-level Page Tables**
- Example architecture
- Class Problem: Multi-Level VM Design
- VM Miscellanea



# How can you organize the page table?

## 2. Option 2: Use a multi-level page table

- Split up single-level page into multiple chunks
- Only allocate space for that chunk if it's non-empty (i.e. program uses some of those physical pages)
- Maintain a top-level page table which contains pointers to each of these chunks

valid	PPN
0	
0	
0	
0	
1	0x1
1	0x2
1	0x3
0	
0	
0	
0	
0	
0	
0	
0	

**Single-level: Tons of wasted space!**

valid	2 <sup>nd</sup> level page table
0	
1	0x1000
0	
0	

valid	PPN
1	0x1
1	0x2
1	0x3
0	

**Multi-level: No need to allocate 2<sup>nd</sup> level entries if all invalid**

# Multi-Level Page Table

- Only allocate second (and later) page tables when needed
- Program starts: everything is invalid, only first level is allocated

valid	2 <sup>nd</sup> level page table
0	
0	
0	
0	

- As we access more, second level page tables are allocated

valid	2 <sup>nd</sup> level page table
1	0x1000
1	0x3500

valid	PPN
1	0x1
1	0x6
1	0x2
1	0x1f

valid	PPN
1	0x9a
1	0x3
0	
1	0xff

Multi-level: Size is proportional to amount of memory used

Common case: most programs use small portion of virtual memory space

# Multi-level page tables

- If program is using much of the virtual address space, multi-level page table could be about the same or larger size than single-level

valid	PPN
1	0x4
1	0x5
0	
0	
1	0x1
1	0x2
1	0x3
0	
0	
1	0x6
0	
0	
1	0x7

**In this case, 2 level is no better**

valid	2 <sup>nd</sup> level page table
1	0x2000
1	0x1000
1	0x3000
1	0x4000

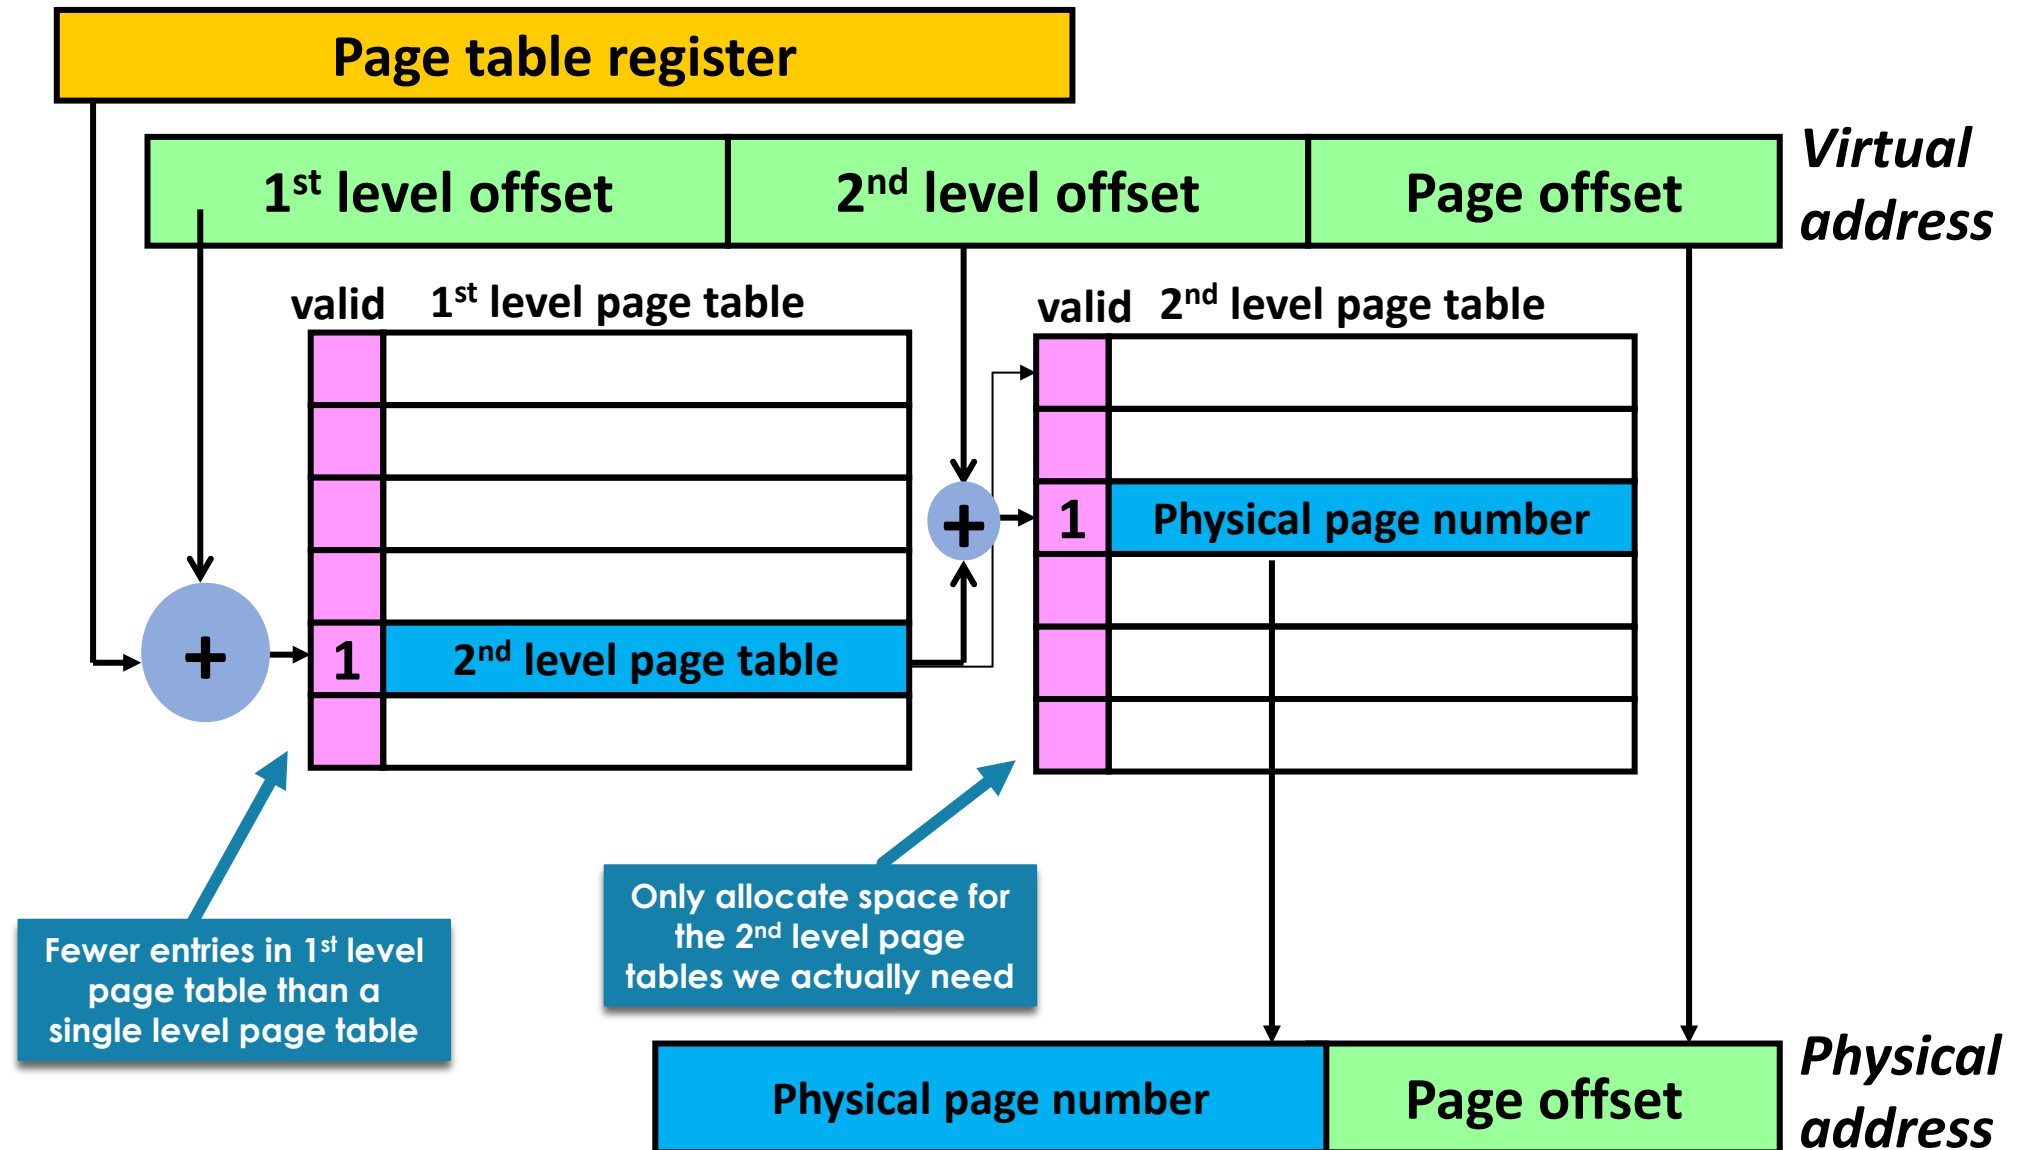
valid	PPN
1	0x1
1	0x2
1	0x3
0	

valid	PPN
1	0x4
1	0x5
0	
0	

valid	PPN
1	0x7
0	
0	
0	

valid	PPN
0	
1	0x6
0	
0	


# Hierarchical page table



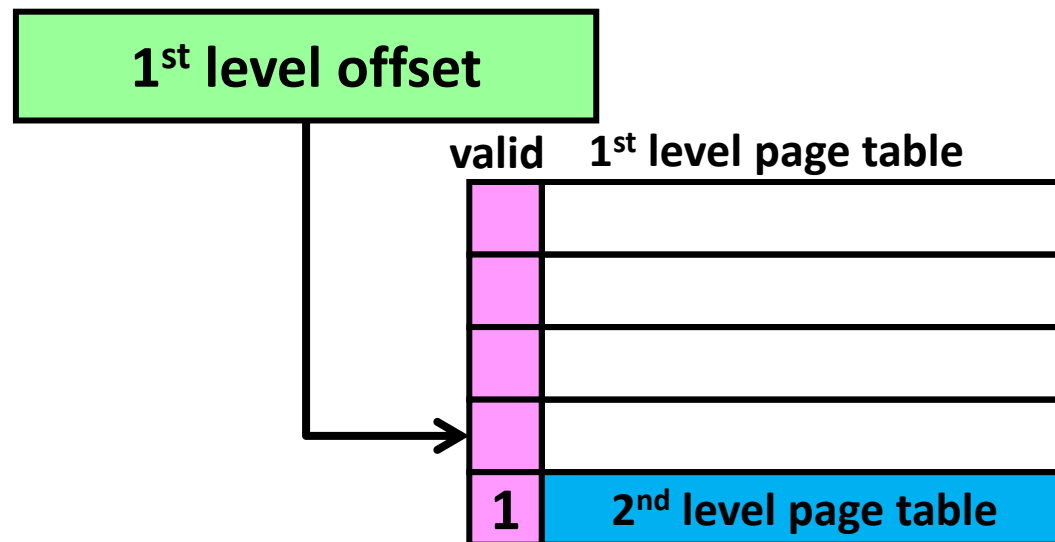
# Agenda

- Motivation for Multi-level Page Tables
- **Example architecture**
- Class Problem: Multi-Level VM Design
- VM Miscellanea

# Hierarchical page table – 32bit Intel x86




- How many bits in the virtual 1<sup>st</sup> level offset field? 10
- How many bits in the virtual 2<sup>nd</sup> level offset field? 10
- How many bits in the page offset? 12
- How many entries in the 1<sup>st</sup> level page table?  $2^{10}=1024$

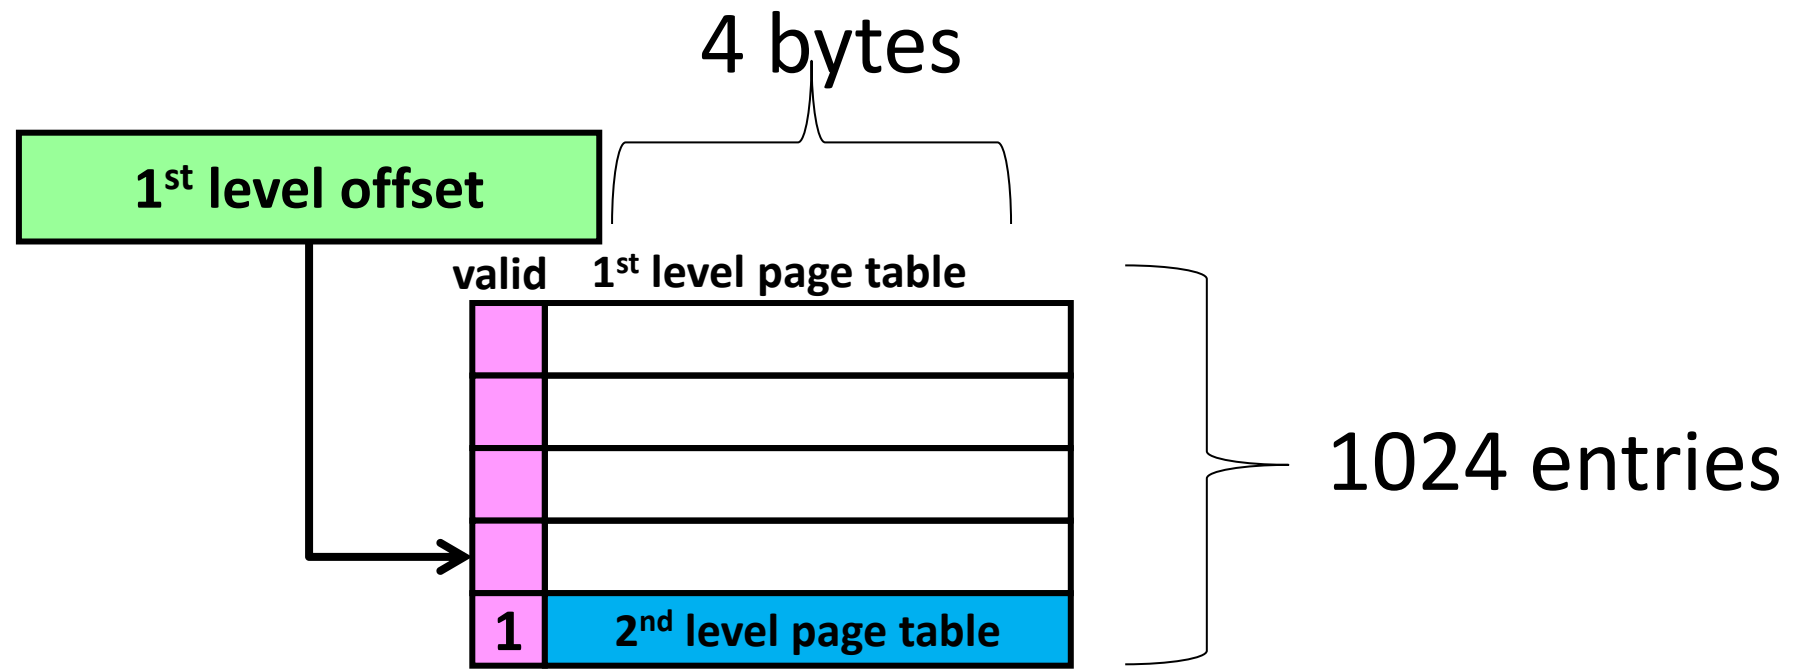




# Hierarchical page table – 32bit Intel x86

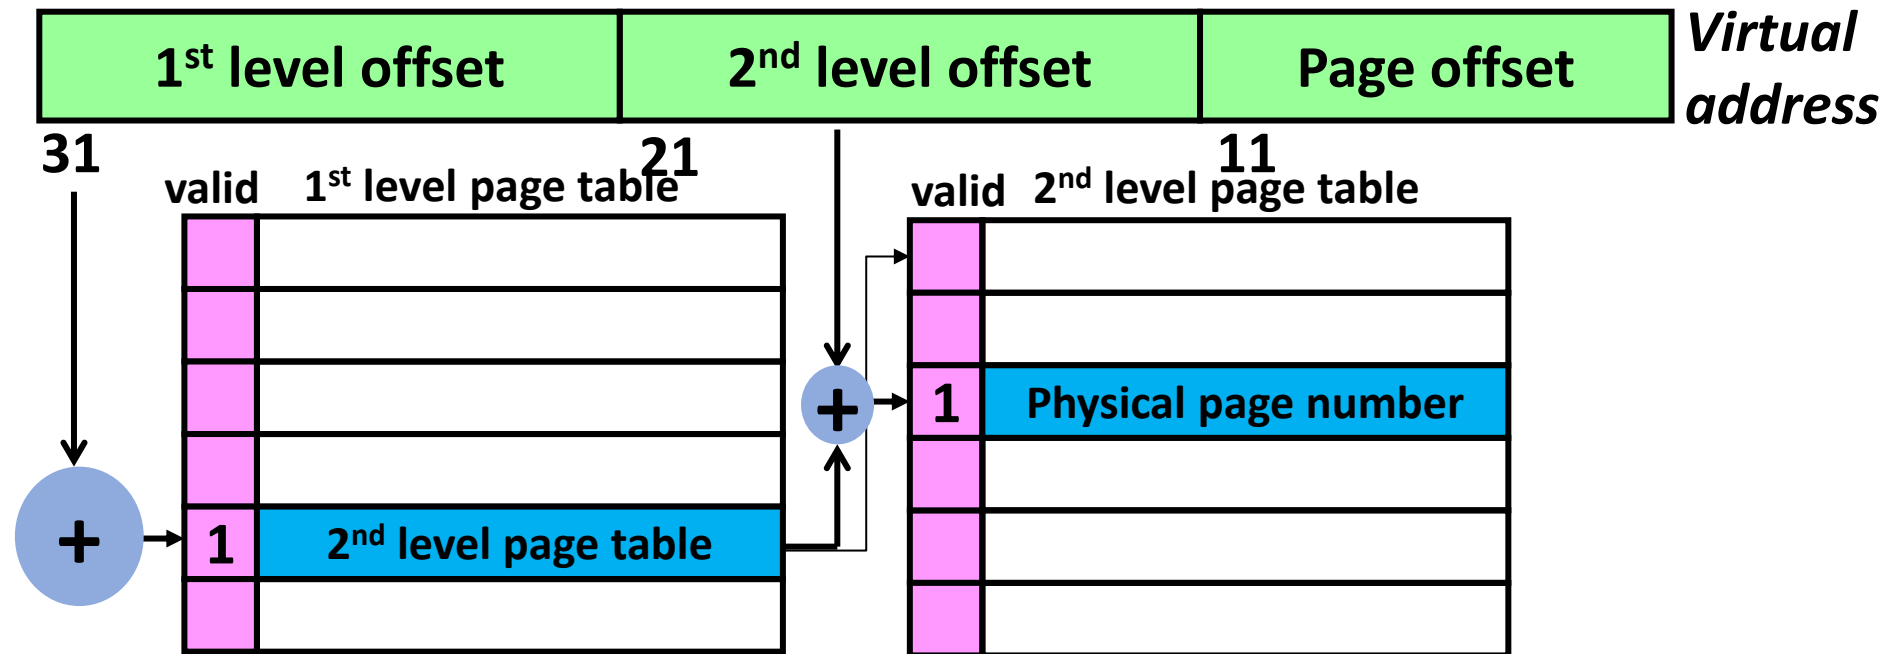


- Let's say physical address size + overhead bits is 4 bytes per entry
- Total size of 1<sup>st</sup> level page table
  - 4 bytes \* 1024 entries = 4 KB




# Hierarchical page table

- How many entries in the 2<sup>nd</sup> level of the page table?
  - $2^{10} = 1024$
- How many bytes for each VPN in a 2<sup>nd</sup> level table?
  - Let's round up to 4 bytes



# Hierarchical page table – 32bit Intel x86



- How many bits in the virtual 1<sup>st</sup> level offset field? 10
- How many bits in the virtual 2<sup>nd</sup> level offset field? 10
- How many bits in the page offset? 12
- How many entries in the 1<sup>st</sup> level page table?  $2^{10}=1024$
- How many bytes for each entry in the 1<sup>st</sup> level page table? 4
- How many entries in the 2<sup>nd</sup> level of the page table?  $2^{10}=1024$
- How many bytes for each entry in a 2<sup>nd</sup> level table?  $\sim 4$
- **What is the total size of the page table?**  
(here  $n$  is number of valid entries in the 1<sup>st</sup> level page table)  **$4K+n*4K$**

# Agenda

- Motivation for Multi-level Page Tables
- Example architecture
- **Class Problem: 32bit Intel x86**
- VM Miscellanea

## Class Problem (32 bit x86)

- What is the least amount of memory that could be used? When would this happen?
- What is the most memory that could be used? When would this happen?
- How much memory is used for this memory access pattern:  
0x00000ABC  
0x00000ABD  
0x10000ABC  
0x20000ABC



# Class Problem (32 bit x86)

- What is the least amount of memory that could be used? When would this happen?
  - 4KB for 1<sup>st</sup> level page table. Occurs when no memory has been accessed (before program runs)
- What is the most memory that could be used? When would this happen?
  - 4KB for 1<sup>st</sup> level page table  
+ 1024\*4KB for all possible 2<sup>nd</sup> level page tables  
= 4100KB (which slightly greater than 4096KB)
  - Occurs when program uses all virtual pages (=  $2^{20}$  pages)
    - But it can also happen under other circumstances. When?

# Class Problem (32 bit x86)

- How much memory is used for this memory access pattern:

0x00000ABC // Page fault

0x00000ABD

0x10000ABC // Page fault

0x20000ABC // Page fault

- 4KB for 1<sup>st</sup> level page table + 3\*4KB for each 2L page table  
= 16 KB

# Agenda

- Motivation for Multi-level Page Tables
- Example architecture
- Class Problem: 32bit Intel x86
- **Class Problem: Multi-Level VM Design**
- VM Miscellanea



# Class Problem – Multi-level VM

- Design a two-level virtual memory system of a byte addressable processor with **24-bit long addresses**. No cache in the system. **256Kbytes of memory** installed, and no additional memory can be added.
  - **Virtual memory page: 512 Bytes**. Each page table entry must be an integer number of bytes, and must be the smallest size required to fit the physical page number + 1 bit to mark valid-entry
  - We want each second-level page table to fit exactly in one memory page, and 1<sup>st</sup> level page table entries are 3 bytes each (a memory address pointer to a 2L page table).
- Compute:
  - Number of entries in each 2<sup>nd</sup> level page table;
  - Number of virtual address bits used to index the 2<sup>nd</sup> level page table;
  - Number of virtual address bits used to index the 1<sup>st</sup> level page table;
  - Size of the 1<sup>st</sup> level page table.

# Class Problem – Multi-level VM

Page Offset: 9 bits (512B page size)

**Physical address = 18b (256KB Mem size)**

**Physical page number = 18b (256KB mem size) – 9b (offset)**

Physical page number = 9b	Page offset = 9b
---------------------------	------------------

2<sup>nd</sup> level page table entry size: 9b (physical page number) + 1b = ~ 2 bytes

2<sup>nd</sup> level page table **fits exactly in 1 page**

#entries in 2<sup>nd</sup> level page table is 512 bytes / 2 bytes = 256

#entries in 2<sup>nd</sup> level page table = 256 → Virtual page bits = 8b

Virtual 1<sup>st</sup> level page bits = 24 – 8 – 9 = 7b

1<sup>st</sup> level page table size =  $2^7 * 3 \text{ bytes} = 384\text{B}$

1 <sup>st</sup> level = 7b	2 <sup>nd</sup> level = 8b	Page offset = 9b
----------------------------	----------------------------	------------------

**Virtual address = 24b**

# Agenda

- Motivation for Multi-level Page Tables
- Example architecture
- Class Problem: Multi-Level VM Design
- **VM Miscellanea**

# Page Replacement Strategies

- Page table indirection enables a fully associative mapping between virtual and physical pages.
- How do we implement LRU in OS?
  - True LRU is expensive, but LRU is a heuristic anyway, so approximating LRU is fine
  - Keep a “***accessed***” ***bit per page***, cleared occasionally by the operating system. Then pick any “unaccessed” page to evict

# Other VM Translation Functions

- Page data location
  - Physical memory, disk, uninitialized data
- Access permissions
  - Read only pages for instructions
  - **This is how your system detects segmentation faults**
- Gathering access information
  - Identifying dirty pages by tracking stores
  - Identifying accesses to help determine LRU candidate

# Next time

- Speeding up virtual memory