

EECS 370

Virtual Memory Basics



Announcements

- Lab
 - Assignment Due Wednesday
 - Pre-lab quiz Thursday
- P4 out
 - Check out simulator on website (not same as project... doesn't cache instrs)

The screenshot displays a simulator interface with three main panels:

- Instructions ?**: Contains a "Load Instructions" dropdown, "Verify", "Step >", and ">>" buttons, and a large empty text area for instruction input.
- Data Cache ?**: Shows a 4-set cache (Set 0 to Set 3) with four "INV" (Invalid) entries per set. Below the cache, configuration options are visible:
 - Cache Configuration ?**: Set to "Custom".
 - # Hits**: 0.
 - # Misses**: 0.
 - Block Size (Words)**: 2.
 - # Sets**: 4.
 - # Blocks/Set**: 4.
- Memory (64 W) ?**: A vertical stack of 64 words, indexed 0 to 5. Each index shows a color-coded bar (green, yellow, red) representing different memory states or data.

Poll: What is the CPI increase over base (1) due to hazards?

Practice Problem: CPI with caches

The *blaster* application run on the LC2k with full data forwarding and all branches predicted not-taken has the following instruction frequencies:

45% R-type 20% Branches 15% Loads 20% Stores

In *blaster*, 40% of branches are taken and 50% of LWs are followed by an immediate use.

The I-cache has a miss rate of 3% and the D-cache has a miss rate of 6% (no overlapping of misses). On a miss, the main memory is accessed and has a latency of 100 ns. The clock frequency is 500 MHz.

Problem Solution

The *blaster* application run on the LC2k with full data forwarding and all branches predicted not-taken has the following instruction frequencies:

45% R-type 20% Branches 15% Loads 20% Stores

In *blaster*, 40% of branches are taken and 50% of LWs are followed by an immediate use.

The I-cache has a miss rate of 3% and the D-cache has a miss rate of 6% (no overlapping of misses). On a miss, the main memory is accessed and has a latency of 100 ns. The clock frequency is 500 MHz.

What is the CPI of *blaster* on the LC2k?

Stalls per cache miss = 100 ns / 2ns = 50 cycles (500 Mhz → 2ns cycle time)

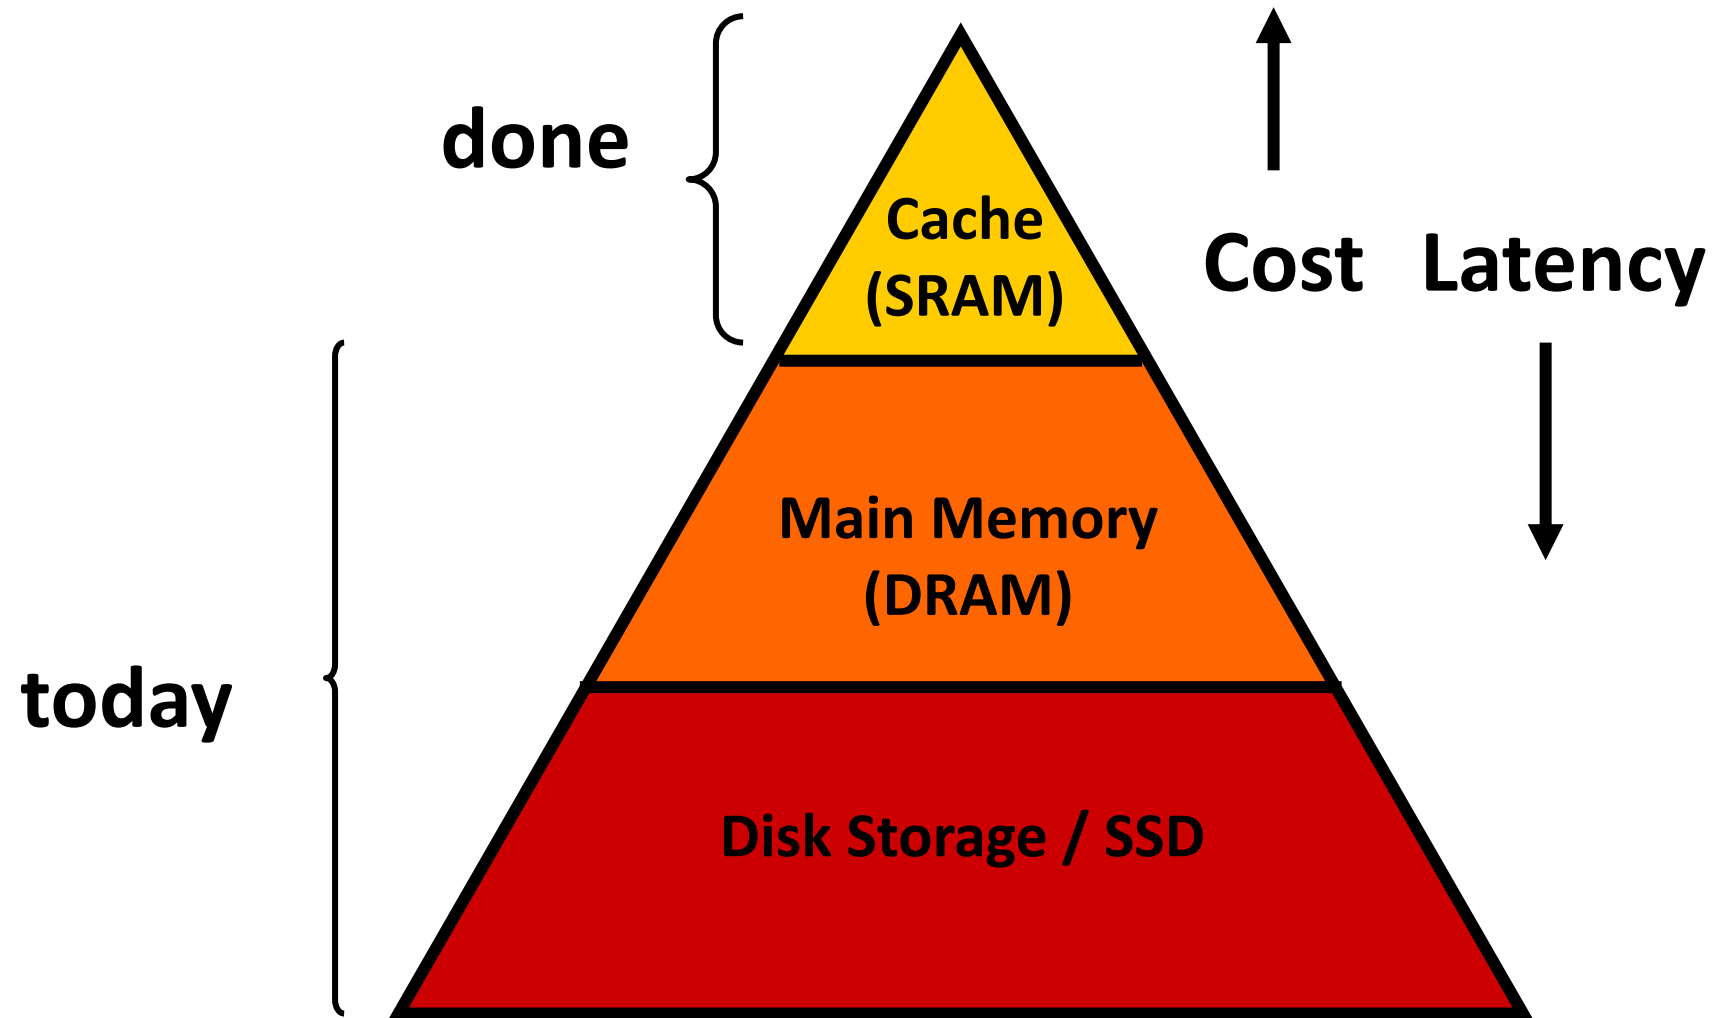
CPI = 1 + data hazard stalls + control hazard stalls + icache stalls + dcache stalls

CPI = 1 + 0.15*0.50*1 + 0.20*0.40*3 + 1*0.03*50 + 0.35*0.06*50

Agenda

- **Motivation**
- Virtual Memory Principles
- Page Tables
- Class Problem

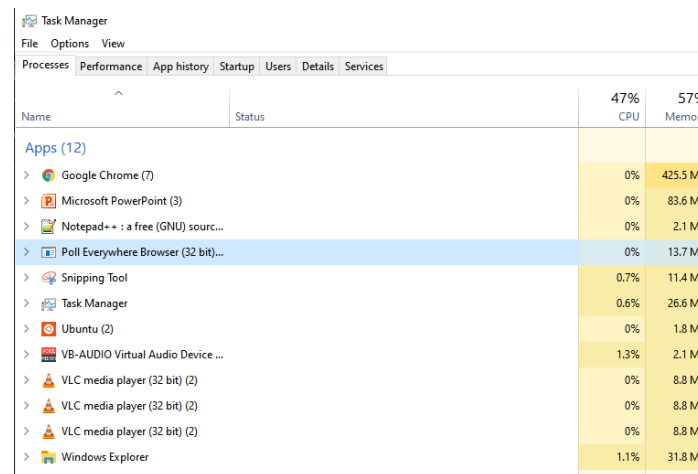
Storage Hierarchy



Memory: Two Issues

1. We've been working with the abstraction that all programs have full, private access to memory

- But in practice, multiple programs run at the same time!

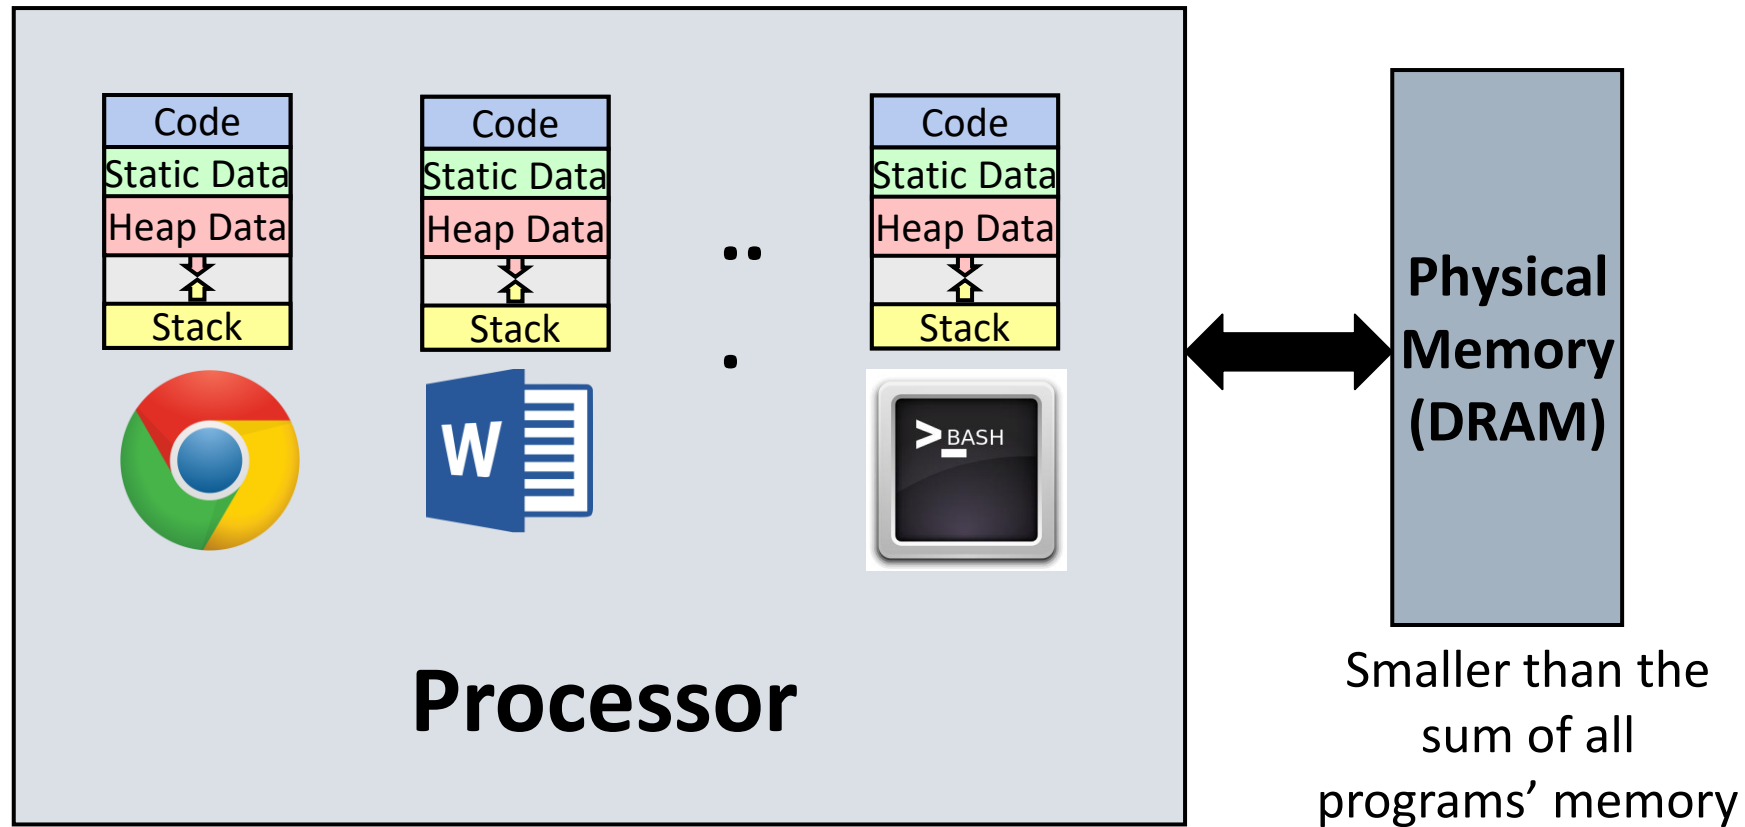


The screenshot shows the Windows Task Manager Performance tab. The 'Memory' section is selected, showing a usage of 57%. Below this, a list of running applications is displayed with their respective CPU and memory usage. The 'Poll Everywhere Browser (32 bit)...' is highlighted.

Name	Status	CPU	Memory
Apps (12)			
> Google Chrome (7)		0%	425.5 MB
> Microsoft PowerPoint (3)		0%	83.6 MB
> Notepad++ : a free (GNU) sourc...		0%	2.1 MB
> Poll Everywhere Browser (32 bit)...		0%	13.7 MB
> Snipping Tool		0.7%	11.4 MB
> Task Manager		0.6%	26.6 MB
> Ubuntu (2)		0%	1.8 MB
> VB-AUDIO Virtual Audio Device ...		1.3%	2.1 MB
> VLC media player (32 bit) (2)		0%	8.8 MB
> VLC media player (32 bit) (2)		0%	8.8 MB
> VLC media player (32 bit) (2)		0%	8.8 MB
> Windows Explorer		1.1%	31.8 MB

- What happens if two programs try to write to the same memory address??

Revisit real system view—multitasking



Memory: Two Issues

2. Even if only one program is running, modern computers have 48-64 bit address spaces!

- No computer actually has 18 exabytes (18 billion GBs)
- What if a computer tries to write to address 0xFFFF...FFFF
- Should it just crash??

Memory: Two Issues


- Modern systems use the same solution for both problems: **virtual memory**
 - In a nutshell, each program thinks it has full, private access to memory (it can safely index any address from 0x0-0xFFF...FFF)
 - Hardware and software transparently maps these addresses to distinct addresses in DRAM and in hard disk / SSD
 - Focus for the next 3 lectures

Agenda

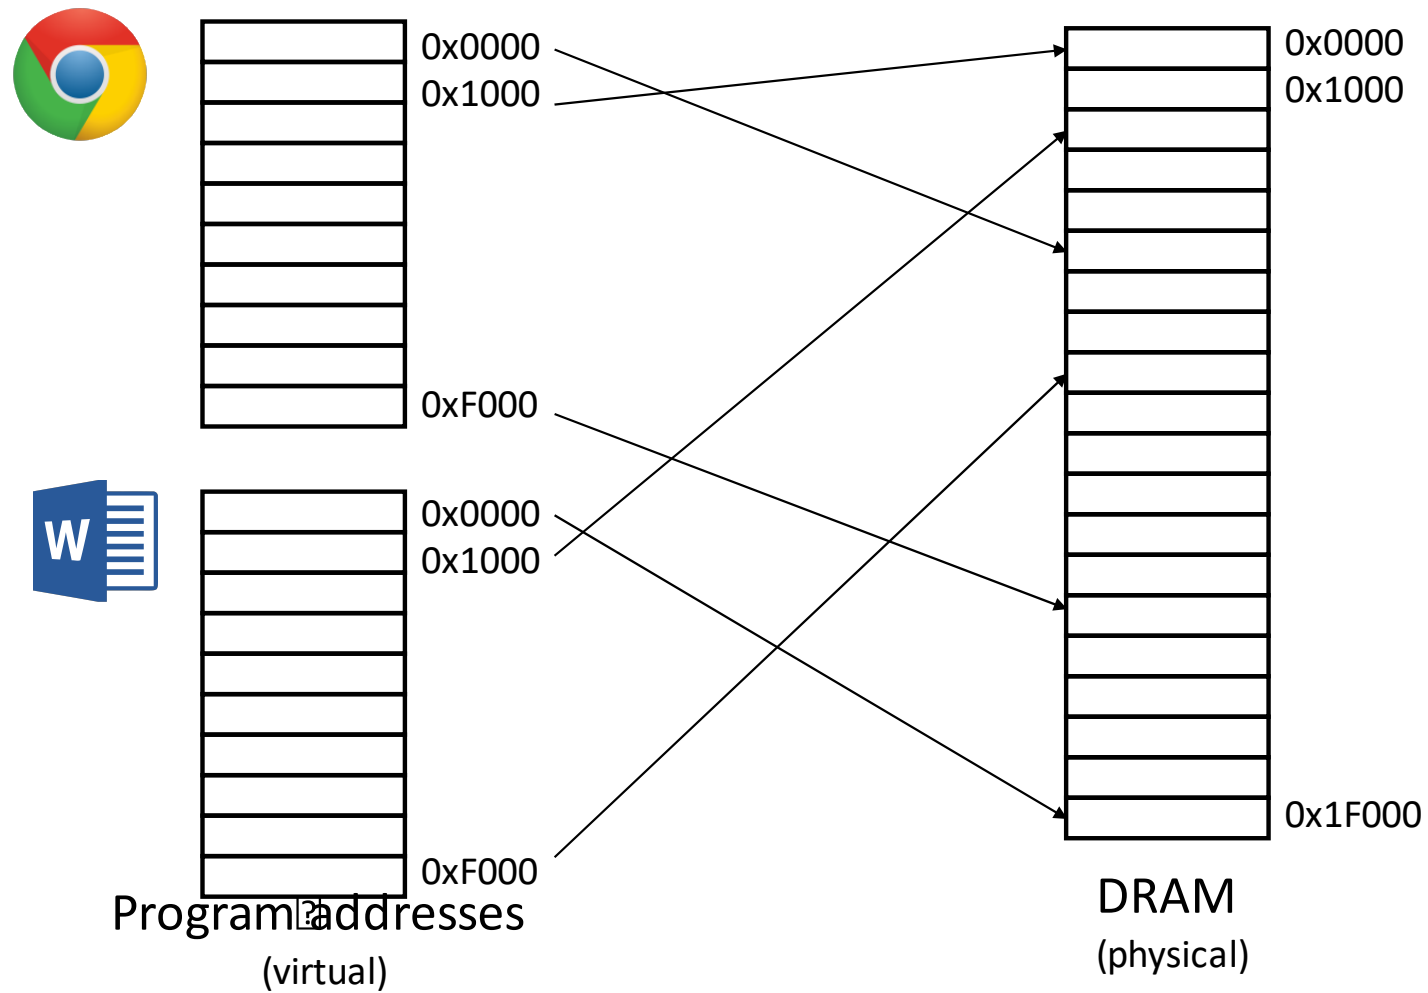
- Motivation
- **Virtual Memory Principles**
- Page Tables
- Class Problem

Basics of Virtual Memory

- Any time you see the word virtual in computer science and architecture it means “using a level of indirection”.
 - Examples?
- Virtual memory hardware changes the virtual address the programmer sees into the physical one the memory chips see.



Virtual Memory




How to Translate Addresses?

- Address Translation is not done entirely in hardware
- We'll get help in software via the **operating system**
- The operating system is a special set of programs
 - Always running (after the system boots)
 - Is in charge of **managing the hardware resources** for all other running programs
 - E.g. initializing memory for a starting program, managing the file system, choosing when a particular program gets to run...
 - ... and translating virtual addresses into physical addresses!
- OS handles address translation by maintaining a data structure in main memory: **the page table**


Virtual memory terminology

- Memory is divided into fixed-size chunks called **Pages** (e.g., 4KB for x86)
 - Size of physical page = size of virtual page
 - A virtual address consists of
 - A virtual page number
 - A page offset field (low order bits of the address)
 - Translating a virtual address into a physical address only requires translating the page numbers
 - The page offset will stay the same

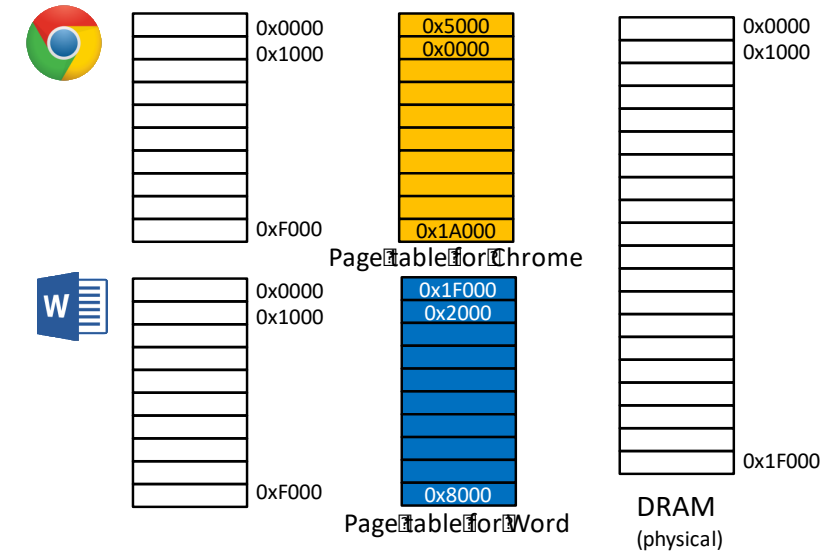
Virtual address



Physical address



Page Table



- Translate page numbers using **page tables**
- Contains address translation information, i.e. virtual page # → physical page #
- Each process has its own page table
 - Maintained by operating system (OS)
- Page tables themselves are kept in memory by OS, and OS knows the physical address of the page tables
 - No address translation is required by the OS for accessing the page tables


Poll: What is the cost of this scheme?

(select all that apply)

- a) Uses up more memory
- b) Takes longer to do loads and stores
- c) Fewer addresses accessible by program
- d) None of the above

Why Pages?

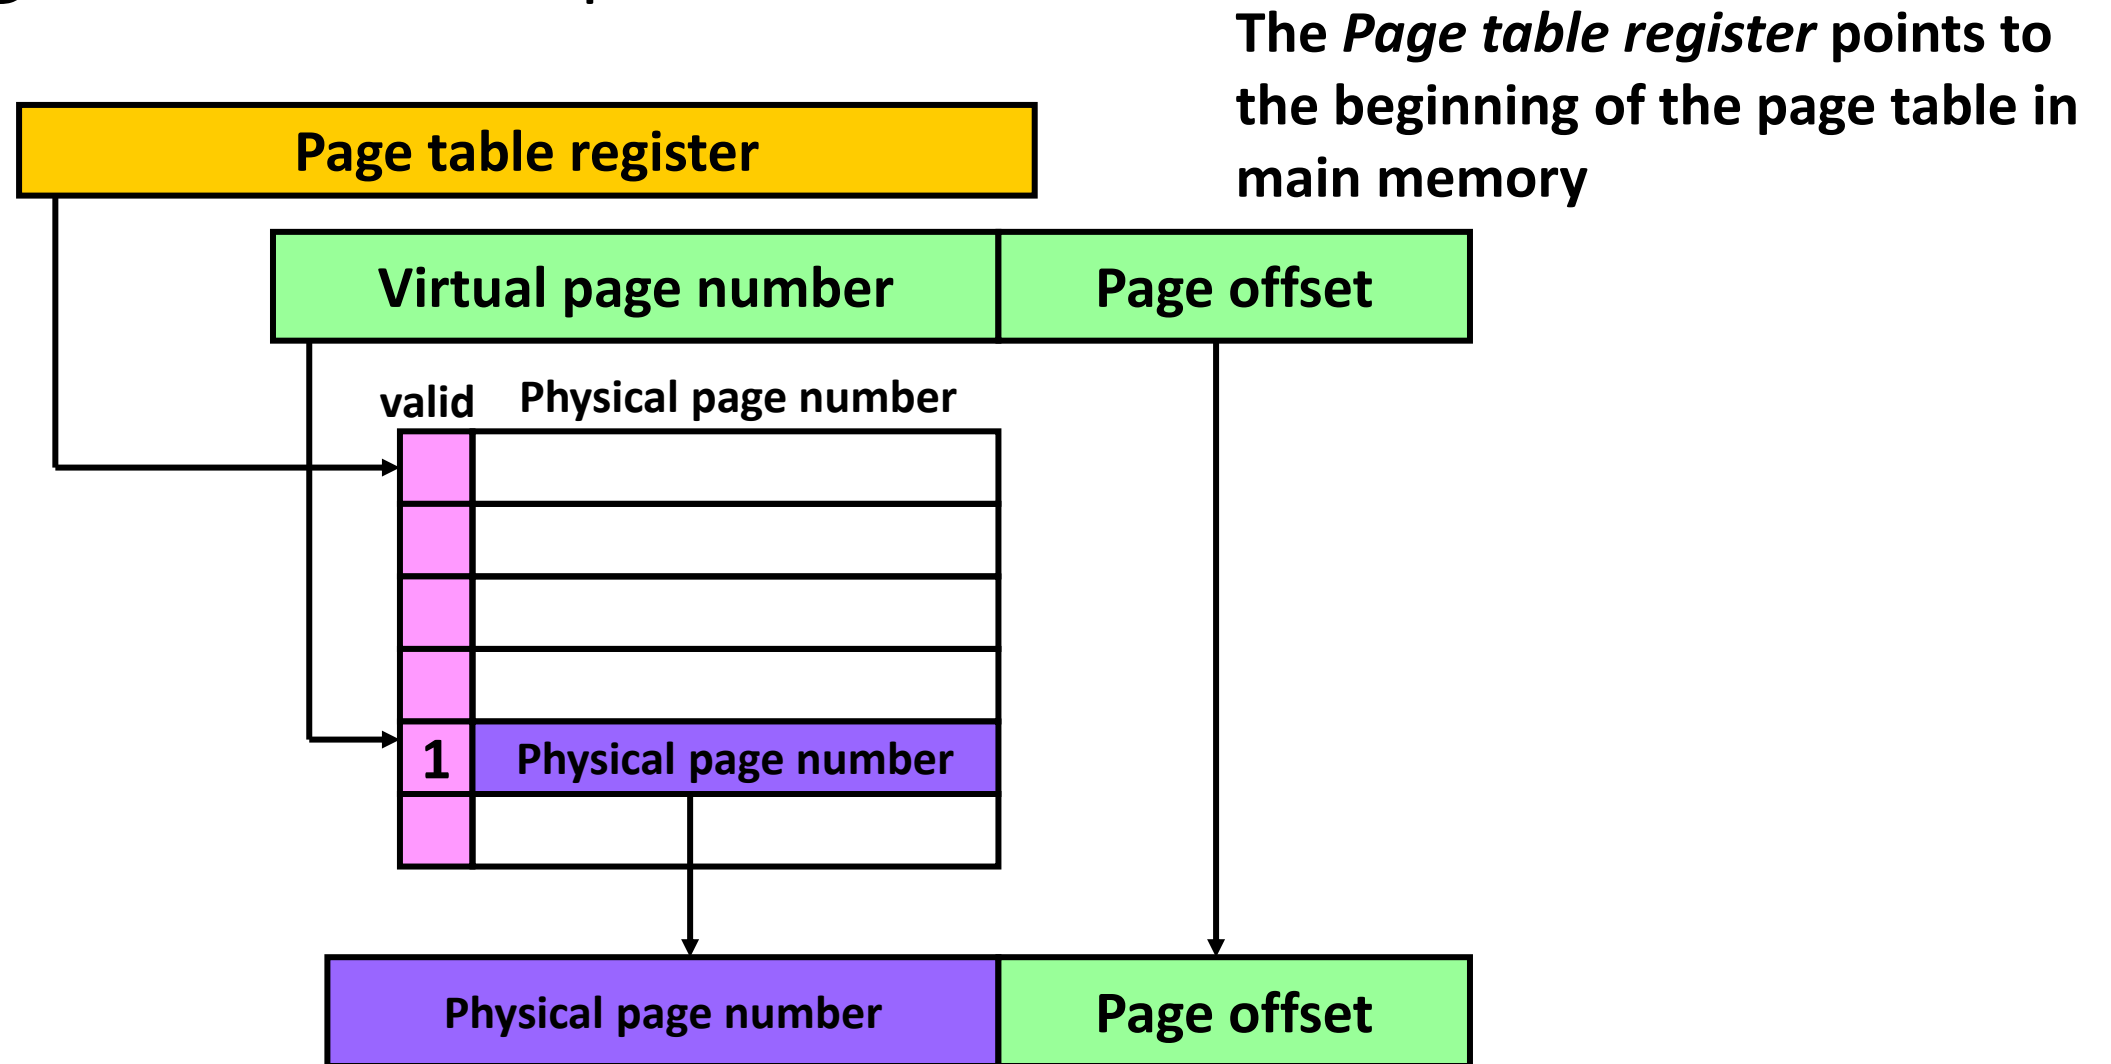
- Why have the idea of "pages"? Why not just have a mapping for each individual address?
 - Equivalent to asking: "why not have pages of size 1 byte?"
 - Otherwise - need a mapping entry for every single element of memory
 - The mapping data would take up as much space as the actual program data!
 - Also would screw up spatial locality of cache blocks (things contiguous in virtual memory wouldn't be contiguous in physical memory)



Agenda

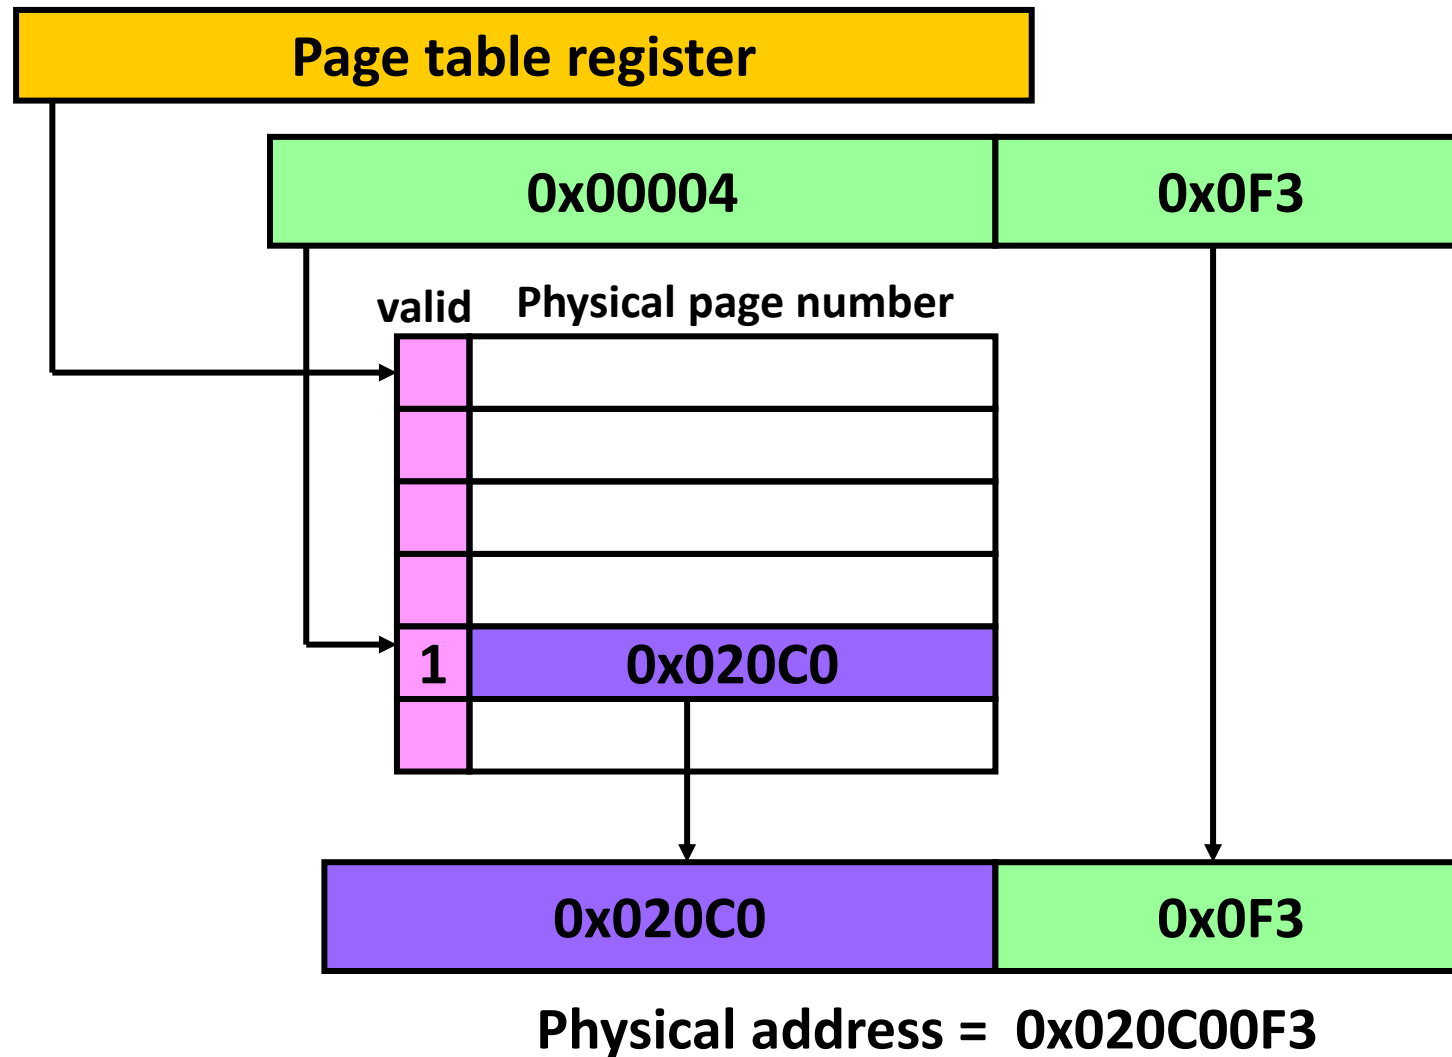
- Motivation
- Virtual Memory Principles
- **Page Tables**
- Class Problem

Page table components



Page table components - Example

Virtual address = 0x000040F3



Virtual Memory Goals

- VM should provide the following 3 capabilities to the programs:

- 1. Transparency**

- Don't need to know how other programs are using memory


- 2. Protection**

- No program can access the data of any other program


- 3. Programs not limited by DRAM capacity**

- Each program can have more data than DRAM size


1-2: How to achieve transparency & protection?



Page Table for Chrome



Page Table for Word



DRAM
(physical)

Ensure that page table entries
across different programs
don't contain the same
physical page numbers

3. How to be **not limited by DRAM capacity**?

- Use disk as temporary space in case memory capacity is exhausted
 - This temporary space in disk is called **swap partition** in Linux-based systems
 - For fun check swap space in a linux system by:
\$: **top**


```
Tasks: 662 total,  1 running, 661 sleeping,  0 stopped,  0 zombie
%Cpu(s):  0.1 us,  0.0 sy,  0.0 ni, 99.8 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
KiB Mem: 32704372 total, 10813444 used, 21890928 free, 1018840 buffers
KiB Swap: 35162108 total,  89248 used, 35072860 free. 7053764 cached Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
60256	nehaag	20	0	25356	3356	2444	R	6.0	0.0	0:00.02	top
1	root	20	0	38424	9040	2780	S	0.0	0.0	1:56.96	init
2	root	20	0	0	0	0	S	0.0	0.0	0:02.21	kthreadd
3	root	20	0	0	0	0	S	0.0	0.0	6:20.75	ksoftirqd/0

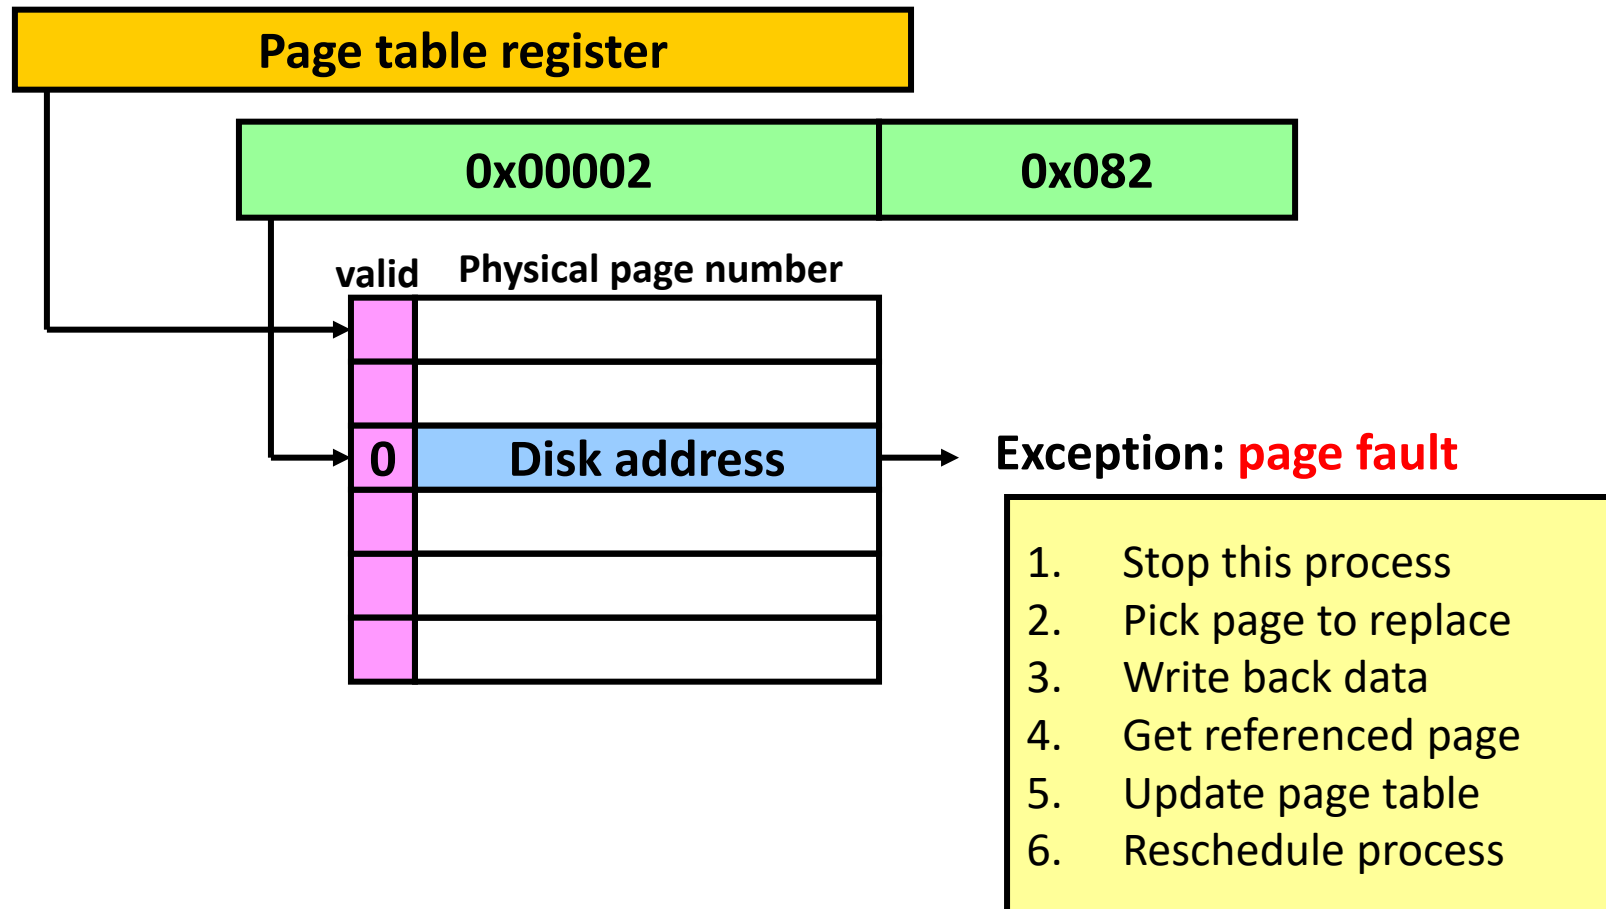
2. How to be **not limited by DRAM capacity?**

- We can mark a page table entry as “Invalid”, indicating that the data for that page doesn’t exist in main memory, but instead is located on the disk
- Looking up a page table entry that corresponds to disk is called a **page fault**

2. How to be **not limited by DRAM capacity?**



Page faults



How do we find it on disk?

- That is not a hardware problem! Go take EECS 482! 😊
- This is the operating system's job. Most operating systems partition the disk into logical devices (C: , D: , /home, etc.)
- They also have a hidden partition to support the disk portion of virtual memory
 - **Swap partition** on UNIX machines
 - You then index into the correct page in the swap partition.

Agenda

- Motivation
- Virtual Memory Principles
- Page Tables
- **Class Problem**

Class Problem

- Given the following:
 - 4KB page size, physical memory of 16KB, page table stored in physical page 0 and can never be evicted, 20 bit, byte-addressable virtual address space.
 - The page table initially has virtual page 0 in physical page 1, virtual page 1 in physical page 2 and no valid data in other physical pages.
- Fill in the table on the next slide for each reference
 - Note: like caches we'll use LRU when we need to replace a page.

4KB page size,
physical memory of 16KB,
page table stored in physical
page 0 and can never be
evicted, 20 bit, byte-
addressable virtual address
space.

The page table initially has
virtual page 0 in physical
page 1, virtual page 1 in
physical page 2 and no valid
data in other physical pages.

Poll: How many hex digits
should the page number
be?

Class Problem (continued)

Virt addr	Virt page	Page fault?	Phys addr
0x00F0C			
0x01F0C			
0x20F0C			
0x00100			
0x00200			
0x30000			
0x01FFF			
0x00200			



Class Problem (continued)

4KB page size,
physical memory of 16KB,
page table stored in physical
page 0 and can never be
evicted, 20 bit, byte-
addressable virtual address
space.

The page table initially has
virtual page 0 in physical
page 1, virtual page 1 in
physical page 2 and no valid
data in other physical pages.

Virt addr	Virt page	Page fault?	Phys addr
0x00F0C	0x0	N	0x1F0C
0x01F0C	0x1	N	0x2F0C
0x20F0C	0x20	Y (into 3)	0x3F0C
0x00100	0x0	N	0x1100
0x00200	0x0	N	0x1200
0x30000	0x30	Y (into 2)	0x2000
0x01FFF	0x1	Y (into 3)	0x3FFF
0x00200	0x0	N	0x1200

Size of the page table

- How big is a page table entry?
 - For 32-bit virtual address:
 - If the machine can support $1\text{GB} = 2^{30}$ bytes of physical memory and we use pages of size $4\text{KB} = 2^{12}$,
 - then the physical page number is $30-12 = 18$ bits.
Plus another **valid bit** + other useful stuff (**read only, dirty, etc.**)
 - Let say about 3 bytes.
- How many entries in the page table?
 - 1 entry per virtual page
 - ARM virtual address is 32 bits – 12 bit page offset = 20
 - Total number of virtual pages = 2^{20}
- Total size of page table = Number of virtual pages
 - * Size of each page table entry
 - = $2^{20} \times 3$ bytes ~ 3 MB

Next time

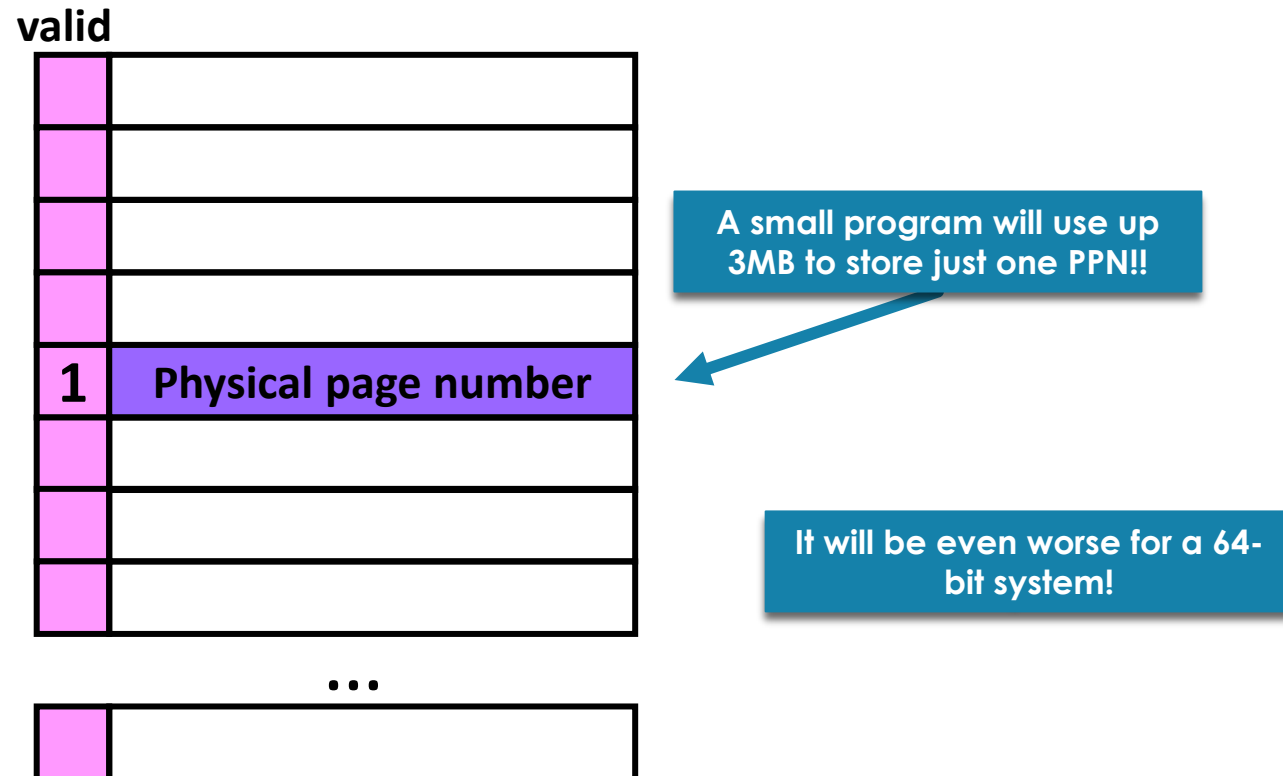
- Improving Virtual Memory Design
 - Multi-level page-tables

Extra Slides



How can you organize the page table?

1. Single-level page table occupies continuous region in physical memory
 - Previous example always takes 3MB regardless of how much virtual memory is used



Multi-Level Page Table

- Only allocate second (and later) page tables when needed
- Program starts: everything is invalid, only first level is allocated

valid	2 nd level page table
0	
0	
0	
0	

- As we access more, second level page tables are allocated

valid	2 nd level page table
1	0x1000
1	0x3500

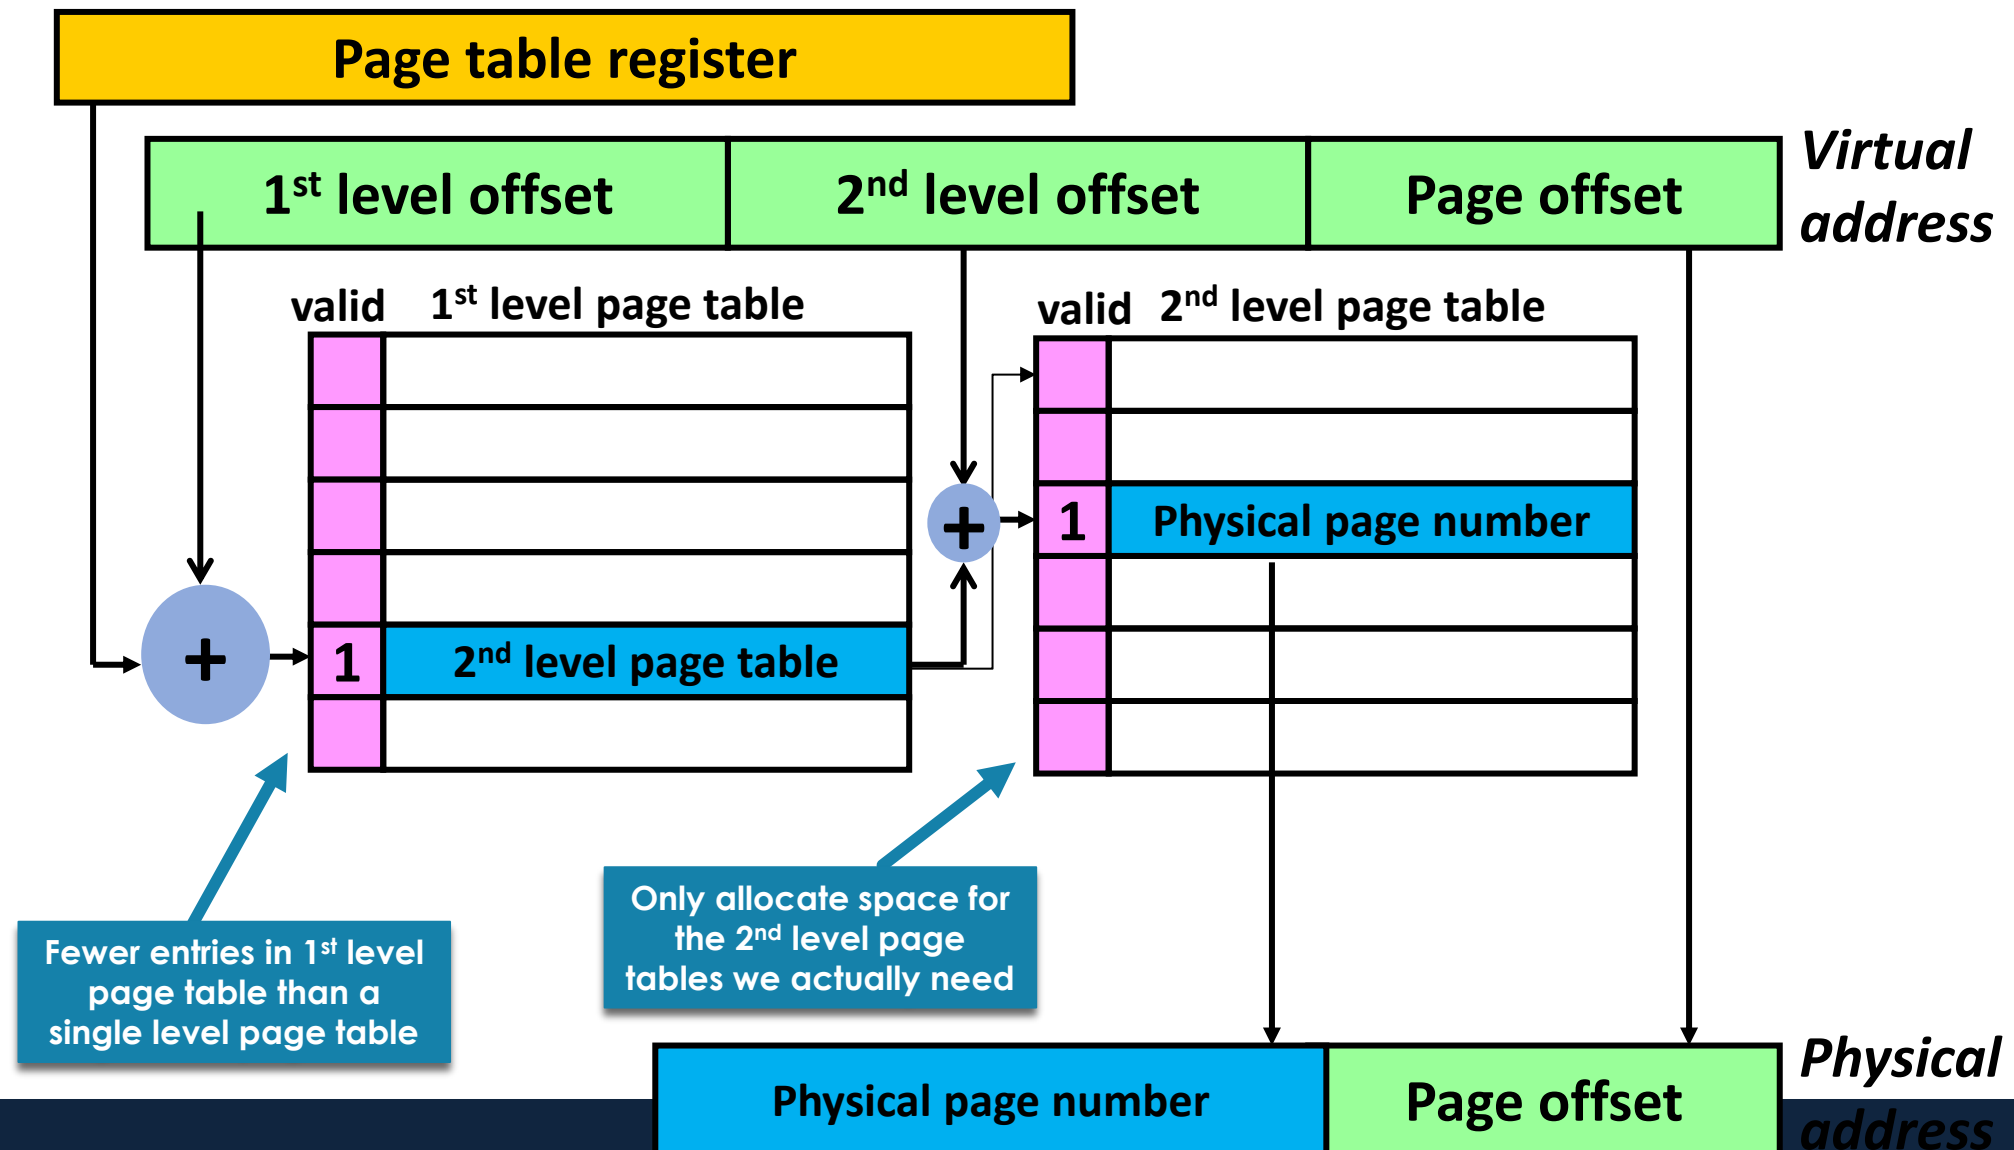
valid	PPN
1	0x1
1	0x6
1	0x2
1	0x1f

valid	PPN
1	0x9a
1	0x3
0	
1	0xff

Multi-level: Size is proportional to amount of memory used

Common case: most programs use small portion of virtual memory space


Hierarchical page table



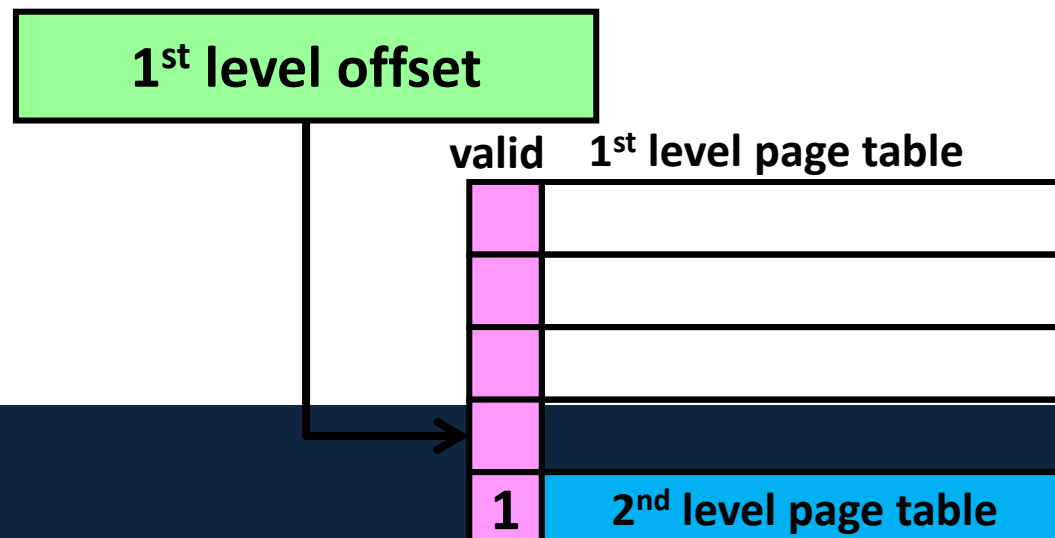
Agenda

- Motivation for Multi-level Page Tables
- **Example architecture**
- Class Problem: 32bit Intel x86
- Class Problem: Multi-Level VM Design
- VM Miscellanea


Hierarchical page table – 32bit Intel x86




- How many bits in the virtual 1st level offset field? 10
- How many bits in the virtual 2nd level offset field? 10
- How many bits in the page offset? 12
- How many entries in the 1st level page table? $2^{10}=1024$



Hierarchical page table – 32bit Intel x86

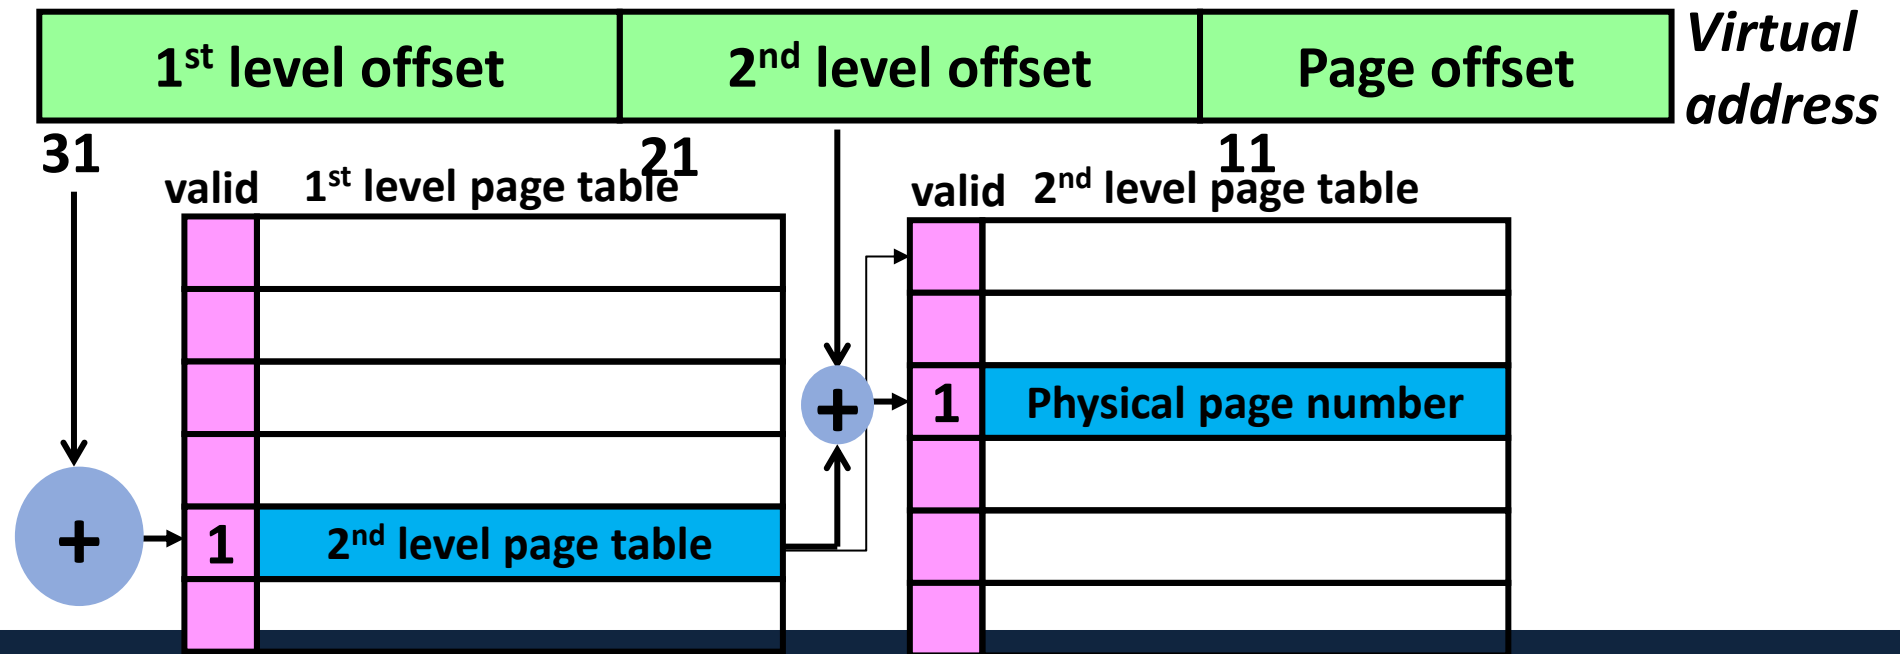


- Let's say physical address size + overhead bits is 4 bytes per entry
- Total size of 1st level page table
 - 4 bytes * 1024 entries = 4 KB




Hierarchical page table

- How many entries in the 2nd level of the page table?
 - $2^{10} = 1024$
- How many bytes for each VPN in a 2nd level table?
 - Let's round up to 4 bytes



Hierarchical page table – 32bit Intel x86



- How many bits in the virtual 1st level offset field? 10
- How many bits in the virtual 2nd level offset field? 10
- How many bits in the page offset? 12
- How many entries in the 1st level page table? $2^{10}=1024$
- How many bytes for each entry in the 1st level page table? 4
- How many entries in the 2nd level of the page table? $2^{10}=1024$
- How many bytes for each entry in a 2nd level table? ~ 4
- **What is the total size of the page table?**
(here n is number of valid entries in the 1st level page table) **$4K+n*4K$**