

**Poll: What's your favorite
Thanksgiving food?**

- a) Turkey
- b) Potatoes
- c) Pie
- d) Sweet, sweet Buckeye
tears

EECS 370

Making Virtual Memory Fast



Announcements

- Lab
 - Last lab!
 - Due next Wed
- P4
 - Last project!
 - Due Thur
- HW 4
 - Last homework!
 - Due Mon
- Final exam
 - ...Last exam!
 - Review lecture on Thursday
 - Wed @ 10:30 am

Review: Organizing the page table

2. Option 2: Use a multi-level page table

- Split up single-level page into multiple chunks
- Only allocate space for that chunk if it's non-empty (i.e. program uses some of those physical pages)
- Maintain a top-level page table which contains pointers to each of these chunks

valid	PPN
0	
0	
0	
0	
1	0x1
1	0x2
1	0x3
0	
0	
Single-level: Tons of wasted space!	
0	
0	
0	

Single-level: Tons of wasted space!

valid	2 nd level page table
0	
1	0x1000
0	
0	

valid	PPN
1	0x1
1	0x2
1	0x3
0	

Multi-level: No need to allocate 2nd level entries if all invalid

Review: Multi-Level Page Table

- Only allocate second (and later) page tables when needed
- Program starts: everything is invalid, only first level is allocated

Single Level

valid	2 nd level page table
0	
0	
0	
0	

Multi-level: Size is proportional to amount of memory used

- As we access more, second level page tables are allocated




valid	2 nd level page table
1	0x1000
1	0x3500

valid	PPN
1	0x1
1	0x6
1	0x2
1	0x1f

valid	PPN
1	0x9a
1	0x3
0	
1	0xff

Common case: most programs use small portion of virtual memory space

Review: Hierarchical page table



Class Problem – Multi-level VM

- Design a two-level virtual memory system of a byte addressable processor with **24-bit long addresses**. No cache in the system. **256Kbytes of memory** installed, and no additional memory can be added.
 - **Virtual memory page: 512 Bytes**. Each page table entry must be an integer number of bytes, and must be the smallest size required to fit the physical page number + 1 bit to mark valid-entry
 - We want each second-level page table to fit exactly in one memory page, and 1st level page table entries are 3 bytes each (a memory address pointer to a 2L page table).
- Compute:
 - Number of entries in each 2nd level page table;
 - Number of virtual address bits used to index the 2nd level page table;
 - Number of virtual address bits used to index the 1st level page table;
 - Size of the 1st level page table.

Class Problem – Multi-level VM

Class Problem – Multi-level VM

Page Offset: 9 bits (512B page size)

Physical address = 18b (256KB Mem size)

Physical page number = 18b (256KB mem size) – 9b (offset)

Physical page number = 9b	Page offset = 9b
---------------------------	------------------

2nd level page table entry size: 9b (physical page number) + 1b = ~ 2 bytes

2nd level page table **fits exactly in 1 page**

#entries in 2nd level page table is 512 bytes / 2 bytes = 256

#entries in 2nd level page table = 256 → Virtual page bits = 8b

Virtual 1st level page bits = 24 – 8 – 9 = 7b

1st level page table size = $2^7 * 3 \text{ bytes} = 384\text{B}$

1 st level = 7b	2 nd level = 8b	Page offset = 9b
----------------------------	----------------------------	------------------

Virtual address = 24b

Agenda

- **Table Look-aside Buffers (TLB)**
- Virtual Memory Walkthrough
- Cache Placement

Virtual Memory Performance

- To translate a virtual address into a physical address, we must first access the page table in physical memory
 - If it's an N-level page table, we must do N total loads before getting the physical page number
- Then we access physical memory again to get the data
 - A load instruction performs at least 2 memory reads
 - A store instruction performs at least 1 read and then a write
- Above lookups are **SLOW**

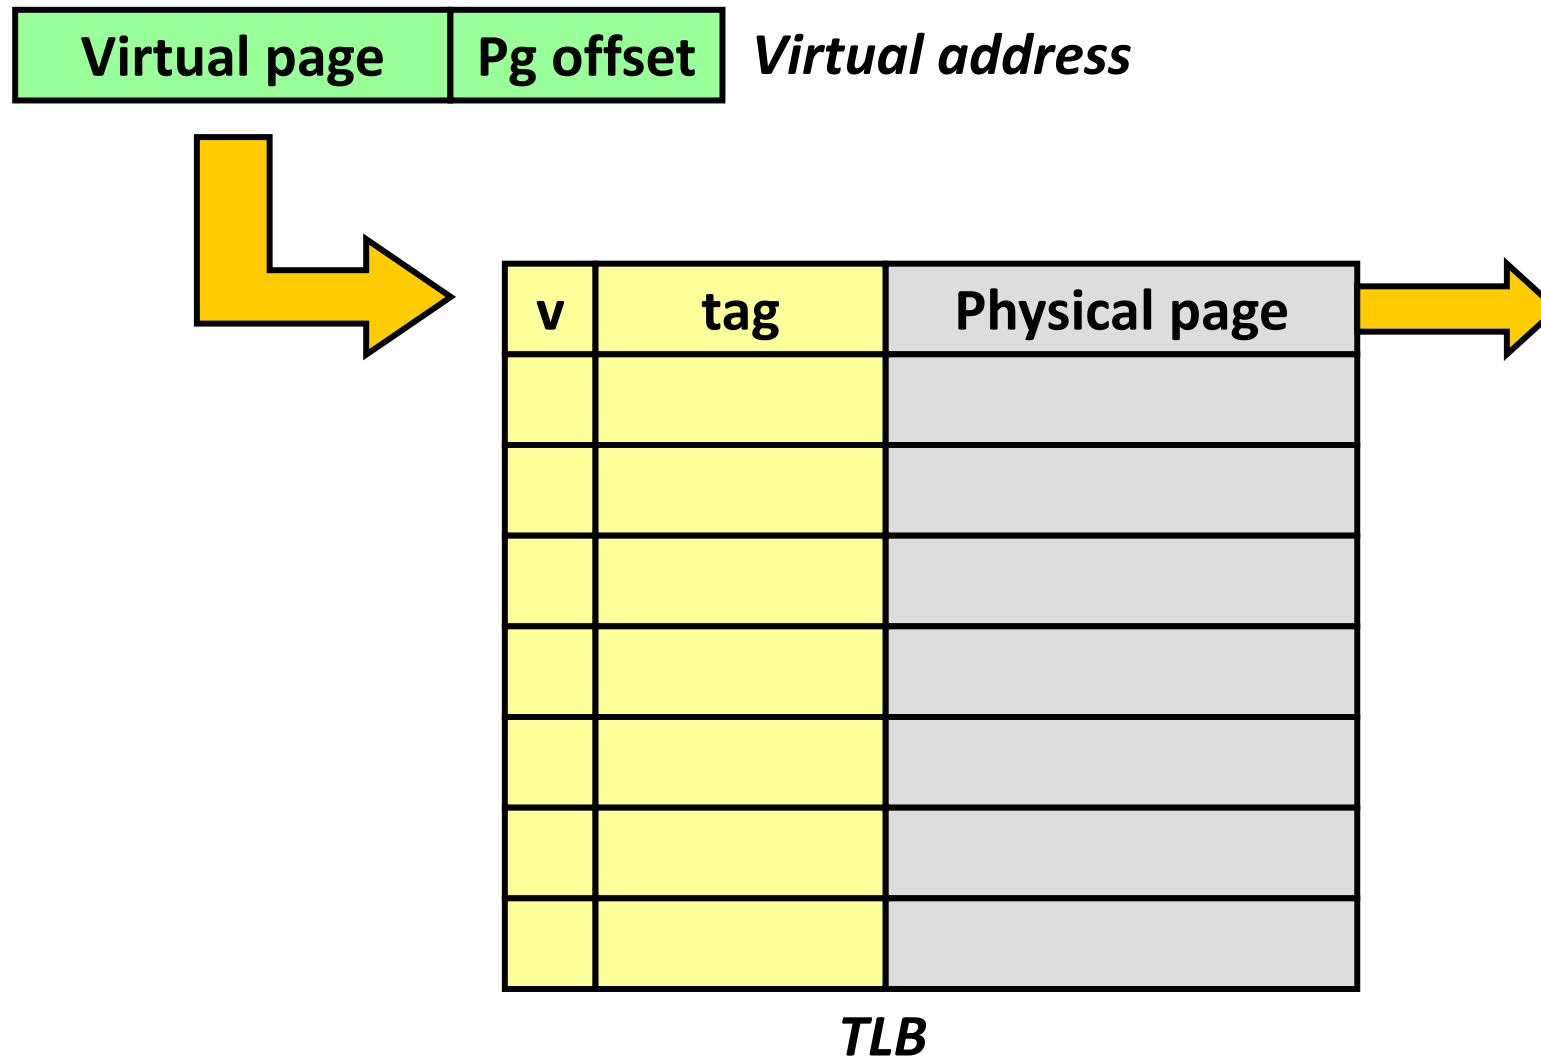
Translation look-aside buffer (TLB)

- We fix this performance problem by avoiding main memory in the translation from virtual to physical pages.
- Buffer common translations in a **Translation Look-aside Buffer (TLB)**, a fast cache memory dedicated to storing a small subset of valid V-to-P translations.
- 16-512 entries common.
- Generally has low miss rate ($< 1\%$).

Poll: Why can we make TLBs smaller than the cache hierarchy?

- a) Addresses are smaller than data
- b) TLB is accessed less frequently than caches
- c) Only need to store info about individual pages

Translation look-aside buffer (TLB)



Putting it all together

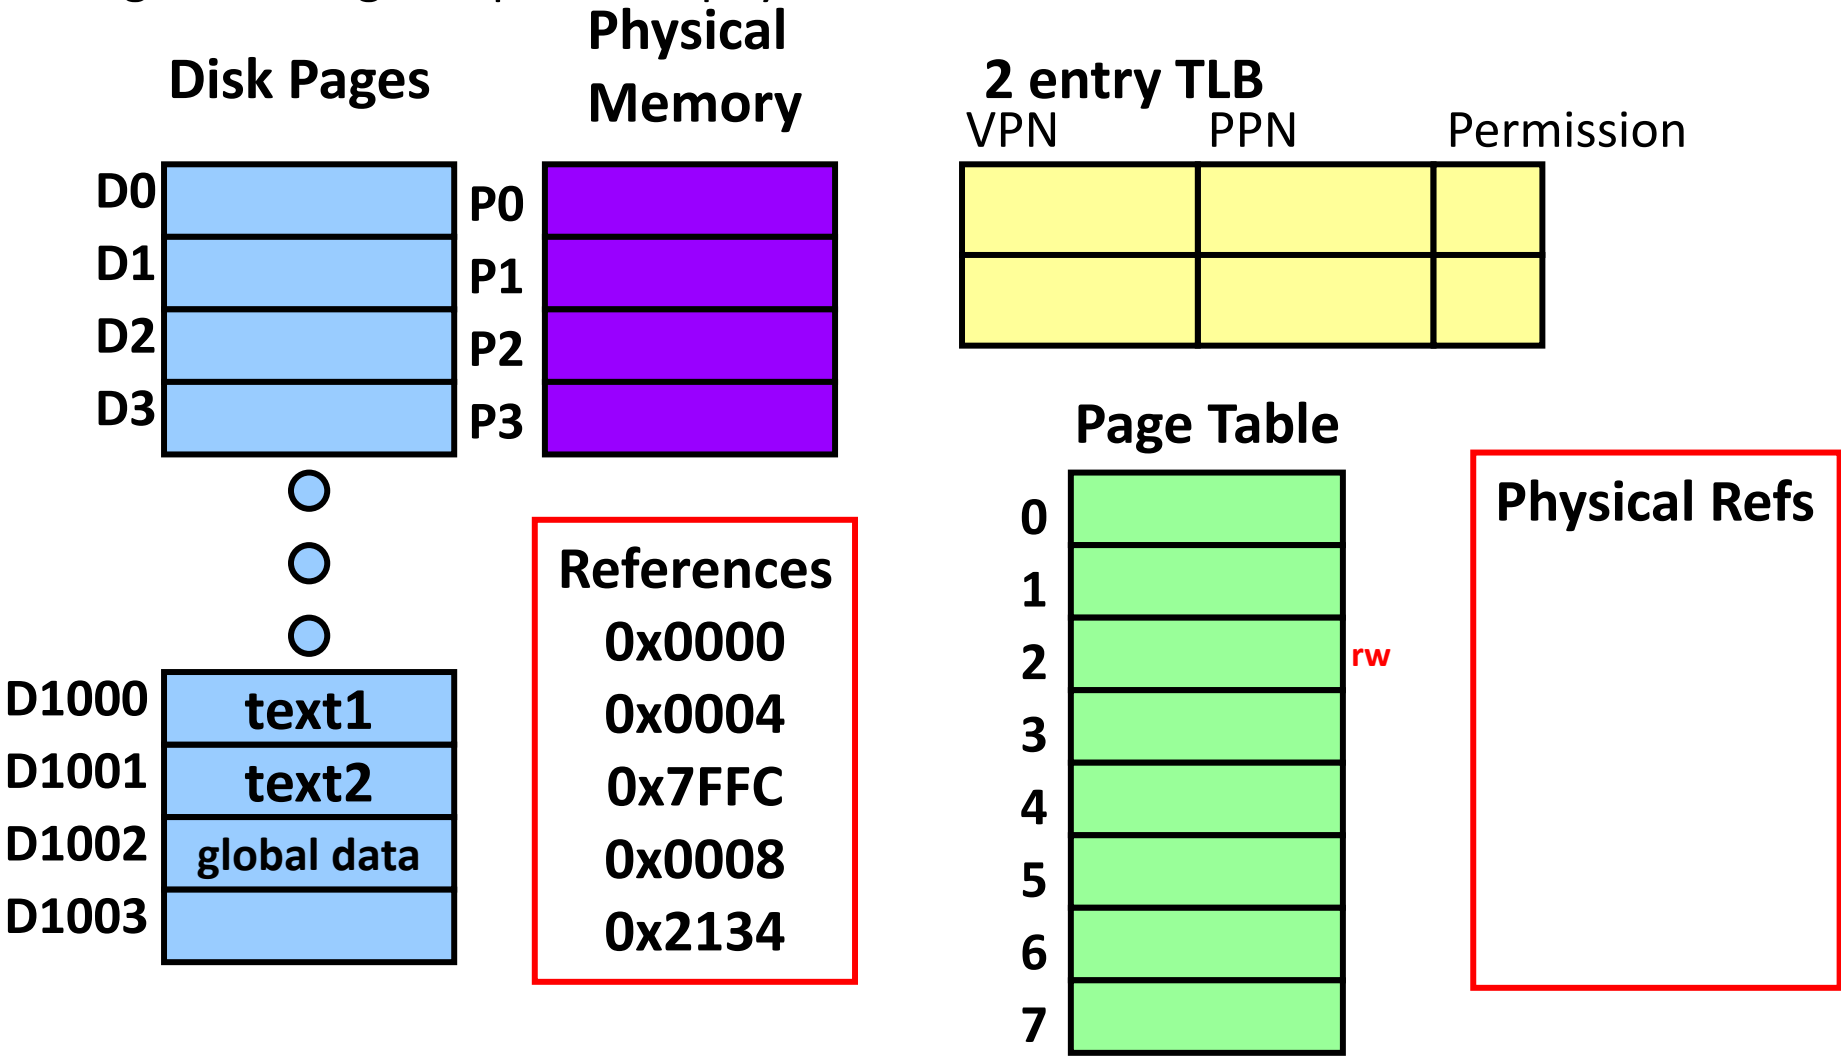
- Loading your program in memory
 - Ask operating system to create a new process
 - Construct a page table for this process
 - Mark all page table entries as invalid with a pointer to the disk image of the program
 - That is, point to the executable file containing the binary.
 - Run the program and get an immediate page fault on the first instruction.

Agenda

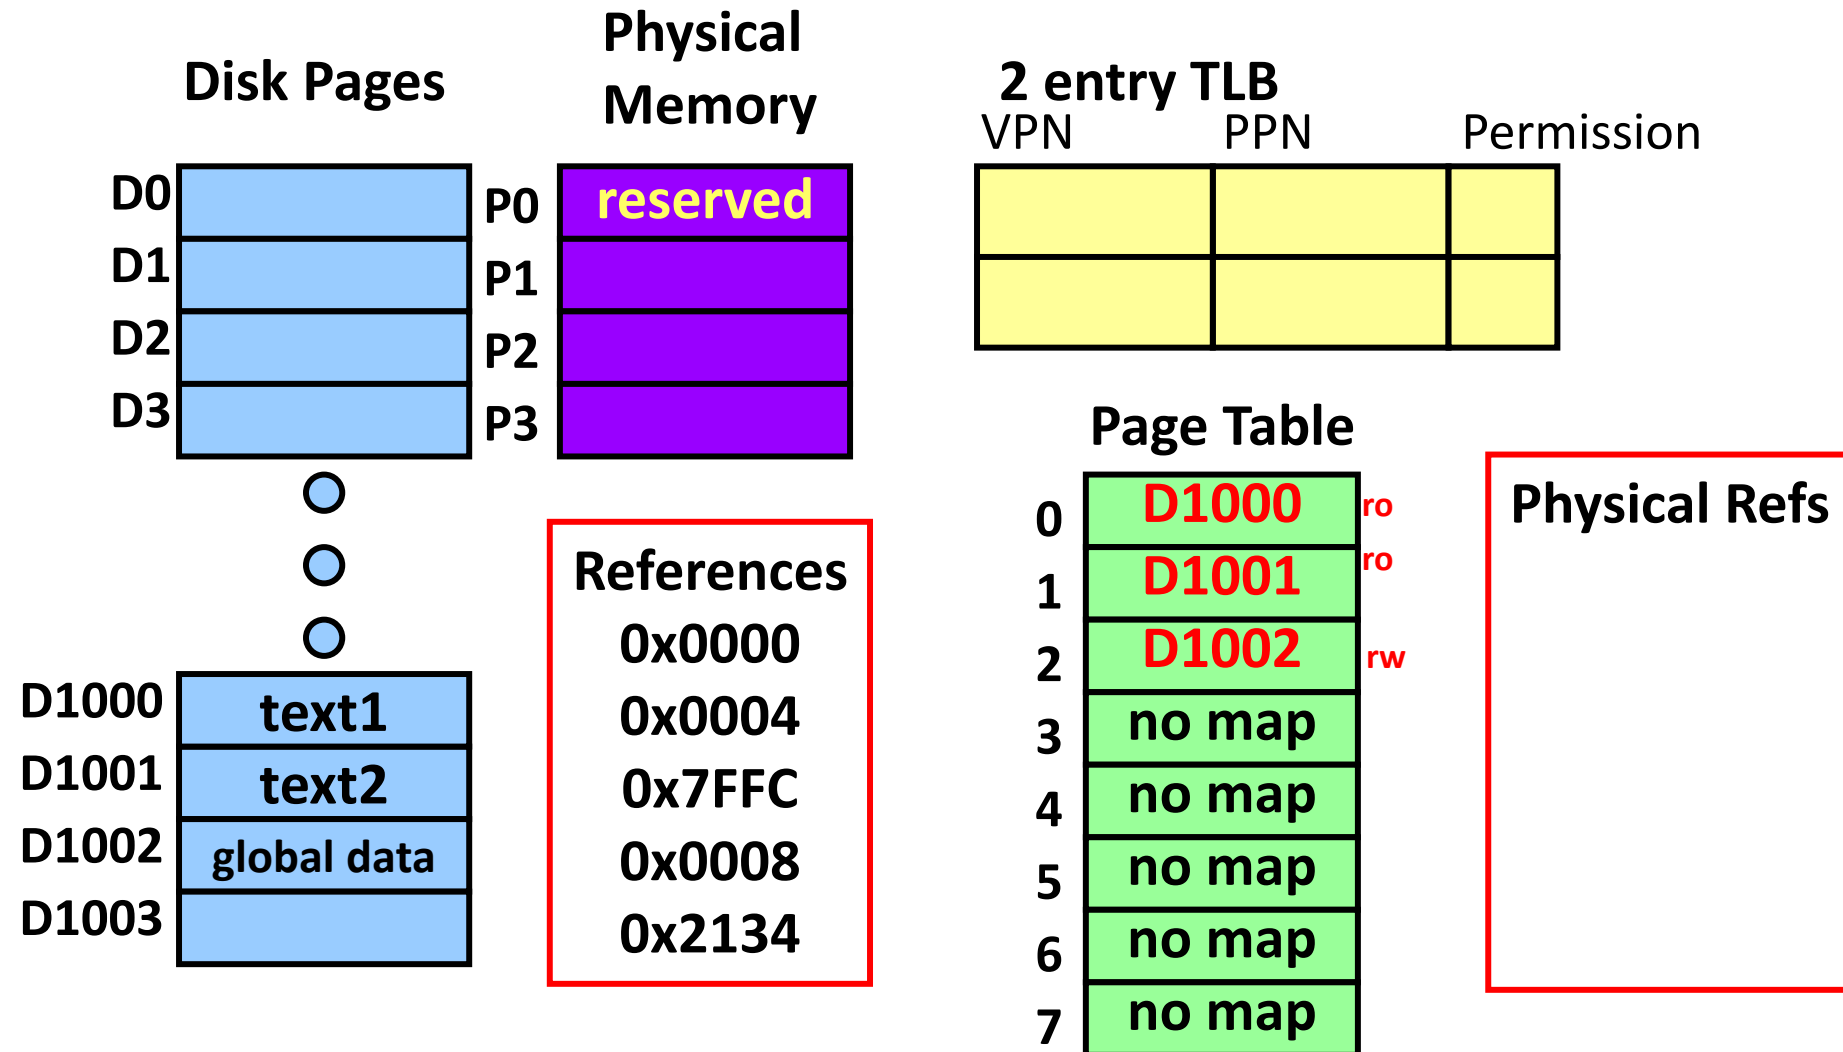
- Table Look-aside Buffers (TLB)
- **Virtual Memory Walkthrough**
- Cache Placement

Loading a program into memory

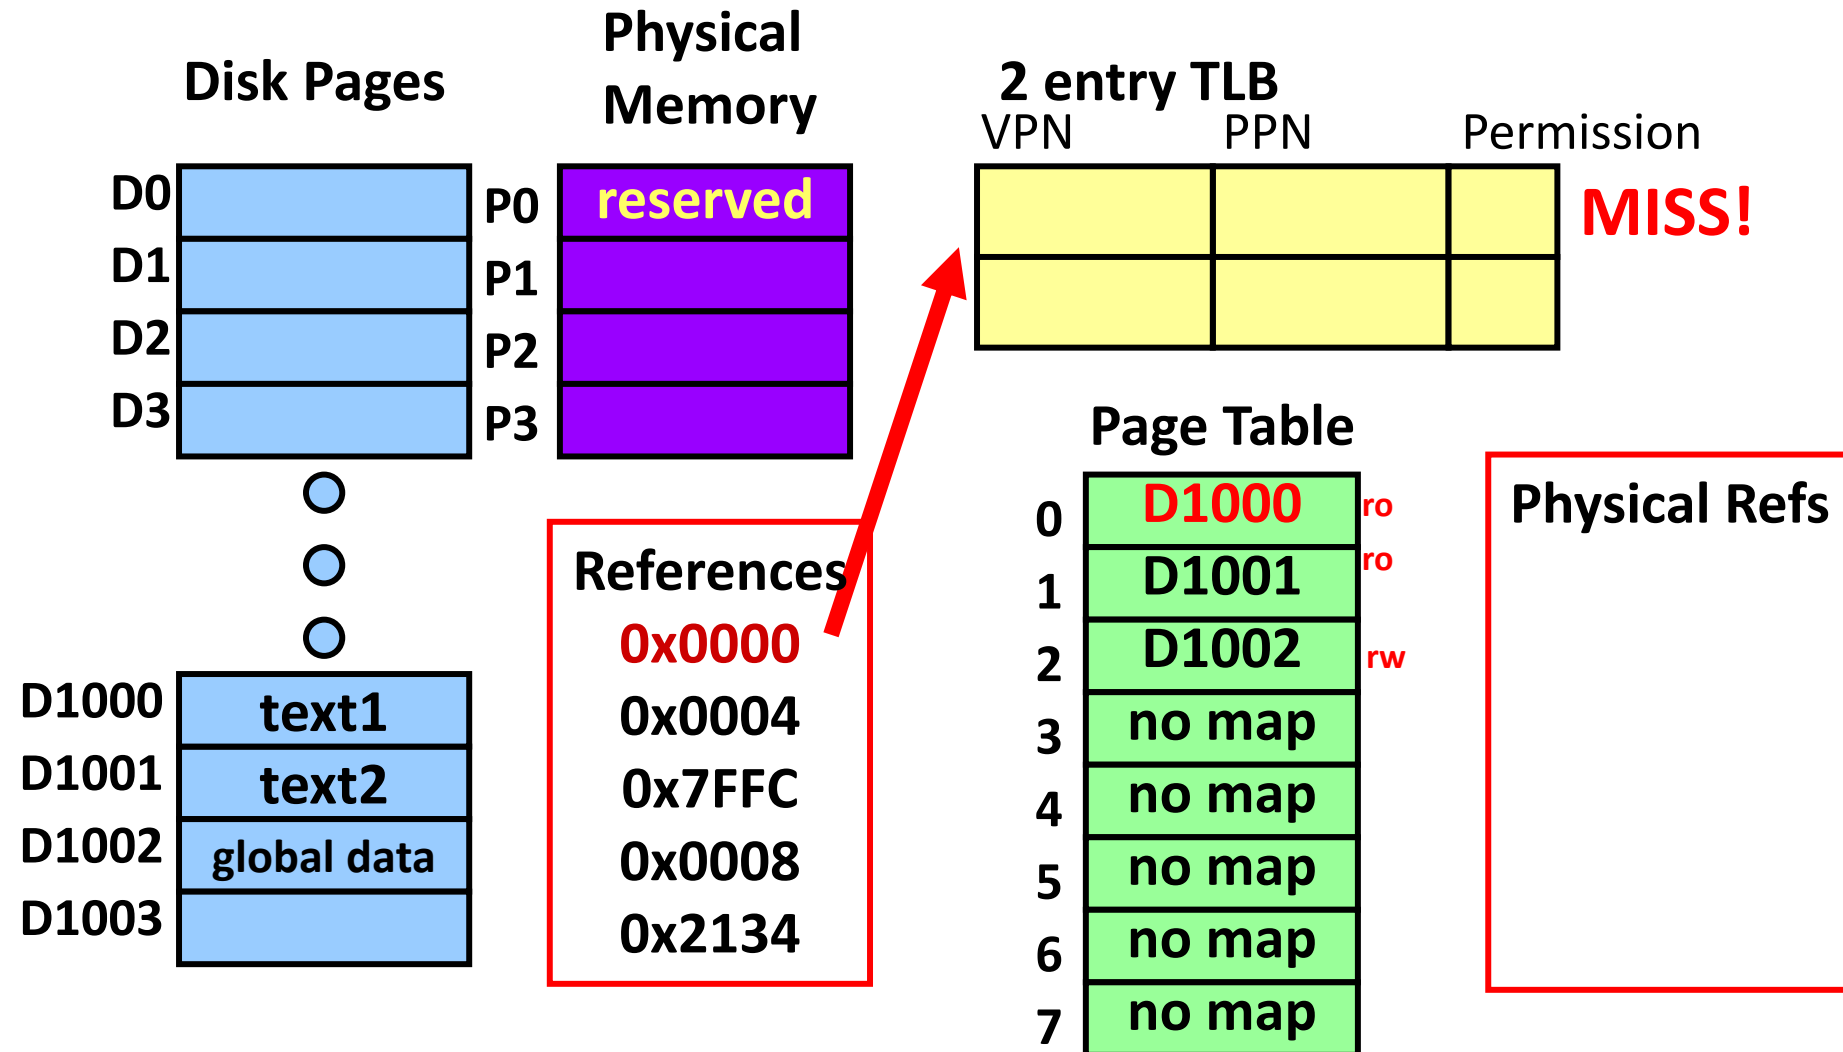
- ❑ Page size = 4 KB, Page table entry size = 4 B
- ❑ Page table register points to physical address 0x0000



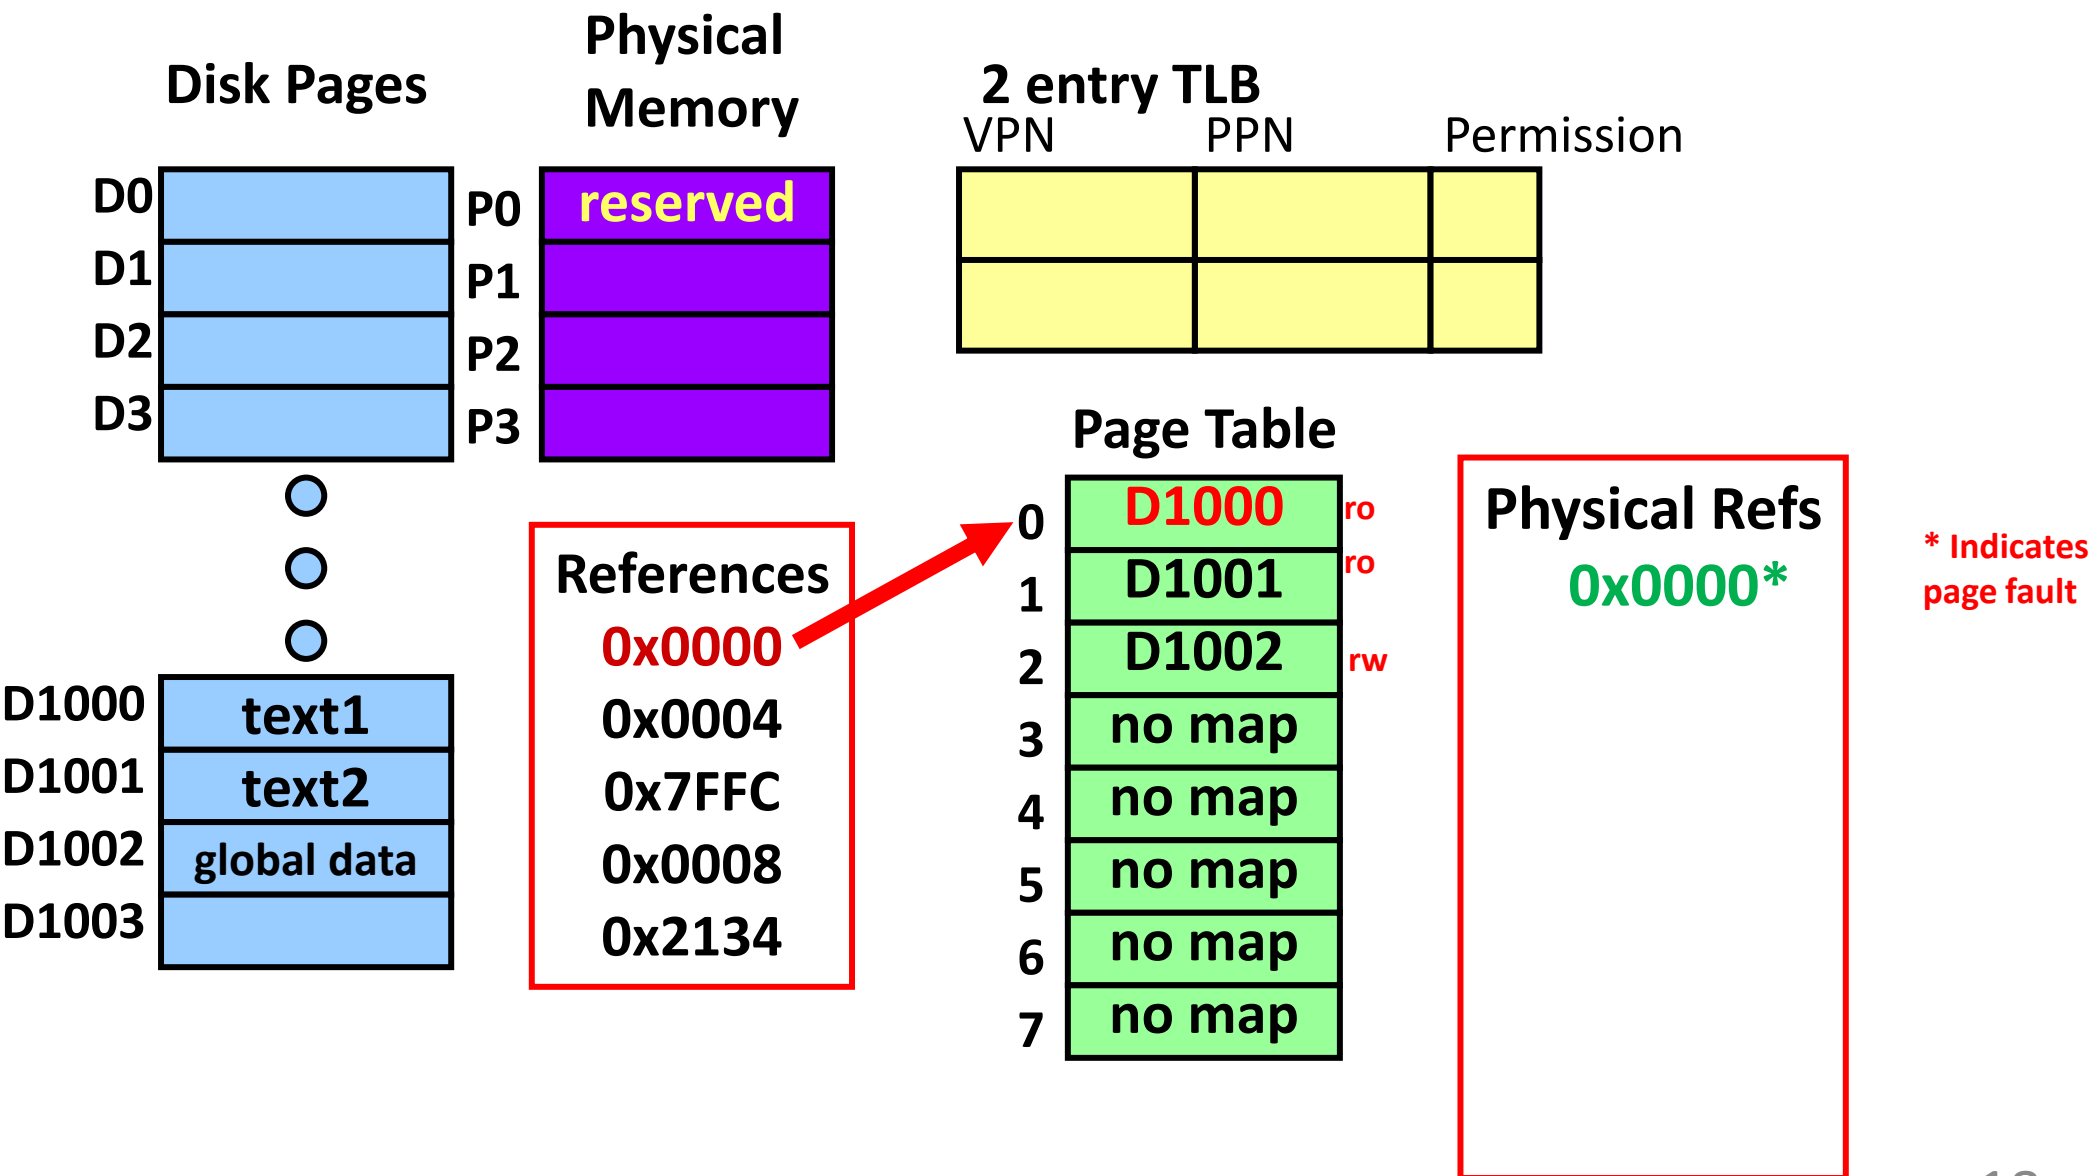
Step 1: Read executable header & initialize page table



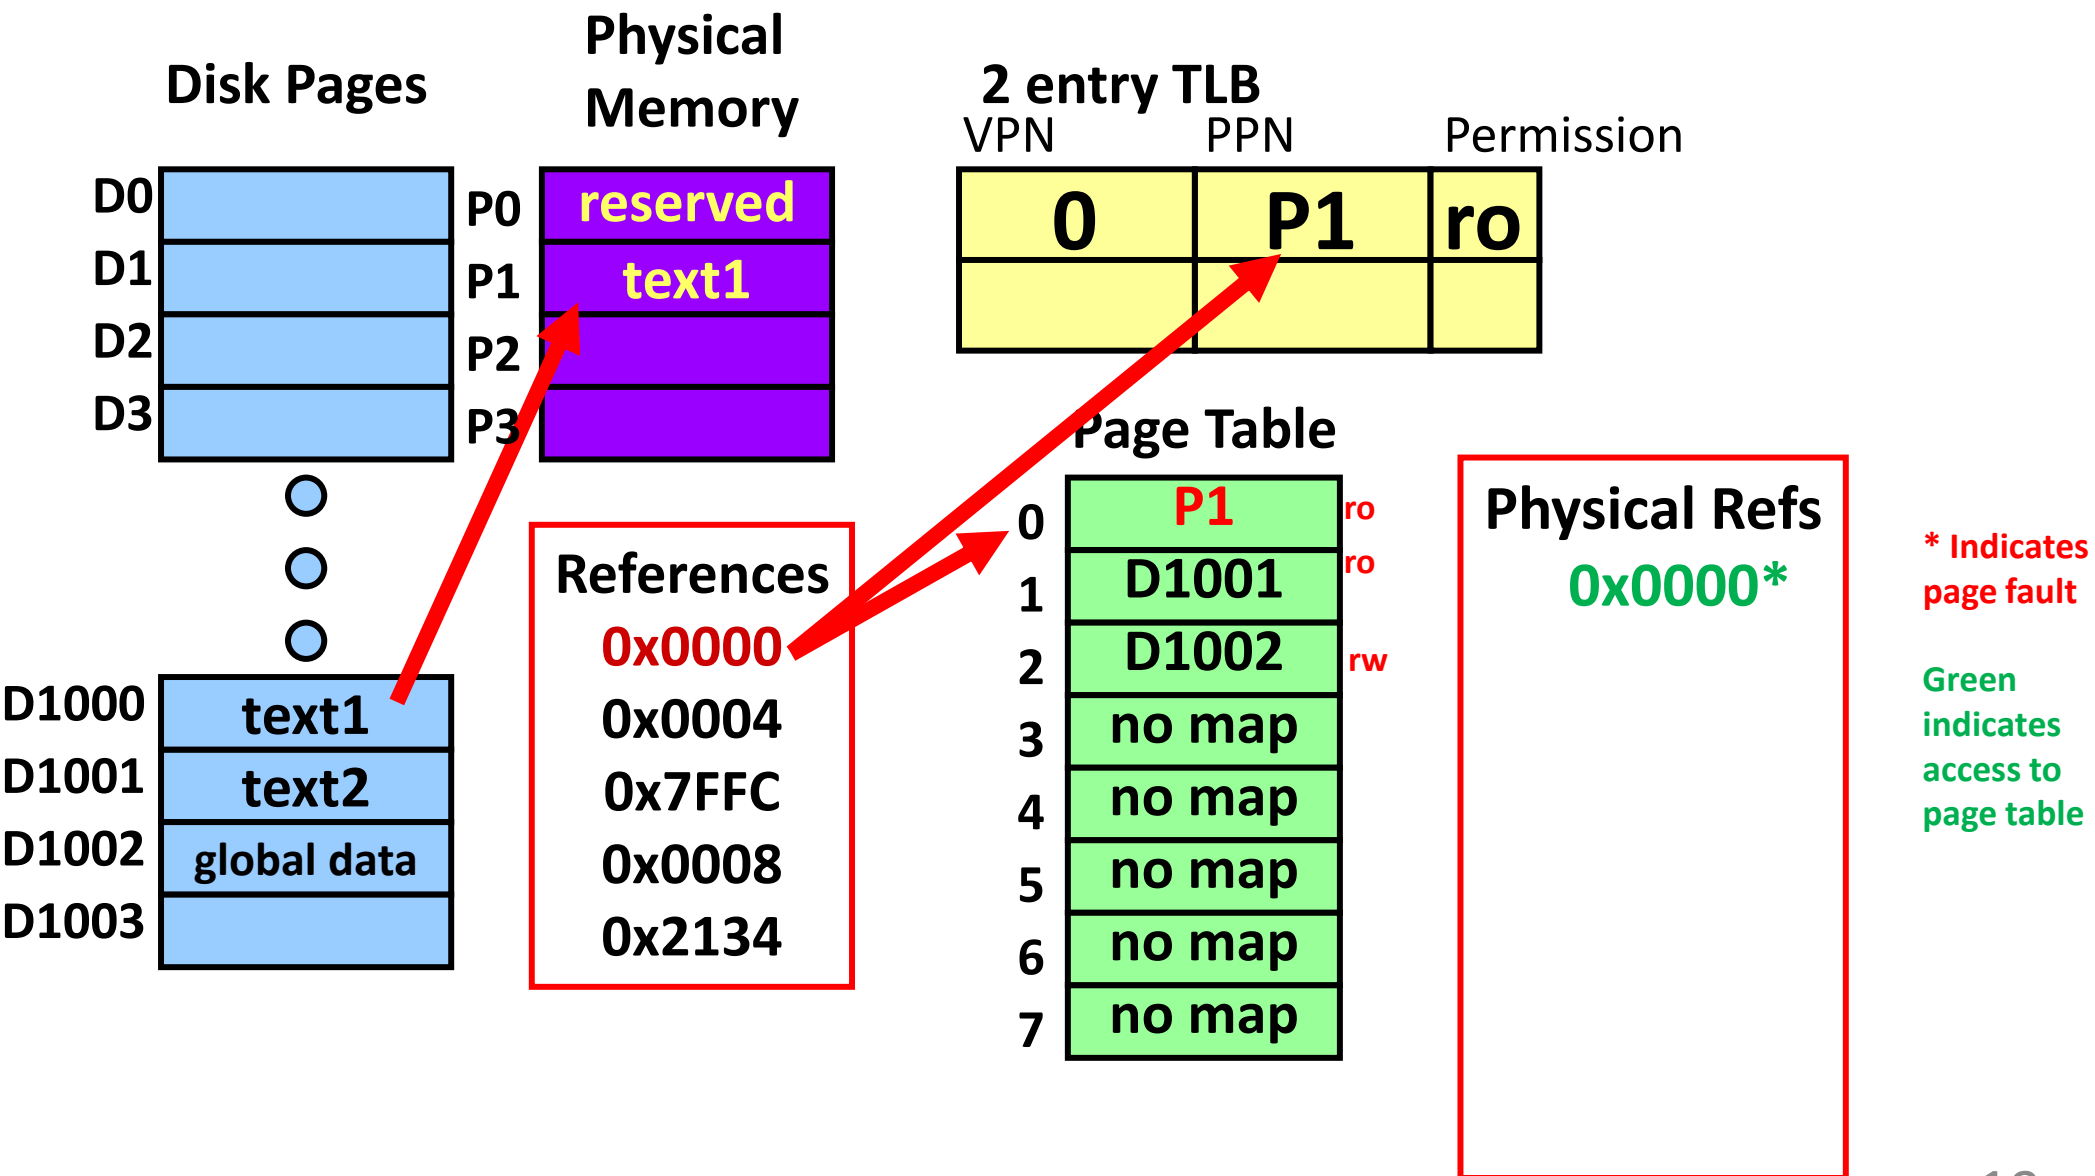
Step 2: Load PC from header & start execution



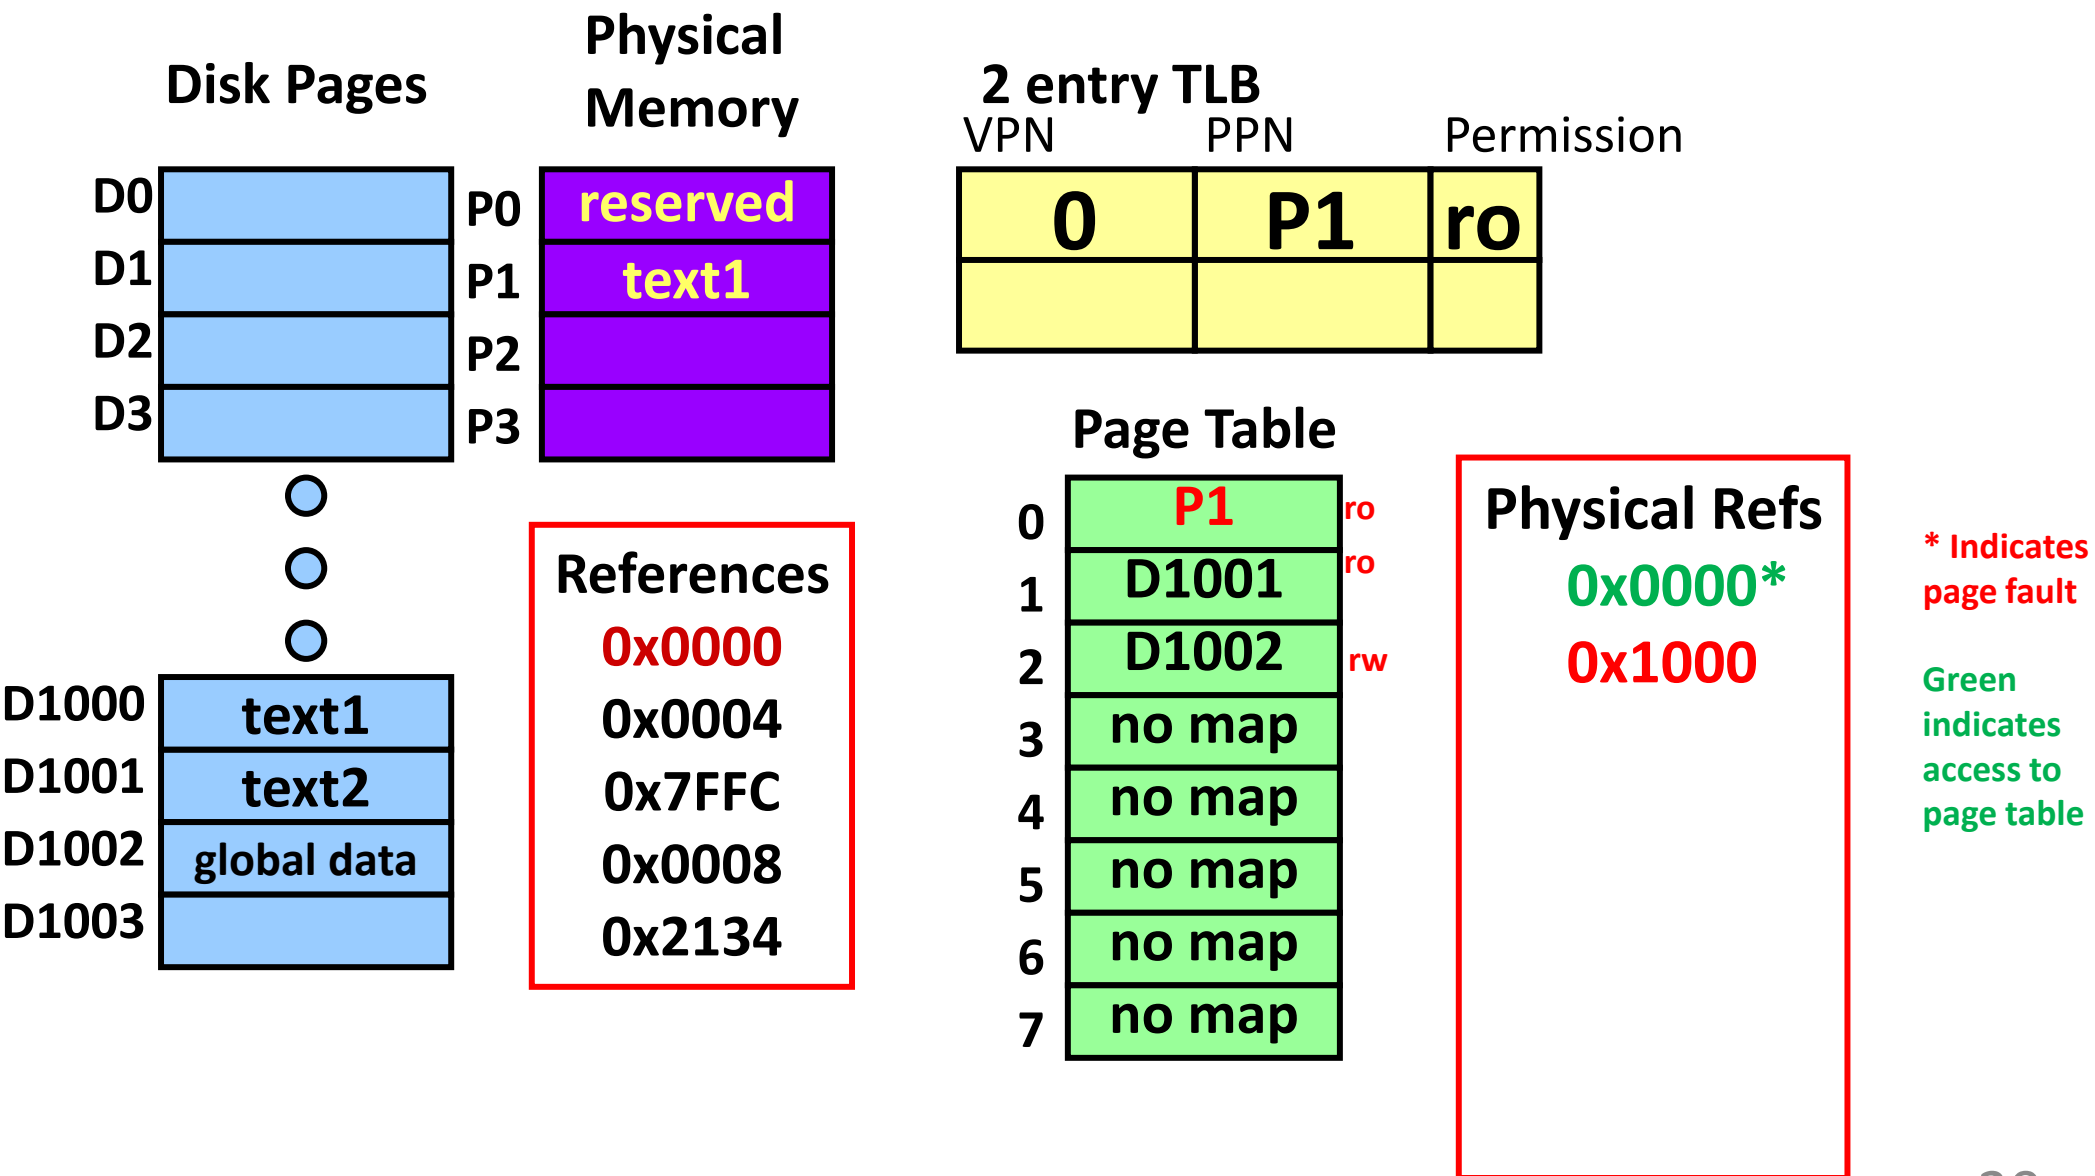
Fetching instruction 0000



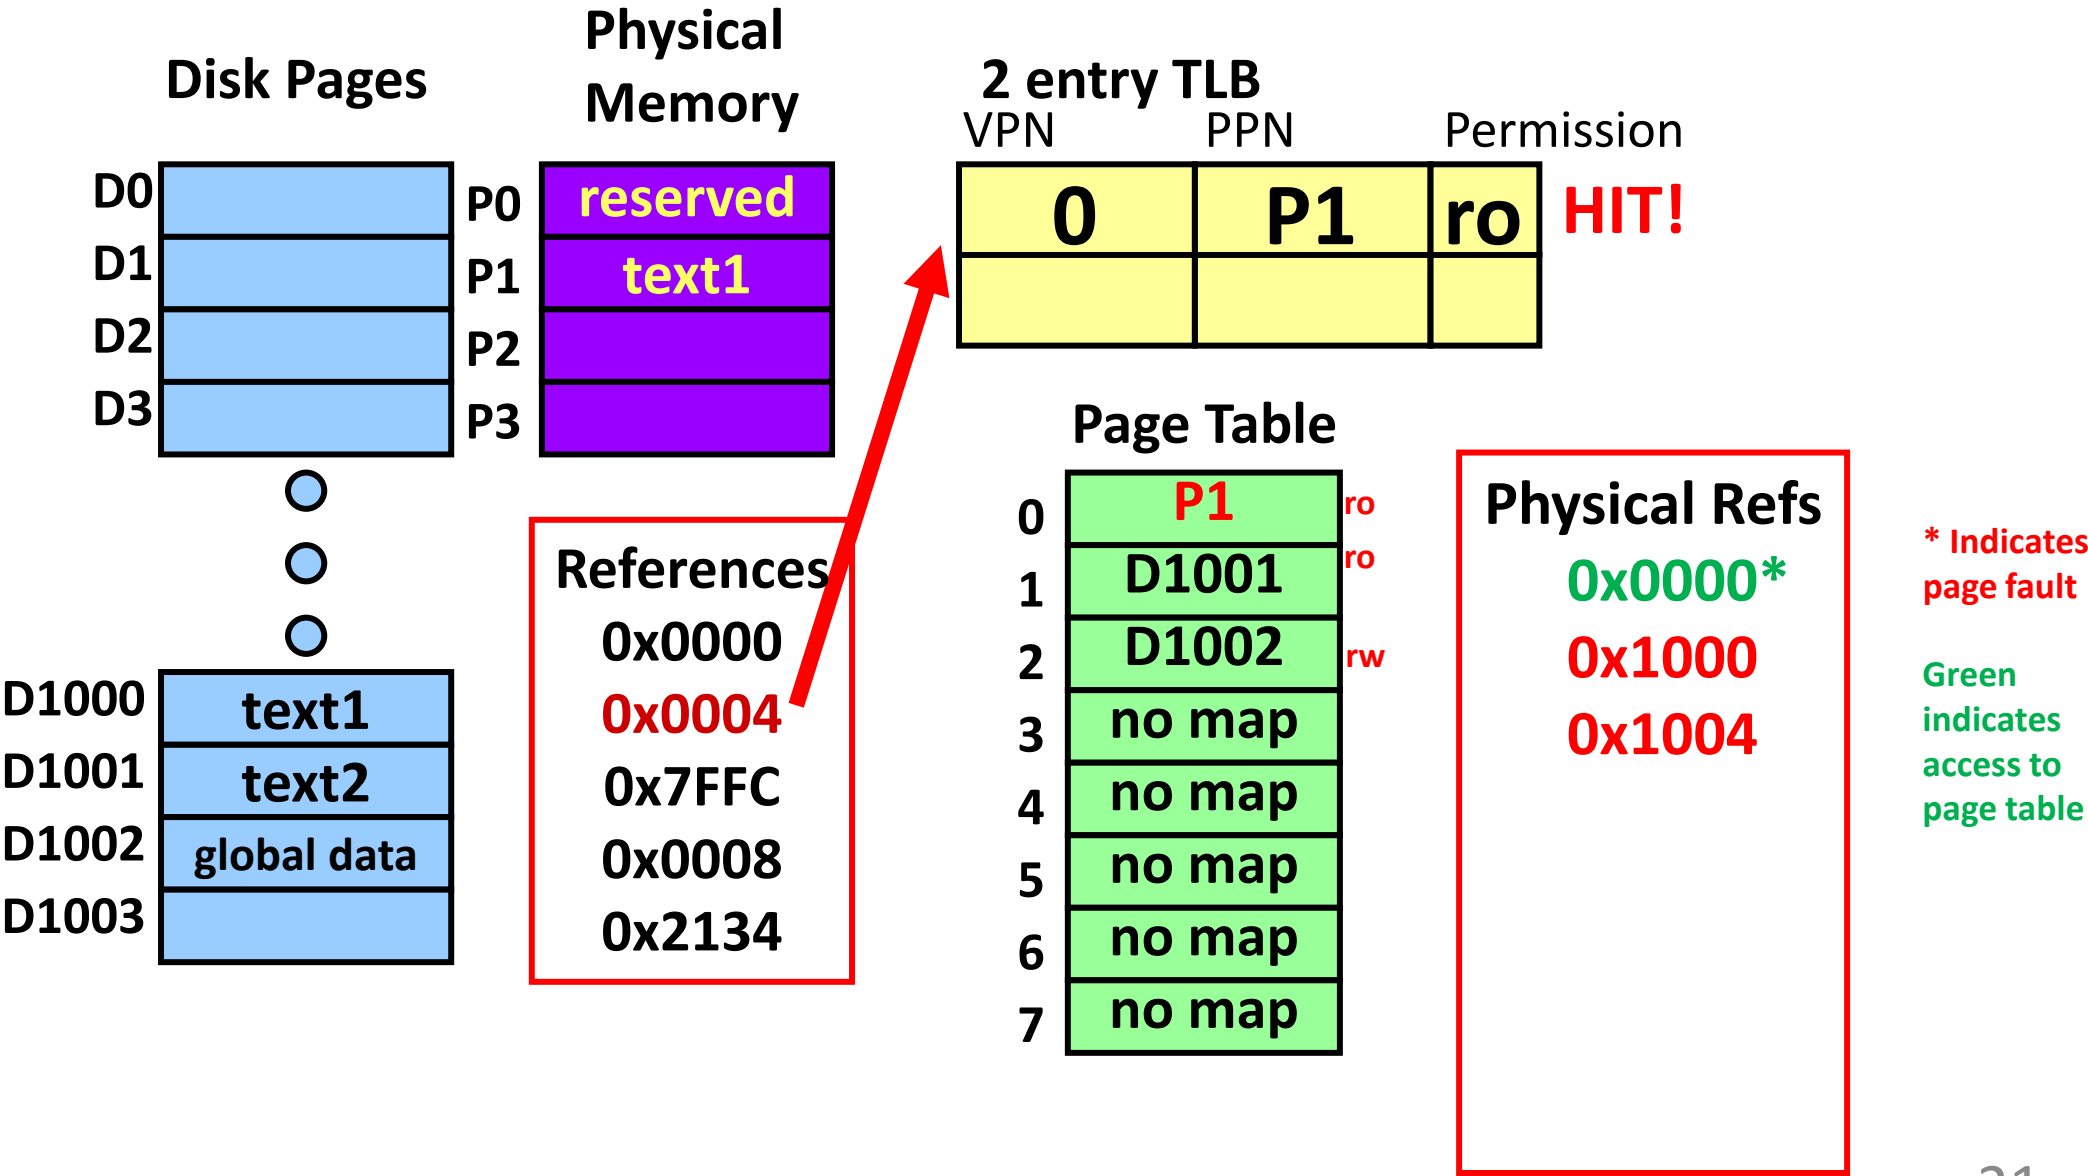
Fetching instruction 0000



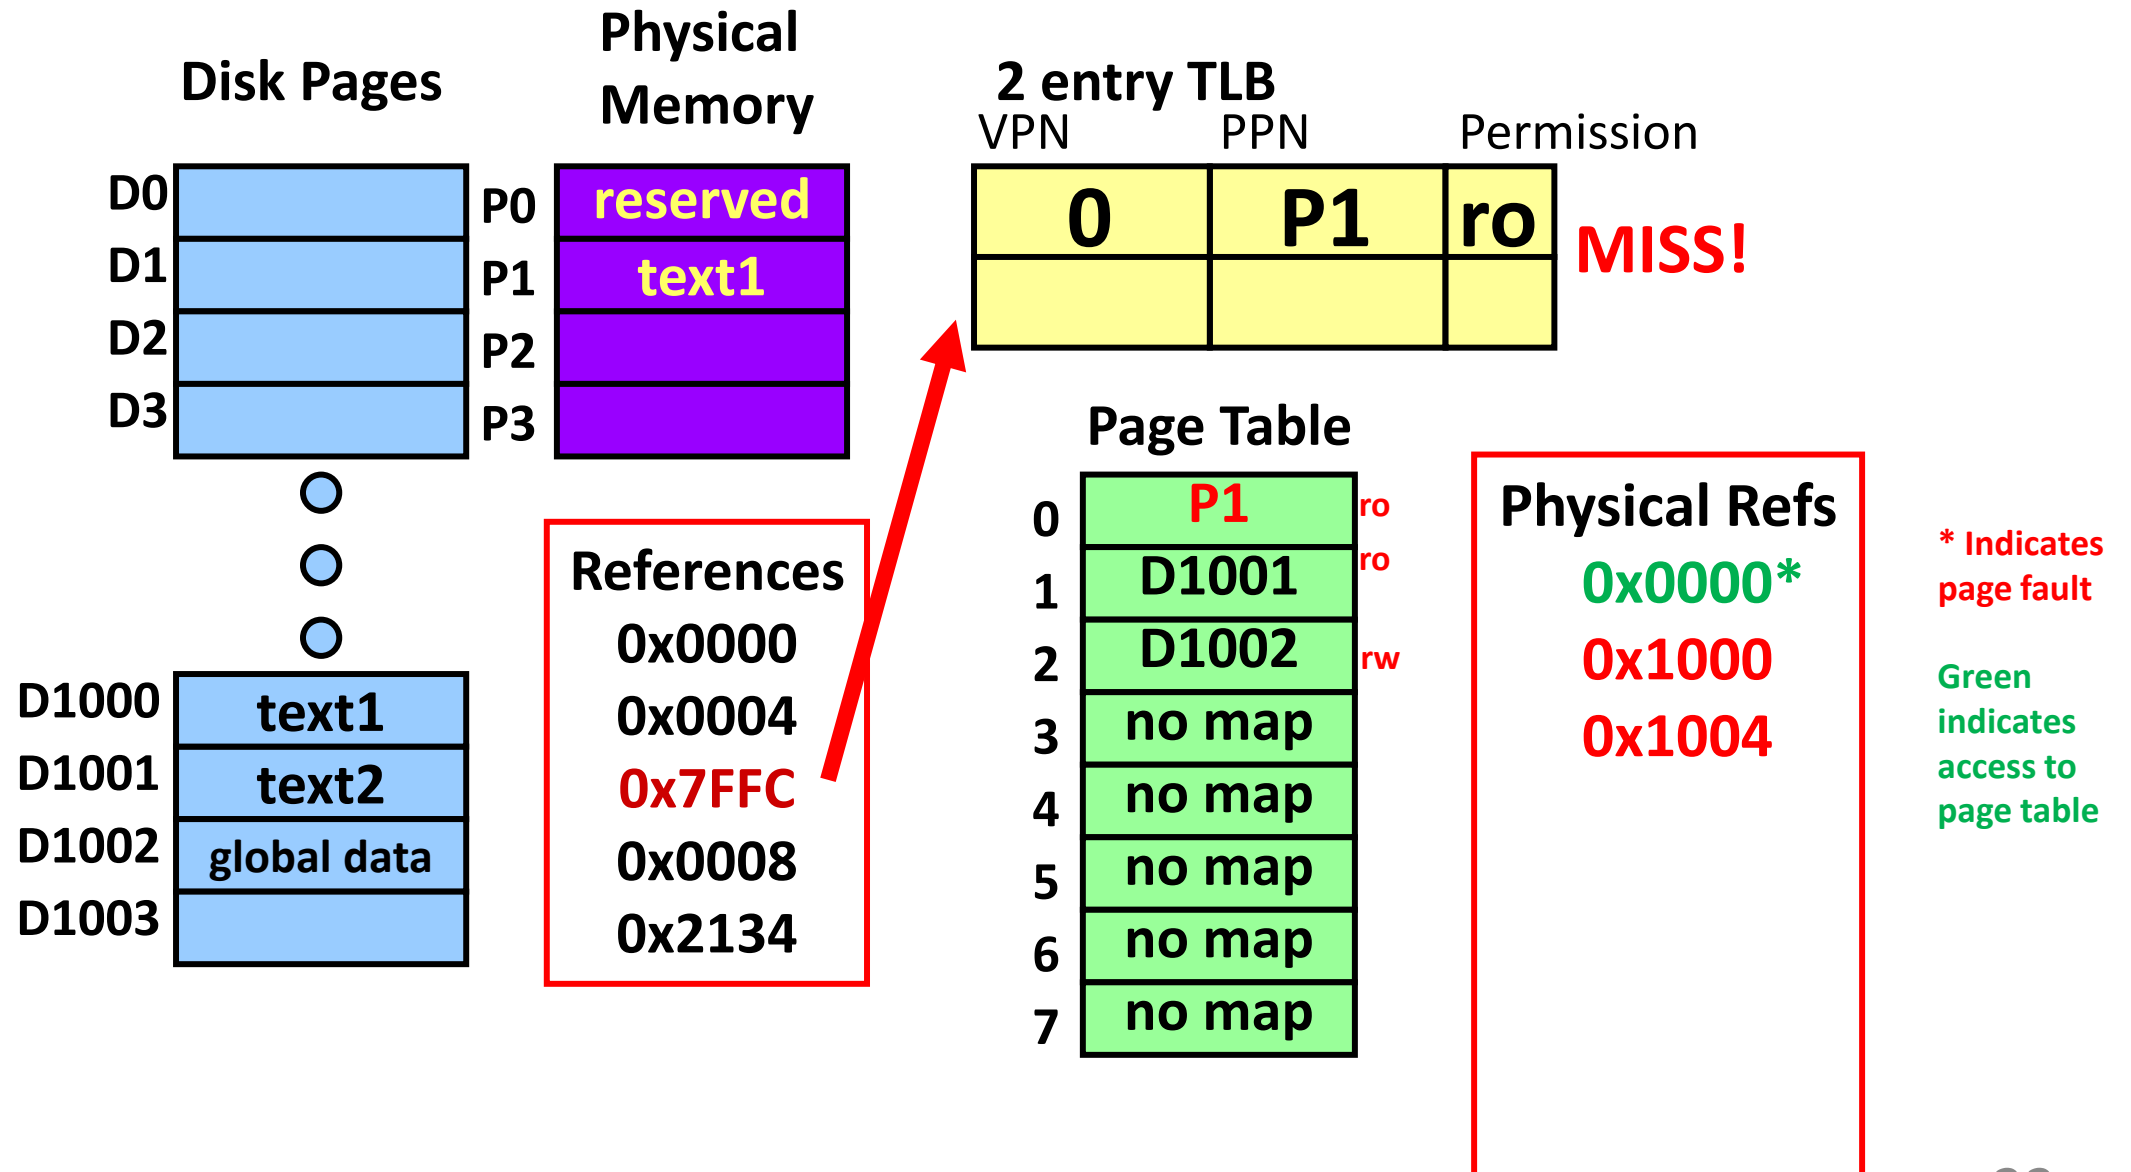
Fetching instruction 0000



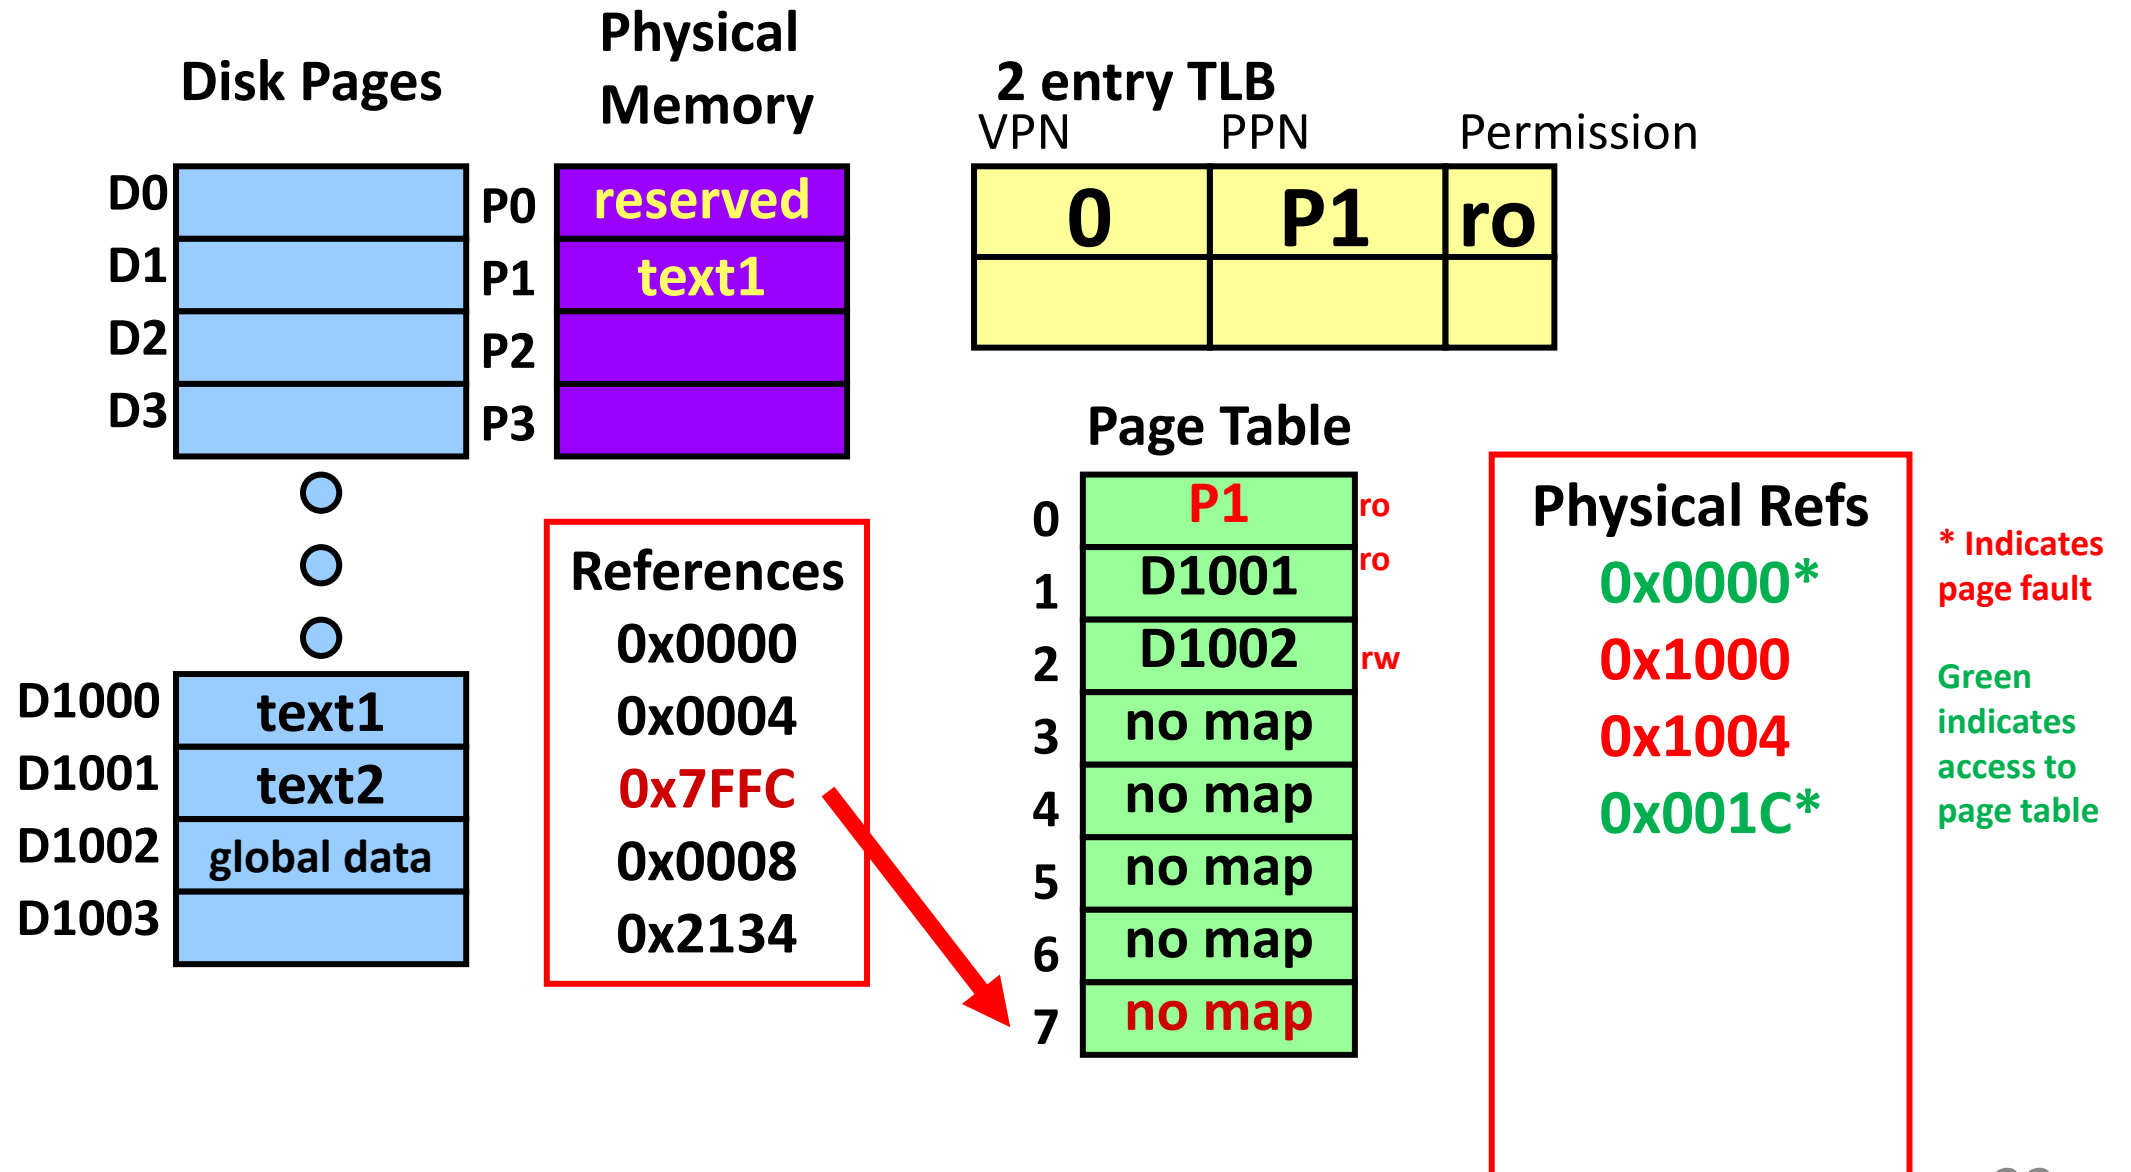
Fetching instruction 0004



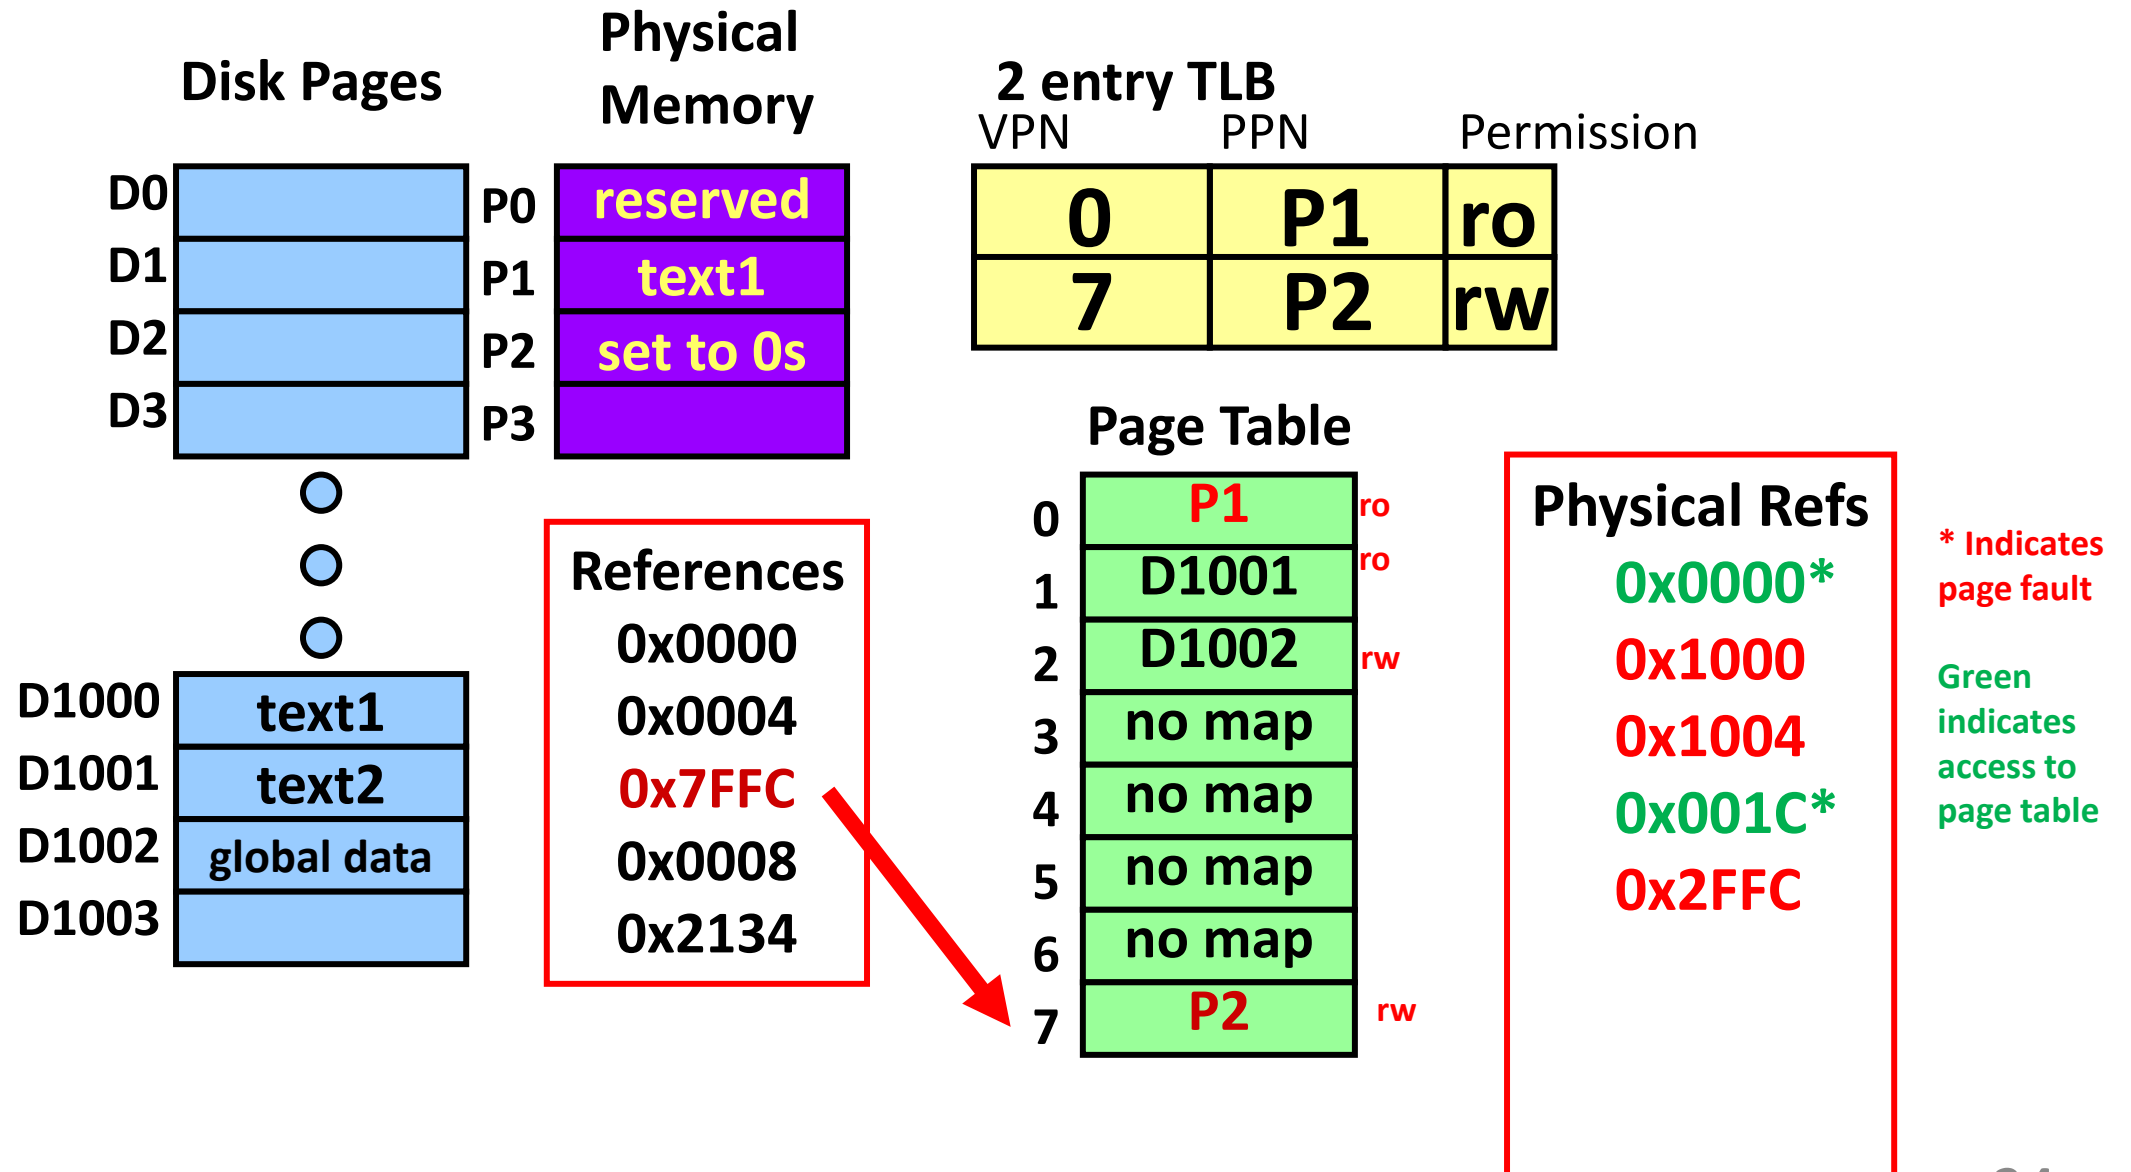
Reference 7FFC



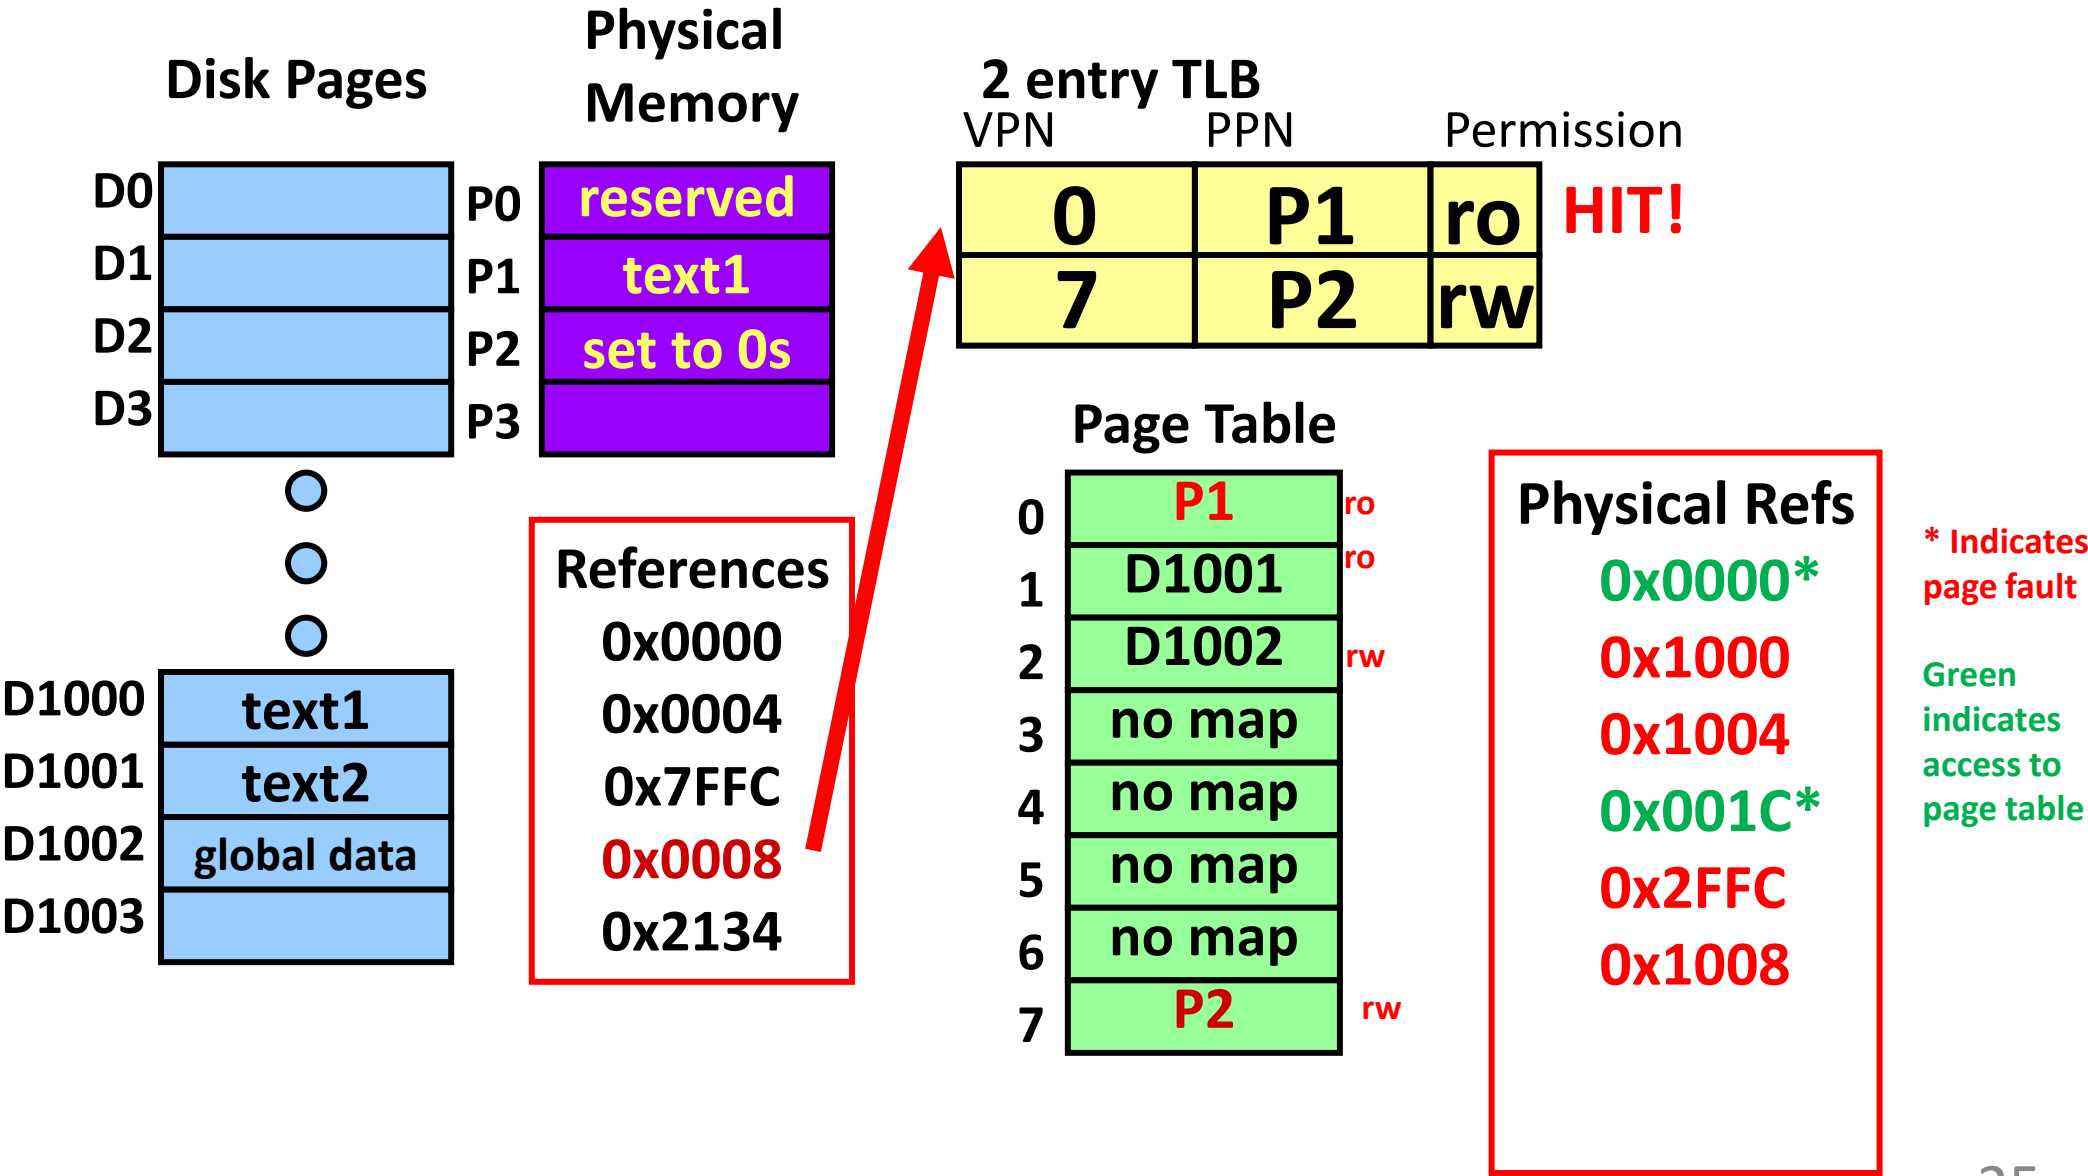
Reference 7FFC



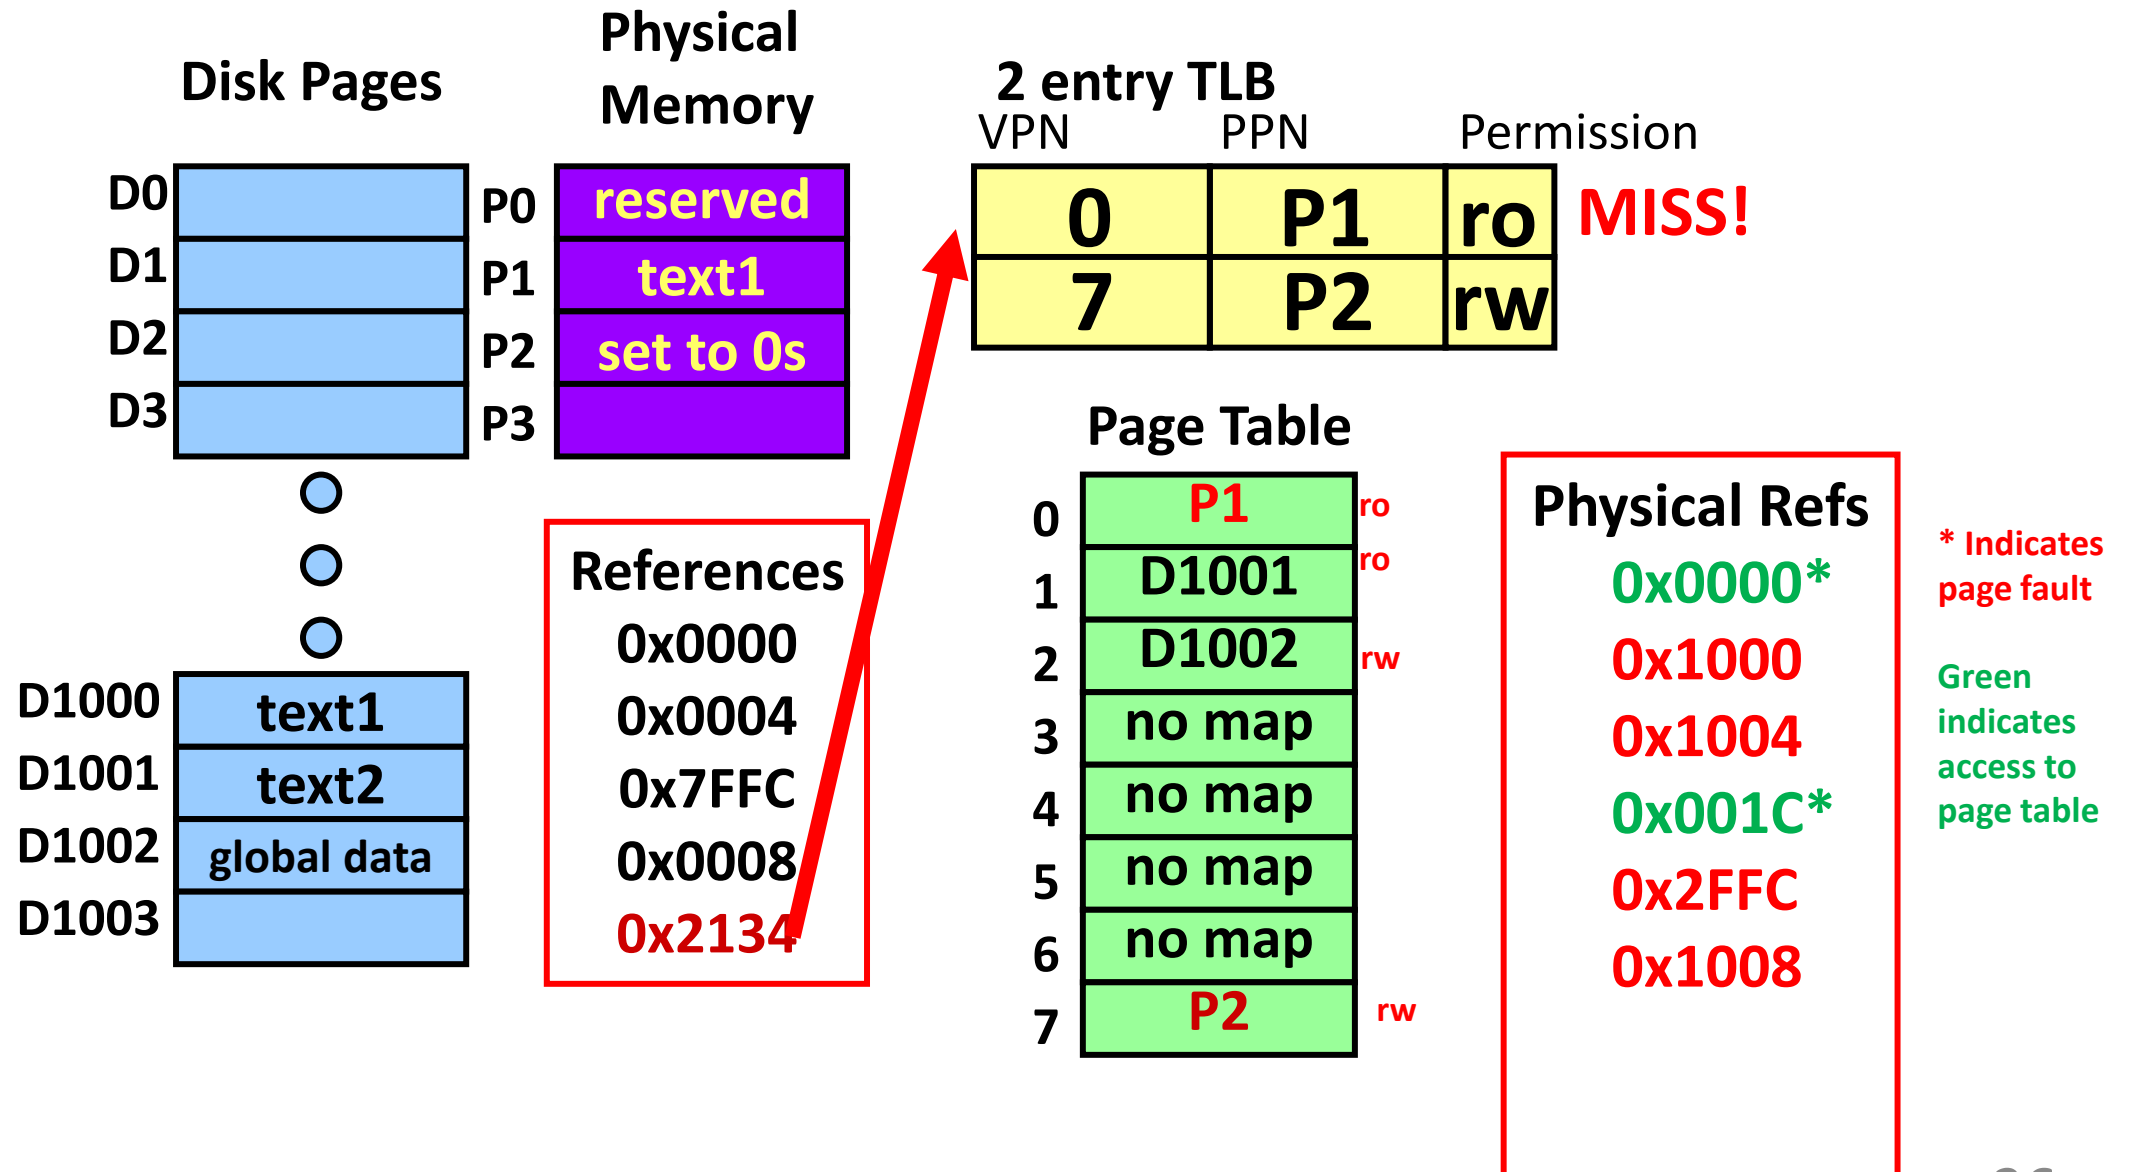
Reference 7FFC



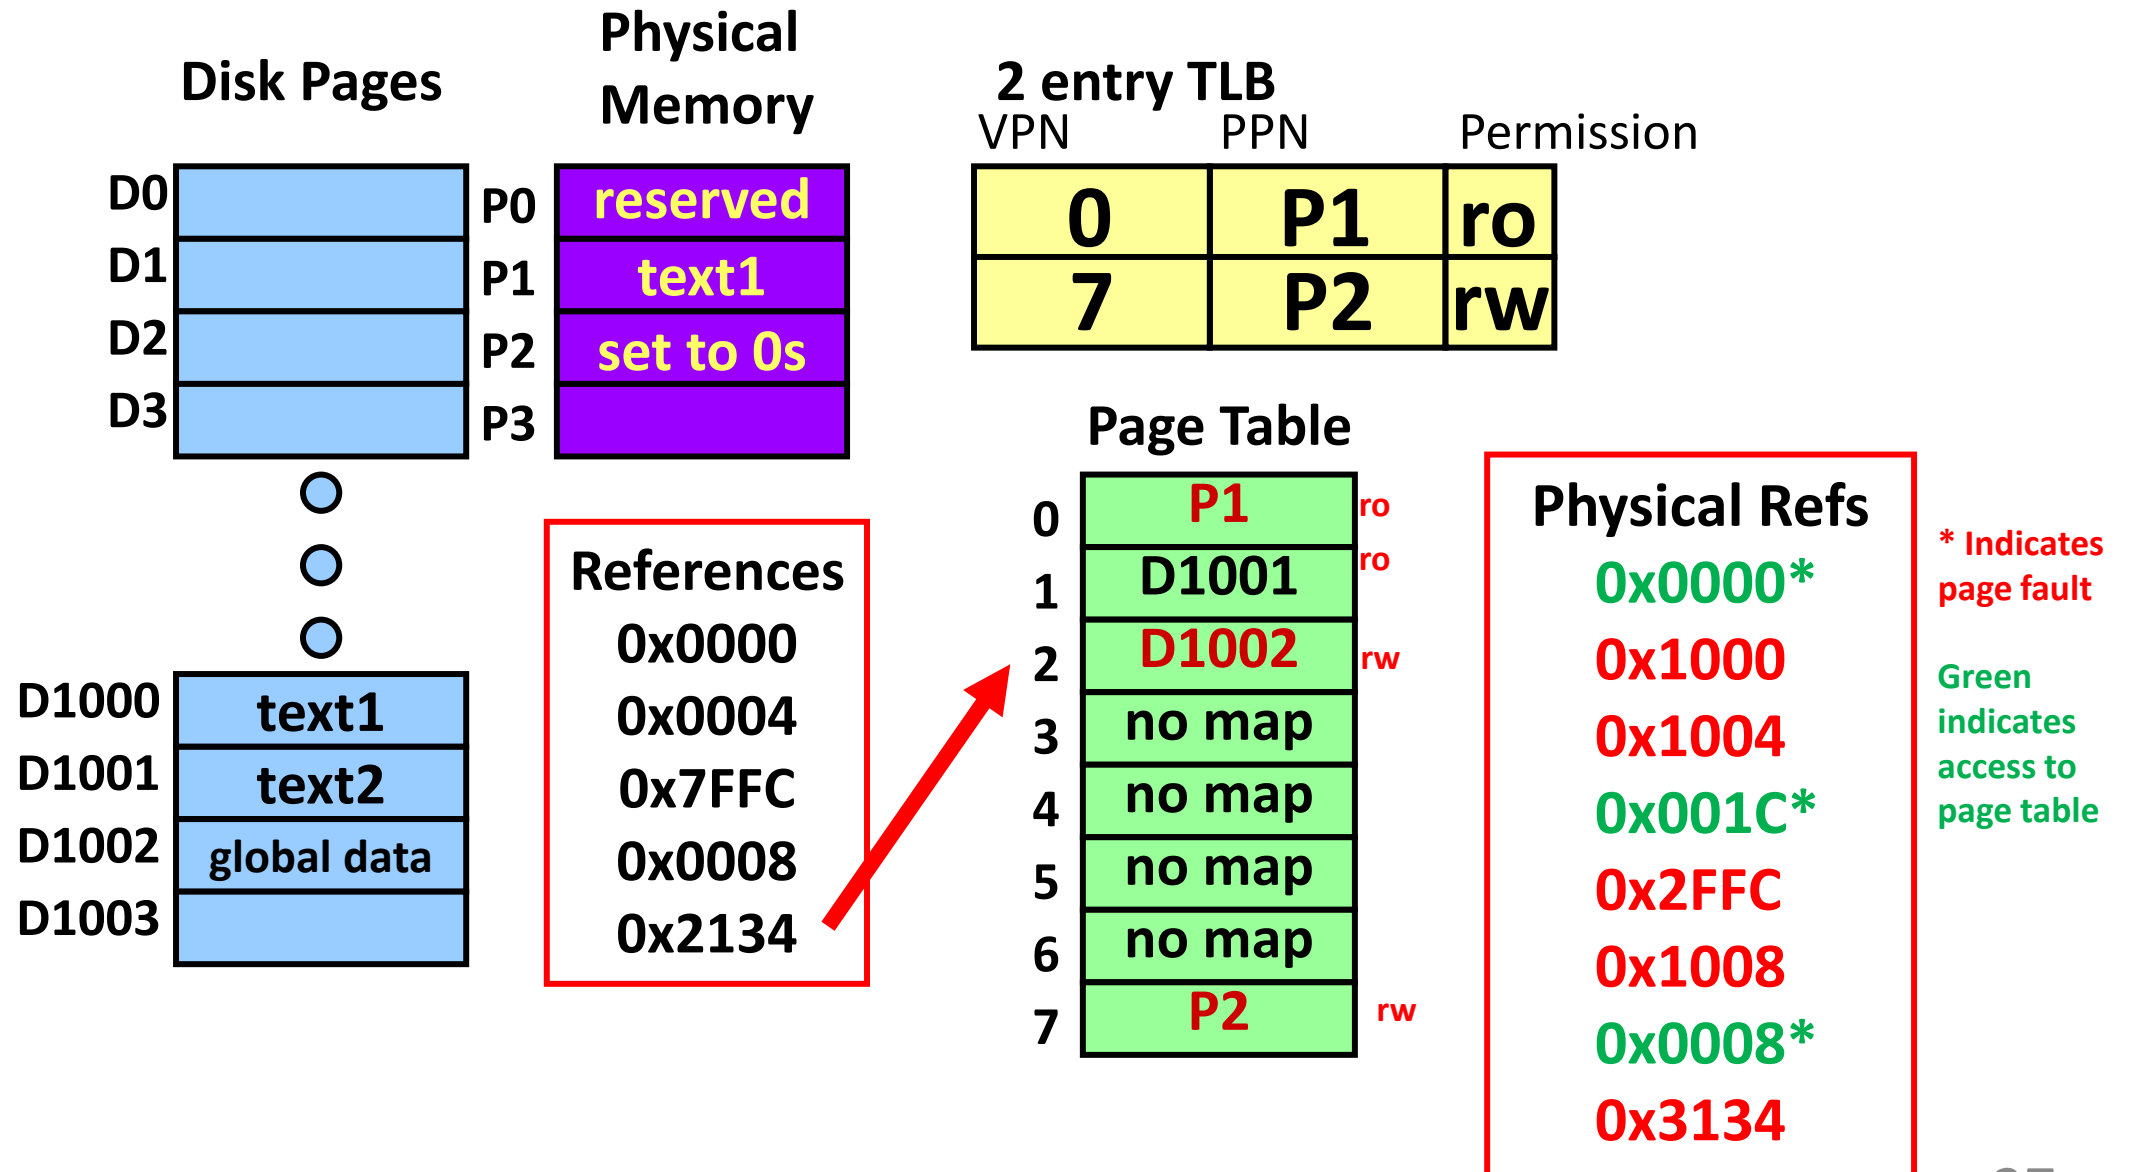
Fetching instruction 0008



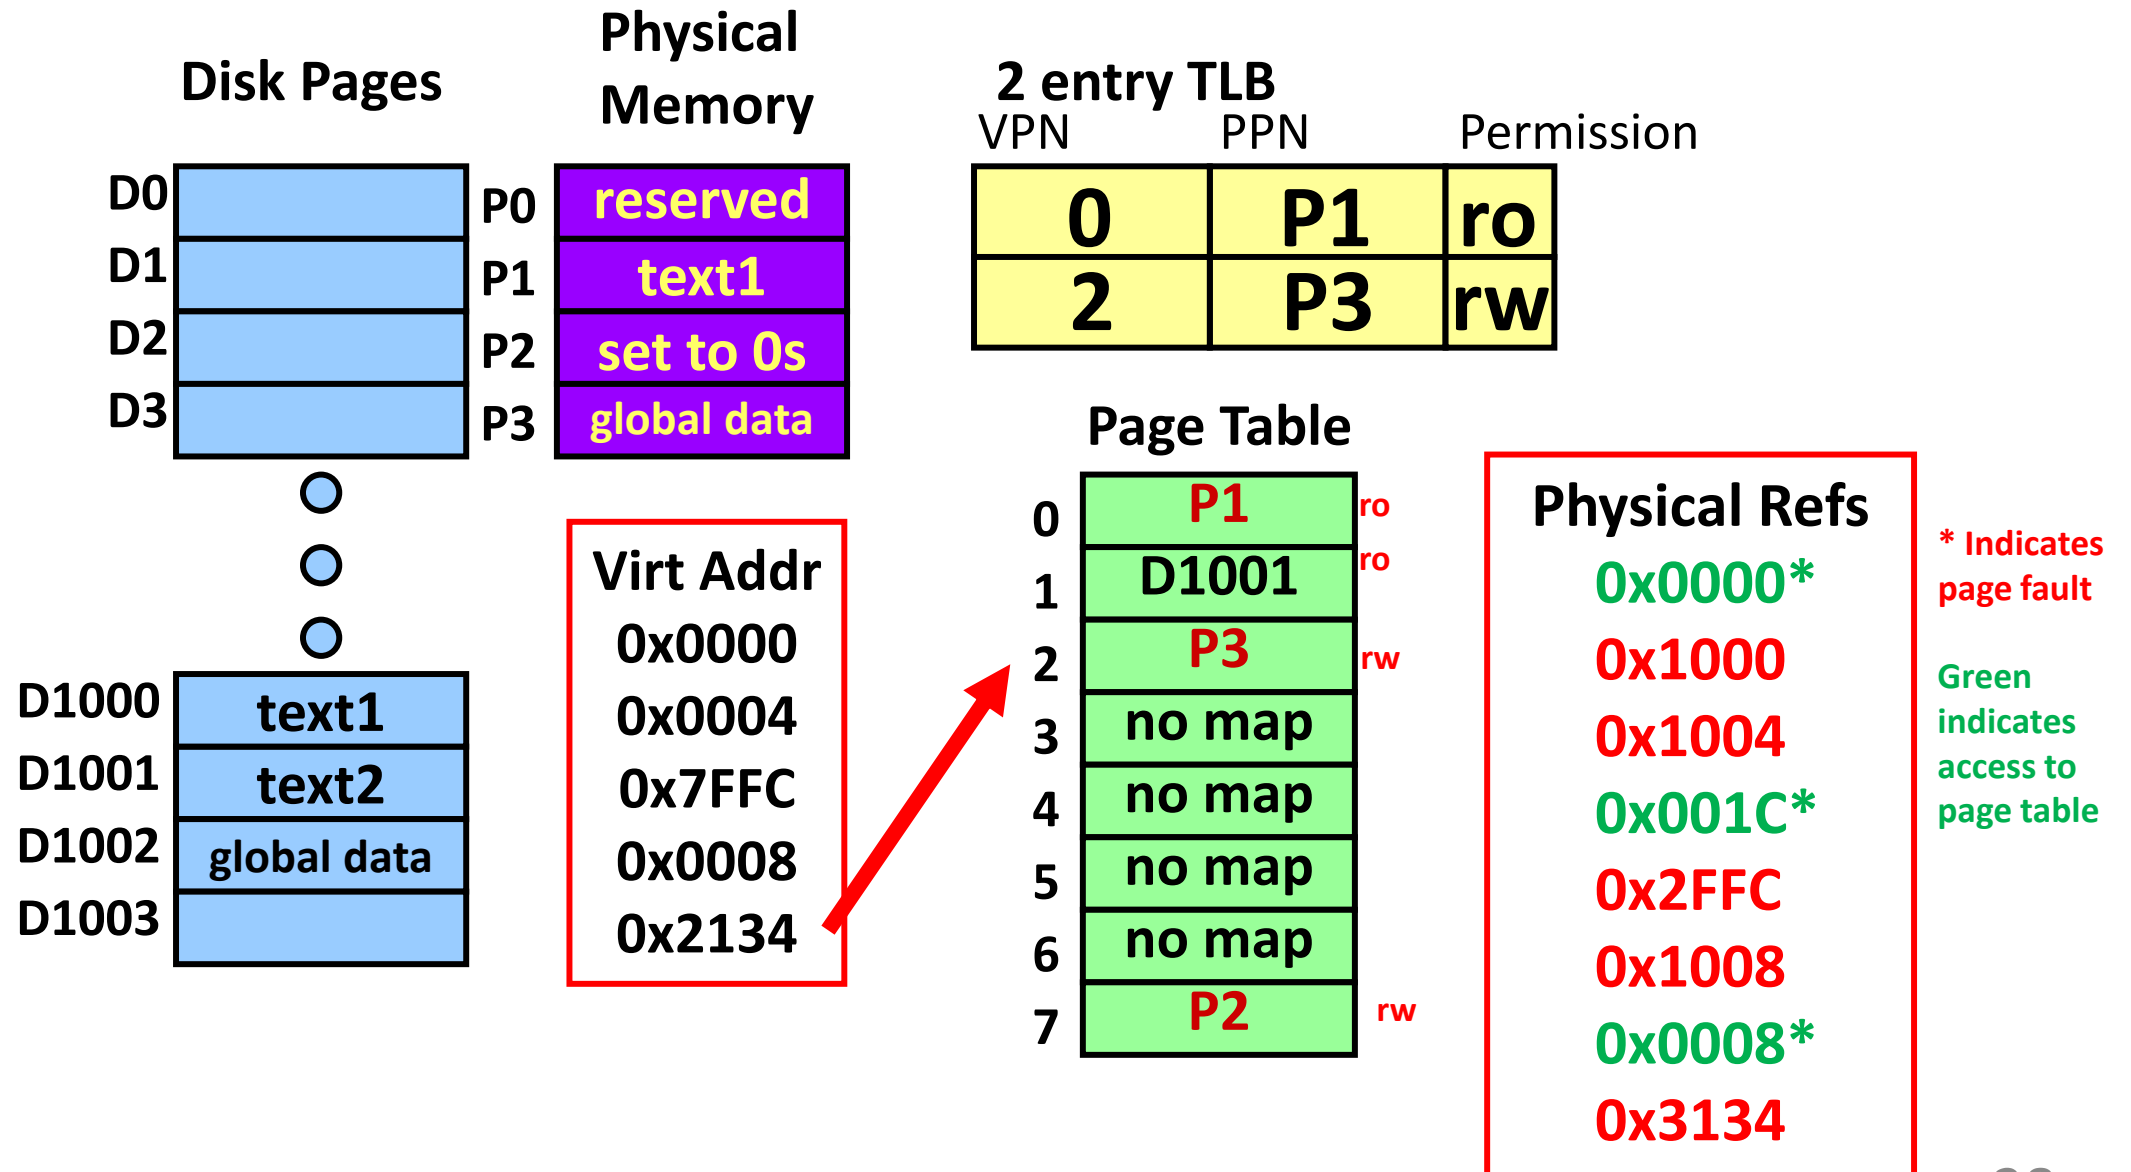
Reference 2134



Reference 2134



Reference 2134



Agenda

- Table Look-aside Buffers (TLB)
- Virtual Memory Walkthrough
- **Cache Placement**

Next topic: Placing Caches in a VM System

- VM systems give us two different addresses:
 - virtual and physical
- Which address should we use to access the data cache?
 - Physical address (after VM translations).
 - We have to wait for the translation; slower.
 - Virtual address (before VM translation).
 - Faster access.
 - More complex.

Cache & VM Organization: Option 1


Physically-addressed Cache



Slower



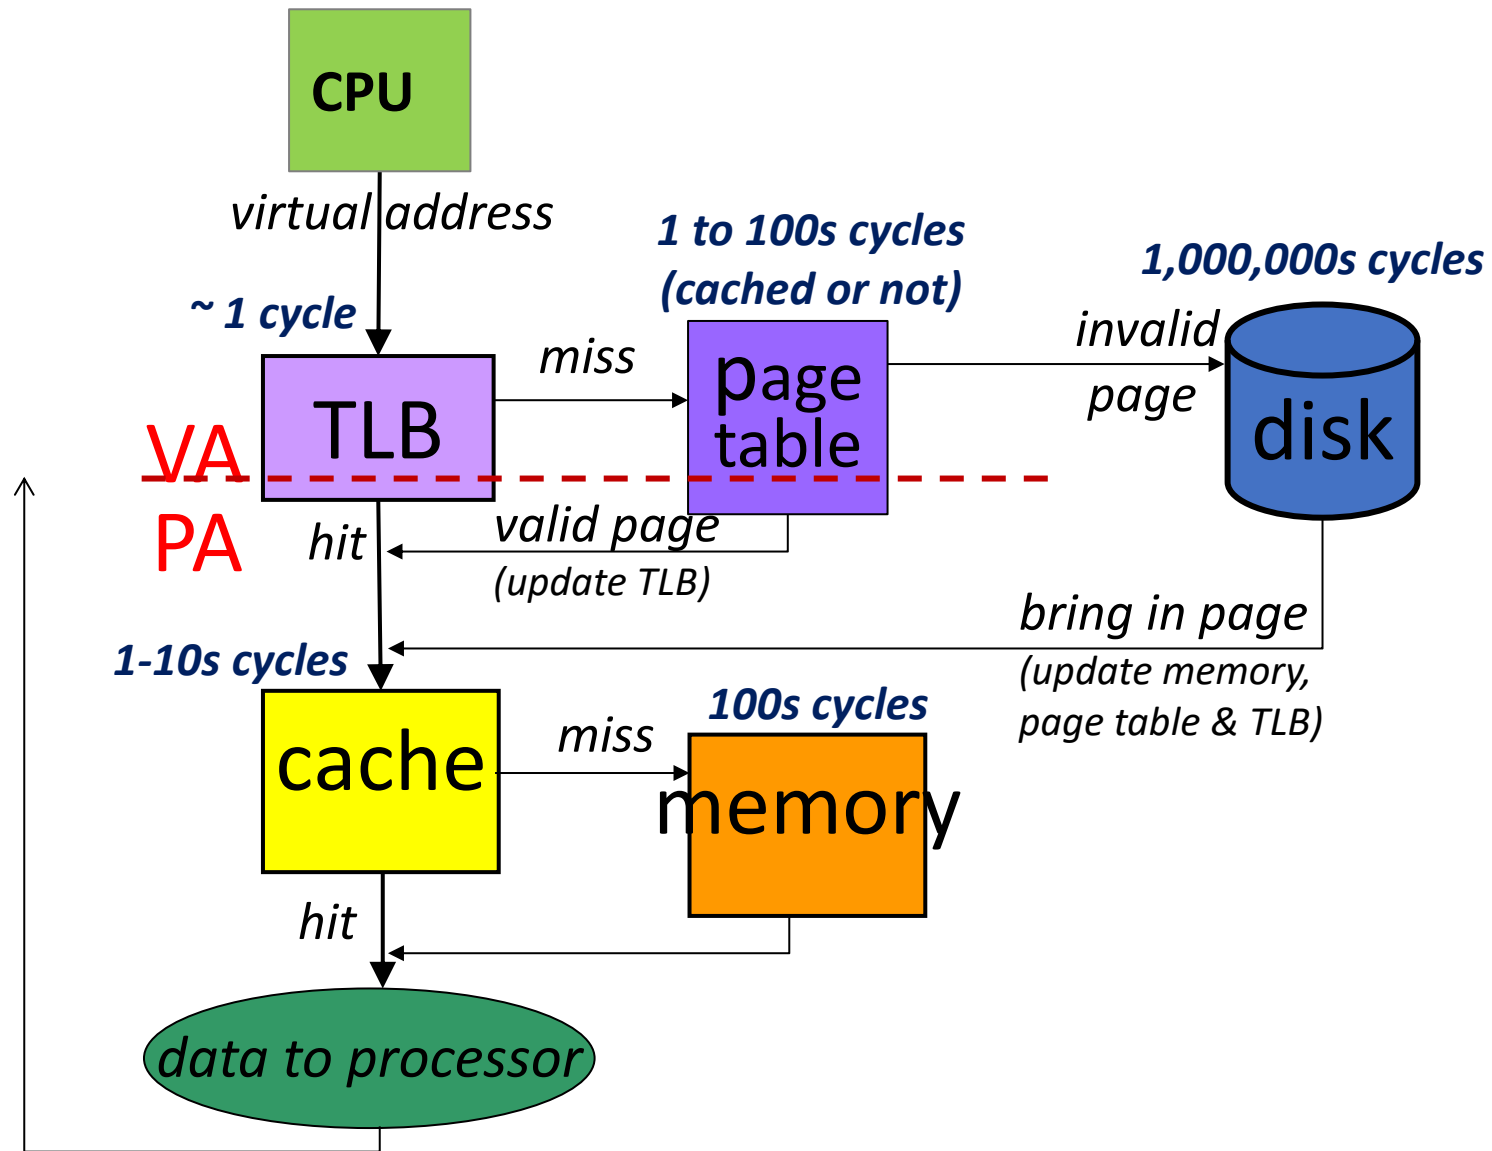
Low complexity



Physically addressed caches

- Perform TLB lookup **before** cache tag comparison
 - Use bits from **physical** address to access cache (tag, set index, and block offset bits)
- Slower access?
 - Tag lookup takes place **after** the TLB lookup
- Simplifies some VM management
 - When switching processes, TLB must be invalidated, but cache OK to stay as is
 - Implications? Might result in fewer cache misses if context switches very common (but they generally are not)


Physically addressed caches: detailed flow



Cache & VM Organization: option 2

Virtually-addressed Cache

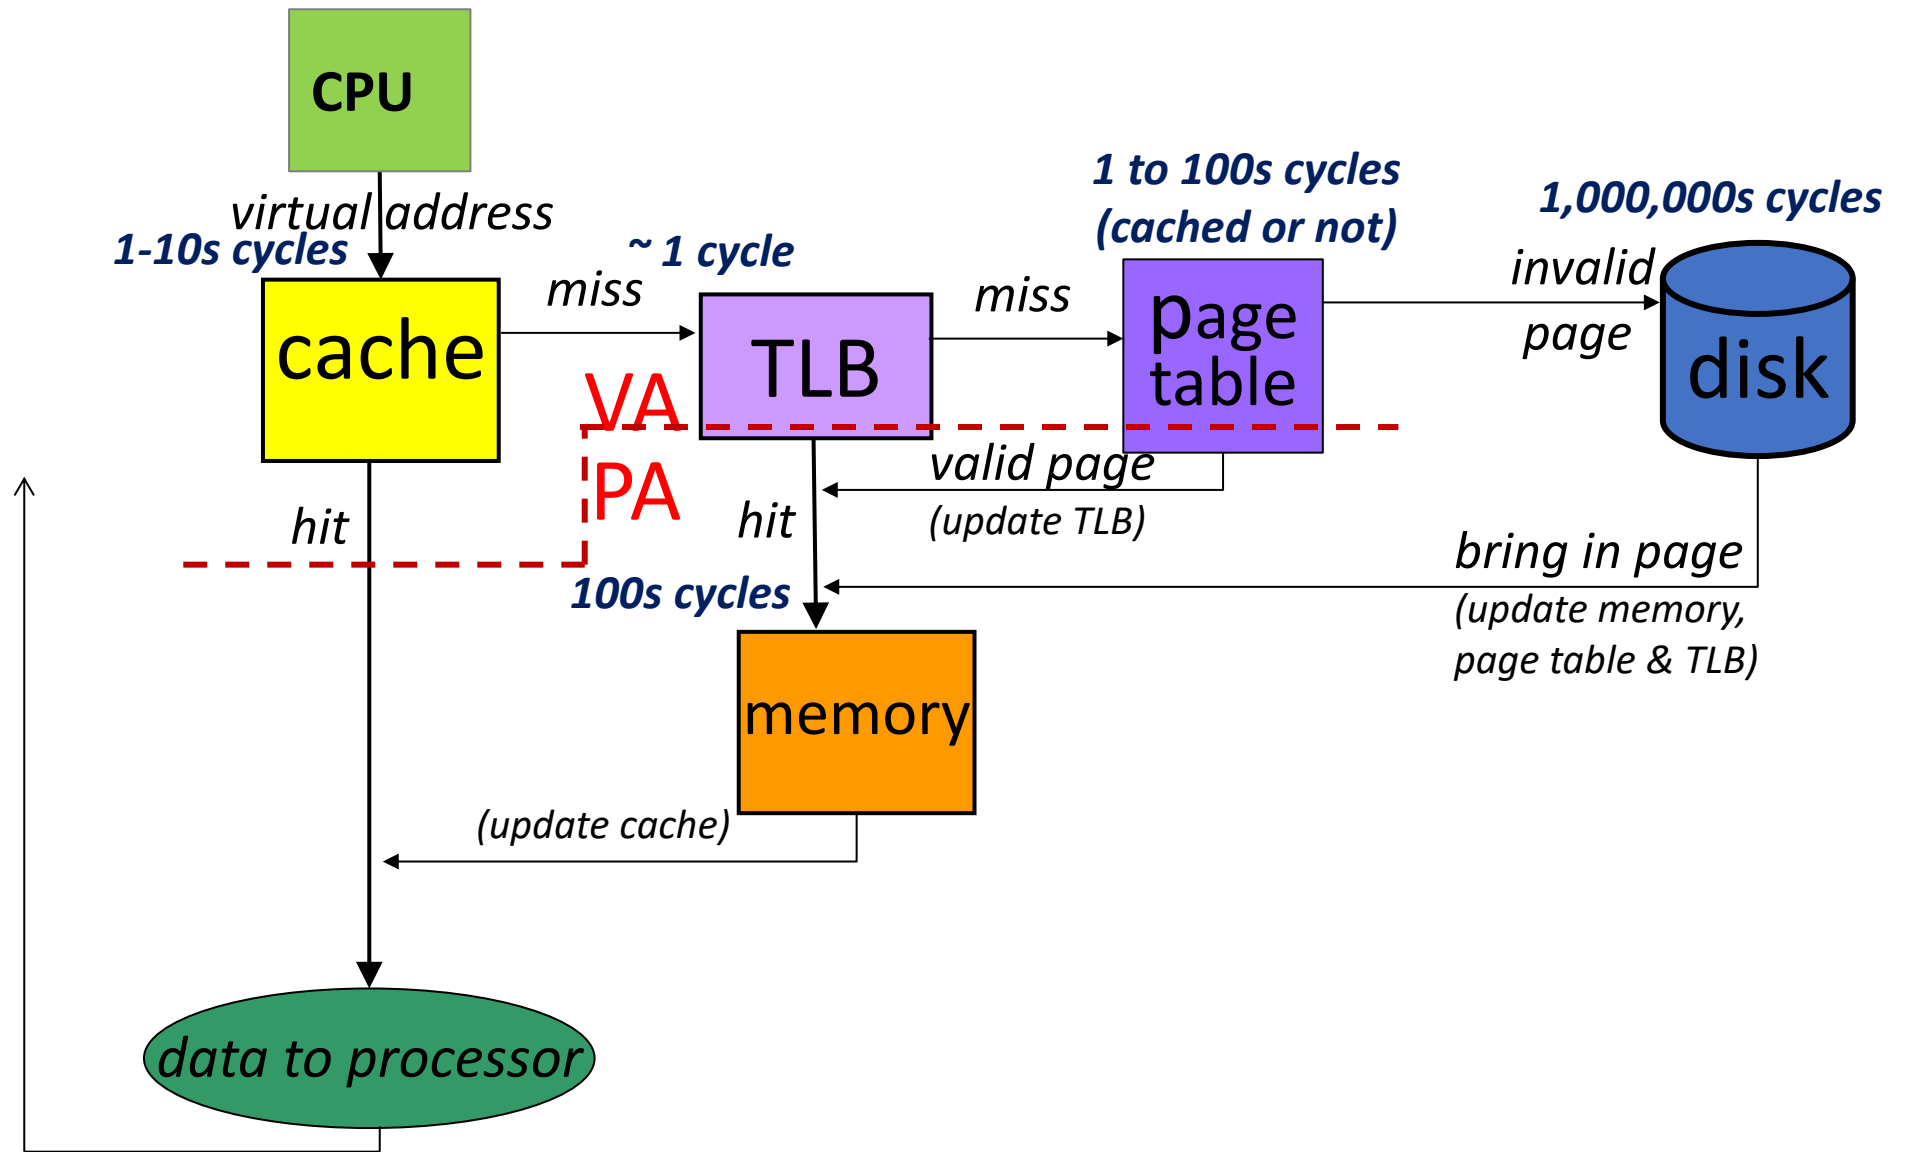
- ✗ High complexity (*aliasing*)
- ✓ Faster



Virtually addressed caches

- Cache uses bits from the virtual address to access cache (tag, set index, and block offset)
- Perform the TLB only if the cache gets a miss.
- Problems:
 - Aliasing: Two processes may refer to the same physical location with different virtual addresses (synonyms)
 - Two processes may have same virtual addresses with different physical addresses (homonyms)
 - When switching processes, TLB must be invalidated, dirty cache blocks must be written back to memory, and cache must be invalidated to solve homonym problem

Virtually addressed caches: detailed flow



OS Support for Virtual Memory

- It must be able to modify the page table register, update page table values, etc.
- To enable the OS to do this, **BUT** not the user program, we have different execution modes for a process.
 - **Executive** (or **supervisor** or **kernel** level) permissions and
 - **User level** permissions.

End of Virtual Memory

Extra Problems



Exercise using previous multi-level VM

Virtual Address	1 st level	2 nd level	Page offset	Page fault?	Physical page num.	Physical Address
0x000F0C						
0x001F0C						
0x020F0C						

1st level = 7b	2nd level = 8b	Page offset = 9b
----------------------------------	----------------------------------	-------------------------

Virtual address = 24b

Physical page number = 9b	Page offset = 9b
----------------------------------	-------------------------

Physical address = 18b

Assume memory for page tables are “somewhere else” in memory

Exercise using previous multi-level VM

Virtual Address	1 st level	2 nd level	Page offset	Page fault?	Physical page num.	Physical Address
0x000F0C	0x00	0x07	0x10C	Y	0x000	0x0010C
0x001F0C						
0x020F0C						

1 st level = 7b	2 nd level = 8b	Page offset = 9b
----------------------------	----------------------------	------------------

Virtual address = 24b

Physical page number = 9b	Page offset = 9b
---------------------------	------------------

Physical address = 18b

Assume memory for page tables are “somewhere else” in memory

Exercise using previous multi-level VM

Virtual Address	1 st level	2 nd level	Page offset	Page fault?	Physical page num.	Physical Address
0x000F0C	0x00	0x07	0x10C	Y	0x000	0x0010C
0x001F0C	0x00	0x0F	0x10C	Y	0x001	0x0030C
0x020F0C						

1 st level = 7b	2 nd level = 8b	Page offset = 9b
----------------------------	----------------------------	------------------

Virtual address = 24b

Physical page number = 9b	Page offset = 9b
---------------------------	------------------

Physical address = 18b

Assume memory for page tables are “somewhere else” in memory

Exercise using previous multi-level VM

Virtual Address	1 st level	2 nd level	Page offset	Page fault?	Physical page num.	Physical Address
0x000F0C	0x00	0x07	0x10C	Y	0x000	0x0010C
0x001F0C	0x00	0x0F	0x10C	Y	0x001	0x0030C
0x020F0C	0x01	0x07	0x10C	Y	0x002	0x0050C

1 st level = 7b	2 nd level = 8b	Page offset = 9b
----------------------------	----------------------------	------------------

Virtual address = 24b

Physical page number = 9b	Page offset = 9b
---------------------------	------------------

Physical address = 18b

Assume memory for page tables are “somewhere else” in memory

What's next for You?!

- What classes should you look at if your interested in this kind of stuff?

More Classes?

- EECS 470 – Computer Architecture
- Picks up where we leave off in 370
- Discuss more sophisticated enhancements to processor design
 - Out-of-order execution
 - Multi-core / multi-threaded processors
- Major Design Experience / Capstone
 - Work in teams to design an actual processor (not simulator) over the course of the semester
 - **LOT'S** of work
- Requires EECS 270
 - For Verilog (a Hardware Description Language)

More Classes?

- EECS 373 – Embedded Systems
 - Learn about parts of computer systems that aren't in the processor or cache
 - Input/output
 - Timers
 - Bus interfaces
 - Followed up by 473 (MDE / Capstone)
- EECS 471 – Applied Parallel Programming with GPUs
 - How are graphics processing units different than normal processors?
 - How do we take advantage of their raw power when writing software?

More Classes?

- EECS 427 – VLSI Design
 - Design a processor at the circuit level
 - Actually layout all the transistors
 - Less architecture – more circuit design
- EECS 570 (Parallel Computer Architecture) + 573 (Microarchitecture)
 - Learn about more current research into architecture
 - Requires 470

More Classes?

- EECS 483 (Front-end compilers) + 583 (Back-end compilers)
 - 583 is more relevant to architecture
 - How to design compilers to write efficient assembly?
 - Requires knowledge of hardware optimizations
- EECS 482 – Operating Systems
 - How do modern systems support multiple threads running simultaneously?
 - How does OS manage virtual memory?
 - How do file systems work?
- EECS 388 – Computer Security
 - How do things like caches "leak" information about a program's data?
 - How can call stacks be tricked into executing arbitrary code?

More Classes?

- EECS 479 – Introduction to Quantum Computing
 - Winter 25
 - How can we redesign the computing stack to take advantage of the bizarre rules of quantum mechanics?
 - Rather than having bits be 0 or 1, have quantum bits be in **superpositions** of 0 and 1
 - Perform massive number of computations simultaneously... sort of
 - Requires 370

Other questions?!