

EECS 388



Introduction to Computer Security

Lecture 3:

Randomness and Pseudorandomness

September 5, 2023

Prof. Halderman



Review: Message Integrity




Problem: Integrity of message from Alice to Bob over an *untrusted channel*

Approach: Alice must append bits to message that only Alice (or Bob) can make

Ideal solution: Random functions

Practical solution: Pseudorandom functions (PRFs)



$f_k()$ is a PRF if it's practically indistinguishable from a random function (unless you know k)

Today's lecture: What are some actual functions that (*we hope*) behave as PRFs?
Where do these random keys k come from?
What else are PRFs useful for?

Cryptographic Hashes



Cryptographic Hash Function

Fixed function $H()$. **No key!**

Input: arbitrary length data

Output: fixed size digest (n bits)

Properties of strong hash functions:

Preimage resistance

Given output h , hard to find input m s.t. $h = H(m)$

Second-preimage resistance

Given input m_1 , hard to find different m_2 s.t. $H(m_1) = H(m_2)$

Collision resistance

Hard to find any pair of inputs m_1, m_2 s.t. $H(m_1) = H(m_2)$

Note: Collisions exist [why?], but should be hard to find

Computing hashes with **OpenSSL**:

```
$ echo "hello world 1" | openssl dgst -sha256
07bafbe63a0e7c57c572aedf1c228022
b537e28785013d7be017fc78731a8cc5
$ echo "hello world 0" | openssl dgst -sha256
8bfa9a398e97152beaaf385847808ad2
d828c1c7251f1a45bc7697723827e7e7
```

Observe how changing even a single bit of the input produces output that appears completely unrelated.

Annoying question:

Are existing hashes *actually* strong?

Annoying answer:

We don't know.

Candidates: MD5, SHA-1, **SHA-256**, SHA-512, SHA-3

Constructing SHA-256



SHA-256 is a widely used hash function that is currently thought to be strong

Input: arbitrary length data

Output: 256-bit digest


Built from a **compression function h** :

Inputs: (256 bits, 512 bits), **Output:** 256 bits

“compresses” 768 bits to 256 bits using a complex internal function (details out of scope for 388)

Uses the **Merkle–Damgård (MD) construction** (illustrated below) to accept arbitrary-length input by repeatedly applying $h()$:

1. Pad input m to the next multiple of 512 bits (adds at least 1 bit, uses fixed algorithm [\[why?\]](#)) and split into 512-bit blocks: b_0, b_1, \dots, b_{n-1}
2. $y_0 := \langle 256\text{-bit constant} \rangle$
 $y_1 := h(y_0, b_0) \dots y_i := h(y_{i-1}, b_{i-1})$
3. Return y_n which is defined to be **SHA-256(m)**



MD Hash Pitfall: Length Extension Attacks



Merkle–Damgård hash functions are susceptible to **length extension attacks**:

Given $y = H(x)$ for some unknown x , attackers can calculate

$$z = H(x \parallel \text{padding} \parallel s)$$

for arbitrary s .

concatenation

That is, given:

$$y = \text{SHA-256}(\text{???})$$

An attacker can produce:

$$z = \text{SHA-256}(\text{???} \parallel \text{original pad} \parallel \text{suffix})$$

Note that this doesn't violate preimage, second-preimage, or collision resistance.

[But why is it a problem?]

Suppose Alice and Bob use this as a verifier:

$$v := \text{SHA-256}(k \parallel m)$$



1. Alice sends $m = \text{"Please go to the bank."}$
2. Mallory doesn't know k , but *can* apply length extension to (m, v) to calculate v' for:
 $m' = \text{"Please go to the bank.} \parallel \text{original_pad} \parallel \text{"Then transfer \$10,000 to Mallory."}$
(*original_pad* is some bytes beyond Mallory's control, but which the recipient might ignore)
3. Since v' is the correct verifier for m' , Bob will accept the modified message as valid

You'll explore how this is done in Project 1!

Practical Solution: HMAC



Message Authentication Code (MAC)

Designed to be used as a secure verifier:

Inputs: **key**, arbitrary length data

Output: fixed size digest (n bits)

HMAC construction turns any secure hash function $H()$ into a MAC:

$$\text{HMAC}_k(m) = H(\underbrace{k \oplus c_1}_{\text{XOR}} \parallel \underbrace{H(\underbrace{k \oplus c_2}_{\text{constant } 363636\dots} \parallel \underbrace{m}_{\text{concatenation}})}_{\text{constant } 5c5c5c\dots})$$

Design protects against length extension!


Example: **HMAC-SHA-256** is an HMAC constructed using **SHA-256** for $H()$

For practical purposes, we believe we can treat **HMAC-SHA-256** as a PRF

Can *reduce* PRF security of HMAC-SHA-256 to a (weaker) security property of SHA-256's compression function

```
$ echo "hi" | openssl dgst -sha256 -hmac Secr3t
8074cdfd007e5cfdc71c2c1cd393a5fe
fa890d7702956a1366a155d79d1cbe77
```

At last, this gives Alice and Bob a suitable approach for protecting message integrity:



Randomness and Pseudorandomness



How should we choose a “secret” key k ?

Select a uniform random value [Why?]

Careful: People are often sloppy about what is “random”

True Randomness

Output of a *physical process* that is inherently unpredictable

Inherent in all physical systems (Heisenberg)

Absent in abstraction of a Turing machine

True randomness in software is scarce.
Getting it requires careful engineering.

Can't just look at the *output* of a process and determine that it was random

Requires assessing the **generation procedure**

```
int getRandomNumber()  
{  
    return 4; // chosen by fair dice roll.  
             // guaranteed to be random.  
}
```

Example: **Extract an unbiased random bit from a biased coin** (Von Neumann's method)

Since true randomness is expensive, often want to take a *small amount of it* and create a longer sequence that's “as good as random”

Pseudorandom Generator

Definition of a PRG



Review: Pseudorandom Function (PRF)

$$f_k: \{0,1\}_m \rightarrow \{0,1\}_n$$

m-bit input and **n**-bit output

Cannot practically distinguish $f_k(x)$
from random without knowing **k**

Pseudorandom Generator (PRG)

$$g_k: \perp \rightarrow \{0,1\}_n \text{ for } n = \text{poly}(|k|)$$

k is a truly random seed

No other inputs

Output is much larger than the input
(Can think of output as *stream of bits*)

Cannot practically distinguish $g_k()$ from a
random stream of bits without knowing **k**

Security definition: (Similar to PRF definition)

1. Let **k** be a secret seed
2. Toss a coin (in secret) to get bit **b**
3. If **b**=0, **s** := a truly random stream
If **b**=1, **s** := $g_k()$
4. Output **s** to Mallory
5. Mallory guesses **b** in *polynomial time**

We say $g()$ is a **secure PRG** if Mallory can't do
meaningfully better than random guessing

* **Usual complexity-theoretic caveats...** no brute-forcing

Annoying question again:

Do PRGs actually exist?

Same annoying answer:

We don't know. (Would imply $P \neq NP$.)

Building a PRG from a PRF



Given a secure PRF $f()$, we can construct a PRG $g()$ and prove that it's secure.*

Construction:

For some random key k and PRF $f()$,
define $g_k() := f_k(0) \parallel f_k(1) \parallel f_k(2) \parallel \dots$

Theorem (informal):

If $f()$ is a secure PRF and
 $g()$ is built from $f()$ by this construction,
then $g()$ is a secure PRG.

You're not responsible for the proof,
but here it is (*at right*), if you're interested:

(*Hard exercise: Build a PRF from a PRG)

Proof by contradiction: If $g()$ as constructed is *not* a secure PRG, $f()$ *can't* be a secure PRF:

1. If $g()$ is not a secure PRG, there exists an algorithm M^* that can distinguish its output from random. (This is from the definition of a secure PRG.)
 2. We can apply M^* to construct an algorithm M to distinguish $f()$ from a random function $r()$:
 - a. Query $h()$ with inputs 0, 1, 2, ...
 - b. Let $s := h(0) \parallel h(1) \parallel h(2) \parallel \dots$
 - c. Apply M^* to s and return the result.
 3. If $h()$ is $r()$, s is a random stream, so M outputs 0. If $h()$ is $f()$, s is $g()$ (by construction), so M outputs 1.
- Thus, M wins the PRF game for $f()$.

This contradicts our assumption that $f()$ is a secure PRF. Therefore, $g()$ must be a secure PRG.

Getting Randomness

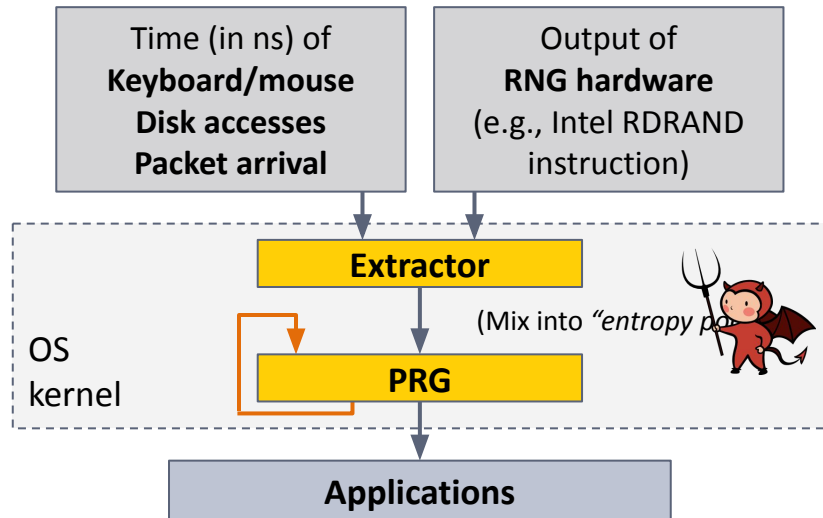


Randomness is **an input** to your program

Typically provided by the OS, via special APIs

OS continuously gathers inputs from physical sources that are hard for adversary to predict, “extracts” uniform bits from them.

[What if attacker can predict *some* of them?]



[What if an attacker learns internal state?]

If compromise is **transient**, can recover by adding more randomness.

How quickly can we recover, though?

Getting Randomness

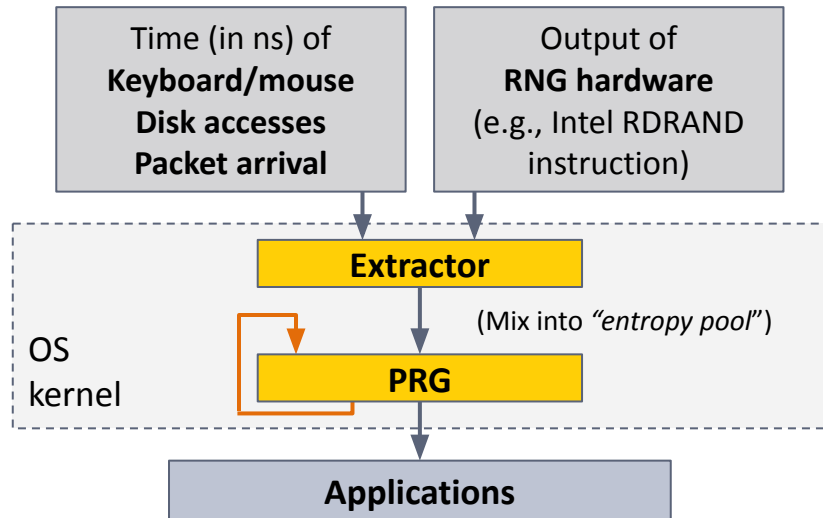


Randomness is **an input** to your program

Typically provided by the OS, via special APIs

OS continuously gathers inputs from physical sources that are hard for adversary to predict, “extracts” uniform bits from them.

[What if attacker can predict *some* of them?]



Generating 16 random bytes with **openssl**:

```
$ openssl rand -hex 16
4345b3cccecb66bef87f9289d72b8d2a

$ openssl rand -hex 16
578fbc36ad4d88ab7d98d7fd16e3737d
```

Caution! Not all “random” APIs are secure (or even unpredictable!) C’s `rand()` function is notoriously bad, as are typical math packages.

Use APIs specified as suitable for cryptography:

C (Linux): `#include <sys/random.h>`
`getrandom(buf, size, 0);`

Python: `import secrets`
`data = secrets.randbits(256);`

JavaScript: `const array = new Uint8Array(32);`
`self.crypto.getRandomValues(array);`

Randomness as an Attack Target

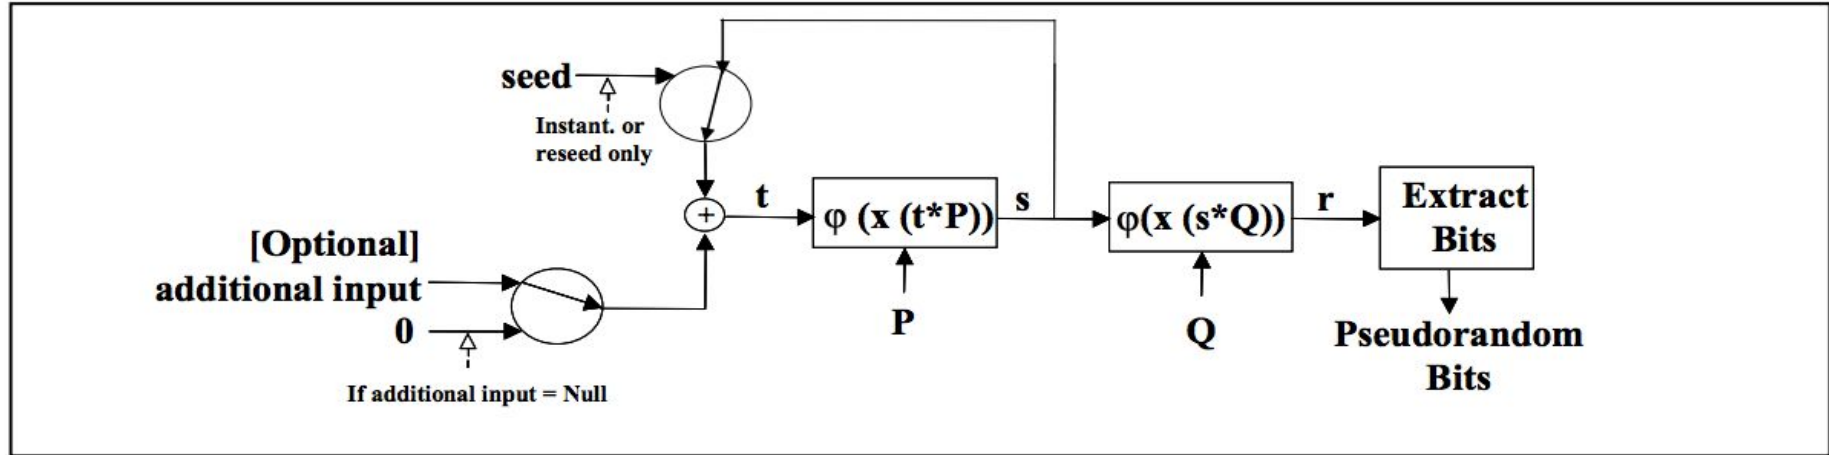


Good randomness is needed everywhere in cryptography.
RNG is very good attack target!

Dual-EC DRBG: 2006 NIST standard that NSA (allegedly)
backdoored. Evidence in Snowden documents.



Construction allows for the existence of a secret backdoor key that can be used to recover the internal RNG state (and determine future output) given knowledge of small amount of past output.



Reminders:

Quiz on Canvas after every lecture

Lab Assignment 1 due Thursday at 6 p.m.

Project 1, Part 1 due Sept. 14 at 6 p.m.

Thursday

Confidentiality

Simple ciphers,
AES,
block cipher modes

Next Week

Wrap Up Crypto Unit

Combining confidentiality and
integrity; Diffie-Hellman,
RSA encryption, digital signatures