

EECS 388



Introduction to Computer Security

Lecture 5:

Combining Confidentiality and Integrity

September 12, 2023

Prof. Halderman



Padding and Block Cipher Modes



Challenge for block ciphers:

How to encrypt **arbitrary-sized messages**?

Padding: Add bytes to end of message to make it a multiple of block size

Flawed approach: add zeros [What's the issue?]

| MM MM MM MM MM 00 00 00 |

Don't know what to remove after decryption!

Better approach (**PKCS7**): Add **n** bytes of value **n**

| MM MM MM MM MM 03 03 03 |

Edge case: Message that ends at block boundary?

| MM MM MM MM MM MM MM MM | 08 08 08 08 08 08 08 08 |

Add an **entire block** of padding

Ensures receiver can **unambiguously** distinguish the padding from the message after decrypting

Cipher modes: Algorithms for applying block ciphers to more than one block

Flawed approach: [What's the issue?]

Encrypted codebook (ECB) mode

Simply encrypt each block independently: $c_i := E_k(p_i)$



Plaintext

Pseudorandom

ECB mode

Cipher Modes



Cipher-block chaining (CBC) mode

“Chains” ciphertexts to obscure later ones

Choose a random **initialization vector IV**

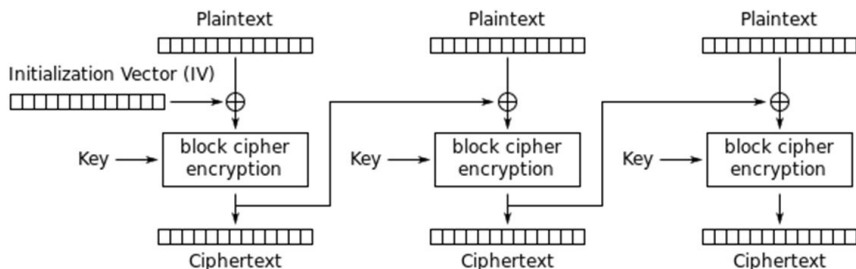
Encrypt: $c_0 := IV; c_i := E_k(p_i \oplus c_{i-1})$

Decrypt: $p_i := D_k(c_i) \oplus c_{i-1}$,

[Why do we need the IV?]

Have to send IV with ciphertext

Can't encrypt blocks in parallel or out of order



Counter (CTR) mode

Turns a block cipher into a stream cipher

Generate **keystream** s for k and unique **nonce**:

$s := E_k(\text{nonce} \parallel 0) \parallel E_k(\text{nonce} \parallel 1) \parallel$

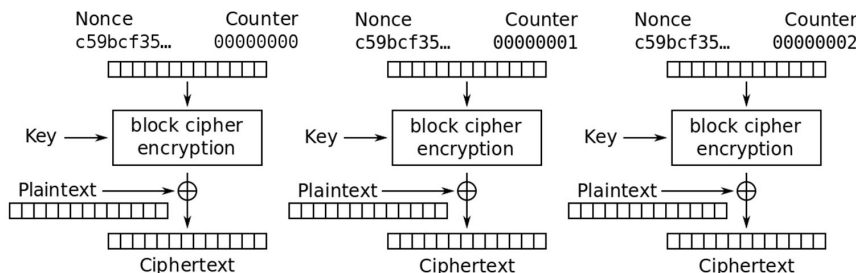
$E_k(\text{nonce} \parallel 2) \parallel \dots$

Encrypt: $c := p \oplus s$ Decrypt: $p := c \oplus s$

Benefits: Doesn't require padding

Efficient parallelism/random access

Caution: Never reuse nonce for same k !



Review: Integrity and Confidentiality



Integrity (tampering)

Let $f()$ be a secure PRF.

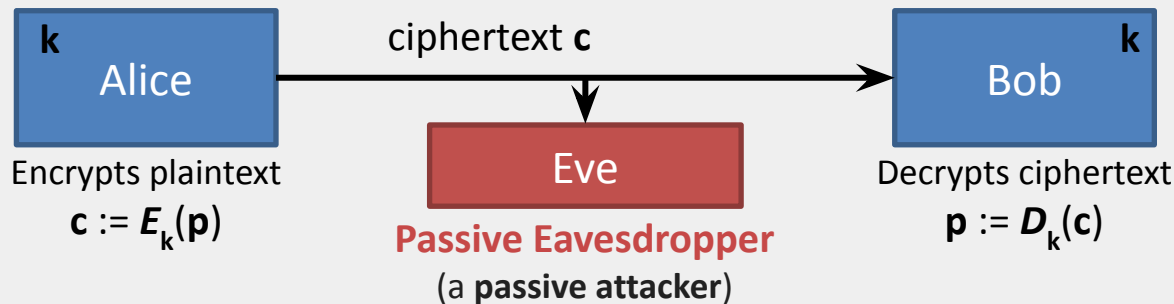
In practice: e.g., **HMAC-SHA-256**



Confidentiality (eavesdropping)

Construct $E()$ and $D()$ from
secure PRG (a stream cipher) **or**
secure PRP (a block cipher) with
appropriate padding/cipher mode

In practice: e.g., **AES-128 in CTR mode**



Today's lecture:

What if we want integrity and confidentiality *at the same time*?

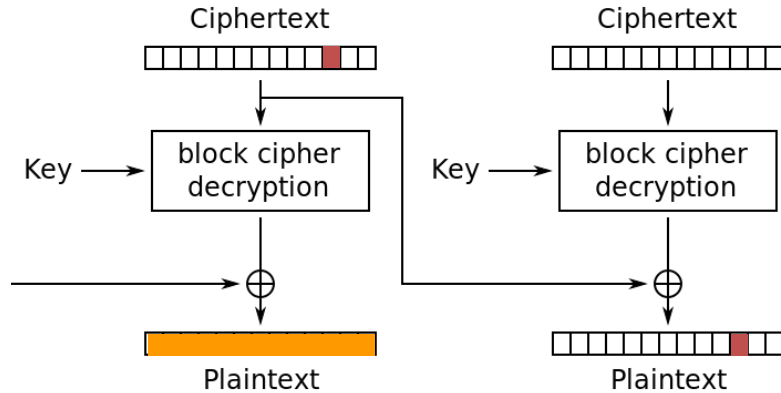
Ciphertext Malleability



Caution: Many encryption methods are **malleable**: can transform a ciphertext into another ciphertext that decrypts to a related plaintext, without knowing the plaintext

Examples:

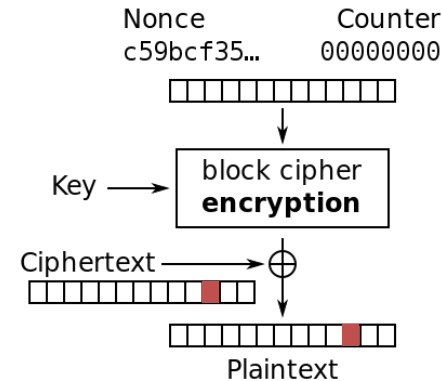
CBC mode decryption



Flipping bits in ciphertext block i will:

- completely corrupt decrypted block i
- flip corresponding bits in decrypted block $i+1$

Counter mode decryption



Flipping bits anywhere in the ciphertext will flip corresponding bits in decrypted plaintext

Need to use other methods to ensure integrity...

Authenticated Encryption (AE)



Two approaches:

1. Generically compose encryption and MAC
2. Build “all-in-one” primitive that does both

Syntax of AE:

$c := E_k(p)$

$p/\text{“fail”} := D_k(c)$



Important difference:

Decryption can fail!

Analogous to Bob rejecting verifier

Security definition:

1. Let k be a secret seed
2. Toss a coin (in secret) to get bit b
3. If $b=0$: $G() := E_k()$; $H() := D_k()$ *
* Rejects previous $E()$ outputs
If $b=1$: $G() :=$ random bits; $H() :=$ “fail”
4. Give Mallory $G()/H()$ oracles
(Mallory gets to repeatedly probe them)
5. Mallory guesses b in *polynomial time*

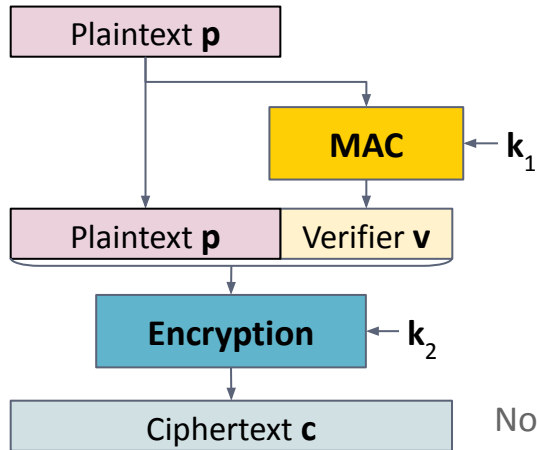
We say AE is **secure** if Mallory can't do meaningfully better than random guessing.

Composing Integrity and Confidentiality

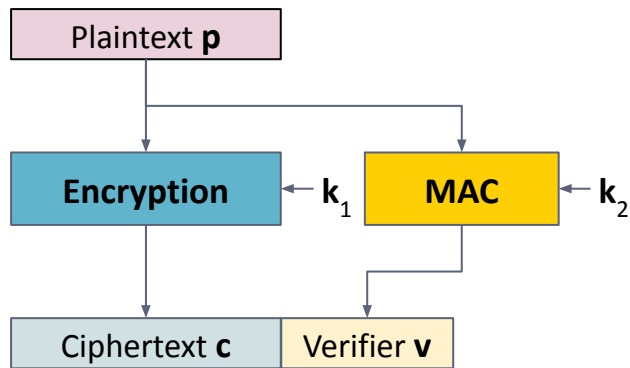


How to **compose** our integrity and confidentiality protocols to achieve both? Three candidates:

MAC-then-Encrypt

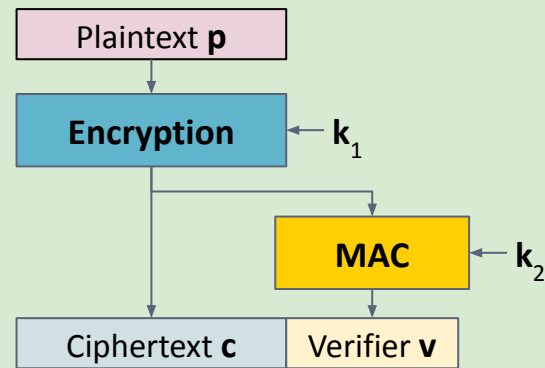


Encrypt-and-MAC



Note: Use separate keys for different purposes.
(Can derive from a single key using PRG.) [Why?]

Encrypt-then-MAC



Safest approach

[Which approach is safest?]

Our encryption methods (so far) only secure against passive eavesdroppers. Only EtM can ensure ciphertext isn't tampered with before decryption.

“Cryptographic Doom Principle”: if you perform any cryptographic operations on a message you've received before verifying the MAC, it will somehow inevitably lead to doom

Example: CBC Padding Oracles



Common flaw when using **MAC-then-Encrypt**

Suppose an implementation uses **CBC mode**.

Decryption involves the following steps:


1. $m \parallel v \parallel \text{pad} := D_k(c)$
2. Check that **pad** is valid **PKCS7**, else raise **PadError**
3. Check that $v = \text{MAC}_k(m)$, else raise **MacError**

This is how TLS 1.0 worked. Seems reasonable?

Any method to distinguish these two error types (even tiny timing differences) **leaks the plaintext!**


Padding oracle: attacker submits any ciphertext and learns if last bytes of plaintext are a valid pad

Example of a **chosen ciphertext attack**



Suppose attacker intercepts c , wants to learn m .

Step 1: Let g be a guess for last byte of block $m[1]$:



Step 2: Send modified ciphertext to padding oracle

If $g = \text{last_byte}$: $g \oplus \text{last_byte} \oplus 0x01 = 0x01$.

Modified plaintext ends in $0x01$, so padding's valid; oracle returns **MacError**

else: Padding is invalid*, oracle returns **PadError**

(*Except for edge cases: e.g., what if $m[1]$ ends in $0x02\ 0x01$ and $g = 0x02$?)

Step 3: Repeat with $g = 0, 1, \dots, 255$ to learn last_byte . Then use a $0x02, 0x02$ pad to learn next byte, etc.

Lesson: Encrypt *then* MAC You'll exploit in P1!

Authenticated Encryption with Associated Data



Preferred modern approach:

Authenticated encryption with associated data (AEAD)

Integrity and encryption in a single primitive:

$\mathbf{c}, \mathbf{v} := \text{Seal}(\mathbf{k}, \mathbf{p}, [\text{associated_data}])$

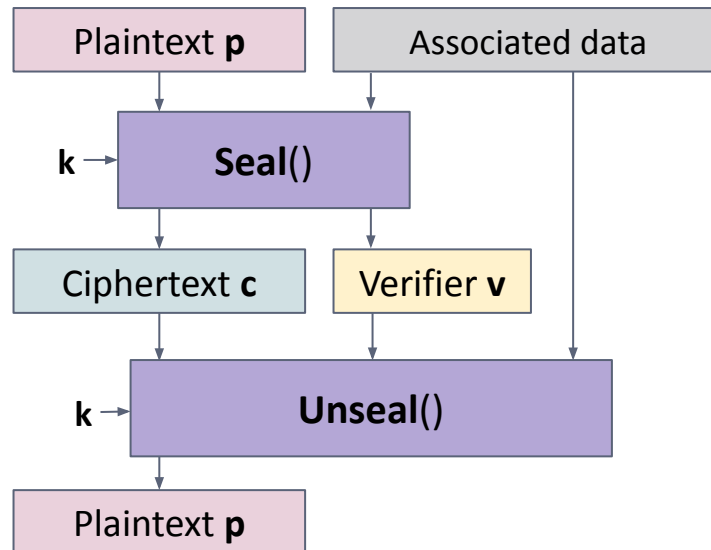
encrypts plaintext \mathbf{p} and returns ciphertext \mathbf{c} and a verifier \mathbf{v} (called a “tag”)

$\mathbf{p}, \text{err} := \text{Unseal}(\mathbf{k}, \mathbf{c}, \mathbf{v}, [\text{associated_data}])$

returns \mathbf{p} or an error if \mathbf{v} does not match the supplied \mathbf{c} and **associated_data**

Optional **associated_data** is covered by verifier but *not encrypted*.

Useful for binding data to its context:
e.g., counter, sender ID, etc.



Examples:

AES-GCM (“Galois Counter Mode”)

hardware accelerated in recent CPUs

ChaCha20-Poly1305, common on mobile

Galois/Counter Mode (GCM)



Galois/Counter Mode (GCM)

Most widely used AEAD cipher mode

Developed by McGrew and Viega in 2004

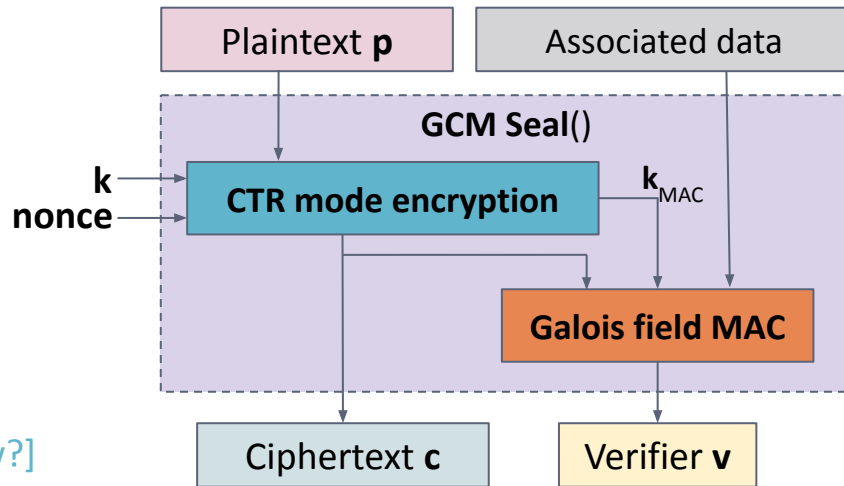
Standardized by IETF, NIST, others

Non-generic composition [What kind?] of AES in CTR mode for encryption and a special MAC based on polynomials over finite (“Galois”) fields.

Note: GCM violates principle of key separation [How?]
(We can prove it’s ok, but it’s delicate.)

Warning: Can construct GCM ciphertexts that decrypt (differently, but without error) under many keys

Prof. Grubbs crafted one that decrypts under 131,072 different keys



Warning: GCM nonce reuse is catastrophic!
Encrypting two ciphertexts with the same $(k, nonce)$ leaks the plaintext **and the MAC key**.

Why? Polynomial root-finding.

Details interesting but beyond our scope in 388.

Parameter Sizes



Issue: How should we set sizes?

Choose $|k|$ to resist brute force attacks,
even as computers become faster.

For ciphers/PRG keys:

- Want to resist exhaustive search for k
- 128 bits considered “classically” safe
($2^{128} \approx 10^{38} \approx$ number of silicon atoms in the earth)
- For quantum-resistance, use 256 bits
(Grover’s algorithm gets attacker “sqrt” speedup)

For hash function outputs:

- Want collision resistance (CR)
- Need 2^n bits of output for n bits of CR,
due to “birthday” attacks
(e.g., SHA-256 has 128 bits of CR)

Estimating what’s feasible to compute?

$$2^{64} \approx 10^{19} \quad 2^{128} \approx 10^{38} \quad 1 \text{ year} \approx 3 \times 10^7 \text{ s}$$

CPU mining: $\approx 10^8$ SHA-256/s

GPU mining: $\approx 10^{11}$ SHA-256/s

ASIC mining: $\approx 10^{14}$ SHA-256/s

Bitcoin miners globally: 10^{20} SHA-256 blocks/s

“Birthday” Attack

Generate random values, look for collision

Requires $2^{|k|/2}$ time, $2^{|k|/2}$ space

[Puzzle: Do it in constant space?]

For 128-bit output, takes 2^{64} steps: doable!

Randomness as an Attack Target

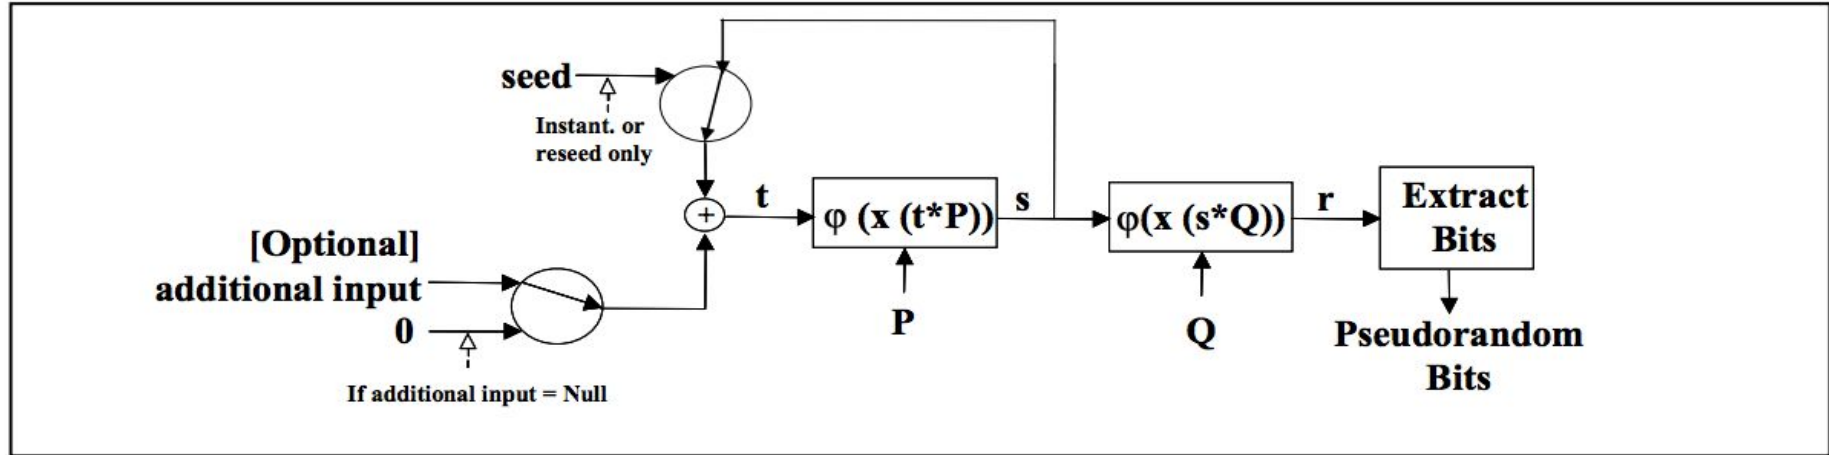


Good randomness is needed everywhere in cryptography.
RNG is very good attack target!

Dual-EC DRBG: 2006 NIST standard that NSA (allegedly)
backdoored. Evidence in Snowden documents.



Construction allows for the existence of a secret backdoor key that can be used to recover the internal RNG state (and determine future output) given knowledge of small amount of past output.



Reminders:

Project 1, Part 1 due Thursday at 6 p.m.

Project 1, Part 2 due 9/21 at 6 p.m.

Quiz on Canvas after every lecture

Thursday

Public Key Cryptography

Diffie-Hellman key exchange,
RSA encryption,
digital signatures

Starting Next Week

Web and Network Security Units