

EECS 388

---



# Introduction to Computer Security

Lecture 18:

## Access Control and Isolation

Nov. 2, 2023

Prof. Ensafi



How do we *control access to resources*?

## Security Model

System abstraction that enables us to discuss and formulate a policy.

## Security Policy

Specifies what actions each subject is allowed to perform on each object.

## Security Mechanism

Implementation of the policy. E.g., OS kernel, encryption

**What we aim for:** **Principle of Least Privilege:** Every program and user should operate using the least *privileges* necessary to complete their job.

- Limits the damage that can result from an accident or error.
- Reduces interactions among privileged programs to minimum for correct operation. Unintentional, unwanted, or improper uses of privilege less likely.
- Minimizes the number of programs that must be audited.
- Example: Classified information concept of “need-to know”.

# Models and Policies



## Security Model

**System abstraction** that enables us to discuss and formulate a policy. Consists of:

### Subjects (who)

UNIX: users

Android: apps

Web: origins

### Objects (what)

UNIX: files, processes

Other: db tables, cookies, device sensors, etc.

### Operations

Ways that subjects can operate on (e.g., read, write, call) objects

## Security Policy

Defines an **access control matrix** that maps subjects and objects to allowed operations.

		Objects			
		File 1	File 2	File 3	File 4
Subjects	Alice	read	read/ write	no access	no access
	Bob	read	read/ write	no access	no access
	Carol	read	write	read/ write	read/ write
	Wendy	read/ write	read/ write	read	read

# Implementing Access Control



What we aim for:

## Principle of Complete Mediation

Every access to every object must be checked for authority by a mediator.

- Primary underpinning of access control protection systems.
- Implies that a foolproof method of identifying the source of every request must be devised.
- Be careful of caching checks!  
If change in authority occurs, cached results must be updated.

**Time-of-check, time-of-use vulnerabilities.**



# Unix Namespace Security Model



## **Subjects (Who?)**

Users

## **Objects (What?)**

Files, directories,  
processes

## **Access Operations**

Read, Write, Execute

# Unix Security Model: Users and Groups



## Subjects (Who?)

Users

Every user has a unique **user name** and numeric user id

```
user@desktop:~$ whoami; id -u  
user  
1000
```

## Objects (What?)

Files, directories,  
processes

A user may belong to several **groups**. Provides **role-based access control**

```
user@desktop:~$ groups  
user adm faculty
```

## Access Operations

Read, Write, Execute

**Superuser** (“root”, id 0) has authority to do anything.  
Least privilege: *Only assume superuser role when necessary.*

Users in a specific group (adm) can temporarily **elevate** to root privileges

```
user@desktop:~$ sudo whoami  
[sudo] password for user:  
root
```

# Unix Security Model: File Permissions



## Subjects (Who?)

Users

## Objects (What?)

Files, directories,  
processes

## Access Operations

Read, Write, Execute

Every file and directory has an owner and group

```
user@desktop:~$ ls -l
total 3
d rwx rwx --- 1 user user 4096 Apr  2 15:56 go
- rw- rw- r-- 1 user user  0 Apr 11 04:15 test.py
d rwx rwx --- 1 user faculty 4096 Dec 28 21:09 courses
```

Permission Bits

Owner

Group

Size

Last Modified

Name

**File permissions bits** specify what role  
can do what to the file:

If user == owner, owner permissions;  
else if user in group, group permissions;  
else other permissions. (If user is root, *override permissions*)

File owner (or root) can change permissions and ownership

**rwX rwX rwX**

Owner

Group

Others

# Unix Security Model: Processes



## Subjects (Who?)

Users

Every process has an **Effective User ID (EUID)** that determines permissions of that process. Processes typically inherit user and group of their parent process, but can be changed by root

## Objects (What?)

Files, directories,  
processes

```
user@desktop:~$ ps -eo user,comm
USER  COMMAND
root  init
user  bash
user  ps
```

## Access Operations

Read, Write, Execute

login and sshd run as root. Authenticate user, then *change* user id and group id to that of user and execute shell (e.g., bash).

Executable files have a **setuid bit**. If set, process has EUID (and privileges) of file's owner, rather than user that executed it.

E.g., passwd command is owned by root and has setuid bit set

**Danger: Any vulnerability in passwd means attacker gets root**



Traditional thinking:

**How to run bad/untrustworthy programs safely?**

- Program from untrusted sites
- Apps exposed to adversarial data
- “Honeypots”

New thinking:

**Be skeptical of all programs, isolate to achieve least privilege**

General goal:

**Confinement:** ensure misbehaving process cannot harm rest of system

Can be implemented at many levels:

- System call interposition:  
(e.g., AppArmor, SELinux)
- Containers:  
isolated userspace instances  
(e.g., Dockers, Kubernetes)
- Virtual machines:  
isolate OSes on a single machine
- Physical isolation:  
E.g., separate hardware, air gaps

# General isolation concept: Confinement



Key component: **reference monitor**

- **Mediates requests** from applications
  - Implements protection policy
  - Enforces isolation and confinement
- Must ***always*** be invoked:
  - Every application request must be mediated
- **Tamperproof:**
  - Reference monitor cannot be killed
  - ... or if killed, then monitored process is killed too
- **Small** enough to be analyzed and validated

# Approach: UNIX chroot system call



**chroot “jails”**: Simple isolation mechanism provided by UNIX kernel

Early use included guest accounts on FTP sites

To use: (must be root)

```
root:/# chroot /tmp/guest    root dir is now /tmp/guest
root:/# su guest              EUID set to guest
guest:/$ ./myapp
```

In this example, myapp is confined such that it can only access files within /tmp/guest for which guest has permission

`open("/etc/passwd", "r")` actually opens `/tmp/guest/etc/passwd`

Application cannot access files outside of jail because it cannot even name them!

# Many ways to evade chroot isolation



- Create device that lets you access raw disk  
Major/minor device number namespace is shared
- Send signals to non-chrooted process  
PID namespace is shared
- Reboot system
- Bind to privileged ports  
Network namespace is shared

# System call interposition



**Observation:** to damage host system (e.g. persistent changes) app must make system calls:

- To delete/overwrite files: **unlink, open, write**
- To do network attacks: **socket, bind, connect, send**

**Idea:** monitor app's system calls and block unauthorized calls

**Implementation options:**

- Completely kernel space (e.g., SE Linux)
- Completely user space (e.g., program shepherding)
- Hybrid (e.g., Systrace)

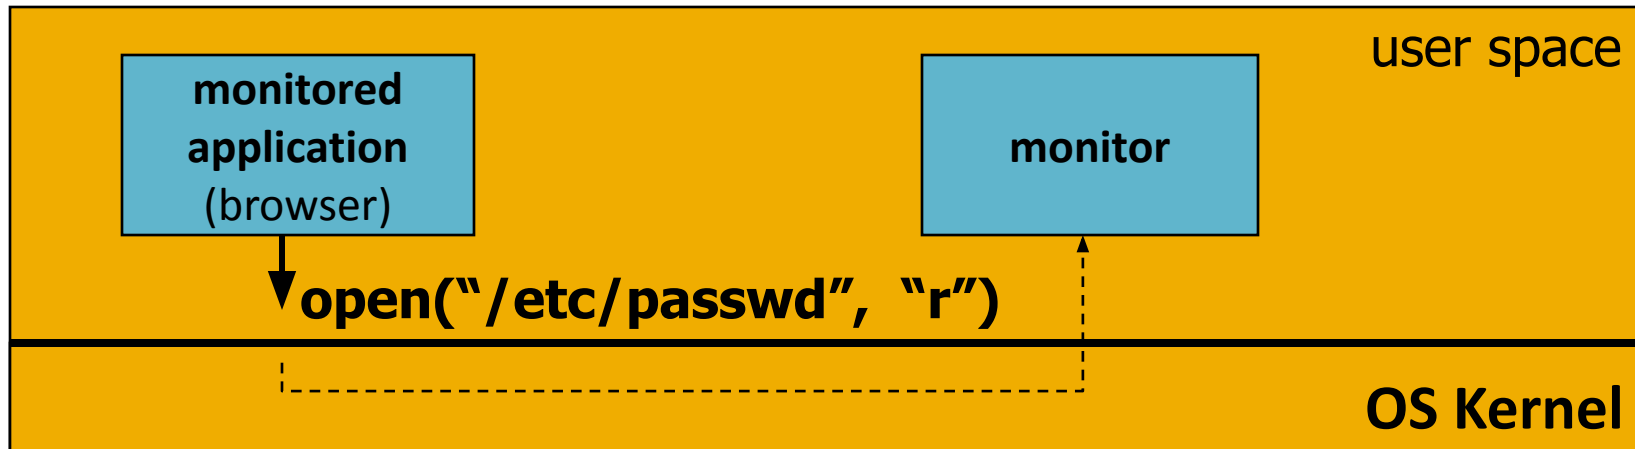
# Implementing System Call Interposition



Linux **ptrace**: process tracing

process calls: **ptrace** (... , pid\_t pid, ...)

and wakes up when **pid** makes a system call.



Monitor checks policy, kills application if request is disallowed.

# System Call Interposition Policies



Sample policy file:

```
path allow  /tmp/*  
path deny  /etc/passwd  
network deny all
```

Manually specifying policy for an app can be difficult:

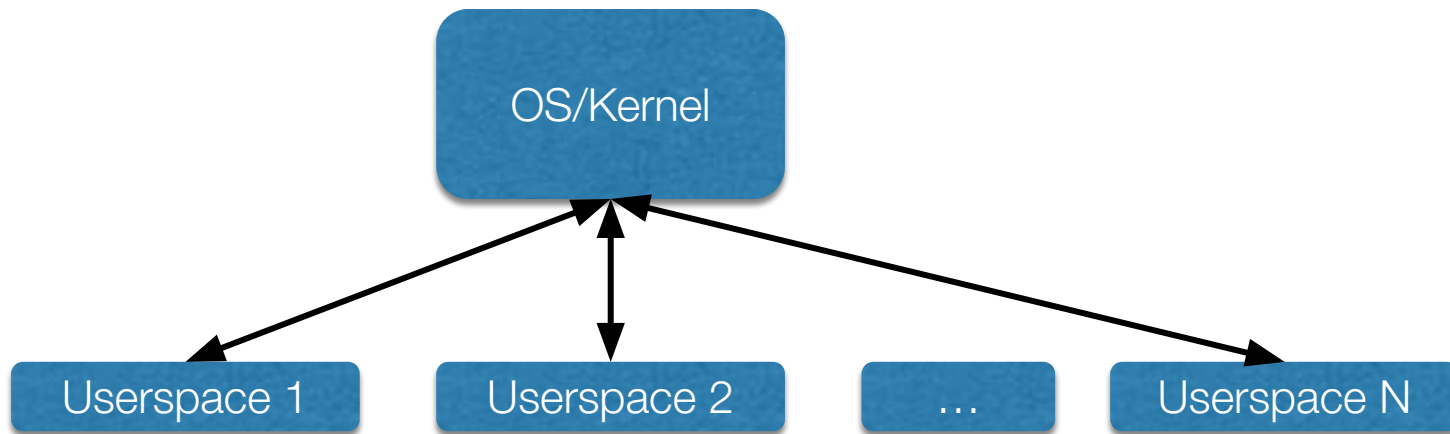
- Can try to auto-generate policy by learning how app behaves on “good” inputs. (But what are “good” inputs?)
- If policy does not cover a specific system call, ask user (But how is the user supposed to decide?)

Difficulty with choosing policy for specific apps (e.g. browser) is the main reason this approach is not widely used

# Containers: Further limiting process interference



A virtualization at the level of the operating system which creates multiple isolated userspace instances (e.g. Docker and Kubernetes)



Restrictions: must be same OS/kernel

Cool things: disk quotas, CPU/RAM limits, root privilege isolation, etc...




# Approach: Virtual Machines



**Confinement:** ensure misbehaving app cannot harm rest of system

Can be implemented at many levels:

**Virtual machines:** isolate OSes on a single machine



## VMM Security assumption:

Malware can infect **guest OS** and guest apps

But malware cannot escape from the infected VM

- Cannot infect **host OS**
- Cannot infect other VMs on the same hardware

Requires that VMM protect itself and is not buggy


- VMM is much simpler than full OS  
... but device drivers run in Host OS

# Problem: Covert Channels



**Covert channel:** unintended communication channel between isolated components

Can be used to leak classified data from secure component to public component



# Covert Channel Example



Both VMs use the same underlying hardware

To send a bit  $b \in \{0,1\}$  malware does:

- $b = 1$ : at 1:00am do CPU intensive calculation
- $b = 0$ : at 1:00am do nothing

At 1:00am listener does CPU intensive calculation and measures completion time

$b = 1 \Rightarrow \text{completion-time} > \text{threshold}$

Many covert channels exist in running system:

- File lock status, cache contents, interrupts, ...
- Difficult to eliminate all


# Approach: Physical Air Gap



**Confinement:** ensure misbehaving app cannot harm rest of system

Can be implemented at many levels:

**Hardware:** run application on isolated hardware (air gap)



⇒ but restrictive, expensive, difficult to manage

# Coming Up



Reminders:

**Lab 4 due today at 6 p.m.**

Quiz on Canvas after each lecture

Tuesday Nov 7

## **Election Cybersecurity**

Vulnerabilities, defenses, policy

Thursday Nov 9

## **Machine Learning Security**

Kexin Pei, University of Chicago