

**EECS 445**

# Introduction to Machine Learning

## Collaborative Filtering

**Prof. Kutty**

# spectral clustering

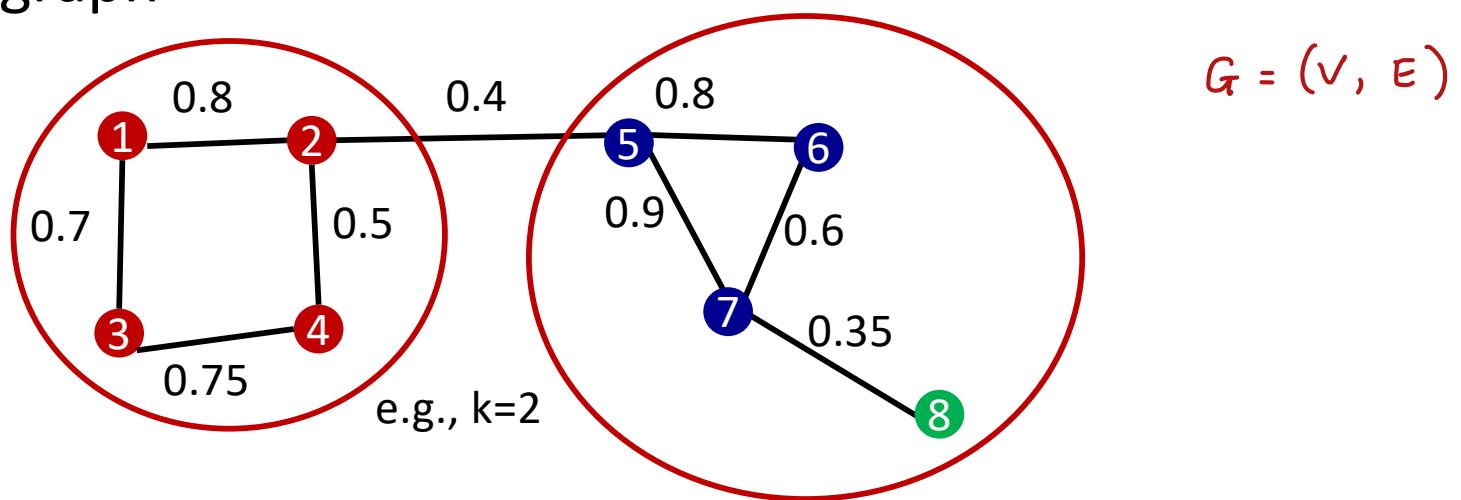
## review

# Spectral Clustering for $k$ partitions

**Input:**  $S_n = \{\bar{x}^{(i)}\}_{i=1}^n$ , valid similarity metric, number of clusters  $k$

Convert to graph

$$\text{e.g. } \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 1 \\ 0 \\ 1 \\ 0 \end{bmatrix}$$



$A_1 = \{4, 5, 7\}$   
solve graph partitioning problem

$$\min_{A_1, \dots, A_k} \text{RatioCut}(A_1, \dots, A_k)$$

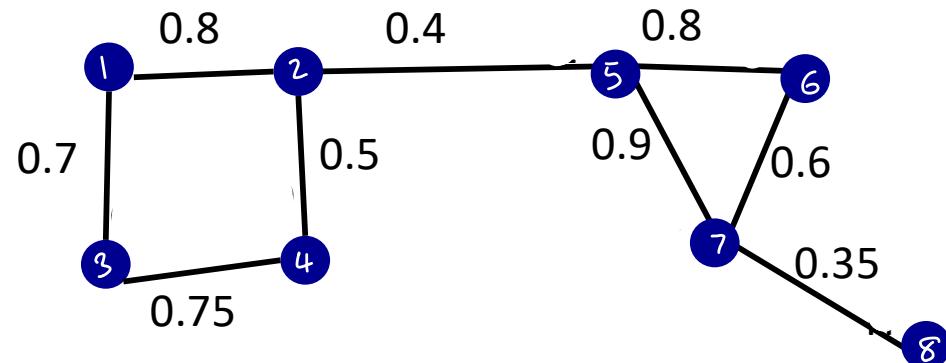
**Output:** clusters assignments

$$\begin{aligned} A_i &\subseteq V \\ i \neq j \quad A_i \cap A_j &= \emptyset \\ \bigcup_{i=1}^k A_i &= V \\ i, j \in \{1, \dots, k\} \end{aligned}$$

# How to find $\min \text{RatioCut}$

- NP hard so solve approximation instead
- Can show that  $\bar{f}^T L \bar{f} \propto \text{RatioCut}(A, \bar{A})$ 
  - graph Laplacian  $L = \mathcal{D} - \mathcal{W}$
  - $k^{th}$  eigenvalue approximates cost of the ratiocut for k clusters
  - $i^{th}$  eigenvector approximates an indicator vector corresponding to the  $i^{th}$  cut
- Note: for more details see Lecture 18, 19 and Discussions 5, 10

# Spectral Clustering Example #2 revisited



Should be able to construct  
W, D, L

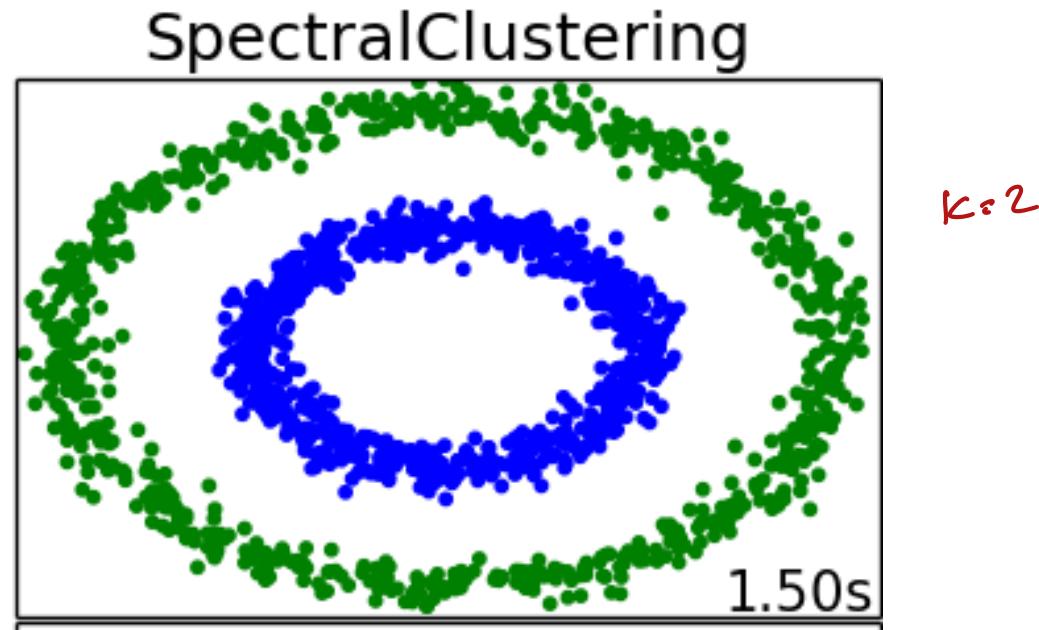
Eigenvectors

1	-0.3536	-0.3467	0.0704	0.6000	-0.2702	0.0594	0.4652	-0.3113
2	-0.3536	-0.2418	-0.0252	0.4060	0.5839	0.0496	-0.3458	0.4337
3	-0.3536	-0.3997	0.1334	-0.2603	-0.6097	-0.1178	-0.4406	0.2233
4	-0.3536	-0.3773	0.1106	-0.6251	0.4141	0.0655	0.3339	-0.2059
5	-0.3536	0.2188	-0.3571	0.0256	0.0968	-0.2652	-0.4338	-0.6566
6	-0.3536	0.2926	-0.4542	-0.1096	-0.1732	0.7125	0.0785	0.1654
7	-0.3536	0.3251	-0.2362	-0.0592	-0.0571	-0.6200	0.4008	0.4021
8	-0.3536	0.5291	0.7582	0.0227	0.0154	0.1160	-0.0583	-0.0509

Eigenvalues

0.0000	0.1349	0.4590	1.2624	1.6492	2.2200	2.7575	3.1170
--------	--------	--------	--------	--------	--------	--------	--------

# Test your understanding: Why does Spectral Clustering work here?



<https://forms.gle/ffiBvNbPjHF8ghi77>

# Thank you for your feedback!

Stay the same:

- Theory + application
- Pacing for in-class notes
- In class problems/exercises and Lecture quizzes

Changes:

- Want more high level overview/recap
- Mid-lecture breaks
- More applications/examples of use cases
- More office hours (?)

# collaborative filtering

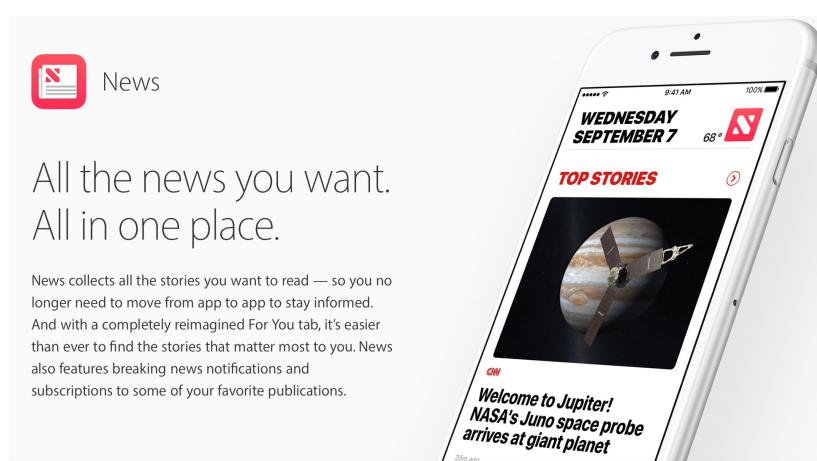
“The most dominant approach for computing recommendations is collaborative filtering...”

Evaluating Recommender Systems: Survey and Framework

Eva Zangerle and Christine Bauer

ACM Computing Surveys 2022

# Recommender systems





## What We Watched: A Netflix Engagement Report

Hours viewed from January to June 2023

Title	Available Globally?	Release Date	Hours Viewed
Blurred background image of a person's face	Yes	2023-03-23	812,100,000
Blurred background image of a person's face	Yes	2023-01-05	665,100,000
Blurred background image of a person's face	Yes	2022-12-30	622,800,000
Blurred background image of a person's face	Yes	2022-11-23	507,700,000
Blurred background image of a person's face	Yes	2023-05-04	503,000,000
Blurred background image of a person's face	Yes	2023-02-09	440,600,000
Blurred background image of a person's face	No	2022-12-30	429,600,000
Blurred background image of a person's face	Yes	2023-02-23	402,500,000
Blurred background image of a person's face	Yes	2021-02-24	302,100,000
Blurred background image of a person's face	Yes	2023-05-25	266,200,000
Blurred background image of a person's face	Yes	2022-11-04	262,600,000
Blurred background image of a person's face	Yes	2023-01-01	252,500,000
Blurred background image of a person's face	Yes	2022-12-02	251,500,000
Blurred background image of a person's face	Yes	2023-05-12	249,900,000
Blurred background image of a person's face	Yes	2023-01-24	235,000,000
Blurred background image of a person's face	Yes	2023-01-14	234,800,000
Blurred background image of a person's face	Yes	2023-03-24	229,700,000
Blurred background image of a person's face	Yes	2023-04-06	221,100,000
Blurred background image of a person's face	Yes	2023-04-20	214,100,000
Blurred background image of a person's face	Yes	2023-03-10	209,700,000
Blurred background image of a person's face	No	2023-05-31	206,500,000
Blurred background image of a person's face	Yes	2023-01-12	205,500,000
Blurred background image of a person's face	Yes	2023-06-16	201,800,000
Blurred background image of a person's face	Yes	2023-05-18	200,700,000
Blurred background image of a person's face	Yes	2023-04-15	194,700,000
Blurred background image of a person's face	Yes	2023-03-16	192,900,000
Blurred background image of a person's face	Yes	2020-04-15	184,000,000

<https://about.netflix.com/en/news/what-we-watched-a-netflix-engagement-report>

# Recommender Systems



## NETFLIX

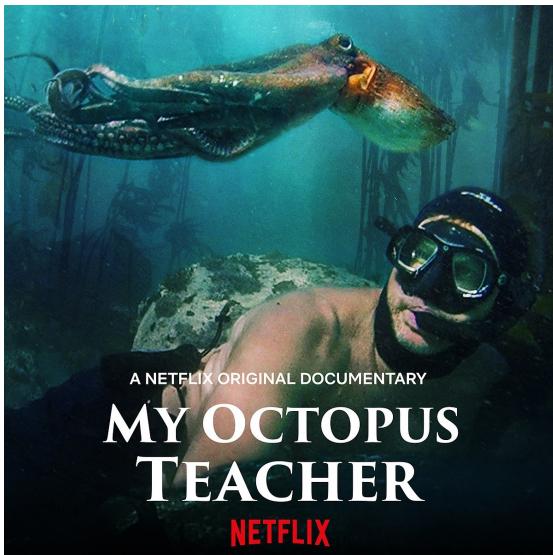


- Launched on January 16, 2007
- ~260.28 million users in more than 190 countries
- data on ~18k titles (representing 99% of all content viewed); record of nearly 100 billion hours viewed
- **Goal:** Predict ratings for titles a user hasn't seen

**Collaborative Filtering** – recommender problems that can be reduced to **matrix completion** (no additional info about the users or about the movies)

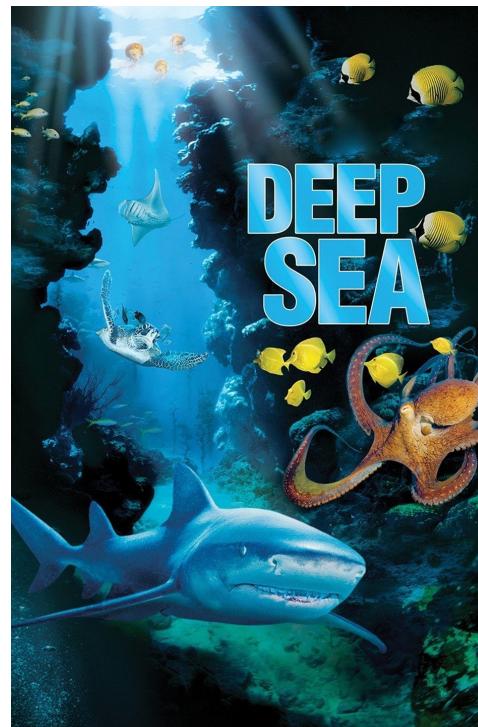
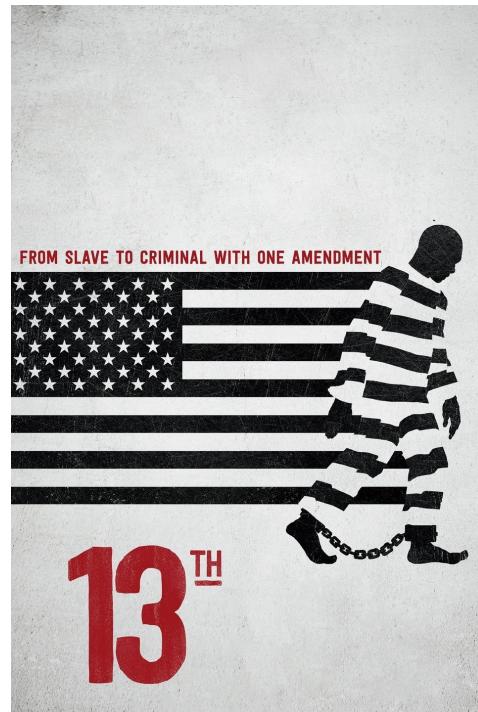
<https://about.netflix.com/en/news/what-we-watched-a-netflix-engagement-report>

if you liked

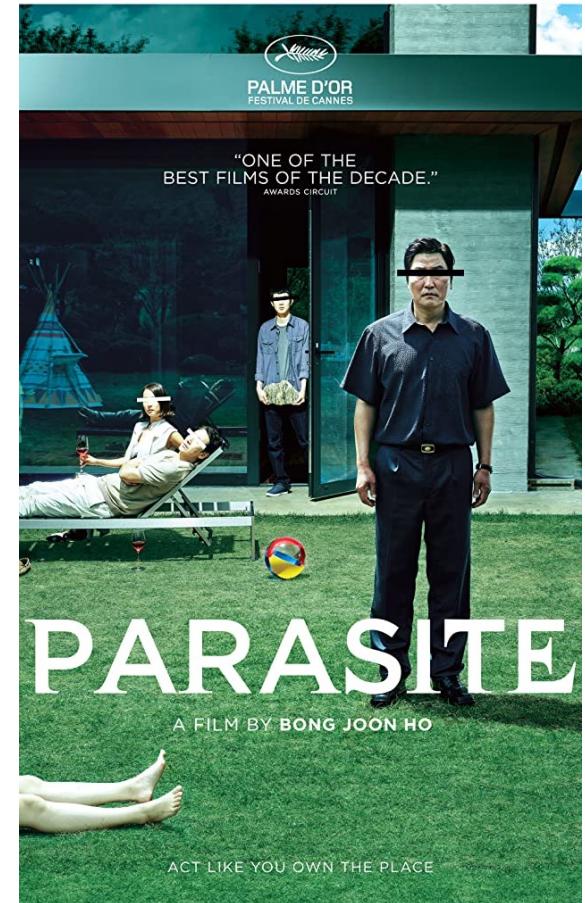


## Steps:

- generate predictions
  - pick movies to present to user



# you might like



# Matrix Completion

$m$  movies

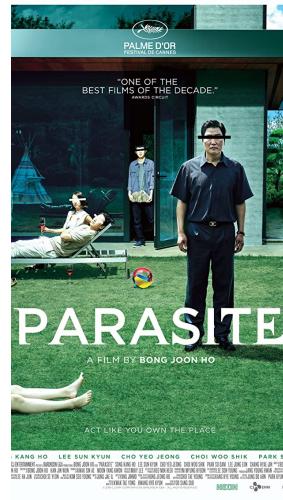
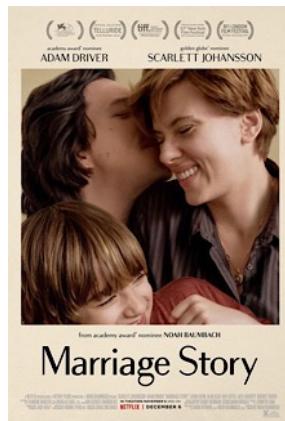
	5				4
		2	3		
	4				
				4	
1					
	2		3		
	5	1			3

call this the utility (or user-item) matrix  $Y$

# How to solve for the missing ratings?

- 1)Matrix factorization
- 2)Nearest neighbor prediction

movies (from our own planet)



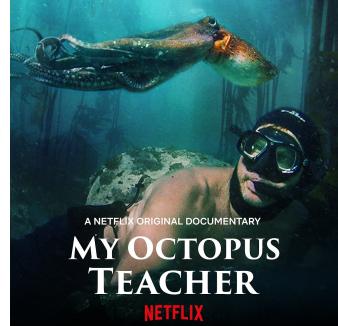
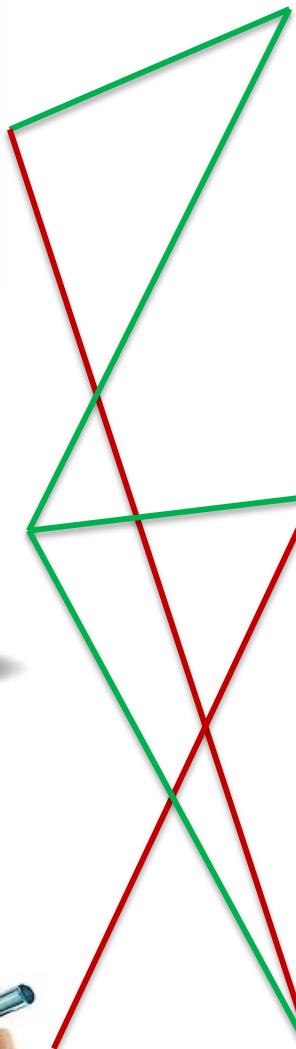
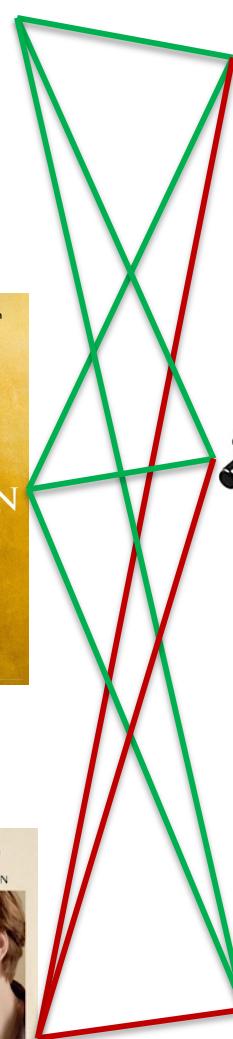
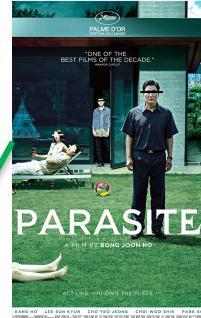
users (in a galaxy far far away...)



↙NN  
prediction



target user



# Approach 2: Nearest Neighbor Prediction

**Key idea:**

Suppose user  $a$  has not rated movie  $i$

To predict the rating

- compute *similarity* between user  $a$  and all other users in the system
- find the  $k$  ‘nearest neighbors’ of user  $a$  who have rated movie  $i$
- compute a prediction based on these users’ ratings of  $i$

# Approach 2: Nearest Neighbor Prediction

**Key idea:**

Suppose user  $a$   has not rated movie  $i$



To predict the rating

- compute *similarity* between user  $a$  and all other users in the system
- find the  $k$  ‘nearest neighbors’ of user  $a$  who have rated movie  $i$
- compute a prediction based on these users’ ratings of  $i$



# Correlation as a measure of similarity

Users:  $a, b$

$R(a, b)$ : set of movies rated by both users a and b

$$\tilde{Y}_{a:b} = \frac{1}{|R(a, b)|} \sum_{j \in R(a, b)} Y_{a_j}$$

average rating of user a for  
movies rated by both users

$\tilde{Y}_{b:a}$

$$sim(a, b) = \frac{\sum_{j \in R(a, b)} (Y_{a_j} - \tilde{Y}_{a:b})(Y_{b_j} - \tilde{Y}_{b:a})}{\sqrt{\sum_{j \in R(a, b)} (Y_{a_j} - \tilde{Y}_{a:b})^2 \sum_{j \in R(a, b)} (Y_{b_j} - \tilde{Y}_{b:a})^2}}$$

↑  
↑

how much the ratings vary  
individually i.e., std dev of the  
ratings of a and b

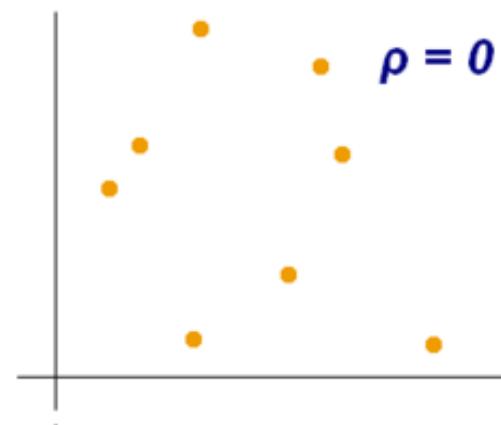
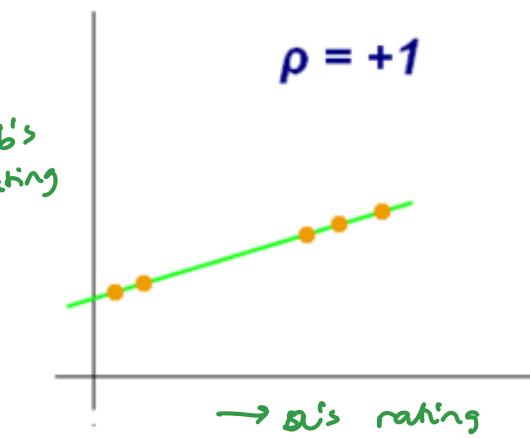
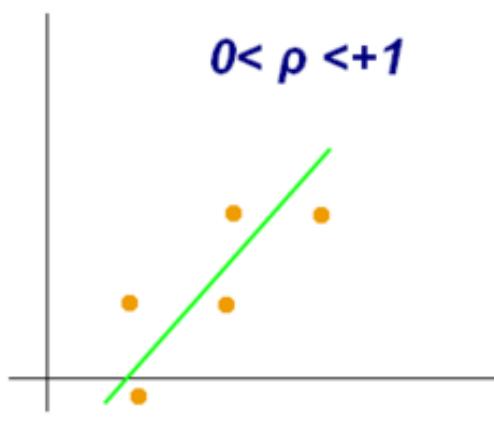
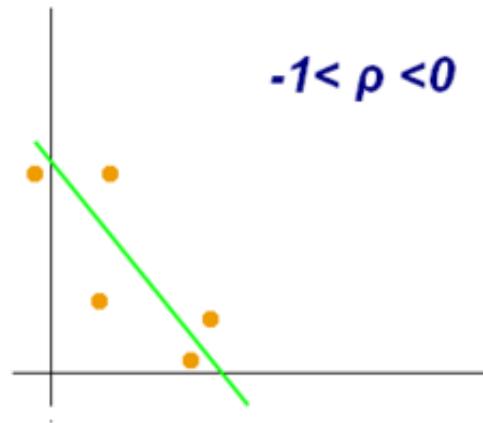
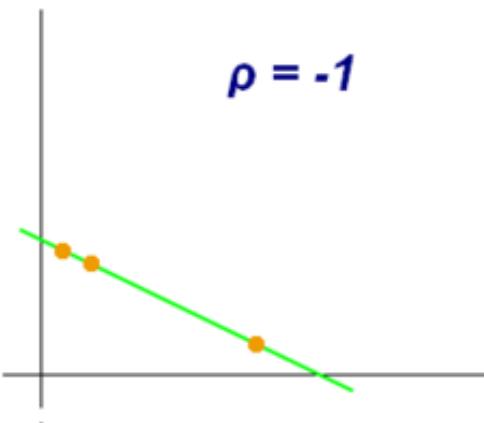
Sample Pearson's correlation coefficient

# Correlation as a measure of similarity

$$sim(a, b) = \frac{\sum_{j \in R(a,b)} (Y_{a_j} - \tilde{Y}_{a:b})(Y_{b_j} - \tilde{Y}_{b:a})}{\sqrt{\sum_{j \in R(a,b)} (Y_{a_j} - \tilde{Y}_{a:b})^2 \sum_{j \in R(a,b)} (Y_{b_j} - \tilde{Y}_{b:a})^2}}$$

- measure of linear relationship between two variables
- ranges from [-1, 1]
  - -1, +1 → perfect linear relationship
  - 0 → no linear correlation
  - absolute value is a measure of degree of correlation

# Correlation (intuition)



# Example: similarity

a	5	?			4
		2	3		
	4				
				4	
1					
	2		3		
b	5	1			3

$$sim(a, b) = \frac{\sum_{j \in R(a,b)} (Y_{a_j} - \tilde{Y}_{a:b})(Y_{b_j} - \tilde{Y}_{b:a})}{\sqrt{\sum_{j \in R(a,b)} (Y_{a_j} - \tilde{Y}_{a:b})^2 \sum_{j \in R(a,b)} (Y_{b_j} - \tilde{Y}_{b:a})^2}}$$

# Example: similarity

a	5	?			4
b	5	1			3

$$sim(a, b) = \frac{\sum_{j \in R(a,b)} (Y_{a_j} - \tilde{Y}_{a:b})(Y_{b_j} - \tilde{Y}_{b:a})}{\sqrt{\sum_{j \in R(a,b)} (Y_{a_j} - \tilde{Y}_{a:b})^2 \sum_{j \in R(a,b)} (Y_{b_j} - \tilde{Y}_{b:a})^2}}$$

<https://forms.gle/ffiBvNbPjHF8ghi77>



$$|R(a,b)| = |\{2,6\}| = 2$$

$$\tilde{y}_{a:b} = \frac{5+4}{2} = 4.5$$

$$\tilde{y}_{b:a} = 4$$

$$\text{sim}(a,b) = \frac{(y_{a2} - \tilde{y}_{a:b})(y_{b2} - \tilde{y}_{b:a}) + (y_{a6} - \tilde{y}_{a:b})(y_{b6} - \tilde{y}_{b:a})}{\sqrt{(y_{a2} - \tilde{y}_{a:b})^2 + (y_{b2} - \tilde{y}_{b:a})^2} \sqrt{(y_{a6} - \tilde{y}_{a:b})^2 + (y_{b6} - \tilde{y}_{b:a})^2}}$$

$$y_{a2} - \tilde{y}_{a:b} = 5 - 4.5 = 0.5$$

$$y_{b2} - \tilde{y}_{b:a} = 5 - 4 = 1$$

$$y_{a6} - \tilde{y}_{a:b} = 4 - 4.5 = -0.5$$

$$y_{b6} - \tilde{y}_{b:a} = 3 - 4 = -1$$

$$\therefore \text{sim}(a,b) = \frac{(0.5)(1) + (-0.5)(-1)}{\sqrt{[0.5^2 + (-0.5)^2][1^2 + (-1)^2]}} = \frac{0.5 + 0.5}{\sqrt{(0.5)(2)}} = 1$$

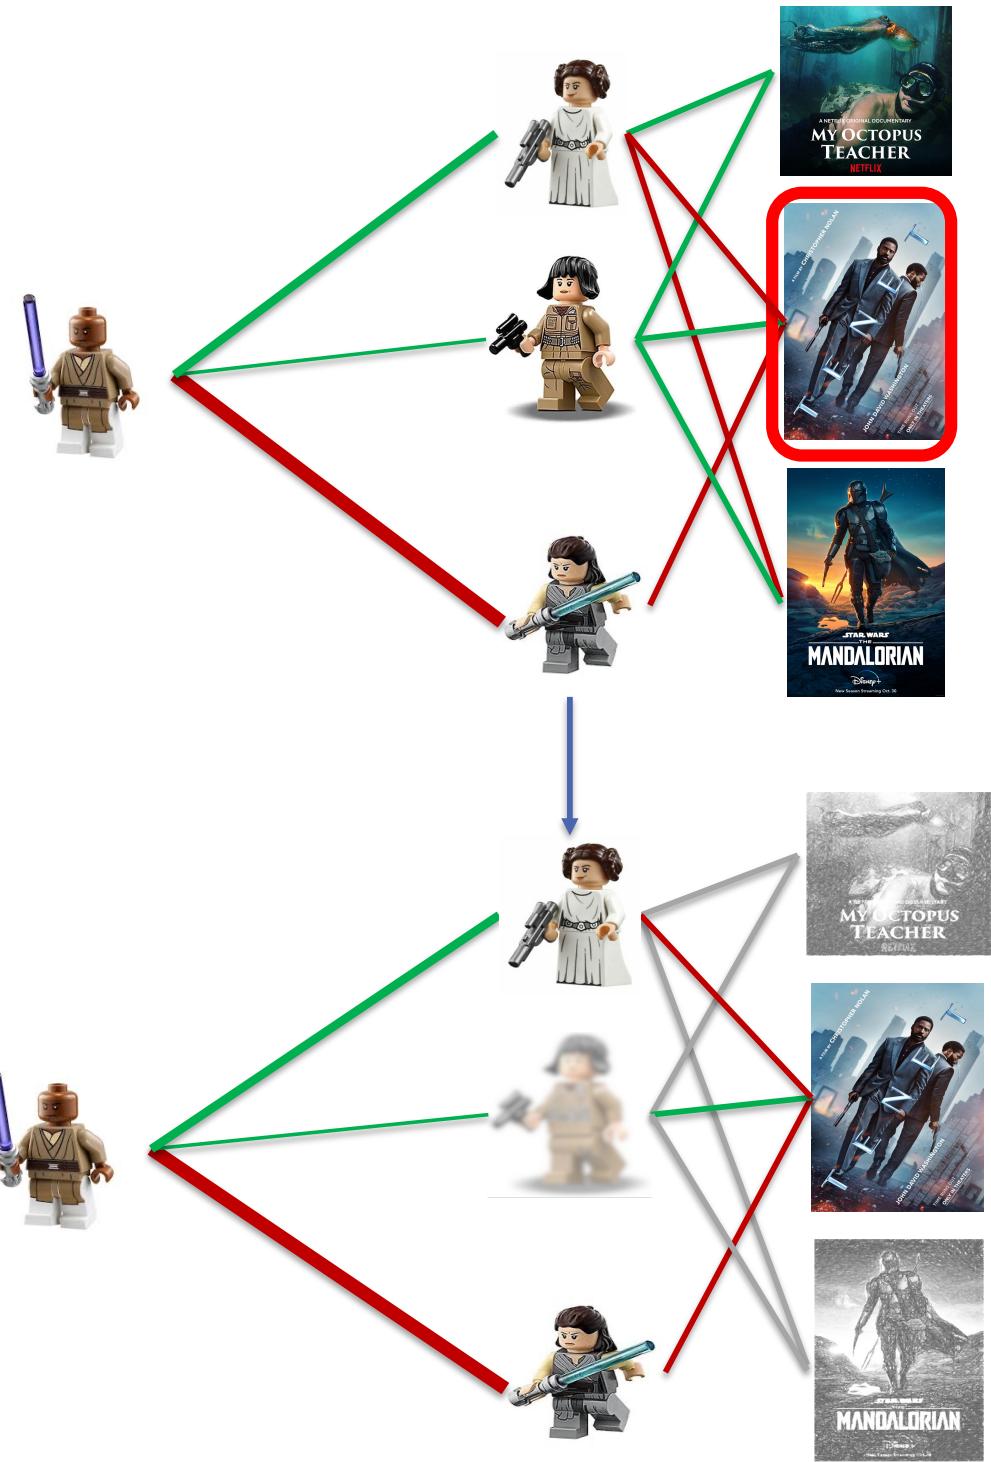
# Approach 2: Nearest Neighbor Prediction

**Key idea:**

Suppose user  $a$  has not rated movie  $i$

To predict the rating

- compute *similarity* between user  $a$  and all other users in the system
- find the  $k$  ‘nearest neighbors’ of user  $a$  who have rated movie  $i$
- compute a prediction based on these users’ ratings of  $i$



ordered in decreasing order of *absolute* value

$\text{sim}(a, \text{user } 1)$

$\text{sim}(a, \text{user } 2)$

$\vdots$

$\text{sim}(a, \text{user } k)$

*non increasing*

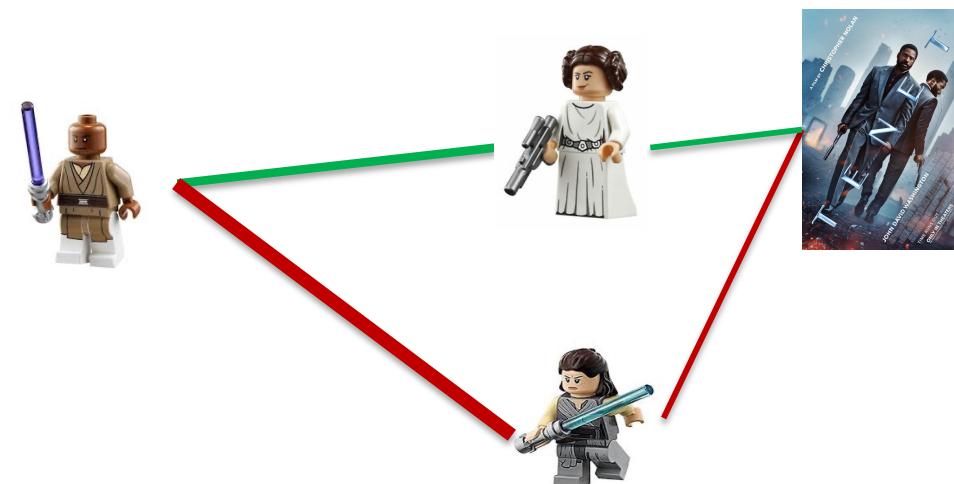
# Approach 2: Nearest Neighbor Prediction

**Key idea:**

Suppose user  $a$  has not rated movie  $i$

To predict the rating

- compute *similarity* between user  $a$  and all other users in the system
- find the  $k$  ‘nearest neighbors’ of user  $a$  who have rated movie  $i$
- compute a prediction based on these users’ ratings of  $i$



# Prediction

$kNN(a, i)$  k “nearest neighbors”  
i.e., the k most similar users to a,  
who have rated movie i

$$\hat{Y}_{a_i} = \bar{Y}_a + \frac{1}{\sum_{b \in kNN(a, i)} |sim(a, b)|} \sum_{b \in kNN(a, i)} sim(a, b)(Y_{b_i} - \bar{Y}_b)$$

# How to choose $k$ ?

- Some heuristics:
  - all users are considered as neighbors
  - Why could limiting neighborhood size result in more accurate predictions?
    - neighbors with low correlation tend to introduce noise
  - Variable  $k$ 
    - varying neighborhoods are selected for each item based on a similarity threshold
  - offline analysis: values in the 20–50 range are a reasonable starting point in many domains.

# How to solve for the missing ratings?

- 1)Matrix factorization
- 2)Nearest neighbor prediction