

# **EECS 482: Introduction to Operating Systems**

## **Lecture 2: Processes**

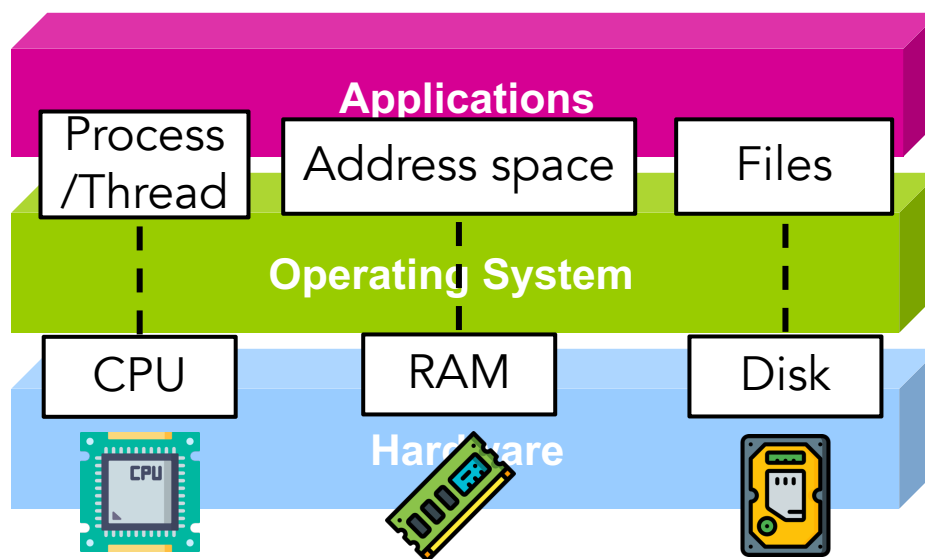
Prof. Ryan Huang

# Administrivia

498-02\* office hours not start until project 2

# OS abstractions

Recap: OS provides abstractions to hide details of hardware from applications

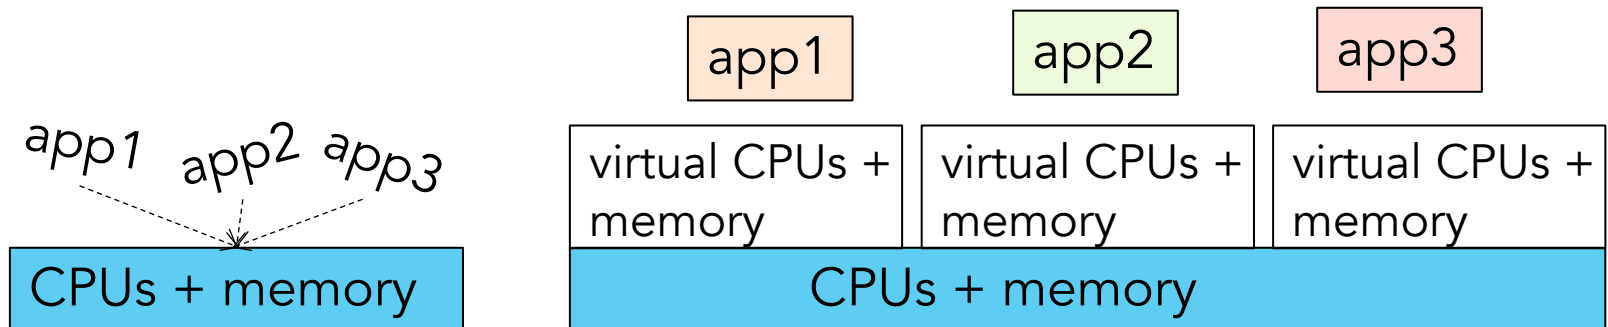


# Process abstraction

Process is the OS abstraction for CPU (execution)

- Sometimes also called a **job** or a **task**

Decompose mix of activities running on a computer into several **independent** tasks



# A process' view

Each process has its own view of the machine

- Its own virtual memory
  - 0xc000 means different thing in P1 & P2
- Its own virtual CPU
- Its own open files

Simplifies programming model

- gcc does not need to care that firefox is running

# What is a process?

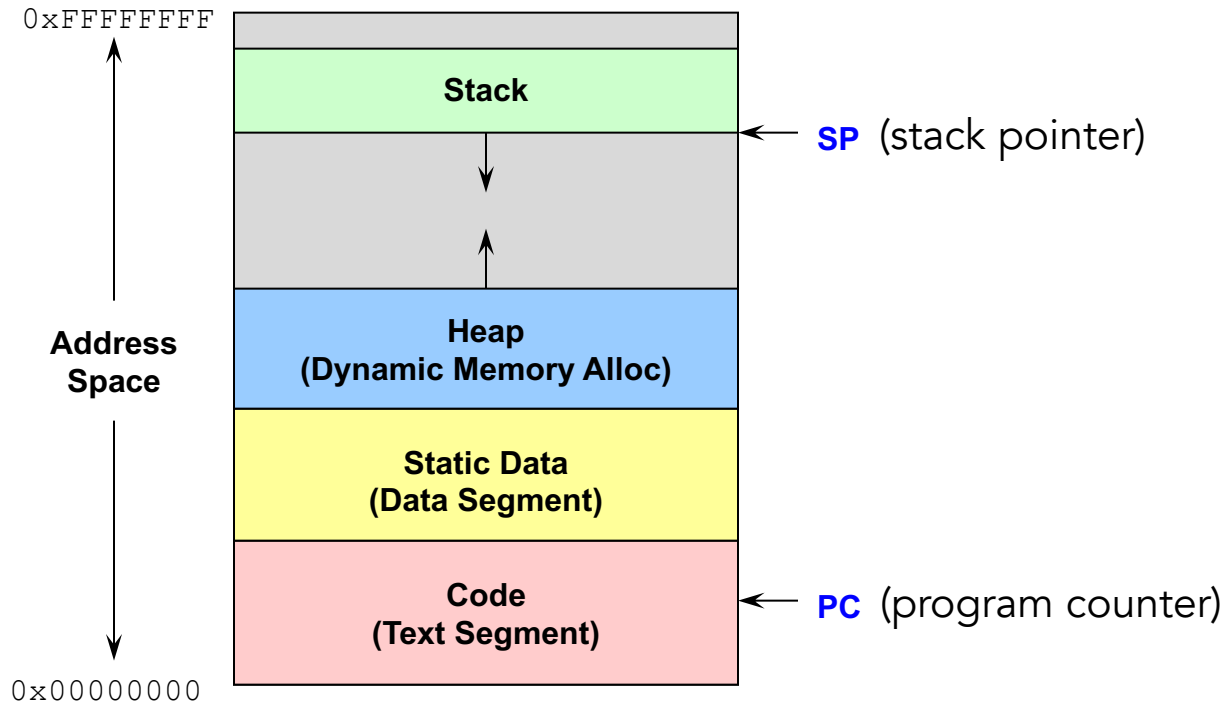
A process is a **program in execution**

- Programs are **static** entities with the potential for execution

It contains all state needed for program execution

- **An address space** (memory)
  - The **code** for the executing program
  - The **data** for the executing program
  - An **execution stack** encapsulating the state of procedure calls
- The program counter (PC) indicating the next instruction
- A set of general-purpose registers with current values
- A set of OS resources
  - Open files, network connections, etc.

# Process address space



All addresses a process can *possibly* use

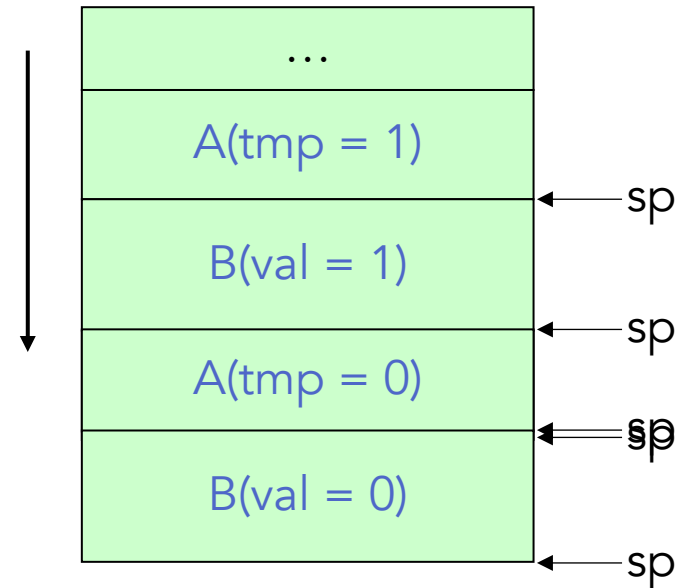
- Note: the addresses are *virtual*

Provides isolation

- Cannot be accessed by another process without permission

# Review: the call stack

```
→ A(int tmp) {  
→     B(tmp);  
→ }  
  
→ B(int val) {  
→     C(val, val + 2);  
→     A(val - 1);  
→ }  
  
→ C(int foo, int bar) {  
→     int v = bar - foo;  
→ }
```



*A stack frame is created for each function invocation and is destroyed after the function call finishes*



# Thread

Modern OSes separate the concepts of processes and **threads**

Thread defines a sequence of **execution stream**

- PC, SP, registers

Process defines the **address space** and **general process attributes** (process ID, open files, etc.)

A thread is bound to a *single* process

But a process can have **multiple threads**

- Sometimes they interact
- Sometimes they work independently

# Per-thread state vs. shared state

## Type of state

- Registers (e.g., `eax`, `ebx`, ...)
- Code
  - + program counter (e.g., `eip`)
- Stack
  - + stack pointer (e.g., `esp`)
- Data segment (heap + static variables)

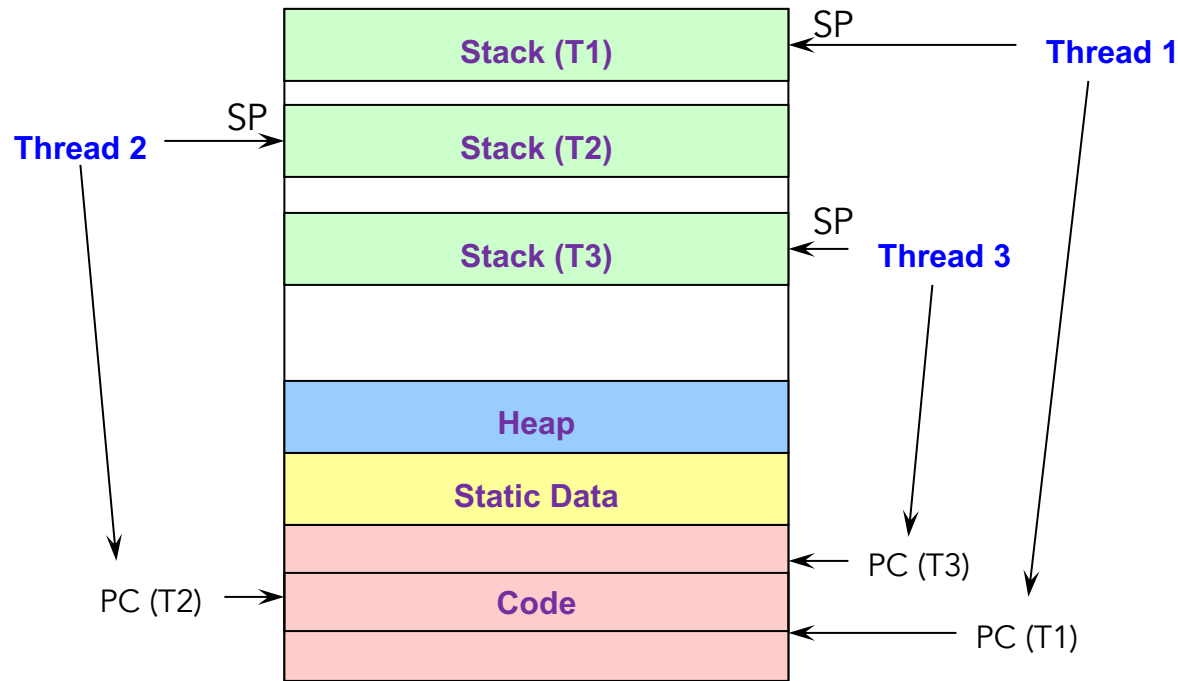
**Per-thread state: which state is needed by (or unique to) each thread as it executes instructions?**

- Stack, Registers, PC, SP

**Other state is shared:**

- Code and data segment

# Threads in a process



*If thread 1 does `p=malloc(..)`; and passes `p` to thread 2, can thread 2 use `p`?*

# Upcoming topics

## Thread: unit of concurrency

- How multiple threads cooperate to accomplish a task
- How multiple threads can share limited number of CPUs

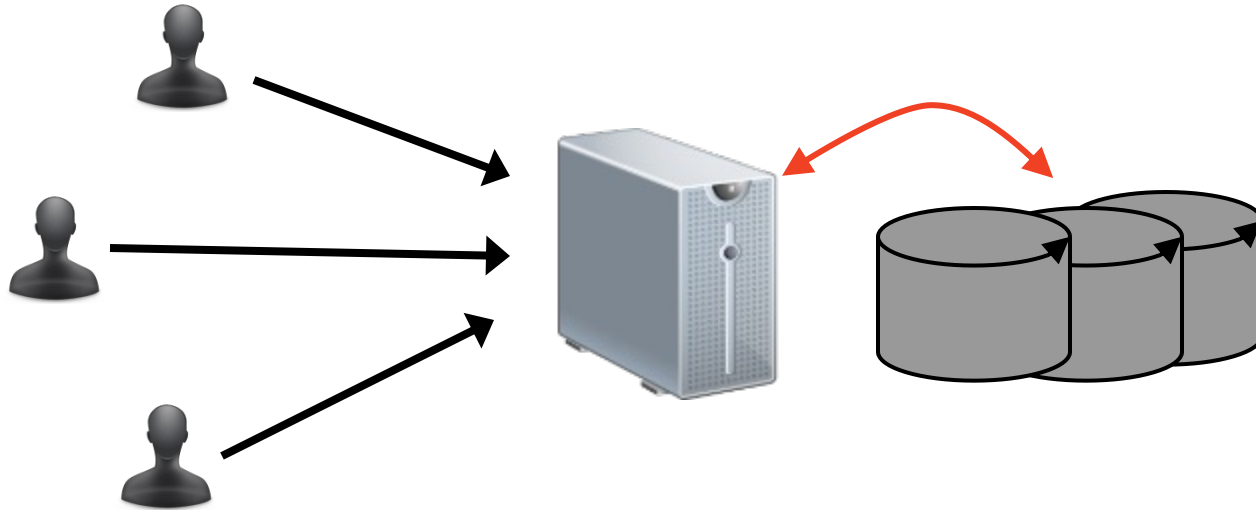
## Address space: unit of state partitioning

- How do address spaces share single physical memory?
  - Efficiently
  - Flexibly
  - Safely

# Why do we need threads?

## Example: a web server application

- Receives multiple simultaneous requests
- Reads web pages from disk to satisfy each request



# Option 1: Handle one request at a time

Request 1 arrives

Server receives request 1

Server starts disk I/O 1a and waits for it to complete..

Request 2 arrives

Server must finish request 1 before handling request 2



**Easy to program, but slow**

- Can't overlap disk requests with computation
- Can't overlap either with network sends and receives

# Option 2: Asynchronous I/O (event-driven)

## Issue I/Os, but don't wait for them to complete

Request 1 arrives

Server receives request 1

Server starts disk I/O 1a (but does not wait for it to complete)

Request 2 arrives

Server receives request 2

Server starts disk I/O 2a

Request 3 arrives

Disk I/O 1a finishes

Continue to process request 1



## Fast, but **difficult to program**

- Lots of extra state to remember at each point
- What requests are being served, what stage they're in
- What disk I/Os are outstanding; which requests they belong to

# Option 3: Multi-threaded web server

## One thread per request

- Thread issues disk (or network) I/O, then waits for it to finish
- Though thread is blocked on I/O, other threads can run
- Where is the state of each request stored?

### Thread 1

Request 1 arrives  
Receive request 1  
Start disk I/O 1a

Disk I/O 1a finishes  
Continue processing  
request 1

### Thread 2

Request 2 arrives  
Receive request 2  
Start disk I/O 2a

### Thread 3

Request 3 arrives  
Receive request 3



# Benefits of threads

Allow multiple things to happen in parallel

Lightweight to create compared to processes

- Mostly allocating a stack

Thread manager takes care of CPU sharing

- One thread can issue blocking I/O, while other threads can still progress
- Private state for each thread

Applications get a “simpler” programming model

- The illusion of a dedicated CPU per thread

# Downsides of threads

## Performance degradation under high concurrency

- Costs of context switching, lock contention, cache misses, scheduling, etc.
- Event-driven model can achieve higher performance
  - Used in high-concurrency software in practice, e.g., Nginx

## Synchronization is hard!

- More on this later..

## Long-standing debate

- *Why threads are a bad idea*, Ousterhout, 1995
- *Why events are a bad idea*, von Behren et al., 2003

# When are threads useful?

Multiple things happening at once

Usually some slow resource

- Network, disk, user, ...

**Examples:**

- Network server (e.g., web server)
- Controlling a physical system (e.g., airplane controller)
- Window system
- Parallel programming

# Ideal scenario

Split computation into threads

Threads run **independently** of each other

- Divide and conquer works best if divided parts are independent

How practical is thread independence?

# Completely independent threads?

## Example 1: Microsoft Word

- One thread formats document
- Another thread spell checks document

## Example 2: Desktop computer

- One thread compiles EECS 482 project
- Another thread runs Minecraft

## Dependence among threads is often inevitable

- i.e., threads are often *cooperating*

Two types of sharing: **application** resource or **hardware** resource

# Non-determinism

## Problem:

- Speed of each processor is unpredictable
- Ordering of events across threads is **non-deterministic**

Thread A - - - - - >  
Thread B - - - - - >  
Thread C ----->

## Consequences:

- **Many** possible global ordering of events
  - Also known as *thread interleaving*
- Different orderings may produce different results

# Non-deterministic ordering → Non-deterministic results

Printing example

Thread 1

print "ABC"

Thread 2

print "123"

What's being shared between these threads?

Possible outputs:

Impossible outputs:

Ordering within thread is sequential

Many ways to merge per-thread order into a global order

# Non-deterministic ordering → Non-deterministic results

Assignment example (x is initially 0)

Thread A

x = 1

Thread B

x = -1

- What's being shared between these threads?
- Possible results:
- Impossible results:



# Non-deterministic ordering → Non-deterministic results

Arithmetic example (y is initially 10)

Thread A

$$x = y + 1$$

Thread B

$$y = y * 2$$

- What's being shared between these threads?
- Possible results:

# Non-deterministic ordering → Non-deterministic results

Assignment example (x is initially 0)

Thread A

x = 1

x = -100  
x += 101

Thread B

x = -1

x = -100  
x += 99

- Possible results:
- Impossible results:

# Atomic operations

Before we can reason at all about cooperating threads, we must know that some operation is **atomic**

- Indivisible, i.e., happens in its entirety or not at all
- No events from other threads can occur in between start and end of an atomic operation

Thread 1: `print "ABC"`      Thread 2: `print "123"`

- What if each print statement were atomic?
- What if printing a single character were **not** atomic?

Need an atomic operation to build a bigger atomic operation

Set of atomic instructions varies by ISA

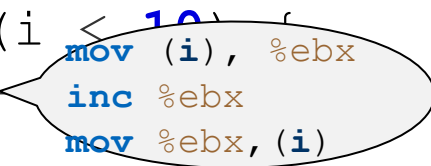
- E.g., memory load is atomic; memory store is atomic; many other instructions are not atomic (e.g., double-precision floating point)

# Example

## Thread A

i=0

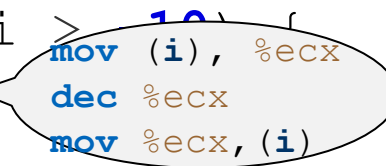
```
while (i < 10) {  
    mov (i), %ebx  
    i++  
    inc %ebx  
    mov %ebx, (i)  
}  
print "A finished"
```



## Thread B

i=0

```
while (i > 10) {  
    mov (i), %ecx  
    i--  
    dec %ecx  
    mov %ecx, (i)  
}  
print "B finished"
```



Which thread will finish first?

Is the winner guaranteed to print first?

Is it guaranteed that someone will win?

Say both threads run at **exactly** the same speed, and start close together

- Is it guaranteed that both threads will loop forever?

# Do you really need to worry about such far-fetched scenarios?

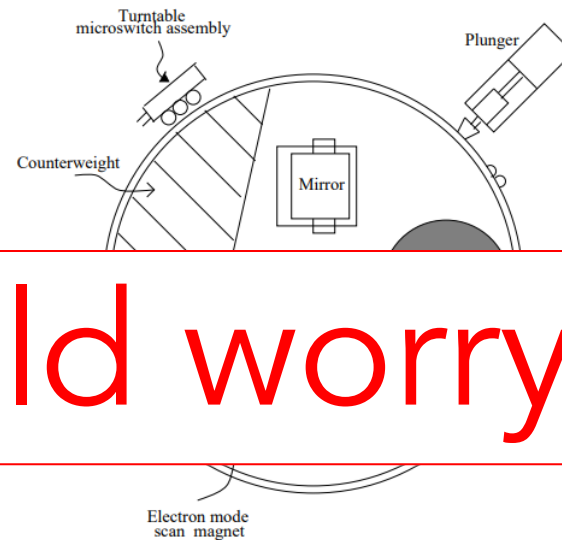


Figure 1: Upper turntable assembly.

## Yes, you should worry!

Northeast blackout of 2003

Therac-25 radiation therapy  
[Leveson95]

Non-deterministic interleaving makes debugging challenging  
**Heisenbug**: a bug that occurs non-deterministically

# Writing correct concurrent programs

Consider and control all possible interleavings of events from multiple threads

- All possible interleavings of atomic operations must result in correct output
- Non-atomic operations to a shared resource must not occur concurrently

Controlling how events from different threads can interleave is called **synchronization**

# Synchronization

## Goals

- **Eliminate interleavings**: all possible interleavings must produce a correct result
  - Trivial solution?
- **Preserve interleavings**: constrain thread interleavings as little as possible

Concurrency only matters for events that access some shared resource

A correct concurrent program should work no matter how fast the CPUs execute the threads