

EECS 482: Introduction to Operating Systems

Lecture 5: Monitors

Prof. Ryan Huang

Administration

Project 1 due this Wednesday

Reminder: honor code

- You may use online resources including ChatGPT, Copilot **only** for general programming help
- **You may NOT use them for help specific to EECS 482 projects**

Monitors

Two mechanisms: one for each type of synchronization

- Locks for mutual exclusion
- Condition variables for ordering constraints

A monitor = a lock + the condition variables associated with that lock

Mesa vs. Hoare monitors

Mesa

```
while (!condition)
    wait(cond_var);
```

- When waiter is woken, the condition might have changed
 - another thread got the lock first
- So it must re-check the condition it was waiting for

Hoare

```
if (!condition)
    wait(cond_var);
```

- Special priority to the woken-up waiter
- Signal thread **immediately** gives up lock
 - reacquires lock after waiter unlocks

We (and most/all Oses) use Mesa monitors

- Waiter is solely responsible for ensuring condition is met
- Always use **while(...)** {wait}

Monitor programming: overall design

Step 1: Think about each thread **independently**

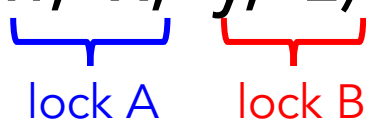
- Each thread should try to make as much forward progress as possible (be greedy)
- A thread should wait only when it is unable to proceed

Step 2: List the shared data needed for the problem (e.g., w, x, y, z)

Monitor programming: mutex

Step 2: List the shared data needed for the problem

(e.g., w, x, y, z)



lock A lock B

Step 3: Assign a lock per group of related, shared data

- lock...unlock around accesses to shared data
- Coarse-grained versus fine-grained locking

Monitor programming: ordering

Step 4: List before-after conditions and assign condition variables to each condition

When is a thread unable to make progress?

Use one condition variable for each before-after condition

- Condition variable belongs to the `lock` that protects the `shared data` that is used to evaluate the condition

Waiting thread uses `while (condition) {cv.wait}`

Signaling thread calls `signal` or `broadcast` when it changes state that may allow another thread to proceed

Typical monitor code

```
lock

// wait if needed

while (condition) {

    wait

}

do stuff

signal or broadcast about the stuff you did

unlock
```


Producer-consumer (bounded buffer)

Producers fill a shared buffer; consumers empty it

Need to synchronize actions of producers and consumers



Why use a shared buffer?

- Allows producers and consumers to operate somewhat independently (more concurrency)

Used in many situations

- Unix pipes (`grep "keyword" foo.txt | wc -l`)
- Project 1

Coke machine with monitors

Problem definition

- Coke machine can hold at most MAX cokes
- Delivery person (producer) adds one coke to machine
- Consumer buys one coke

Step 1: Think about threads independently

- Delivery threads add cokes, consumer threads remove cokes

Step 2: Identify shared state

- State of coke machine: `int numCokes`

Step 3: Assign locks

- One lock (`mutex cokeLock`) to protect all shared data

Coke machine with monitors

Step 4: List before-after conditions and assign a condition variable to each condition

- Consumers wait while there are no cokes (`numCokes == 0`)
 - `cv waitingConsumers`
- Delivery persons wait while the cokes are full (`numCokes == MAX`)
 - `cv waitingProducers`

Coke machine with monitors

Consumer ()

```
// take coke out of machine
```

Producer ()

```
// add coke to machine
```

Coke machine with monitors

```
int numCokes
mutex cokeLock
cv waitingConsumers, waitingProducers
```

Consumer()

```
cokeLock.lock()

while (numCokes == 0) {
    waitingConsumers.wait(cokeLock)
}

// take coke out of machine
numCokes--

waitingProducers.signal()

cokeLock.unlock()
```

Producer()

```
cokeLock.lock()

while (numCokes == MAX) {
    waitingProducers.wait(cokeLock)
}

// add coke to machine
numCokes++

waitingConsumers.signal()

cokeLock.unlock()
```

Reducing number of signals?

```
int numCokes
mutex cokeLock
cv waitingConsumers, waitingProducers
```

Consumer()

```
cokeLock.lock()

while (numCokes == 0) {
    waitingConsumers.wait(cokeLock)
}

// take coke out of machine
numCokes--

waitingProducers.signal()

cokeLock.unlock()
```

Producer()

```
cokeLock.lock()

while (numCokes == MAX) {
    waitingProducers.wait(cokeLock)
}

// add coke to machine
numCokes++

if (numCokes == 1)
    waitingConsumers.signal()

cokeLock.unlock()
```

Would this still work?

Reducing number of signals

`numCokes = 0`

C1 and C2 waiting on `waitingConsumers`

P1 increments `numCokes` to 1 and **signals**

P2 increments `numCokes` to 2, but **does not signal**

Only one of C1 and C2 wakes up!

Reducing condition variables?

```
int numCokes
mutex cokeLock
cv waiting
```

Consumer()

```
cokeLock.lock()

while (numCokes == 0) {
    waiting.wait(cokeLock)
}

// take coke out of machine
numCokes--

waiting.signal()

cokeLock.unlock()
```

Producer()

```
cokeLock.lock()

while (numCokes == MAX) {
    waiting.wait(cokeLock)
}

// add coke to machine
numCokes++

waiting.signal()

cokeLock.unlock()
```

Would this still work?

Reducing condition variables

Say $\text{MAX} = 1$, and $\text{numCokes} = 0$

C1 and C2 wait

P1 increments numCokes to 1 and signals

- Wakes up C1

P2 waits because $\text{numCokes} = \text{MAX}$

C1 decrements numCokes to 0 and signals

- May wake up C2!

Why is this bad?

Reader-writer locks

Example: multiple threads would like to read or write Wolverine Access course catalog

Threads need to acquire lock to read or write shared data

Student:

`lock()`

`browse course catalog`

`unlock()`

Administrator:

`lock()`

`change course catalog`

`unlock()`

Any concern with this code?

How to safely allow more concurrency?

How to safely allow more concurrency?

Idea: threads disclose whether they intend to read or write the shared data through separate read/write lock methods

Student:

```
lock() readerLock()  
browse course catalog  
unlock() readerUnlock()
```


Administrator:

```
lock() writerLock()  
change course catalog  
unlock() writerUnlock()
```

Allow **multiple concurrent readers**, if no threads are writing data


Allow a **single writer**, if no other threads are reading or writing

Another level of abstraction



Applications

Even higher-level
synchronization operations
Concurrent programs
(readerLock, readerUnlock,
writerLock, writerUnlock)



Operating System

Higher-level synchronization
operations
(lock, monitor, semaphore)



Hardware

Atomic operations
(load/store, interrupt enable/
disable, test&set)

Reader-writer lock with monitors

Step 1: Think about threads independently

Step 2: Identify shared state (to implement reader-writer lock methods)

- `numReaders, numWriters`

Step 3: Assign locks to shared state:

- `mutex rwLock`

Step 4: List the before-after conditions and assign a condition variable for each condition:

- readerLock waits `while numWriters>0`
 - `cv waitingReaders`
- writerLock waits `while numReaders>0 || numWriters>0`
 - `cv waitingWriters`

Reader-writer lock with monitors

```
int numReaders, numWriters
mutex rwLock
cv waitingReaders, waitingWriters
```

```
readerLock () {
    rwLock.lock()
    while (numWriters>0) {
        waitingReaders.wait(rwlock)
    }
    numReaders++
    rwLock.unlock()
}
```

```
writerLock() {
    rwLock.lock()
    while (numReaders>0 || numWriters>0) {
        waitingWriters.wait(rwLock)
    }
    numWriters++
    rwLock.unlock()
}
```

```
readerUnlock () {
    rwLock.lock()
    numReaders--
    waitingWriters.signal()
    rwLock.unlock()
}
```

```
writerUnlock () {
    rwLock.lock()
    numWriters--
    waitingReaders.broadcast()
    waitingWriters.signal()
    rwLock.unlock()
}
```

Reader-writer lock with monitors

What will happen if a writer finishes and there are several waiting readers and writers?

How long will a writer wait?

How to give priority to a waiting writer?