

EECS 482: Introduction to Operating Systems

Lecture 6: Building Threads

Prof. Ryan Huang

Administration

Project 1 due today

- Soft deadline
 - You may continue to submit to AG after the deadline but the extra points won't receive the hand-grading multiplier

Project 2 will be posted today

- Implement a thread library

For ungrouped students, we're in the process of matching

Threads can interact in two ways

Multiple threads share data to cooperate on task

- Coordinate via synchronization operations, e.g., locks, condition variables, semaphores

We've been assuming that each thread has its own processor (of unpredictable speed). But actually...

Multiple threads can share a single processor

What is a thread?

A thread is a sequence of **executing** instructions

So what's a **non-running** thread?

- A non-running thread is a **paused** execution
- Can pausing a thread break a correct concurrent program?

How to **pause** a thread and **resume** it later?

- Called "**context switch**"

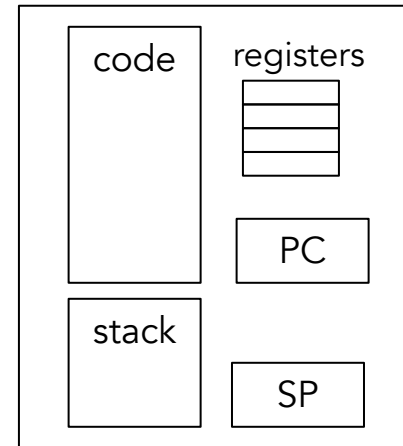
Per-thread state

What state is needed by a thread as it runs?

- Code
 - + program counter (e.g., `eip`)
- Stack
 - + stack pointer (e.g., `esp`)
- Registers (e.g., `eax`, `ebx`, ...)

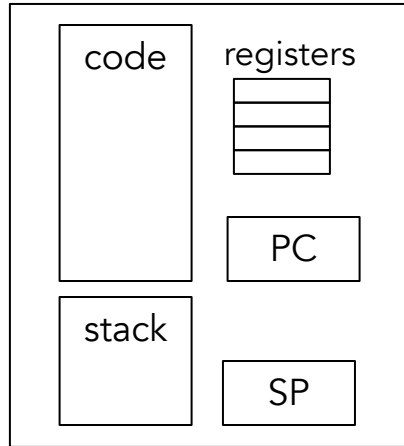
CPU

Thread context

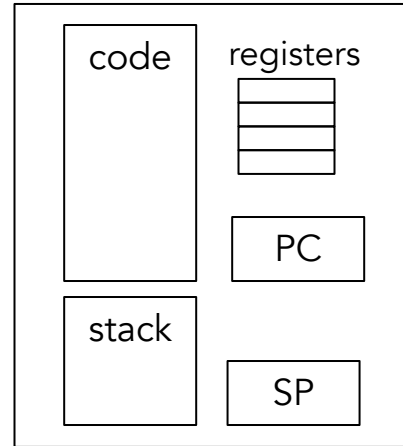


Sharing the CPU

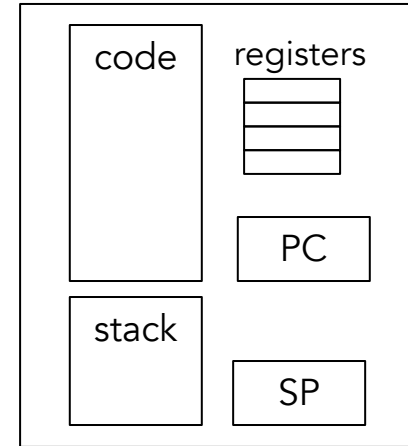
Thread 1 context



Thread 2 context



Thread 3 context

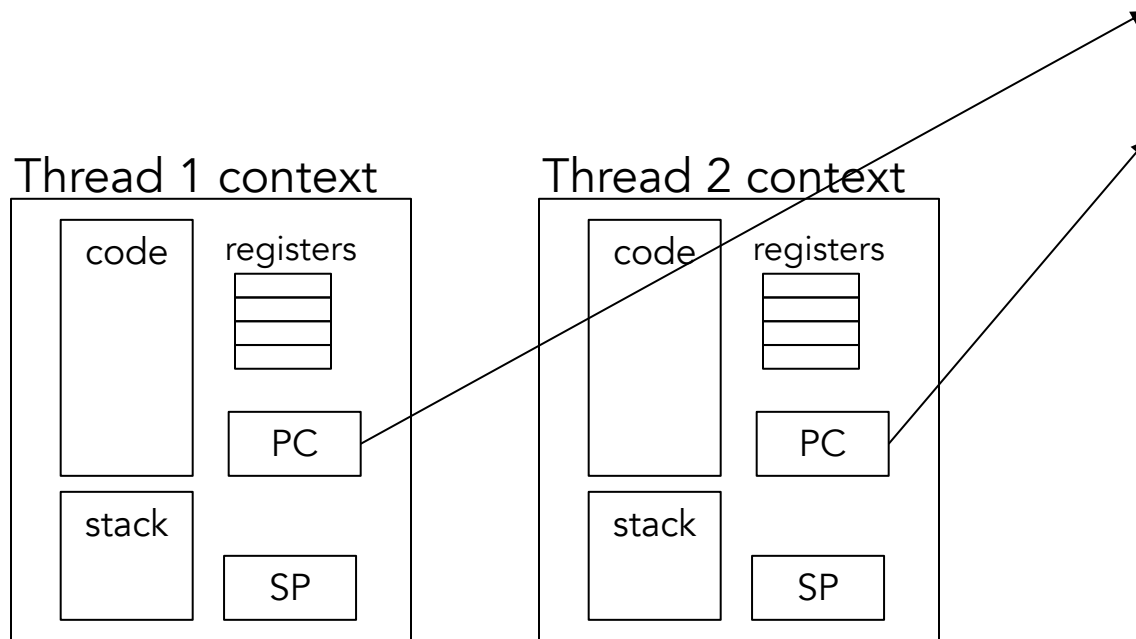


CPU

Optimizations

How to avoid copying code back and forth?

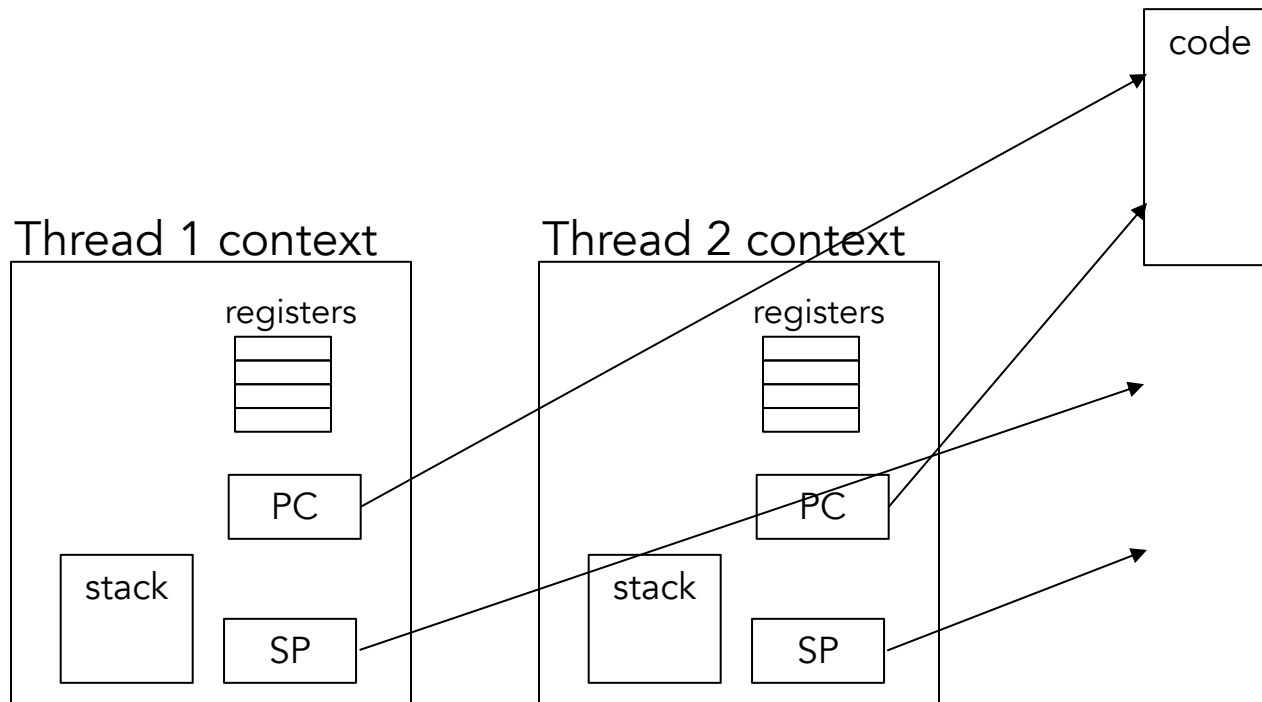
- Store separately from thread context (PC is still part of context)
- Threads use same code



Optimizations

How to avoid copying the stack?

- Store separately from thread context (SP is still part of context)



Thread control block (TCB)

OS data structure to store thread info

- In memory
- Record execution context (SP, PC, registers)
 - when the thread is not running (paused)
- Why is TCB needed?
 - What if we do not record the execution context?

Allocated when a new thread is created

- Typically a thread stack is allocated together
 - e.g., in one 4KB page
- Destroyed when a thread is finished

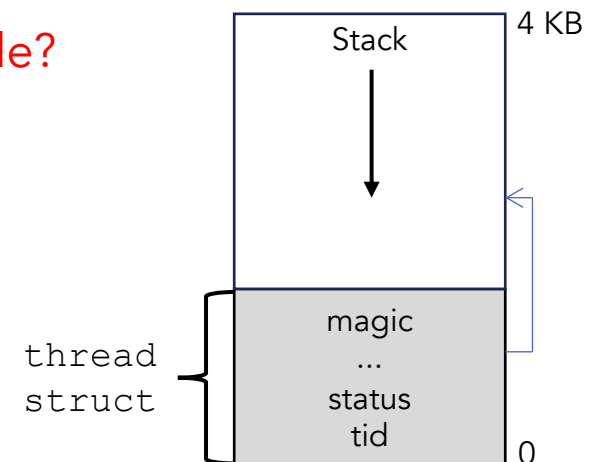
TCB example

An example TCB struct

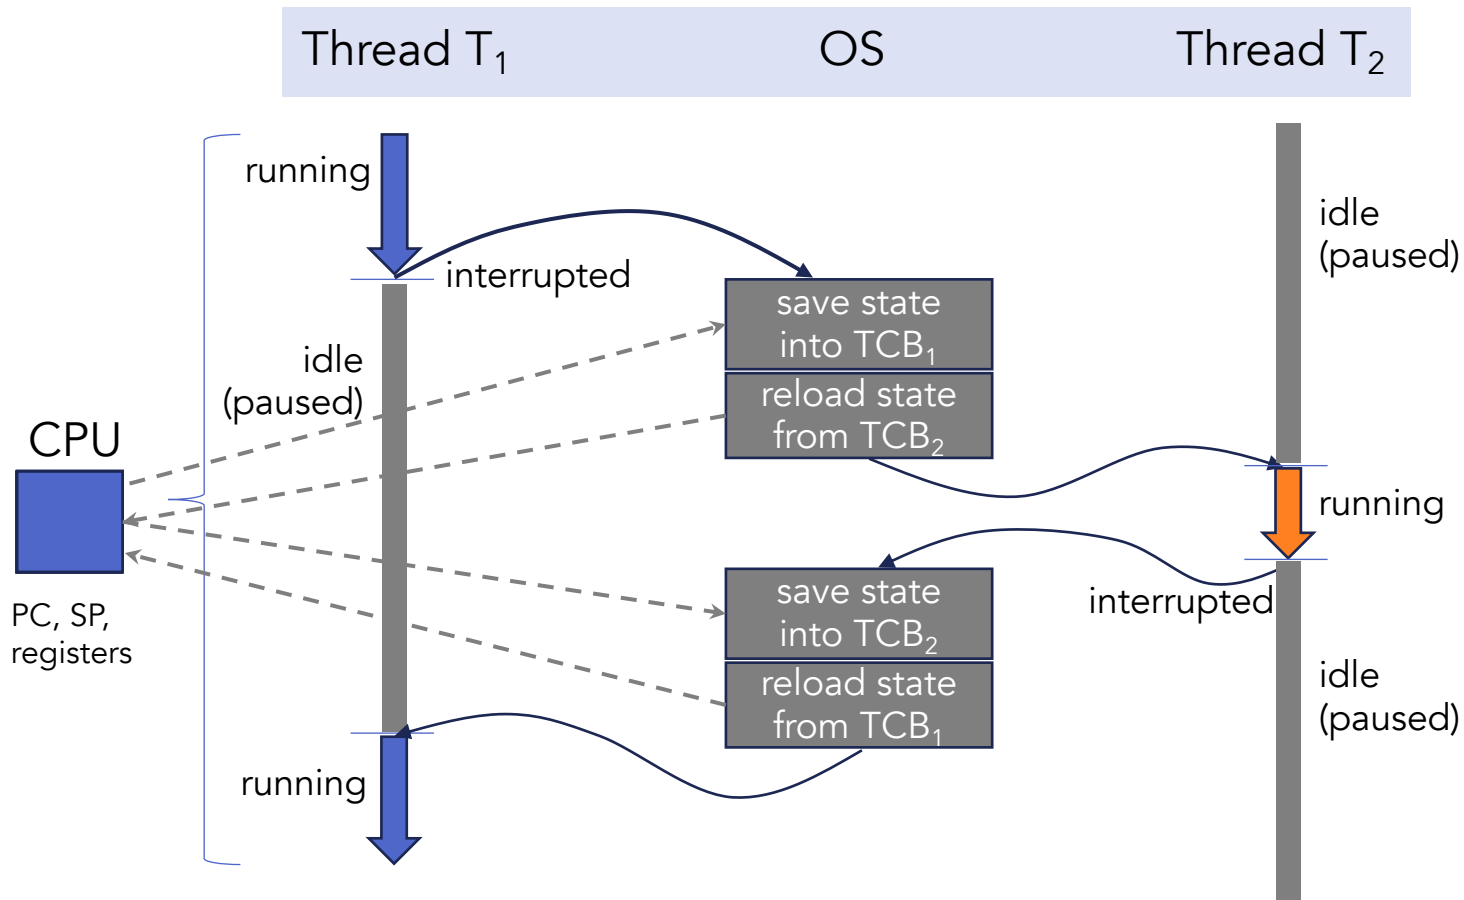
```
struct thread {  
    tid_t tid; /* Thread identifier. */  
    enum thread_status status; /* Thread state. */  
    char name[16]; /* Name (for debugging purposes). */  
    uint8_t *sp; /* Saved stack pointer. */  
    int priority; /* Priority. */  
    struct list_elem allelem; /* List element for all threads list. */  
    struct list_elem elem; /* List element. */  
    unsigned magic; /* Detects stack overflow. */  
};
```

Where are the PC and registers stored in this example?

They are pushed onto the stack



Sharing the CPU



Two perspectives when threads are sharing a processor

Thread view:

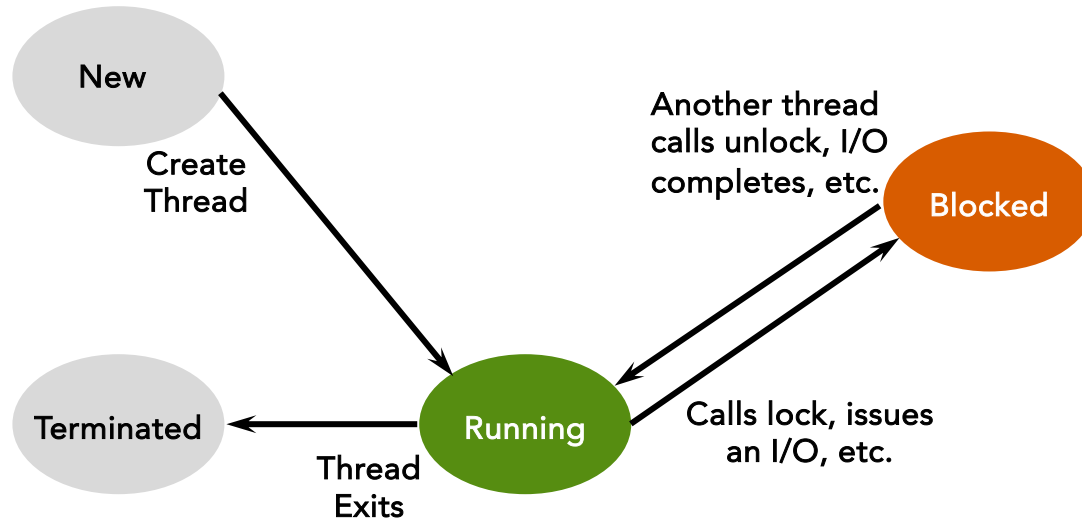
- Running → Paused → Running

CPU view:

- Thread 1 → Thread 2 → Thread 1

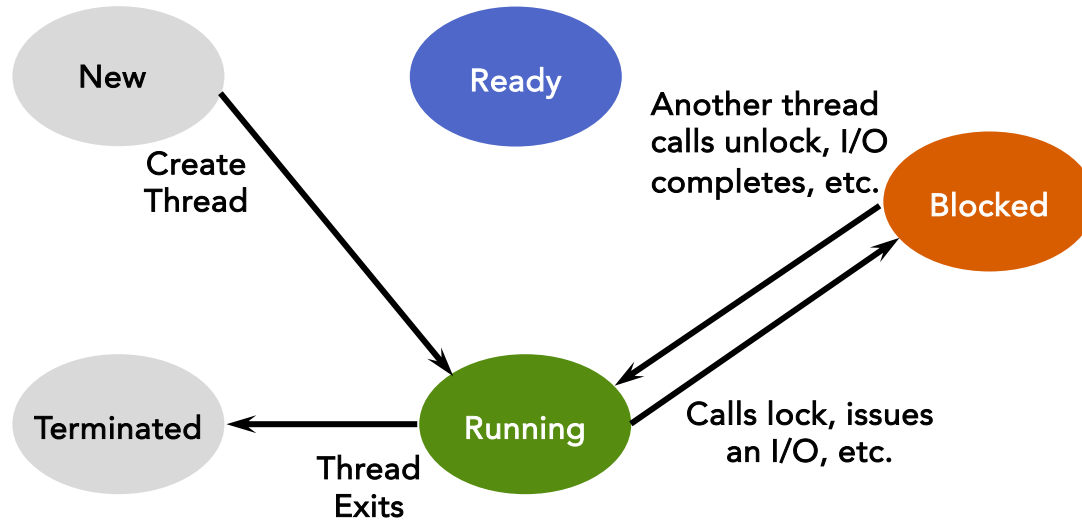
First, let's take the thread view

States of a thread



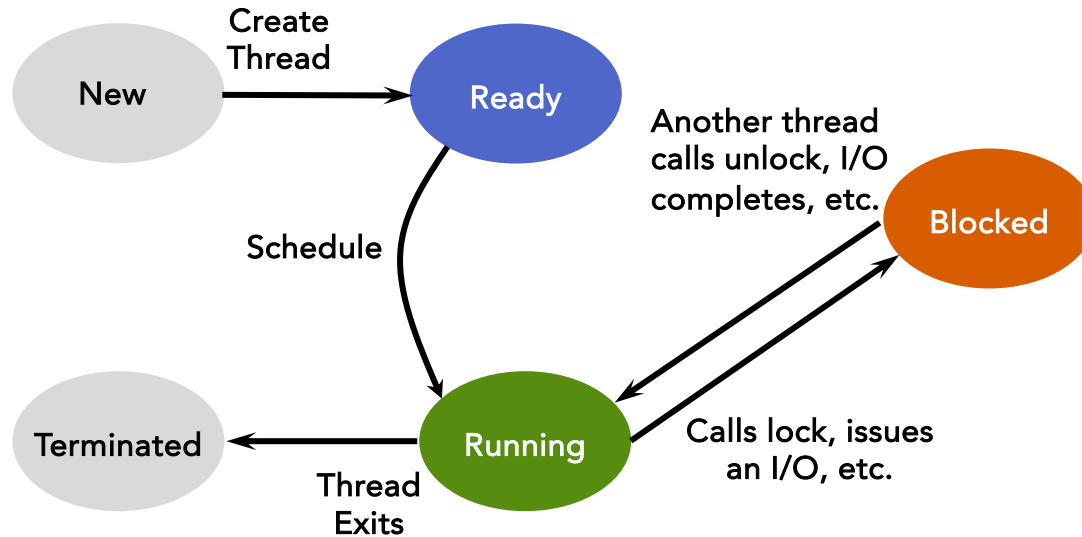
What if there are more threads than CPUs?

States of a thread



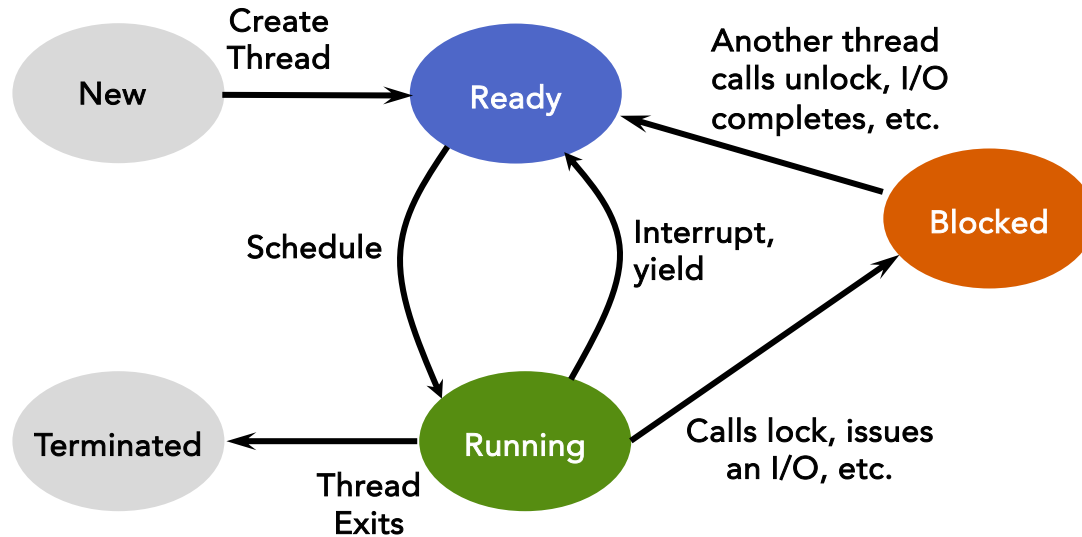
What if there are more threads than CPUs?

States of a thread



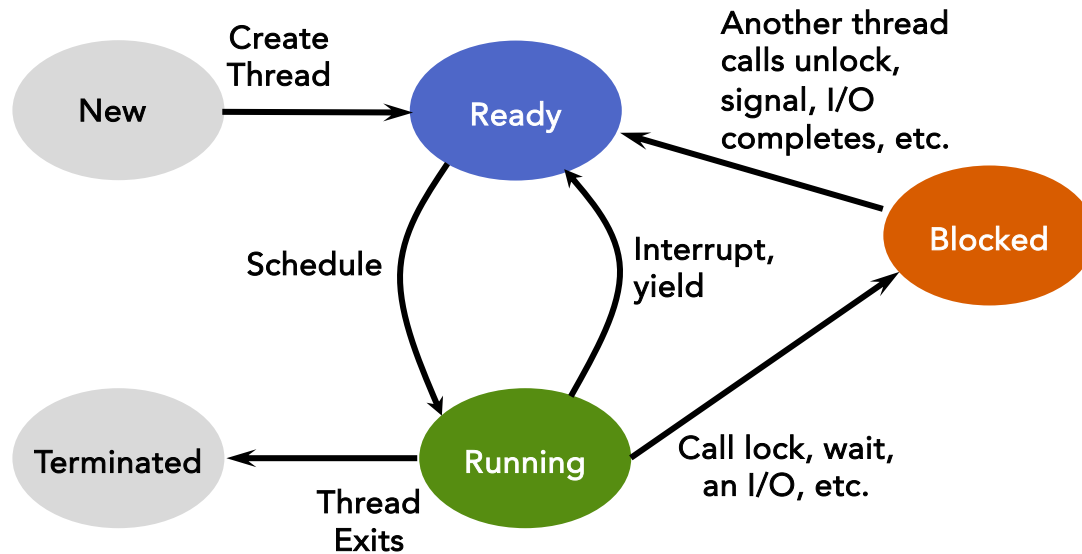
What if there are more threads than CPUs?

States of a thread



What if there are more threads than CPUs?

States of a thread



Why no transition from Ready to Blocked?

Why no transition from Blocked to Running?

State queues

How does the OS keep track of threads?

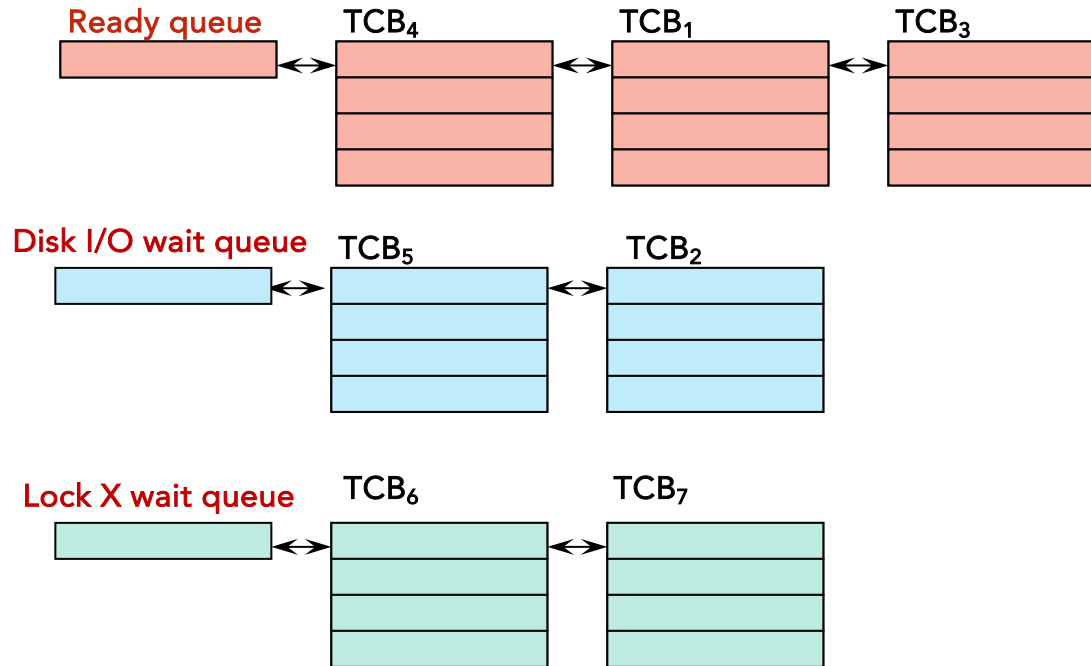
Simple approach: all threads list

- How to find out threads in the ready state?
 - Iterate through the list
- **Problem: slow!**

Improvement: partition list based on states

- OS maintains a collection of queues
 - Typically one queue for each state: ready, waiting, etc.
- A TCB is queued on a state queue according to its current state
 - Queue element is typically a pointer to a TCB
- As a thread changes state, its TCB is moved from one queue into another

State queues



- There may be many wait queues, one for each type
- of wait (lock, disk, console, timer, network, etc.)
-

What about a running queue?

How many threads can be in the running state simultaneously?

Two perspectives when threads are sharing a processor

Thread view:

- Running → Paused → Running

CPU view:

- Thread 1 → Thread 2 → Thread 1

Next, let's take the CPU view

Switching threads

- 1. Current thread returns control to OS
2. OS chooses new thread to run
3. OS saves current thread state: CPU to TCB
4. OS loads context of next thread: TCB to CPU
5. OS runs next thread

How does thread return control to OS?

Internal events

- Initiated by the current thread itself

Examples

- Thread calls `lock()`, `wait()`, `down()`, etc.
- Thread requests OS to do some work (e.g., I/O)
- Thread voluntarily gives up CPU with `yield()`

Thread yielding

Threads voluntarily give up the CPU with `yield`

Ping Thread

```
while (true) {  
    printf("ping\n");  
    yield();  
}
```

Pong Thread

```
while (true) {  
    printf("pong\n");  
    yield();  
}
```

What is the output of running these two threads?

What does it mean for `yield` to return?

Cooperative thread scheduling

How does thread return control to OS?

Internal events

Are internal events sufficient?

External events

- Initiated by something outside the current thread
- **Interrupts**: a hardware event that transfers control from thread to OS interrupt handler

Interrupts

An interrupt is an event raised by some hardware component (e.g., an I/O device)

- *Asynchronous*: originating outside the thread
 - not caused by the current executing instruction
- Indicates that some device needs services
- All interrupts are assigned a number, typically [0, 255]

Processor checks for pending interrupt

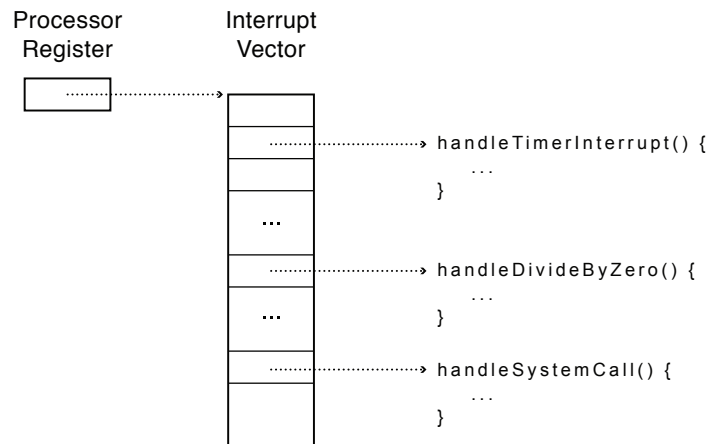
When a pending interrupt is detected

1. CPU stops executing current code
2. CPU calls an OS *interrupt handler*

Interrupt Vector Table (IVT)

A data structure to associate interrupts with handlers

- Each entry is called an **interrupt vector**
 - specifies the address of the interrupt handler
- Programmed by the OS, *i.e.*, software interface for interrupts



CPU invokes the corresponding interrupt handler in the *IVT* based on the pending interrupt number

Example: timer interrupt

Timer is critical for an OS

- Fallback mechanism for OS to reclaim control
- OS could set timer to go off every 10 ms
- When timer expires, it generates an interrupt
 - Guarantees that OS will get control back in ≤ 10 ms

Timer interrupt handler forces the current thread to "call" yield

- Achieves preemptive scheduling
- Causes an involuntary context switch