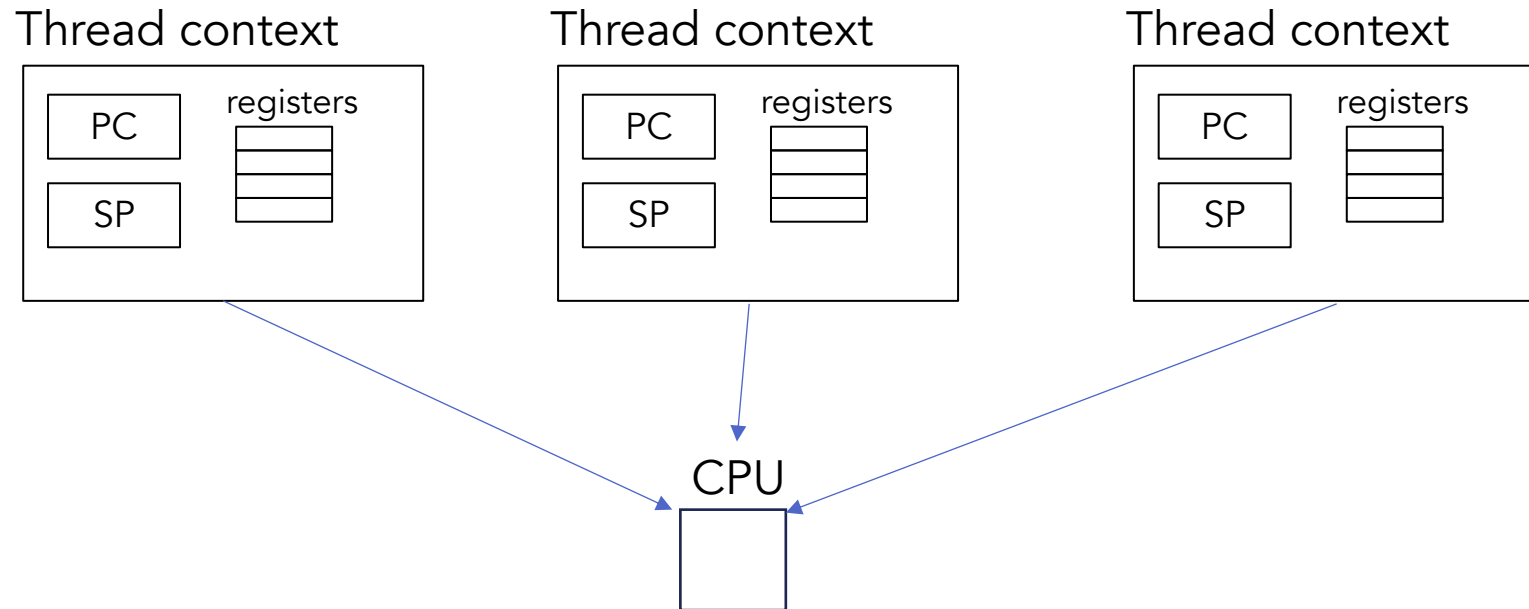


# **EECS 482: Introduction to Operating Systems**

## **Lecture 7: Switching Threads**

Prof. Ryan Huang

# Recap: threads share the CPU



**Each thread has a thread control block (TCB)**

- Record execution context when the thread is not running

# Recap: threads return control to OS

## Internal events

- Thread calls `lock()`, `wait()`, `down()`, etc.
- Thread requests OS to do some work (e.g., I/O)
- Thread voluntarily gives up CPU with `yield()`

## External events

- Initiated by something outside the thread
- Interrupts: a hardware event that transfers control from thread to OS interrupt handler
- Example: timer interrupt

# Switching threads

- 1. Current thread returns control to OS
2. OS chooses new thread to run
3. OS saves current thread state: CPU to TCB
4. OS loads context of next thread: TCB to CPU
5. OS runs next thread

# Choosing next thread to run

1 ready thread: run the ready thread


- What if there's only one thread, and that thread calls yield?

>1 ready thread: need to make a decision

- CPU's scheduling policy
- Lots of options: FIFO, priority, round robin, etc.
- Will discuss in a later lecture

0 ready threads: what should CPU do?

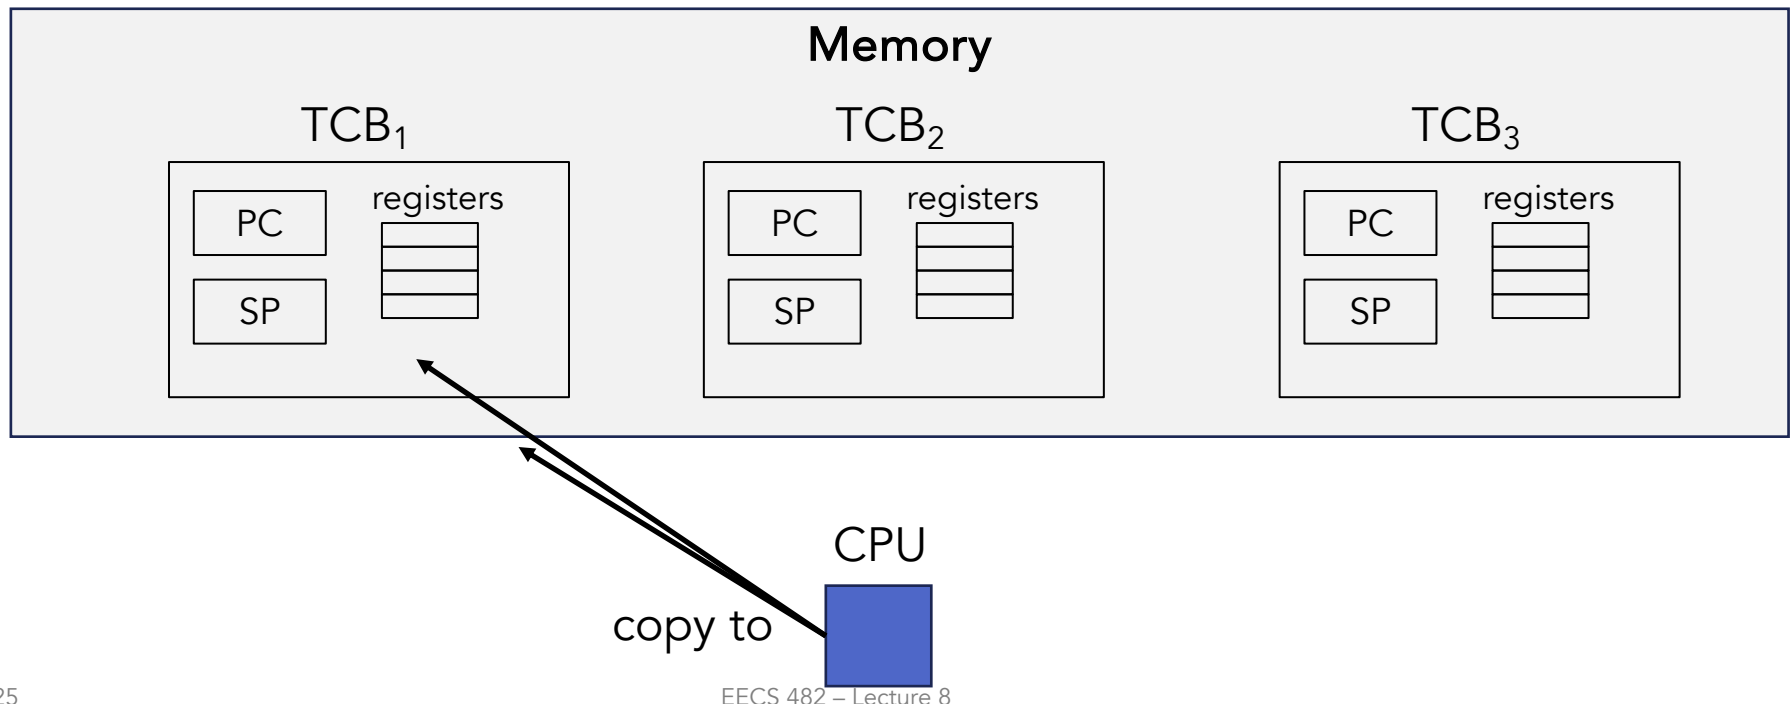
# Switching threads

1. Current thread returns control to OS
-  2. OS chooses new thread to run
3. OS saves current thread state: CPU to TCB
4. OS loads context of next thread: TCB to CPU
5. OS runs next thread

# Save state of the current thread

## Save registers, stack pointer, PC

- Copy their values from CPU to memory
- Where in memory?
- Which instructions (from EECS 370)?



# Save state of the current thread

## Non-negligible cost

- E.g., saving floating point registers is expensive
  - optimization: only save if a thread uses floating point
- May also require switching address space (if the new thread is a different process)

## Very machine dependent

- What machine state needs to be saved and how to save them depends on the ISA




# Save state of the current thread

Tricky to get right!

- Why won't the following code work?

```
100    save PC
101    switch to next thread
```



Context switch (saving + restoring) needs to be implemented in assembly language

- It requires manipulating physical registers
- It works at the level of the *calling convention* (standard for how functions should be implemented)
  - E.g., how arguments are passed, return address on the stack

# Save state of the current thread


Good news: handy library functions available

- Allow user-level context switching
- Use them to implement user-level thread library (Project 2)

**Glibc API:** `getcontext(ucontext_t *ucp)`

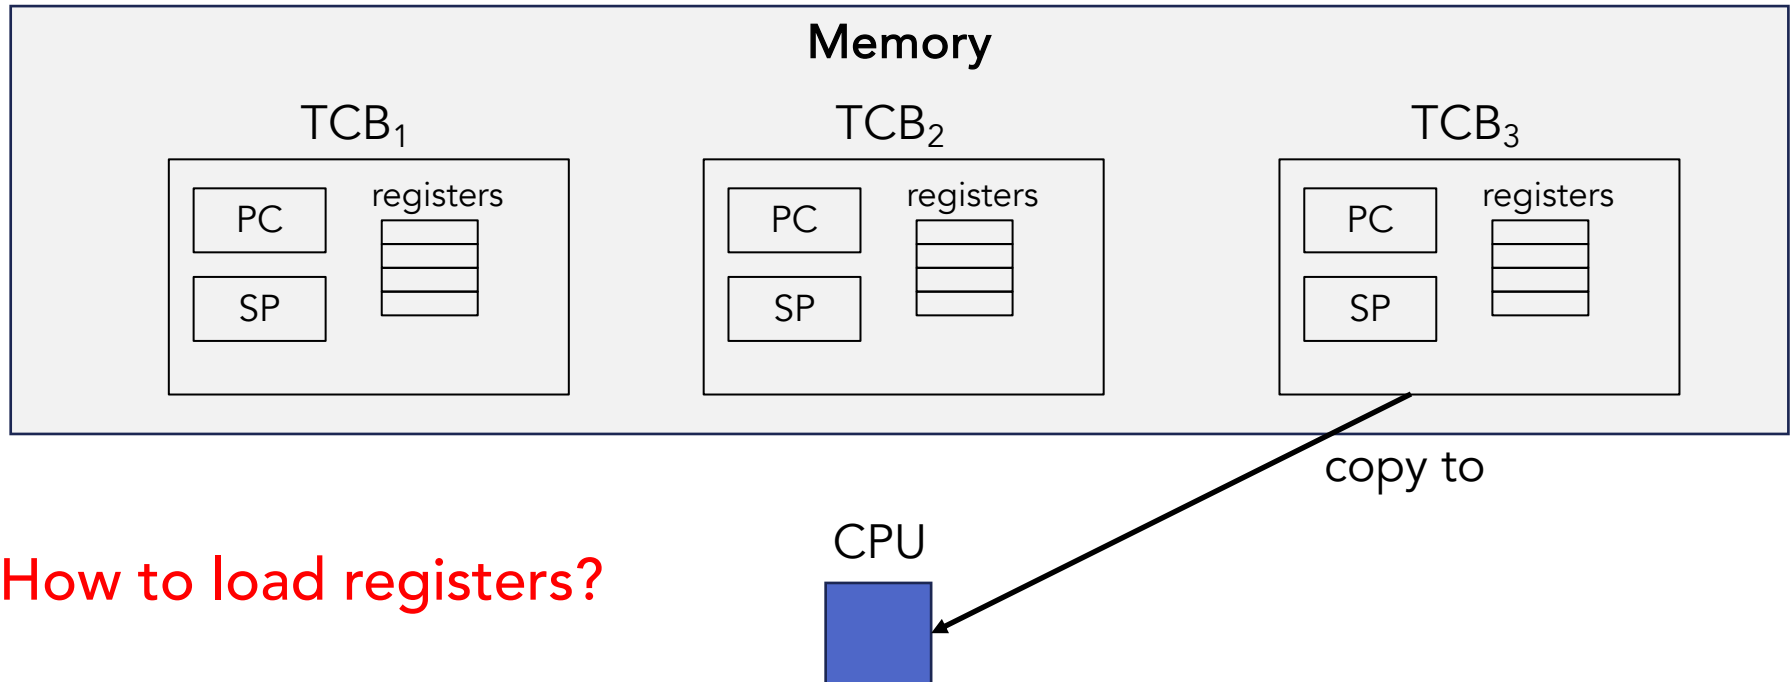
- Its implementation takes care of properly saving the current execution state into the context at `ucp`.

# Switching threads

1. Current thread returns control to OS
2. OS chooses new thread to run
-  3. OS saves current thread state: CPU to TCB
4. OS loads context of next thread: TCB to CPU
5. OS runs next thread

# Load state of the next thread

Copy from memory to CPU



How to load registers?

How to load stack pointer?

How to resume execution?

# Load state of the next thread

**Glibc API:** `setcontext(const ucontext_t *ucp)`

- Copies the context at `ucp` from memory to CPU

How does the thread that is giving up control run again?

# Switching threads

`getcontext()` copies context from CPU to memory

`setcontext()` copies context from memory to CPU

Does the following code work?

## Thread A

```
10 do stuff
11 getcontext // saves state of thread A
12 setcontext to thread B
13 do more stuff
```

## Thread B

```
do stuff
getcontext // saves state of thread B
setcontext to thread A
```

After B's `setcontext()`, where does A resume execution?

Where do we want A to resume execution?

# Switching threads

**Glibc API:** `swapcontext(ucontext_t *oucp, ucontext_t *ucp)`

- Correctly combines `getcontext()` and `setcontext()`

## Thread A

```
10 do stuff
11 swapcontext(A's context, B's context)
12 do more stuff
13 ...
```

## Thread B

```
do stuff
swapcontext(B's context, A's context)
do more stuff
...
```

After B's `swapcontext()`, where does A resume execution?

When will B's `swapcontext()` return?

# Who is carrying out these steps?

E.g., switching from thread A to B?

The CPU?

The OS?

Which thread is carrying out these steps?



# Example of thread switching

## Thread 1

```
print "start thread 1"  
yield()  
print "end thread 1"
```

## Thread 2

```
print "start thread 2"  
yield()  
print "end thread 2"
```

## yield()

```
print "start yield: thread %d"  
switch to next thread (swapcontext)  
print "end yield: thread %d"
```

### Thread 1 output

```
start thread 1  
start yield: thread 1
```

```
end yield: thread 1  
end thread 1
```

### Thread 2 output

```
start thread 2  
start yield: thread 2
```

```
end yield: thread 2  
end thread 2
```

# Example of thread switching

## Thread 1

```
print "start thread 1"  
yield()  
print "end thread 1"
```

## Thread 2

```
print "start thread 2"  
yield()  
print "end thread 2"
```

## yield()

```
print "start yield: thread %d"  
switch to next thread (swapcontext)  
print "end yield: thread %d"
```

### Thread 1 output

start thread 1  
start yield: thread 1

end yield: thread 1  
end thread 1

### Thread 2 output

start thread 2  
start yield: thread 2

end yield: thread 2  
end thread 2



# Example of thread switching

## Thread 1

```
print "start thread 1"  
yield()  
print "end thread 1"
```

## Thread 2

```
print "start thread 2"  
yield()  
print "end thread 2"
```

## yield()

```
print "start yield: thread %d"  
switch to next thread (swapcontext)  
print "end yield: thread %d"
```

## Thread 2 output

```
start thread 1  
start yield: thread 1  
start thread 2  
start yield: thread 2  
end yield: thread 1  
end thread 1  
end yield: thread 2  
end thread 2
```

# Creating a new thread

Create a running thread? Seems challenging

Instead, create **ready** thread

- Key idea: make it look like it was running, put on ready queue
- Then just wait for it to be scheduled!

Construct TCB in the state it would be in if it paused at start of its initial function

# How to create a thread

1. Allocate TCB
2. Allocate stack
3. Initialize context in TCB
  - New thread should look like it was about to call a function
  - Set PC to start of a function
  - Set general-purpose registers and stack to func parameters
  - Linux: `makecontext()` initializes context and stack
4. Add TCB to ready queue

# Administration

You can now do most of Project 2

- `cpu::cpu`
- Thread create, yield, join
- Mutex lock, unlock
- CV wait, signal, broadcast

This is a very hard project *conceptually*


Urge you to start early!

# A brief detour back to writing concurrent programs...

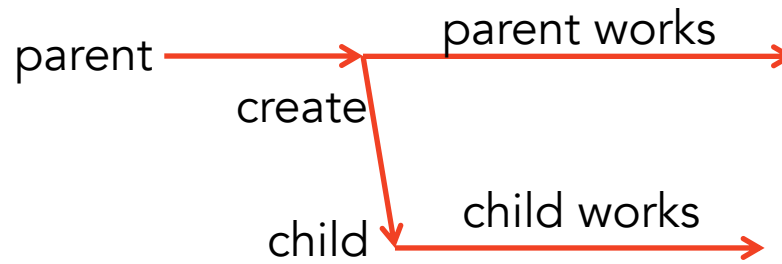
# How to use new thread

Creating a thread is similar to a procedure call

Synchronous procedure call (EECS 280)



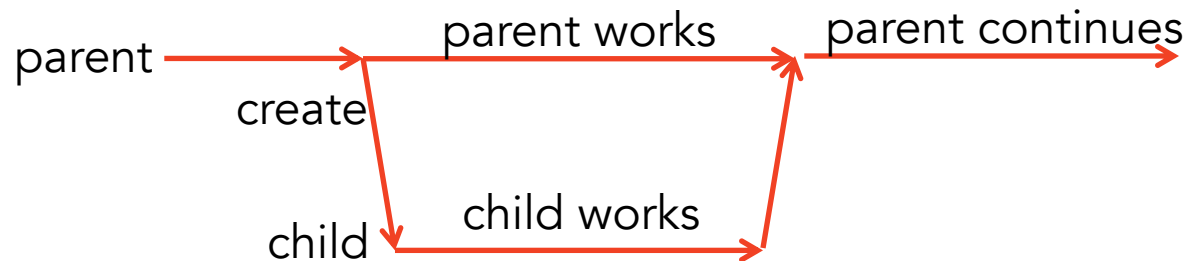
Asynchronous procedure call (EECS 482)





# Synchronizing with child

What if parent wants to work for a while, then wait for child to finish?

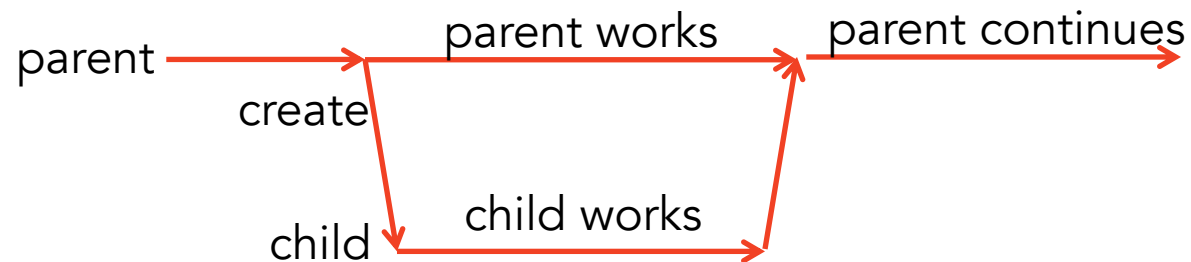


# Synchronizing with child

```
parent()  
    create child thread  
    print "parent works"  
    print "parent continues"
```

```
child()  
    print "child works"
```

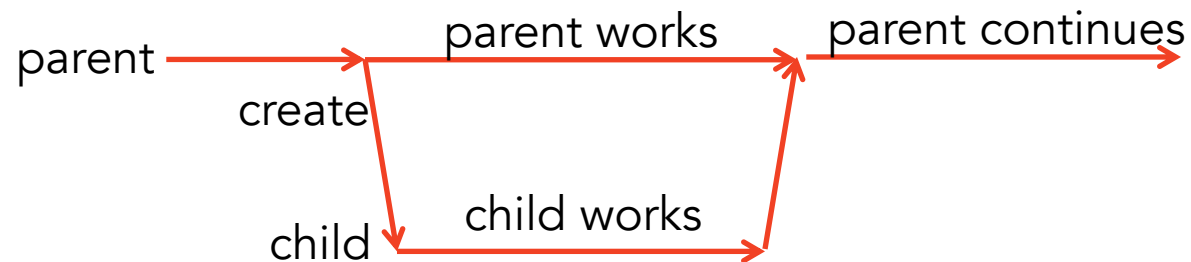
Would this work?



# Synchronizing with child: yield?

```
parent()  
    create child thread  
    print "parent works"  
yield?  
    print "parent continues"
```

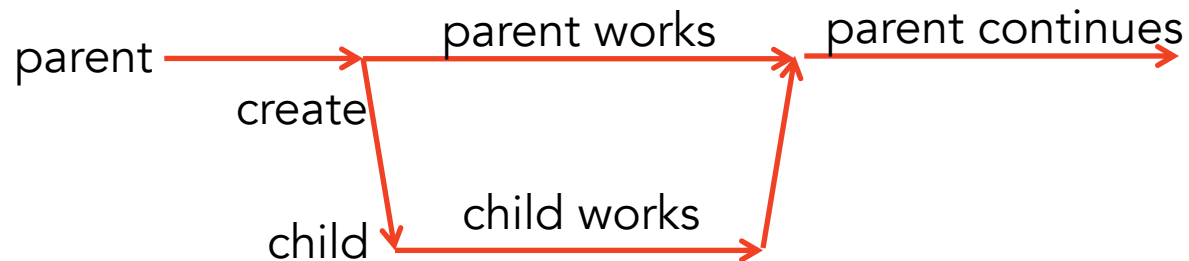
```
child()  
    print "child works"
```



# Synchronizing with child: monitors?

```
parent()  
    create child thread  
    lock  
    print "parent works"  
    wait?  
    print "parent continues"  
    unlock
```

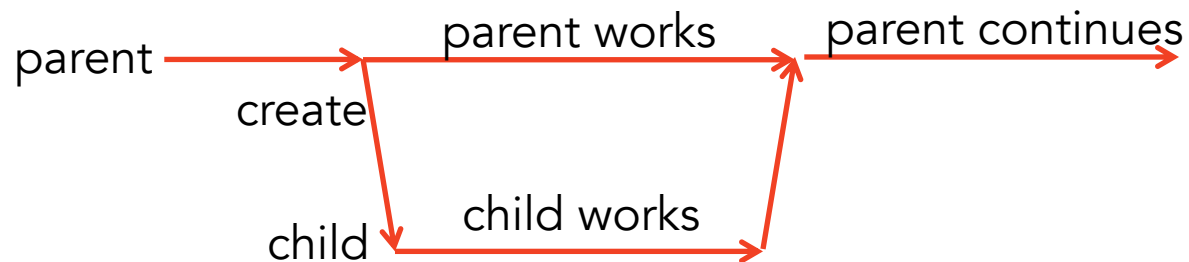
```
child()  
    lock  
    print "child works"  
    signal?  
    unlock
```



# Synchronizing with child: monitors?

```
parent()  
    create child thread  
    lock  
    print "parent works"  
    while (!childDone)  
        wait  
    print "parent continues"  
    unlock
```

```
child()  
    lock  
    print "child works"  
    childDone = 1  
    signal  
    unlock
```



# Synchronizing with child: join

```
parent()  
    create child thread  
    print "parent works"  
    child.join  
    print "parent continues"
```

```
child()  
    print "child works"
```

