

EECS 482: Introduction to Operating Systems

Lecture 9: Semaphores

Prof. Ryan Huang

Review: writing concurrent programs

Multiple threads work on shared data

- Must synchronize to ensure correct results

Synchronization

- Two types: mutual exclusion, ordering
- Difficult with low-level mechanism (e.g., atomic load, atomic store)

Monitors

- **Locks** for mutual exclusion
- **Condition variables** for ordering constraints

Semaphores

An abstract “integer” (initialized to some value N)

- Proposed by Dijkstra in the “THE” system in 1968
- Generalized lock

Provides two operations

- `down()` : wait for semaphore value to become positive, then decrement it by 1
 - Originally called `P()` , also `wait()`
- `up()` : increment semaphore value by 1
 - Originally called `V()` , also `signal()`

```
while (1) { Atomic
    if (value > 0) {
        value--
        break
    }
}
```

Atomic

```
value++
```

No other operations – not even just reading its value!

Semaphores (cont'd)

Represent a resource with N available units or a resource that allows concurrent accesses

- E.g., a parking lot with N available spots

`down ()` will return only N more times than `up ()`

Special case: **binary semaphore**

- Semaphore value is 0 or 1
- `up ()` atomically **sets** value to 1
- Versus **counting semaphore**

Is binary semaphore equivalent to lock?

Semaphores can enforce mutual exclusion and ordering

Mutual exclusion

```
semaphore sem(1)  
sem.down()  
critical section  
sem.up()
```

Ordering constraints

- Example: ensure that Thread A's task is done before Thread B's task

```
semaphore sem(0)
```

Thread A

```
task A  
sem.up()
```

Thread B

```
sem.down()  
task B
```

Recap: producer-consumer

Producers fill a shared buffer; consumers empty it

Coke machine problem

- Coke machine can hold at most MAX cokes
- Delivery person (producer) adds one coke to machine
- Consumer buys one coke

Producer-consumer with semaphores

As always, think about shared data, mutual exclusion, and before-after constraints

Semaphores:

- *mtx*: for exclusive access to coke machine
- *fullSlots*: counts number of cokes in machine
- *emptySlots*: counts number of empty slots in machine

Before-after constraints

- Consumer must wait if no cokes in machine
- Producer must wait if machine is full

Producer-consumer with semaphores

```
semaphore mtx(1);           // mutual exclusion to shared data
semaphore fullSlots(0);     // number of full slots
semaphore emptySlots(MAX);  // number of empty slots
```

```
producer {
    mtx.down()

    emptySlots.down()

    // add coke to machine

    fullSlots.up()

    mtx.up()

}
```

```
consumer {
    mtx.down()

    fullSlots.down()

    // take coke out of machine

    emptySlots.up()

    mtx.up()

}
```


Producer-consumer with semaphores


Why do we need different semaphores for fullSlots and emptySlots?

What if there's 1 full slot, and multiple consumers call down() at the same time?

Producer-consumer with semaphores

```
semaphore mtx(1);           // mutual exclusion to shared data
semaphore fullSlots(0);     // number of full slots
semaphore emptySlots(MAX);  // number of empty slots
```

```
producer {
    mtx.down()

     emptySlots.down()

    // add coke to machine

    fullSlots.up()

    mtx.up()

}
```

```
consumer {
    mtx.down()

    fullSlots.down()

    // take coke out of machine

    emptySlots.up()


    mtx.up()

}
```

What's wrong with this solution?

Producer-consumer with semaphores

```
semaphore mtx(1);           // mutual exclusion to shared data
semaphore fullSlots(0);     // number of full slots
semaphore emptySlots(MAX);  // number of empty slots
```



```
producer {
    mtx.down()

    emptySlots.down()

    // add coke to machine

    fullSlots.up()

    mtx.up()

}
```

```
consumer {
    mtx.down()

    fullSlots.down()

    // take coke out of machine

    emptySlots.up()

    mtx.up()

}
```

Any concerns?

Producer-consumer with semaphores


What about the state of the system between
between `emptySlots.down()` and `mtx.down()`?

Does the order of `up()` matter?

What if producer called `fullSlots.up()` earlier?

```
semaphore mtx(1);           // mutual exclusion to shared data
semaphore fullSlots(0);     // number of full slots
semaphore emptySlots(MAX);  // number of empty slots
```

```
producer {
    emptySlots.down()

    mtx.down()

    // add coke to machine

    mtx.up()

    fullSlots.up()

}
```

```
consumer {
    fullSlots.down()

    mtx.down()

    // take coke out of machine

    mtx.up()

    emptySlots.up()

}
```

Monitors vs. semaphores

Monitors

- Locks for mutual exclusion
- Condition variables for ordering

Semaphores: one mechanism for both mutual exclusion and ordering

- Elegantly minimal
- Can be difficult to use

Mutual exclusion (locks vs. semaphores)

lock = binary semaphore (initialized to 1)

lock() = down()

unlock() = up()

Mutex	Semaphore
<pre>mutex m m.lock() <critical section> m.unlock()</pre>	<pre>semaphore m(1) m.down() <critical section> m.up()</pre>

Ordering (condition variables vs. semaphores)

Condition variable	Semaphore
<code>while (condition) { wait(m) }</code>	<code>down()</code>
Waiting condition is in user code and uses user variables (more flexible)	Waiting condition is in semaphore code and uses semaphore's value (wait if $v \neq 0$)
User variables are protected by mutex	Semaphore's value is protected by semaphore's implementation of <code>up()</code> , <code>down()</code>
Must hold lock when calling wait	Must not hold "lock" (semaphore) when calling <code>down()</code> for ordering

Ordering (condition variables vs. semaphores)

Condition variable	Semaphore
No memory of past signals: <u>Thread A</u> wait() <u>Thread B</u> signal()	Remembers past up () calls <u>Thread A</u> down() <u>Thread B</u> up()
What happens if wait(), then signal()?	What happens if down(), then up()?
What happens if signal(), then wait()?	What happens if up(), then down()?
Why is this ok?	

Implementing custom waiting condition with semaphores

Semaphores work best if the shared integer and waiting condition ($\text{value} == 0$) map naturally to problem domain

How to implement custom waiting condition with semaphores?

What's a condition variable?

A place for threads to wait

"A place for threads"

- A set of waiting threads

"to wait"

- To wait, create a semaphore(0), add the semaphore to the waiting set, then call down() on it
- To signal, call up() on the waiting thread's semaphore

```
mutex cokeLock
cv waitingProducers
cv waitingConsumers
```

Producer

```
cokeLock.lock()

while (numCokes == MAX) {
    waitingProducers.wait(cokeLock)
}

numCokes++

waitingConsumers.signal()

cokeLock.unlock()
```

Consumer

```
cokeLock.lock()

while (numCokes == 0) {
    waitingConsumers.wait(cokeLock)
}

numCokes--

waitingProducers.signal()

cokeLock.unlock()
```

```
sem cokeLock(1)
set<sem*> waitingConsumers
set<sem*> waitingProducers
```

Producer

```
cokeLock.down()
while (numCokes == MAX) {
waitingProducers.wait(cokeLock)
    sem *s = new sem(0)
    waitingProducers.insert(s)
    cokeLock.up()
    s->down()
    cokeLock.down()
}
numCokes++
waitingConsumers.signal()
if (!waitingConsumers.empty()) {
    waitingConsumers.begin()->up()
    waitingConsumers.erase(
        waitingConsumers.begin())
}
cokeLock.up()
```

Consumer

```
cokeLock.down()
while (numCokes == 0) {
    waitingConsumers.wait(cokeLock)
}

numCokes--

waitingProducers.signal()

cokeLock.up()
```

```
sem cokeLock(1)
set<sem*> waitingConsumers
set<sem*> waitingProducers
```

Producer


```
cokeLock.down()
while (numCokes == MAX) {
waitingProducers.wait(cokeLock)
    sem *s = new sem(0)
    waitingProducers.insert(s)
    cokeLock.up()
    s->down()
    cokeLock.down()
}
numCokes++
waitingConsumers.signal()
if (!waitingConsumers.empty()) {
    waitingConsumers.begin()->up()
    waitingConsumers.erase(
        waitingConsumers.begin())
}
cokeLock.up()
```

Consumer

```
cokeLock.down()
while (numCokes == 0) {
waitingConsumers.wait(cokeLock)
    sem *s = new sem(0)
    waitingConsumers.insert(s)
    cokeLock.up()
    s->down()
    cokeLock.down()
}
numCokes--
waitingProducers.signal()
if (!waitingProducers.empty()) {
    waitingProducers.begin()->up()
    waitingProducers.erase(
        waitingProducers.begin())
}
cokeLock.up()
```

```
sem cokeLock(1)
set<sem*> waitingConsumers
set<sem*> waitingProducers
```

Producer

```
cokeLock.down()
while (numCokes == MAX) {
    sem *s = new sem(0)
    waitingProducers.insert(s)
    cokeLock.up()
     s->down()
    cokeLock.down()
}
numCokes++
if (!waitingConsumers.empty()) {
    waitingConsumers.begin()->up()
    waitingConsumers.erase(
        waitingConsumers.begin())
}
cokeLock.up()
```

Consumer

```
cokeLock.down()
while (numCokes == 0) {
    sem *s = new sem(0)
    waitingConsumers.insert(s)
    cokeLock.up()
    s->down()
    cokeLock.down()
}
numCokes--
if (!waitingProducers.empty()) {
    waitingProducers.begin()->up()
    waitingProducers.erase(
        waitingProducers.begin())
}
cokeLock.up()
```

What happens if we get interrupted after releasing the cokeLock but before downing the semaphore?