

# **EECS 482: Introduction to Operating Systems**

## **Lecture 14: Page Replacement**

Prof. Ryan Huang

# Administration

Fill in midterm teaching evaluation (due tomorrow)

## Midterm grade stats

Mean	61.85
Median	63.25
Standard deviation	16.0

## Midterm review session

- This Thursday 6-7:30 pm, CHRYS 220
- Regrade requests open after review session for a week

## Project 3 is out

- Due in about three weeks (March 28<sup>th</sup>)

# Page Replacement

When a page fault occurs, the OS needs a physical page to load the faulted page from disk into

What if all the page frames are in use?

The OS must choose a page frame to evict

- Free it up for use

**Page replacement algorithm** determines this

- Goal: minimize page faults
- Greatly affect performance of paging (virtual memory)
- Also called page eviction policies

# Locality

All paging schemes depend on locality

- Processes reference pages in localized patterns

## Temporal locality

- Locations referenced recently likely to be referenced again

## Spatial locality

- Locations near recently referenced locations are likely to be referenced soon

Processes usually exhibit both kinds of locality,  
making paging practical despite its costs

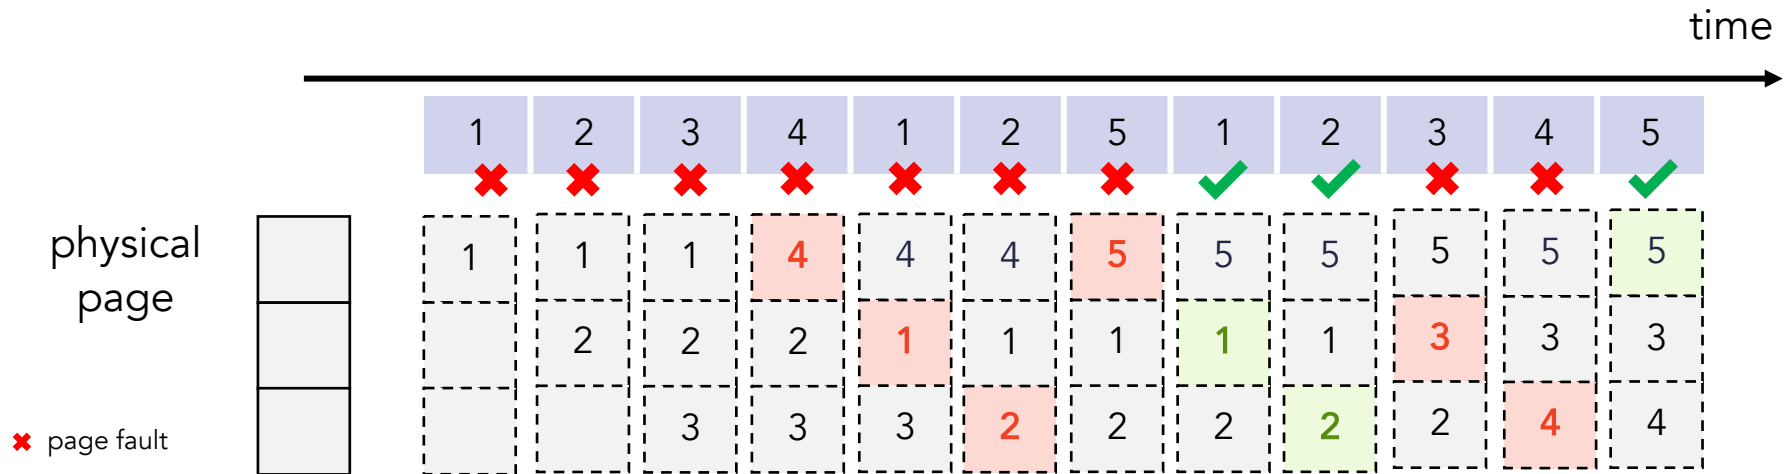
# First-In First-Out (FIFO)

## Evict "oldest" page

- Brought into memory longest time ago

## Example

- Access trace (virtual page #): 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3 physical pages



9 page faults

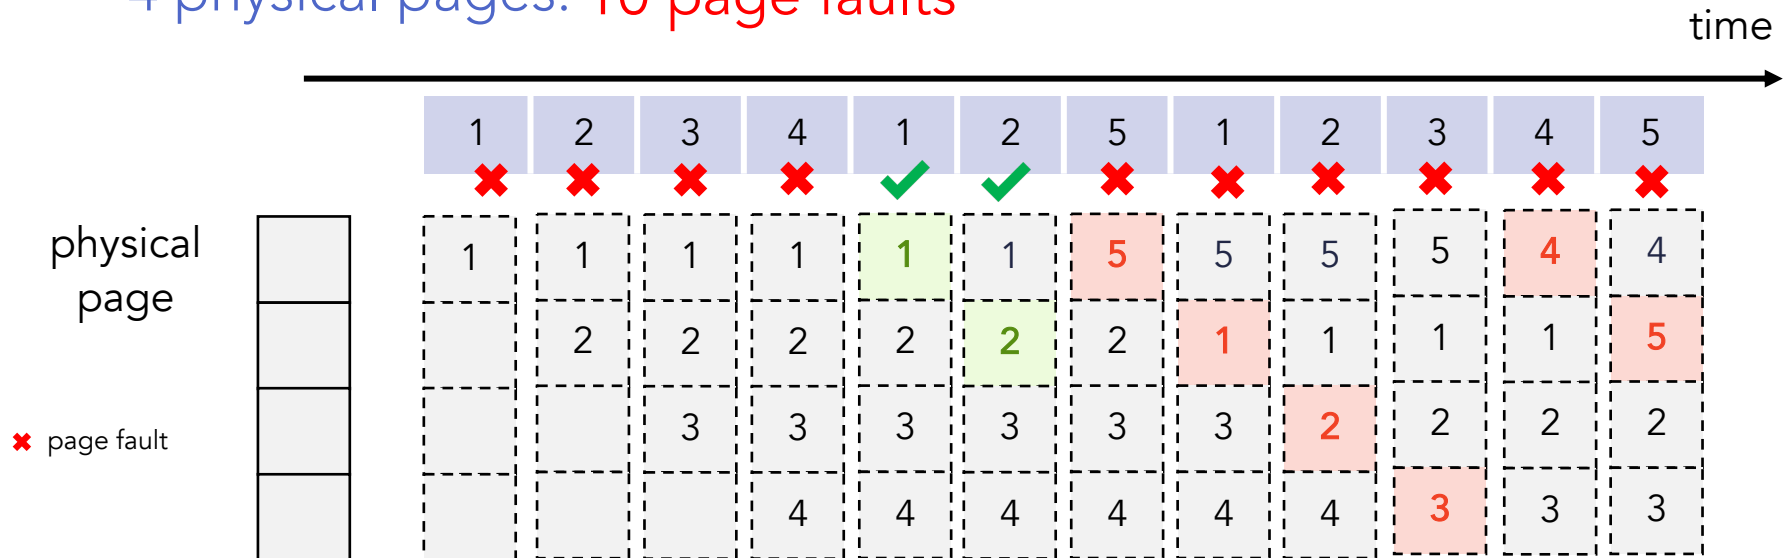
# First-In First-Out (FIFO)

## Evict "oldest" page


- Brought into memory longest time ago

## Example

- Access trace (virtual page #): 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3 physical pages: 9 page faults
- 4 physical pages: 10 page faults



# Belady's Anomaly



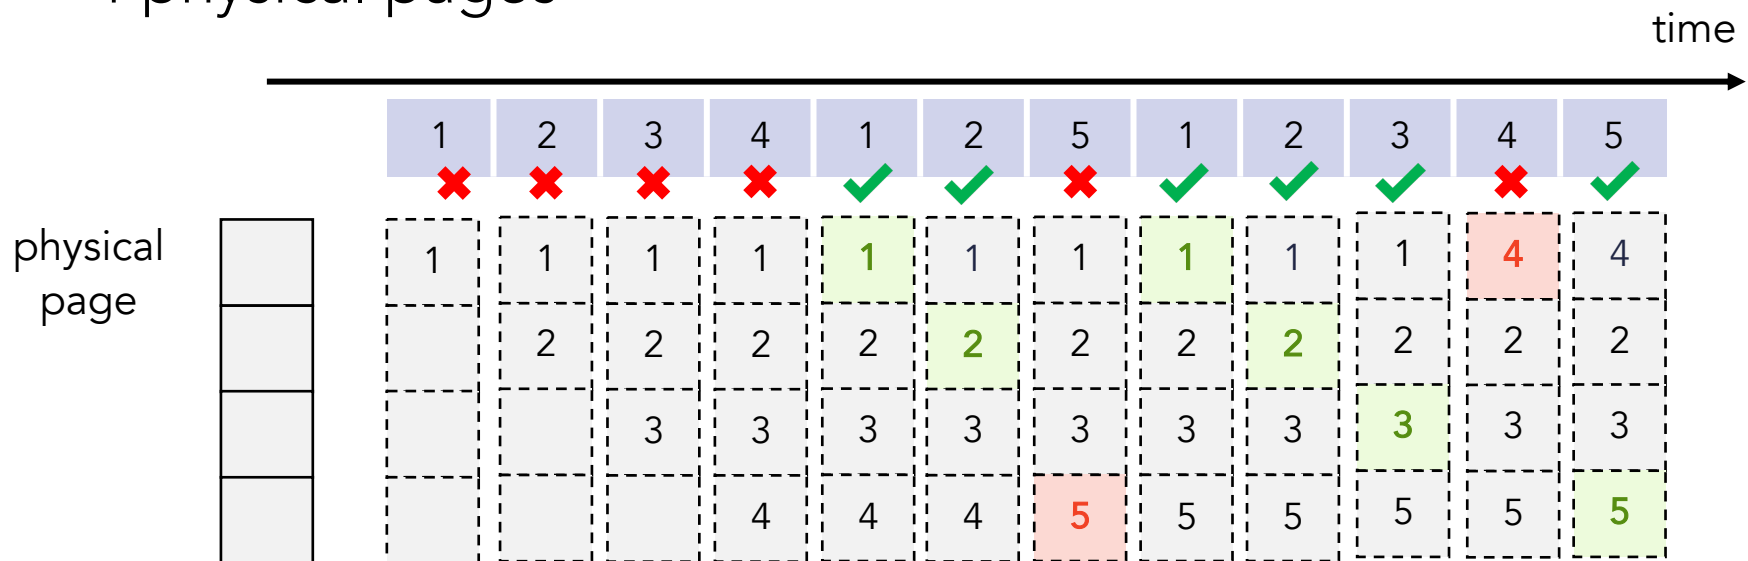
More physical memory doesn't always mean fewer faults!

# Optimal replacement (OPT)

Replace page that will not be used for longest time in the future

## Example

- Access trace (virtual page #): 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 4 physical pages



6 page faults



# Optimal replacement (OPT)

## Belady's Algorithm

- Property: minimal page faults
- Rationale: the best page to evict is the one never touched again
  - Never is a long time, so picking the page closest to "never" is the next best thing
- Proved by Belady

## Problem?

- Requiring knowing the future!

## Why is Belady's algorithm useful then?

- Use it as a yardstick
- Compare page replacement algorithms with the optimal
  - If optimal is not much better, then algorithm is pretty good
  - If optimal is much better, then algorithm could use some work

# Least recently used (LRU)

Replace the page that was last used longest ago

- If page hasn't been used for a while, it probably won't be used for a long time in the future

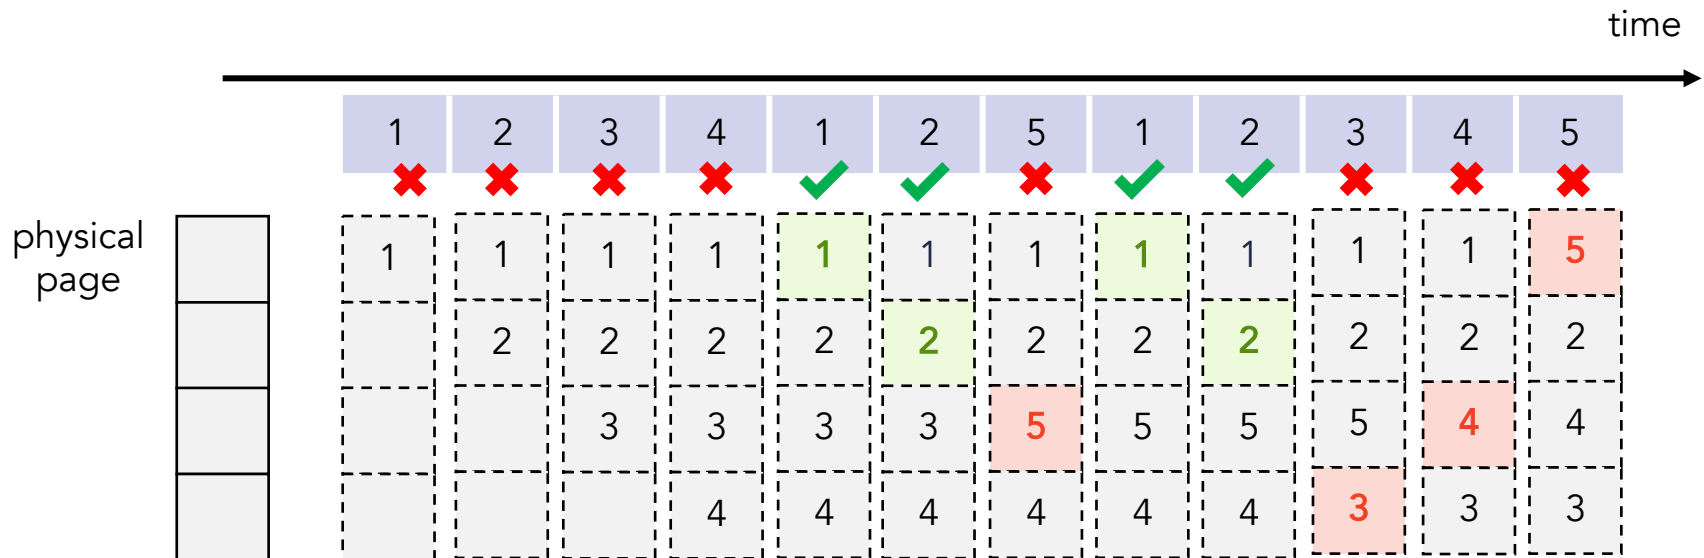
Approximates OPT

- **Temporal locality**: the future tends to reflect the past

# Least recently used (LRU)

## Example

- Access trace (virtual page #): 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 4 physical pages 8 page faults



## Problems

- Can be pessimal – example?
  - Looping over memory (then want MRU eviction)
- How to implement?

# Strawman LRU Implementations

## Stamp PTEs with timer value?

- E.g., CPU has cycle counter
- Automatically writes value to PTE on each page access
- Scan page table to find oldest counter value = LRU page
- Problem?
  - Would double memory traffic!

## Keep doubly-linked list of pages?

- On access remove page, place at tail of list
- Problem: again, very expensive

## What to do?

- Approximate LRU, don't try to do it exactly

# Approximating LRU

Most MMUs maintain an **accessed/referenced** bit

- In PTEs
- Set by MMU when a page is read or written
- Can be cleared by OS

How to use access bit to identify pages that have *not been used for a while?*



# Clock algorithm


Do FIFO but skip accessed pages

Keep **resident pages** in circular list


Scan: inspect page pointed by clock hand

- If access bit is set, this page has been accessed “recently”
  - Clear the bit, skip this page, advance clock hand
- If access bit is not set, this page has not been accessed “recently”
  - Evict

A.k.a. second-chance replacement



# Clock example



What if all pages were referenced since last sweep?

# Page eviction

What to do with data from evicted page? Why?

When do you not need to write page to disk?

This is a “write-back” cache

How to tell?

- Most MMUs provide a “dirty” bit for each resident page
- MMU sets “dirty” bit when page is written
- “dirty” means “content different from disk”

Why not write to disk on every store (write-through cache)?



# Optimizing when to do work

do work earlier: may  
reduce waiting later



do work later: may  
reduce total work




time



When would you not need to write a new data value to disk?

Project 3 memory manager should be as lazy as possible  
(with CLOCK replacement)

# MMU algorithm



```
if (virtual page is invalid or non-resident or protected) {
```

```
    trap to OS fault handler
```

```
    retry access
```

```
} else {
```

```
    physical page # = pageTable[virtual page #].physPageNum
```

```
    pageTable[virtual page #].referenced = true
```

```
    if (access is write) {
```

```
        pageTable[virtual page #].dirty = true
```

```
    }
```

```
    access physical memory at {physical page #}{offset}
```

```
}
```

# Page table contents

PTE

physical page #

valid

resident

protection

dirty

referenced

```
if (virtual page is invalid or non-resident or protected) {  
    trap to OS fault handler  
    ...  
}
```

Other information about page is stored in OS data structure (e.g., location on disk)

**Do we *have to* keep a valid bit in PTE?**

- Mark invalid pages as non-resident, then sort out specifics in OS after fault

**Do we *have to* keep a resident bit in PTE?**

- Mark non-resident pages as non-readable/non-writable, then sort out specifics in OS after fault

# Page table contents

## minimalist

PTE

physical page #

protection

dirty

referenced

## Do we have to keep the dirty bit?

- Have OS (not MMU) maintain dirty bit in its own data structure
- Naive solution: trap on every store & mark dirty
  - improvement: which store instructions change "dirty" bit?

## Do we have to keep the reference bit?

## General pattern:

- What information are you maintaining?
- What accesses change this information?
- Set protection to trap on these accesses

H/W typically do keep more info than the minimal

31	12	11	10	9	8	7	6	5	4	3	2	1	0
Physical page number		Ignored			G	A T	D	A	C D	W T	U	W	P

Diagram annotations: Red dashed arrows point from the labels 'dirty', 'accessed (reference)', and 'present (resident)' to bits 6, 5, and 0 respectively in the table above.