

EECS 482: Introduction to Operating Systems

Lecture 16: Managing multiple address spaces

Prof. Ryan Huang

Physical memory allocation

How to divide physical memory among processes?

- Fairness vs. efficiency

Variable space

- A process' set of pages grows and shrinks dynamically
- **Global replacement**: can evict pages from any process
 - One process can ruin it for the rest

Fixed space

- Each process gets a fixed number of physical pages
- **Local replacement**: can only evict a process' own pages
 - Some processes may do well while others suffer

Virtual memory performance

Physical memory is a cache for address space

- Cache miss is a page fault
- Swap page from disk to memory, slow (disk I/O)

Performance degrades rapidly as miss rate goes up

- Avg access time = hit rate * hit time + miss rate * miss time
- E.g., hit time = .0001 ms; miss time = 10 ms
 - Average access time (100% hit rate) = .0001 ms
 - Average access time (99% hit rate) = .100099 ms
 - Average access time (90% hit rate) = 1.0009 ms

High miss rate → thrashing

- OS spent most time in paging data back and forth from disk
- Little time spent doing useful work (making progress)


Reasons for thrashing

The system is **overcommitted**

- No idea which pages should be in memory to reduce faults

Common reasons

- Access pattern has no locality
- Each process fits individually, but too many processes



Solutions to thrashing?

Buy more DRAM!

Change scheduler to reduce thrashing

- E.g., system has 1 GB memory
- How to run processes A, B, C, D (each 500 MB)?
- Why does this help?

Working set

Thrashing depends on portion of address space being **actively used** by each process

What do we mean by “actively using”?

A **working set** of a process models the dynamic locality of its memory usage

- $WS(t, w) = \{P \mid \text{Page } P \text{ was referenced in the time interval } (t - w, t) \}$
 - w – working set window (in absolute time or units of mem references)

WS size: # of unique pages referenced in $(t-w, t)$

- The working set size changes with program locality
- During periods of poor locality, the working set size becomes larger

Working set (cont'd)

Intuition: want the working set to be in memory to prevent heavy faulting

- *Expect to run for w time without getting a fault*
- Sum of all working sets should fit in memory
- Only run subset of processes that fit in memory

Difficult questions

- How do we determine w ?
- How to handle changes in working set?

Not a practical model, but the intuition is still valid

- When people ask, "*How much memory does Firefox need?*", they are in effect asking for the size of Firefox's working set

Creating a process

A process is created by another process

- Parent is creator, child is created (Unix: ps "PPID" field)
- What about the first process?
 - Unix: init process (PID 1), ancestors of all other processes

OS provides process creation system calls

After creating a child process

- the parent process may either wait for the child to finish its task or continue in parallel

Process creation: Unix

In Unix, processes are created using `fork()`

```
int fork()
```

1. Creates and initializes a new Process Control Block (PCB)
2. Creates a new address space
3. Initializes the address space with a **copy** of the address space of the parent
4. Initializes the kernel resources to point to the parent's resources (e.g., open files)
5. Places the PCB on the ready queue

`fork()` **returns twice**

- Huh?
- Returns the child's PID to the parent, "0" to the child

Fork() example

```
#include <stdio.h>
#include <unistd.h>
int main(int argc, char *argv[])
{
    char *name = argv[0];
    int child_pid = fork();
    if (child_pid == 0) {
        printf("Child of %s is %d\n", name, getpid());
        return 0;
    } else {
        printf("My child is %d\n", child_pid);
        return 0;
    }
}
```

What does this program print?

Example output

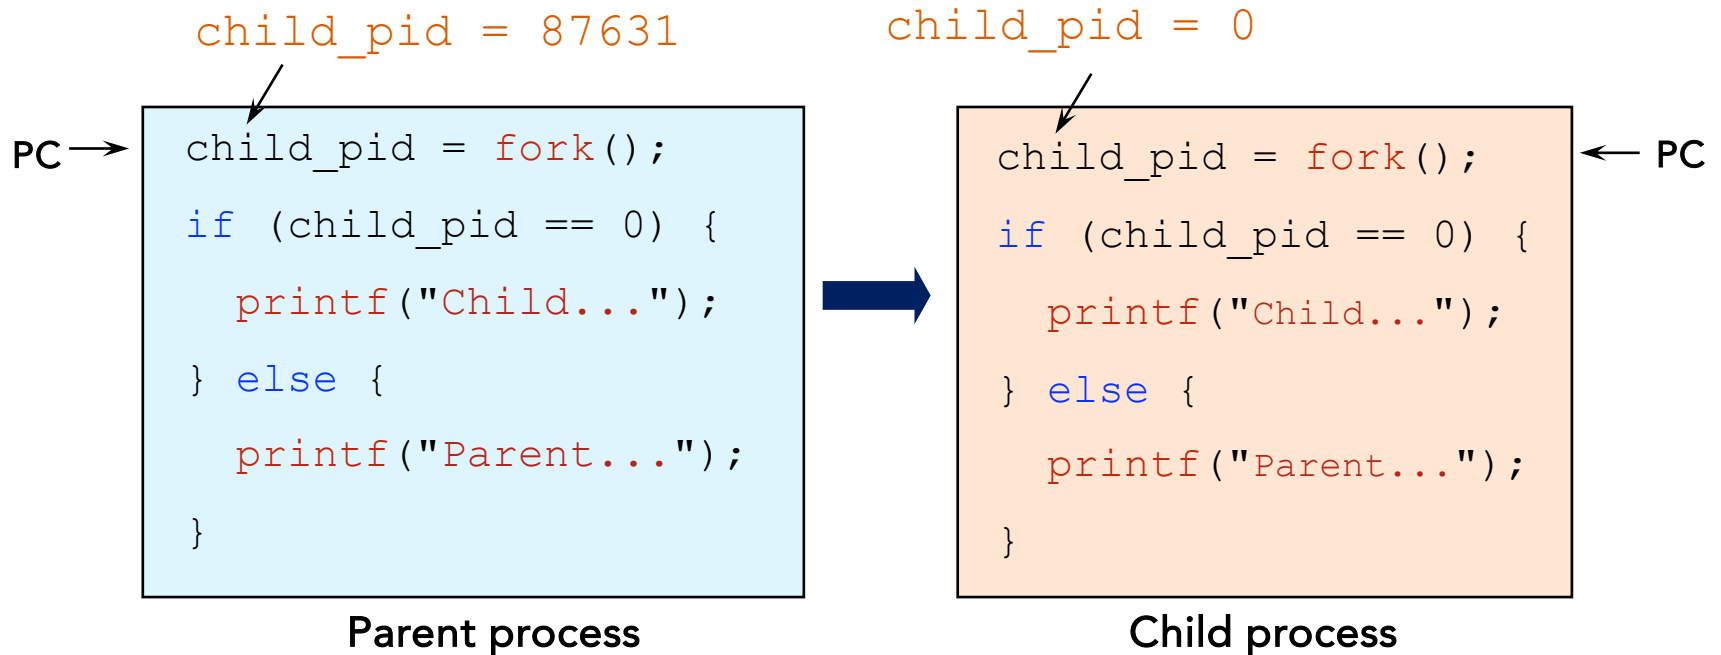
```
$ gcc -o test test.c
```

```
$ ./test
```

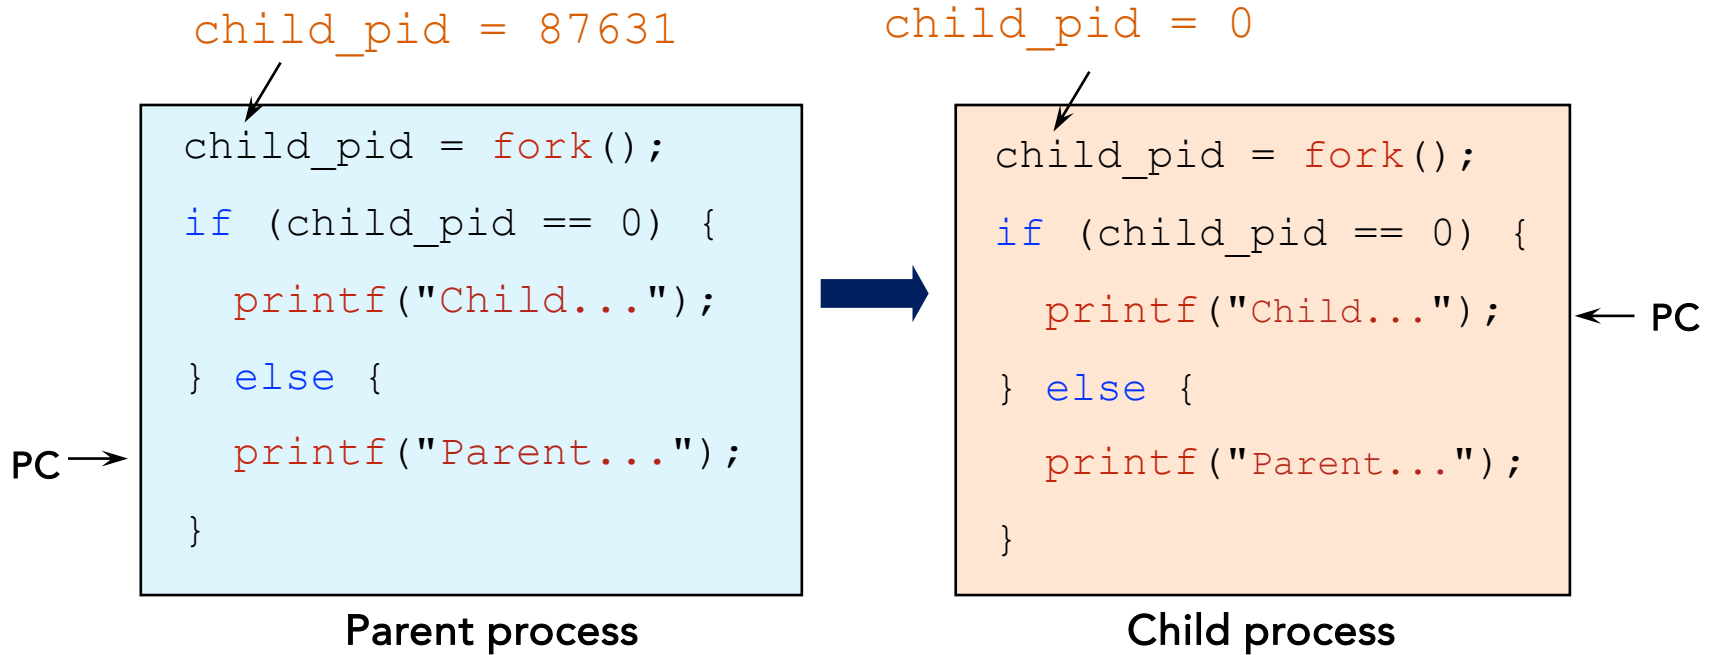
My child is 87631

Child of ./test is 87631

Duplicating address spaces



Divergence

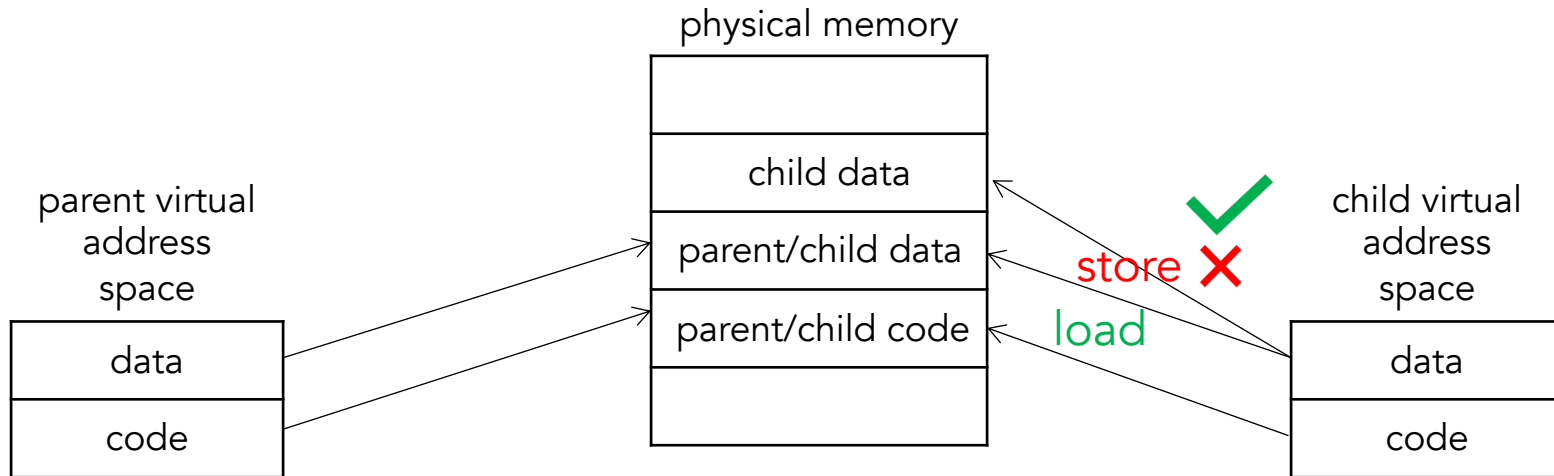


Avoiding work on fork()

Copying entire address space is expensive

Instead, make it **look** like it's copied, but do the work later

This is called **copy-on-write**



What if parent process does a store to data?

Process creation: Unix (2)

How do we actually start a new program?

```
int execv(char *prog, char *argv[])  
int execve(const char *filename, char *const argv[], char  
*const envp[])
```

`execv()`

1. Stops the current process
2. **Replace** the process' address space with "prog"
3. Initializes hardware context and args for the new program
4. Places the PCB onto the ready queue

Note: `exec` **does not** create a new process

What does it mean for `exec` to return?

Why `fork()`?

Most calls to `fork` followed by `exec`

- could also combine into one `spawn` system call

Very useful when the child...

- Is cooperating with the parent
- Relies upon the parent's data to accomplish its task

Example: web server

Example: shell

```
while (1) {  
    int sock = accept();  
    if ((child_pid = fork()) == 0) {  
        // Handle client request  
    } else {  
        // Close socket  
    }  
}
```


Why `fork()`?

Most calls to `fork` followed by `exec`

- could also combine into one `spawn` system call

Very useful when the child...

- Is cooperating with the parent
- Relies upon the parent's data to accomplish its task

Real win is simplicity of interface

- Tons of things you might want to do to child process:
 - manipulate file descriptors, set environment variables, reduce privileges, ...
- Yet **`fork()`** requires no arguments at all

Process creation: Windows

Without `fork`, needs tons of different options for creating a new process

Example: Windows `CreateProcess` system call

- Combines both `fork()` and `exec()`

```
BOOL WINAPI CreateProcess(  
    _In_opt_ LPCTSTR lpApplicationName,  
    _Inout_opt_ LPTSTR lpCommandLine,  
    _In_opt_ LPSECURITY_ATTRIBUTES lpProcessAttributes,  
    _In_opt_ LPSECURITY_ATTRIBUTES lpThreadAttributes,  
    _In_ BOOL bInheritHandles,  
    _In_ DWORD dwCreationFlags,  
    _In_opt_ LPVOID lpEnvironment,  
    _In_opt_ LPCTSTR lpCurrentDirectory,  
    _In_ LPSTARTUPINFO lpStartupInfo,  
    _Out_ LPPROCESS_INFORMATION lpProcessInformation  
);
```

- Also `CreateProcessAsUser`, `CreateProcessWithLogonW`, `CreateProcessWithTokenW`, ...