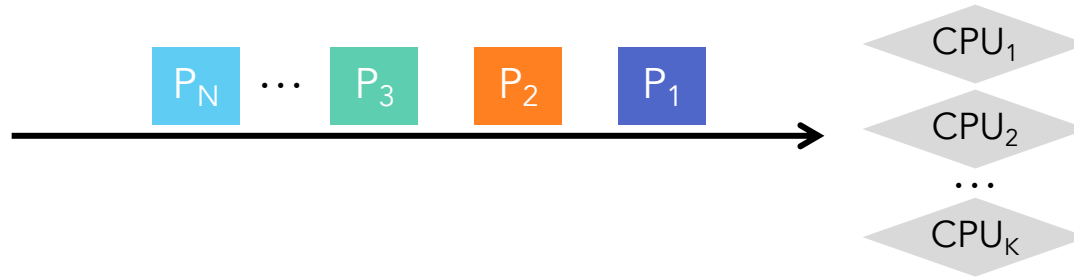


# **EECS 482: Introduction to Operating Systems**

## **Lecture 17: CPU Scheduling**

Prof. Ryan Huang

# CPU scheduling



## The scheduling problem:

- Have  $N$  jobs ready to run
- Have  $K \geq 1$  CPUs

## Choose which job to run on which CPU, for how long?

- We'll refer to schedulable entities as **jobs** – could be processes, threads, people, etc.

## How to choose

- Project 2 and 3: FIFO
- Many other policies possible

# Scheduling criteria

## Why do we care?

- How do we measure the effectiveness of a scheduling algorithm?

## Two broad goals

- Performance
- Fairness

# Performance

Throughput – # of processes that complete per unit time

- $\# \text{ jobs/time}$  (higher is better)

Turnaround time – time for each process to complete

- $T_{\text{finish}} - T_{\text{start}}$  (lower is better)

Response time – time from request to *first* response

- $T_{\text{response}} - T_{\text{request}}$  (lower is better)
- e.g., key press to echo (not launch to exit)

Above criteria are affected by secondary criteria

- *Waiting time* –  $\text{Avg}(T_{\text{wait}})$  time each process waits in the ready queue
- *CPU utilization* –  $\%CPU$  fraction of time CPU doing productive work

# Fairness

Second goal of CPU scheduling: share CPU among threads in “fair” manner

What does “fair” mean?

Fairness may conflict with performance

Starvation = extremely unfair

# First-come, first-served (FCFS)

Run jobs in order that they arrive

- E.g., Say  $P_1$  needs 24 sec, while  $P_2$  and  $P_3$  need 3.
- Say  $P_2, P_3$  arrived immediately after  $P_1$ , get:



Throughput: 3 jobs / 30 sec = 0.1 jobs/sec

Turnaround Time:  $P_1 : 24, P_2 : 27, P_3 : 30$

- Average TT:  $(24 + 27 + 30) / 3 = 27$

Can we do better?

# FCFS continued

Suppose we scheduled  $P_2$ ,  $P_3$ , then  $P_1$

- Would get:



Throughput: 3 jobs / 30 sec = 0.1 jobs/sec

Turnaround Time:  $P_1 : 30$ ,  $P_2 : 3$ ,  $P_3 : 6$

- Average TT:  $(30 + 3 + 6) / 3 = 13$  – much less than 27

**Lesson:** scheduling algorithm can reduce TT

Can a scheduling algorithm improve throughput?

- Yes, if jobs require both computation and I/O


# FCFS limitations

## FCFS algorithm is non-preemptive in nature

- Once CPU time has been allocated to a process, other processes can get CPU time only after the current process has finished or gets blocked.

## Short jobs can be stuck behind long ones

- This property of FCFS scheduling is called *Convoy Effect*





# Shortest job first (SJF)

## Shortest Job First (SJF)

- Choose the job with the smallest expected CPU burst
  - Person with smallest # of items in shopping cart checks out first

## Example

- Three jobs,  $P_1$  needs 8 sec,  $P_2$  needs 4 sec,  $P_3$  needs 2 sec









Average Waiting Time:  $(0 + 2 + 6) / 3 = 2.67$

# SJF is optimal

SJF has *provably* optimal (minimum) *average waiting time (AWT)*

Previous example:  $P_1$  8 sec,  $P_2$  4 sec,  $P_3$  2 sec

- How many possible schedules?

schedule 1		$AWT = (0+8+12)/3 = 6.67$
schedule 2		$AWT = (0+8+10)/3 = 6$
schedule 3		$AWT = (0+4+12)/3 = 5.33$
schedule 4		$AWT = (0+4+6)/3 = 3.33$
schedule 5		$AWT = (0+2+10)/3 = 4$
SJF		$AWT = (0+2+6)/3 = 2.67$

# SJF limitations

Can potentially lead to unfairness or starvation

Impossible to know size of CPU burst ahead of time

- Like choosing person in line without looking inside cart

How can you make a reasonable guess?

- Estimate CPU burst length based on past
- E.g., exponentially weighted average


# Round robin (RR)

## FIFO + periodic preemption

- Each job is given a time slice called a **quantum**
- Preempt job after duration of quantum
- When preempted, move to back of FIFO queue

## Example:

- Three jobs,  $P_1$  needs 24 sec, while  $P_2$  and  $P_3$  need 3 sec
- Quantum is 1 sec



- Average TT:  $(30 + 8 + 9) / 3 = 15.67$ 
  - Compared to FCFS: 27


# Round robin vs. FCFS

Round robin achieved lower average turnaround time than FCFS in last example

- Is it always the case?

Another example:

- Two jobs, both arrive at time 0, both take 100 sec



- What would the average turnaround time be with RR?
  - Turnaround Time:  $P_1=199$ ,  $P_2=200$
  - Average TT:  $(199 + 200) / 2 = 199.5$
- How does it compare to FCFS?
  - Turnaround Time:  $P_1=100$ ,  $P_2=200$
  - Average TT:  $(100+200) / 2 = 150$

# Round robin trade-offs

## Advantages:

- Fair allocation of CPU across jobs
- Low average waiting time when job lengths vary
- Good for responsiveness if small number of jobs

## Disadvantages

- Longer average turnaround time when job lengths are uniform
- Context switches are frequent and need to be very fast

# Round robin time quantum

## How to pick quantum?

- What if time slice is too long?
- What if time slice is too short?
- Want much larger than context switch cost
- Majority of CPU bursts should be less than quantum
- But not so large system reverts to FCFS

Typical values: 1–100 msec

# Priority

## Priority scheduling

- Associate a numeric priority with each process
  - E.g., smaller number means higher priority (Unix/BSD)
- Give CPU to the process with highest priority

## Problem?

- Starvation: low-priority jobs can wait indefinitely

## Solution?

- "Aging"
  - Increase priority as a function of waiting time
  - Decrease priority as a function of CPU consumption



# Scheduling summary

Many different policies: FCFS, round robin, SJF, priority, deadline, proportional share, etc.

- OS schedulers use different policies in different settings
  - Interactive, CPU-bound, batch, etc.
- Still actively researched today!

Combining scheduling algorithms to optimize for multiple objectives

- Have multiple queues
- Use a different algorithm for each queue
- Move processes among queues
- Example: Multiple-level feedback queues (MLFQ)