

EECS 482: Introduction to Operating Systems

Lecture 21: Networks

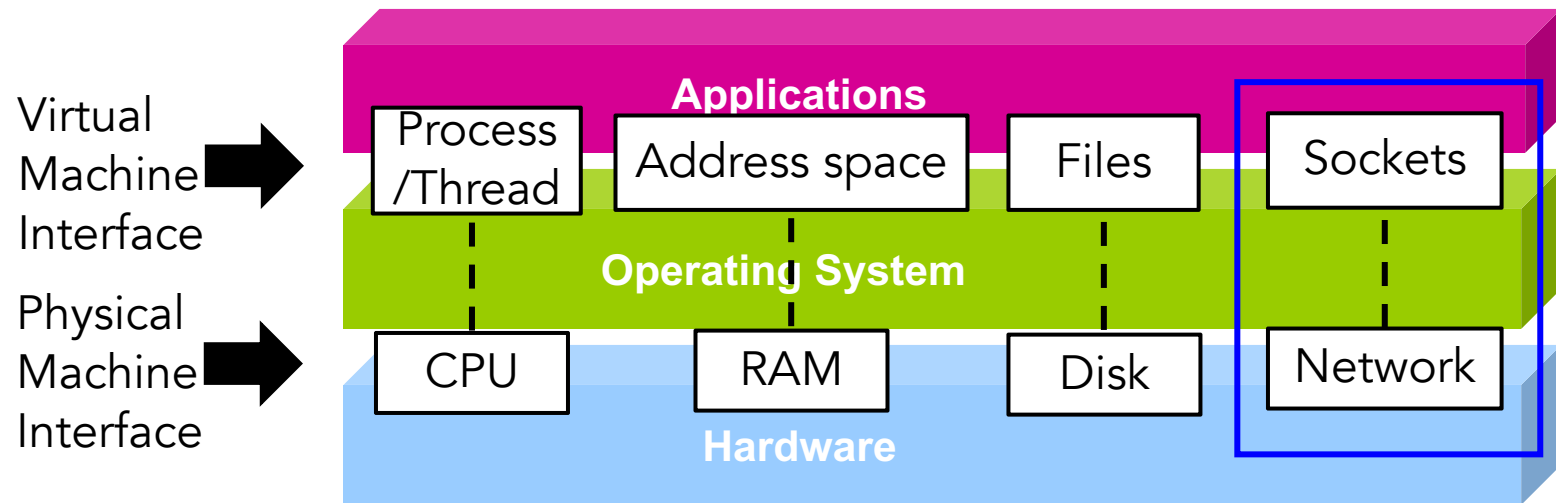
Prof. Ryan Huang

Administration

Student + staff social #2

- Thursday, April 3rd, from 19:00 to 21:00 in GGBL 2147
- Board games and card games
- RSVP by like the Piazza post @1265

OS abstractions




Network abstractions

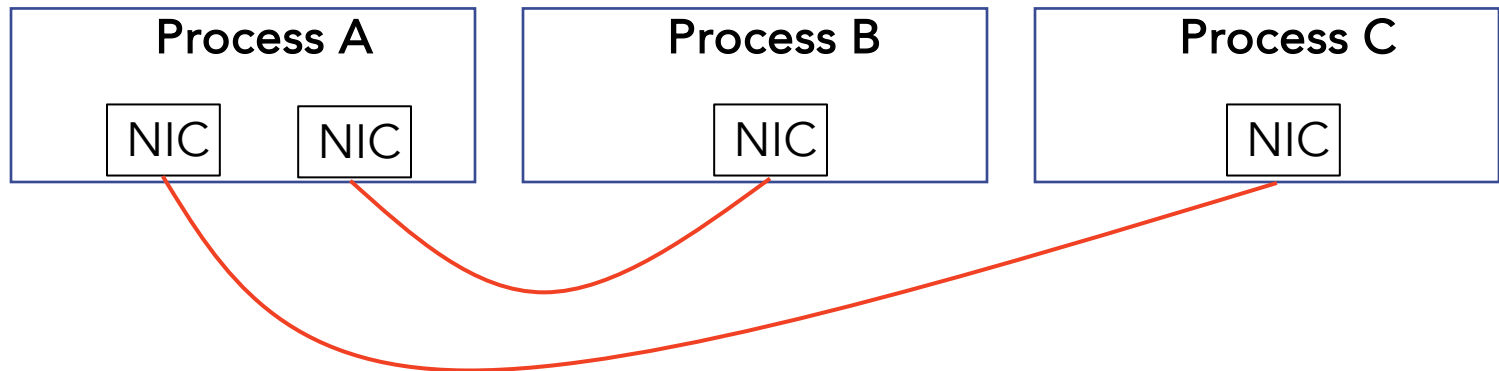
Hardware reality	OS abstraction
Destination named by MAC address of NIC	Destination named by symbolic hostname
Deliver to computer on LAN	Route to final destination across multiple networks
Machine-to-machine	Process-to-process
Unordered, unreliable	Ordered, reliable
Finite-size messages	Byte streams
Multiple distinct computers, each with own resources	Single computer with combined resources

OS abstraction of network

Hardware reality



OS abstraction



Machine-to-machine communication

➔ Process-to-process communication

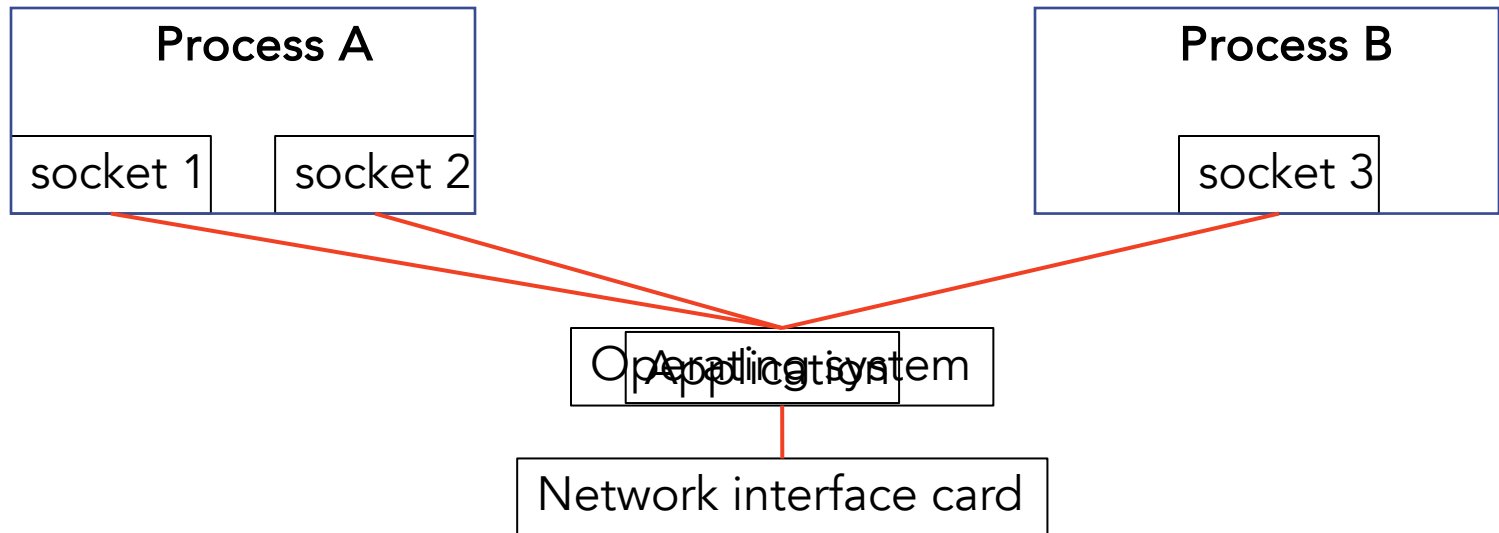
Machine ➔ process

- Processors ➔ threads
- DRAM ➔ address spaces
- Disks ➔ files and directories
- Network interface card ➔ sockets

Socket

- Virtual network interface card
- Endpoint for communication
- NIC named by MAC/IP address; socket named by “port number” (via bind)
- Programming interface: BSD sockets

Socket abstraction



Apartment analogy

Apartment building	Computer

Types of sockets

UDP (user datagram protocol): IP + sockets

TCP (transmission control protocol): IP + sockets + reliable, ordered byte streams

How to provide **reliable, ordered byte streams** on top of **unreliable, unordered packets**?

Unordered → ordered messages

Hardware reality: messages can be re-ordered

- Sender: msg A, msg B, msg C, msg D, msg E
- Receiver: msg A, msg B, msg D, msg C, msg E

OS abstraction: messages received in order sent

How to provide this abstraction?

- Assign sequence numbers to successive messages
- What to do once an arrived message is out of order?

• E.g., #0, #1, #3, #2 #4 →

- Drop #3?
- Drop #2?
- Save #3, deliver after #2 is delivered?

Unordered → ordered messages

Message ordering requires a notion of network "connection"

- Should not enforce ordering for unrelated messages

TCP: connects two sockets with a point-to-point, bidirectional, ordered channel

- Server calls `accept` to accept incoming connection
- Client calls `connect` to establish connection with server

Sequence # is specific to a socket-to-socket connection

Unreliable → reliable messages

Hardware interface: messages can be

- Dropped
- Duplicated
- Corrupted

OS abstraction: each message is delivered exactly once (and in the right order)

Unreliable → reliable messages

How to detect and fix a **dropped message**?

How to fix (once you've detected that a message was dropped)?

How does sender know the message was dropped?

Extra communication needed

Is the sender's conclusion always correct?

Unreliable → reliable messages

How to deal with duplicate messages?

How to deal with corrupted messages?

Solving problems via transformation

- Corrupted messages → dropped messages
- Potential dropped messages → potential duplicates
- Solve duplicates by adding sequence # and dropping duplicates

Byte streams

Hardware interface: Send/receive distinct (finite-sized) messages

OS abstraction: Send/receive infinite stream of bytes


How to implement the byte stream abstraction?

- Sender:
- Receiver:

Using byte streams

TCP has no message boundaries (unlike UDP)

- Sender sends bytes; receiver receives bytes



```
size = recv(sock, buf, 150, 0);
```

- Blocks until can receive > 0 bytes (or error)
- May return $[1, 150]$ bytes
- To receive > 1 byte, must call `recv` in loop
 - or use `MSG_WAITALL` to receive a known number of bytes

How to send message over byte streams?

How to know # of bytes to receive?

- Convention (size specified by protocol)
- Size specified in header
- End-of-message delimiter
- Sender closes connection

Project 4

Read manual pages

Understand C strings

- What are you passing when you call `fs_writeblock(..., "hello world", ...)` ?

What data should be a C string, and what data need not be a C string?

- File data?
- `str` functions assume data is a C string. Caller must guarantee this pre-condition.

What data is trusted, and what is not trusted?

- Data sent by client?
- Initial file system image?