

EECS 482: Introduction to Operating Systems

Lecture 22: Client-Server

Prof. Ryan Huang

Client-server

Common way to structure a distributed application:

- **Server** provides some centralized service
- **Client** makes request to server, waits for response

Example: **web server**

- Server stores and returns web pages
- Clients run Web browsers, make GET/POST requests

Example: **coke machine**

- Server manages state associated with coke machine
- Clients call `client_produce()` or `client_consume()`, which send request to the server and return when done
- Server responds to client after request has been satisfied

Client-server example

```
client_produce() {  
    send produce request to server  
    receive response  
}  
  
server() {  
    while (1) {  
        receive request  
        if (produce request) {  
            add coke to machine  
        } else {  
            take coke out of machine  
        }  
        send response  
    }  
}
```

Problems?

Example with waiting

```
client_produce() {  
    send produce request to server  
    receive response  
}  
  
server() {  
    while (1) {  
        receive request  
        if (produce request) {  
            while (machine is full) wait  
            add coke to machine  
            signal  
        } else {  
            ...  
        }  
        send response  
    }  
}
```

Problems?

Example with threads

```
server() {  
    while (1) {  
        receive request  
        if (produce request) {  
            create thread that calls server_produce()  
        } else {  
            create thread that calls server_consume()  
        }  
    }  
}
```

Problems?

```
server_produce() {  
    lock  
    while (machine is full) {  
        wait  
    }  
    add coke to machine  
    signal  
    send response  
    unlock  
}
```

Example with threads

```
server() {  
    while (1) {  
        wait for new connection from client  
        create thread that calls handle_request()  
    }  
}
```

```
handle_request() {  
    receive request  
    call server_produce() or server_consume()  
}
```

```
server_produce() {  
    lock  
    while (machine is full) {  
        wait  
    }  
    add coke to machine  
    signal  
    unlock  
    send response  
}
```

Alternatives

How to lower overhead of creating threads?

How to structure the server that avoid blocking for slow operations

- Synchronous/blocking
 - Threads
- Asynchronous/non-blocking
 - Polling (via `select`)
 - Events

What are the slow operations in producer/consumer?

What are the slow operations in Project 4?

Client-server interface

Message send/receive

Problem: too low-level and tiresome

- Need to worry about message formats
- Must wrap up information into message at source
- Must decide what to do with message at destination
- Have to pack and unpack data from messages
- May need to sit and wait for multiple messages to arrive

Messages are not a very natural programming model

Alternative?

Procedure Calls

Procedure calls are a more natural way to communicate

- Every language supports them
- Semantics are well-defined and understood
- Natural for programmers to use

Idea: let servers export procedures that can be called by client programs

- Similar to module interfaces, class definitions, etc.
- Clients just do a procedure call as if they were directly linked with the server
- Under the cover, the procedure call is converted into a message exchange with the server

Remote Procedure Calls (RPC)

Use procedure call as a model for distributed (remote) communication

Lots of issues

- How do we make this invisible to the programmer?
- What are the semantics of parameter passing?
- How do we bind (locate, connect to) servers?
- How do we support heterogeneity (OS, arch, language)?
- How do we make it perform well?

RPC model

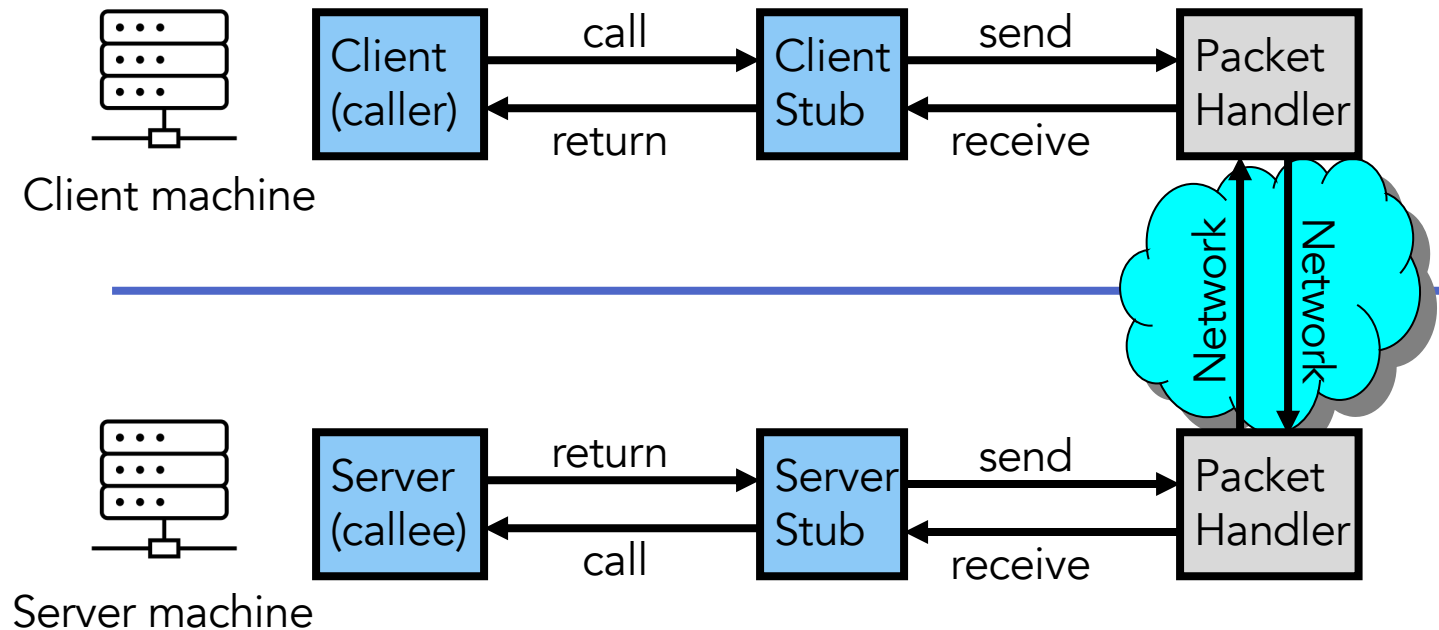
Two **stubs** for each server procedure: client and server

Server programmer implements the server procedures and **links** them with server-side stubs

Client programmer implements the client program and **links** it with client-side stubs

Stubs are the “glues” for managing all details of the remote communication between client and server

RPC information flow



RPC stubs

Client stub:

1. Constructs message with function name and parameters
2. Sends request message to server
7. Receives response from server
8. Returns response to client

Server stub:

3. Receives request message
4. Invokes correct function with specified parameters
5. Constructs response message with return value
6. Sends response to client stub

Producer-consumer using RPC

Client stub

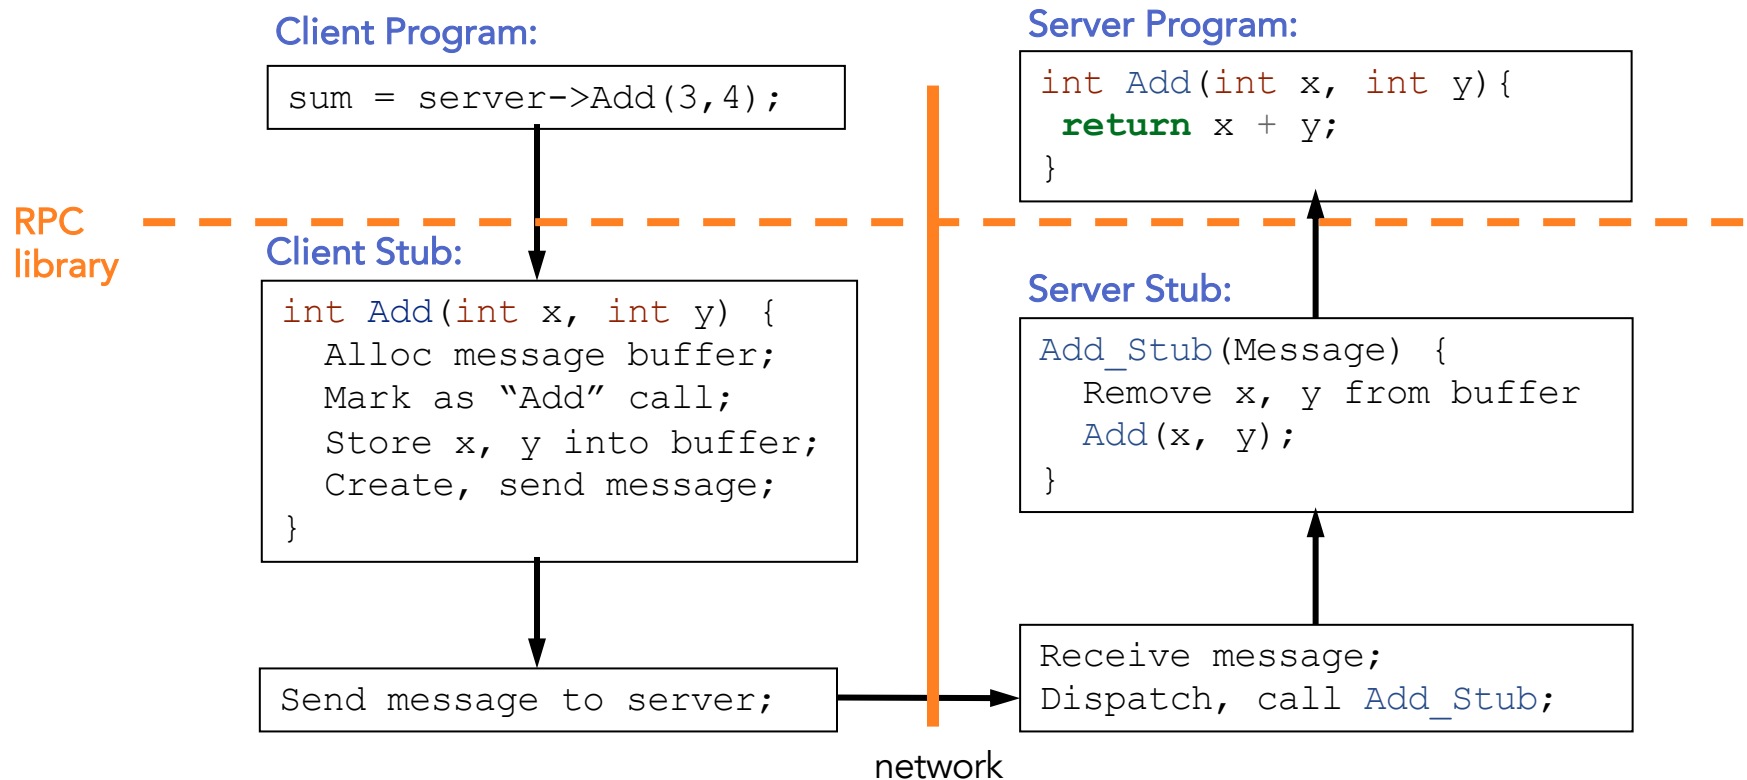
```
int produce (int n) {  
    int status;  
    → send (sock, &n, sizeof(n));  
    → recv (sock, &status, sizeof(status)); // add MSG_WAITALL or loop  
    → return(status);  
}
```

Server stub

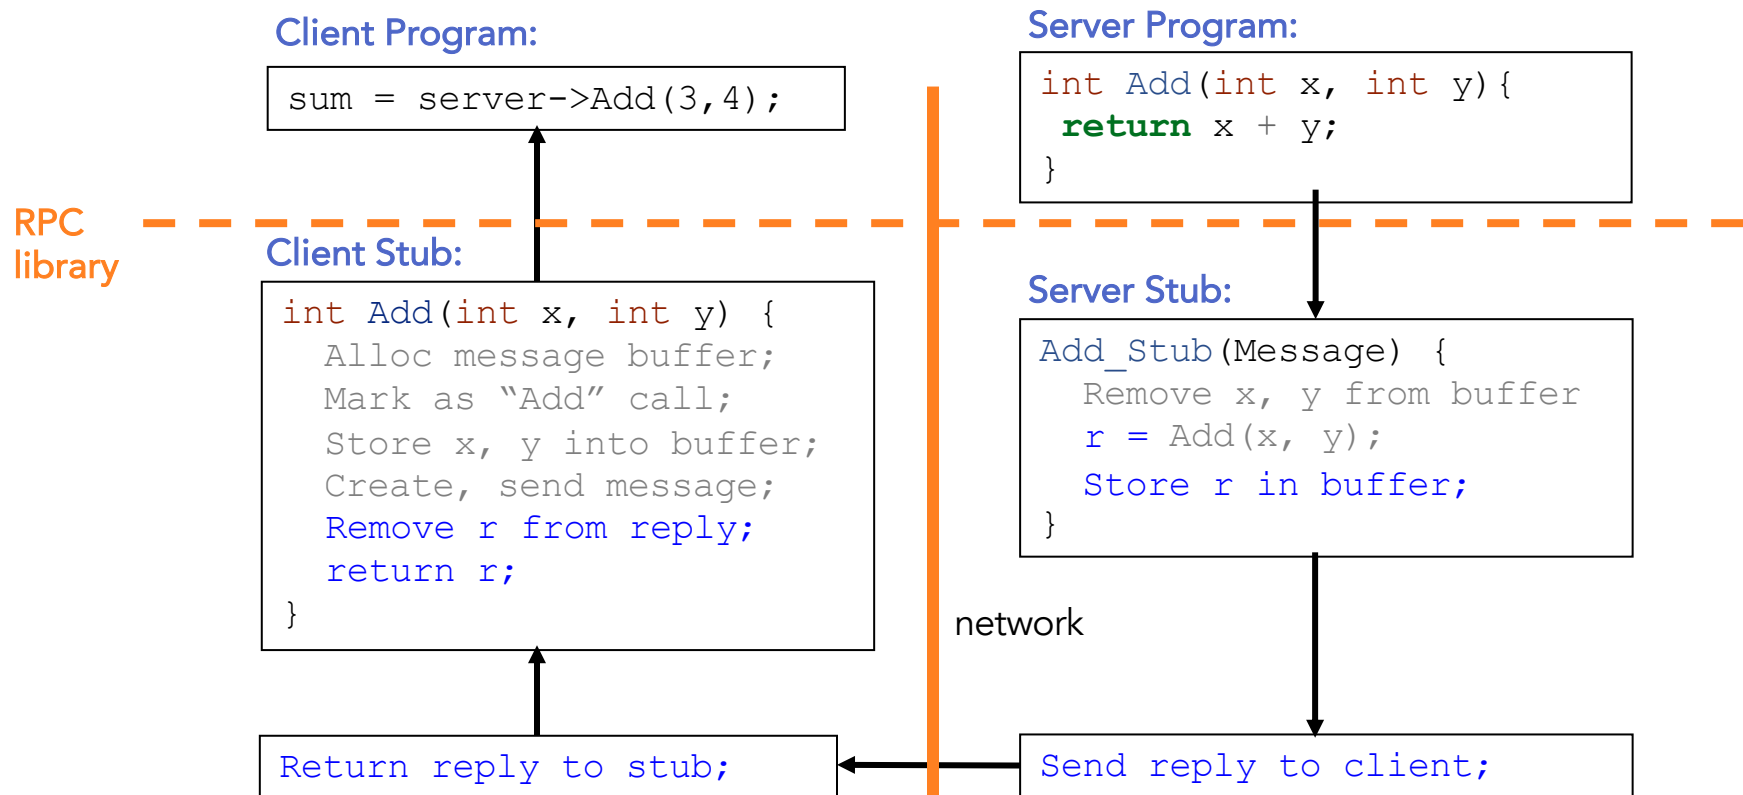
```
void produce_stub () {  
    int n;  
    int status;  
    → recv (sock, &n, sizeof(n)); // add MSG_WAITALL or loop  
    → status = produce(n);  
    → send (sock, &status, sizeof(status));  
}
```

→ invoke actual produce function in server

Another RPC example: call



Another RPC example: return



Generation of stubs

Stubs can be generated automatically

What do we need to know to do this?

Interface description

- Server defines an interface using an *interface definition language* (IDL)
- The IDL specifies the names, parameters, and return types for client-callable server procedures

Stub compiler (e.g., rpcgen on Linux) reads the IDL and produces the server and client stubs

RPC transparency

RPC aims to be as transparent as possible

- Make remote procedure calls look like local procedure calls

What factors might break the transparency?

- or make it difficult to achieve transparency

Pass a pointer to remote system?

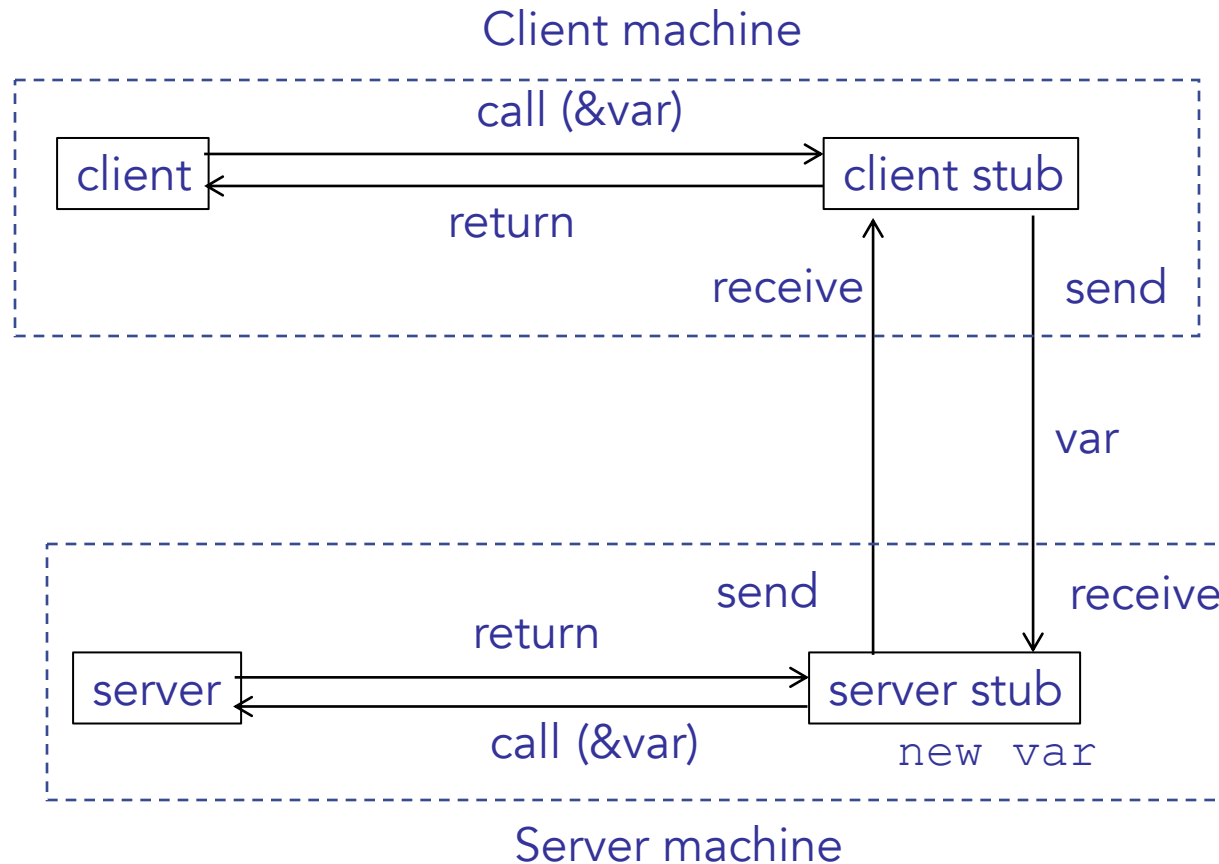
Client stub

- Client calls client stub with pointer
- Client stub sends ...

Server stub

- Server stub receives ...
- Server stub calls server with pointer

Pass a pointer to remote system?



Data representation

Client and server must agree on how to represent data

E.g., endianness of multi-byte integers

- In a 4-byte integer, which is the least significant byte?

Byte 3	DE
Byte 2	AD
Byte 1	BE
Byte 0	EF

Is this 0xDEADBEEF or
0xEFBEADDE?

E.g., character representation

E.g., C-string

RPC binding

Binding is the process of connecting client to server

- Static: fixed at compile time
- Dynamic: performed at runtime

The server exports its interface when it starts up

- Identifies itself to a network name server
- Tells RPC runtime that it's alive and ready to accept calls


The client imports the server before issuing any calls

- RPC runtime uses the name server to find the location of a server and establish a connection

The import and export operations are explicit in the server and client programs

RPC failures

RPC can have more complex failure modes than than local procedure calls



What does a failure look like to the client RPC library?

- Client never sees a response from the server
- Client does *not* know if the server saw the request
 - Maybe server/net failed just before sending reply

RPC failure semantics

Simplest scheme: **at-least-once** behavior

- RPC library waits for response for time T , if none arrives, re-send the request
- Repeat this a few times
- Still no response \rightarrow return an error to the application

Problems?

- E.g., request is "deduct \$100 from bank account"

When is at-least-once behavior *OK*?

- If it's ok to repeat an operation, e.g., `get(key)` ;
- If the application has its own way of dealing with duplicates

RPC failure semantics (2)

Another (better) RPC behavior: **at-most-once**

- Server RPC code detects duplicate requests returns previous reply instead of re-running handler
- How to detect a duplicate request?
 - client includes unique ID (XID) with each request and uses the same XID for re-send
 - server checks an incoming XID in a table, if an entry is found, directly returns the reply

RPC failure semantics (3)

What if an at-most-once server crashes and re-starts?

- If duplicate info is in memory, server will forget and accept duplicate requests after re-start
- It could write the duplicate info to disk
- What if the server fails to restart?
- Replica server could also replicate duplicate info

What about "exactly-once"?

- at-most-once plus unbounded retries plus fault-tolerant service

Project 4

Use assertions to catch errors early

- # of free disk blocks matches file system contents?
- Are you unlocking a lock that you hold?
- Verify initial file system is not malformed

Use showfs to verify that contents of file system match your expectations

Test cases: cover all states of data structures (e.g., direntries array)