

# REST APIs



Slides by Andrew DeOrio

# Agenda

- HTTP review
- REST API intro
- JSON
- Tools
- Collections and pagination
- Verbs and status codes
- Design principles

# Review: HTTP request methods

- Request method indicates server action
- GET: request a resource
  - Example: load a page
- HEAD: identical to GET, but without response body
  - Example: see if page has changed
- POST: send data to server
  - Example: web form
- Others we will cover later in the REST API lecture
  - TODAY

# Review: HTTP request headers

- Headers accompany request

- Most are optional

```
$ curl --verbose http://cse.eecs.umich.edu/ >  
index.html  
* Connected to cse.eecs.umich.edu  
(141.212.113.143) port 80 (#0)  
> GET / HTTP/1.0  
> Host: cse.eecs.umich.edu  
> User-Agent: curl/7.54.0  
> Accept: */*
```

- Host distinguishes between DNS names sharing a single IP address
  - Required as of HTTP/1.1
- User-Agent: which browser is making the request
- Accept: which content ("file") types the client will accept

# Review: HTTP status code

- Response starts with a status code
  - 1XX: Informational
  - 2XX: Successful
  - 3XX: Client Error
  - 4XX: Server Error
- ```
$ curl --verbose http://cse.eecs.umich.edu/  
> GET / HTTP/1.0  
< HTTP/1.0 200 OK
```
- ```
$ curl --verbose http://cse.eecs.umich.edu/asdf  
> GET /asdf HTTP/1.0  
< HTTP/1.0 404 Not Found
```

# Review: HTTP response headers

- Headers accompany a response

- Most are optional

```
$ curl --verbose http://cse.eecs.umich.edu/  
* Connected to cse.eecs.umich.edu  
> GET / HTTP/1.0  
...  
< HTTP/1.0 200 OK  
< Date: Tue, 12 Sep 2017 20:04:20 GMT  
< Server: Apache/2.2.15 (Red Hat)  
< Accept-Ranges: bytes  
< Connection: close  
< Transfer-Encoding: chunked  
< Content-Type: text/html; charset=UTF-8
```

# Review: HTTP content type

- Content type describes the "file" type and encoding

```
$ curl --verbose http://cse.eecs.umich.edu/  
* Connected to cse.eecs.umich.edu  
> GET / HTTP/1.0
```

```
...
```

```
< HTTP/1.0 200 OK
```

```
...
```

```
< Content-Type: text/html; charset=UTF-8
```

```
<!doctype html><html lang="en">
```

```
...
```

```
</html>
```

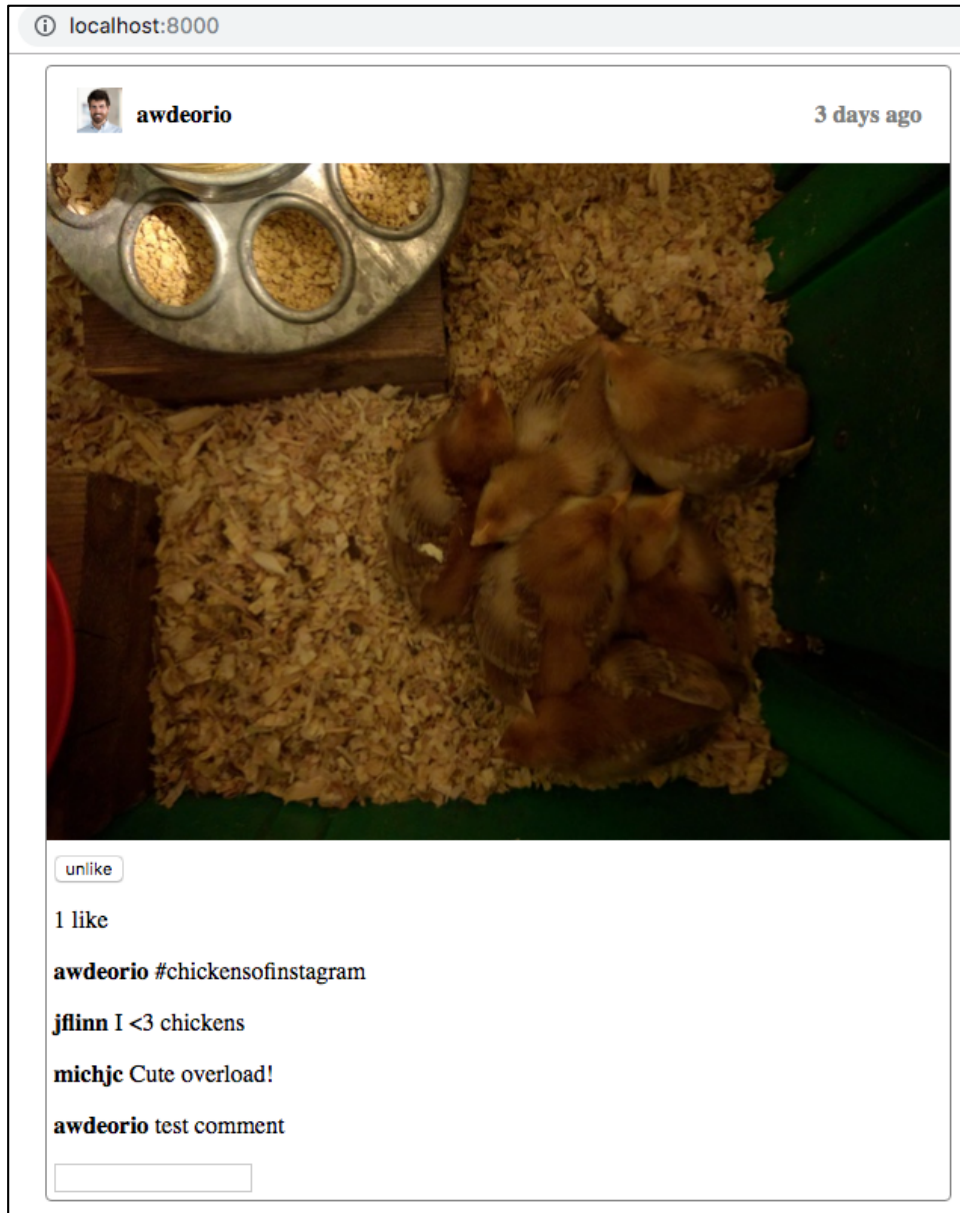
- In past lectures, GET requests have returned HTML content
- In today's lecture, we'll return JSON data

# Agenda

- HTTP review
- **REST API intro**
- JSON
- Tools
- Collections and pagination
- Verbs and status codes
- Design principles



# Not a REST API: human-readable



# REST API: machine-readable

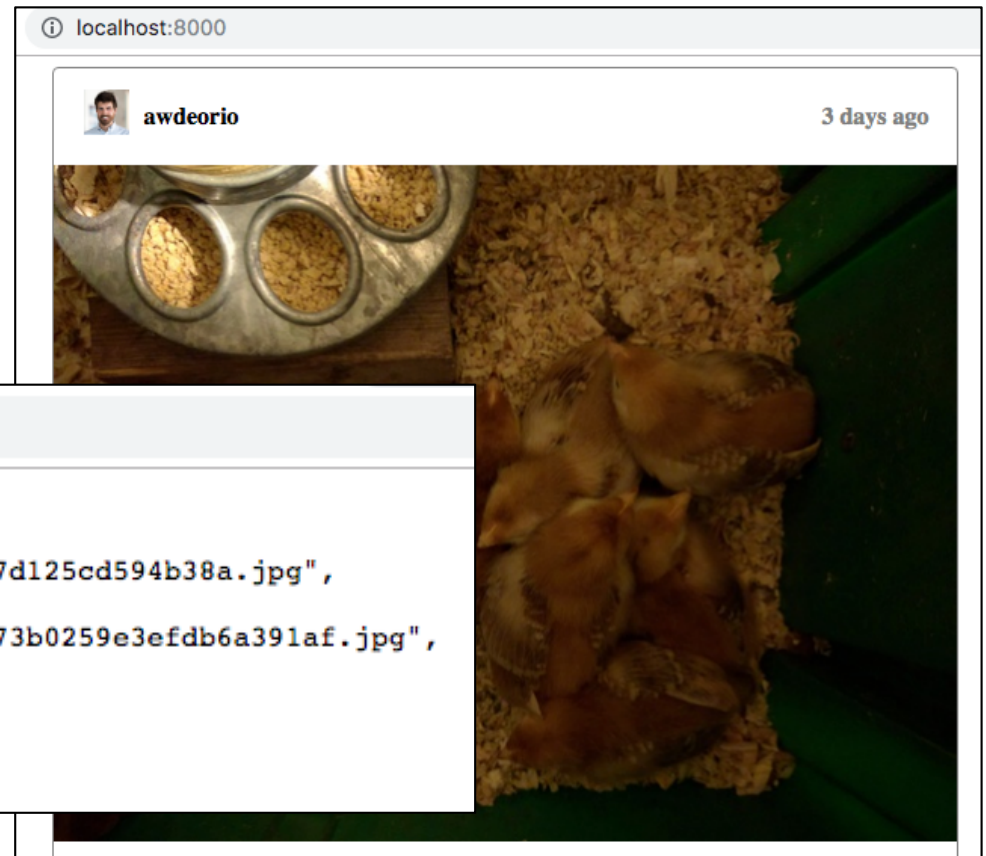


A screenshot of a web browser window. The address bar shows the URL `localhost:8000/api/v1/p/3/`. The main content area displays a JSON object representing a REST API response. The JSON is formatted with syntax highlighting, showing keys in quotes and values in quotes or as booleans. The response includes fields for age, image URL, owner, owner image URL, owner show URL, post show URL, and the full API URL.

```
{
  "age": "2019-09-20 17:28:59",
  "img_url": "/uploads/9887e06812ef434d291e4936417d125cd594b38a.jpg",
  "owner": "awdeorio",
  "owner_img_url": "/uploads/e1a7c5c32973862ee15173b0259e3efdb6a391af.jpg",
  "owner_show_url": "/u/awdeorio/",
  "post_show_url": "/p/3/",
  "url": "/api/v1/p/3/"
}
```

# REST API: machine-readable

- JavaScript runs in browser
- JavaScript receives JSON data from server
- JavaScript renders data on page
- Also: Two servers can communicate via JSON.



# REST APIs use HTTP

- HTTP request includes a method
- HTTP response includes a status code and JSON data

```
$ curl --verbose localhost:8000/api/v1/p/1/
> GET /api/v1/p/1/ HTTP/1.0
< HTTP/1.0 200 OK
< Content-Type: application/json
{
  "age": "2019-09-20 17:28:59",
  "img_url": "/uploads/122a7d27ca1d7420a1072f695d9290fad4501a41.jpg",
  "owner": "awdeorio",
  "owner_img_url": "/uploads/e1a7c5c32973862ee15173b0259e3efdb6a391af.jpg",
  ...
}
```

# REST APIs

- REST: Representational State Transfer
- Interoperability between different web systems
- REST is not ...
  - A standard
  - A language
- REST is ...
  - A collection of principles
  - Some best practices
  - Usually uses HTTP and JSON
- Originally defined by Roy Fielding in his doctoral dissertation

# Agenda

- HTTP review
- REST API intro
- **JSON**
- Tools
- Collections and pagination
- Verbs and status codes
- Design principles

# JSON

- JSON: JavaScript Object Notation
- Lightweight data-interchange format
- Based on JavaScript syntax
  - Uses conventions familiar to programmers in many languages
- Commonly used to send data from a server to a web client
  - Client parses JSON using JavaScript and displays content
- Ubiquitous with REST APIs

# JSON structures

- Object: a collection of name/value pairs
  - In other languages: object, record, struct, dictionary, hash table, keyed list, or associative array

```
{ "name": "DeOrion", "num_chicken": 4 }
```
- Array: an ordered list of values
  - In other languages: array, vector, list, or sequence

```
[ "Marilyn", "Maude", "Myrtle II", "Mabel"]
```
- A value is:
  - string
  - number
  - true
  - false
  - null
  - Object
  - Array



# Valid JSON

- Validate JSON

```
$ curl -s https://api.github.com/users/awdeorio | jsonlint
```

- Pitfall: **no trailing commas allowed!**

```
{  
  "login": "awdeorio",  
  "id": 7503005,  
  ...
```

```
  "updated_at": "2017-12-12T19:11:17Z" /
```

```
}
```

- More details: <http://www.json.org/>

# Agenda

- HTTP review
- REST API intro
- JSON
- **Tools**
- Collections and pagination
- Verbs and status codes
- Design principles

# curl

- REST API at the command line
- HTTP GET request returns a JSON-formatted string

```
$ curl https://api.github.com/users/awdeorio
{
  "login": "awdeorio",
  "id": 7503005,
  ...
```

# jq and python

- Pretty-print JSON using [jq](#)

```
$ curl -s https://api.github.com/users/awdeorio | jq
{
  "login": "awdeorio",
  "id": 7503005,
  ...
}
```

- Pretty-print JSON using Python

```
$ curl -S https://api.github.com/users/awdeorio | python -m json.tool
{
  "login": "awdeorio",
  "id": 7503005,
  ...
}
```

# Httpie

- Improved CLI and color coding with [httpie](#)

```
$ http https://api.github.com/users/awdeorio
HTTP/1.0 200 OK
{
  "login": "awdeorio",
  "id": 7503005,
  ...
}
```

# httpbin.org

- <https://httpbin.org> is an echo server
  - Responds with whatever you sent to it

```
$ http POST httpbin.org/anything hello=world
...
{
    ...
    "json": {
        "hello": "world"
    },
    "method": "POST",
    "url": "http://httpbin.org/anything"
}
```

# Agenda

- HTTP review
- REST API intro
- JSON
- Tools
- **Collections and pagination**
- Verbs and status codes
- Design principles

# Detail view AKA item view

- Our previous example showed a REST API *detail view* or *item view*
- It returns one object from the database
- Notice the the *id* part of the URL
  - Also called a "slug"

```
$ curl localhost:8000/api/v1/p/1/
{
  "age": "2019-09-20 17:28:59",
  "img_url": "/uploads/122a7d27ca1d7420a1072f695d9290fad4501a41.jpg",
  "owner": "awdeorio",
  "owner_img_url": "/uploads/e1a7c5c32973862ee15173b0259e3efdb6a391af.jpg",
  "...": "...",
  "url": "/api/v1/p/1/"
}
```



# List view AKA collection view

- REST APIs often expose collections of items

```
$ curl localhost:8000/api/v1/p/
```


```
{  
  "results": [  
    {  
      "postid": 3,  
      "url": "/api/v1/p/3/"  
    },  
    {  
      "postid": 2,  
      "url": "/api/v1/p/2/"  
    },  
    ...  
  ]  
}
```

# Pagination

- Pagination from the UI perspective
- REST API enables this
- Instagram et. al use REST API pagination for infinite scroll



...



# Pagination

- List views should return a limited number of items
  - What if there were 10 million posts?
- Sensible default, e.g., 10 posts
  - `$ curl localhost:8000/api/v1/p/`
- Get the next 10 results
  - `$ curl localhost:8000/api/v1/p/?page=1`
- Customizable size
  - `$ curl localhost:8000/api/v1/p/?size=20`

# Agenda

- HTTP review
- REST API intro
- JSON
- Tools
- Collections and pagination
- **Verbs and status codes**
- Design principles

# REST API verbs and status codes

- All our examples so far have been GET requests
  - GET request
  - 200 OK response
- This takes care of reading data. What about create, modify and delete?

# REST API verbs

- GET: return datum
  - Example: return a post
- POST: create new datum
  - Example: create a new post
- PATCH: update part of a datum
  - Example: modify part of an existing post
- PUT: replace the entire datum
  - Example: replace an existing post
- DELETE: delete datum
  - Example: remove a post

# POST request

- POST creates an object
- Request includes JSON body

```
POST localhost:8000/api/v1/p/ HTTP/1.0
{
  "img_url": "122a7d27ca1d7420a1072f695d9290fad4501a41.jpg",
  "owner": "awdeorio",
  ...
}
```

# POST response

- POST returns 201 CREATED on success
- Response includes a copy of the created object
  - Object usually includes a link to itself

```
POST localhost:8000/api/v1/p/ HTTP/1.0
```

```
...
```

```
HTTP/1.0 201 CREATED
```

```
{  
  "img_url": "122a7d27ca1d7420a1072f695d9290fad4501a41.jpg",  
  "owner": "awdeorio",  
  ...  
  "url": "/api/v1/p/1/"  
}
```



# PATCH request

- PATCH modifies part of an existing object
- Request URL includes an ID
- Request includes JSON body
- Example: change the picture in a post
- Notice that the JSON body is short, and only contains the field that should be modified

```
PATCH localhost:8000/api/v1/p/1/ HTTP/1.0
{
  "img_url": "ad7790405c539894d25ab8dcf0b79eed3341e109.jpg",
}
```

# PATCH response

- PATCH returns 200 OK on success
- Response includes a copy of the **entire** modified object

```
PATCH localhost:8000/api/v1/p/1/ HTTP/1.0
```

```
...
```

```
HTTP/1.0 200 OK
```

```
{  
  "img_url": "ad7790405c539894d25ab8dcf0b79eed3341e109.jpg",  
  "owner": "awdeorio",  
  ...  
  "url": "/api/v1/p/1/"  
}
```

# PUT request

- PUT replaces an entire existing object
- Request URL includes an ID
- Request includes JSON body
- Example: replace an entire post
- The JSON body is long, and contains a replacement value for every field

```
PUT localhost:8000/api/v1/p/1/ HTTP/1.0
{
  "img_url": ...,
  "owner": "jflinn",
  "owner_img_url": ...,
  ...
}
```

# PUT response

- PUT returns 200 OK on success
- Response includes a copy of the **entire** modified object

```
PUT localhost:8000/api/v1/p/1/ HTTP/1.0
```

```
...
```

```
HTTP/1.0 200 OK
```

```
{
```

```
  "img_url": ...,
```

```
  "owner": "jflinn",
```

```
  "owner_img_url": ...,
```

```
...
```

```
}
```

# DELETE request

- DELETE removes an object
- Request URL includes an ID
- No body in request

```
DELETE localhost:8000/api/v1/p/1/ HTTP/1.0
```

# DELETE response

- DELETE **returns** 204 NO CONTENT **on success**
- No body in response

```
DELETE localhost:8000/api/v1/p/1/ HTTP/1.0
```

```
...
```

```
HTTP/1.0 204 NO CONTENT
```

# Not found response

- GET a deleted item, receive a 404 response

```
DELETE localhost:8000/api/v1/p/1/ HTTP/1.0  
HTTP/1.0 204 NO CONTENT
```

```
GET localhost:8000/api/v1/p/1/ HTTP/1.0  
HTTP/1.0 404 NOT FOUND
```

# REST API status codes

- 200 OK
- 201 Created
  - Successful creation after POST
- 204 No Content
  - Successful DELETE
- 304 Not Modified
  - Used for conditional GET calls to reduce band-width usage
- 400 Bad Request
  - General error
- 401 Unauthorized
  - Missing or invalid authentication
- 403 Forbidden
  - User is not authorized
- 404 Not Found
  - Resource could not be found
- 409 Conflict
  - E.g., duplicate entries and deleting root objects when cascade-delete is not supported
- 500 Internal Server Error
  - General catch-all for server-side exceptions

<http://www.restapitutorial.com/httpstatuscodes.html>



# Agenda

- HTTP review
- REST API intro
- JSON
- Tools
- Collections and pagination
- Verbs and status codes
- **Design principles**

# REST design principles

- Uniform interface
  - Resource-based
  - Manipulation of resources through representations
  - Self-descriptive messages
  - Hypermedia as the engine of application state (HATEOAS)
- Client-server architecture
- Stateless
- Cacheable
- Layered
- Code on demand (optional)

# Uniform interface: resource-based

- Individual resources are identified in requests using URIs as resource identifiers.
- Think of a URI like a pointer
- ID in the URL

```
GET /api/v1/p/1/ HTTP/1.0
```

# Uniform interface: manipulation of resources through representations

- When a client holds a representation of a resource, including any metadata attached, it has enough information to modify or delete the resource on the server.
- Object usually contains a link to itself

```
GET /api/v1/p/1/ HTTP/1.0
```

```
HTTP/1.0 200 OK
```

```
{
```

```
    ...'
```

```
    "url": "/api/v1/p/1/"
```

```
}
```

# Uniform interface: self-descriptive messages

- Each message includes enough information to describe how to process the message
- Content-Type and charset

```
GET /api/v1/p/1/ HTTP/1.0
```

```
HTTP/1.0 200 OK
```

```
> Content-Type:application/json; charset=utf-8
```

```
...
```

```
{
```

```
...
```

```
}
```

# Uniform interface: HATEOAS

- HATEOAS: Hypermedia as the Engine of Application State
- Clients deliver state via body contents, query-string parameters, request headers and the requested URI (the resource name).
- Services deliver state to clients via body content, response codes, and response headers. This is technically referred-to as hypermedia (or hyperlinks within hypertext).
- Everything you need is in the request

```
POST /api/v1/p/ HTTP/1.0
{
  "img_url": "122a7d27ca1d7420a1072f695d9290fad4501a41.jpg",
  "owner": "awdeorio",
  ...
}
```

# Uniform interface: HATEOAS

- Links are contained in the returned body (or headers) to supply the URI for retrieval of the object itself or related objects.

```
GET /api/v1/p/ HTTP/1.0
```

```
HTTP/1.0 200 OK
```

```
{  
  "results": [  
    {  
      "postid": 3,  
      "url": "/api/v1/p/3/"  
    },  
    {  
      "postid": 2,  
      "url": "/api/v1/p/2/"  
    },  
    ...  
  ]  
}
```

# Client-server architecture

- The uniform interface separates clients from servers
- *Abstraction* between client and server
- Can change the server without modifying the client
- Can change the client without modifying the server
- Example: database is too slow, replace portions of it with key-value store like Redis



# Stateless

- Everything needed to handle the request is in the request itself
  - URI, query-string parameters, body, or headers
- After the server does it's processing, the appropriate state, or the piece(s) of state that matter, are communicated back to the client via headers, status and response body

```
POST /api/v1/p/ HTTP/1.0
```

```
HTTP/1.0 201 Created
{
  // contents of created object
}
```

# Cacheable

- Clients cache some kinds of responses to eliminate requests
  - Example: cache an image so you don't load it every time
- Responses must implicitly or explicitly define themselves as cacheable
- Example: Last-Modified header

# Layered

- A client cannot ordinarily tell whether it is connected directly to the end server, or to an intermediary along the way
- No need to connect to a specific machine
- Just need the data from this URI

## Code on demand (optional)

- Servers are able to temporarily extend or customize the functionality of a client by transferring logic to it that it can execute
- Examples of this may include compiled components such as Java applets and client-side scripts such as JavaScript
- Not common

# Summary

- A client and a server can communicate via a REST API
- Two servers can communicate via a REST API
- REST APIs use HTTP
- REST APIs are machine-readable
- REST APIs usually return JSON data

# Public APIs

- GitHub  
<https://developer.github.com/v3/>
- LinkedIn  
<https://developer.linkedin.com/>
- Facebook  
<https://developers.facebook.com/docs/graph-api>
- Twitter  
<https://dev.twitter.com/rest/public>