

EECS 489

Computer Networks

Winter 2025

Mosharaf Chowdhury

Material with thanks to Aditya Akella, Sugih Jamin, Philip Levis, Sylvia Ratnasamy, Peter Steenkiste, and many other colleagues.

Recap: Inserting RR into DNS

- ❑ Registrar inserts RR pairs into the .com TLD server:
 - (foobar.com, dns1.foobar.com, NS)
 - (dns1.foobar.com, 212.44.9.129, A)

- ❑ Human readability
 - Name of the nameserver doesn't change when it's IP changes

- ❑ Load balancing
 - + (dns1.foobar.com, 212.44.9.128, A)
 - Pick one at random for load balancing

- ❑ Support for IPv6
 - + (dns1.foobar.com, 2001:db8::1, AAAA)
 - Pick between IPv4 and IPv6

Agenda

- Video streaming
- Datacenter applications

Why is video important?

- Dominates the global Internet traffic landscape
 - About 65%, i.e., every 2 of 3 bytes in 2022!
- Major sources
 - Netflix
 - YouTube
 - TikTok
 - ...

The video medium

- ❑ Video is a sequence of images/frames displayed at a constant rate (moving pictures)
- ❑ Digital image is an array of pixels, each pixel represented by bits
- ❑ Examples:
 - Single frame image encoding: 1024x1024 pixels, 24 bits/pixel \Rightarrow 3 MB/image
 - Movies: 24 frames/sec \Rightarrow 72 MB/sec
 - TV: 30 frames/sec \Rightarrow 90 MB/sec

The video medium (cont'd)

- ❑ Compression is key
 - Lots of algorithms to compress
- ❑ The same video can be (and typically is) compressed to multiple quality levels
 - E.g., 480p, 720p, 1080p, 4K
- ❑ Why multiple resolutions?
 - Adapt to user network conditions

How to watch a video?

1. Download and watch

- Often **too large** to send in one GET
- Doesn't even make sense even if its possible
 - » Users must **wait too long**
 - » Users **may skip forward!** ⇒ bandwidth waste
 - » User's **connection quality may change** (e.g., switching from WiFi to LTE) ⇒ lower resolution to match bandwidth

❓ Our focus is *not* live video

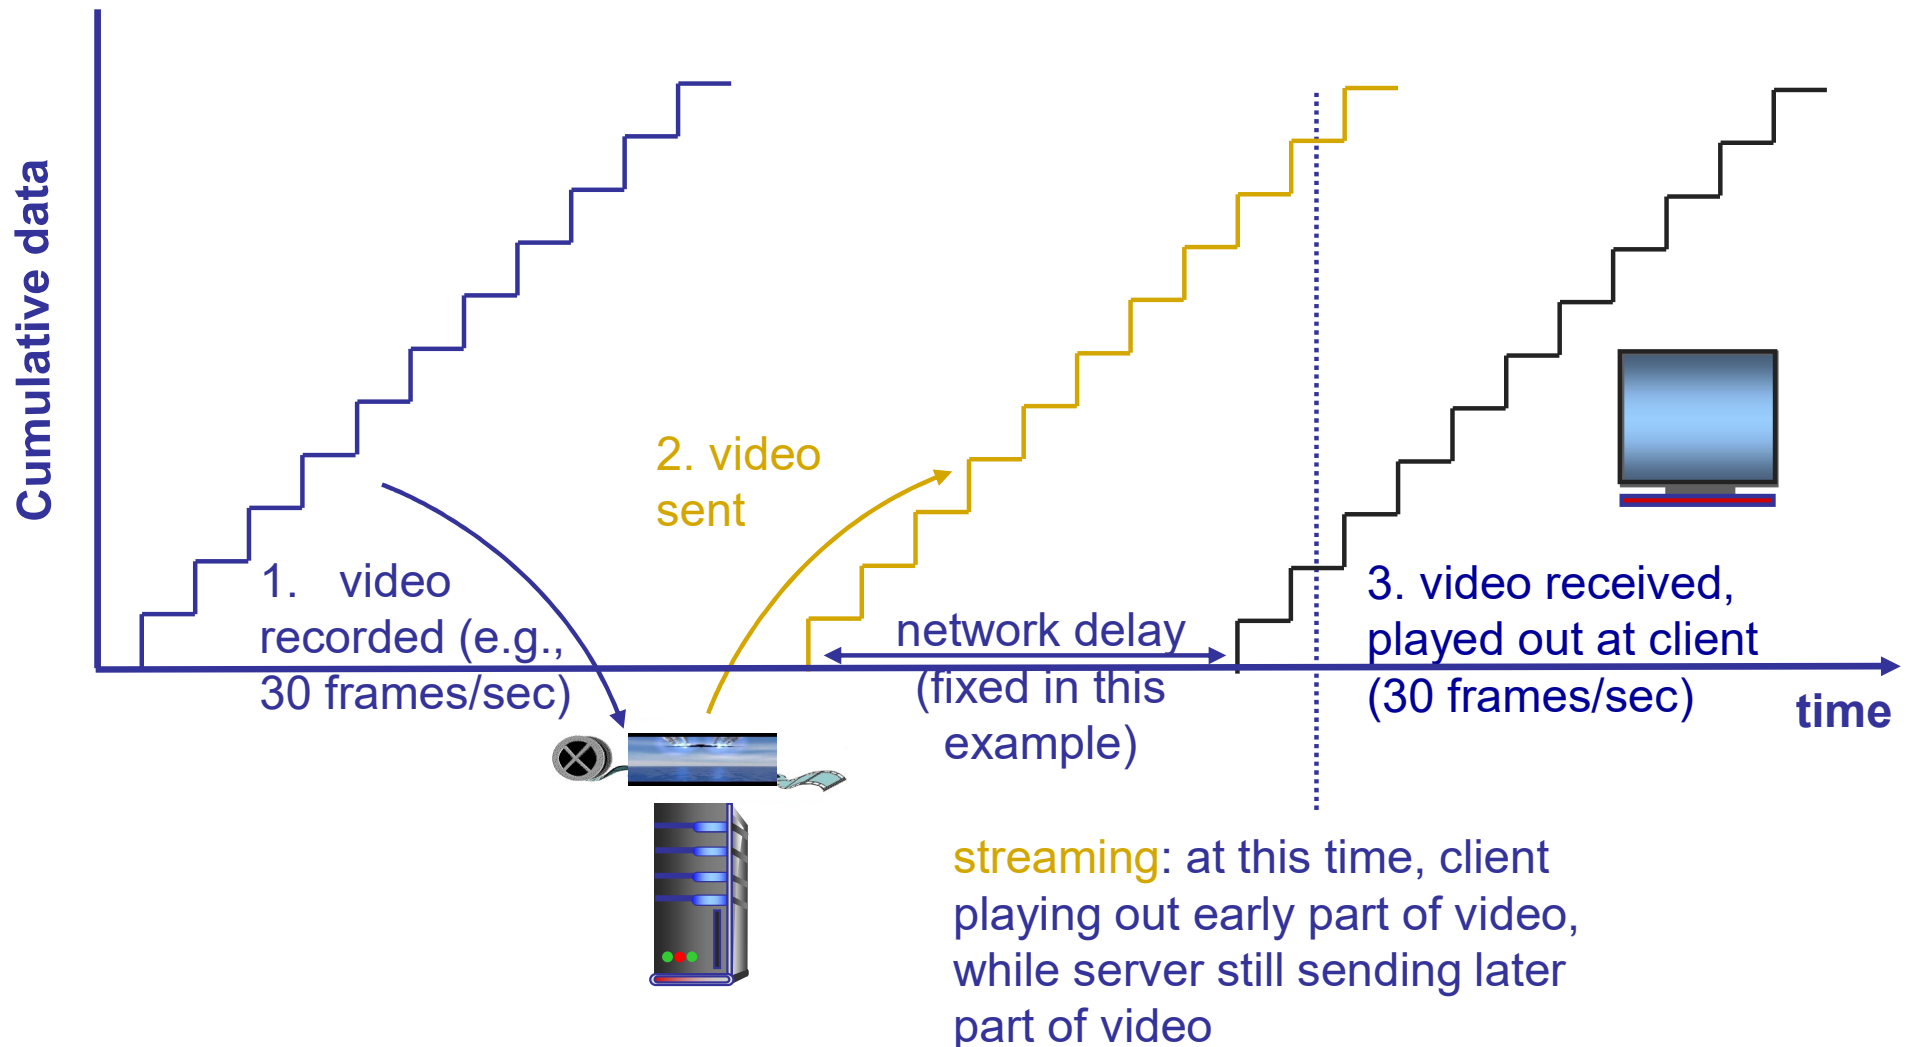
How to watch a video?

2. Video streaming over HTTP

HTTP streaming

- Video is stored at an HTTP server with a URL
- Clients send a GET request for the URL
- Server sends the video file as a stream

HTTP streaming



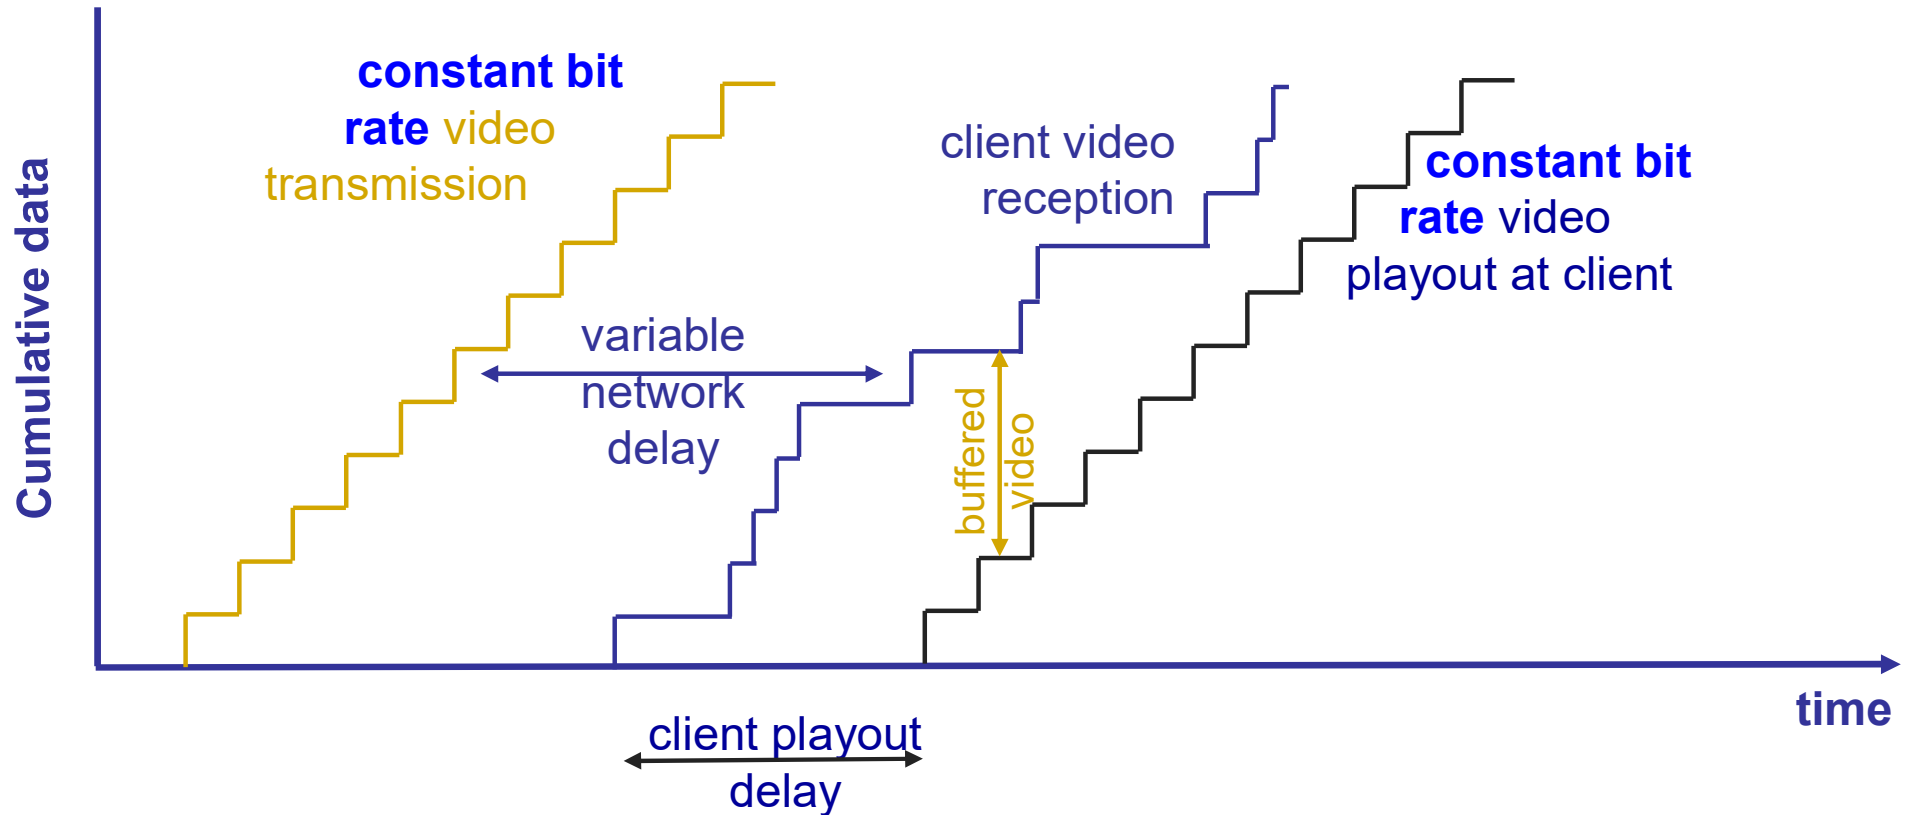
What if the latency is high?

- ❑ Client first buffers for a while
 - To minimize interruptions later
- ❑ Once the buffer reaches a threshold
 - The video plays in the **foreground**
 - More frames are downloaded in the **background**

Challenges

- ❑ What if network delay is not fixed?
 - i.e., How to absorb delay variations?
- ❑ What if users jump forward, fast-forward, rewind, pause?
 - i.e., How to handle user interactions?
- ❑ Handle packet loss, retransmission etc.

HTTP streaming: Revisited



- **Client-side buffering and playout delay:** compensate for network-added delay, delay jitter

Issues with HTTP streaming


- ❑ Same bitrate for all clients
 - Clients can have very different network conditions
 - Clients network conditions can change over time
- ❑ Cannot dynamically adapt to conditions

DASH : Dynamic Adaptive Streaming over HTTP


- ❑ Keep **multiple resolutions** of the same video
 - Stored in a manifest file in the HTTP server
- ❑ Client asks for the manifest file first to learn about the options
- ❑ Asks for chunks at a time and measures available bandwidth while they are downloaded
 - Low bandwidth \Rightarrow switch to lower bitrate
 - High bandwidth \Rightarrow switch to higher bitrate
- ❑ **Adaptive bit rate (ABR)**

CLOUD SYSTEMS

Who's serving Web services?



Who's serving Web services?




Cloud datacenters run the world




Source: Google

Cloud datacenters run the world



Cloud datacenters run the world



Source: Facebook

How big is a datacenter (DC)?

- ❑ 1M servers/site [Microsoft/Amazon/Google]
- ❑ >\$1B to build one site [UMich Datacenter!]
- ❑ >\$20M/month/site operational costs [MS'09]
- ❑ Data center hardware spending grew to a record \$282 billion in 2024. [Synergy Research Group report]
- ❑ But only $O(100)$ sites

Implications (1)

□ Scale

- Need scalable designs
- Low-cost designs: e.g., use commodity technology
- High utilization (efficiency): e.g., >60% avg. utilization
 - » **Contrast:** avg. utilization on Internet links often ~30%
- Tolerate frequent failure
 - » Large number of (low cost) components
- Automate

Implications (2)


- Service model: clouds / multi-tenancy
 - Performance guarantees
 - Isolation guarantees
 - Portability

Applications

- ❑ Common theme: **parallelism**
 - Applications decomposed into tasks
 - Running in parallel on different machines
- ❑ Two common paradigms
 - Partition-Aggregate
 - Map-Reduce

Partition-Aggregate

eeecs 489



Partition-Aggregate

eeecs 489



GitHub - mosharaf/eeecs489: EECS 489: Computer Networks @ the ...

<https://github.com/mosharaf/eeecs489> ▼

EECS 489: Computer Networks (F'18) ... **EECS 489** takes a top-down approach to explore how networks operate and how network applications are written. ... Kurose and Ross, Computer Networking: A Top-Down Approach, 7th.

UM EECS 489: Computer Networks

www.eecs.umich.edu/courses/eeecs489/ ▼

Lecture: MWF 9:30 - 10:30 in 1500 **EECS**. Discussion/Lab: W 12:30 - 1:30 in 2166 DOW or W 4:30 - 5:30 in 1014 DOW. The discussion sessions will mostly be ...

EECS 489

www.eecs.umich.edu/courses/eeecs489/f99/ ▼

News group umich.eecs.class.489. Everything posted here will be automatically forwarded to the **eeecs489staff@eecs.umich.edu** mailing list hourly.

EECS 489 - EECS @ Michigan - University of Michigan

<https://www.eecs.umich.edu/eeecs/academics/courses/eeecs-489.html> ▼

Course Homepage: <http://www.eecs.umich.edu/courses/eeecs489/w10/>. Coverage We study how networks operate and how network applications are written.

Revamping EECS 489: A Retrospective | Mosharaf Chowdhury

<https://www.mosharaf.com/blog/2017/05/07/revamping-eeecs-489-a-retrospective/> ▼

May 7, 2017 - A couple of weeks ago, we wrapped up the Spring 2017 offering of the **EECS 489**:

End-to-end response time

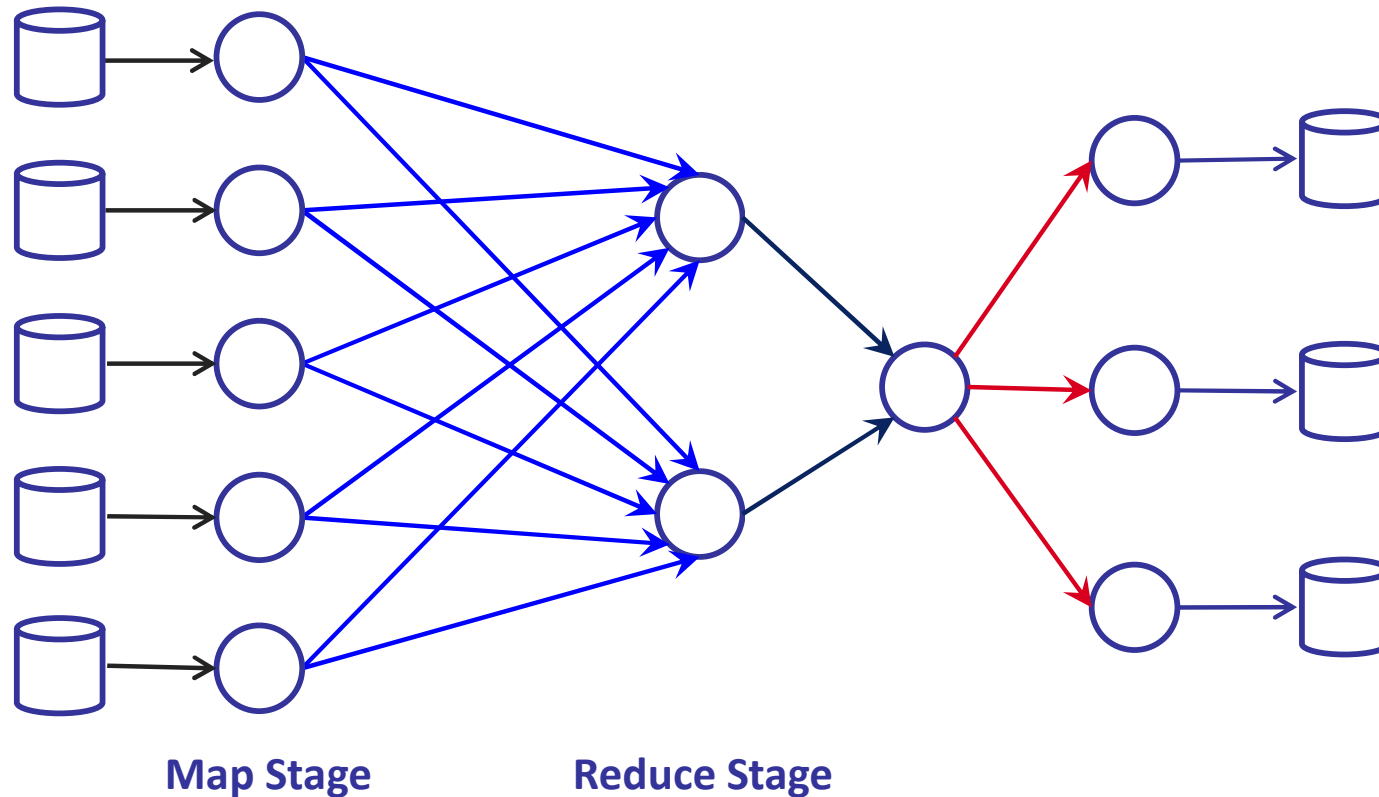
- **Less than 200 milliseconds** between receiving user query in the browser and displaying the results
 - RTT = $O(10)$ to 100 milliseconds
 - What remains?
- Next time, when the page is not loading fast enough, think about the poor servers working for you 😊

5-MINUTE BREAK!

Applications


- Common theme: parallelism
 - Applications decomposed into tasks
 - Running in parallel on different machines
- Two common paradigms
 - Partition-Aggregate
 - Map-Reduce

Map-Reduce




The most popular software that follows this paradigm is **Apache Spark**


Datacenter networks



Datacenter networks (Cont.)



Datacenter traffic



East-West traffic

- ❑ Traffic **between servers** in the datacenter
- ❑ Communication within “big data” computations
- ❑ Traffic may shift on small timescales (< minutes)


Datacenter traffic characteristics

- ❑ Two key characteristics
 - Most flows are small
 - Most bytes come from large flows
- ❑ Applications want
 - High bandwidth (large flows)
 - Low latency (small flows)


High bandwidth

- ▣ Ideal: Each server can talk to any other server at its full access link rate
- ▣ Conceptually: Datacenter network as one giant switch


Datacenter network as one giant switch



Datacenter network as one giant switch



Datacenter network as one giant switch



High bandwidth

- ❑ Ideal: Each server can talk to any other server at its full access link rate
- ❑ Conceptually: Datacenter network as one giant switch
 - Would require a 10 Pbits/sec switch!
 - » 1M ports (one port/server)
 - » 10Gbps per port
- ❑ **Practical approach:** build a network of switches (“fabric”) with high “bisection bandwidth”
 - Each switch has practical #ports and link speeds

Bisection bandwidth

- Partition a network into two equal parts
- Minimum bandwidth between the partitions is the bisection bandwidth
- **Full bisection bandwidth**: bisection bandwidth in an N node network is $N/2$ times the bandwidth of a single link
 - Nodes of any two halves can communicate at full speed with each other

Achieving full bisection bandwidth

Scale up


- Make links fatter toward the core of the network

Problem: Scaling up a traditional tree topology is expensive!


- Requires non-commodity / impractical / link and switch components

Solutions?

- **Over-subscribe** (i.e., provision less than full BBW)
- **Better topologies**



Oversubscription




Need techniques to **avoid congesting oversubscribed links!**


Oversubscription

- ❑ Not enough bandwidth
 - **Oversubscription**: Less bandwidth in the ToR-Agg links than all the servers' bandwidth in the rack
 - **Oversubscription ratio**: Ratio between bandwidth underneath and bandwidth above
- ❑ Not enough paths between server pairs
 - Load balancing issues
 - Failure recovery issues

Better topologies




Better topologies




Clos topology

- ❑ Multi-stage network
- ❑ k pods, where each pod has two layers of $k/2$ switches
 - $k/2$ ports up and $k/2$ down
- ❑ All links have the same b/w
- ❑ At most $k^3/4$ machines
- ❑ Example
 - $k = 4$
 - 16 machines
- ❑ For $k=48$, 27648 machines



Challenges in scale-out designs?

- Topology offers high bisection bandwidth
- All other system components must be able to exploit this available capacity
 - Routing must use all paths
 - Transport protocol must fill all pipes (fast)



Summary

❑ Video streaming

- Too large to send, so stream it
- Dynamically adapt to the network and users

❑ Cloud systems

- Forms the backend of modern web services
- Runs in datacenters where **all** the processing happens