

EECS 489

Computer Networks

Winter 2025

Mosharaf Chowdhury

Material with thanks to Aditya Akella, Sugih Jamin, Philip Levis, Sylvia Ratnasamy, Peter Steenkiste, and many other colleagues.

Agenda

- From reliable data transfer to TCP
- TCP connection setup and teardown

Recap: Designing a reliable transport protocol

- Stop and Wait vs Sliding Window
- Sliding Window
 - Acknowledgements: Cumulative vs Selective
 - Resending packets: Go-Back-N vs Selective Repeat

TCP: TRANSMISSION CONTROL PROTOCOL



The TCP Abstraction

- ❑ TCP delivers a reliable, in-order, byte stream
- ❑ **Reliable**: TCP resends lost packets (recursively)
 - Until it gives up and shuts down connection
- ❑ **In-order**: TCP only hands consecutive chunks of data to application
- ❑ **Byte stream**: TCP assumes there is an incoming stream of data, and attempts to deliver it to app

What does TCP use from what we've seen so far?

- Most of what we've seen
 - Checksums
 - Sequence numbers are byte offsets
 - Sender and receiver maintain a sliding window
 - Receiver sends cumulative acknowledgements (like GBN)
 - » Sender maintains a single retransmission timer
 - Receivers buffer out-of-sequence packets (like SR)
- Few more: fast retransmit, timeout estimation algorithms etc.


Build the TCP header



Build the TCP header



Computed
over pseudo-header
and data





What does TCP do?

- Most of what we've seen
 - Checksum
 - Sequence numbers are byte offsets


Build the TCP header

Byte offsets
(NOT packet id),
because TCP is a
byte stream



TCP “stream of bytes” service...

Application @ Host A




Application @ Host B

... provided using TCP

“segments”

Host A




TCP Data

Segment sent when:


1. Segment full (Max Segment Size),
2. Not full, but times out

TCP Data

Host B



TCP segment



? IP packet

- No bigger than **Maximum Transmission Unit (MTU)**
- E.g., up to 1500 bytes with Ethernet


? TCP packet

- IP packet with a TCP header and data inside
- TCP header ≥ 20 bytes long

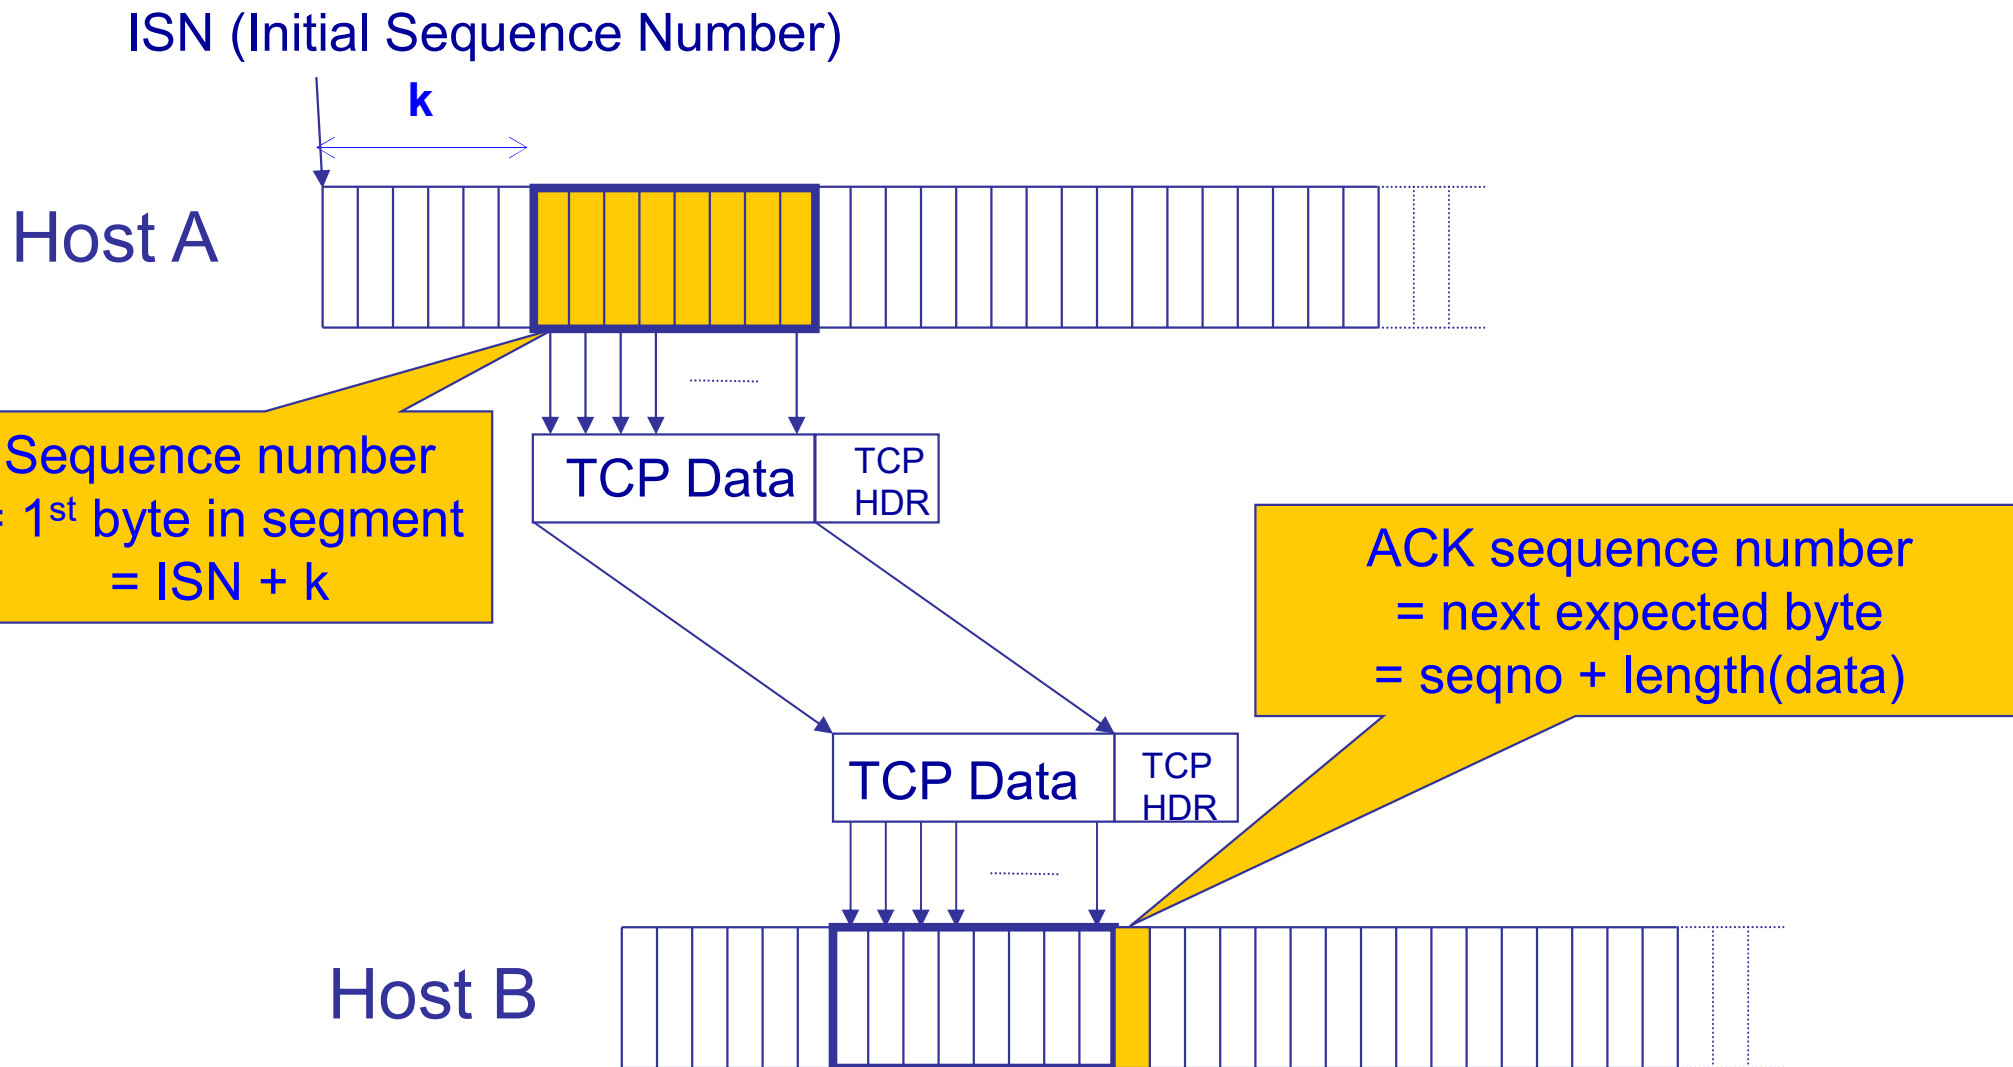
? TCP segment

- No more than **Maximum Segment Size (MSS)** bytes
- E.g., up to 1460 consecutive bytes from the stream
- $MSS = MTU - (IP \text{ header}) - (TCP \text{ header})$

Sequence numbers





Sequence numbers



Build the TCP header

Starting byte
offset of data
carried in this
segment



What does TCP do?

- Most of what we've seen
 - Checksum
 - Sequence numbers are byte offsets
 - Receiver sends **cumulative acknowledgements** (like GBN)

ACKs and sequence numbers

- Sender sends packet
 - Data starts with sequence number X
 - Packet contains B bytes $[X, X+1, X+2, \dots, X+B-1]$
- Upon receipt of packet, receiver sends an ACK
 - If all data prior to X already received:
 - » ACK acknowledges $X+B$ (because that is next expected byte)
 - If highest in-order byte received is Y s.t. $(Y+1) < X$
 - » ACK acknowledges $Y+1$
 - » Even if this has been ACKed before


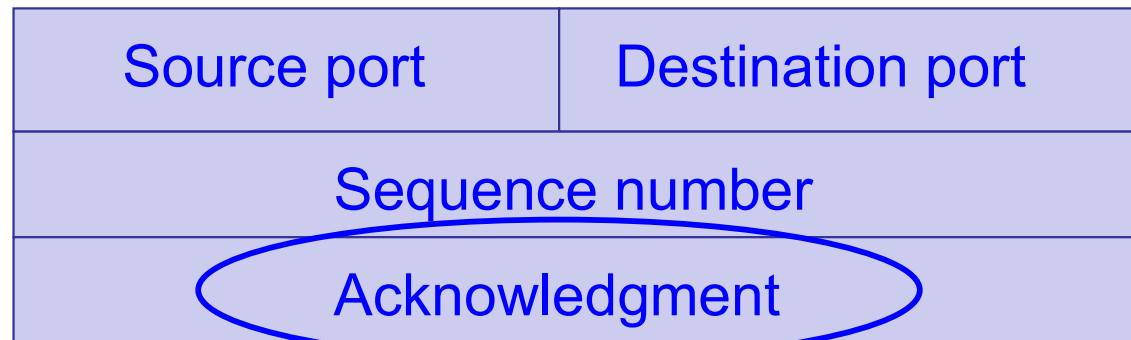
Typical operation

- | Sender: $\text{seqno} = X$, $\text{length} = B$
- | Receiver: $\text{ACK} = X + B$
- | Sender: $\text{seqno} = X + B$, $\text{length} = B$
- | Receiver: $\text{ACK} = X + 2B$
- | Sender: $\text{seqno} = X + 2B$, $\text{length} = B$

- | Seqno of next packet is same as last ACK field

Build the TCP header

Acknowledgment
gives seqno just
beyond highest
seqno received
in order



What does TCP do?

- | Most of what we've seen
 - Checksum
 - Sequence numbers are byte offsets
 - Receiver sends cumulative acknowledgements (like GBN)
 - Receivers **can buffer out-of-sequence packets** (like SR)

Loss with cumulative ACKs

- | Sender sends packets with 100B and seqnos.:
 - 100, 200, 300, 400, 500, 600, 700, 800, ...
- | Assume the fifth packet (seqno 500) is lost, but no others
- | Stream of ACKs will be:
 - 200, 300, 400, 500, 500, 500, 500, ...

What does TCP introduce?

- | Most of what we've seen
 - Checksum
 - Sequence numbers are byte offsets
 - Receiver sends cumulative acknowledgements (like GBN)
 - Receivers can buffer out-of-sequence packets (like SR)
- | Introduces **fast retransmit**: duplicate ACKs trigger early retransmission

Loss with cumulative ACKs

- | Duplicate ACKs are a sign of an isolated loss
 - The lack of ACK progress means 500 hasn't been delivered
 - Stream of ACKs means some packets are being delivered
- | Trigger retransmission upon receiving k duplicate ACKs
 - » TCP uses $k=3$
 - » Faster than waiting for timeout

Loss with cumulative ACKs

- | Two choices after resending
 - Send missing packet and move sliding window by the number of dup ACKs
 - »Speeds up transmission, but might be wrong
 - Send missing packet, and wait for ACK to move sliding window
 - »Is slowed down by single dropped packets

- | Which should TCP do?
 - Choose correctness

5-MINUTE BREAK!

Announcements

- | A2 is out!
 - Due on Feb 28



What does TCP introduce?

- | Most of what we've seen
 - Checksum
 - Sequence numbers are byte offsets
 - Receiver sends cumulative acknowledgements (like GBN)
 - Receivers buffer out-of-sequence packets (like SR)
- | Introduces fast retransmit: duplicate ACKs trigger early retransmission
- | Sender maintains a **single retransmission timer** (like GBN) and retransmits on timeout

Retransmission timeout

- | If the sender hasn't received an ACK by timeout, **retransmit the first packet** in the window
- | Challenge: **How do we pick a timeout value?**

Timing illustration




Retransmission timeout

- | If the sender hasn't received an ACK by timeout, retransmit the first packet in the window
- | How to set timeout?
 - Too long: connection has low throughput
 - Too short: retransmit packet that was just delayed
- | Solution: make timeout proportional to RTT
 - But how do we measure RTT?

RTT estimation



- Exponential weighted average of RTT samples

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$



Problem: Ambiguous measurements

- How do we differentiate between the real ACK, and ACK of the retransmitted packet?



Solution: Ambiguous measurements

- | Don't use SampleRTT from retransmissions
 - Once retransmitted, ignore that segment in the future

Karn/Partridge algorithm

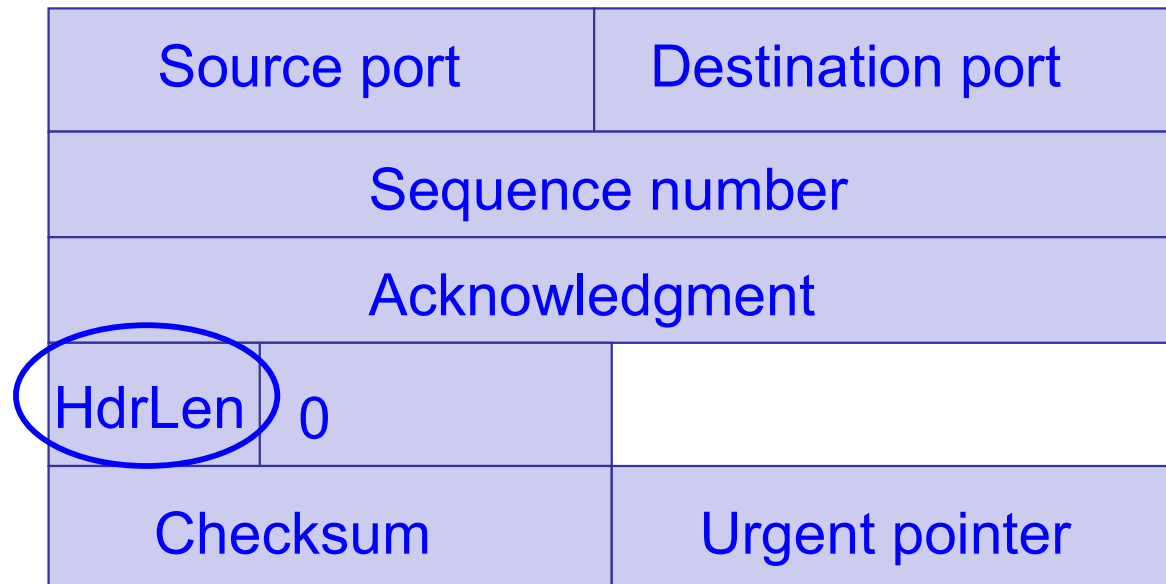
- | Computes EstimatedRTT using $\alpha = 0.125$
- | Timeout value (RTO) = $2 \times \text{EstimatedRTT}$
 - Every time RTO timer expires, set $\text{RTO} \leftarrow 2 \cdot \text{RTO}$
 - » Employs exponential backoff (Up to ≥ 60 sec)
 - Every time new measurement comes in (i.e., successful original transmission), reset RTO back to $2 \times \text{EstimatedRTT}$
- | Sensitive to RTT variations

Jacobson/Karels algorithm

- | **Problem:** need to better capture variability in RTT
 - **Solution:** Directly measure deviation
- | Deviation = | SampleRTT – EstimatedRTT |
- | DevRTT: exponential average of Deviation
- | **RTO = EstimatedRTT + 4 x DevRTT**

Build the TCP header

Number of 4-
byte words in the
header;
5: No options



Data


TCP CONNECTION ESTABLISHMENT

Initial Sequence Number (ISN)

- | Why not just use $ISN = 0$?
 - IP addresses and port #s uniquely identify a connection; however, these port #s get used again
 - » Small chance an old packet is still in flight
 - Predictable; makes spoofing connection easier
- | Hosts exchange ISNs when establishing connection

Establishing a TCP connection

- | **Three-way handshake** to establish connection
 - Host A sends a SYN (open; “synchronize sequence numbers”) to host B
 - Host B returns a SYN acknowledgment (SYN ACK)
 - Host A sends an ACK to acknowledge the SYN ACK



Build the TCP header

Flags:

SYN

ACK

FIN

RST

PSH

URG

| | | | |
|-----------------|---|------------------|--|
| Source port | | Destination port | |
| Sequence number | | | |
| Acknowledgment | | | |
| HdrLen | 0 | Flags | |
| Checksum | | Urgent pointer | |

Data

Step 1: A's initial SYN packet

A tells B to open a connection

| A's port | | | B's port |
|-----------------------------|---|-----|----------------|
| A's Initial Sequence Number | | | |
| N/A | | | |
| 5 | 0 | SYN | |
| Checksum | | | Urgent pointer |

Step 1: B's SYN-ACK packet

B tells it accepts
and is ready to
accept next
packet

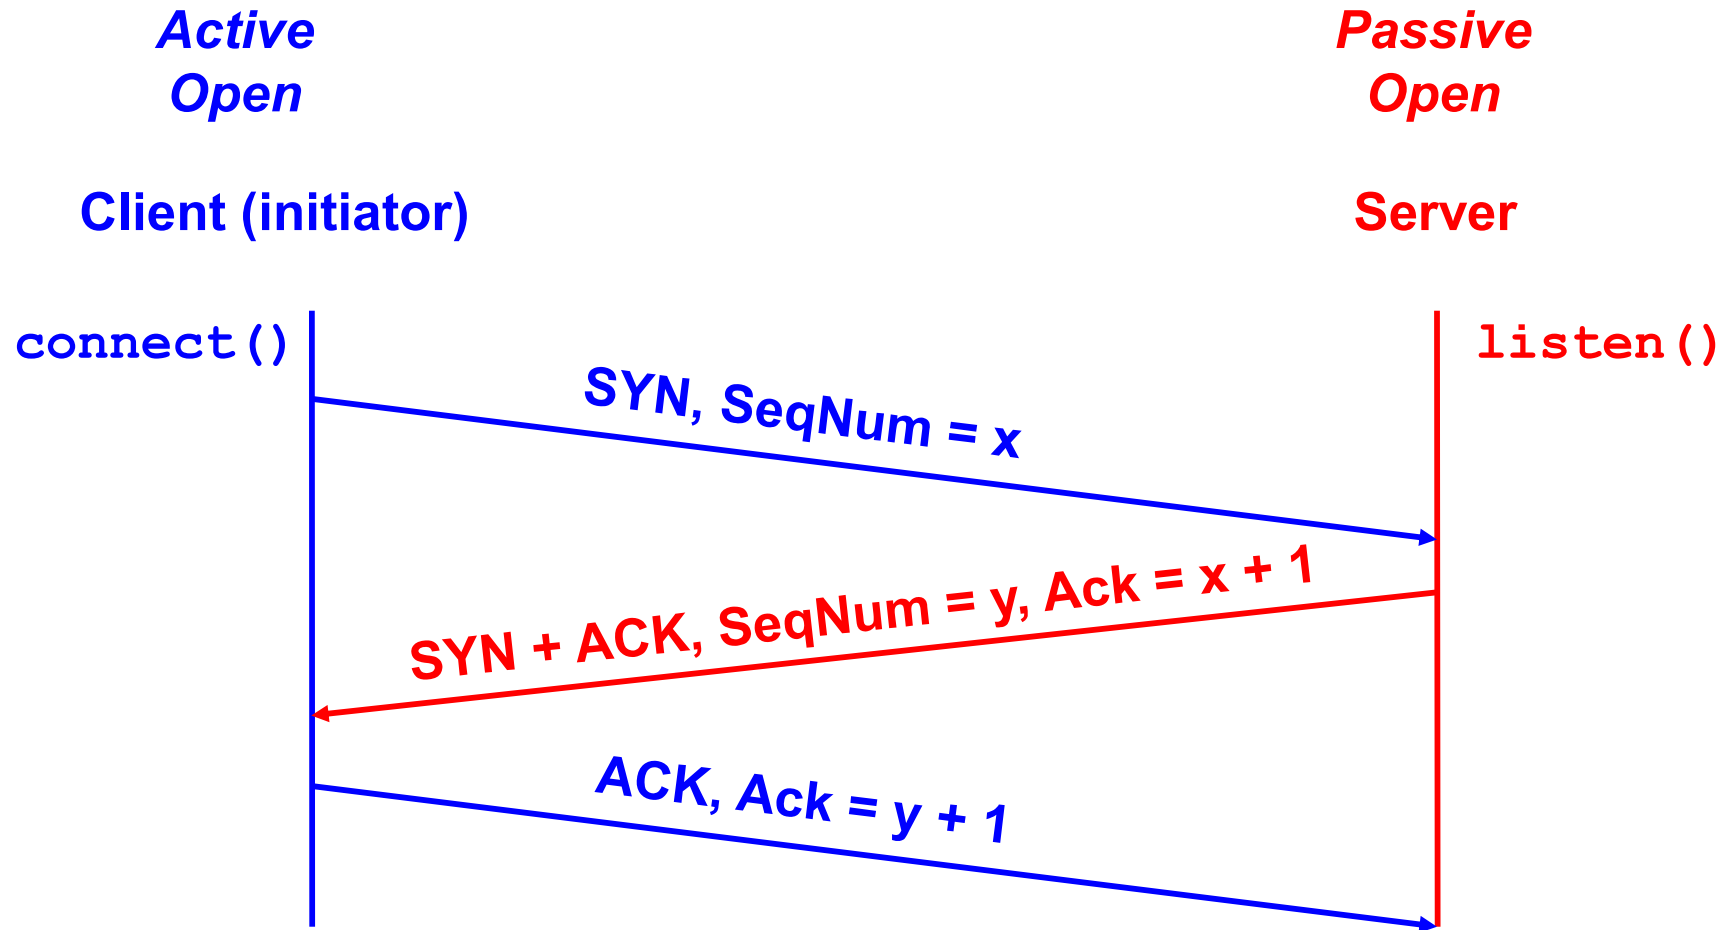
| | | | |
|-----------------------------|---|----------------|--|
| B's port | | A's port | |
| B's Initial Sequence Number | | | |
| ACK=A's ISN+1 | | | |
| 5 | 0 | SYN ACK | |
| Checksum | | Urgent pointer | |

Step 1: A's ACK to SYN-ACK

A tells B to open a connection

| A's port | | | B's port |
|---------------------------------|---|-----|----------------|
| A's Initial Sequence Number + 1 | | | |
| ACK=B's ISN+1 | | | |
| 5 | 0 | ACK | |
| Checksum | | | Urgent pointer |

TCP's 3-Way handshaking




What if the SYN Packet Gets Lost?

- | Suppose the SYN packet gets lost
 - Packet dropped by the network or server is busy
- | Eventually, no SYN-ACK arrives
 - Sender retransmits the SYN on timeout
- | How should the TCP sender set the timer?
 - Sender has no idea how far away the receiver is
 - Default is 3 seconds

TCP CONNECTION TEARDOWN

Normal termination, one side at a time




- Finish (FIN) to close and receive remaining bytes
 - FIN occupies one byte in the sequence space
- Other host acks the byte to confirm
- Closes A's side of the connection, but not B's
 - Until B likewise sends a FIN
 - Which A then acks

TIME_WAIT:

Avoid reincarnation


B will retransmit FIN if ACK is lost

Normal termination, both together



- | Same as before, but B sets FIN with their ack of A's FIN

Abrupt termination



- | A sends a RESET (RST) to B
 - E.g., because application process on A crashed
- | That's it
 - B does not ack the RST
 - Thus, RST is not delivered reliably, and any data in flight is lost
 - But: if B sends anything more, will elicit another RST

Summary

- | Reliability is not easy!
- | Next
 - Flow control
 - LOTs of congestion control