

# **EECS 489**

# **Computer Networks**

**Winter 2025**

Mosharaf Chowdhury

*Material with thanks to Aditya Akella, Sugih Jamin, Philip Levis, Sylvia Ratnasamy, Peter Steenkiste, and many other colleagues.*

# Agenda

---

- Link-state routing
- Distance-vector routing

# Recap: Least-cost path routing

---

- **Given:** router graph & link costs
- **Goal:** find least-cost path
  - From each source router to each destination router
- **Easy way to avoid loops**
  - No reasonable cost metric is minimized by traversing a loop

# Recap:

## Dijkstra's algorithm

---

- Network topology, link costs known to all nodes
  - All nodes have same info
- Each node (“src”) computes least-cost paths to all other nodes
  - After  $k$  iterations, know least-cost path to  $k$  destinations

# From routing algorithm to protocol

---

- ❑ Dijkstra's is a local computation!
  - Computed by a node given complete network graph
- ❑ Possibilities:
  - Option#1: a separate machine runs the algorithm
  - Option#2: every router runs the algorithm
- ❑ The Internet currently uses Option#2

# Link-state routing

---

- Every router knows its local “link state”
  - Router u: “(u,v) with cost=2; (u,x) with cost=1”
- Each router floods its **local link state to all other routers** in the network
  - Does so periodically or when its link state changes
- Every router learns the entire network graph
  - Each runs Dijkstra’s Shortest-Path First (SPF) algorithm locally to compute forwarding table

# Flooding link state

---

## ❑ Flooding

- A node sends its link-state info out all of its links
- The next node forwards the info on all of its links except the one the information arrived at

## ❑ When to initiate flooding?

- Topology change (e.g., link/node failure/recovery)
- Configuration change (e.g., link cost change)
- Periodically
  - » To refresh link-state information (soft states)
  - » Typically (say) every 30 minutes

# Convergence

---

## □ Why flood?

- To get all the nodes in the network to **converge** to the new topology

## □ Upon convergence, all nodes will have **consistent routing information** and can **compute consistent forwarding**:

- All nodes have the same link-state database
- All nodes forward packets on shortest paths
- The next router on the path forwards to the expected next hop




# Convergence delay


---

- Time to achieve convergence
- Sources of convergence delay
  - Time to detect failure
  - Time to flood link-state information
  - Time to re-compute forwarding tables
- What happens if it takes too long to converge?

# Loop from convergence delay



**u** and **w** think that the path to **y** goes through **v**



**v** thinks that the path to **y** goes through **w**

# Performance during convergence period

---

- Looping packets
- Lost packets due to black holes
- Out-of-order packets reaching the destination

# Link-state routing

---

## ▣ Scalability?

- $O(NE)$  messages
- $O(N^2)$  computation time
- $O(\text{Network diameter})$  convergence delay
- $O(N)$  entries in forwarding table

# Link-state routing protocols

---

- ❑ **OSPF**: Open Shortest Path First
- ❑ **IS-IS**: Intermediate System to Intermediate System
  - Similar to OSPF

# OSPF:

## Open Shortest-Path First

---

- ❑ Open: publicly available
- ❑ Uses link-state algorithm
  - Link-state packet dissemination
  - Topology map at each node
  - Route computation using Dijkstra's algorithm
- ❑ Router floods OSPF link-state advertisements to all other routers in entire AS
  - Carried in OSPF messages **directly over IP** (rather than TCP or UDP)
    - » Requires reliable transmission

# Distance-vector protocol

---

- Link-state routing protocol
  - Each node **broadcasts** its **local** information
- Distance-vector routing protocol
  - The opposite (sort of)
  - Each node **tells its neighbors** about its **global** view

# Bellman-Ford equation

---

□ Let

➤  $d_x(y) :=$  cost of least-cost path from  $x$  to  $y$

□ Then

➤  $d_x(y) = \min_v \{ c(x, v) + d_v(y) \}$


cost from neighbor  $v$  to destination  $y$

cost to neighbor  $v$

*min* taken over all neighbors  $v$  of  $x$



# Bellman-Ford example



$$d_x(z) = 9, d_w(z) = 9, d_v(z) = 6$$

B-F equation says:

$$\begin{aligned} d_u(z) &= \min \{c(u,x) + d_x(z), \\ &\quad c(u,w) + d_w(z), \\ &\quad c(u,v) + d_v(z) \} \\ &= \min \{5 + 9, \\ &\quad 3 + 9, \\ &\quad 7 + 6\} = 12 \end{aligned}$$

Neighbor achieving the minimum (w) is next hop in shortest path, used in forwarding table

# Distance vector algorithm

---

- $D_x(y)$  is the estimate of least cost from  $x$  to  $y$ 
  - $x$  maintains its own distance vector  $\mathbf{D}_x = [D_x(y): y \in N]$
- Node  $x$ :
  - Knows cost to each neighbor  $v$ :  $c(x,v)$
  - Maintains its neighbors' distance vectors
    - » For each neighbor  $v$ ,  $x$  has  $\mathbf{D}_v = [D_v(y): y \in N]$

# Distance vector algorithm

---

- From time-to-time, each node sends its own distance vector estimate to neighbors
- When  $x$  receives new DV estimate from neighbor, it updates its own DV using B-F equation
  - $D_x(y) \leftarrow \min_v \{c(x,v) + D_v(y)\}$  for each node  $y \in N$
- Eventually, the estimate  $D_x(y)$  **may** converge to the actual least cost  $d_x(y)$

# Distance vector algorithm

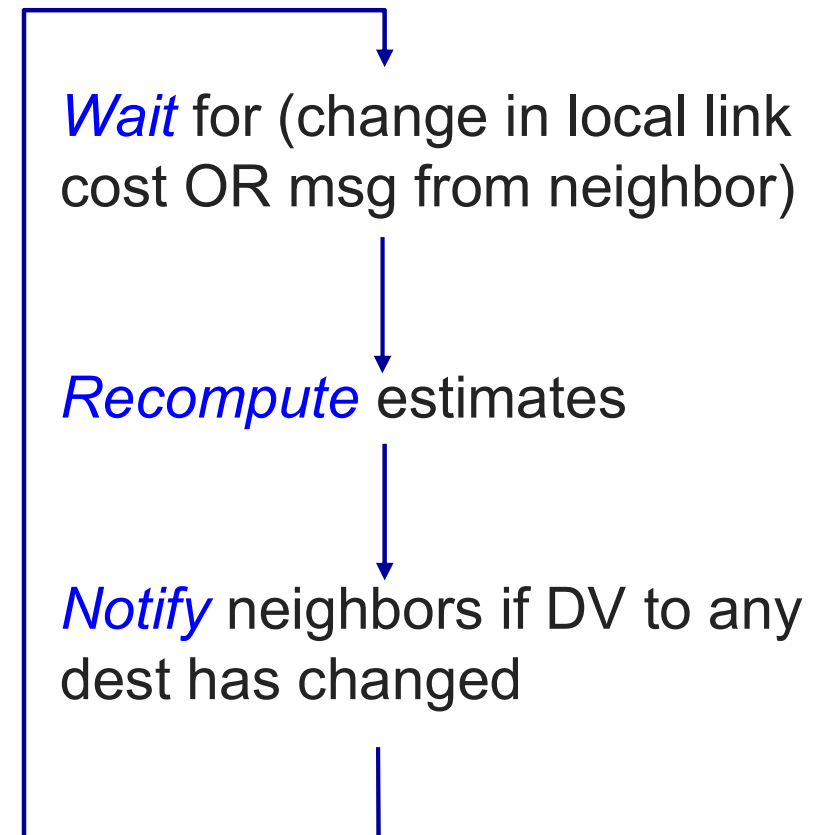
## ? Iterative, asynchronous

- Local iterations caused by
  - » Local link cost change
  - » DV update message from neighbor

## ? Distributed

- Each node notifies neighbors only when its DV changes
  - » Neighbors then notify their neighbors if necessary

## @each node:



---

**5-MINUTE BREAK!**


# Thanks for the midterm eval!

---

- ▣ **Projects:** Better organization and completion of project materials before their release, including resolving issues with project specifications and the autograder.
- ▣ **Pace and Clarity of Lectures:** Need for a slower, clearer delivery of lecture material.
- ▣ **Resources and Practice Materials:** More resources such as practice exams were noted to aid in studying.

# Example

	x	y	z
y	2	0	1
z	3	1	0




	x	y	z
y	2	0	1
z	3	1	0

# Example

	x	y	z
y	4	0	1
z	3	1	0

routing loop!




	x	y	z
y	2	0	1
z	3	1	0



# Example

	x	y	z
y	4	0	1
z	3	1	0

routing loop!




	x	y	z
y	4	0	1
z	3	1	0

# Example

	x	y	z
y	4	0	1
z	3	1	0

routing loop!




	x	y	z
y	4	0	1
z	5	1	0

# Example

	x	y	z
y	4	0	1
z	5	1	0

routing loop!




	x	y	z
y	4	0	1
z	5	1	0

# Example

	x	y	z
y	6	0	1
z	5	1	0

routing loop!




	x	y	z
y	4	0	1
z	5	1	0

# Example

	x	y	z
y	6	0	1
z	5	1	0

routing loop!




	x	y	z
y	6	0	1
z	5	1	0

# Example

	x	y	z
y	6	0	1
z	5	1	0

routing loop!




	x	y	z
y	6	0	1
z	7	1	0

# Example

	x	y	z
y	6	0	1
z	7	1	0

routing loop!




	x	y	z
y	6	0	1
z	7	1	0

# Example

	x	y	z
y	8	0	1
z	7	1	0

routing loop!




	x	y	z
y	6	0	1
z	7	1	0



# Example

	x	y	z
y	8	0	1
z	7	1	0

routing loop!




	x	y	z
y	8	0	1
z	7	1	0

# Example

	x	y	z
y	8	0	1
z	7	1	0

routing loop!



Count-to-infinity  
scenario

	x	y	z
y	8	0	1
z	7	1	0

# Problems with Bellman-Ford

---


- ❑ Routing loops
  - z routes through y, y routes through x
  - y loses connectivity to x
  - y decides to route through z
- ❑ Can take a very long time to resolve
  - Count-to-infinity scenario

# Poisoned reverse

---


- One **heuristic** to avoid count-to-infinity
  - If z routes to x through y,
    - » z advertises to y that its cost to x is infinite
  - y never decides to route to x through z

	x	y	z
y	2	0	1
z	$\infty$	1	0




	x	y	z
y	2	0	1
z	3	1	0

	x	y	z
y	$\infty$	0	1
z	$\infty$	1	0




	x	y	z
y	2	0	1
z	3	1	0

	x	y	z
y	$\infty$	0	1
z	$\infty$	1	0



	x	y	z
y	$\infty$	0	1
z	3	1	0


	x	y	z
y	$\infty$	0	1
z	$\infty$	1	0



	x	y	z
y	$\infty$	0	1
z	7	1	0




	x	y	z
y	$\infty$	0	1
z	7	1	0



	x	y	z
y	$\infty$	0	1
z	7	1	0

	x	y	z
y	8	0	1
z	7	1	0



	x	y	z
y	$\infty$	0	1
z	7	1	0

# Poisoned reverse

---

- ❑ One **heuristic** to avoid count-to-infinity
  - If z routes to x through y,
    - » z advertises to y that its cost to x is infinite
  - y never decides to route to x through z
- ❑ **Not guaranteed**
- ❑ **Loop-free routing** examples include
  - Path vector
  - Source tracing

# Distance-vector routing

---

## ❑ Scalability?

- Requires fewer messages than Link-State
- $O(N)$  update time on arrival of a new DV from neighbor
- $O(\text{network diameter})$  convergence time
- $O(N)$  entries in forwarding table

## ❑ RIP is a protocol that implements DV (IETF RFC 2080)

# Comparison of LS and DV routing

---

## Messaging complexity

- LS: with  $N$  nodes,  $E$  links,  $O(NE)$  messages sent
- DV: exchange between neighbors only

## Speed of convergence

- LS: relatively fast
- DV: convergence time varies
  - Count-to-infinity problem

## Robustness: what happens if router malfunctions?

- LS:
  - Node can advertise incorrect **link** cost
  - Each node computes its *own* table
- DV:
  - Node can advertise incorrect **path** cost
  - Each node's table used by others (error propagates)

# Similarities between LS and DV routing

---

- Both are shortest-path based routing
  - Minimizing cost metric (link weights) a common optimization goal
    - » Routers share a common view as to what makes a path “good” and how to measure the “goodness” of a path
- Due to shared goal, commonly used inside an organization
  - RIP and OSPF are mostly used for **intra**-domain routing

# Summary

---

- Intra-AS routing
  - Link-state routing
  - Distance-vector routing