

The background is a dark blue gradient with a subtle pattern of small white dots. Overlaid on this are several white geometric elements: a large circular scale on the left with degree markings from 150 to 260, and several concentric circles of varying sizes, some with arrows indicating clockwise or counter-clockwise rotation.

ENGR 101 – Chapter 5

Working With Images

Laura Alford, James Juett, Rick Niciejewski

9/10/2020

Functions for Manipulating Matrices

- MATLAB includes many functions for messing with the layout and shape of data within matrices.
- Here are a few, for your reference:
 - **rot90** – rotate the data in a matrix by 90 degrees counterclockwise
 - **fliplr** – flip the matrix "left to right" (i.e. horizontal flip)
 - **flipud** – flip the matrix "up to down" (i.e. vertical flip)
 - **reshape** – keep the same data, but pack it into a different number of rows/columns to "reshape" the matrix
 - **repmat** – create a larger matrix by replicating (i.e. "tiling") this one
- See the MATLAB documentation online for more details!

Transposing a Matrix

- Transposing a matrix takes the rows and turns them into columns:
- Row 1 becomes column 1, Row 2 becomes column 2, etc.

`b = a'` % b is the transpose of a

This single quote is the transpose operator.

WARNING

It's easy to miss this small operator!
Include a comment to remind yourself you transposed data.

a	1	2	3
	4	5	6

b	1	4
	2	5
	3	6

Grayscale Image Representation



- We'll start with grayscale images (i.e. no color).
- Each pixel is simply a single **intensity value**.
 - The higher the value, the closer to white.

- There are two ways to represent intensity:

- An **integer** between 0 and 255, inclusive.
- A real number (a **double**) between 0 and 1, inclusive.

Internally, MATLAB considers integers and doubles to be different "types" of data.



Grayscale Image Representation

- A grayscale image is just a grid of intensity values.
- In MATLAB, this is just a matrix of numbers!

0	0	255	0	0
255	76	76	76	255
76	0	226	0	76
226	226	0	226	226
226	0	40	0	226


- That's it, really!



This is supposed to look like a dog. See it?

Images are Just Numbers in a Matrix

□ Today we'll see how to perform a variety of image processing operations just by manipulating matrices.



□ For example:

□ `grayImg(1, 1) = 255;`
 `grayImg(1, 5) = 255;`

□ `grayImg(grayImg == 226) = 120;`

255	0	255	0	255
255	76	76	76	255
76	0	120	0	76
120	120	0	120	120
120	0	40	0	120

`grayImg`

File Input/Output for Images

- To load an image from a file, use the **imread** function.

```
img = imread('filename.jpg')
```

- To save an image to a file, use the **imwrite** function.

```
imwrite(img, 'filename.jpg')
```

- MATLAB can handle most common image file formats:

- .jpg, .png, .gif, .bmp, .ppm, etc.

The imshow Function

- First, load the file using `imread`:

```
cat_gray = imread('cat_gray.jpg');
```

If you forget this semicolon, MATLAB will try to print out a giant image matrix.



If it does, hit ctrl-c to tell it to stop.

- Now, you can use the `imshow` function to display the image.

```
imshow(cat_gray);
```

MATLAB will open another window to display the image.



What's wrong with this image?



cat_gray.jpg

It has very poor contrast. The cat kind of fades into the background.

Question: What are the max/min intensities used in this image?

`min(min(cat_gray))`

71

`max(max(cat_gray))`

190

0

255



Exercise: Contrast Stretching

- We can improve the image by using more of the possible intensity values.
- This is a *linear interpolation* problem: stretch the range $[71, 190]$ to be $[0, 255]$.
- To do this:
 - Subtract 71 from all pixels
 - Then multiply all pixels by 2.14



cat_gray.jpg



Solution: Contrast Stretching










```
cat_gray = (cat_gray - 71) .* 2.14
```



RGB Color Image Representation

- To represent a color, we need three different values for the amounts of the primary colors red, green, and blue.¹

(R, G, B)

		
(255, 0, 0)	(0, 255, 0)	(0, 0, 255)
		
(0, 0, 0)	(255, 255, 255)	(100, 100, 100)
		
(101, 151, 183)	(124, 63, 63)	(163, 73, 164)


¹ These are the primary colors of light. You may also be familiar with the primary colors of pigment, which are magenta, yellow, and cyan.

RGB Color Image Representation


- In MATLAB, a color image is represented as three different color **channels**.



0	0	255	0	0
255	126	126	126	255
126	0	255	0	126
255	255	0	255	255
255	0	134	0	255



0	0	255	0	0
255	66	66	66	255
66	0	219	0	66
219	219	0	219	219
219	0	0	0	219




0	0	250	0	0
250	0	0	0	250
0	0	183	0	0
183	183	0	183	183
183	0	0	0	183

1 These are the primary colors of light. You may also be familiar with the primary colors of pigment, which are magenta, yellow, and cyan.

RGB Color Image Representation

□ We could store these color channels as three individual matrices, but then we have more things to keep track of...

□ Instead, we can store them on top of each other in a 3D array!



0	0	250	0	0
250	0	0	0	250
0	0	255	0	0
0	0	183	0	0
255	66	66	66	255
183	183	205	183	183
66	0	219	0	66
183	106	106	106	183
219	219	0	219	219
126	0	255	0	126
219	0	0	0	219
255	255	0	255	255
255	0	134	0	255

Accessing Parts of a 3D Array

- **Layers** in a 3D array are controlled by a 3rd dimension.
- Row/column indexing becomes row/column/layer indexing

➡ `mat3d(:, :, 1)`

➡ `mat3d(:, :, 3)`

➡ `mat3d(4, :, 2)`

➡ `mat3d(:, [2, 5], 3)`

➡ `mat3d(3, end, 1)`

0	0	250	0	0
250	0	0	0	250
0	0	255	0	0
0	0	183	0	0
255	66	66	66	255
103	103	255	103	103
66	0	219	0	66
255	126	126	126	255
219	219	0	219	219
126	0	255	0	126
219	0	0	0	219
255	255	0	255	255
255	0	134	0	255

Working With Images as 3D Arrays

□ There are two main modes of operation...

"I want the whole image."

`img(____, ____, :)`

Use the `:` to select all channels.

"I want a single channel."

`img(____, ____, 2)`

Select only the channel you want.

Working With the Whole Image



```
cat = imread('cat_color.jpg');  
imshow(cat);
```

```
% Vertical Flip
```

```
vf = cat([end:-1:1], :, :);  
imshow(vf);
```

```
% Crop the image
```

```
cropped = cat(:, [200:600], :);  
imshow(cropped);
```

We want to flip
all the channels.

Again, select all
channels to be cropped.

Select ONLY columns
200 through 600.

Working With a Single Channel (General Pattern)

□ A pattern for working with single channels:



red

```
% Pull out the red channel  
% to work with it individually  
red = cat(:, :, 1);
```



red


```
% Make changes to the red channel  
red(:) = 255;
```

```
% Put the red channel back in  
cat(:, :, 1) = red;
```

IMPORTANT

red is a **copy** of the red channel. You need this assignment to copy the changes back in.

cat



Working With a Single Channel (Shortcut)

- A shortcut – just work with the channel in place:

```
cat(:,:,1) = 0;    % Set all values in the red channel to 0
cat(:,:,2) = 2 .* cat(:,:,2); % Multiply green channel values by 2
```

- This doesn't always work. For example, to use logical indexing, you need the general pattern on the previous slide.

- e.g. Set all red channel values less than 50 to be 0.

```
cat(:,:,1)(cat(:,:,1) < 50) = 0;
```

DOES NOT WORK

MATLAB doesn't allow a logical indexing operation right after a regular indexing operation. Instead, assign the channel to a variable, then use the logical indexing.


This looks artsy. Let's try it. Any ideas?



This looks artsy. Let's try it. Any ideas?



```
green = cat(:,:,2); % copy green channel  
green(:) = 0;      % set to zero  
cat(:,:,2) = green; % copy back into image  
% It doesn't work 0.0
```



HSV Color Image Representation

- RGB is only one of several image representations.
- HSV is an alternate that works well for certain applications.
 - **Hue**: "which color?"
 - **Saturation**: "how strong is the color?"
 - **Value**: "how bright?"



Hue




HSV Color Image Representation


- HSV images are also stored as a 3D array.
- However, in MATLAB HSV channel values range from 0 to 1.
 - (In MATLAB, RGB values range between 0 and 255)
- To convert between RGB and HSV, use built-in functions:
 - % convert img from RGB to HSV
`hsvImg = rgb2hsv(img);`
 - % convert hsvImg from HSV back to RGB
`img = hsv2rgb(hsvImg);`

Hue

- The hue channel encodes a color as a number between 0 and 1



Can we use HSV to do this?



Exercise: Removing a Color

- First, convert to HSV:

```
cat = imread('cat_color.jpg');  
hsv = rgb2hsv(cat);
```



- Next, make copies of the hue and saturation channels to work with.

```
hue = hsv(:, :, 1);  
sat = hsv(:, :, 2);
```

- Now, the tricky part. Find all locations with a hue between 0.14 and 0.5, and set the saturation to 0 (i.e. meaning color "strength" of 0). *Hint: Use logical indexing.*

You do this part...

- Finally, copy the saturation channel back in (it's the one we changed), convert back to rgb format, and display using `imshow`.

```
hsv(:, :, 2) = sat;  
imshow(hsv2rgb(hsv));
```

Solution: Removing a Color

- First, convert to HSV:

```
cat = imread('cat_color.jpg');  
hsv = rgb2hsv(cat);
```



- Next, make copies of the hue and saturation channels to work with.

```
hue = hsv(:, :, 1);  
sat = hsv(:, :, 2);
```

- Now, the tricky part. Find all location with a hue between 0.14 and 0.5, and set the saturation to 0 (i.e. meaning color "strength" of 0). *Hint: Use logical indexing.*

```
sat(0.14 < hue & hue < 0.5) = 0;
```

- Finally, copy the saturation channel back in (it's the one we changed), convert back to rgb format, and display using `imshow`.

```
hsv(:, :, 2) = sat;  
imshow(hsv2rgb(hsv));
```