

The background is a dark blue gradient with a subtle pattern of small white dots. Overlaid on this are several faint, light blue circular and semi-circular patterns. A prominent feature is a large circular scale on the left side, with markings from 150 to 260 in increments of 10. There are also several concentric circles and arcs with arrows indicating direction, some solid and some dashed.

ENGR 101 – Chapter 16

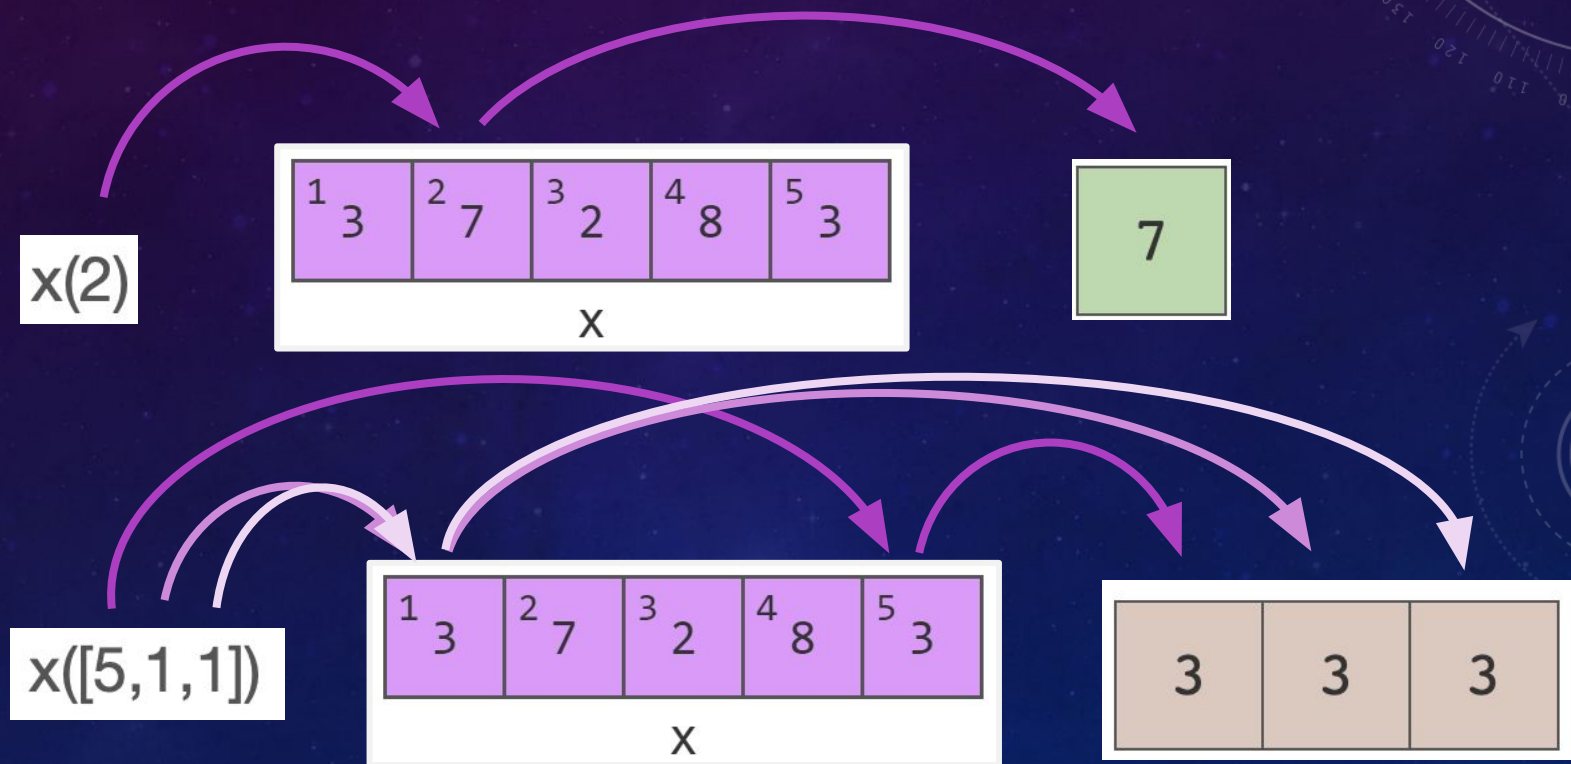
Vectors

3/13/21

INTRO

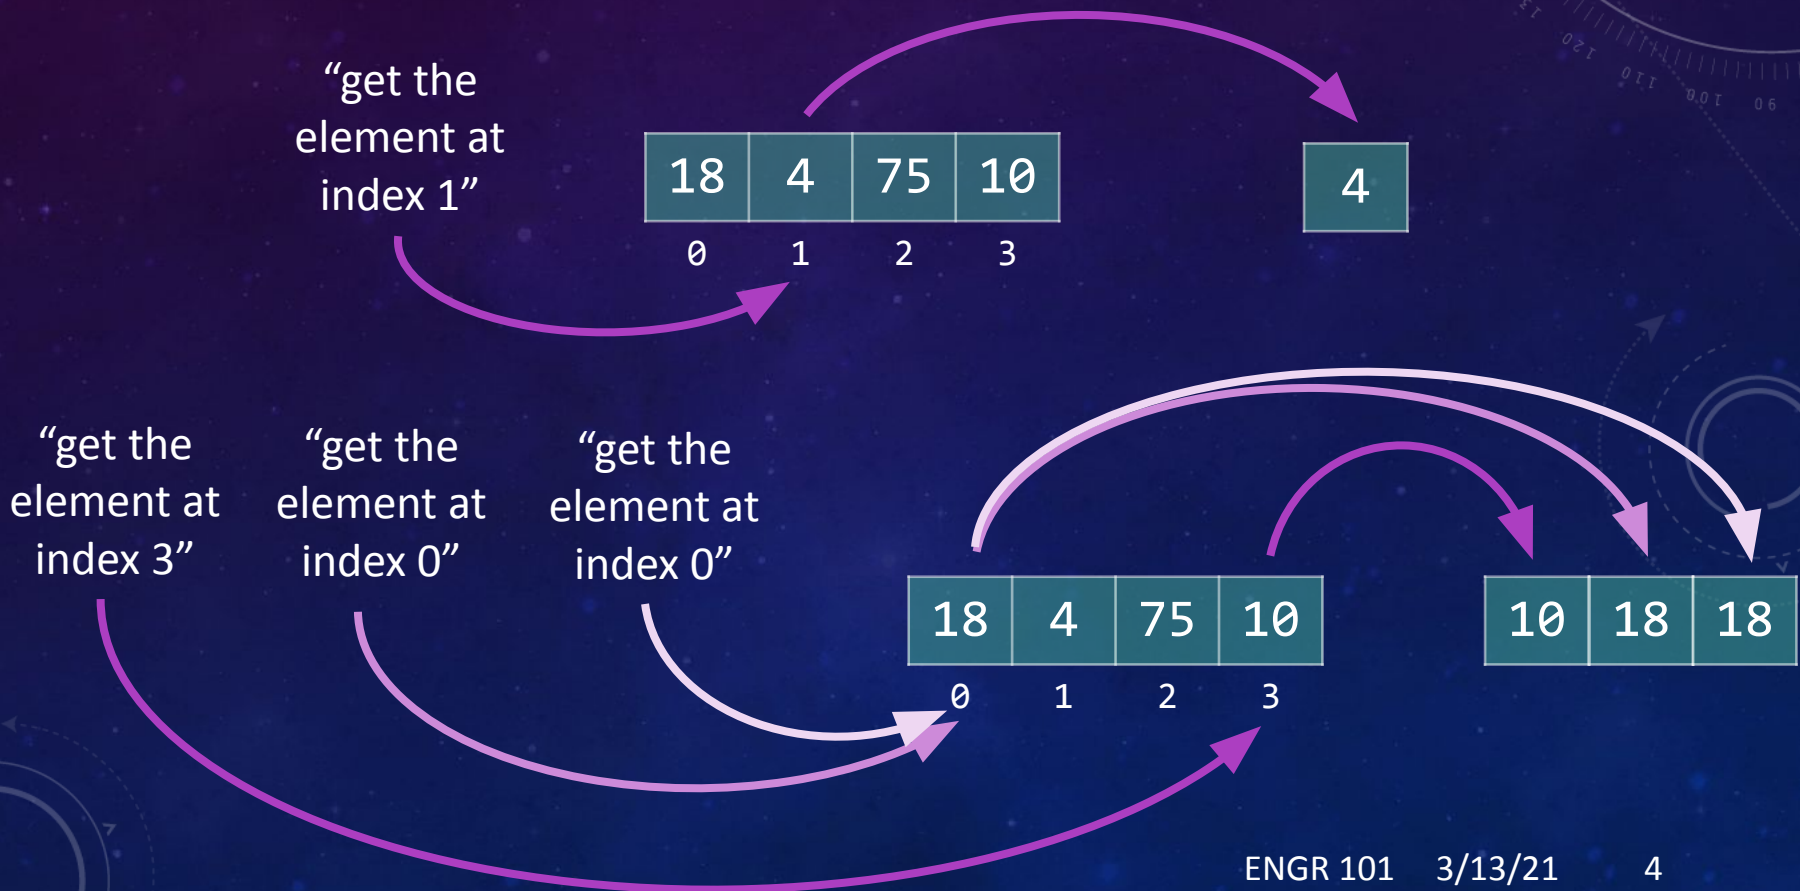
Recall: Using Indexing to Access Data in a Vector

- Indexing in MATLAB uses **random access** to read and write data. The indices start at **1**.



Indexing in C++

- Indexing in C++ also uses **random access** to read and write data, but the indices start at **0** (just like with strings in C++).



DECLARING AND INITIALIZING VECTORS

vectors in C++

- In C++, a vector is used to store a sequence of elements.
 - The elements must be **homogenous** – i.e. all of the same type.
- This is conceptually similar to a vector in MATLAB, but the details are different.
- To use vectors, first include the vector library:

```
#include <vector>  
using namespace std;
```

vectors in C++

- Declare a vector like this:

```
vector<int> someInts;
```

The name someInts will refer to the vector as a whole

In addition to the base type of vector, provide the type of elements it will hold.

- A vector can store elements of any type, as long as they match the type with which it is declared.

```
vector<double> someDoubles;  
vector<bool> someBools;
```

Initializing a vector


- By default, a vector is empty.

```
vector<int> someInts;
```



- Specify one number to allocate an initial number of elements. Their values aren't initialized, though.

```
vector<int> someInts(5);
```




- You can provide a default initial value for all elements.

```
vector<int> someInts(4, 42);
```



- You can provide different initial values for the elements.

```
vector<int> someInts{1, 5, 7};
```



This style needs C++11 to work. You may need to compile using `g++ -std=c++11`

INDEXING INTO VECTORS IN C++

vector indexing

- Can use **indexing** with `[]` to work with individual elements.
- As with strings, indices start with 0 and go up to `length - 1`.

```
vector<int> someInts(4,42);
```

42	42	42	42
0	1	2	3

```
someInts[2] = 5;
```

42	42	5	42
0	1	2	3

```
cout << "index 2: " << someInts[2] << endl;
```

```
index 2: 5
```

vector indexing

- Can use **indexing** with `[]` to work with individual elements.
- As with strings, indices start with 0 and go up to `length - 1`.

```
vector<int> someInts(4,42);
```

42	42	42	42
0	1	2	3

```
someInts[2] = 5;
```

42	42	5	42
0	1	2	3

```
cout << "index 3: " << someInts[3] << endl;
```

```
index 3: 42
```

Indexing Out Of Bounds

- As with strings, it's possible to index off the end of a vector, which results in undefined behavior at runtime.
- Basically, this goes to whatever memory happens to be next to the vector.
 - Maybe you get "lucky" and this memory wasn't important.
 - Maybe you mess up another variable that happens to be there.
 - Maybe your program isn't allowed to use that chunk of memory!
 - This causes a crash called a **segmentation fault** (*aka seg fault*).
 - Maybe it catches on fire¹. Who knows!

¹ We're just kidding. This won't actually happen.

The at Function

- Again, you have the option to use the **.at function** rather than indexing with the square brackets.
- This contains an implicit check to make sure the index is valid.
- The tradeoff is that **at** is slightly slower than **[]**.

```
vector<int> someInts(4,42);
```

42	42	42	42
0	1	2	3

```
someInts[4] = 43;
```

Causes undefined behavior.
(Harder to debug.)

```
someInts.at(4) = 43;
```

Causes a runtime error with a
nice message. (Easier to debug.)

VECTOR FUNCTIONS

Vector functions

- Many different functions can be called on a vector, using the dot notation. Here are a few common ones:

<code>size</code>	Returns the number of elements.
<code>front</code>	Returns a reference to the first element.
<code>back</code>	Returns a reference to the last element.
<code>at</code>	Works like indexing, but does bounds checking.
<code>empty</code>	Returns whether the vector is empty (as a <code>bool</code>).
<code>clear</code>	Removes all elements from the vector.
<code>push_back</code>	Adds a new element to the back of the vector.
<code>pop_back</code>	Removes the last element in the vector.

A Quick Technicality...

- The `vector size()` function returns an “unsigned” `int`.
 - Some compilers will give warnings about comparing this to a regular `int`.
- You can fix this by adding an explicit conversion to an `int`.

```
// Returns the value of the maximum element in the vector
int max_element(const vector<int> &vec) {

    int max_so_far = vec[0]; // assume first is largest

    // iterate through vector, looking for any larger
    for (int i = 0; i < int(vec.size()); ++i) {
        if (vec[i] > max_so_far) {
            max_so_far = vec[i];
        }
    }
    return max_so_far;
}
```

Tells the compiler to convert the result of `vec.size()` to a regular `int`.

TRAVERSING A VECTOR

Traversing a vector

- Use a loop to **traverse** through each element in a vector:

```
vector<int> vec(4,42);  
vec[2] = 5;
```

42	42	5	42
0	1	2	3

i is called the
index variable.

Keep going as long as
i is within bounds.

Increment to next
index on each iteration.

```
for (int i = 0; i < vec.size(); ++i) {  
    cout << vec.at(i) << endl;  
}
```

On each iteration, access
the element at index *i*.

vectors have a
.size() function.



3 min

Exercise: print

printVectorOfInts.cpp

- Write a function called `print`, that will print out a `vector<int>` in the format { _ _ _ _ }, where _ represents an element.

```
// Write the print function here
```

```
int main() {  
    vector<int> someInts(4,42);  
    someInts[2] = 5;  
    print(someInts); // prints { 42 42 5 42 }  
}
```

42	42	5	42
0	1	2	3

If you want to use C++11, you may need to compile thus:
`g++ -std=c++11 printVectorOfInts.cpp -o printInts`

Solution: print

printVectorOfInts.cpp

- Write a function called `print`, that will print out a `vector<int>` in the format { _ _ _ _ }, where _ represents an element.

```
void print(vector<int> vec) {  
    cout << "{ ";  
    for (int i = 0; i < vec.size(); ++i) {  
        cout << vec[i] << " ";  
    }  
    cout << "}" << endl;  
}
```

This function is inefficient...why?

The pass by value copy is expensive!

```
int main() {  
    vector<int> someInts(4,42);  
    someInts[2] = 5;  
    print(someInts); // prints { 42 42 5 42 }  
}
```

42	42	5	42
0	1	2	3

ADDING/REMOVING ELEMENTS FROM A VECTOR


Adding/Removing from a vector

□ To add an element to the back, use **push_back**:

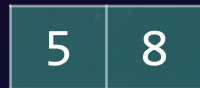
```
vector<int> vec;
```

 size: 0

```
vec.push_back(5);
```

 size: 1
0

```
vec.push_back(8);
```

 size: 2
0 1


```
for (int i = 0; i < 3; ++i) {  
    vec.push_back(i);  
}
```

 size: 5
0 1 2 3 4


Adding/Removing from a vector

□ To remove an element from the back, use **pop_back**:


...
...
...



`vec.pop_back();`




`vec.push_back(0);`



`vec.pop_back();`

`vec.pop_back();`

`vec.pop_back();`





Exercise: Emptying a vector

□ Both of these look good, but only one works. Which one?

1

```
vector<int> vec;  
...  
for (int i = 0; i < vec.size(); ++i) {  
    vec.pop_back();  
}
```

5	8	0	1	2
0	1	2	3	4

size: 5

size: 0

2

```
vector<int> vec;  
...  
while (vec.size() > 0) {  
    vec.pop_back();  
}
```

5	8	0	1	2
0	1	2	3	4

size: 5

size: 0

Solution: Emptying a vector

□ Both of these look good, but only one works. Which one?

~~1~~

```
vector<int> vec;
```

...

The size changes!

5	8	0	1	2
0	1	2	3	4

 size: 5

```
for (int i = 0; i < vec.size(); ++i) {  
    vec.pop_back();  
}
```

5	8
0	1

 size: 2

The loop stops when
i is 3 and size is 2.

2

```
vector<int> vec;
```

...

5	8	0	1	2
0	1	2	3	4

 size: 5

```
while (vec.size() > 0) {  
    vec.pop_back();  
}
```

--

 size: 0

THE .erase FUNCTION

vector erase

□ To remove elements from a vector, use the **erase** function:

```
int main() {  
    vector<int> vec;  
    // Put 8, 6, 7, 5, 3, 0, 9 into the vector  
  
    // this would remove the first element (e.g. the 8)  
    vec.erase(vec.begin());  
}
```

The `.begin()` function refers to the first element in the vector.
When using the erase function, you have to specify indices relative to `.begin()` – it's just how the `erase()` function works.

vector erase

□ To remove elements from a vector, use the **erase** function:

```
int main() {  
    vector<int> vec;  
    // Put 8, 6, 7, 5, 3, 0, 9 into the vector  
  
    // this would remove the first element (e.g. the 8)  
    // vec.erase(vec.begin());  
  
    // this would remove the element at index 2 (e.g. the 7)  
    vec.erase(vec.begin() + 2);  
  
}
```

specify indices relative to .begin()

vector erase

□ To remove elements from a vector, use the **erase** function:

```
int main() {  
    vector<int> vec;  
    // Put 8, 6, 7, 5, 3, 0, 9 into the vector  
  
    // this would remove the first element (e.g. the 8)  
    // vec.erase(vec.begin());  
  
    // this would remove the element at index 2 (e.g. the 7)  
    // vec.erase(vec.begin() + 2);  
  
    // this version would remove a range of indices (e.g. 7, 5, 3)  
    vec.erase(vec.begin() + 2, vec.begin() + 5);  
  
}
```

Note that the upper bound is "exclusive", so the 5th element here wouldn't be removed.

vector erase

□ To remove elements from a vector, use the **erase** function:

```
int main() {  
    vector<int> vec;  
    // Put 8, 6, 7, 5, 3, 0, 9 into the vector  
  
    // this would remove the first element (e.g. the 8)  
    // vec.erase(vec.begin());  
  
    // this would remove the element at index 2 (e.g. the 7)  
    // vec.erase(vec.begin() + 2);  
  
    // this version would remove a range of indices (e.g. 7, 5, 3)  
    // vec.erase(vec.begin() + 2, vec.begin() + 5);  
  
    // this version erases all the way to the end  
    vec.erase(vec.begin() + 2, vec.end());  
}
```

PASSING VECTORS TO FUNCTIONS

Pass by Reference


- Pass by reference gets rid of the expensive copy...
- ...but it's also risky, because it allows modification of the parameter.

- Consider this code:

Pass by reference: work with the original, not a copy.

```
bool anyZeros(vector<int> &vec) {  
    for (int i = 0; i < vec.size(); ++i) {  
        if ( vec[i] = 0 ) { return true; }  
    }  
    return false;  
}
```

Oops. We meant
for this to be ==.



```
int main() {  
    vector<int> someInts(4,42); // { 42 42 42 42 }  
    cout << anyZeros(someInts) << endl; // prints false  
    print(someInts); // prints { 0 0 0 0 }  
}
```


Pass by const Reference

□ To prevent accidental modification, add the **const** keyword.

□ `const` is short for "constant".

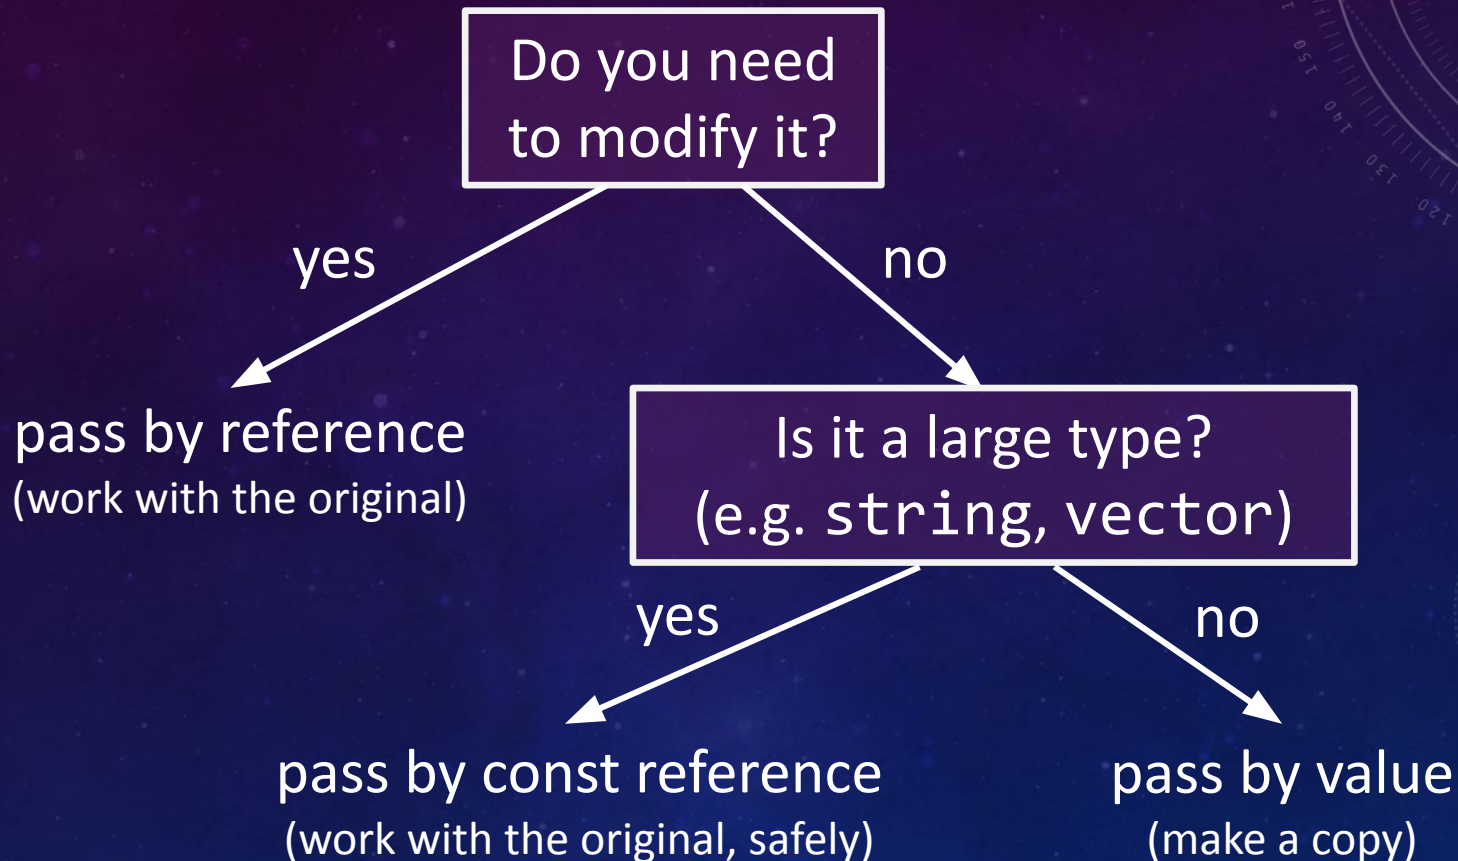
□ You'll get a compiler error to warn you if you accidentally change it.

```
bool anyZeros(const vector &vec) {  
    for (int i = 0; i < vec.size(); ++i) {  
        if ( vec[i] = 0 ) { return true; }  
    }  
    return false;  
}
```

Error! vec is const.

```
int main() {  
    vector<int> someInts(4,42);  
    cout << anyZeros(someInts) << endl; // prints false  
    print(someInts); // prints { 0 0 0 0 }  
}
```

Parameter Passing



COMMON PATTERNS

Common Pattern: "Make space, then fill"

- If you know ahead of time how many elements you need:
 - First allocate enough elements.
 - Then fill in values.
- Example: Put the first N odd numbers in a vector

```
int N = 7;  
vector<int> vec(N);
```

?	?	?	?	?	...	?
0	1	2	3	4	...	N-1

```
for (int i = 0; i < N; ++i) {  
    vec[i] = 2 * i + 1;  
}
```

1	3	5	7	9	...	13
0	1	2	3	4	...	N-1

Common Pattern: "Fill as you go"

- If you don't know ahead of time how many elements you need:
 - Just add them as you go by using `push_back`.
 - The vector will grow as needed to accommodate everything.
- Example: Read values into a vector from a data file.

Now `double` is
our element type.

```
vector<double> data;  
ifstream fileIn("sensor.dat");  
double value;  
while (fileIn >> value) {  
    data.push_back(value);  
}  
fileIn.close();  
// Now we can use the data vector,  
// and more convenient than constantly opening the file
```

size: 0

2	8	-	1	6	
7	7	4	1	4	
.	.	3	.	.	
5	4	.	8	1	...
2	2	9	6	1	
		7			

size: 109

Common Pattern: Using an “Accumulator”

- Task: Compute the result of combining a set of elements
 - Example: Find the sum or product of elements in a vector
- Strategy: *Start a “running total” with the identity for the operation you’re using. Then add elements one at a time.*

```
// Returns the sum of elements in a vector
int sum(const vector<int> &vec) {

    int sum = 0; // start at 0 because it's the additive identity

    // iterate through vector, adding each one to running total
    for (int i = 0; i < vec.size(); ++i) {
        sum += vec[i];
    }
    return sum;
}
```

Common Pattern: Finding the “Best” Element

- Task: Find the “best” element according to some criteria.
 - Example: Find the maximum or minimum
- Strategy: *Keep track of the best so far, compare it to each element, and replace if you find a better one.*

```
// Returns the value of the maximum element in the vector
int max_element(const vector<int> &vec) {

    int max_so_far = vec[0]; // assume first is largest

    // iterate through vector, looking for any larger
    for (int i = 0; i < vec.size(); ++i) {
        if (vec[i] > max_so_far) {
            max_so_far = vec[i];
        }
    }
    return max_so_far;
}
```

Note: This function can't be used on an empty vector.



3 min

Exercise: Find the Minimum

analyze.cpp

- Write a function called `minVal`, that will return the minimum value of a vector of doubles

```
minVal(           ) { // complete this line

// Write the minVal function here
}

int main() {
// code provided to read data from sensor.dat
cout << "The lowest sensor reading was: "
cout << minVal(data) << endl; // prints -98.44
}
```

If you want to use C++11, you may need to compile thus:
`g++ -std=c++11 analyze.cpp -o analyze`

Solution: Find the Minimum

analyze.cpp

```
double minVal(const vector<double> &vec) {  
  
    // assume first is smallest  
    double min_so_far = vec.at(0);  
  
    // iterate through vector, looking for any smaller  
    for (int i = 0; i < vec.size(); ++i) {  
        if (vec.at(i) < min_so_far) {  
            min_so_far = vec.at(i);  
        }  
    }  
  
    return min_so_far;  
}
```

Question

We can't use this approach to find the smallest element and change it. Why not?

If you want to use C++11, you may need to compile thus:
g++ -std=c++11 analyze.cpp -o analyze

Common Pattern: Finding the Index of the “Best” Element

- Find the index of the minimum element rather than its value.
- This is useful for looking the element up later to change it.

```
// Returns the value of the minimum element in the vector
int index_of_min_element(const vector<int> &vec) {

    int index_of_min = 0; // start at first index

    for (int i = 0; i < vec.size(); ++i) {
        if (vec[i] < vec[index_of_min]) {
            index_of_min = i;
        }
    }

    return index_of_min;
}
```

We keep track of the index of the best element instead of the value.

Recall: Parallel Vectors

□ states

This is a column vector containing a list of the states of Earth in ascending alphabetic order. This variable actually contains a different "type" of data called a string that is used to represent words.

□ populations

A parallel column vector to states that stores the population of the corresponding state.

- The index number links the two vectors together

`states.at(3)` □ `populations.at(3)`

Afghanistan	33369945
Albania	2903700
Algeria	40375954
American Samoa	55602
Andorra	69165
Angola	25830958
Anguilla	14763
Antigua and Barbuda	92738
Argentina	43847277
Armenia	3026048
⋮	⋮

statespopulations

Common Pattern: Accessing Parallel Vectors

- To get data that is “parallel”, access each vector using the same index number
- Example: Displaying state names and populations.

```
vector<string> states;  
vector<string> populations;  
...  
// Display first state  
cout << "The first state is: " << states.at(0);  
cout << " -- population " << populations.at(0) << endl;  
  
// Display 10th state  
cout << "The first state is: " << states.at(9);  
cout << " -- population " << populations.at(9) << endl;  
  
// Display last state  
cout << "The first state is: " << states.at(states.size() - 1);  
cout << " -- population " << populations.at(states.size() - 1) << endl;
```

*As we will soon see, an alternative often used in C++ is to create a custom data type that encapsulates both a state's name and population.

Common Pattern: Any/All

- Task: Check if any (or all) element(s) match some criteria
 - Example: Are there any zeros? Are all elements positive?
- Strategy: *Always frame it as an “any” question. Then use a loop with **early termination** to check for any such element.*

```
// Returns whether there are any zeros in the vector
bool any_zeros(const vector<int> &vec) {

    // iterate and check for any zeros
    for(int i = 0; i < vec.size(); ++i) {
        if ( vec[i] == 0 ) {
            return true;
        }
    }
    return false;
}
```

Checking for any element that matches.
If one is found, return immediately! (No
need to check the rest.)

If we make it to here,
there weren't any.

Common Pattern: Any/All

- Task: Check if any (or all) element(s) match some criteria
 - Example: Are there any zeros? Are all elements positive?
- Strategy: *Always frame it as an “any” question. Then use a loop with **early termination** to check for any such element.*
 - *Use negation to turn an “all” question into an “any” question.*

```
// Returns whether all the elements in the vector are positive
bool all_positive(const vector<int> &vec) {

    // iterate and check for any non-positives
    for(int i = 0; i < vec.size(); ++i) {
        if ( !(vec[i] > 0) ) {
            return false;
        }
    }
    return true;
}
```

Checking for counterexamples. If one is found, return immediately!

If we make it to here, they must have all been ok.

Common Pattern: Searching for a Value

- Task: Find the location of a particular element.
 - Example: At what index does the value 0 first occur?
- Strategy: *This is similar to an “any” question, but we return the current index instead of true.*

```
// Returns the index at which the given value first occurs in
// the vector. If it is not present, returns -1
int find(const vector<int> &vec, int value) {

    // iterate and check for the value
    for(int i = 0; i < vec.size(); ++i) {
        if ( vec[i] == value ) {
            return i;
        }
    }
    return -1;
}
```

Checking for any element that matches.
If one is found, return the index!

If we make it to here, there weren't any.
Use -1 to represent “no elements found”.